

System Administration for the Solaris 10 OS Part 1

Student Guide

SA-200-S10 Rev C.1

D61734GC11
Edition 1.1
D65080 and D65081

ORACLE®

Copyright © 2009, 2010, Oracle and/or its affiliates. All rights reserved.

Disclaimer

This document contains proprietary information, is provided under a license agreement containing restrictions on use and disclosure, and is protected by copyright and other intellectual property laws. You may copy and print this document solely for your own use in an Oracle training course. The document may not be modified or altered in any way. Except as expressly permitted in your license agreement or allowed by law, you may not use, share, download, upload, copy, print, display, perform, reproduce, publish, license, post, transmit, or distribute this document in whole or in part without the express authorization of Oracle.

The information contained in this document is subject to change without notice. If you find any problems in the document, please report them in writing to: Oracle University, 500 Oracle Parkway, Redwood Shores, California 94065 USA. This document is not warranted to be error-free.

Sun Microsystems, Inc. Disclaimer

This training manual may include references to materials, offerings, or products that were previously offered by Sun Microsystems, Inc. Certain materials, offerings, services, or products may no longer be offered or provided. Oracle and its affiliates cannot be held responsible for any such references should they appear in the text provided.

Restricted Rights Notice

If this documentation is delivered to the U.S. Government or anyone using the documentation on behalf of the U.S. Government, the following notice is applicable:

U.S. GOVERNMENT RIGHTS

The U.S. Government's rights to use, modify, reproduce, release, perform, display, or disclose these training materials are restricted by the terms of the applicable Oracle license agreement and/or the applicable U.S. Government contract.

Trademark Notice

Oracle and Java are registered trademarks of Oracle and/or its affiliates. Other names may be trademarks of their respective owners.

AMD, Opteron, the AMD logo, and the AMD Opteron logo are trademarks or registered trademarks of Advanced Micro Devices. Intel and Intel Xeon are trademarks or registered trademarks of Intel Corporation. All SPARC trademarks are used under license and are trademarks or registered trademarks of SPARC International, Inc. UNIX is a registered trademark licensed through X/Open Company, Ltd.

This page intentionally left blank.

Daniel Rodriguez Sanchez (daniel.r.rodriguez@ericsson.com) has a
non-transferable license to use this Student Guide.

This page intentionally left blank.

Daniel Rodriguez Sanchez (daniel.r.rodriguez@ericsson.com) has a
non-transferable license to use this Student Guide.

Table of Contents

About This Course	Preface-i
Course Goals.....	Preface-i
Topics Not Covered.....	Preface-iii
How Prepared Are You?.....	Preface-iii
Introductions	Preface-iii
How to Use the Course Materials.....	Preface-iv
Conventions	Preface-iv
 Introducing the Solaris™ 10 OS Directory Hierarchy	1-1
Objectives	1-1
Introducing File Systems	1-2
Default Solaris File Systems	1-2
Introducing / (root) Subdirectories.....	1-2
Introducing Important System Directories	1-3
Introducing Important In-Memory System Directories	1-4
Introducing File Components	1-9
File Names.....	1-9
Inodes.....	1-9
Data Blocks.....	1-10
Identifying File Types.....	1-11
Directories	1-12
Regular Files	1-12
Symbolic Links	1-12
Device Files	1-15
Introducing Hard Links	1-18
Creating New Hard Links	1-19
Removing Hard Links	1-20
Exercise: Identifying File Types	1-21
Preparation.....	1-21
Task	1-21
Exercise Summary.....	1-23
Exercise Solutions: Identifying File Types	1-24

Managing Local Disk Devices.....	2-1
Objectives	2-1
Introducing the Basic Architecture of a Disk.....	2-2
Physical Disk Structure	2-2
Data Organization on Disk Platters.....	2-3
Disk Labels and Partition Tables	2-4
x86/x64 Partitions and Slices	2-8
Disk Device Naming Conventions	2-10
Introducing Solaris OS Device Naming Conventions	2-17
Logical Device Names.....	2-17
Physical Device Names	2-19
Instance Names	2-22
Listing System Devices.....	2-23
The /etc/path_to_inst File.....	2-23
The prtconf Command.....	2-26
The format Utility	2-27
The prtdiag command.....	2-28
Reconfiguring Devices	2-31
Performing a Reconfiguration Boot.....	2-31
Using the devfsadm Command	2-32
Partitioning a Hard Disk.....	2-34
Introducing Disk Partitioning	2-34
Recognizing Disk Space and Undesirable Conditions	2-35
Recognizing Wasted Disk Space.....	2-37
Recognizing Overlapping Disk Slices	2-38
The format Utility and Disk Partition Tables	2-39
Using the format Utility	2-42
Defining Slices on a SPARC System.....	2-43
Using the format Utility to Create fdisk Partitions.....	2-51
Defining Slices on an x86/x64 System.....	2-55
EFI Label and VTOC Label Comparison	2-63
EFI Disk Label Restrictions	2-64
Viewing the VTOC.....	2-64
Replacing a SPARC or x86/x64 VTOC on Disk	2-69
Replacing fdisk Partition Tables on x86/x64 Systems	2-71
Extended Volume Table of Contents	2-73
Exercise: Managing Local Disk Devices	2-75
Preparation.....	2-75
Exercise Summary.....	2-76
Task 1 – Identifying Device Files	2-76
Task 2 – Adding a New Disk or Tape Device	2-78
Task 3 - Working With Disks and Partitions on SPARC	2-79
Task 4 - Working With Disks and Partitions on x86/x64.....	2-83
Exercise Summary.....	2-88

Exercise Solutions - Managing Local Disk Devices	2-89
Task 1 – Identifying Device Files	2-89
Task 2 – Adding a New Disk or Tape Device	2-92
Task 3 - Working With Disks and Partitions on SPARC	2-93
Task 4 - Working With Disks and Partitions on x86/x64.....	2-107
Interface Configurations	3-1
Objectives	3-1
Controlling and Monitoring Network Interfaces.....	3-2
Displaying the MAC Address.....	3-2
Displaying the IP Address.....	3-3
Marking an Ethernet Interface as Down.....	3-3
Sending ICMP ECHO_REQUEST Packets.....	3-4
Capturing and Inspecting Network Packets.....	3-5
Configuring IPv4 Interfaces at Boot Time	3-6
Introducing IPv4 Interface Files.....	3-6
Changing the System Host Name	3-10
Exercise: The Solaris OS Network Commands.....	3-13
Preparation.....	3-13
Task	3-13
Exercise Summary	3-15
Exercise Solutions: The Solaris OS Network Commands	3-16
Preparation.....	3-16
Task and Solutions.....	3-17
Managing the Solaris OS UFS	4-1
Objectives	4-1
Introducing Solaris OS File Systems	4-2
Disk-Based File Systems.....	4-2
Distributed File Systems	4-3
Pseudo File Systems	4-3
Creating a New ufs File System.....	4-5
Viewing the Solaris OS ufs File System	4-5
Partition Tables and Reserved Space	4-8
ufs File System Structures.....	4-9
Using the newfs Command.....	4-15
Checking the File System Using the fsck Command	4-18
Data Inconsistencies Checked by the fsck Command	4-18
Non-Interactive Mode	4-19
Interactive Mode	4-20
Resolving File System Inconsistencies	4-21
Reconnecting an Allocated Unreferenced File.....	4-21
Adjusting a Link Counter	4-22
Salvaging the Free List	4-22
Using Backup Superblocks	4-23

Monitoring File System Use	4-28
Using the df Command	4-28
Using the du Command	4-31
Using the quot Command.....	4-33
Exercise: Creating and Maintaining ufs File Systems	4-35
Preparation.....	4-35
Task	4-35
Exercise Summary.....	4-40
Exercise Solutions: Creating and Maintaining ufs File Systems.....	4-41
Task	4-41
ZFS File System Introduction	5-1
Objectives	5-1
What Is ZFS?	5-2
ZFS Pooled Storage.....	5-2
Transactional Semantics.....	5-2
Checksums and Self-Healing Data	5-3
Scalability	5-3
Snapshots	5-3
Administration	5-4
ZFS Terminology.....	5-5
Creating a Basic ZFS File System.....	5-6
ZFS Storage Pool Components.....	5-7
Using Disks in a ZFS Storage Pool	5-7
Using Files in a ZFS Storage Pool	5-9
Virtual Devices in a Storage Pool	5-10
Using Cache Devices in a ZFS Storage Pool.....	5-11
ZFS Storage Pool Replication Features	5-12
Mirrored Storage Pool Configuration.....	5-12
RAID-Z Storage Pool Configuration	5-12
Self-Healing Data in a Replicated Configuration.....	5-13
Dynamic Striping in a Storage Pool	5-14
Creating and Destroying ZFS Storage Pools.....	5-15
Creating a ZFS Storage Pool.....	5-15
Destroying ZFS Storage Pools.....	5-16
Querying ZFS Storage Pool Status	5-17
Getting Basic ZFS Storage Pool Information.....	5-17
Double Parity RAID-Z.....	5-18
ZFS Storage Pool HEALTH Status	5-18
Creating and Destroying ZFS File Systems.....	5-21
Creating a ZFS File System	5-21
Destroying a ZFS File System.....	5-22
Renaming a ZFS File System	5-24
ZFS Properties	5-25
Read-Only ZFS Properties	5-25

Settable ZFS Properties	5-26
ZFS Property Enhancements	5-32
Querying ZFS Information	5-35
Listing Basic ZFS Information.....	5-35
Sending and Receiving ZFS Data	5-36
Creating and Destroying ZFS Snapshots.....	5-37
Exercise: ZFS File System Introduction	5-38
Preparation.....	5-38
Exercise Summary.....	5-38
Task 1 – Creating and Destroying ZFS Storage Pools	5-39
Task 2 – Create, Rename, and Destroy ZFS File Systems.....	5-40
Task 3 – Working with ZFS Properties	5-42
Exercise Summary.....	5-45
Exercise Solution: ZFS File System Introduction	5-46
Preparation.....	5-46
Exercise Summary.....	5-46
Task 1 – Creating and Destroying ZFS Storage Pools	5-47
Task 2 – Creating, Renaming, and Destroying ZFS File Systems.....	5-52
Task 3 – Working With ZFS Properties	5-55
Performing Mounts and Unmounts	6-1
Objectives	6-1
Working With Mounting Basics	6-2
Determining Which File Systems Are Currently Mounted.....	6-4
Mounting a File System Automatically	6-4
Introducing the Virtual File System Table: /etc/vfstab	6-5
Introducing the /etc/mnttab File System Table.....	6-8
Performing Mounts	6-11
Mounting a Local File System Manually.....	6-11
Using the mount Command Options	6-12
Mounting All File Systems Manually	6-14
Mounting a New File System	6-15
Mounting Different Types of File Systems.....	6-16
Managing ZFS Mount Points	6-19
Automatic Mount Points.....	6-19
Legacy Mount Points.....	6-20
Mounting a ZFS File System	6-21
Performing Unmounts	6-23
Unmounting a File System	6-23
Unmounting All File Systems	6-23
Unmounting a Busy File System.....	6-24
Unmounting a ZFS File System	6-26

Repairing Important Files if Boot Fails on a SPARC System.....	6-27
Repairing Files if Boot Fails on x86/x64 Systems	6-28
Accessing Mounted Diskettes, CD-ROMs, or DVDs.....	6-30
Using the Volume Management Service	6-31
Restricting Access to Mounted Diskettes, CD-ROMs, or DVDs.....	6-33
Stopping Volume Management (vold)	6-33
Troubleshooting Volume Management Service	6-34
Accessing a Diskette, CD-ROM, or DVD Without VMS	6-35
Using the mount Command.....	6-35
Exercise: Mounting File Systems	6-36
Preparation.....	6-36
Task	6-36
Exercise Summary.....	6-41
Exercise Solutions: Mounting File Systems.....	6-42
Task	6-42
Performing Solaris 10 OS Package Administration.....	7-1
Objectives	7-1
Introducing the Fundamentals of Package Administration	7-2
Software Packages	7-2
The /var/sadm/install/contents File	7-2
Package Formats	7-4
Administering Packages From the Command Line	7-6
Translating Package Formats	7-6
Displaying Installed Software Packages Information	7-7
Adding a Software Package	7-9
Checking a Package Installation	7-13
Removing a Software Package.....	7-14
Adding Packages by Using a Spool Directory.....	7-16
Streaming One or More Packages	7-17
Reviewing Package Administration.....	7-19
Exercise 1: Adding a Software Package.....	7-20
Preparation.....	7-20
Task	7-20
Exercise Summary.....	7-21
Exercise 1 Solutions: Adding a Software Package	7-22
Preparation.....	7-22
Task	7-22
Exercise 2: Manipulating Software Packages.....	7-23
Preparation.....	7-23
Task	7-23
Exercise Summary.....	7-26
Exercise 2 Solutions: Manipulating Software Packages	7-27
Preparation.....	7-27

Task	7-27
Managing Software Patches on the Solaris 10 OS.....	8-1
Objectives	8-1
Patch Administration Preparation	8-2
Introducing Solaris OS Patches.....	8-2
Best Practices	8-3
Accessing Patches and Patch Information	8-4
Accessing Patches Using SunSolve	8-5
Accessing Patch Documents.....	8-6
Latest Patch Utility Patches	8-8
Downloading Patches Using SunSolve.....	8-10
Checking Patch Levels	8-14
Preparing Patches for Installation	8-15
Installing and Removing Patches	8-17
Installing a Patch.....	8-17
Removing a Patch	8-18
Installing Patch Clusters	8-19
Patch Administration From the Command Line (CLI)	8-24
Using the smpatch Command	8-25
Phases for Applying Updates	8-25
Applying an Update In Three Steps.....	8-26
Applying an Update In One Step	8-28
Configuring the Patch Management Environment.....	8-29
Examples for Configuring Patch	
Management Environment.....	8-29
Using the Update Policy for Applying Updates	8-31
Example of Using the Update Policy	8-33
Working With Multiple Updates.....	8-37
Authorization and Authentication	8-38
Exercise: Maintaining Patches.....	8-39
Preparation.....	8-39
Task	8-39
Exercise Summary.....	8-41
Exercise Solutions: Maintaining Patches	8-42
Task	8-42
Using Boot PROM Commands	9-1
Objectives	9-1
Introducing Boot PROM Fundamentals	9-2
Goal of the OpenBoot™ Architecture Standard.....	9-3
Boot PROM	9-3
System Configuration Information	9-5
Disabling the Abort Sequence	9-8
Displaying POST to the Serial Port	9-9
Using Basic Boot PROM Commands	9-11
Identifying the System Boot PROM Version	9-12

Booting the System	9-13
Accessing More Detailed Information	9-15
Listing NVRAM Parameters	9-16
Changing NVRAM Parameters	9-17
Restoring Default NVRAM Parameters	9-18
Displaying Devices Connected to the System	9-19
Identifying the System Boot Device	9-23
The show-devs Command	9-25
The devalias Command	9-27
The boot-device Parameter	9-28
Creating and Removing Custom Device Aliases	9-29
The nvalias Command.....	9-29
The nvunalias Command	9-32
Viewing and Changing NVRAM Parameters From the OS	9-34
Using the eeprom Command	9-34
Interrupting an Unresponsive System	9-36
Aborting an Unresponsive System.....	9-36
Exercise: Using Boot PROM Commands.....	9-37
Preparation.....	9-37
Task	9-37
Exercise Summary.....	9-41
Exercise Solutions: Using Boot PROM Commands	9-42
Task	9-42
Using the Grand Unified Bootloader.....	10-1
Objectives	10-1
Introduction to GRUB	10-2
GRUB Terminology	10-2
GRUB-Based Booting.....	10-3
How GRUB-Based Booting Works	10-4
GRUB Device Naming Conventions	10-6
GRUB Functional Components.....	10-6
The GRUB Menu	10-7
Modifying x86 System Boot Behavior in the Solaris OS.....	10-9
Setting Solaris Boot Parameters Using eeprom	10-9
Modifying Boot Behavior Using the kernel Command	10-12
The findroot Command for x86	10-16
GRUB Boot Archives	10-17
Managing the Primary Boot Archive	10-19
Booting a System in the GRUB-Based Boot Environment	10-20
Booting a System to Default Run Level (Multiuser).....	10-20
How to Boot a System to Run Level S (Single-User Level).....	10-21
How to Boot a System Interactively	10-22
Stopping an x86 System for Recovery.....	10-23
Booting the Failsafe Archive for Recovery	10-24

Interrupting an Unresponsive System.....	10-26
Aborting an Unresponsive System.....	10-26
Exercise: Using GRUB	10-27
Preparation.....	10-27
Task	10-27
Exercise Summary.....	10-30
Exercise Solutions: Using GRUB.....	10-31
Task	10-31
Performing Legacy Boot and Shutdown Procedures	11-1
Objectives	11-1
Solaris SPARC Boot Design Overview	11-2
New Design	11-2
Features	11-3
System Requirements	11-4
Legacy Boot and Shutdown.....	11-5
Identifying Boot Process Phases	11-6
SPARC Boot PROM Phase.....	11-8
SPARC Boot Programs Phase.....	11-9
x86/x64 Boot PROM Phase	11-10
x86/x64 Boot Programs Phase	11-10
The kernel Initialization Phase	11-11
The init Phase.....	11-17
Legacy Run Level Fundamentals	11-19
Determining the Current Run Level of a System	11-20
Changing Run Levels	11-20
Controlling Legacy Boot Processes	11-21
The /sbin Directory	11-21
The /etc/rc#.d Directories	11-23
Start Run Control Scripts	11-24
Stop Run Control Scripts	11-24
The /etc/init.d Directory.....	11-25
Performing System Shutdown Procedures	11-27
The /usr/sbin/init Command	11-29
The /usr/sbin/shutdown Command	11-29
“Ungraceful” Shutdown Commands.....	11-31
Service Management Facility	12-1
Objectives	12-1
Introduction to SMF.....	12-2
Features	12-2
Service Definition.....	12-3
Service Identifiers.....	12-4
Service States	12-5
Milestones	12-6
The svc.startd Daemon.....	12-9
The inetd Restarter.....	12-10

Boot Archive Service	12-11
The Service Configuration Repository.....	12-11
The Service Repository Database.....	12-12
Stopping and Starting Services Using SMF	
Commands.....	12-14
Using svcs to Determine Why Services are Not	
Running	12-18
Creating New Service Scripts.....	12-18
Secure By Default.....	12-26
Installation and Secure by Default	12-26
Custom Configurations.....	12-26
Limited Networking Profile	12-27
Opening Selected Services	12-29
Exercise: Managing Services With SMF.....	12-34
Preparation.....	12-34
Task	12-34
Exercise Summary.....	12-38
Exercise Solutions: Managing Services With SMF	12-39
Preparation.....	12-39
Task	12-39
Performing User Administration	13-1
Objectives	13-1
Introducing User Administration.....	13-2
Main Components of a User Account.....	13-2
System Files That Store User Account Information.....	13-3
Setting Password Aging on a User Account	13-14
Managing User Accounts.....	13-15
Introducing Command-Line Tools.....	13-15
Creating a User Account	13-16
Modifying a User Account	13-21
Deleting a User Account	13-23
Creating a Group Entry	13-24
Modifying a Group Entry	13-25
Deleting a Group Entry	13-27
Troubleshooting Login Issues	13-28
Login Problem in the GNOME Desktop.....	13-31
Managing Initialization Files.....	13-32
Introducing System-Wide Initialization Files	13-32
Introducing User Initialization Files	13-33
Customizing the User's Work Environment.....	13-34
Exercise: Adding User and Group Accounts	13-37
Preparation.....	13-37
Task 1 – Disabling the Solaris OS Registration	
Window	13-38
Task 2 – Adding Group Entries	13-38

Task 3 – Adding User Accounts	13-39
Task 4 – Examining Configuration Files.....	13-41
Task 5 – Establishing Password Aging.....	13-42
Task 6 – Modifying User Accounts and Group Entries	13-43
Exercise Summary.....	13-45
Exercise Solutions: Adding User and Group Accounts	13-46
Task 1 – Disabling the Solaris OS Registration Window	13-47
Task 2 – Adding Group Entries	13-47
Task 3 – Adding User Accounts	13-48
Task 4 – Examining Configuration Files.....	13-51
Task 5 – Establishing Password Aging	13-52
Task 6 – Modifying User Accounts and Group Entries	13-53
Controlling System Processes.....	14-1
Objectives	14-1
Viewing System Processes.....	14-2
Using the prstat Command	14-2
Killing Frozen Processes	14-5
Using the kill and pkill Commands	14-5
Performing a Remote Login	14-7
Scheduling an Automatic One-Time Execution of a Command.....	14-8
Using the at Command	14-8
Controlling Access to the at Command.....	14-10
Scheduling an Automatic Recurring Execution of a Command.....	14-11
Introducing the crontab File Format	14-11
Using the crontab Command	14-13
Controlling Access to the crontab Command.....	14-15
Exercise: Using Process Control.....	14-17
Preparation.....	14-17
Task	14-17
Exercise Summary.....	14-19
Exercise Solutions: Using Process Control	14-20
Task	14-20
Solaris™ 10 Operating System Installation Requirements	15-1
Objectives	15-1
Identifying the Hardware Requirements for Solaris 10 OS Installation	15-2
Identifying the Fundamentals of Solaris 10 OS Installations.....	15-3
Solaris 10 OS Installation Options	15-3
Solaris 10 OS Upgrade Options	15-6

Troubleshooting Solaris 10 OS Upgrades.....	15-7
Identifying the Solaris 10 OS Software Components.....	15-8
Software Packages	15-8
Software Clusters	15-8
Solaris OS Software Groups	15-9
Defining Guidelines for Installing Solaris 10 OS	
From a DVD	15-12
Secure by Default Installation Enhancement.....	15-14
Installation and Secure by Default	15-14
Simulation: Solaris 10 Operating System Installation	15-15
Backing Up a Mounted File System With a UFS Snapshot.....	A-1
Objectives	A-1
Creating a UFS Snapshot	A-2
Using the fssnap Command	A-2
Limiting the Size of the Backing-Store File	A-4
Displaying Information for a ufs File System Snapshot	A-4
Backing Up the UFS Snapshot File	A-6
Performing a Backup of a UFS Snapshot.....	A-6
Performing an Incremental Backup Using a UFS Snapshot.....	A-7
Restoring Data From a UFS Snapshot Backup	A-10
Deleting a UFS Snapshot.....	A-10
Exercise: Working With UFS Snapshots	A-11
Preparation.....	A-11
Task	A-11
Exercise Summary.....	A-12
Exercise Solutions: Working With UFS Snapshots	A-13
Task	A-13
Performing File System Backups.....	B-1
Objectives	B-1
Introducing Backup Fundamentals.....	B-2
Routine File System Backups	B-2
Tape Drive Naming	B-2
Tape Drive Control	B-3
Scheduled Backup Strategies	B-4
Backing Up an Unmounted File System.....	B-8
The ufsdump Command.....	B-8
ufsdump Command Options	B-8
The ufsrestore Utility	B-9
Tape Backups.....	B-9
Remote Backups to a Tape.....	B-11
Backups to Disk.....	B-12
Saving ZFS Data With Other Backup Products	B-15
Exercise: Backing Up a File System to Tape.....	B-16

Preparation.....	B-16
Task	B-16
Exercise Summary.....	B-18
Exercise Solutions: Backing Up a File System to Tape	B-19
Task	B-19
Exercise: Backing Up a File System to an Extra Disk Drive.....	B-21
Preparation.....	B-21
Task	B-21
Exercise Summary.....	B-23
Exercise Solutions: Backing Up File System to Extra Disk Drive	B-24
Task	B-24
Performing File System Restores.....	C-1
Objectives	C-1
Restoring a ufs File System.....	C-2
Restoring a Regular File System	C-2
Restoring the /usr File System.....	C-4
Performing Special Case Recovery of / (Root) on SPARC	C-6
Performing Special Case Recovery of / (Root) on x86/x64.....	C-7
Invoking an Interactive Restore	C-9
Performing an Incremental Restore	C-12
Exercise: Restoring Files and File Systems	C-16
Preparation.....	C-16
Exercise Summary.....	C-16
Task 1- Back Up the Root (/) File System to Tape.....	C-17
Task 2 - Back up the / (Root) File System to a File on Disk.....	C-17
Task 3 – Restore /etc/inet/hosts File From Tape or File on Disk	C-18
Task 4 – Destroy and Restore the Root (/) File System.....	C-20
Exercise Summary.....	C-23
Exercise Solutions: Restoring Files and File Systems	C-24
Task 1- Back Up the Root (/) File System to Tape.....	C-24
Task 2 - Back up / (Root) File System to a File on Disk	C-25
Task 3 – Restore /etc/inet/hosts From Tape or File on Disk.....	C-27
Task 4 – Destroy and Restore the Root (/) File System.....	C-29

Unauthorized reproduction or distribution prohibited. Copyright© 2011, Oracle and/or its affiliates.

Daniel Rodriguez Sanchez (daniel.r.rodriguez@ericsson.com) has a
non-transferable license to use this Student Guide.

Preface

About This Course

Course Goals

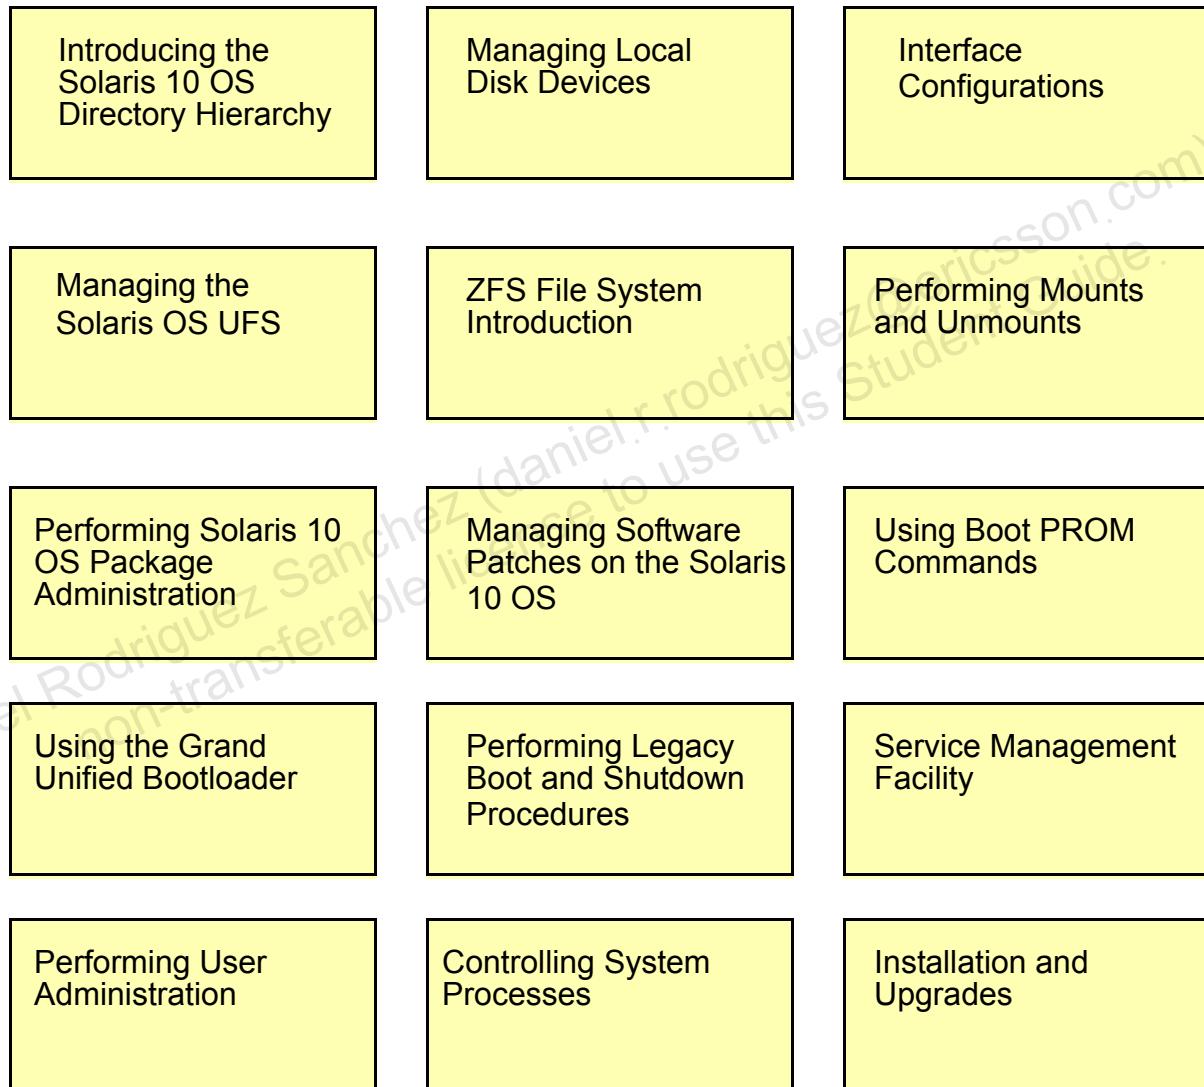
Upon completion of this course, you should be able to:

- Manage file systems
- Manage disk labels and system devices
- Control and monitor network interfaces
- Perform mounts and unmounts
- Administer packages
- Obtain and use patches
- Perform system boot procedures
- Use SMF administrative commands
- Perform user administration
- Control System Processes
- Install software

Course Map

The following course map enables you to see what you have accomplished and where you are going in reference to the course goals.

Outline



Topics Not Covered

This course does not cover the following topics. Many of these topics are covered in other courses offered by Sun Educational Services.

- Basic UNIX® commands – Covered in SA-100-S10: *UNIX® Essentials Featuring the Solaris™ 10 Operating System*
- The vi editor – Covered in SA-100-S10: *UNIX® Essentials Featuring the Solaris™ 10 Operating System*
- Basic UNIX file security – Covered in SA-100-S10: *UNIX® Essentials Featuring the Solaris™ 10 Operating System*
- Network File System (NFS) environment configuration – Covered in SA-202-S10: *System Administration for the Solaris™ 10 Operating System. Part 2*
- All the new features in Solaris 10 – Covered in SA-225-S10: *Solaris™ 10 for Experienced System Administrators*
- Naming services – Covered in SA-202-S10: *System Administration for the Solaris™ 10 Operating System, Part 2*
- Troubleshooting – Covered in ST-350: *Sun™ Systems Fault Analysis Workshop*
- System tuning – Covered in SA-400: *Solaris™ System Performance Management*

Refer to the Sun Educational Services catalog for specific information and registration.

How Prepared Are You?

To succeed in this course you should be able to create procedural application programs.

Introductions

Please introduce yourself to the other students and to the instructor.
Please include the following when you introduce yourself:

- Your name, title, and job responsibility
- Your company affiliation
- Your experience as it relates to topics presented in this course
- Your reasons for enrolling in this course
- Your expectations of this course.

How to Use the Course Materials

The course materials contain learning modules that are composed of the following components:

- Goals – You should be able to accomplish the goals after finishing this course and meeting all of its objectives.
- Objectives – You should be able to accomplish the objectives after completing a portion of instructional content. Objectives support goals and can support other higher-level objectives.
- Lecture – The instructor presents information specific to the objective of the module. This information helps you obtain the knowledge and skills necessary to succeed with the activities.
- Activities include exercises, self-checks, discussions, and demonstrations. Activities help you master an objective.
- Visual aids – The instructor might use visual aids to convey a concept, such as a process, in a visual form. Visual aids commonly contain graphics, animation, and video.

Conventions

The following conventions are used in the course materials.

Icons



Additional resources – Indicates other references that provide additional information on the topics described in the module.



Note – Indicates additional information that can help you but is not crucial to your understanding of the concept being described. You should be able to understand a concept or complete a task without this information. Examples of notational information include keyword shortcuts and minor system adjustments.



Caution – Indicates a risk of personal injury from a nonelectrical hazard, or risk of irreversible damage to data, software, or the operating system. A caution indicates that the possibility of a hazard (as opposed to certainty) might happen, depending on the action of the user.

Typographical Conventions

Typeface	Meaning	Example
AaBbCc123	command, file, method, class and directory names, computer output	Edit the .login file
AaBbCc123	What you type and to highlight topic discussions	% su 1 import java.io.*; 2 import javax.servlet.*; 3 import javax.servlet.http.*; The javax.servlet interface is imported to access life cycle methods (line 2).
AaBbCc123	placeholder, replace with real name	The command to remove a file is rm <i>filename</i>
AaBbCc123	What you enter in class activities	Type chmod a+rwx filename to grant read, write, and execute rights for <i>filename</i> to world, group, and users.
<i>Palatino italics</i>	book titles, new terms, emphasis	Read Chapter 6 in the <i>User's Guide</i> . These are called <i>class</i> options.

Module 1

Introducing the SolarisTM 10 OS Directory Hierarchy

Objectives

Upon completion of this module, you should be able to:

- Describe file systems
- Identify file components
- Identify file types
- Describe hard links

Introducing File Systems

A *file system* is a structure of directories you use to organize and store files. The term *file system* describes the following:

- A particular type of file system: disk-based, network-based, or virtual
- An entire file tree, beginning with the root (/) directory
- The data structure of a disk slice or other media storage device
- A portion of a file tree structure that is attached to a mount point on the main file tree so that the files are accessible

Usually, you know from the context which meaning is intended.

The Solaris™ Operating System (OS) uses the virtual file system (VFS) architecture, which provides a standard interface for different file system types. The VFS architecture enables the kernel to handle basic operations, such as reading, writing, and listing files. The VFS architecture also makes it easier to add new file systems.

Default Solaris File Systems

The Solaris Unix file system (UFS) is hierarchical. It starts with the root directory (/) and continues down through a number of directories. The Solaris installation process enables you to install a default set of directories and use a set of conventions to group similar file types together.

The Solaris OS includes an alternate file system called the Z file system (ZFS). You can use the ZFS with the UFS or as the primary file system. The ZFS uses storage pools to manage physical storage. The ZFS is also a hierarchical file system, starting with the root directory (/). See the *Solaris ZFS Administration Guide* for more information on ZFS.

Introducing / (root) Subdirectories

The directory hierarchy of the Solaris OS is organized for administrative convenience. Branches within this directory tree segregate directories that are used for different purposes. For example, directories exist to hold files that are private to the local system, files to share with other systems, and home directories.



Logically, all directories fall below the / (*root*) directory. Physically, however, directories can be located on a single file system or divided among multiple file systems. Every Solaris OS must have a root file system, but can also have other file systems attached at points within the directory hierarchy. Most file systems are structures created on disk slices that contain files and directories. Some file systems reside in areas of virtual memory and are managed by the Solaris kernel.

Note – Refer to man -s5 filesystem for information on file system organization.

Introducing Important System Directories

The Solaris OS consists of a hierarchy of critical system directories and files that are necessary for the operating system to function properly. Table 1-1 lists of some of the critical, disk-based, system directories and subdirectories that are found in the Solaris OS.

Table 1-1 Critical Directories

/	The root of the overall file system namespace.
/bin	A symbolic link to the /usr/bin directory. It is the directory location for the binary files of standard system commands.
/dev	The primary directory for logical device names. The contents of this directory are symbolic links that point to device files in the /devices directory.
/etc	The directory that holds host-specific configuration files and databases for system administration.
/export	The default directory for commonly shared file systems, such as users' home directories, application software, or other shared file systems.
/home	The default directory or mount point for a user's home directory.
/kernel	The directory of platform-independent loadable kernel modules that are required as part of the boot process.
/lib	The contents of this directory are shared executable files and Service Management Facility executables.

Table 1-1 Critical Directories (Continued)

/mnt	A convenient, temporary mount point for file systems.
/opt	The default directory or mount point for add-on application packages.
/platform	The directory of platform-dependent loadable kernel modules.
/sbin	The single-user bin directory that contains essential executables that are used during the booting process and in manual system-failure recovery.
/usr	The directory that contains programs, scripts, and libraries that are used by all system users.
/var	<p>The directory for varying files, which usually includes temporary, logging, or status files.</p> <p>Following the introduction of the Service Management Facility in the Solaris 10 OS, the /var directory hierarchy is more heavily used than in previous releases.</p> <p>It is important that the /var directory has sufficient disk space available to store software package information, log files, spool files, and so on.</p>

Introducing Important In-Memory System Directories

Table 1-2 lists of some of the important in-memory system directories and subdirectories that are found in the Solaris 10 OS.

Table 1-2 In-Memory System Directories

/dev/fd	The directory that contains special files relating to current file-descriptors in use by the system.
/devices	The primary directory for physical device names.
/etc/mnttab	A memory-based file, in its own file system, that contains details of current file system mounts.

Table 1-2 In-Memory System Directories (Continued)

/etc/svc/volatile	The directory that contains log files and reference files relating to the current state of system services.
/proc	The directory that stores current process-related information. Every process has its own set of subdirectories below the /proc directory.
/system/contract	<p>CTFS (the contract file system) is the interface for creating, controlling, and observing contracts. A contract enhances the relationship between a process and the system resources it depends on by providing richer error reporting and (optionally) a means of delaying the removal of a resource.</p> <p>The service management facility (SMF) uses process contracts to track the processes which compose a service, so that a failure in a part of a multi-process service can be identified as a failure of that service.</p> <p>The contract file system supports all the SMF services.</p>
/system/object	The OBJFS (object) file system describes the state of all modules currently loaded by the kernel. This file system is used by debuggers to access information about kernel symbols without having to access the kernel directly. It is used primarily for Dtrace activity.
/tmp	The directory for temporary files. This directory is cleared during the boot sequence.
/var/run	The directory that contains lock files, special files, and reference files for a variety of system processes and services.



Note – These in-memory directories are maintained by the kernel and system services. You should never manually create, alter, or remove files from these directories.

The following tables list primary subdirectories under key directories.

Table 1-3 Primary Subdirectories Under the /dev Directory

Subdirectory	Description
/dev/dsk	Block disk devices
/dev/fd	File descriptors
/dev/md	Logical volume management metadisk devices
/dev/pts	Pseudo terminal devices
/dev/rdsk	Raw disk devices
/dev/rmt	Raw magnetic tape devices
/dev/term	Serial devices

Table 1-4 Primary Subdirectories Under the /etc Directory

Subdirectory	Description
/etc/acct	Configuration information for the accounting system
/etc/cron.d	Configuration information for the cron utility
/etc/default	Default information for various programs
/etc/inet	Configuration files for network services
/etc/init.d	Scripts for starting and stopping services
/etc/lib	Dynamic linking libraries needed when the /usr file system is not available
/etc/lp	Configuration information for the printer subsystem
/etc/mail	Configuration information for the mail subsystem
/etc/nfs	Configuration file for NFS server logging
/etc/opt	Configuration information for optional packages
/etc/rc#.d	Legacy scripts that are executed when entering or leaving a specific run level

Table 1-4 Primary Subdirectories Under the /etc Directory (Continued)

Subdirectory	Description
/etc/security	Control files for Role Based Access Control and security privileges
/etc/skel	Default shell initialization files for new user accounts
/etc/svc	The Service Management Facility database and log files

Table 1-5 Contents of the /usr Directory

Subdirectory	Description
/usr/bin	Standard system commands
/usr/ccs	C-compilation programs and libraries
/usr/demo	Demonstration programs and data
/usr/dt	Directory or mount point for Java Desktop System (JDS) software
/usr/include	Header files (for C programs, and so on)
/usr/jdk	Directories that contain Java™ technology programs and libraries
/usr/kernel	Platform-independent loadable kernel modules that are not generally required during the boot process
/usr/lib	Architecture-dependent databases, various program libraries, and binaries that are not invoked directly by the user
/usr/opt	Configuration information for optional packages
/usr/sbin	System administration commands
/usr/spool	Symbolic link to the /var/spool directory

Table 1-6 Primary Subdirectories Under the /var Directory

Subdirectory	Description
/var/adm	Log files (for syslog, system accounting, and so on).
/var/crash	For storing crash dump files following a catastrophic system failure. Files from this directory can be analyzed by Help Desk staff to determine the cause of the system crash.
/var/spool	Spoooled files (for mail, print services, and so on).
/var/svc	Service Management Facility control files and logs.
/var/tmp	Long-term storage of temporary files across a system reboot, as an alternative to the /tmp directory.

Introducing File Components

Many files in the Solaris OS make use of a file name and a record called an inode. Most files also make use of data blocks. In general, a file name is associated with an inode, and an inode provides access to data blocks.

Figure 1-1 illustrates the relationship between Filenames, Inodes, and Data Blocks.

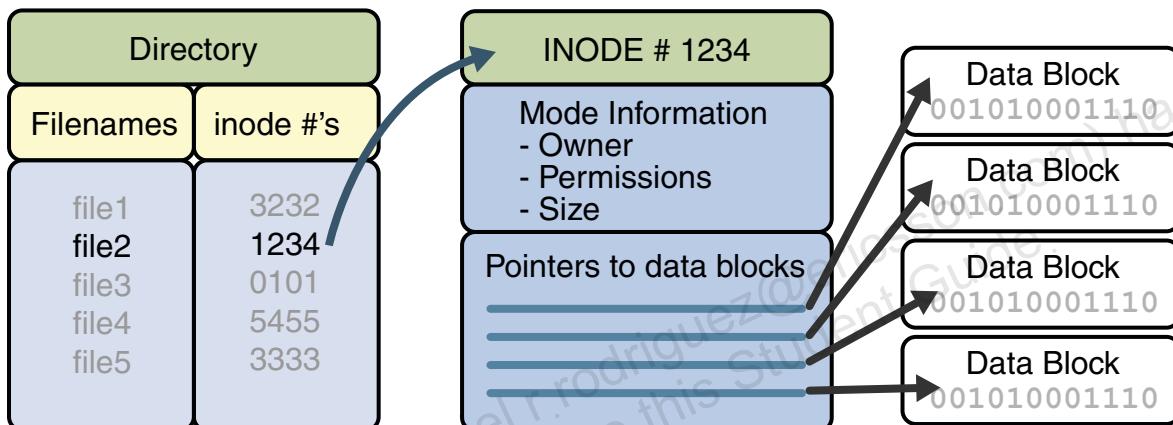


Figure 1-1 Filename, Inode, and Data Block Relationship

File Names

File names are the objects most often used to access and manipulate files. A file must have a name that is associated with an inode.

Inodes

Inodes are the objects that the Solaris OS uses to record information about a file. In general, inodes contain two parts:

1. *File information*, including the file owner, permissions, and size.
2. *Pointers to data blocks* associated with the file content.

Inodes are numbered, and each file system contains its own inode list. When you create a new UFS, a complete new inode list is also created for that UFS.

Data Blocks

Data blocks are disk space units used to store data. Regular files, directories, and symbolic links make use of data blocks. Device files do not hold data.

ZFS uses the following data structures:

- *Dnode data structure*: this data structure contains pointers to data blocks.
- *Znode data structure*: this data structure contains information about the file, including its owner, permissions, and size.

Dnode and znode data structures are called metadata and are stored dynamically by the ZFS. ZFS implements the *ZFS POSIX layer* (ZPL) which is a primary interface for ZFS. The ZPL allows commands you use with the UFS to work with ZFS. ZFS emulates the UFS inode number.

Identifying File Types

The Solaris OS supports a standard set of file types that are found in nearly all UNIX®-based operating systems. In general, files provide a means of storing data, activating devices, or allowing inter-process communication. Of the different types of files that exist in the Solaris OS, there are four main file types:

- Regular or ordinary files
- Directories
- Symbolic links
- Device files

Regular files, directories, and symbolic links all store one or more types of data. Device files do not store data. Instead, device files provide access to devices.

Use the `ls` command to distinguish different file types from one another. The character in the first column of information that the `ls -l` command displays indicates the file type.

The following examples, taken from an Ultra™ 20 and Sun Fire™ V40z server, show partial listings of directories that contain a variety of different file types:

```
# cd /etc
# ls -l
total 618
drwxr-xr-x  2 adm      adm          512 Dec 17 19:08 acct
-rw-r--r--  1 root     sys          253 Nov  9 09:48 aggregation.conf
lrwxrwxrwx  1 root     root         14 Dec 17 19:04 aliases ->
./mail/aliases
drwxr-xr-x  6 root     other        512 Dec 17 19:02 amd64
drwxr-xr-x  7 root     bin          512 Dec 17 19:07 apache
drwxr-xr-x  2 root     bin          512 Dec 17 19:09 apache2
(output truncated)

# cd /devices/pci@0,0/pci-ide@7,1/ide@1
# ls -l
total 2
drwxr-xr-x  2 root     sys          512 Dec 29 19:18 sd@0,0
brw-r----  1 root     sys          28,  0 Dec 30 18:45 sd@0,0:a
crw-r----  1 root     sys          28,  0 Dec 30 18:30 sd@0,0:a,raw
(output truncated)
```

The character in the first column identifies each file type, as follows:

Table 1-7 Characters in the First Column of `ls -l` Command

- Regular files
- d Directories
- l Symbolic links
- b Block-special device files
- c Character-special device files

Directories

Directories store information that associates file names with inode numbers. Unlike regular files, which can hold many different types of data, directories only hold file name-to-inode associations.

A directory contains entries for files of all types that are logically found within that directory.

Regular Files

Perhaps the most common file types found in the Solaris OS are regular files, which enable the user to store many different types of data. Regular files can hold American Standard Code for Information Interchange (ASCII) text or binary data, including image data, database data, application-related data, and more.

There are many ways to create regular files. For example, a user could use the vi editor to create an ASCII text file, or a user could use a compiler to create a file that contains binary data. As another example, a user could use the touch command with a non-existent file name to create a new, empty, regular file.

Symbolic Links

A symbolic link is a file that points to another file. Like directories, which contain only directory information, symbolic links contain only one type of data.

A symbolic link contains the path name of the file to which it points. Because symbolic links use path names to point to other files, they can point to files in other file systems.

The size of a symbolic link always matches the number of characters in the path name it contains.

In the following example, the symbolic link called /bin points to the directory ./usr/bin. The size of the symbolic link is 9 bytes because the path name ./usr/bin contains nine characters.

```
# cd /
# ls -l bin
lrwxrwxrwx 1 root      root          9 Sep 19 15:41 bin -> ./usr/bin
```

Symbolic links can point to regular files, directories, other symbolic links, and device files. They can use absolute or relative path names.

The **ln** command with the **-s** option creates a symbolic link.

```
# ls -l
total 0
-rw-r--r-- 1 root      other         0 Oct 26 10:15 file1
#
# ln -s file1 link1
# ls -l
total 16
-rw-r--r-- 1 root      other         0 Oct 26 10:15 file1
lrwxrwxrwx  1 root      other        5 Oct 26 10:15 link1 -> file1
#
```

Symbolic links direct read and write operations to the file to which they point. The listing above shows that using **link1** as a command-line argument would cause a command to refer to the file called **file1**.

Figure 1-2 shows a symbolic link file called `link1`. As shown in the following figure, the `link1` file is associated with inode number 3561. The data block for the `link1` file contains the path name (`./file1`) to `file1`.

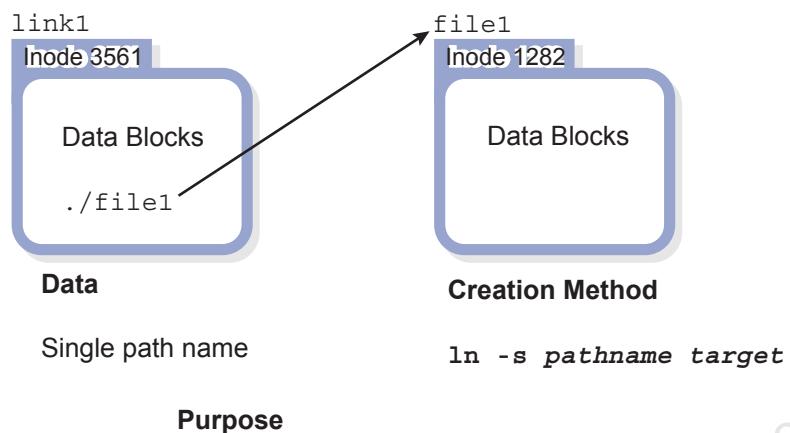


Figure 1-2 Symbolic Links

Device Files

A device file provides access to a device. Unlike regular files, directories, and symbolic links, device files do not use data blocks. Instead, the inode information of device files holds numbers that refer to devices. Use the `ls -l` command to display these numbers.

For example, a long listing of a regular file shows the file's size in the fifth field of output.

```
# cd /etc
# ls -al |more
total 599
drwxr-xr-x  77 root      sys          4096 Sep 23 08:36 .
drwxr-xr-x  26 root      root         1024 Sep 23 08:40 ..
-rw-r--r--  1 root      root        2236 Sep 23 08:36 .cpr_config
drwxr-xr-x  3 root      bin          512  Sep 19 16:39 .java
-rw-r--r--  1 root      sys          524  Sep 19 15:41 .login
-rw-r--r--  1 root      other        18   Sep 19 16:30 .sysidconfig.apps
-rw-r--r--  1 root      other        284  Sep 19 16:00 .sysIDtool.state
(output truncated)
```

A long listing of a device file shows two numbers, separated by a comma, where the file size details would normally have been displayed. These two numbers are called major and minor device numbers. In the following example, the device file `dad@0, 0:a` refers to major device number 136 and minor device number 8.

```
# cd /devices/pci@1f,0/pci@1,1/ide@3
# ls -l dad@0*
total 4
drwxr-xr-x  2 root      sys          512  Sep 19 20:13 dad@0,0
brw-r----  1 root      sys        136,  8 Sep 23 08:35 dad@0,0:a
crw-r----  1 root      sys        136,  8 Sep 23 12:51 dad@0,0:a,raw
brw-r----  1 root      sys        136,  9 Sep 23 08:35 dad@0,0:b
crw-r----  1 root      sys        136,  9 Sep 23 12:51 dad@0,0:b,raw
brw-r----  1 root      sys        136, 10 Sep 23 12:51 dad@0,0:c
crw-r----  1 root      sys        136, 10 Sep 23 12:51 dad@0,0:c,raw
(output truncated)
```

A major device number identifies the specific device driver required to access a device. A minor device number identifies the specific unit of the type that the device driver controls.

A reconfiguration boot creates device files and symbolic links to the device files automatically. In the Solaris OS, use the `devfsadm` command to create new device files.

A relationship exists between the device file and the device it controls. The major and minor device numbers contained in the inode establish the relationship.

Figure 1-3 shows the relationship between the device file `dad@0, 0:a` and the disk device it controls. The inode information for `dad@0, 0:a` contains major number 136 and minor number 8. Major number 136 identifies the `dad` device driver. The `dad` device driver controls integrated device electronics (IDE) disk drives. Minor number 8, in this case, identifies Slice 0.

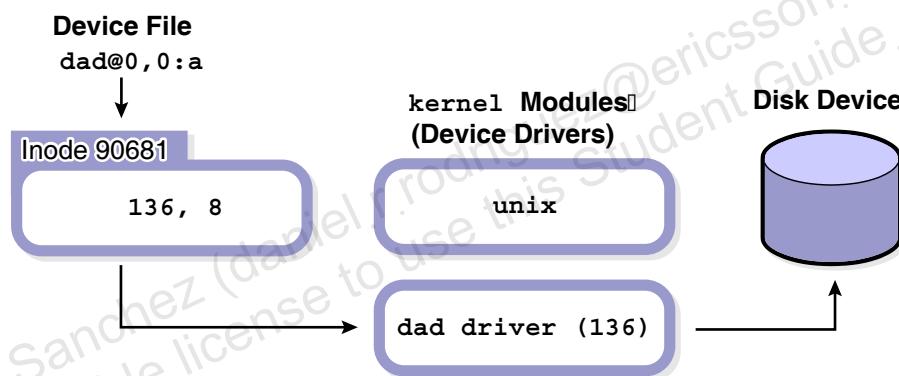


Figure 1-3 Device File Example

You can verify that the device driver is available as a kernel module:

```
# modinfo -w | grep -w dad
21 122e118 7b48 136 1 dad (DAD Disk Driver 1.86)
# modinfo -c | grep -w dad
21          1 dad                               LOADED/INSTALLED
```

Device files fall into two categories: character-special devices and block-special devices. Character-special devices are also called character or raw devices. Block-special devices are often called block devices. Device files in these two categories interact with devices differently.

Character-Special Device Files

The file type “c” identifies character-special device files. Data is accessed as a data stream.

The following example shows a character-special device file.

```
crw-r----- 1 root      sys        136,   8 Sep 23 12:51 dad@0,0:a,raw
```

Block-Special Device Files

The file type “b” identifies block-special device files. For disk devices, block-special device files call for I/O operations based on a defined block size. The block size depends on the particular device.

The following example shows a block-special device file.

```
brw-r----- 1 root      sys        136,   8 Sep 23 08:35 dad@0,0:a
```

Data transferred between a process and a block-special device is first stored in a kernel-managed memory-based cache. This provides better performance when data is being accessed from block-special devices in a repetitive manner. Also, block devices allow random seeks to be performed, and character devices do not.

Introducing Hard Links

A hard link is the association between a file name and an inode. A hard link is not a separate type of file. Every type of file uses at least one hard link. Every entry in a directory constitutes a hard link. Think of every file name as a hard link to an inode. When you create a file, using the `touch` command, for example, a new directory entry is created that links the file name you specified with a particular inode. In this way, creating a new file creates a hard link.

In Figure 1-4, the file called `file1` is listed in the directory `dir1`. In `dir1`, the name `file1` is associated with inode number 1282. The hard link is the association between the name `file1`, and inode number 1282.

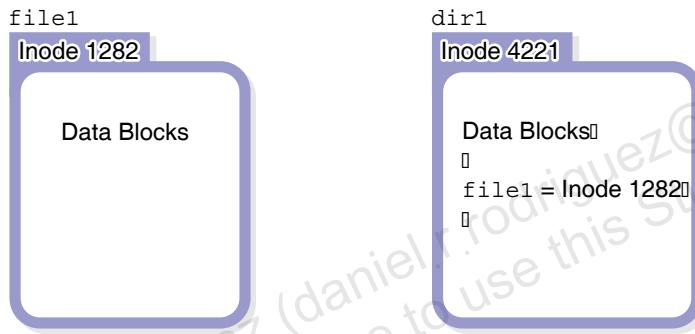


Figure 1-4 Hard Link

Information in each inode keeps count of the number of file names associated with it. This is called a link count. In the output from the `ls -l` command, the link count appears between the column of file permissions and the column identifying the owner. In the following example, the file called `file1` uses one hard link.

```
# cd dir1
# touch file1
# ls -l
total 0
-rw-r--r--    1 root      root          0 Sep 23 13:19 file1
```

Creating New Hard Links

A new hard link for a file name increments the link count in the associated inode.

In the following example, inode 1282 now has two hard links, one for file1 and the other for file2. The ls -li command lists the inode number in the left-most column. The find -inum command locates files and directories that have the same inode numbers.

```
# ln file1 file2
# ls -l
total 0
-rw-r--r--  2 root      root          0 Sep 23 13:19 file1
-rw-r--r--  2 root      root          0 Sep 23 13:19 file2
# ls -li
total 0
    1282 -rw-r--r--  2 root      root          0 Sep 23 13:19 file1
    1282 -rw-r--r--  2 root      root          0 Sep 23 13:19 file2
# find . -inum 1282
./file1
./file2
```

The ln command creates new hard links to regular files.

For example, the ln file1 file2 command creates a new directory entry called file2. The file2 file name is associated with the same inode that is associated with file1.

Figure 1-5 shows the result of the ln command. In dir1, two file names are associated with inode number 1282. Unlike symbolic links, hard links cannot span file systems.

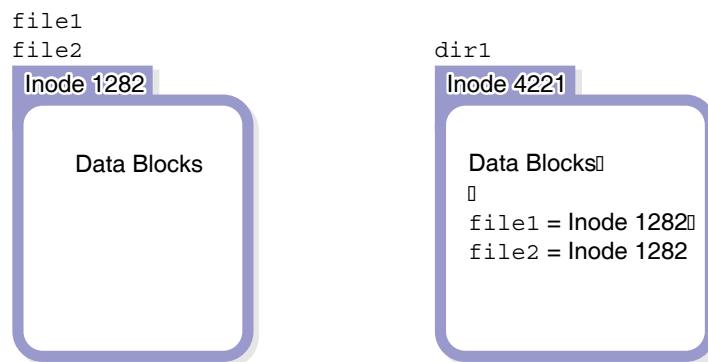


Figure 1-5 File Names Associated With an Inode Number

Removing Hard Links

Deleting one of the files has no effect on the other file. The link count decrements accordingly.

The following example shows how deleting `file1` from the previous example has no effect on `file2`, except to decrement its link count.

```
# rm file1
# ls -li
total 0
 1282 -rw-r--r--    1 root      root          0 Sep 23 13:19 file2
#
```

Exercise: Identifying File Types

In this exercise, you complete the following tasks:

- Use basic Solaris commands to identify and work with the four major file types in the Solaris OS.

Preparation

Refer to the lecture notes as necessary to perform the following steps and answer the following questions.

Task

Complete the following steps:

1. Log in as the root user, and open a terminal window. In the / (root) directory, perform a long listing, and record the name of the first symbolic link listed.

2. What is the size in bytes of the link you found in Step 1? How many characters are there in the name of the file to which this link points?

3. Change to the /dev/dsk directory. Record the file types that you find in this directory.

4. Use the appropriate options of the ls command to display information for the files referenced by the files in the /dev/dsk directory. Record the file types reported.

5. Change to the /dev/pts directory, and use the same commands you used in Steps 3 and 4 for the /dev/dsk directory. Record the file types you find.

Exercise: Identifying File Types

6. Change to the /etc/init.d directory, and identify the type of file in this directory.
-

7. How many hard links are associated with the /etc/init.d/apache file? What is the inode number associated with this file?
-

Find the number of files in the /etc directory or below that have the same inode number as that used by the /etc/init.d/apache file.

8. Create a new directory called /testdir. Create a file in this directory called file1. Create a symbolic link called link1 that points to file1.

9. List file1 and the link1 symbolic link. Do these files use the same or different inodes?
-

10. In the /testdir directory, create a new directory called newdir. What is the number of hard links associated with the newdir directory? What is the inode number associated with the newdir directory?
-

11. List all files, including hidden files, that exist in the newdir directory. Which of these files uses the same inode as the newdir directory?
-

12. Create a new directory called dir2 below the newdir directory. What happens to the link count for the newdir directory?
-

13. Use the ls command with appropriate options to find the new file name that uses the same inode as the newdir directory. Record the name of the new file.
-

Exercise Summary



Discussion – Take a few minutes to discuss what experiences, issues, or discoveries you had during the lab exercises.

- Experiences
- Interpretations
- Conclusions
- Applications

Exercise Solutions: Identifying File Types

This section provides task steps and their solutions.

Task

Complete the following steps:

1. Log in as the root user, and open a terminal window. In the / (root) directory, perform a long listing, and record the name of the first symbolic link listed.

```
# cd /
# ls -l
total 599
lrwxrwxrwx 1 root      root          9 Apr 27
17:51 bin -> ./usr/bin/
drwxr-xr-x 3 root      sys           512 Apr 27
17:52 boot/
drwxr-xr-x 2 root      root          512 Apr 28
08:36 Desktop/
drwxr-xr-x 18 root     sys          5120 Apr 27
18:43 dev/
drwxr-xr-x 2 root      sys           512 Apr 27
18:39 devices/
drwxr-xr-x 2 root      root          512 Apr 28
08:36 Documents/
drwxr-xr-x 87 root     other         4608 Apr 27
18:45 etc/
drwxr-xr-x 2 root      sys           512 Apr 27
18:03 export/
dr-xr-xr-x 1 root      root          1 Apr 27
18:43 home/
drwxr-xr-x 15 root     sys          512 Apr 27
18:16 kernel/
```

In the italicized text, "The /bin symbolic link ... directory.", remove '/' from '/bin'.

The /bin symbolic link should be the first link listed in the / (root) directory.

2. What is the size in bytes of the link you found in Step 1? How many characters are there in the name of the file to which this link points?

The bin symbolic link contains 9 bytes of data and points to ./usr/bin.

3. Change to the /dev/dsk directory. Record the file types that you find in this directory.

```
# cd /dev/dsk
# ls -l
total 126
lrwxrwxrwx 1 root      root          46 Dec 17 19:16 c0t0d0p0 ->
.../..../devices/pci@0,0/pci-ide@6/ide@0/sd@0,0:q
lrwxrwxrwx 1 root      root          46 Dec 17 19:16 c0t0d0p1 ->
.../..../devices/pci@0,0/pci-ide@6/ide@0/sd@0,0:r
lrwxrwxrwx 1 root      root          46 Dec 17 19:16 c0t0d0p2 ->
.../..../devices/pci@0,0/pci-ide@6/ide@0/sd@0,0:s
...
```

The /dev/dsk directory contains symbolic links.

4. Use the appropriate options of the ls command to display information for the files referenced by the files in the /dev/dsk directory. Record the file types reported.

```
# ls -ll
total 0
brw-r---- 1 root      sys           33, 16 Dec 31 15:33 c0t0d0p0
brw-r---- 1 root      sys           33, 17 Dec 31 15:33 c0t0d0p1
brw-r---- 1 root      sys           33, 18 Dec 31 15:33 c0t0d0p2
...
```

The symbolic links in the /dev/dsk directory point to block-special device files.

5. Change to the /dev/pts directory, and use the same commands you used in Steps 3 and 4 for the /dev/dsk directory. Record the file types you find.

```
# cd /dev/pts
# ls -l
total 66
lrwxrwxrwx 1 root      root          28 Dec 17 19:16 0 ->
.../..../devices/pseudo/pts@0:0
lrwxrwxrwx 1 root      root          28 Dec 17 19:16 1 ->
.../..../devices/pseudo/pts@0:1
lrwxrwxrwx 1 root      root          29 Dec 17 19:23 10 ->
.../..../devices/pseudo/pts@0:10
...
# ls -ll
total 0
crw--w---- 1 root      tty           24,  0 Dec 18 07:42 0
crw--w---- 1 root      tty           24,  1 Dec 17 20:49 1
crw-r--r-- 1 root      sys           24, 10 Dec 31 16:36 10
...
```

Exercise Solutions: Identifying File Types

The /dev/pts directory contains symbolic links.

The symbolic links in the /dev/pts directory point to character-special device files.

6. Change to the /etc/init.d directory, and identify the type of file in this directory.

```
# cd /etc/init.d ; ls -l
total 128
-rwxr--r-- 1 root      sys          625 Nov  9 09:48 acct
-rwxr--r-- 1 root      sys          397 Nov  9 09:48 acctadm
-rwxr--r-- 6 root      sys         1649 Sep  5 08:02 apache
-rwxr--r-- 5 root      sys          824 Aug 16 09:47 appserv
...
...
```

The /etc/init.d directory contains regular files.

7. How many hard links are associated with the /etc/init.d/apache file? What is the inode number associated with this file?

```
# ls -li apache
10142 -rwxr--r-- 6 root      sys          1649 Sep  5 08:02 apache
#
```

The /etc/init.d/apache file has six hard links associated with it. The inode number varies among different systems.

8. Find the number of files in the /etc directory or below that have the same inode number as that used by the /etc/init.d/apache file. In this example, the inode number is 10142.

```
# ls -i /etc/init.d/apache
10142 /etc/init.d/apache
# find /etc -inum 10142 -exec ls -i {} \;
10142 /etc/init.d/apache
10142 /etc/rc0.d/K16apache
10142 /etc/rc1.d/K16apache
10142 /etc/rc2.d/K16apache
10142 /etc/rc3.d/S50apache
10142 /etc/rcS.d/K16apache
#
#
```

Six files, including /etc/init.d/apache, use the same inode number.

9. Create a new directory called /testdir. Create a file in this directory called file1. Create a symbolic link called link1 that points to file1.

```
# mkdir /testdir
# cd /testdir
# touch file1
# ln -s file1 link1
```

10. List file1 and the link1 symbolic link. Do these files use the same or different inodes?

```
# ls -li
total 2
  187998 -rw-r--r--  1 root      root          0 Dec 31 16:40 file1
  188110 lwxrwxrwx  1 root      root          5 Dec 31 16:40 link1 ->
file1
#
```

These two files use two different inodes.

11. In the /testdir directory, create a new directory called newdir. What is the number of hard links associated with the newdir directory? What is the inode number associated with the newdir directory?

```
# mkdir newdir
# ls -ldi newdir
  188152 drwxr-xr-x  2 root      root          512 Dec 31 16:43 newdir
#
```

The link count for the newdir directory is two. The inode number varies among different systems.

12. List all files, including hidden files, that exist in the newdir directory. Which of these files uses the same inode as the newdir directory?

```
# ls -lia newdir
total 4
  188152 drwxr-xr-x  2 root      root          512 Dec 31 16:43 .
  187997 drwxr-xr-x  3 root      root          512 Dec 31 16:43 ..
#
```

The file called dot (.) uses the same inode as the newdir directory.

13. Create a new directory called dir2 below the newdir directory. What happens to the link count for the newdir directory?

```
# mkdir newdir/dir2
# ls -ldi newdir
  188152 drwxr-xr-x  3 root      root          512 Dec 31 16:45 newdir
#
```

The link count increases from two to three.

Exercise Solutions: Identifying File Types

14. Use the `ls` command with appropriate options to find the new file name that uses the same inode as the `newdir` directory. Record the name of the new file.

```
# ls -laRi newdir
newdir:
total 6
    188152 drwxr-xr-x    3 root      root          512 Dec 31 16:45 .
    187997 drwxr-xr-x    3 root      root          512 Dec 31 16:43 ..
    188154 drwxr-xr-x    2 root      root          512 Dec 31 16:45 dir2

newdir/dir2:
total 4
    188154 drwxr-xr-x    2 root      root          512 Dec 31 16:45 .
    188152 drwxr-xr-x    3 root      root          512 Dec 31 16:45 ..
#
#
```

The newdir/dir2/.. file uses the same inode as the newdir directory

Module 2

Managing Local Disk Devices

Objectives

Upon completion of this module, you should be able to:

- Describe the basic architecture of a disk
- Manage disk labels
- Describe the naming conventions for devices
- Describe support for iSCSI target devices
- List system devices
- Reconfigure devices
- Perform hard disk partitioning
- Describe EFI disk labels

Introducing the Basic Architecture of a Disk

A disk device has physical components and logical components. The physical components include disk platters and read/write heads. The logical components include disk slices, cylinders, tracks, and sectors.

Physical Disk Structure

A disk is physically composed of a series of flat, magnetically coated platters that are stacked on a spindle. The spindle turns while the read/write heads move radially as a single unit, reading and writing data on the platters.

Figure 2-1 identifies the parts of a disk.

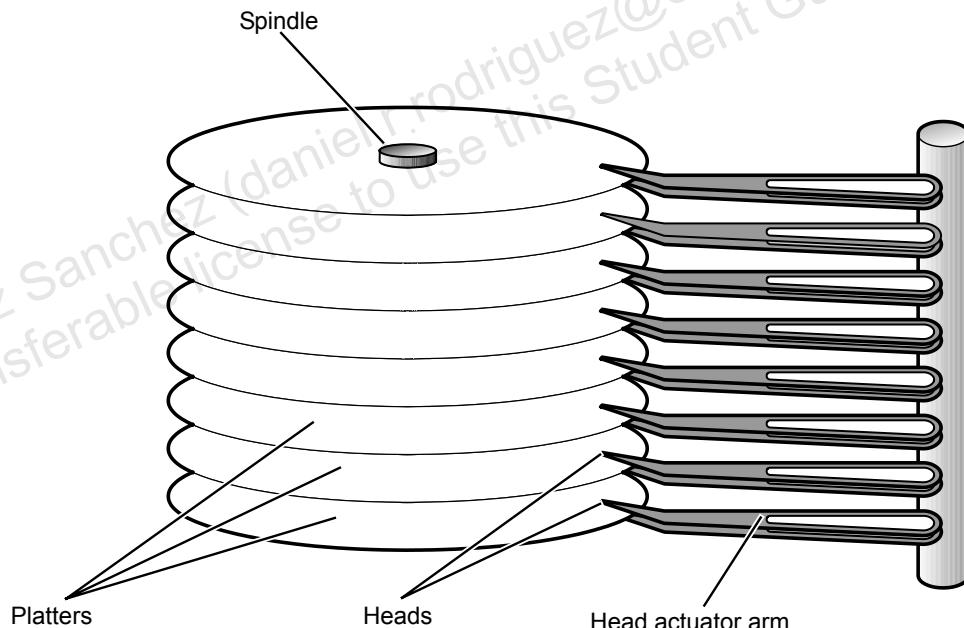


Figure 2-1 Components of a Disk

The following list describes the physical components of a disk:

- The disk storage area is composed of one or more platters.
- The platters rotate.
- The head actuator arm moves the read/write heads as a unit radially.
- The read/write heads read and write data on the magnetic surface on both sides of the platters.

Data Organization on Disk Platters

Figure 2-2 shows the logical components of a disk platter.

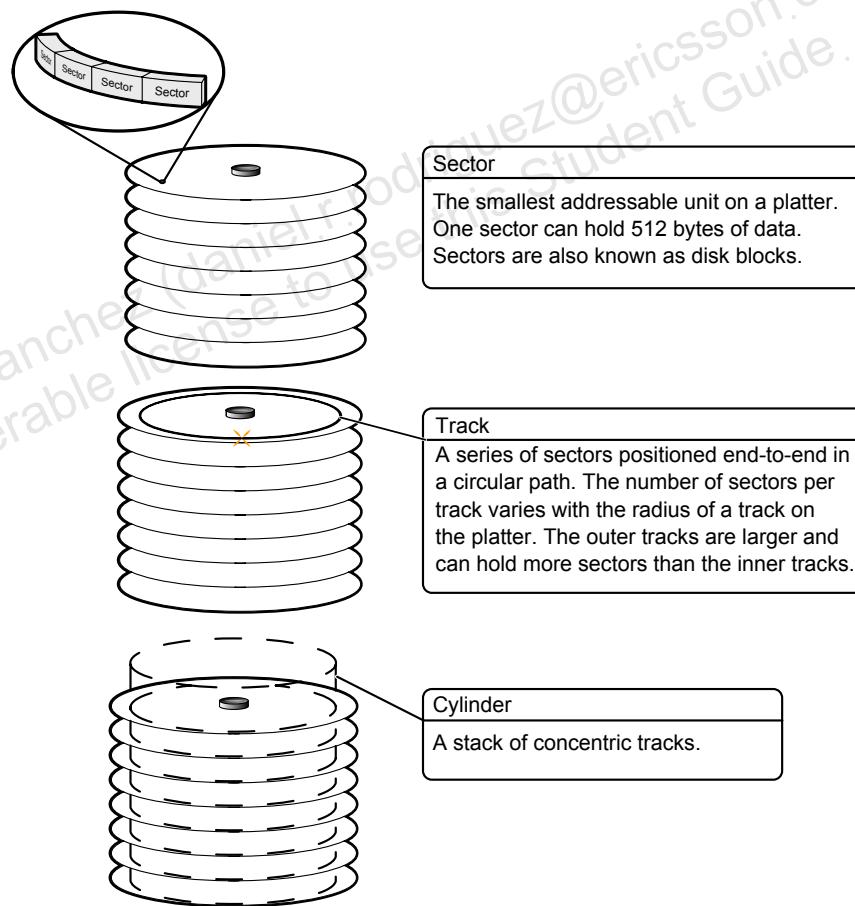


Figure 2-2 Data Organization on Disk Platters

A disk platter is divided into sectors, tracks, and cylinders.

Sector	The smallest addressable unit on a platter. One sector can hold 512 bytes of data. Sectors are also known as disk blocks.
Track	A series of sectors positioned end-to-end in a circular path.
Cylinder	A stack of tracks.

The number of sectors per track varies with the radius of a track on the platter. The outermost tracks are larger and can hold more sectors than the inner tracks. Disks present a fixed number of sectors per track to the Solaris OS, even though physically the number of sectors per track varies by track location. Because a disk spins continuously and the read/write heads move as a single unit, the most efficient seeking occurs when the sectors to be read from or written to are located in a single cylinder.

Disk Labels and Partition Tables

Disk slices are groups of cylinders that are commonly used to organize data by function. For example, one slice can store critical system files and programs while another slice on the same disk can store user-created files. Grouping cylinders into slices is done to organize data, facilitate backups, and provide swap space. A starting cylinder and an ending cylinder define each slice. These cylinder boundaries determine the size of a slice.

An area of every disk is set aside for storing information about the disk controller, geometry, and slices. This information is called the disk label. Partition, cylinder, and slice information is stored in partition tables within disk labels. SPARC®-based systems maintain one partition table on each disk. To label a disk means to write slice information onto the disk. You usually label a disk after you change its slices.

The Solaris OS supports the following two disk labels:

- The Sun Microsystems, Inc. (SMI) disk label:
 - Is typically used in SPARC-based Solaris OSs.
 - Is the SPARC volume table of contents (VTOC) label for disks and is often called the SMI VTOC disk label.
 - Provides support for disks that are less than 1 terabyte.
 - Occupies the first sector of a disk in SPARC systems

- Includes a partition table in which you can define up to eight (0 through 7) disk partitions (slices). A starting cylinder and an ending cylinder define each slice. Whole-cylinder boundaries determine the sizes of the slices. Slice 2 represents the entire disk. Slice 2 maintains important data about the entire disk, such as the size of the actual disk and the total number of cylinders available for the storage of files and directories.
- Must be used for SPARC-based and Solaris x86/x64-based systems boot disks.
- Solaris OS systems for the x86/x64 platforms maintain two partition tables on each disk. The first sector of disks on x86/x64 systems contains a fixed disk (fdisk) partition table. The second sector of the Solaris fdisk partition on an x86/x64 system holds the partition table that defines slices within the Solaris fdisk partition. The label that holds this table is known as the x86/x64 VTOC.
- The fdisk partition table defines up to four fdisk partitions. One of these fdisk partitions may be used for the Solaris OS. Provision is made for up to sixteen slices within a Solaris fdisk partition, but generally only ten of these are used (eight, plus two used for platform-specific purposes).
- The extensible firmware interface (EFI) disk label:
 - Provides support for disks that are larger than 1 terabyte on systems that run a 64-bit Solaris kernel.
 - Provides support for virtual disk volumes.
 - Includes a partition table in which you can define up to ten (0 through 9) disk partitions (slices).
 - Is compatible with the UFS. You can create a UFS that is larger than 1 terabyte.
 - Solaris OS systems do not currently boot from disks that use EFI labels.

Use either SMI or EFI labels if a disk is not being used as a boot disk. Use either the format or prtvtoc command to check whether a disk has an SMI VTOC or EFI label.

Figure 2-3 shows how disk slices might reside on a disk.

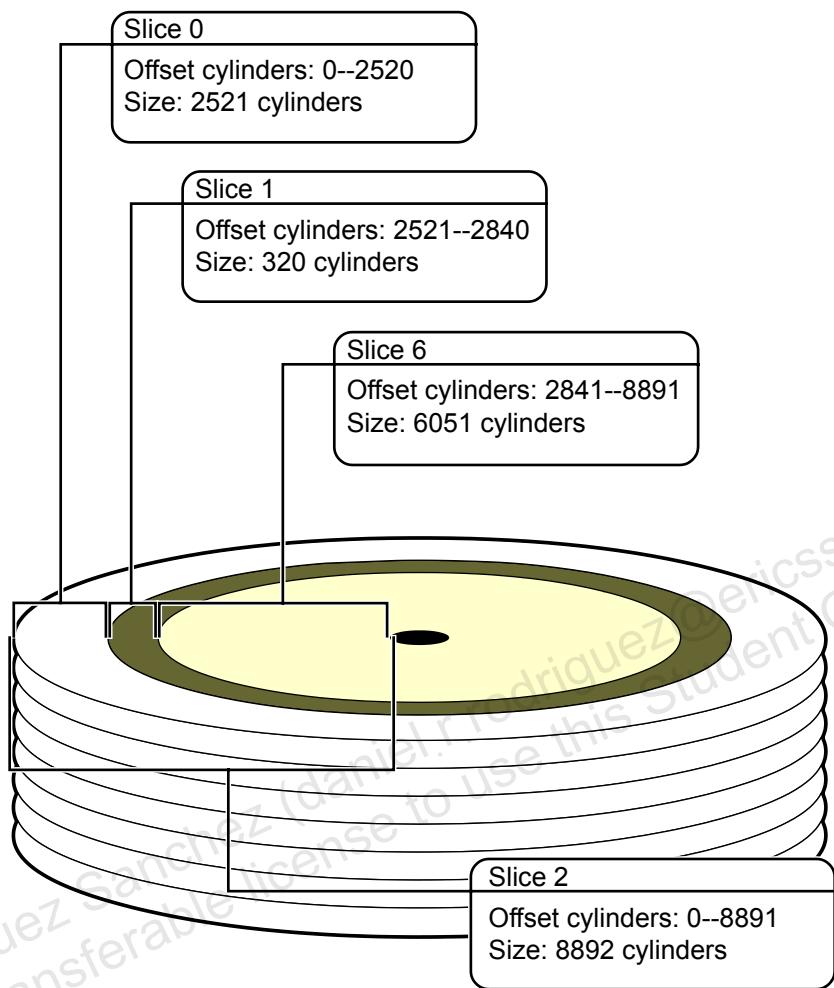


Figure 2-3 Cylinders and Slices

Table 2-1 shows disk slices and the different file systems they could hold.

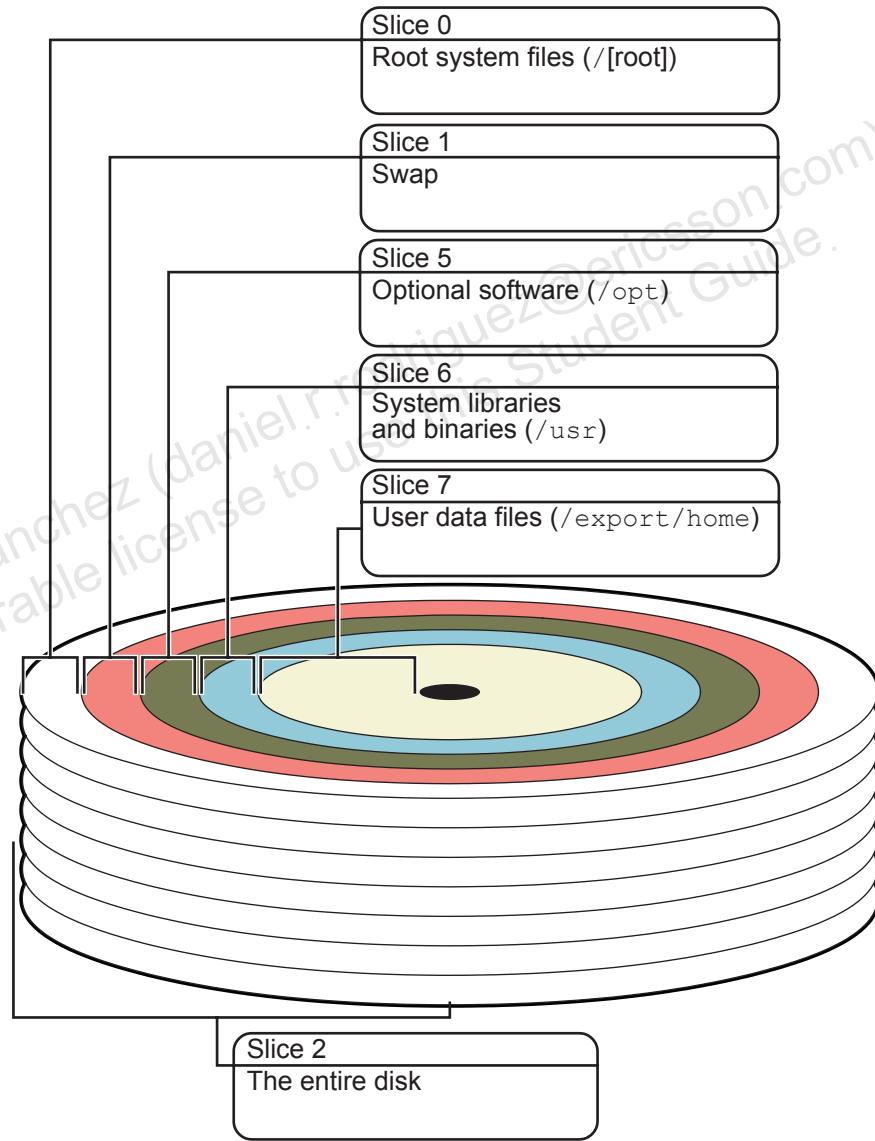
Table 2-1 Disk Slices

Slice	Name	Function
0	/	The root directory's system files
1	swap	Swap area
2		Entire disk
5	/opt	Optional software
6	/usr	System executables and programs

Table 2-1 Disk Slices (Continued)

Slice	Name	Function
7	/export/home	User files and directories

Figure 2-4 shows a possible configuration convention for organizing data. The example disk is divided into slices that logically organize the data on the boot disk.

**Figure 2-4** Top View of Five Configured Disk Slices

x86/x64 Partitions and Slices

On x86/x64 platform Solaris OS systems, disks typically use SMI labels. The boot disk must use an SMI label. Disks on x86/x64 systems maintain an fdisk partition table and a partition table for the Solaris fdisk partition.

Systems based on x86/x64 platforms use the fdisk partition table to identify parts of the disk reserved for different operating systems, and to identify the partition from which the system will boot. Systems boot from the fdisk partition that you designate the *active* fdisk partition. The sections that follow assume that the fdisk partition that holds the Solaris OS is the active fdisk partition.

The fdisk partition table defines up to four fdisk partitions. Disks on x86/x64 systems require at least one defined fdisk partition. You can assign only one fdisk partition on a disk to be used for Solaris.

The fdisk menu in the format (1M) utility is typically used to create or modify the fdisk partition table. The fdisk (1M) command also allows you to directly manipulate the fdisk partition table. It is recommended that you use the fdisk menu in the format utility to create or modify the fdisk partition table.

The second sector of the Solaris fdisk partition contains a partition table that defines slices within the Solaris fdisk partition. The label that contains this partition table is known as the x86/x64 VTOC.

Using the format utility, a Solaris fdisk partition can be divided into ten slices, labeled slice 0 through slice 9. Slices 0 through 7 are used for the same purposes as disk slices found on disks that use SPARC SMI labels. Slice 2 represents all of the space within the Solaris fdisk partition. Slices 8 and 9 are used for purposes specific to x86/x64 hardware.

The format utility automatically creates slices 8 and 9 when you define an x86/x64 VTOC within the Solaris fdisk partition. Slice 8 is always required. Slice 9 will exist, but remains unassigned, on SCSI or Fibre-Channel disks. The format utility prevents modifying slices 8 and 9.

Solaris OS systems on the x86/x64 platform create device files to support sixteen slices for each disk, numbered 0 through 15. The format utility, however, limits its operation to slices 0 through 7. Solaris x86/x64 systems also create device files that represent the four possible fdisk partitions.

Figure 2-5 describes a set of slices within a Solaris fdisk partition. For all slices marked with shaded title bars, cylinder values are expressed relative to the first cylinder in the Solaris fdisk partition. The three objects marked with unshaded title bars describe cylinder values relative to the first cylinder of the entire disk.

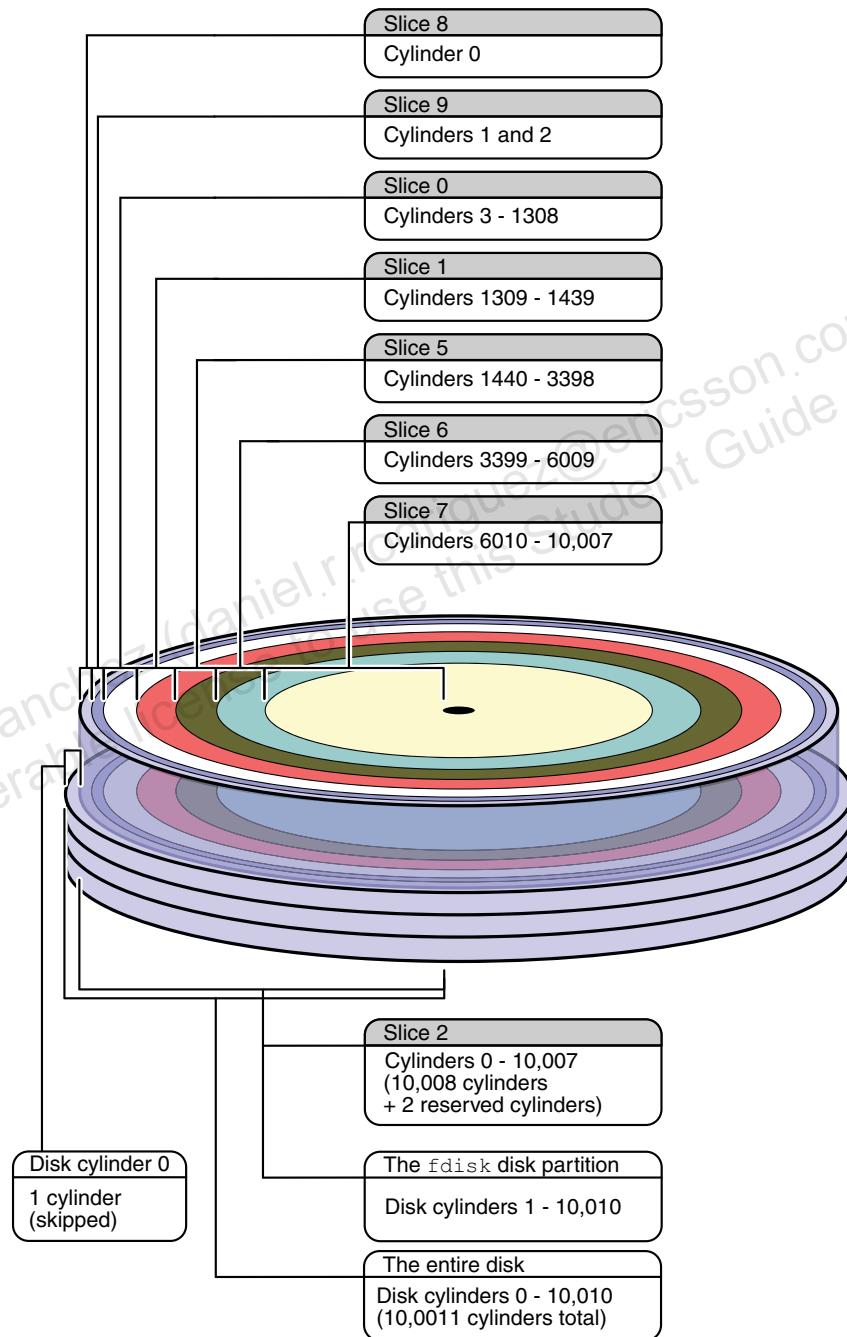


Figure 2-5 Solaris OS Slices on a Disk with a Single fdisk Partition

By default, slice 8 is the boot slice, and contains the GRUB stage1 program in sector 0, the Solaris disk label and VTOC in sectors 1 and 2, and GRUB stage2 program beginning at sector 50. Slice 8 occupies the first cylinder (cylinder 0) of the Solaris fdisk partition.

By convention on IDE and SATA disk drives, Slice 9 is tagged as the alternates slice. If it is defined, slice 9 occupies the second and third cylinders (cylinders 1 and 2) of the Solaris fdisk partition, and contains blocks used to store bad block information.

For an IDE or SATA disk attached to an x86/x64 system, the partition table you can display using the `format` utility shows that slices 8 and 9 are both defined, and that they use one and two cylinders respectively. For example:

Part	Tag	Flag	Cylinders	Size	Blocks
(output omitted)					
8	boot	wu	0 - 0	7.84MB	(1/0/0) 16065
9 alternates		wm	1 - 2	15.69MB	(2/0/0) 32130

For a SCSI disk attached to an x86/x64 system, the partition table shows that slice 8 is defined and uses one cylinder (cylinder 0), but slice 9 has no space assigned. For example:

Part	Tag	Flag	Cylinders	Size	Blocks
(output omitted)					
8	boot	wu	0 - 0	7.84MB	(1/0/0) 16065
9 unassigned		wm	0	0	(0/0/0) 0

Disk Device Naming Conventions

Solaris OS systems use a standard naming convention to identify disk devices. Disk device names identify a series of components that together specify a particular disk device. The disk device name includes the controller, target, disk (or LUN), fdisk partition, and slice that a particular disk device uses.

Examples of disk device names include the following:

- `c1t0d0s0` – A SCSI disk device name, that specifies controller 1, target 0, disk 0, and slice 0
- `c1d0p0` – An IDE or SATA disk name on an x86/x64 system that specifies controller 1, disk 0, and fdisk partition 0

- c1d0s0 – An IDE or SATA disk name that specifies controller 1, disk 0, and slice 0
- c5t12d0p0 – A SCSI disk device name on an x86/x64 system, that specifies controller 5, target 12, disk 0, and fdisk partition 0
- c5t12d0s0 – A SCSI disk device name, that specifies controller 5, target 12, disk 0, and slice 0
- c3t266000C0FFF7C140d31s2 – A Fibre-Channel attached LUN name that specifies controller 3, WWN 266000C0FFF7C140, LUN 31, and slice 2

Disk device names represent the full name of a slice or fdisk partition. These names identify the following components:

Controller number	Identifies the host bus adapter (HBA), which controls communication between the system and disk unit. The HBA takes care of transferring both commands and data to and from the device. Controller numbers are assigned in sequential order, such as c0, c1, c2, and so on.
Target number	Target numbers, such as t0, t1, t2, and t3, identify a unique hardware address assigned to the SCSI target controller of a disk, tape, or CD-ROM. Some external disk drives have an address switch located on the rear panel. Some internal disks have address pins that are jumpered to assign that disk's target number. Names for IDE or SATA disks typically do not use target numbers, although on some systems they do. Fibre-Channel attached disks may use a World Wide Name (WWN) instead of a target number.
Disk number	The disk number is also known as the logical unit number (LUN). This number may vary from d0 if more than one disk can be attached at the target location.
fdisk Partition number	The fdisk partition number ranging from p0 to p4. The number p0 represents the whole physical disk. Device names that include fdisk partition IDs are found only on x86/x64 systems.
Slice number	A slice number ranging from s0 to s7 on SPARC systems, and from s0 to s15 on x86/x64 systems.

SCSI Disk Device Names

SCSI specifications allow multiple disks to be attached to a single target controller. In many cases however, a SCSI disk and target controller are combined into one physical unit. SCSI disks that are constructed this way are described as using an *embedded* SCSI configuration. The disk number is always set to d0 with embedded SCSI disks.

Figure 2-6 shows the string that represents the full name of a SCSI disk slice.

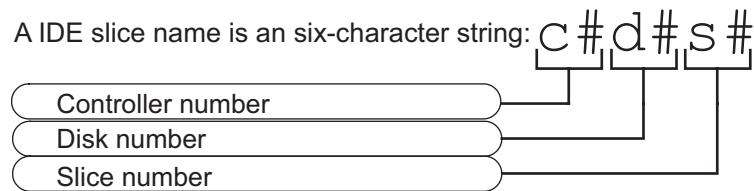


Figure 2-6 SCSI Disk Slice Naming Convention

Figure 2-7 shows the embedded SCSI configuration.

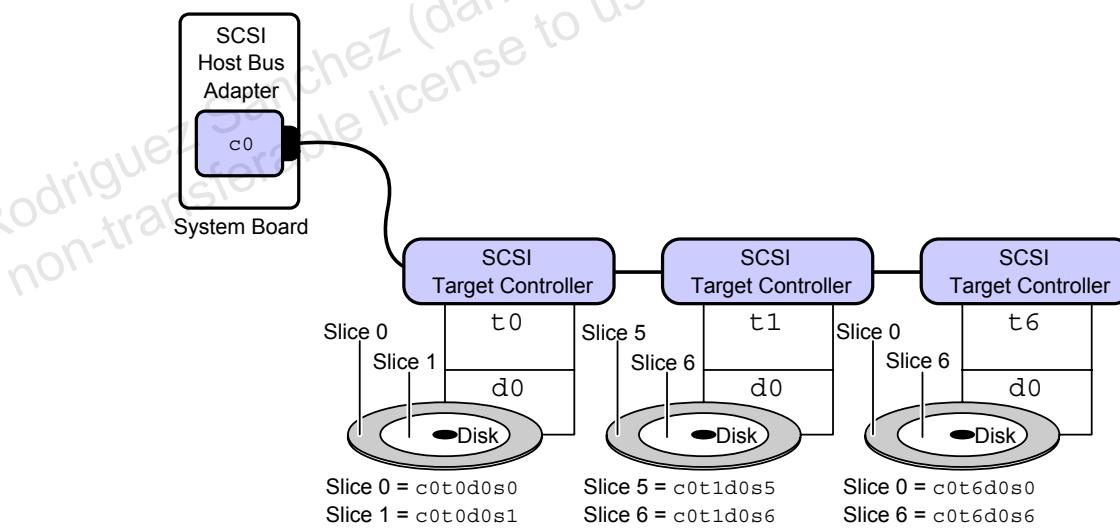


Figure 2-7 Embedded SCSI Configuration

Some external SCSI devices have multiple disks attached to a single target controller. Solaris OS systems use disk numbers to represent these disks, and so may present disk numbers other than d0 (such as d1, d2) in their device names.

Disk arrays that are capable of hardware RAID configurations typically present multiple LUNs for the target values they present to the Solaris OS. LUNs from arrays like these are also represented by `d0 - dn` disk numbers in their device names.

iSCSI Target Devices

The internet small computer system interface (iSCSI) is an internet protocol (IP) based storage networking standard for linking data storage subsystems. By carrying SCSI commands over IP networks, the iSCSI protocol enables you to access block devices from across the network as if they were connected to the local system.

The Solaris 10 8/07 release provides support for iSCSI target devices. These target devices can be disk or tape. Releases prior to Solaris 10 8/07 provided support for iSCSI initiators. The advantage of setting up Solaris iSCSI targets is that you might have existing fibre-channel devices that can be connected to clients without the cost of fibre-channel host bus adapters (HBAs). In addition, systems with dedicated arrays can now export replicated storage with the ZFS or the UFS.

Setting Up iSCSI Target Devices

You can use the `iscsitadm` command to set up and manage your iSCSI target devices. For the disk device that you select as your iSCSI target, you'll need to provide an equivalently sized ZFS or UFS as the backing store for the iSCSI daemon. After the target device is set up, use the `iscsiadm` command to identify your iSCSI targets, which will discover and use the iSCSI target device.

Advantages of Using iSCSI Targets and Initiators

The benefits of using Solaris iSCSI targets and initiators are:

- The iSCSI protocol runs across existing Ethernet networks.
 - You can use any supported network interface card (NIC), Ethernet hub or switch.
 - One IP port can handle multiple iSCSI target devices.
 - You can use existing infrastructure and management tools for IP networks.

- You might have existing fibre-channel devices that can be connected to clients without the cost of fibre-channel HBAs. In addition, systems with dedicated arrays can now export replicated storage with ZFS or UFS file systems.
- There is no upper limit on the maximum number of configured iSCSI target devices.
- The protocol can be used to connect to Fibre Channel or iSCSI Storage Area Network (SAN) environments with the appropriate hardware.

IDE and SATA Disk Device Names

Integrated device electronics (IDE) and serial advanced technology attachment (SATA) disks do not use target controllers. Device names for these disks consequently represent the controller, disk, and slice or fdisk partition used by the IDE or SATA drive.

Figure 2-8 illustrates the IDE and SATA disk slice naming conventions.

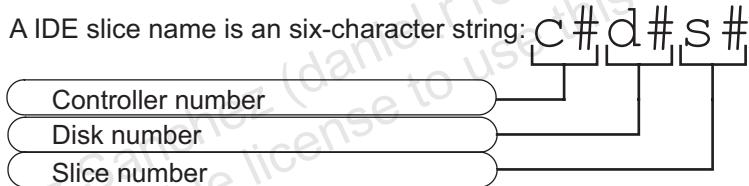


Figure 2-8 IDE and SATA Disk Slice Naming Conventions

For example, Ultra 20 systems support two internal SATA disks, where each disk is attached to a separate IDE controller.

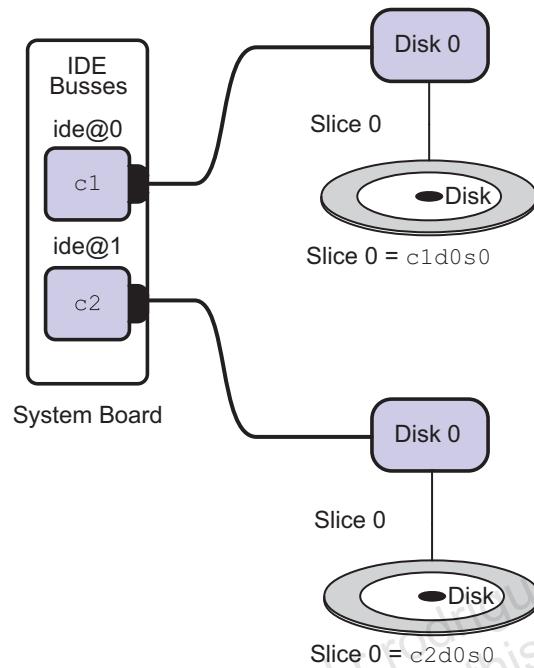
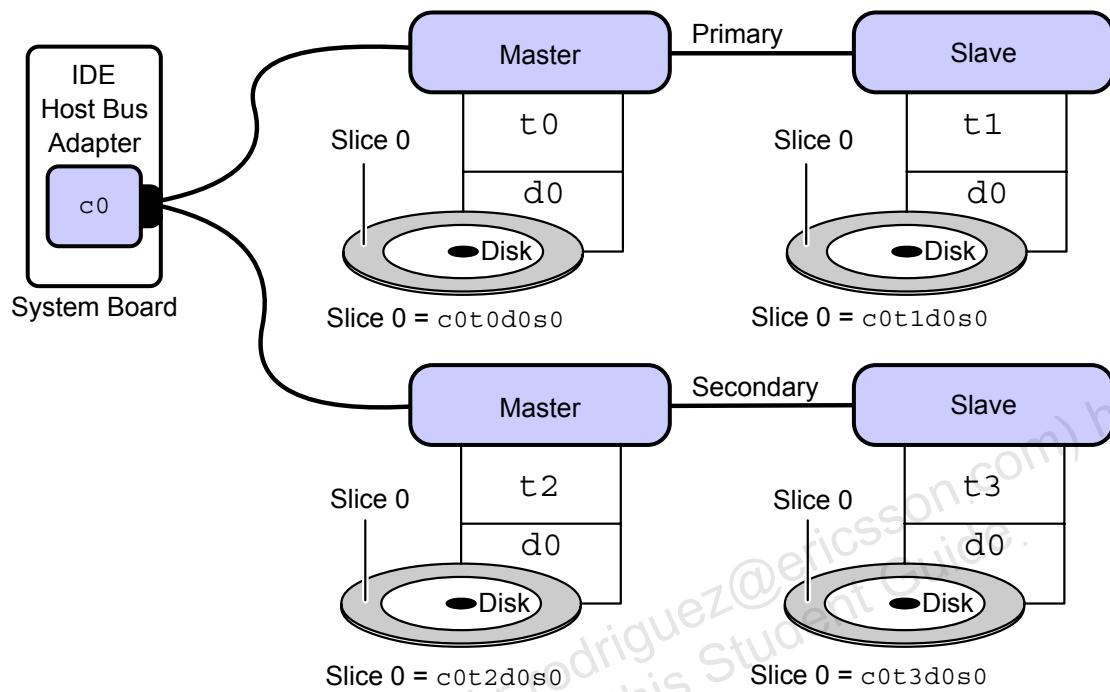


Figure 2-9 Ultra 20 IDE Configuration

Ultra 10 systems present a different example. Although IDE disks do not use target controllers, Ultra 10 systems use a target (t_n) value to represent the identity of the disks on its primary and secondary IDE busses. Target values on these systems denote the following:

- t_0 - master device on primary IDE bus
- t_1 - slave device on primary IDE bus
- t_2 - master device on secondary IDE bus
- t_3 - slave device on secondary IDE bus

Figure 2-10 shows the configuration of the IDE architecture on an Ultra



10.

Figure 2-10 Ultra 10 IDE Configuration

Introducing Solaris OS Device Naming Conventions

In the Solaris OS, all devices are represented by three different types of names, depending on how the device is being referenced:

- Logical device names
- Physical device names
- Instance names

Logical Device Names

Logical disk device names are symbolic links to the physical device names in the /devices directory. Logical device names are used primarily to specify a device when you enter commands on the command line. All logical device names are kept in the /dev directory. Logical device names contain the controller number, target number if required, disk number, fdisk partition number or slice number.

Every disk device has entries in the /dev/dsk and /dev/rdsk directories for its block and character disk devices, respectively. To display the entries in the /dev/dsk directory, perform the command:

```
# ls /dev/dsk
c0t0d0s0  c0t0d0s4  c0t2d0s0  c0t2d0s4  c1t1d0s0  c1t1d0s4
c0t0d0s1  c0t0d0s5  c0t2d0s1  c0t2d0s5  c1t1d0s1  c1t1d0s5
c0t0d0s2  c0t0d0s6  c0t2d0s2  c0t2d0s6  c1t1d0s2  c1t1d0s6
c0t0d0s3  c0t0d0s7  c0t2d0s3  c0t2d0s7  c1t1d0s3  c1t1d0s7
```

- c0t0d0s0 through c0t0d0s7 – Identifies the device names for disk slices 0 through 7 for a disk that is attached to Controller 0, at Target 0, on Disk Unit 0.
- c0t2d0s0 through c0t2d0s7 – Identifies the device names for disk slices 0 through 7 for a disk that is attached to Controller 0, at Target 2, on Disk Unit 0.
- c1t1d0s0 through c1t1d0s7 – Identifies the device names for disk slices 0 through 7 for a disk that is attached to Controller 1, at Target 1, on Disk Unit 0.

On an x86/x64 system, the list of logical device names for a disk will include names for fdisk partitions, and names for sixteen slices. For example:

```
# ls c1d0*
c1d0p0  c1d0p3  c1d0s1  c1d0s12  c1d0s15  c1d0s4  c1d0s7
c1d0p1  c1d0p4  c1d0s10  c1d0s13  c1d0s2   c1d0s5  c1d0s8
c1d0p2  c1d0s0  c1d0s11  c1d0s14  c1d0s3   c1d0s6  c1d0s9
#
```

- c1d0p0 through c1d0p4 – Identify the device names for fdisk partitions 0 through 4, for a disk that is attached to Controller 1, as Disk Unit 0.
- c1d0s0 through c1d0s15 – Identify the device names for slices 0 through 14, for a disk that is attached to Controller 1, as Disk Unit 0.

On an x86/x64 system, the list of logical device names for a LUN from a Fibre-Channel attached storage device will include names for fdisk partitions, and names for sixteen slices. For example:

```
# ls c3t266000C0FFF7C140d31*
c3t266000C0FFF7C140d31p0  c3t266000C0FFF7C140d31s10
c3t266000C0FFF7C140d31s3
c3t266000C0FFF7C140d31p1  c3t266000C0FFF7C140d31s11
c3t266000C0FFF7C140d31s4
c3t266000C0FFF7C140d31p2  c3t266000C0FFF7C140d31s12
c3t266000C0FFF7C140d31s5
c3t266000C0FFF7C140d31p3  c3t266000C0FFF7C140d31s13
c3t266000C0FFF7C140d31s6
c3t266000C0FFF7C140d31p4  c3t266000C0FFF7C140d31s14
c3t266000C0FFF7C140d31s7
c3t266000C0FFF7C140d31s0  c3t266000C0FFF7C140d31s15
c3t266000C0FFF7C140d31s8
c3t266000C0FFF7C140d31s1
c3t266000C0FFF7C140d31s9
#
```

- c3t266000C0FFF7C140d31p0 through c3t266000C0FFF7C140d31p4 – Identify the device names for fdisk partitions 0 through 4, for a LUN that is attached to Controller 3, through an array's Fibre-Channel port 266000C0FFF7C140 as LUN 31.
- c3t266000C0FFF7C140d31s0 through c3t266000C0FFF7C140d31s15 – Identify the device names for slices 0 through 15, for a LUN that is attached to Controller 3, through an array's Fibre-Channel port 266000C0FFF7C140 as LUN 31.

Physical Device Names

Physical device names uniquely identify the physical location of hardware devices on the system and are maintained in the `/devices` directory.

A physical device name contains the hardware information, represented as a series of node names, separated by slashes, that indicate the path through the system's physical device tree to the device. Listing the symbolic links that exist in the `/dev/dsk` and `/dev/rdsk` directories displays the names of the physical device files they point to.

To display a physical device name, use the `ls -l` command to list a logical device name. This example is from a Sun Blade™ 1500 workstation:

```
# ls -l /dev/dsk/c0t0d0s0
lrwxrwxrwx 1 root      root          43 Jan 23 16:06 /dev/dsk/c0t0d0s0 -
> ../../devices/pci@1e,600000/ide@d/dad@0,0:a
#
```

On an Ultra 20, listing the logical device `/dev/dsk/c1d0s0` displays the following physical device name information:

```
# ls -l /dev/dsk/c1d0s0
lrwxrwxrwx 1 root      root          48 Dec 17 19:16 /dev/dsk/c1d0s0 ->
../../devices/pci@0,0/pci-ide@7/ide@0/cmdk@0,0:a
#
```

Adding the `-L` option to the `ls -l` command displays the file information for the physical device file referenced by the logical device file. For example:

```
# ls -lL /dev/dsk/c1d0s0
brw-r---- 1 root      sys          102,   0 Dec 17 20:49 /dev/dsk/c1d0s0
#
```

Some systems use internal Fibre Channel Arbitrated Loop (FC-AL) disks. Their physical device names will appear slightly different because they include a World Wide Name (WWN). The following example was taken from a Sun™ Enterprise 3500 server:

```
# ls -l /dev/rdsk/c0t0d0s0
lrwxrwxrwx 1 root      root          78 Jun 16 2000 /dev/rdsk/c0t0d0s0 ->
../../devices/sbus@2,0/SUNW,socal@d,10000/sf@0,0/ssl@w21000020375b9ab6,0:
a,raw
```

In this example, the string `21000020375b9ab6` is the WWN.

Physical device names for external Fibre-Channel attached storage also display a World Wide Name (WWN). The following example was taken from a Sun™ V20z server attached to a Sun 3510 Fibre-Channel array:

```
# ls -l /dev/rdsk/c3t266000C0FFF7C140d31s2
lrwxrwxrwx 1 root      root          89 Jan 15 17:51
/dev/rdsk/c3t266000C0FFF7C140d31s2 ->
.../..../devices/pci@0,0/pci1022,7450@a/pci1077,10a@5/fp@0,0/disk@w266000c0f
ff7c140,1f:c,raw
#
```

In this example, the string 266000c0fff7c140 is the WWN. These physical device name examples describe the device locations in a system device tree.

Figure 2-11 shows a partial device tree from an Ultra 20 workstation.

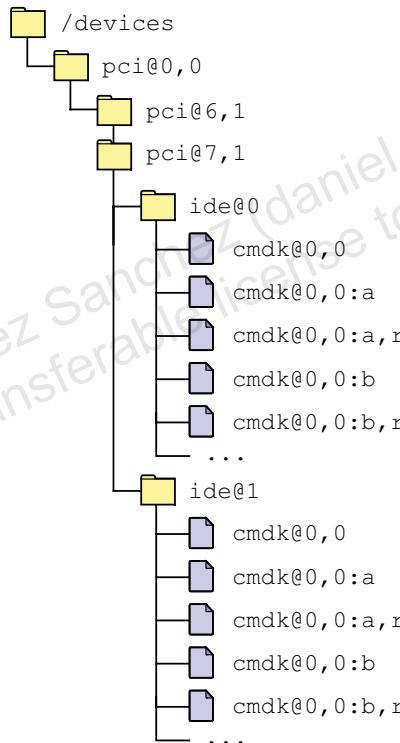


Figure 2-11 The /devices Directory Structure

The device tree can be thought of as existing on two levels. First, there is the device tree as recognized by the hardware at boot time. This device tree consists of all devices that were powered on and accessible to the hardware testing at power-on time.

The second device tree is that known to the Solaris OS kernel. This tree is managed using the various software controls available to the Solaris OS. The devices in this tree must have appropriate device files. If no device file exists for a powered-on device that is physically attached to the system, the kernel does not recognize this device in its device tree until a new device probe is initiated.

Note – Different hardware platforms may have different device trees.



The top-most directory in the hierarchy is called the root node of the device tree. The bus nexus nodes and the leaf nodes below the root object have device drivers associated with them.

A device driver is the software that communicates with the device. This software must be available to the kernel so that the system can use the device.

During system initialization, the kernel identifies the physical location of a device. The kernel associates a node with an address, `nodename@address`, which is the physical device name. In Figure 2-11, `cmdk@0` is the common disk driver disk device at address 0. One of these disk device instances is attached to both `ide@0` and `ide@1` IDE busses.

Instance Names

Instance names are abbreviated names assigned by the kernel for each system device. Device-related messages that display on the system console or that are typically saved in the /var/adm/messages log file often use instance names to identify devices. An instance name is a shortened name for the physical device name. Four instance name examples are shown:

- `sdn`
where `sd` (SCSI disk) is the disk name and `n` is the number, such as `sd0` for the first SCSI disk device.
- `cmdkn`
where `cmdk` (common disk driver) is the disk name, and `n` is the number, such as `cmdk0` for the first SATA disk device in an Ultra 20
- `dadn`
where `dad` (direct access device) is the disk name and `n` is the number, such as `dad0` for the first IDE disk device.
- `atan`
where `ata` (Advanced Technology Attachment), also known as Integrated Drive Electronics (IDE), is the disk name and `n` is the number, such as `ata0` for the first IDE disk device

Listing System Devices

In the Solaris OS, there are several ways to list a system devices, including:

- Using the /etc/path_to_inst file
- Using the prtconf command
- Using the format utility

The /etc/path_to_inst File

For each device, the system records its physical name and instance name in the /etc/path_to_inst file. These names are used by the kernel to identify every possible device. This file is read only at boot time.

The /etc/path_to_inst file is maintained by the kernel, and it is generally not necessary, nor is it advisable, for the system administrator to change this file.

The following example from an Ultra 20 shows entries in the /etc/path_to_inst file. The text within the parentheses indicates what device the entry identifies, and does not appear in the actual file.

```
# cat /etc/path_to_inst
#
# Caution! This file contains critical kernel state
#
"/pseudo" 0 "pseudo"
"/options" 0 "options"
"/xsvc" 0 "xsvc"
"/objmgr" 0 "objmgr"
"/scsi_vhci" 0 "scsi_vhci"
"/isa" 0 "isa"
"/isa/asy@1,3f8" 0 "asy"
"/ramdisk" 0 "ramdisk"
"/pci@0,0" 0 "npe"
"/pci@0,0/pci108e,5347@a" 0 "nge"
"/pci@0,0/pci108e,5347@2" 0 "ohci"
"/pci@0,0/pci108e,5347@2(mouse@1" 0 "hid"
"/pci@0,0/pci108e,5347@2/hub@2" 0 "hubd"
"/pci@0,0/pci108e,5347@2/hub@2/keyboard@4" 1 "hid"
"/pci@0,0/pci108e,5347@2/hub@1" 1 "hubd"
"/pci@0,0/pci108e,5347@2/hub@1/keyboard@4" 2 "hid"
```

Listing System Devices

```

"/pci@0,0/pci108e,5347@2/mouse@2" 3 "hid"
"/pci@0,0/pci108e,5347@2,1" 0 "ehci"
"/pci@0,0/pci108e,5347@4" 0 "audio810"
"/pci@0,0/pci10de,5d@e" 3 "pcie_pci"
"/pci@0,0/pci10de,5d@e/display@0" 0 "vgatext"
"/pci@0,0/pci-ide@6" 0 "pci-ide"
"/pci@0,0/pci-ide@6/ide@0" 0 "ata"
"/pci@0,0/pci-ide@6/ide@0/sd@0,0" 0 "sd" (DVD/CDROM)
"/pci@0,0/pci-ide@7" 1 "pci-ide"
"/pci@0,0/pci-ide@7/ide@0" 2 "ata"
"/pci@0,0/pci-ide@7/ide@0/cmdk@0,0" 0 "cmdk" (Disk 0)
"/pci@0,0/pci-ide@7/ide@1" 3 "ata"
"/pci@0,0/pci-ide@7/ide@1/cmdk@0,0" 1 "cmdk" (Disk 1)
"/pci@0,0/pci-ide@8" 2 "pci-ide"
"/pci@0,0/pci10de,5c@9" 0 "pci_pci"
"/pci@0,0/pci10de,5c@9/pci1106,3044@6" 0 "hci1394"
"/pci@0,0/pci1022,1101@18,1" 0 "mc-amd"
"/pci@0,0/pci1022,1102@18,2" 1 "mc-amd"
"/pci@0,0/pci1022,1103@18,3" 0 "amd64_gart"
"/iscsi" 0 "iscsi"
"/agpgart" 0 "agpgart"
#

```

Instance names correlate to values in the last two fields in /etc/path_to_inst entries. For example, the entry for disk 0 on an Ultra 20 contains this information:

```

"/pci@0,0/pci-ide@7/ide@0/cmdk@0,0" 0 "cmdk"

```

This line and the instance name cmdk0 identify the same device. The field values in bold match the disk name and number in the instance name. The physical device path /pci@0,0/pci-ide@7/ide@0/cmdk@0,0 identifies the path through the device tree to the device.

Different systems have different physical device paths. Entries in the /etc/path_to_inst file reflect those differences. This partial example is from a Sun Blade 1500 workstation:

```

# cat /etc/path_to_inst
#
# Caution! This file contains critical kernel state
#
"/pseudo" 0 "pseudo"
"/scsi_vhci" 0 "scsi_vhci"
"/options" 0 "options"
"/pci@1e,600000" 0 "pcisch"

```

```

"/pci@1e,600000/pmu@6" 0 "pmubus"
"/pci@1e,600000/pmu@6/ppm@0,b3" 0 "m1535ppm"
"/pci@1e,600000/pmu@6/beep@0,b2" 0 "grbeep"
"/pci@1e,600000/pmu@6/i2c@0,0" 0 "smbus"
"/pci@1e,600000/pmu@6/i2c@0,0/card-reader@40" 0 "scmi2c"
"/pci@1e,600000/isa@7" 0 "ebus"
"/pci@1e,600000/isa@7/i2c@0,320" 0 "pcf8584"
"/pci@1e,600000/isa@7/i2c@0,320/clock-generator@0,d2" 0 "ics951601"
"/pci@1e,600000/isa@7/i2c@0,320/hardware-monitor@0,5c" 0 "adm1031"
"/pci@1e,600000/isa@7/i2c@0,320/motherboard-fru-prom@0,a8" 0 "seeprom"
"/pci@1e,600000/isa@7/i2c@0,320/dimm-spd@0,a0" 1 "seeprom"
"/pci@1e,600000/isa@7/i2c@0,320/dimm-spd@0,a2" 2 "seeprom"
"/pci@1e,600000/isa@7/power@0,800" 0 "power"
"/pci@1e,600000/isa@7/serial@0,3f8" 0 "su"
"/pci@1e,600000/isa@7/serial@0,2e8" 1 "su"
"/pci@1e,600000/isa@7/dma@0,0" 0 "isadma"
"/pci@1e,600000/isa@7/dma@0,0/parallel@0,378" 0 "ecpp"
"/pci@1e,600000/usb@a" 0 "ohci"
"/pci@1e,600000/usb@b" 1 "ohci"
"/pci@1e,600000/usb@b/keyboard@1" 2 "hid"
"/pci@1e,600000/usb@b/mouse@2" 0 "hid"
"/pci@1e,600000/pci@2" 0 "pci_pci"
"/pci@1e,600000/pci@2/usb@8,2" 0 "ehci"
"/pci@1e,600000/pci@2/usb@8" 2 "ohci"
"/pci@1e,600000/pci@2/usb@8,1" 3 "ohci"
"/pci@1e,600000/pci@2/firewire@b" 0 "hci1394"
"/pci@1e,600000/sound@8" 0 "audiots"
"/pci@1e,600000/ide@d" 0 "uata"
"/pci@1e,600000/ide@d/sd@2,0" 1 "sd" (DVD/CDROM)
"/pci@1e,600000/ide@d/dad@0,0" 1 "dad" (Disk 0)
"/pci@1e,600000/ide@d/dad@1,0" 0 "dad" (Disk 1)

```



Note – This example illustrates that the numeric values in instance names may not necessarily match the values found in logical or physical device names. On the Sun Blade 1500, the disk 0 (c0t0d0) is referenced by an entry in /etc/path_to_inst for the dad1 instance name.

The following example is from an /etc/path_to_inst file with an entry for an FC-AL disk:

```
"/sbus@2,0/SUNW,socal@d,10000/sf@0,0/ssd@w21000020375b9ab6,0" 0 "ssd"
```

The following example is from an /etc/path_to_inst file with an entry for an external Fibre-Channel attached LUN:

```
"/pci@0,0/pci1022,7450@a/pci1077,10a@5/fp@0,0/disk@w266000c0fff7c140,1f"  
212 "sd"
```

The prtconf Command

Use the prtconf command to display the system's configuration information, including the total amount of memory installed and the configuration of system peripherals, which is formatted as a device tree.

The prtconf command lists all possible instances of devices, whether the device is attached or not attached to the system. Use the prtconf and grep commands to list only the attached devices on the system. This example is from an Ultra 20 system:

```
# prtconf | grep -v not  
System Configuration: Sun Microsystems i86pc  
Memory size: 4095 Megabytes  
System Peripherals (Software Nodes):  
  
i86pc  
scsi_vhci, instance #0  
isa, instance #0  
    asy, instance #0  
pci, instance #0  
    pci108e,5347, instance #0  
        hub, instance #1  
            keyboard, instance #2  
            mouse, instance #3  
    pci108e,5347, instance #0  
    pci108e,5347, instance #0  
    pci-ide, instance #0  
        ide, instance #0  
            sd, instance #0  
    pci-ide, instance #1  
        ide, instance #2  
            cmdk, instance #0  
        ide, instance #3  
            cmdk, instance #1
```

```
pci-ide, instance #2
pci10de,5c, instance #0
    pci1106,3044, instance #0
    pci108e,5347, instance #0
(output truncated)
```



Note – The grep -v not command is used to omit all lines containing the word “not” from the output (such as driver not attached).

The format Utility

Use the format utility to display both logical and physical device names for all currently available disks. This example output is from a Sun Blade 1500 system:

```
# format
Searching for disks...done
```

AVAILABLE DISK SELECTIONS:

0. c0t0d0 <HDS722512VLAT80 cyl 57459 alt 2 hd 16 sec 255>
 /pci@1e,600000/ide@d/dad@0,0
1. c0t1d0 <HDS722512VLAT80 cyl 57459 alt 2 hd 16 sec 255>
 /pci@1e,600000/ide@d/dad@1,0

Specify disk (enter its number): ^D
#



Note – Press Control-D to exit the format utility without selecting a disk.

Here is an example of the `format` utility output from an Ultra 20 system:

```
# format
Searching for disks...done

AVAILABLE DISK SELECTIONS:
 0. c1d0 <DEFAULT cyl 39887 alt 2 hd 64 sec 63>
    /pci@0,0/pci-ide@7/ide@0/cmdk@0,0
 1. c2d0 <DEFAULT cyl 10008 alt 2 hd 255 sec 63>
    /pci@0,0/pci-ide@7/ide@1/cmdk@0,0
Specify disk (enter its number) : ^D
#
```

The `prtdiag` command

Use the `prtdiag` command to display system configuration and diagnostic information. See the `prtdiag(1M)` man page for more information.

```
# prtdiag
System Configuration: Sun Microsystems sun4v Sun Fire T200
Memory size: 32640 Megabytes
===== Virtual CPUs =====
=====

CPU ID Frequency Implementation Status
----- -----
0      1200 MHz SUNW,UltraSPARC-T1 on-line
1      1200 MHz SUNW,UltraSPARC-T1 on-line
2      1200 MHz SUNW,UltraSPARC-T1 on-line
3      1200 MHz SUNW,UltraSPARC-T1 on-line
4      1200 MHz SUNW,UltraSPARC-T1 on-line
5      1200 MHz SUNW,UltraSPARC-T1 on-line
6      1200 MHz SUNW,UltraSPARC-T1 on-line
7      1200 MHz SUNW,UltraSPARC-T1 on-line
8      1200 MHz SUNW,UltraSPARC-T1 on-line
9      1200 MHz SUNW,UltraSPARC-T1 on-line
10     1200 MHz SUNW,UltraSPARC-T1 on-line
11     1200 MHz SUNW,UltraSPARC-T1 on-line
12     1200 MHz SUNW,UltraSPARC-T1 on-line
13     1200 MHz SUNW,UltraSPARC-T1 on-line
14     1200 MHz SUNW,UltraSPARC-T1 on-line
15     1200 MHz SUNW,UltraSPARC-T1 on-line
16     1200 MHz SUNW,UltraSPARC-T1 on-line
```

17	1200	MHz	SUNW,UltraSPARC-T1	on-line
18	1200	MHz	SUNW,UltraSPARC-T1	on-line
19	1200	MHz	SUNW,UltraSPARC-T1	on-line
20	1200	MHz	SUNW,UltraSPARC-T1	on-line
21	1200	MHz	SUNW,UltraSPARC-T1	on-line
22	1200	MHz	SUNW,UltraSPARC-T1	on-line
23	1200	MHz	SUNW,UltraSPARC-T1	on-line
24	1200	MHz	SUNW,UltraSPARC-T1	on-line
25	1200	MHz	SUNW,UltraSPARC-T1	on-line
26	1200	MHz	SUNW,UltraSPARC-T1	on-line
27	1200	MHz	SUNW,UltraSPARC-T1	on-line
28	1200	MHz	SUNW,UltraSPARC-T1	on-line
29	1200	MHz	SUNW,UltraSPARC-T1	on-line
30	1200	MHz	SUNW,UltraSPARC-T1	on-line
31	1200	MHz	SUNW,UltraSPARC-T1	on-line

===== Physical Memory Configuration =====

Segment Table:

Base Address	Segment Size	Interleave Factor	Bank Size	Contains Modules
0x0	32 GB	4	4 GB	MB/CMP0/CH0/R0/D0 MB/CMP0/CH0/R0/D1
			4 GB	MB/CMP0/CH0/R1/D0 MB/CMP0/CH0/R1/D1
			4 GB	MB/CMP0/CH1/R0/D0 MB/CMP0/CH1/R0/D1
			4 GB	MB/CMP0/CH1/R1/D0 MB/CMP0/CH1/R1/D1
			4 GB	MB/CMP0/CH2/R0/D0 MB/CMP0/CH2/R0/D1
			4 GB	MB/CMP0/CH2/R1/D0 MB/CMP0/CH2/R1/D1
			4 GB	MB/CMP0/CH3/R0/D0 MB/CMP0/CH3/R0/D1
			4 GB	MB/CMP0/CH3/R1/D0 MB/CMP0/CH3/R1/D1

===== IO Configuration =====

Location Model	Type	Slot	Path	IO	Name

Listing System Devices

```
-----  
IOBD/NET0      PCIE IOBD          /pci@780/pci@0/pci@1/network@0  
network-pciex8086,105e  
IOBD/NET1      PCIE IOBD          /pci@780/pci@0/pci@1/network@0,1  
network-pciex8086,105e  
IOBD/PCIE      PCIE IOBD          /pci@780/pci@0/pci@9/scsi@0  
scsi-pciex1000,56 LSI,1064E  
IOBD/NET2      PCIE IOBD          /pci@7c0/pci@0/pci@2/network@0  
network-pciex8086,105e  
IOBD/NET3      PCIE IOBD          /pci@7c0/pci@0/pci@2/network@0,1  
network-pciex8086,105e  
IOBD/PCIX      PCIX IOBD          /pci@7c0/pci@0/pci@1/pci@0/isa@2  
isa  
IOBD/PCIX      PCIX IOBD          /pci@7c0/pci@0/pci@1/pci@0/usb@5  
usb-pciiclass,0c0310  
IOBD/PCIX      PCIX IOBD          /pci@7c0/pci@0/pci@1/pci@0/usb@6  
usb-pciiclass,0c0310  
IOBD/PCIX      PCIX IOBD          /pci@7c0/pci@0/pci@1/pci@0/ide@8  
ide-pci10b9,5229
```

===== Environmental Status

Fan sensors:

All fan sensors are OK.

Temperature sensors:

All temperature sensors are OK.

Current sensors:

All current sensors are OK.

Current indicators:

All current indicators are OK.

Voltage sensors:

All voltage sensors are OK.

===== FRU Status =====

Location	Name	Status
0651NNN1R6:CH	MB	disabled
0651NNN1R6:CH/PS1	PS	disabled

Reconfiguring Devices

You can cause a system to recognize a newly-added device in a number of ways, including invoking a reconfiguration boot, or running the `devfsadm` command.

Performing a Reconfiguration Boot

The reconfiguration boot process adds new device entries to a system's `/etc/path_to_inst` file, and new symbolic links and device files to the `/dev` and `/devices` directories.

For example, the following steps reconfigure a system to recognize a new external disk.

1. Create the `/reconfigure` file. This file causes the system to check for the presence of any newly installed devices the next time it is powered on or booted.

```
# touch /reconfigure
#
```
2. Use the `init 5` command to shut down the system. This command safely powers off the system, allowing for addition or removal of devices. (If the device is already attached to your system, you can use the `init 0` command to shut down without powering off.)

```
# init 5
```
3. Install the disk device. Make sure that the address of the device you add does not conflict with the address of other devices on the system.
4. Turn on the power to all external devices.
5. Power on and boot the system.
6. Log in, and use either the `prtconf` command or the `format` utility to verify that the disk device has been added.

After the disk is recognized by the system, begin the process of defining disk slices.



Note – On a SPARC-based system, if the `/reconfigure` file was not created before the system was shut down, you can invoke a manual reconfiguration boot with the open boot PROM (OBP) level command: `boot -r`. On an x86/x64 system that boots using GRUB, you can edit the kernel command associated with the desired boot selection to add the `-r` option, and then boot the system.

Using the `devfsadm` Command

Many systems run critical applications on a 24-hour, 7-day-a-week basis. It might not be possible to perform a reconfiguration boot on these systems. In this situation, you can use the `devfsadm` command.

The `devfsadm` command attempts to load every driver in the system and attach all possible device instances. For new devices it finds, the `devfsadm` command then creates the appropriate physical device files in the `/devices` directory and symbolic links in the `/dev` directory. In addition to managing these files, the `devfsadm` command also maintains the `/etc/path_to_inst` file.

```
# devfsadm
```

To restrict the operation of the `devfsadm` command to a specific device class, use the `-c` option.

```
# devfsadm -c device_class
```

The values for `device_class` include `disk`, `tape`, `port`, `audio`, and `pseudo`. For example, to restrict the `devfsadm` command to the `disk` device class, perform the command:

```
# devfsadm -c disk
```

Use the `-c` option more than once on the command line to specify multiple device classes. For example, to specify the `disk`, `tape`, and `audio` device classes, perform the command:

```
# devfsadm -c disk -c tape -c audio
```

To restrict the use of the **devfsadm** command to configure only devices for a named driver, use the **-i** option, and the name of the specific driver.

```
# devfsadm -i driver_name
```

The following examples use the **-i** option.

- To configure only those disks supported by the **sd** driver, perform the command:

```
# devfsadm -i sd
```

- To configure devices supported by the **st** driver, perform the command:

```
# devfsadm -i st
```

You can add the **-v** option to **devfsadm** commands to display verbose output of changes to the device tree.

```
# devfsadm -v
```

To invoke cleanup routines that remove unreferenced symbolic links and device files for devices no longer attached to the system, perform the command:

```
# devfsadm -C
```

Partitioning a Hard Disk

Use the `format` utility to prepare hard disk drives for use with the Solaris OS. Although the `format` utility performs other disk management activities, the main function of the `format` utility is to manage disk partitions and slices.

The procedures used to partition disks on x86/x64 systems require the additional steps to create a Solaris `fdisk` partition before you can create slices within that partition. SPARC systems do not use `fdisk` partitions.

The same principles of allocating disk space apply to creating slices within a Solaris `fdisk` partition on an x86/x64 system as apply to creating slices on a disk in a SPARC system. This section describes disk space concepts, and possible undesirable conditions you can create when you define slices using the `format` utility.

Note – You do not need to partition the disk before you install the Solaris OS. Disk partitioning is part of the installation process.



Introducing Disk Partitioning

To divide a disk into `fdisk` partitions or slices:

1. Identify the correct disk.
2. Plan the layout of the disk.
3. On x86/x64 systems, use the `fdisk` menu in the `format` utility to create the `fdisk` partitions you require, and save the `fdisk` partition table to disk.
4. Use the partition menu in the `format` utility to divide the disk or Solaris `fdisk` partition into slices.
5. Label the disk or Solaris `fdisk` partition with new slice information.

Users with privileges can use the `format` utility. If a regular user runs the `format` utility, the following error message displays:

```
$ /usr/sbin/format
Searching for disk...done
No permission (or no disk found) !
```

Recognizing Disk Space and Undesirable Conditions

Slices are defined by an offset and a size, expressed in whole cylinders. On SPARC systems, the offset is the distance from the disk's first cylinder, cylinder 0. On x86/x64 systems, slice offsets are specified from the first cylinder of the Solaris fdisk partition.

Figure 2-12 shows an example of disk slice sizes and offsets on a SPARC system.

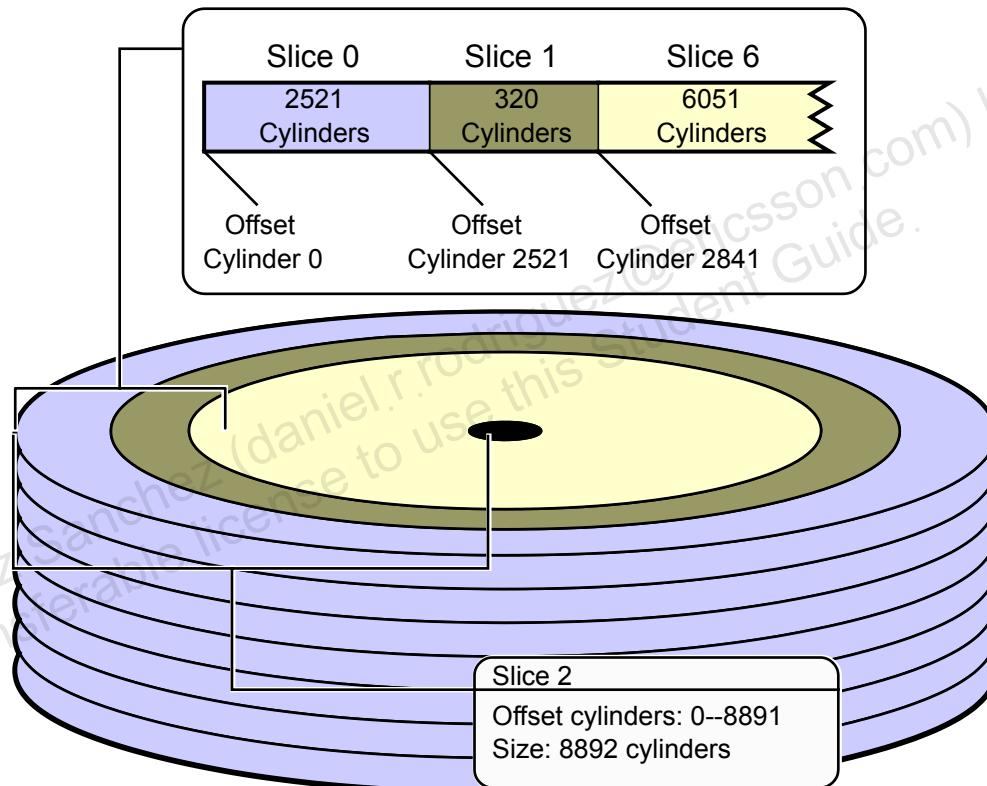


Figure 2-12 Offsets and Sizes for Disk Partitions on SPARC Systems

The offset for Slice 0 is 0 cylinders, and its size is 2521 cylinders. Slice 0 begins on Cylinder 0 and ends on Cylinder 2520.

The offset for Slice 1 is 2521 cylinders, and its size is 320 cylinders. Slice 1 begins on Cylinder 2521 and ends on Cylinder 2840.

The offset for Slice 6 is 2841 cylinders, and its size is 6051 cylinders. Slice 6 begins on Cylinder 2841 and ends on the last available cylinder, which is Cylinder 8891.

Partitioning a Hard Disk

As viewed with the format utility, slices within Solaris fdisk partitions on an x86/x64 systems are also offset from cylinder 0. On these systems however, the format utility presents the first cylinder of the Solaris fdisk partition as cylinder 0.

On IDE or SATA disks in x86/x64 systems, slice 8 occupies cylinder 0, and slice 9 occupies cylinders 1 and 2 of the Solaris fdisk partition. The format utility prevents you from changing these two slice definitions.

Figure 2-13 shows an example of disk slice sizes and offsets on an x86/x64 system.

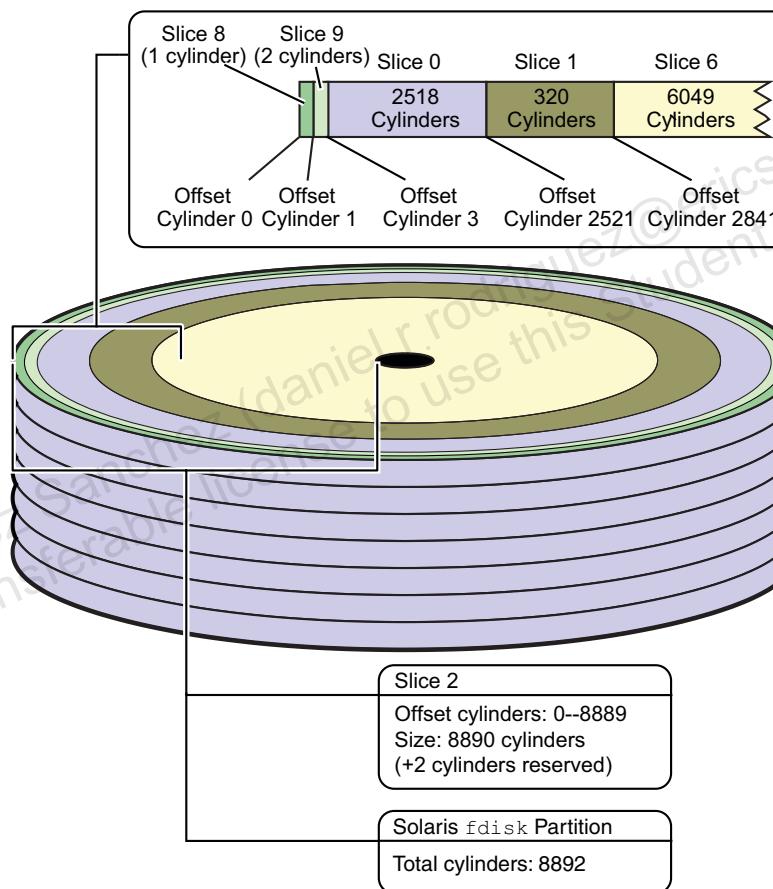


Figure 2-13 Offsets and Sizes for Disk Partitions on x86/x64 Systems

Additional slices you specify within the Solaris fdisk partition use the same offset and cylinder count method to specify their positions and sizes as used on SPARC systems. Because slice 8 and 9 occupy the first three cylinders of the Solaris fdisk partition, the first available cylinder for additional slice definitions is cylinder 3.

Slice 2 starts at cylinder 0 and ends on the last available cylinder of the Solaris fdisk partition. Typically two cylinders are reserved from the total number of cylinders that define the Solaris fdisk partition size.

Recognizing Wasted Disk Space

Wasted disk space occurs when one or more cylinders are not allocated to any disk slice. This undesirable space allocation can exist on SPARC and x86/x64 systems.

Figure 2-14 shows a disk with cylinders that are not allocated.

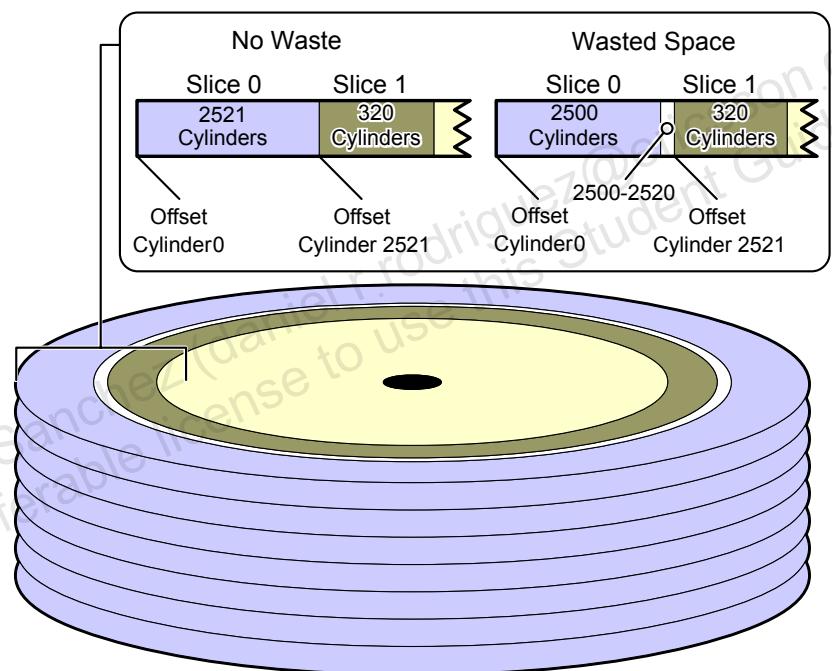


Figure 2-14 A Disk Slice With Wasted Space

Because the cylinders are not allocated to the disk slice, Cylinders 2500 through 2520 are unusable.

Wasted disk space occurs during partitioning when one or more cylinders have not been allocated to a disk slice. This might happen intentionally or accidentally. If there are unallocated slices available, then wasted space can possibly be assigned to a slice later on.

Recognizing Overlapping Disk Slices

Overlapping disk slices occur when one or more cylinders are allocated to more than one disk slice. This undesirable allocation of space can exist on both SPARC and x86/x64 systems.

Figure 2-15 shows a disk with cylinders allocated to more than one disk slice.

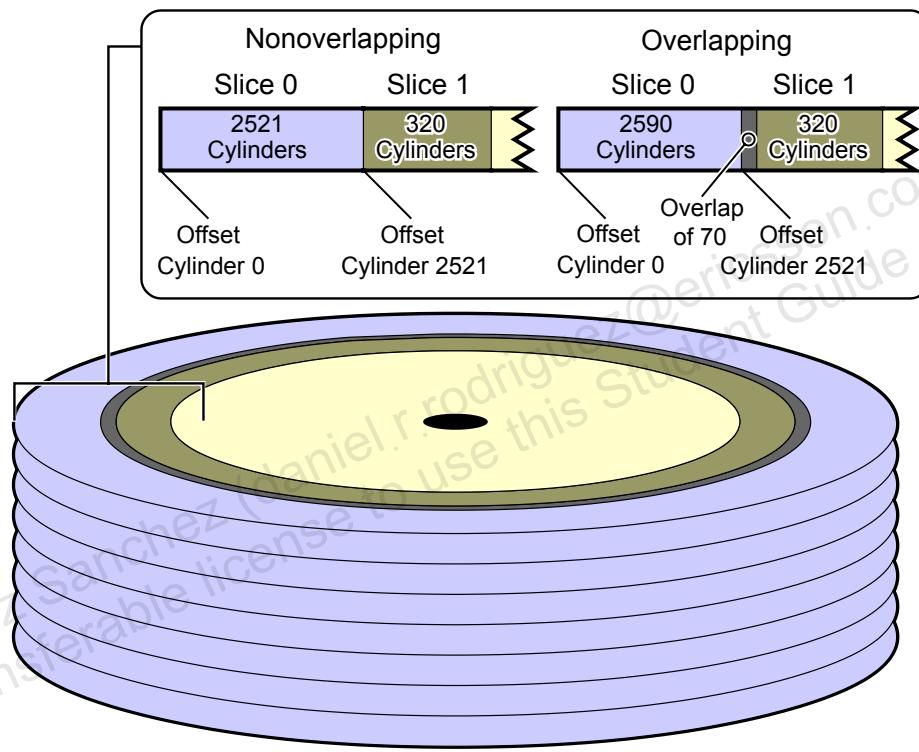


Figure 2-15 Disk Slices With Overlapping Cylinders

In Figure 2-15, Cylinders 2521 through 2590 are overlapping two disk slices.

This condition occurs when the size of one slice is increased and the starting cylinder number of the next slice is not adjusted. Only the `format` utility `modify` command warns you of overlapping slices.

```
partition> modify
Select partitioning base:
 0. Current partition table (unnamed)
 1. All Free Hog
Choose base (enter number) [0]? 0
```



Warning: Overlapping partition (1) in table.
Warning: Fix, or select a different partition table.

Caution – Do not change the size of slices that are currently in use. When a disk with existing slices is repartitioned and re-labeled, any existing data can become inaccessible. Copy existing data to backup media before the disk is repartitioned, and restore the data to the disk after the disk is re-labeled and contains a new file system.



Note – If two slices overlap, when data is saved into one of the slices, data could be overwritten in the other slice located on the tracks in the same disk cylinder.

The format Utility and Disk Partition Tables

The format utility allows you to modify two partition tables on disk;

- fdisk partition tables
- Solaris OS partition tables (SPARC VTOC and x86/x64 VTOC)

fdisk Partition Tables

The fdisk partition table defines up to four fdisk partitions on disk. Solaris OS systems on the x86/x64 platform use the fdisk partition table to identify parts of the disk reserved for different operating systems, and to identify the partition from which the x86/x64 system will boot.

Only x86/x64 systems use fdisk partition tables. You use the fdisk menu in the format utility to view and modify fdisk partition tables.

Solaris OS Partition Tables

The SPARC VTOC and x86 and x64 VTOC define the slices that the Solaris OS will use on a disk. You use the partition menu in the format utility to view and modify these partition tables.

As the root user, when you use the format utility and select a disk to modify, a copy of the VTOC is read into memory, and is displayed as the current partition table in the partition menu.

The source of this partition table differs between SPARC and x86/x64 systems:

- SPARC systems read the VTOC from the first sector of the disk.
- x86/x64 systems read the VTOC from the second sector of the Solaris fdisk partition.

The partition menu in the format utility allows you to define slices and then save the slice definitions into the VTOC on disk.

Note – The partition menu describes writing the VTOC as *labeling* the disk. To label a disk means to write slice information into the VTOC. Keep in mind that the partition menu affects the VTOC that defines Solaris OS slices, and does not affect the fdisk partition table.

If you fail to label the disk or Solaris fdisk partition after defining slices, the slice information is lost.

The format utility also works with a file called /etc/format.dat, which is read when you run the format utility.

The /etc/format.dat file is a table of available disk types and a set of predefined partition tables that you can use to partition a disk quickly.

Note – The partition menu uses the term *partition* exclusively, instead of *slice*. In describing working with this menu, the terms *partition* and *slice* mean the same thing.

Figure 2-16 shows the relationship among the VTOC, the current label in memory, and the predefined partition table in the /etc/format.dat file.

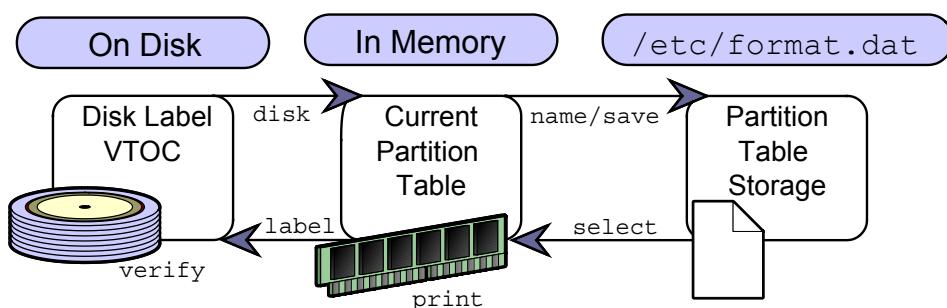


Figure 2-16 Solaris OS Partition Table Locations

The save function in the format utility writes partition table information into `./format.dat` by default. To save partition table information into the `/etc/format.dat` file, you must either specify the absolute pathname of the `/etc/format.dat` file when prompted, or run the format utility while located in the `/etc` directory and accept the default file name presented by the save function. A `format.dat` file exists in the `/etc` directory in the Solaris OS by default.

Table 2-2 describes terms presented by the partition menu in the format utility.

Table 2-2 Terms in the partition Menu

Term	Description
Part	The slice number. The format utility only allows you to modify slice numbers 0 through 7.
Tag	A value that indicates how the slice is being used. This field is no longer used by Solaris, but Veritas Volume Manager assigns specific tags to slices it uses. 0 = unassigned 1 = boot 2 = root 3 = swap 4 = usr 5 = backup 6 = stand 8 = home 9 = alternates Veritas Volume Manager array tags: 14 = public region 15 = private region
Flag	This field is no longer used by Solaris. Historical values are defined as follows: 00 <code>wm</code> = The disk slice is writable and mountable. 01 <code>wu</code> = The disk slice is writable and unmountable. <i>This is the default state of slices dedicated for swap areas.</i> 10 <code>rm</code> = The disk slice is read-only and mountable. 11 <code>ru</code> = The disk slice is read-only and unmountable.
Cylinders	The starting and ending cylinder numbers for the slice.
Size	The slice size: Mbytes (MB), Gbytes (GB), blocks (b), or cylinders (c).

Table 2-2 Terms in the partition Menu

Blocks	The space assigned to the slice, expressed as cylinders/tracks/sectors assigned to the slice. This value typically only reflects a whole number of cylinders.
--------	---

Using the format Utility

The format utility presents a menu-driven, command-line interface. that initially presents the FORMAT menu. Commands in the FORMAT menu allow you to, among other functions, select a disk, select the fdisk menu, select the partition menu, save new disk and partition definitions to a file, and label the disk. The FORMAT menu displays the `format>` prompt.

Typing the name of a command or menu runs the command or accesses the menu you specify. Commands and menus indicate how to return to the FORMAT menu. For most, entering `quit` or `q` returns to the FORMAT menu.

For example, from the `format>` prompt, entering `partition`, or just `part` takes you to the PARTITION menu. The `partition>` prompt displays. This menu allows you to define Solaris OS slices, print the existing partition table, and write the VTOC to disk. Entering `quit` or `q` returns you to the FORMAT menu.

Defining Slices on a SPARC System



Caution – Do not change the size of disk slices that are currently in use.

The following steps demonstrate how to divide a disk into slices:

1. As the root user, run the `format` utility.

```
# format
```

Searching for disks...done

AVAILABLE DISK SELECTIONS:

- 0. c0t0d0 <HDS722512VLAT80 cyl 57459 alt 2 hd 16 sec 255>
/pci@1e,600000/ide@d/dad@0,0
- 1. c0t1d0 <HDS722512VLAT80 cyl 57459 alt 2 hd 16 sec 255>
/pci@1e,600000/ide@d/dad@1,0

Specify disk (enter its number):

The `format` utility searches for all attached disks that are powered on. For each disk it finds, the `format` utility displays the logical device name, Sun marketing name, physical parameters, and physical device name.

2. Choose the second disk by selecting the number located to the left of that disk's logical device name. From the preceding display, the number chosen is 1. The `format` utility's main menu appears.

Specify disk (enter its number): **1**

selecting c0t1d0

[disk formatted, no defect list found]

FORMAT MENU:

disk	- select a disk
type	- select (define) a disk type
partition	- select (define) a partition table
current	- describe the current disk
format	- format and analyze the disk
repair	- repair a defective sector
show	- translate a disk address
label	- write label to the disk
analyze	- surface analysis
defect	- defect list management
backup	- search for backup labels
verify	- read and display labels
save	- save new disk/partition definitions

Partitioning a Hard Disk

```
volname      - set 8-character volume name  
!<cmd>      - execute <cmd>, then return  
quit  
format>
```

The specific menu selections that you can use to view, change, or commit slices include the following:

partition	Displays the Partition menu
label	Writes the current slice definition to the VTOC
verify	Reads and displays the disk label
quit	Exits the format utility

3. Type partition at the format> prompt. The PARTITION menu displays.

```
format> partition
```

PARTITION MENU:

0	- change `0' partition
1	- change `1' partition
2	- change `2' partition
3	- change `3' partition
4	- change `4' partition
5	- change `5' partition
6	- change `6' partition
7	- change `7' partition
select	- select a predefined table
modify	- modify a predefined partition table
name	- name the current table
print	- display the current table
label	- write partition map and label to the disk
!<cmd>	- execute <cmd>, then return
quit	

```
partition>
```

The PARTITION menu enables you to perform the following functions:

0–7	Specify the offset and size of up to eight slices
select	Choose a predefined partition table from the /etc/format.dat file
modify	Change the current partition table in memory
name	Provide a means to identify the partition table in the /etc/format.dat file
print	Display the current partition table in memory
label	Write the current partition table to the VTOC
!<cmd>	Escape from the format utility and execute a command from the shell

4. Type print at the partition prompt to display the VTOC that was copied to random access memory (RAM) when the format utility was invoked.

```
partition> print
```

Current partition table (original):

Total disk cylinders available: 57459 + 2 (reserved cylinders)

Part	Tag	Flag	Cylinders	Size	Blocks
0	unassigned	wm	0	0	(0/0/0) 0
1	unassigned	wm	0	0	(0/0/0) 0
2	backup	wm	0 - 57458	111.79GB	(57459/0/0) 234432720
3	unassigned	wm	0	0	(0/0/0) 0
4	unassigned	wm	0	0	(0/0/0) 0
5	unassigned	wm	0	0	(0/0/0) 0
6	unassigned	wu	0	0	(0/0/0) 0
7	unassigned	wu	0	0	(0/0/0) 0

```
partition>
```

Partitioning a Hard Disk

The partition table name appears in parentheses in the first line of the table.

The columns of the table have the following meanings:

Part	The disk slice number
Tag	The predefined, optional tag
Flag	The predefined, optional flag
Cylinders	The starting and ending cylinder numbers for the slice
Size	The slice size in blocks (b), cylinders (c), Mbytes (MB), or Gbytes (GB)
Blocks	The number of cylinders/tracks/sectors, and the total number of sectors in the slice

5. Select Slice 0 (zero) by entering 0.

```
partition> 0
Part      Tag    Flag      Cylinders      Size          Blocks
 0 unassigned   wm        0            0      (0/0/0)
0
```

Enter partition id tag[unassigned] :

6. When prompted for the ID tag, type a question mark (?), and press Return to list the available choices. You can change a tag by entering a new tag name.

Enter partition id tag[unassigned] : ?

Expecting one of the following: (abbreviations ok) :

unassigned	boot	root	swap
usr	backup	stand	var
home	alternates	reserved	

Enter partition id tag[unassigned] :

7. Type the tag alternates, and press Return.

Enter partition id tag[unassigned] : **alternates**

Enter partition permission flags[wm] :

8. When prompted for the permission flags, type a question mark (?), and press Return to list the available choices. You can change a flag by entering the new flag name.

Enter partition permission flags[wm] : ?

Expecting one of the following: (abbreviations ok) :

wm	- read-write, mountable
wu	- read-write, unmountable
rm	- read-only, mountable
ru	- read-only, unmountable

Enter partition permission flags[wm] :

9. Press Return to accept the default flag.

Enter partition permission flags[wm] : <return>

Enter new starting cyl[0] :

10. Press Return to accept the starting cylinder of 0 (zero).

Enter new starting cyl[0] : <return>

Enter partition size[0b, 0c, 0e, 0.00mb, 0.00gb] :

11. Enter 980mb for the new partition size for Slice 0.

Enter partition size[0b, 0c, 0e, 0.00mb, 0.00gb] : **980mb**

partition>

12. Type print, and press Return. The Partition table appears.

partition> **print**

Current partition table (unnamed) :

Total disk cylinders available: 57459 + 2 (reserved cylinders)

Part	Tag	Flag	Cylinders	Size	Blocks
0	alternates	wm	0 - 491	980.16MB	(492/0/0) 2007360
1	unassigned	wm	0	0	(0/0/0) 0
2	backup	wm	0 - 57458	111.79GB	(57459/0/0) 234432720
3	unassigned	wm	0	0	(0/0/0) 0
4	unassigned	wm	0	0	(0/0/0) 0
5	unassigned	wm	0	0	(0/0/0) 0
6	unassigned	wu	0	0	(0/0/0) 0
7	unassigned	wu	0	0	(0/0/0) 0

partition>

The current partition table shows the change to Slice 0.

Now adjust the starting cylinder for Slice 1.

Partitioning a Hard Disk

13. Select slice number 1 by typing **1**.

```
partition> 1
Part      Tag    Flag     Cylinders      Size          Blocks
  1 unassigned   wm        0            0      (0/0/0)
0
```

Enter partition id tag[unassigned] :

14. Type the tag **swap**, and press Return.

Enter partition id tag[unassigned] : **swap**

Enter partition permission flags [wm] :

15. Type **wu** at the permission flags selection, and press Return.

Enter partition permission flags [wm] : **wu**

Enter new starting cyl[0] :

16. Enter the new starting cylinder for Slice 1. Specify the cylinder that follows the last cylinder used by slice 0.

Enter new starting cyl[0] : **492**

Enter partition size[0b, 0c, 492e, 0.00mb, 0.00gb] :

17. Enter the new partition size for Slice 1.

Enter partition size[0b, 0c, 492e, 0.00mb, 0.00gb] : **512mb**

partition>

18. Type **print**, and press Return.

partition> **print**

Current partition table (unnamed) :

Total disk cylinders available: 57459 + 2 (reserved cylinders)

Part	Tag	Flag	Cylinders	Size	Blocks
0	alternates	wm	0 - 491	980.16MB	(492/0/0) 2007360
1	swap	wu	492 - 749	513.98MB	(258/0/0) 1052640
2	backup	wm	0 - 57458	111.79GB	(57459/0/0) 234432720
3	unassigned	wm	0	0	(0/0/0) 0
4	unassigned	wm	0	0	(0/0/0) 0
5	unassigned	wm	0	0	(0/0/0) 0
6	unassigned	wu	0	0	(0/0/0) 0
7	unassigned	wu	0	0	(0/0/0) 0

partition>

The current partition table shows the change to Slice 1.

The new starting cylinder for Slice 1 is one greater than the ending cylinder for Slice 0.

Now adjust the starting cylinder for Slice 7.

19. Type **7** to select Slice 7.

```
partition> 7
Part      Tag     Flag      Cylinders          Size            Blocks
    7 unassigned   wu           0                 0      (0/0/0)
0
```

Enter partition id tag[unassigned] :

20. Type the tag **home**, and press Return.

Enter partition id tag[unassigned] : **home**

Enter partition permission flags[wu] :

21. Press Return to select the default flag.

Enter partition permission flags[wu] : <**return**>

Enter new starting cyl[0] :

22. Type the new starting cylinder for Slice 7.

Enter new starting cyl[0] : **750**

Enter partition size[0b, 0c, 750e, 0.00mb, 0.00gb] :

23. Specify the new partition size for Slice 7 by typing a dollar (\$) sign.

Enter partition size[0b, 0c, 750e, 0.00mb, 0.00gb] : **\$**

partition>

Note – Enter a dollar (\$) sign as a value for the last partition size to automatically assign the remaining space on the disk to this slice.



24. Type **print** to display the partition table.

partition> **print**

Current partition table (unnamed) :

Total disk cylinders available: 57459 + 2 (reserved cylinders)

Part	Tag	Flag	Cylinders	Size	Blocks
0	alternates	wm	0 - 491	980.16MB	(492/0/0) 2007360
1	swap	wu	492 - 749	513.98MB	(258/0/0) 1052640
2	backup	wm	0 - 57458	111.79GB	(57459/0/0) 234432720
3	unassigned	wm	0	0	(0/0/0) 0
4	unassigned	wm	0	0	(0/0/0) 0
5	unassigned	wm	0	0	(0/0/0) 0
6	unassigned	wu	0	0	(0/0/0) 0
7	home	wu	750 - 57458	110.33GB	(56709/0/0) 231372720

partition>

Add up the cylinders in the Blocks column for Slice 0, Slice 1, and Slice 7. The number should equal the total number of cylinders contained in Slice 2.

Partitioning a Hard Disk

25. After checking the partition table to ensure that there are no errors, label the disk by typing **label**.

```
partition> label
```

```
Ready to label disk, continue? y
```

```
partition>
```

26. Enter **quit** or **q** to exit the PARTITION menu and exit the **format** utility.

```
partition> q
```

FORMAT MENU:

disk	- select a disk
type	- select (define) a disk type
partition	- select (define) a partition table
current	- describe the current disk
format	- format and analyze the disk
repair	- repair a defective sector
show	- translate a disk address
label	- write label to the disk
analyze	- surface analysis
defect	- defect list management
backup	- search for backup labels
verify	- read and display labels
save	- save new disk/partition definitions
volname	- set 8-character volume name
!<cmd>	- execute <cmd>, then return
quit	

```
format> q
```

```
#
```

Using the format Utility to Create fdisk Partitions

On x86/x64 systems, you must establish at least one fdisk partition on each disk you intend to use. This procedure is not necessary or available on SPARC-based systems.

On x86/x64 systems, the format utility provides access to the fdisk menu. The fdisk menu allows you to create fdisk partitions, to specify a Solaris fdisk partition, and choose an fdisk partition to use as the active (boot) partition. Only x86/x64 systems require fdisk partitions.

The fdisk menu allows you to create up to four fdisk partitions, however, only one Solaris OS fdisk partition can exist on a disk.

To create fdisk partitions, run the format utility and select an appropriate disk (in this example c2d0):

```
# format
Searching for disks...done

AVAILABLE DISK SELECTIONS:
 0. c1d0 <DEFAULT cyl 39887 alt 2 hd 64 sec 63>
   /pci@0,0/pci-ide@7/ide@0/cmdk@0,0
 1. c2d0 <DEFAULT cyl 10008 alt 2 hd 255 sec 63>
   /pci@0,0/pci-ide@7/ide@1/cmdk@0,0
Specify disk (enter its number): 1
```

The format menu displays. To select another disk to view or modify, enter disk while in the FORMAT menu.

```
Controller working list found
[disk formatted, defect list found]
```

FORMAT MENU:

- | | |
|-----------|-------------------------------------|
| disk | - select a disk |
| type | - select (define) a disk type |
| partition | - select (define) a partition table |
| current | - describe the current disk |
| format | - format and analyze the disk |
| fdisk | - run the fdisk program |
| repair | - repair a defective sector |
| show | - translate a disk address |
| label | - write label to the disk |
| analyze | - surface analysis |
| defect | - defect list management |

Partitioning a Hard Disk

```
backup      - search for backup labels
verify      - read and display labels
save        - save new disk/partition definitions
volname     - set 8-character volume name
!<cmd>     - execute <cmd>, then return
quit
format>
```

Enter the fdisk command to display the fdisk menu. If the disk has no fdisk partition defined, the fdisk menu presents the option to create a single fdisk partition that uses the entire disk, and that is identified as the Solaris fdisk partition on this disk. Type n to edit the fdisk partition table.

Note – Accepting this option makes the single fdisk partition it creates the *active* fdisk partition. This probably is not desirable on most systems.



format> **fdisk**

No fdisk table exists. The default partition for the disk is:

a 100% "SOLARIS System" partition

Type "y" to accept the default partition, otherwise type "n" to edit the partition table.

n

The fdisk menu displays information about the disk size, cylinder size, and current fdisk partitions. In this example, no fdisk partition is yet defined on the disk:

```
Total disk size is 10011 cylinders
Cylinder size is 16065 (512 byte) blocks
```

Partition	Status	Type	Start	End	Length	%
=====	=====	=====	====	==	=====	==

WARNING: no partitions are defined!

SELECT ONE OF THE FOLLOWING:

1. Create a partition
2. Specify the active partition
3. Delete a partition

4. Change between Solaris and Solaris2 Partition IDs
5. Exit (update disk configuration and exit)
6. Cancel (exit without updating disk configuration)

Enter Selection:

To create a single fdisk partition that uses the whole disk, select option 1
 - Create a partition. Enter the number that selects the type of disk partition that you would like to create. Select option 1 to create a SOLARIS2 fdisk partition.

Enter Selection: **1**

Select the partition type to create:

1=SOLARIS2	2=UNIX	3=PCIXOS	4=Other
5=DOS12	6=DOS16	7=DOSEXT	8=DOSBIG
9=DOS16LBA	A=x86 Boot	B=Diagnostic	C=FAT32
D=FAT32LBA	E=DOSEXTLBA	F=EFI	0=Exit? 1



Note – Earlier Solaris OS releases may have listed the SOLARIS partition type. The SOLARIS2 partition type simply uses a different identifying value in the disk label than the SOLARIS type did. This change prevents potential conflicts with other partition types.

The fdisk menu then prompts you to enter the percentage of the disk you would like to use for this fdisk partition. Enter 100 to use the whole disk:

Specify the percentage of disk to use for this partition (or type "c" to specify the size in cylinders). **100**

The fdisk menu then asks if this should be the active fdisk partition. Only the fdisk partition used to boot the system should be marked as the active fdisk partition. Because this is going to be a non-bootable partition, enter n for no.

Should this become the active partition? If yes, it will be activated each time the computer is reset or turned on.

Please type "y" or "n". **n**

Partitioning a Hard Disk

The fdisk menu displays the list of fdisk partitions. The empty Status column indicates that the Solaris2 fdisk partition is not active.

Total disk size is 10011 cylinders
Cylinder size is 16065 (512 byte) blocks

Partition	Status	Type	Cylinders				%
			Start	End	Length	====	
1		Solaris2	1	10010	10010	100	

SELECT ONE OF THE FOLLOWING:

1. Create a partition
2. Specify the active partition
3. Delete a partition
4. Change between Solaris and Solaris2 Partition IDs
5. Exit (update disk configuration and exit)
6. Cancel (exit without updating disk configuration)

Enter Selection:

Select option 5 to save the fdisk partition configuration.

Enter Selection: **5**

format>

 **Note** – To install multiple operating systems on a disk, create multiple fdisk partitions using fdisk option 1.

The following is an example of a disk with three fdisk partitions supporting three different operating systems:

Total disk size is 10011 cylinders
Cylinder size is 16065 (512 byte) blocks

Partition	Status	Type	Cylinders				%
			Start	End	Length	====	
1		Win95 FAT32	1	3303	3303	33	
2		UNIX System	3304	6606	3303	33	
3		Solaris2	6607	9909	3303	33	

Defining Slices on an x86/x64 System

Caution – Do not change the size of disk slices that are currently in use.



The following steps demonstrate how to divide a disk into slices:

1. As the root user, type `format` at the prompt, and press Return.

```
# format
```

```
Searching for disks...done
```

AVAILABLE DISK SELECTIONS:

- 0. c1d0 <DEFAULT cyl 39887 alt 2 hd 64 sec 63>
 /pci@0,0/pci-ide@7/ide@0/cmdk@0,0
- 1. c2d0 <DEFAULT cyl 10008 alt 2 hd 255 sec 63>
 /pci@0,0/pci-ide@7/ide@1/cmdk@0,0

Specify disk (enter its number):

The `format` utility searches for all attached disks that are powered on. For each disk it finds, the `format` utility displays the logical device name, Sun marketing name, physical parameters, and physical device name.

2. Choose the second disk by selecting the number located to the left of that disk's logical device name. From the preceding display, the number chosen is 1. The `format` utility's main menu displays.

Specify disk (enter its number): **1**

selecting c2d0

Controller working list found

[disk formatted, defect list found]

FORMAT MENU:

disk	- select a disk
type	- select (define) a disk type
partition	- select (define) a partition table
current	- describe the current disk
format	- format and analyze the disk
fdisk	- run the fdisk program
repair	- repair a defective sector
show	- translate a disk address
label	- write label to the disk
analyze	- surface analysis
defect	- defect list management
backup	- search for backup labels
verify	- read and display labels

Partitioning a Hard Disk

```
save      - save new disk/partition definitions
volname   - set 8-character volume name
!<cmd>   - execute <cmd>, then return
quit
format>
```

The specific menu selections that you can use to view, change, or commit disk slices include the following:

partition	Displays the PARTITION menu
label	Writes the current partition definition to the VTOC of the Solaris fdisk partition
verify	Reads and displays the disk label
quit	Exits the format utility

3. Type **partition** at the format prompt. The PARTITION menu displays.

```
format> partition
```

PARTITION MENU:

0	- change `0' partition
1	- change `1' partition
2	- change `2' partition
3	- change `3' partition
4	- change `4' partition
5	- change `5' partition
6	- change `6' partition
7	- change `7' partition
select	- select a predefined table
modify	- modify a predefined partition table
name	- name the current table
print	- display the current table
label	- write partition map and label to the disk
!<cmd>	- execute <cmd>, then return
quit	

```
partition>
```

The Partition menu enables you to perform the following functions:

0–7	Specify the offset and size of up to eight slices
select	Choose a predefined partition table from the /etc/format.dat file
modify	Change the current partition table in memory
name	Provide a means to identify the partition table in the /etc/format.dat file
print	Display the current partition table in memory
label	Write the current partition table to the VTOC of the Solaris fdisk partition
!<cmd>	Escape from the utility and execute a command from the shell

- Type print at the partition prompt to display the disk label that was copied to random access memory (RAM) when the format utility was invoked. This example is from a SATA disk drive in an Ultra 20.

```
partition> print
Current partition table (original):
Total disk cylinders available: 10008 + 2 (reserved cylinders)
```

Part	Tag	Flag	Cylinders	Size	Blocks
0	unassigned	wm	0	0	(0/0/0) 0
1	unassigned	wm	0	0	(0/0/0) 0
2	backup	wu	0 – 10007	76.67GB	(10008/0/0) 160778520
3	unassigned	wm	0	0	(0/0/0) 0
4	unassigned	wm	0	0	(0/0/0) 0
5	unassigned	wm	0	0	(0/0/0) 0
6	unassigned	wm	0	0	(0/0/0) 0
7	unassigned	wm	0	0	(0/0/0) 0
8	boot	wu	0 – 0	7.84MB	(1/0/0) 16065
9	alternates	wm	1 – 2	15.69MB	(2/0/0) 32130

```
partition>
```

The partition table name appears in parentheses in the first line of the table. Slice 8 is the boot slice. The boot slice uses the first cylinder (cylinder 0) of the Solaris OS fdisk partition. Slice 9 uses the next two cylinders, and the alternates tag. Slice 9 contains blocks used to store bad block information. This slice is not found or used for this purpose on SCSI or Fibre disks.

Partitioning a Hard Disk

The columns of the table have the following meanings:

Part	The slice number within the Solaris fdisk partition
Tag	The predefined, optional tag
Flag	The predefined, optional flag
Cylinders	The starting and ending cylinder number for the slice
Size	The slice size in blocks (b), cylinders (c), Mbytes (MB), or Gbytes (GB)
Blocks	The number of cylinders/tracks/sectors, and the total number of sectors in the slice

5. Select Slice 0 (zero) by entering 0.

```
partition> 0
Part Tag Flag Cylinders Size Blocks
 0 unassigned  wm      0        0   (0/0/0)    0
```

6. When prompted for the ID tag, type a question mark (?), and press Return to list the available choices. You can change a tag by entering a new tag name.

```
Enter partition id tag[unassigned] : ?
Expecting one of the following: (abbreviations ok) :
  unassigned  boot      root      swap
  usr         backup     stand     var
  home        alternates reserved
```

```
Enter partition id tag[unassigned] :
```

7. Type the tag reserved, and press Return.

```
Enter partition id tag[unassigned] : reserved
```

```
Enter partition permission flags [wm] :
```

8. When prompted for the permission flags, type a question mark (?), and press Return to list the available choices. You can change a flag by entering the new flag name.

```
Enter partition permission flags [wm] : ?
```

```
Expecting one of the following: (abbreviations ok) :
```

wm	- read-write, mountable
wu	- read-write, unmountable
rm	- read-only, mountable
ru	- read-only, unmountable

```
Enter partition permission flags [wm] :
```

9. Press Return to accept the default flag.

Enter partition permission flags[wm] : <**return**>

10. Enter 3 to accept the starting cylinder of 3. Slices 8 and 9 are already using cylinders 0-2.

Enter new starting cyl[3] : <**return**>

Enter partition size[0b, 0c, 3e, 0.00mb, 0.00gb] :

11. Enter 400mb for the new partition size for Slice 0.

Enter partition size[0b, 0c, 3e, 0.00mb, 0.00gb] : **400mb**
partition>

12. Type print, and press Return. The Partition table appears.

partition> **print**

Current partition table (unnamed) :

Total disk cylinders available: 10008 + 2 (reserved cylinders)

Part	Tag	Flag	Cylinders	Size	Blocks
0	reserved	wm	3 - 53	400.06MB (51/0/0)	819315
1	unassigned	wm	0	0 (0/0/0)	0
2	backup	wu	0 - 10007	76.67GB (10008/0/0)	160778520
3	unassigned	wm	0	0 (0/0/0)	0
4	unassigned	wm	0	0 (0/0/0)	0
5	unassigned	wm	0	0 (0/0/0)	0
6	unassigned	wm	0	0 (0/0/0)	0
7	unassigned	wm	0	0 (0/0/0)	0
8	boot	wu	0 - 0	7.84MB (1/0/0)	16065
9	alternates	wm	1 - 2	15.69MB (2/0/0)	32130

partition>

The current partition table shows the change to Slice 0. Adjust the starting cylinder for Slice 1.

13. Select slice number 1 by typing **1**.

partition> **1**

Part	Tag	Flag	Cylinders	Size	Blocks
1	unassigned	wm	0	0	(0/0/0)
0					

Enter partition id tag[unassigned] :

14. Type the tag swap, and press Return.

Enter partition id tag[unassigned] : **swap**

Enter partition permission flags[wm] :

15. Type wu at the permission flags selection, and press Return.

Enter partition permission flags[wm] : **wu**

Enter new starting cyl[3] :

Partitioning a Hard Disk

16. Enter the new starting cylinder for Slice 1. Specify the cylinder that follows the ending cylinder for slice 0.

Enter new starting cyl[3]: **54**

Enter partition size[0b, 0c, 54e, 0.00mb, 0.00gb]:

17. Enter the new partition size for Slice 1.

Enter partition size[0b, 0c, 54e, 0.00mb, 0.00gb]: **512mb**

partition>

18. Type print, and press Return.

partition> **print**

Current partition table (unnamed):

Total disk cylinders available: 10008 + 2 (reserved cylinders)

Part	Tag	Flag	Cylinders	Size	Blocks
0	reserved	wm	3 - 53	400.06MB	(51/0/0) 819315
1	swap	wu	54 - 119	517.72MB	(66/0/0) 1060290
2	backup	wu	0 - 10007	76.67GB	(10008/0/0) 160778520
3	unassigned	wm	0	0	(0/0/0) 0
4	unassigned	wm	0	0	(0/0/0) 0
5	unassigned	wm	0	0	(0/0/0) 0
6	unassigned	wm	0	0	(0/0/0) 0
7	unassigned	wm	0	0	(0/0/0) 0
8	boot	wu	0 - 0	7.84MB	(1/0/0) 16065
9	alternates	wm	1 - 2	15.69MB	(2/0/0) 32130

partition>

The current partition table shows the change to Slice 1. The new starting cylinder for Slice 1 is one greater than the ending cylinder for Slice 0. Adjust the starting cylinder for Slice 7.

19. Type 7 to select Slice 7.

partition> **7**

Part	Tag	Flag	Cylinders	Size	Blocks
7	unassigned	wm	0	0	(0/0/0)
0					

Enter partition id tag[unassigned]:

20. Type the tag **home**, and press Return.

Enter partition id tag[unassigned]: **home**

Enter partition permission flags [wm]:

21. Press Return to select the default flag.

Enter partition permission flags [wm]: <**return**>

Enter new starting cyl[3]:

22. Type the new starting cylinder for Slice 7. Specify the cylinder that follows the last cylinder for slice 1.

Enter new starting cyl[3]: **120**

Enter partition size[0b, 0c, 120e, 0.00mb, 0.00gb]:

23. Specify the new partition size for Slice 7 by typing a dollar (\$) sign.

Enter partition size[0b, 0c, 120e, 0.00mb, 0.00gb]: **\$**

partition>



Note – Enter a dollar (\$) sign as a value for the last partition size to automatically assign the remaining space on the disk to this slice.

24. Type **print** to display the partition table.

partition> **print**

Current partition table (unnamed):

Total disk cylinders available: 10008 + 2 (reserved cylinders)

Part	Tag	Flag	Cylinders	Size	Blocks
0	reserved	wm	3 - 53	400.06MB	(51/0/0) 819315
1	swap	wu	54 - 119	517.72MB	(66/0/0) 1060290
2	backup	wu	0 - 10007	76.67GB	(10008/0/0) 160778520
3	unassigned	wm	0	0	(0/0/0) 0
4	unassigned	wm	0	0	(0/0/0) 0
5	unassigned	wm	0	0	(0/0/0) 0
6	unassigned	wm	0	0	(0/0/0) 0
7	home	wm	120 - 10007	75.75GB	(9888/0/0) 158850720
8	boot	wu	0 - 0	7.84MB	(1/0/0) 16065
9	alternates	wm	1 - 2	15.69MB	(2/0/0) 32130

partition>

Add the cylinders in the Blocks column for Slice 0, Slice 1, Slice 7, Slice 8, and Slice 9. The number should equal the total number of cylinders contained in Slice 2.

25. After checking the partition table, to ensure that there are no errors, label the Solaris fdisk partition by typing **label**.

partition> **label**

Ready to label disk, continue? **y**

partition>

Partitioning a Hard Disk

26. Enter quit or q to exit the PARTITION menu and exit the format utility.

partition> **q**

FORMAT MENU:

disk	- select a disk
type	- select (define) a disk type
partition	- select (define) a partition table
current	- describe the current disk
format	- format and analyze the disk
fdisk	- run the fdisk program
repair	- repair a defective sector
show	- translate a disk address
label	- write label to the disk
analyze	- surface analysis
defect	- defect list management
backup	- search for backup labels
verify	- read and display labels
save	- save new disk/partition definitions
volname	- set 8-character volume name
!<cmd>	- execute <cmd>, then return
quit	

format> **q**

 **Note** – To select another disk to view or modify, enter disk in the FORMAT menu.

EFI Label and VTOC Label Comparison

The EFI disk label differs from the VTOC disk label as follows:

- The EFI label provides support for disks greater than 1 terabyte.
- The EFI label provides usable slices 0-6, where slice 2 is just another slice.
- Partitions (or slices) cannot overlap with the primary or backup label, nor with any other partition. The EFI label size is usually 34 sectors, so partitions start at sector 34. This feature means that no partition can start at sector zero.
- No cylinder, head, or sector information is stored in the EFI label. Sizes are reported in blocks.
- Information that was stored in alternate cylinder areas (the last two cylinders of the disk) is now stored in slice 8.
- If you use the format utility to change partition sizes, the unassigned partition tag is assigned to partitions with sizes equal to zero. By default, the format utility assigns the usr partition tag to any partition with a size greater than zero. Use the partition change menu to reassign partition tags after the partitions are changed. You cannot change a partition with a non-zero size to the unassigned partition tag.

Note – For SMI/VTOC labeled disks, you have 8 slices (0 to 7). For EFI labeled disks, you have 10 slices (0 to 9). How you use slices depends on your requirements. You can allocate some space to all the slices, leave some slices blank, or allocate one slice as the entire disk space.

The default behavior of both EFI and SMI/VTOC labeled disks are that slice 2 always occupies the entire disk space. For EFI labeled disks, you can resize slice 2. You cannot do this with SMI/VTOC labeled disks.

EFI Disk Label Restrictions

If you have disks in your environment that contain more than 1 terabyte, you must consider the following restrictions:

- Layered software products intended for systems with VTOC-labelled disks might be incapable of accessing a disk with an EFI disk label.
- A disk with an EFI label is not recognized on systems running previous Solaris releases.
- You cannot boot from a disk with an EFI disk label.
- You cannot use the `fdisk` command on a disk with an EFI label that is greater than 1 terabyte.
- The EFI specification prohibits overlapping slices. The entire disk is represented by `cxtydz`.
- The EFI disk label provides information about disk or partition sizes in sectors and blocks, but not in cylinders and heads.

The following format options are either not supported or are not applicable on disks with EFI labels:

- The `save` option is not supported because disks with EFI labels do not need an entry in the `format.dat` file.
- The `backup` option is not applicable because the disk driver finds the primary label and writes it back to the disk.

Viewing the VTOC

You can use two methods to view the SPARC or x86/x64 VTOC on disk:

- Use the `verify` command in the `format` utility
- Run the `prtvtoc` command from the command line

Note – The VTOC on SPARC systems is found in the first sector on the disk, and the VTOC on x86/x64 systems is found in the second sector of the Solaris `fdisk` partition on the disk.



Reading a Disk VTOC Using the verify Command

The **verify** command enables you to view a SPARC or x86/x64 VTOC from within the **format** utility. To read a disk's VTOC, perform the following steps:

1. Run the **format** utility and select the disk you want to use.
2. At the **format>** prompt, enter the **verify** command, and press Return.

This example is from a Sun Blade 1500:

```
format> verify
```

Primary label contents:

Volume name	= < >				
ascii name	= <HDS722512VLAT80 cyl 57459 alt 2 hd 16 sec 255>				
pcyl	= 57461				
ncyl	= 57459				
acyl	= 2				
nhead	= 16				
nsect	= 255				
Part	Tag	Flag	Cylinders	Size	Blocks
0 alternates	wm		0 – 491	980.16MB	(492/0/0)
2007360					
1 swap	wu		492 – 749	513.98MB	(258/0/0)
1052640					
2 backup	wm		0 – 57458	111.79GB	(57459/0/0)
234432720					
3 unassigned	wm		0	0	(0/0/0)
0					
4 unassigned	wm		0	0	(0/0/0)
0					
5 unassigned	wm		0	0	(0/0/0)
0					
6 unassigned	wm		0	0	(0/0/0)
0					
7 home	wm		750 – 57458	110.33GB	(56709/0/0)
231372720					

```
format>
```

Partitioning a Hard Disk

This example is from an Ultra 20:

format> **verify**

Primary label contents:

```
Volume name = <           >
ascii name  = <DEFAULT cyl 10008 alt 2 hd 255 sec 63>
pcyl        = 10010
ncyl        = 10008
acyl        =    2
bcyl        =    0
nhead       = 255
nsect       =   63
Part         Tag      Flag    Cylinders          Size        Blocks
  0  reserved   wu      3 - 53    400.06MB  (51/0/0)
819315
  1  swap       wu      54 - 119   517.72MB  (66/0/0)
1060290
  2  backup     wu      0 - 10007  76.67GB   (10008/0/0)
160778520
  3 unassigned  wu      0          0          (0/0/0)
0
  4 unassigned  wu      0          0          (0/0/0)
0
  5 unassigned  wu      0          0          (0/0/0)
0
  6 unassigned  wu      0          0          (0/0/0)
0
  7  home       wu      120 - 10007 75.75GB   (9888/0/0)
158850720
  8  boot       wu      0 - 0      7.84MB    (1/0/0)
16065
  9 alternates wu      1 - 2      15.69MB   (2/0/0)
32130
```

format>

3. Type quit or q and press Return to exit the format utility.

Reading a Disk VTOC Using the `prtvtoc` Command

The `prtvtoc` command enables you to view a SPARC or x86/x64 VTOC from the command line.

Use the `prtvtoc` command as described in these examples to view the VTOC. This example is from a Sun Blade 1500:

```
# prtvtoc /dev/rdsk/c0t1d0s2
* /dev/rdsk/c0t1d0s2 partition map
*
* Dimensions:
*      512 bytes/sector
*      255 sectors/track
*      16 tracks/cylinder
*     4080 sectors/cylinder
*    57461 cylinders
*   57459 accessible cylinders
*
* Flags:
*   1: unmountable
*  10: read-only
*
*          First       Sector      Last
* Partition Tag  Flags    Sector      Count    Sector Mount Directory
*            0    9  00        0  2007360  2007359
*            1    3  01  2007360  1052640  3059999
*            2    5  00        0 234432720 234432719
*            7    8  00  3060000 231372720 234432719
```

Partitioning a Hard Disk

This example is from an Ultra 20 workstation:

```
# prtvtoc /dev/rdsck/c2d0s2
* /dev/rdsck/c2d0s2 partition map
*
* Dimensions:
*      512 bytes/sector
*      63 sectors/track
*     255 tracks/cylinder
*   16065 sectors/cylinder
*   10010 cylinders
*  10008 accessible cylinders
*
* Flags:
*   1: unmountable
*  10: read-only
*
*
* Partition Tag Flags      First Sector    Sector Count  Last Sector Mount Directory
* 0       11  00          48195    819315      867509
* 1       3   01          867510   1060290     1927799
* 2       5   01           0 160778520 160778519
* 7       8   00          1927800 158850720 160778519
* 8       1   01           0 16065      16064
* 9       9   00          16065    32130       48194
#
#
```

The VTOC information includes the following fields:

Dimensions	Describes the logical dimensions of the disk.
Partition	A number that identifies the slice.
Tag	A value that indicates how the slice is being used. The Solaris OS no longer uses this field.
Flags	The Solaris OS no longer uses this field.
First Sector	Defines the first sector of the slice.
Sector Count	Defines the total number of sectors in the slice.
Last Sector	Defines the last sector number in the slice.
Mount Directory	If the field is empty, the slice is currently not mounted.

Replacing a SPARC or x86/x64 VTOC on Disk

Use the `prtvtoc` command to save VTOC information to a file. This enables you to use the `fmthard` command to replace the SPARC or x86/x64 VTOC to recover from events such as these:

- The VTOC on the disk was destroyed.
- You accidentally changed the slice information on the disk and did not save a backup label in the `/etc/format.dat` file.



Note – The `prtvtoc` and `fmthard` commands read and write the SPARC or x86/x64 VTOC on disks. They do not affect the `fdisk` partition table on x86/x64 systems. Use the `fdisk` menu in the `format` utility to establish `fdisk` partitions on x86/x64 systems.

To save VTOC information to a file, use the `prtvtoc` command as in the following examples. Use the device name for the correct disk, and save the `prtvtoc` output to a file named as you wish.

```
# prtvtoc /dev/rdsd/c0t1d0s2 > /var/tmp/c0t1d0s2.vtoc
# prtvtoc /dev/rdsd/c2d0s2 > /var/tmp/c2d0s2.vtoc
```

The `fmthard` Command

To replace a SPARC or x86/x64 VTOC on disk, you can use the saved output from the `prtvtoc` command as input to the `fmthard` command. In this syntax example, you would replace `datafile` with the name of the file that contains the `prtvtoc` command output. Use slice 2 in the name of the target disk or Solaris `fdisk` partition.

```
fmthard -s datafile /dev/rdsd/c#t#d#s2
```



Caution – The `fmthard` command cannot write a VTOC to an unlabeled disk. Use the `format` utility for this purpose.

If the need to replace a SPARC or x86/x64 VTOC arises and the VTOC was previously saved to a file, the following options are available:

- Run `format`, select the disk, and label it with the default partition table, or define slices and label the disk as required.
- Use the `fmthard` command to write saved VTOC information back to the disk.

This example is from a Sun Blade 1500:

```
# fmthard -s /var/tmp/c0t1d0s2.vtoc /dev/rdsk/c0t1d0s2
fmthard: New volume table of contents now in place.
#
```

This example is from an Ultra 20:

```
# fmthard -s /var/tmp/c2d0s2.vtoc /dev/rdsk/c2d0s2
fmthard: New volume table of contents now in place.
#
```

- Use the fmthard command to initialize the VTOC of a disk.

```
# fmthard -s /dev/null /dev/rdsk/c1t3d0s2
fmthard: New volume table of contents now in place.
#
```

x86/x64 VTOC Information and Solaris fdisk Partitions

On x86/x64 systems, the Solaris fdisk partition size must match the size described in the saved VTOC information. If the Solaris fdisk partition size differs from the size described in the prtvtoc command output, fmthard reports errors similar to these:

```
# fmthard -s /var/tmp/c2d0s2.vtoc /dev/rdsk/c2d0s2
fmthard: Partition 2 specifies the full disk and is not equal
full size of disk. The full disk capacity is 80373195 sectors.
fmthard: Partition 2 specified as 160778520 sectors starting at 0
      does not fit. The full disk contains 80373195 sectors.
fmthard: Partition 7 specified as 158850720 sectors starting at 1927800
      does not fit. The full disk contains 80373195 sectors.
fmthard: New volume table of contents now in place.
#
```

These messages serve only as a warning. The erroneous partition information is written to the x86/x64 VTOC.

To correct this situation, you can:

- Use the partition menu in the format utility to re-define the slices in the Solaris fdisk partition.
- Use the fdisk menu in the format utility to re-define the fdisk partitions to match the size defined in the saved prtvtoc output, and use fmthard to write the x86/x64 VTOC information to disk.
- Use the fdisk command to replace the fdisk partition table with saved fdisk partition information, and use fmthard to write the x86/x64 VTOC information to disk.

Replacing fdisk Partition Tables on x86/x64 Systems

The `fdisk` command enables you to view and modify `fdisk` partition tables on disk. The `fdisk` command allows you to save a disk's `fdisk` partition table to a file. You can then use the `fdisk` command to replace the `fdisk` partition table on disk with the saved information. This procedure only pertains to x86/x64 systems.

To view an `fdisk` partition table, use the `fdisk` command as in the following example from an Ultra 20 workstation:

```
# fdisk -W - /dev/rdsk/c2d0p0

* /dev/rdsk/c2d0p0 default fdisk table
* Dimensions:
*   512 bytes/sector
*   63 sectors/track
*   255 tracks/cylinder
* 10011 cylinders
*
* systid:
*   1: DOSOS12
*   2: PCIXOS
*   4: DOSOS16
*   5: EXTDOS
*   6: DOSBIG
*   7: FDISK_IFS
*   8: FDISK_AIXBOOT
*   9: FDISK_AIXDATA
*  10: FDISK_0S2BOOT
*  11: FDISK_WINDOWS
*  12: FDISK_EXT_WIN
*  14: FDISK_FAT95
*  15: FDISK_EXTLBA
*  18: DIAGPART
*  65: FDISK_LINUX
*  82: FDISK_CPM
*  86: DOSDATA
*  98: OTHEROS
*  99: UNIXOS
* 101: FDISK_NOVELL3
* 119: FDISK_QNX4
* 120: FDISK_QNX42
* 121: FDISK_QNX43
* 130: SUNIXOS
* 131: FDISK_LINUXNAT
```

Partitioning a Hard Disk

```
* 134: FDISK_NTFSVOL1
* 135: FDISK_NTFSVOL2
* 165: FDISK_BSD
* 167: FDISK_NEXTSTEP
* 183: FDISK_BSDIFS
* 184: FDISK_BSDISWAP
* 190: X86BOOT
* 191: SUNIXOS2
* 238: EFI_PMBR
* 239: EFI_FS
*
*   Id      Act    Bhead   Bsect   Bcyl      Ehead   Esect   Ecyl      Rsect      Numsect
  191      0       0        1        1      254       63     1023     16065
160810650
#
```

The fdisk command lists the disk geometry, all of the possible partition types and their corresponding ID values, and the currently defined fdisk partitions. In this example, only one fdisk partition is defined. Its type 191 indicates that it is a Solaris OS (SUNIXOS2 = SOLARIS2) partition.

To save fdisk partition table information to a file, use the fdisk command as in the following example. Replace the argument to the -W option with the name of the file you wish to use, and specify fdisk partition 0 of the correct disk.

```
# fdisk -W /var/tmp/c2d0p0.fdisk /dev/rdsk/c2d0p0
#
```

To replace the fdisk partition table on disk with the saved information, use the fdisk command as in the following example. Replace the argument to the -F option with the name of the file that contains the saved partition information, and specify fdisk partition 0 of the correct disk.

```
# fdisk -F /var/tmp/c2d0p0.fdisk /dev/rdsk/c2d0p0
#
```

The -F option will zero out the VTOC on the Solaris partition if the fdisk partition table changes.

Extended Volume Table of Contents

Prior to the Solaris 10 10/09 release, the Sun disk label limited the size of the bootable disk to less than 1 Tbyte. Now, installing and booting from disks in the 1 to 2 Tbytes range is supported with the 64-bit Solaris kernel.

The Extended VTOC feature is available after a standard installation of Solaris 10 10/09 release, and can be used immediately as it is on by default.

Note – There is currently no support for disks over 1 Tbyte on 32-bit kernels.



Previously, an Extensible Firmware Interface (EFI) label for a disk that is larger than 1 Tbyte was always used. Now the Volume Table of Contents (VTOC) label is extended to any size disk, but the addressable space by the VTOC is limited to 2 Tbytes.

Features of the new Extended VTOC

- The `fdisk` utility can now be used on a disk that is greater than 1 Tbyte on x86 systems. Support is added for up to 2 Tbyte partitions in the Master Boot Record (MBR) for non-EFI partition types. Other non-EFI partitions may be subject to a limit depending on partition type.
- The `format -e` utility can be used to label a disk of any size with a VTOC label, but the addressable space is limited to 2 Tbytes.
- In addition to using a VTOC label, the `format -e` command can be used to label a disk that is 2 Tbytes in size or less with an EFI label.
- When the `fdisk` utility is run on a disk that is greater than 2 Tbytes in size, a warning message is displayed indicating that a non-EFI partition greater than 2 Tbytes cannot be created.

Note – The Solaris Volume Manager software in the Solaris 10/09 release is modified to create meta-devices that support physical disks with VTOC labels up to 2 Tbytes in size.





Extended VTOC Requirements and Dependencies

- The Extended VTOC is supported on both x86 (64-bit) and SPARC platforms, and utilizes the Newboot feature.
- The Extended VTOC works with target and HBA drivers that support disks larger than 1 Tbyte.
- The DAD and IDE interfaces on SPARC, and the PCMCIA device on x86 systems do not support disks that are larger than 1 Tbyte.

Note – Disks with existing EFI labels continue to work as in previous Solaris releases.

Exercise: Managing Local Disk Devices

In this exercise, you complete the following tasks:

- Identify logical, physical, and instance names for disk devices
- Add a new disk or tape drive to a system
- Create and manage disk slices on Solaris systems on the SPARC platform
- Create and manage fdisk partitions and disk slices on x86/x64 Solaris systems
- Use Solaris Management Console to manage local disks on SPARC-based systems

Preparation

This exercise requires one unused spare disk.

This exercise applies to both SPARC and x86/x64 systems. Keep in mind that these different systems typically use different disk device names. For example:

- The boot disk of a Sun Blade 1500 (SPARC) is typically c0t0d0.
- The boot disk of an Ultra 20 (x86/x64) is typically c1d0.

Throughout this exercise, be certain to use the correct device names for your particular system.

Task 2 in this exercise is optional. Task 2 requires a system that is configured with an external disk or tape drive. In order to perform Task 2, one of these two conditions must be met:

- The external device was turned off during system installation, so no device files for the device will exist in the /dev structure.
- Device files for the external device are removed before starting Task 2. To remove device files:
 - Shut down the system to run state 0 or 5, and power off the external device.
 - Boot the system to the default run state.
 - Login as root and run devfsadm -C to clear any device files related to the external device.

Exercise Summary

This exercise is composed of five tasks that illustrate how to manage local disk devices. The tasks in this exercise describe the following concepts:

- Identifying Device Files - demonstrates the relationships among logical and physical device files, and device instance names.
- Adding a New Disk or Tape Device - describes the procedure required to build device files for new devices added to your system. This task is optional, depending on the peripherals available on your system.
- Working With Disks and Partitions on SPARC Systems - describes using the format utility, and the prtvtoc and fmthard commands, to create and manage disk slices on disks attached to Solaris OS systems on the SPARC platform.
- Working With Disks and Partitions on x86/x64 Systems - describes using the format utility, and the fdisk, prtvtoc and fmthard commands, to create and manage fdisk partitions and disk slices on disks attached to x86/x64 Solaris systems.

Task 1 – Identifying Device Files

Complete the following steps:

1. Log in as the root user and complete the following steps:
 - a. Open a terminal window.
 - b. Expand the window so that it occupies the entire screen area.
 - c. Change to the /dev/dsk directory.
 2. List the files in this directory and complete the following steps:
 - a. Identify the files related to the boot disk of your system.
Systems typically use c0t0d0 or c1d0, but this may vary.
 - b. Locate the file that represents slice 0 on this disk and display a long listing of it.
- Which type of file did you just locate? The file type indicator is the first character on the left side of the long listing.
-
- c. Record the full path name to which this file points.
-

3. Issue a long listing command of the path name you recorded.

Which type of file is this?

The `ls -lL` command displays the same information, but shows only the link file name, not the real device file name.

4. Change to the `/dev/rdsk` directory and complete the following steps:

- a. Display a long listing of the same file name you selected in Step 2.

Which type of file is this?

- b. Record the full path name to which this file points.
-

5. Issue a long listing command of the path name you recorded.
-

Which type of file is this?

The `ls -lL` command displays the same information, but shows only the link file name, not the real device file name.

6. Change to the `/etc` directory and display the contents of the `path_to_inst` file.
 7. Use the information from the previous steps to locate and record the entry for your boot disk.
-

What is the instance name for the device listed in this step?

Task 2 – Adding a New Disk or Tape Device

Complete the following steps:



Note – This task is optional. Complete the steps in this task if an external disk or tape device is available on your system.

1. In the /dev/dsk or /dev/rmt directories, confirm that no files exist for your external disk or tape device; for example, /dev/dsk/c1t0d0s0 or /dev/rmt/0.
If files for the external device do exist, ask your instructor for guidance. Do not simply remove device files that exist. Refer to the Preparation section for advice about removing device files.
2. If you are using a SPARC system, shut it down to run state 0.
3. If you are using an x86/x64 system, shut it down to run state 5. This shuts down the system and turns off the power.
4. Power on the external disk or tape drive attached to your system.
5. If you are using an x86/x64 system, power on your system and allow it to boot to its default run state.
6. If you are using a SPARC system, boot the system to its default run state.
7. Log in as the root user and complete the following steps:
 - a. Open a terminal window.
 - b. Run the devfsadm command with the -v option to create new links and device files for the new disk or tape drive.
 - c. Observe the messages that the devfsadm command displays.
8. Confirm that new links and device files exist in the /dev/dsk and /dev/rdsk directories for disks or /dev/rmt for tape drives.
Use the devfsadm -v output to determine what files to look for.

Task 3 - Working With Disks and Partitions on SPARC

Complete the following steps:

1. Log in as the root user and complete the following steps:
 - a. Open a terminal window.
 - b. Run the format utility.
2. Complete the following steps:
 - a. Record the list of disks presented by the format utility. For example, c0t0d0 and c0t1d0.
 - b. Press **Control-D** to exit the format utility.
3. Use the prtvtoc command to list the VTOC for each of the disks you found in the previous step and complete the following steps:
 - a. Examine the Mount Directory field in the information that the prtvtoc command displays.
 - b. Record the name of a disk that has no mount directory listed.
For your classroom environment, this is an unused disk.

Note – Systems in production can use unmounted slices to store production data. Do not assume that unmounted slices are unused on systems in production. Systems will display various lists of partitions, so the list you see may differ from the list the examples.



Unused disk:

-
4. Run the format utility again and select the unused disk from the list of disks presented.
 5. Display the partition menu and complete the following steps:
 - a. Print the current partition table.
 - b. Record the number of gigabytes (Gbytes) assigned to Slice 2.

Gbytes:



Note – The list of defined partitions will vary, depending on how the spare disk was last used.

6. Divide the number of Gbytes by 4 and complete the following steps:

- a. Use the result as the number of Gbytes to assign as disk space to four slices.
- b. Round *down* to the next whole Gbyte if the result includes a fraction.

For example, if the slice 2 is 111.79 Gbytes, $111.79 / 4 = 27.94$ Gbytes. In this case use 27 Gbytes as the size for the four slices.

Gbytes/4:

7. Display the Partition menu again and complete the following steps:

- a. Select Slice 0.
- b. Accept the defaults for tags and flags.
- c. Start this first slice on Cylinder 0.
- d. Enter the resulting number of Gbytes from the previous step for the slice size.
- e. Print the partition table again to verify the change.

8. Set the sizes of Slices 1, 3, and 4 so that they are the same as Slice 0.

You should begin each successive slice on the cylinder that follows the ending cylinder of the previous slice.

9. Set Slices 5, 6, and 7 to start at Cylinder 0, and then assign them 0 Mbytes.

10. Complete the following steps:

- a. Print the partition table.

Except for slice 2, is there any overlap of ending and beginning cylinders for any of the slices listed?

- b. Proceed to the following steps to introduce this problem.

11. Complete the following steps:

- a. Add 1 to the number Gbytes/4 value listed in Step 6.

(Gbytes/4) + 1:

- b. Change Slice 0 so that it uses the new size.

The Cylinders column in the partition table should now indicate that Slice 0 ends after Slice 1 begins.

12. Complete the following steps:

- Use the modify command from the partition menu to attempt to fix this problem.
- Select Item 0 to modify the current partition table.

Which warnings display?

13. Modify the partition table and select Item 1 to use the All Free Hog option.

The partition table displays.

14. Observe the Cylinders and Size columns, and then verify that they are all zero except for slice 2.

15. Respond to the prompts to continue the process and complete the following steps:

- Select Slice 4 as the All Free Hog partition.
- Use the size listed in Step 6 for Slices 0, 1, and 3.
- Enter Return to set the other slices to size 0.

At the end of this process, you should have three slices of equal size, where Slice 4 takes up any extra room if it exists.

16. Name the partition table "MYDISKpartition", and then label the disk.

17. Save your new partition table to the /etc/format.dat file and complete the following steps:

- Carefully read the message that the format utility displays, and then enter the correct file name.
- Quit the format utility when you have finished.
- Use the tail command to view the contents of the /etc/format.dat file.
- Note that your new partition information is appended to the file.

18. Verify your new partition table with the prtvtoc command.

19. Create a directory called /vtoc.

Exercise: Managing Local Disk Devices

20. Use the `prtvtoc` command to print the partition table that you just created and complete the following steps:
 - a. Save its output to a file in the `/vtoc` directory.
 - b. Name the file so that it corresponds with the disk you are examining.
 - c. Use the `cat` command to verify that valid information exists in the file that you create.
21. Use the following `dd` command to destroy the disk label and complete the following steps:
 - a. Be certain to specify the correct disk device name for the `of=` argument.
 - b. Enter all other arguments exactly as listed.

```
# dd if=/dev/zero of=/dev/rdsk/c0t1d0s2 bs=512 count=1
1+0 records in
1+0 records out
#
```

22. Attempt to read the label from the same disk by using the `prtvtoc` command.

What happens?

23. Complete one of the following steps:
 - If the `prtvtoc` command reported an "Unable to read Disk geometry" message, use the `format` utility to place a default label on the disk for which you destroyed the label earlier.
 - If the `prtvtoc` command reports that only slice 2 exists, or slice 0 and 2 exist on the disk, proceed to step 24.
24. Use the `fmthard` command to write to the disk the label information you saved earlier.
25. Attempt to read the label from the same disk.

Was this successful?

Task 4 - Working With Disks and Partitions on x86/x64

Complete the following steps:

1. Log in as the root user and complete the following steps:
 - a. Open a terminal window.
 - b. Run the `format` utility.
2. Complete the following steps:
 - a. Record the list of disks presented by the `format` utility. For example, `c1d0` and `c2d0`.
 - b. Press **Control-D** to exit the `format` utility.
3. Use the `prtvtoc` command to list the VTOC for each of the disks you found in the previous step, and complete the following steps:
 - a. If the `prtvtoc` command displays an error similar to this example, you have identified the unused disk:

```
# prtvtoc /dev/rdsk/c2d0s2
prtvtoc: /dev/rdsk/c2d0s2: Unable to read Disk geometry errno = 0x5
#
```

If this error displays, record the name of this disk and skip steps b and c.

- b. If the `prtvtoc` command displays a partition table for both disks, examine the `Mount Directory` field in the command output.
- c. Record the name of a disk that has no mount directory listed. For your classroom environment, this is an unused disk.

Note – Systems in production can use unmounted slices to store production data. Do not assume that unmounted slices are unused on systems in production.



Unused disk:

-
4. Run the `format` utility again, and select the unused disk from the list of disks presented.
 - a. If the disk has no `fdisk` partition table defined, `format` asks you to select a drive type. Select the `DEFAULT` entry. The `FORMAT` menu displays.

- Unauthorized reproduction or distribution prohibited. Copyright© 2011, Oracle and/or its affiliates.
- b. If the disk has a defined fdisk partition table, the FORMAT menu displays:
 - 5. Display the fdisk menu and complete the following steps to create one Solaris fdisk partition that uses the whole disk:
 - a. If the disk has no defined fdisk partition table, the fdisk menu asks if you want to accept the default Solaris partition. Enter n to decline, and then edit the partition table.
 - b. If the disk has a defined fdisk partition table, the table displays when you enter the fdisk menu.
 - c. Delete all fdisk partitions that exist.
 - 6. Use the fdisk menu to create one Solaris fdisk partition that uses the whole disk.
 - a. Select the SOLARIS2 partition type
 - b. Enter 100 to specify using 100% of the disk
 - c. Enter n to prevent this partition from being the active partition
 - d. Select item 5 from the fdisk menu to update the disk configuration, and exit the fdisk menu.
 - 7. Display the partition menu and complete the following steps:
 - a. Print the current partition table.
 - b. Record the number of gigabytes (Gbytes) assigned to Slice 2.

Note – Various slices may already be defined. This task replaces all slices that already exist.



Gbytes:

-
- 8. Divide the number of Gbytes by 4 and complete the following steps:
 - a. Use the result as the number of Gbytes to assign as disk space to four slices.
 - b. Round *down* to the next whole Gbyte if the result includes a fraction.

For example, if the slice 2 is 76.67 Gbytes, $76.67 / 4 = 19.16$ Gbytes. In this case use 19 Gbytes as the size for the four slices.

Gbytes/4:

- Unauthorized reproduction or distribution prohibited. Copyright© 2011, Oracle and/or its affiliates.
-
9. Display the PARTITION menu again and complete the following steps:
 - a. Select Slice 0.
 - b. Accept the defaults for tags and flags.
 - c. Use the default starting cylinder. Typically this is cylinder 3.
 - d. Enter the resulting number of Gbytes from the previous step for the slice size.
 - e. Print the partition table again to verify the change.
 10. Attempt to select slice 8 and 9 for modification. What happens?

-
11. Set the sizes of Slices 1, 3, and 4 so that they are the same as Slice 0.

You should begin each successive slice on the cylinder that follows the ending cylinder of the previous slice. For example:

12. Set Slices 5, 6, and 7 to start at cylinder 0, and then assign them 0 Mbytes.

13. Complete the following steps:

- a. Print the partition table.

Except for slice 2, is there any overlap of ending and beginning cylinders for any of the slices listed?

-
- b. Proceed to the following steps to introduce this problem.

14. Complete the following steps:

- a. Add 1 to the number Gbytes/4 value listed in Step 6.

(Gbytes/4) + 1:

-
- b. Change Slice 0 so that it uses the new size.

The Cylinders column in the partition table should now indicate that Slice 0 ends after Slice 1 begins.

15. Complete the following steps:

- a. Use the modify command from the partition menu to attempt to fix this problem.

- b. Select Item 0 to modify the current partition table.

Which warnings display?

16. Modify the partition table and select Item 1 to use the All Free Hog option.
The partition table displays.
17. Observe the Cylinders and Size columns, and then verify that they are all zero except for slice 2, 8, and possibly 9.
18. Respond to the prompts to continue the process and complete the following steps:
 - a. Select Slice 4 as the All Free Hog partition.
 - b. Use the size listed in Step 8 for Slices 0, 1, and 3.
 - c. Enter Return to set the other slices to size 0.

At the end of this process, you should have three slices of equal size, where Slice 4 takes up any extra room if it exists.
19. Name the partition table "MYDISKpartition", and then label the disk.
20. Save your new partition table to the /etc/format.dat file and complete the following steps:
 - a. Carefully read the message that the format utility displays, and then enter the correct file name.
 - b. Quit the format utility when you have finished.
 - c. Use the tail command to view the contents of the /etc/format.dat file.
 - d. Note that your new partition information is appended to the file.
21. Verify your new partition table with the prtvtoc command.
22. Create two directories, one called /vtoc and another called /fdisk.
23. Use the prtvtoc command to print the partition table that you just created and complete the following steps:
 - a. Save its output to a file in the /vtoc directory.
 - b. Name the file so that it corresponds with the disk you are examining.
 - c. Use the cat command to verify that valid information exists in the file that you create.
24. Verify the fdisk partition table for your spare disk with the fdisk command.

- Unauthorized reproduction or distribution prohibited. Copyright© 2011, Oracle and/or its affiliates.
-
25. Use the fdisk command to print the fdisk partition table and complete the following steps:
 - a. Save its output to a file in the /fdisk directory.
 - b. Name the file so that it corresponds with the disk you are examining.
 - c. Use the cat command to verify that valid information exists in the file that you create.
 26. Use the following dd command to destroy the x86/x64 VTOC and complete the following steps:
 - a. Be certain to specify the correct disk device name for the of= argument.
 - b. Enter all other arguments exactly as listed.

```
# dd if=/dev/zero of=/dev/rdsk/c2d0s2 bs=512 count=2
2+0 records in
2+0 records out
#
```

27. Attempt to read the label from the same disk by using the prtvtoc command. What happens?

28. Run the format utility, and perform the following steps:
 - a. Select your spare disk.
 - b. Display the fdisk menu.
 - c. Delete the defined Solaris fdisk partition
 - d. Save the fdisk partition table, exit the fdisk menu, and quit the format utility.
29. Attempt to use the fmthard command to write to the disk the x86/x64 VTOC information you saved earlier. What happens?

30. Use the fdisk command to replace the fdisk partition table with the fdisk information you saved earlier.
31. Use the fmthard command to write to the disk the label information you saved earlier.
32. Attempt to read the label from the same disk.
Was this successful?

Exercise Summary

Discussion – Take a few minutes to discuss what experiences, issues, or discoveries you had during the lab exercises.



- Experiences
- Interpretations
- Conclusions
- Applications

Exercise Solutions - Managing Local Disk Devices

This section contains solutions to the exercise.

Task 1 – Identifying Device Files

Complete the following steps:

1. Log in as the root user and complete the following steps:
 - a. Open a terminal window.
 - b. Expand the window so that it occupies the entire screen area.
 - c. Change to the /dev/dsk directory.

```
# cd /dev/dsk
```

2. List the files in this directory and complete the following steps:

- a. Identify the files related to the boot disk of your system.

Systems typically use c0t0d0 or c1d0, but this may vary.

- b. Locate the file that represents slice 0 on this disk and display a long listing of it.

```
# ls
```

Sun Blade 1500 example:

```
# ls -l c0t0d0s0
```

Ultra 20 example:

```
# ls -l c1d0s0
```

Which type of file did you just locate? The file type indicator is the first character on the left side of the long listing.

Files in this directory are symbolic links. The letter l in the left-most column identifies a symbolic link.

- c. Record the full path name to which this file points.

Systems that use PCI bus architectures list path names similar to the following:

Sun Blade 1500 example:

```
.../.../devices/pci@1e,600000/ide@d/dad@0,0:a
```

Ultra 20 example:

```
.../.../devices/pci@0,0/pci-ide@7/ide@0/cmdk@0,0:a
```

3. Issue a long listing command of the path name you recorded.

Sun Blade 1500 example:

```
# ls -l ../../devices/pci@1e,600000/ide@d/dad@0,0:a
```

Ultra 20 example:

```
# ls -l ../../devices/pci@0,0/pci-ide@7/ide@0/cmdk@0,0:a
```

Which type of file is this?

Files in this directory are device files. The b character in the left-most column identifies a block-special device file.

The `ls -1L` command displays the same information, but shows only the link file name, not the real device file name.

4. Change to the `/dev/rdsk` directory and complete the following steps:

- a. Display a long listing of the same file name you selected in Step 2.

```
# cd /dev/rdsk
```

Sun Blade 1500 example:

```
# ls -l c0t0d0s0
```

Ultra 20 example:

```
# ls -l c1d0s0
```

Which type of file is this?

Files in this directory are symbolic links. The letter l in the left-most column identifies a symbolic link.

- b. Record the full path name to which this file points.

Systems that use PCI bus architectures list path names similar to the following:

Sun Blade 1500 example:

```
../../../../devices/pci@1e,600000/ide@d/dad@0,0:a,raw
```

Ultra 20 example:

```
../../../../devices/pci@0,0/pci-ide@7/ide@0/cmdk@0,0:a,raw
```

5. Issue a long listing command of the path name you recorded.

Sun Blade 1500 example:

```
# ls -l ../../devices/pci@1e,600000/ide@d/dad@0,0:a,raw
```

Ultra 20 example:

```
# ls -l ../../devices/pci@0,0/pci-ide@7/ide@0/cmdk@0,0:a,raw
```

Which type of file is this?

Files in this directory are device files. The c character in the left-most column identifies a character-special device file.

The `ls -lL` command displays the same information, but shows only the link file name, not the real device file name.

6. Change to the `/etc` directory and display the contents of the `path_to_inst` file.

```
# cd /etc
# more path_to_inst
```

7. Use the information from the previous steps to locate and record the entry for your boot disk.

A Sun Blade 1500 workstation, for example, would typically use `c0t0d0` as its boot disk. This relates to the device file called `dad@0,0` and is listed in the `/etc/path_to_inst` file.

Systems that use PCI bus architectures may have entries in `/etc/path_to_inst` similar to the following example from a Sun Blade 1500:

```
"/pci@1e,600000/ide@d/dad@0,0" 1 "dad"
```

Instance names are composed of the last two fields of each line in `/etc/path_to_inst`. Examples include the `dad`, `cmdk`, or `sd` tags, and the number that precedes them in `/etc/path_to_inst`.

What is the instance name for the device listed in this step?

The instance name for the device could be `dad1`, `dad0`, `cmdk0`, `sd3`, or `sd0`, depending on the system architecture.



Task 2 – Adding a New Disk or Tape Device

Complete the following steps:

Note – This task is optional. Perform the steps in this task if an external disk or tape device is available on your system.

1. In the /dev/dsk or /dev/rmt directories, confirm that no files exist for your external disk or tape device; for example, /dev/dsk/c1t0d0s0 or /dev/rmt/0.
If files for the external device do exist, ask your instructor for guidance. Do not simply remove device files that exist. Refer to the Preparation section for advice about removing device files.
2. If you are using a SPARC system, shut it down to run state 0.
init 0
3. If you are using an x86/x64 system, shut it down to run state 5. This shuts down the system and turns off the power.
init 5
4. Power on the external disk or tape drive attached to your system.
5. If you are using an x86/x64 system, power on your system and allow it to boot to its default run state.
6. If you are using a SPARC system, boot the system to its default run state.
ok boot
7. Log in as the root user and complete the following steps:
 - a. Open a terminal window.
 - b. Run the devfsadm command with the -v option to create new links and device files for the new disk or tape drive.
devfsadm -v
 - c. Observe the messages that the devfsadm command displays.
8. Confirm that new links and device files exist in the /dev/dsk and /dev/rdsk directories for disks or /dev/rmt for tape drives.
Use the devfsadm -v output to determine what files to look for.

Task 3 - Working With Disks and Partitions on SPARC

Complete the following steps:

1. Log in as the root user and complete the following steps:
 - a. Open a terminal window.
 - b. Run the **format** utility.

```
# format
```

Searching for disks...done

AVAILABLE DISK SELECTIONS:

0. c0t0d0 <HDS722512VLAT80 cyl 57459 alt 2 hd 16 sec 255>
/pci@1e,600000/ide@d/dad@0,0
1. c0t1d0 <DEFAULT cyl 57459 alt 2 hd 16 sec 255>
/pci@1e,600000/ide@d/dad@1,0

Specify disk (enter its number):

2. Complete the following steps:
 - a. Record the list of disks presented by the **format** utility. For example, c0t0d0 and c0t1d0.
 - b. Press **Control-D** to exit the **format** utility.

Specify disk (enter its number): ^D

```
#
```

3. Use the **prtvtoc** command to list the VTOC for each of the disks you found in the previous step and complete the following steps:
 - a. Examine the Mount Directory field in the information that the **prtvtoc** command displays.
 - b. Record the name of a disk that has no mount directory listed.
- For your classroom environment, this is an unused disk.



Note – Systems in production can use unmounted slices to store production data. Do not assume that unmounted slices are unused on systems in production. Systems will display various lists of partitions, so the list you see may differ from the list the examples.

Exercise Solutions - Managing Local Disk Devices

```
# prtvtoc /dev/rdsk/c0t1d0s2
* /dev/rdsk/c0t1d0s2 partition map
*
* Dimensions:
*      512 bytes/sector
*      255 sectors/track
*      16 tracks/cylinder
*     4080 sectors/cylinder
*    57461 cylinders
*   57459 accessible cylinders
*
* Flags:
*   1: unmountable
*  10: read-only
*
*          First        Sector        Last
* Partition Tag  Flags    Sector    Count    Sector Mount Directory
*      0      2    00           0    265200    265199
*      1      3    01    265200    265200    530399
*      2      5    01           0  234432720  234432719
*      6      4    00    530400  233902320  234432719
#
#
```

Unused disk: *The disk you identify will depend on your specific system.*

4. Run the format utility again and select the unused disk from the list of disks presented.

```
# format
```

Searching for disks...done

AVAILABLE DISK SELECTIONS:

0. c0t0d0 <HDS722512VLAT80 cyl 57459 alt 2 hd 16 sec 255>
/pci@1e,600000/ide@d/dad@0,0
1. c0t1d0 <DEFAULT cyl 57459 alt 2 hd 16 sec 255>
/pci@1e,600000/ide@d/dad@1,0

Specify disk (enter its number) :**1**

selecting c0t1d0

[disk formatted, no defect list found]

FORMAT MENU:

- | | |
|-----------|-------------------------------------|
| disk | - select a disk |
| type | - select (define) a disk type |
| partition | - select (define) a partition table |
| current | - describe the current disk |

```

format      - format and analyze the disk
repair      - repair a defective sector
show        - translate a disk address
label       - write label to the disk
analyze     - surface analysis
defect      - defect list management
backup      - search for backup labels
verify      - read and display labels
save        - save new disk/partition definitions
volname     - set 8-character volume name
!<cmd>    - execute <cmd>, then return
quit

format>

```

5. Display the partition menu and complete the following steps:
 - a. Print the current partition table.
 - b. Record the number of gigabytes (Gbytes) assigned to Slice 2.

```
format> part
```

PARTITION MENU:

```

0      - change '0' partition
1      - change '1' partition
2      - change '2' partition
3      - change '3' partition
4      - change '4' partition
5      - change '5' partition
6      - change '6' partition
7      - change '7' partition
select - select a predefined table
modify - modify a predefined partition table
name   - name the current table
print   - display the current table
label   - write partition map and label to the disk
!<cmd> - execute <cmd>, then return
quit

```

Exercise Solutions - Managing Local Disk Devices

```
partition> print
```

```
Current partition table (original) :
```

```
Total disk cylinders available: 57459 + 2 (reserved cylinders)
```

Part	Tag	Flag	Cylinders	Size	Blocks
0	root	wm	0 - 64	129.49MB	(65/0/0) 265200
1	swap	wu	65 - 129	129.49MB	(65/0/0) 265200
2	backup	wu	0 - 57458	111.79GB	(57459/0/0) 234432720
3	unassigned	wm	0	0	(0/0/0) 0
4	unassigned	wm	0	0	(0/0/0) 0
5	unassigned	wm	0	0	(0/0/0) 0
6	usr	wm	130 - 57458	111.53GB	(57329/0/0) 233902320
7	unassigned	wm	0	0	(0/0/0) 0

```
partition>
```

Note – The list of defined partitions will vary, depending on how the spare disk was last used.

Gbytes: *Your entry will depend on the disks present in your system.*

6. Divide the number of Gbytes by 4 and complete the following steps:
 - a. Use the result as the number of Gbytes to assign as disk space to four slices.
 - b. Round *down* to the next whole Gbyte if the result includes a fraction.
For example, if the slice 2 is 111.79 Gbytes, $111.79 / 4 = 27.94$ Gbytes. In this case use 27 Gbytes as the size for the four slices.
Gbytes/4: Your entry will depend on your system.
7. Display the Partition menu again and complete the following steps:
 - a. Select Slice 0.
 - b. Accept the defaults for tags and flags.
 - c. Start this first slice on Cylinder 0.
 - d. Enter the resulting number of Gbytes from the previous step for the slice size.

- e. Print the partition table again to verify the change.

```
partition> 0
Part      Tag     Flag      Cylinders      Size          Blocks
  0      root    wm        0 -      64      129.49MB      (65/0/0)
265200
```

```
Enter partition id tag[root]: <Return>
Enter partition permission flags[wm]: <Return>
Enter new starting cyl[0]: 0
Enter partition size[265200b, 65c, 64e, 129.49mb, 0.13gb]: 27gb
partition> print
Current partition table (unnamed):
Total disk cylinders available: 57459 + 2 (reserved cylinders)
```

Part	Tag	Flag	Cylinders	Size	Blocks
0	root	wm	0 - 13878	27.00GB	(13879/0/0) 56626320
1	swap	wu	65 - 129	129.49MB	(65/0/0) 265200
2	backup	wu	0 - 57458	111.79GB	(57459/0/0) 234432720
3	unassigned	wm	0	0	(0/0/0) 0
4	unassigned	wm	0	0	(0/0/0) 0
5	unassigned	wm	0	0	(0/0/0) 0
6	usr	wm	130 - 57458	111.53GB	(57329/0/0) 233902320
7	unassigned	wm	0	0	(0/0/0) 0

```
partition>
```

8. Set the sizes of Slices 1, 3, and 4 so that they are the same as Slice 0.

You should begin each successive slice on the cylinder that follows the ending cylinder of the previous slice. For example:

```
partition> 1
Part      Tag     Flag      Cylinders      Size          Blocks
  1      swap    wu        65 -      129      129.49MB      (65/0/0)
265200
```

```
Enter partition id tag[swap]: <Return>
Enter partition permission flags[wu]: <Return>
Enter new starting cyl[65]: 13879
Enter partition size[265200b, 65c, 13943e, 129.49mb, 0.13gb]: 27gb
partition> print
Current partition table (unnamed):
Total disk cylinders available: 57459 + 2 (reserved cylinders)
```

Exercise Solutions - Managing Local Disk Devices

```
Part      Tag    Flag     Cylinders      Size          Blocks
 0        root   wm      0 - 13878    27.00GB      (13879/0/0)  56626320
 1        swap   wu      13879 - 27757  27.00GB      (13879/0/0)  56626320
 2        backup  wu      0 - 57458   111.79GB      (57459/0/0)  234432720
 3 unassigned  wm      0           0           (0/0/0)       0
 4 unassigned  wm      0           0           (0/0/0)       0
 5 unassigned  wm      0           0           (0/0/0)       0
 6        usr    wm      130 - 57458  111.53GB      (57329/0/0)  233902320
 7 unassigned  wm      0           0           (0/0/0)       0
```

partition>

9. Set Slices 5, 6, and 7 to start at Cylinder 0, and then assign them 0 Mbytes.

partition> **6**

```
Part      Tag    Flag     Cylinders      Size          Blocks
 6        usr    wm      130 - 57458  111.53GB      (57329/0/0).
233902320
```

Enter partition id tag[usr]: <Return>

Enter partition permission flags[wm]: <Return>

Enter new starting cyl[130]: **0**

Enter partition size[233902320b, 57329c, 57328e, 114210.12mb, 111.53gb]:

0

partition>

10. Complete the following steps:

- a. Print the partition table.

partition> **print**

Current partition table (unnamed):

Total disk cylinders available: 57459 + 2 (reserved cylinders)

```
Part      Tag    Flag     Cylinders      Size          Blocks
 0        root   wm      0 - 13878    27.00GB      (13879/0/0)  56626320
 1        swap   wu      13879 - 27757  27.00GB      (13879/0/0)  56626320
 2        backup  wu      0 - 57458   111.79GB      (57459/0/0)  234432720
 3 unassigned  wm      27758 - 41636  27.00GB      (13879/0/0)  56626320
 4 unassigned  wm      41637 - 55515  27.00GB      (13879/0/0)  56626320
 5 unassigned  wm      0           0           (0/0/0)       0
 6 unassigned  wm      0           0           (0/0/0)       0
 7 unassigned  wm      0           0           (0/0/0)       0
```

partition>

Except for slice 2, is there any overlap of ending and beginning cylinders for any of the slices listed?

No.

- b. Proceed to the following steps to introduce this problem.
11. Complete the following steps:
 - a. Add 1 to the number Gbytes/4 value listed in Step 6.
(Gbytes/4) + 1: Your entry will depend on your system.
 - b. Change Slice 0 so that it uses the new size.

```
partition> 0
Part      Tag    Flag     Cylinders          Size            Blocks
  0       root   wm      0 - 13878        27.00GB        (13879/0/0)
56626320
```

```
Enter partition id tag[root] : <Return>
Enter partition permission flags[wm] : <Return>
Enter new starting cyl[0] : 0
Enter partition size[56626320b, 13879c, 13878e, 27649.57mb, 27.00gb] :
28gb
partition> print
Current partition table (unnamed):
Total disk cylinders available: 57459 + 2 (reserved cylinders)
```

Part	Tag	Flag	Cylinders	Size	Blocks
0	root	wm	0 - 14392	28.00GB	(14393/0/0) 58723440
1	swap	wu	13879 - 27757	27.00GB	(13879/0/0) 56626320
2	backup	wu	0 - 57458	111.79GB	(57459/0/0) 234432720
3	unassigned	wm	27758 - 41636	27.00GB	(13879/0/0) 56626320
4	unassigned	wm	41637 - 55515	27.00GB	(13879/0/0) 56626320
5	unassigned	wm	0	0	(0/0/0) 0
6	unassigned	wm	0	0	(0/0/0) 0
7	unassigned	wm	0	0	(0/0/0) 0

```
partition>
```

The Cylinders column in the partition table should now indicate that Slice 0 ends after Slice 1 begins.

12. Complete the following steps:
 - a. Use the modify command from the partition menu to attempt to fix this problem.
 - b. Select Item 0 to modify the current partition table.

```
partition> ?
Expecting one of the following: (abbreviations ok):
  0      - change '0' partition
  1      - change '1' partition
  2      - change '2' partition
  3      - change '3' partition
```

Exercise Solutions - Managing Local Disk Devices

4 - change '4' partition
5 - change '5' partition
6 - change '6' partition
7 - change '7' partition
select - select a predefined table
modify - modify a predefined partition table
name - name the current table
print - display the current table
label - write partition map and label to the disk
!<cmd> - execute <cmd>, then return
quit

partition> **modify**

Select partitioning base:

0. Current partition table (unnamed)
1. All Free Hog

Choose base (enter number) [0]? **0**

Warning: Overlapping partition (1) in table.

Warning: Fix, or select a different partition table.

partition>

Which warnings display?

Warning: Overlapping partition (1) in table.

Warning: Fix, or select a different partition table.

13. Modify the partition table and select Item 1 to use the All Free Hog option.

partition> **modify**

Select partitioning base:

0. Current partition table (unnamed)
1. All Free Hog

Choose base (enter number) [0]? **1**

Part	Tag	Flag	Cylinders	Size	Blocks	
0	root	wm	0	0	(0/0/0)	0
1	swap	wu	0	0	(0/0/0)	0
2	backup	wu	0 - 57458	111.79GB	(57459/0/0)	234432720
3	unassigned	wm	0	0	(0/0/0)	0
4	unassigned	wm	0	0	(0/0/0)	0
5	unassigned	wm	0	0	(0/0/0)	0
6	usr	wm	0	0	(0/0/0)	0
7	unassigned	wm	0	0	(0/0/0)	0

Do you wish to continue creating a new partition table based on above table[yes]?

The partition table displays.

14. Observe the Cylinders and Size columns, and then verify that they are all zero except for slice 2.
15. Respond to the prompts to continue the process and complete the following steps:
- Select Slice 4 as the All Free Hog partition.
 - Use the size listed in Step 6 for Slices 0, 1, and 3.
 - Enter Return to set the other slices to size 0.

table based on above table[yes]? **y**

Free Hog partition[6]? **4**

```
Enter size of partition '0' [0b, 0c, 0.00mb, 0.00gb]: 27gb
Enter size of partition '1' [0b, 0c, 0.00mb, 0.00gb]: 27gb
Enter size of partition '3' [0b, 0c, 0.00mb, 0.00gb]: 27gb
Enter size of partition '5' [0b, 0c, 0.00mb, 0.00gb]: <Return>
Enter size of partition '6' [0b, 0c, 0.00mb, 0.00gb]: <Return>
Enter size of partition '7' [0b, 0c, 0.00mb, 0.00gb]: <Return>
```

Part	Tag	Flag	Cylinders	Size	Blocks
0	root	wm	0 - 13878	27.00GB	(13879/0/0) 56626320
1	swap	wu	13879 - 27757	27.00GB	(13879/0/0) 56626320
2	backup	wu	0 - 57458	111.79GB	(57459/0/0) 234432720
3	unassigned	wm	27758 - 41636	27.00GB	(13879/0/0) 56626320
4	unassigned	wm	41637 - 57458	30.78GB	(15822/0/0) 64553760
5	unassigned	wm	0	0	(0/0/0) 0
6	usr	wm	0	0	(0/0/0) 0
7	unassigned	wm	0	0	(0/0/0) 0

Okay to make this the current partition table[yes]?

At the end of this process, you should have three slices of equal size, where Slice 4 takes up any extra room if it exists.

Exercise Solutions - Managing Local Disk Devices

16. Name the partition table "MYDISKpartition", and then label the disk.

Okay to make this the current partition table [yes]? **y**
Enter table name (remember quotes): "**MYDISKpartition**"

Ready to label disk, continue? **y**

partition> **q**

FORMAT MENU:

disk	- select a disk
type	- select (define) a disk type
partition	- select (define) a partition table
current	- describe the current disk
format	- format and analyze the disk
repair	- repair a defective sector
show	- translate a disk address
label	- write label to the disk
analyze	- surface analysis
defect	- defect list management
backup	- search for backup labels
verify	- read and display labels
save	- save new disk/partition definitions
volname	- set 8-character volume name
!<cmd>	- execute <cmd>, then return
quit	

format>

17. Save your new partition table to the /etc/format.dat file and complete the following steps:
- Carefully read the message that the format utility displays, and then enter the correct file name.
 - Quit the format utility when you have finished.
 - Use the tail command to view the contents of the /etc/format.dat file.

- d. Note that your new partition information is appended to the file.

```
format> save
Saving new disk and partition definitions
Enter file name["./format.dat"]: /etc/format.dat
format> quit
# tail /etc/format.dat
disk_type = "DEFAULT" \
: ctlr = ata : ncyl = 57459 : acyl = 2 : pcyl = 57461 \
: nhead = 16 : nsect = 255 : rpm = 5400

partition = "MYDISKpartition" \
: disk = "DEFAULT" : ctlr = ata \
: 0 = 0, 56626320 : 1 = 13879, 56626320 : 2 = 0, 234432720 \
: 3 = 27758, 56626320 : 4 = 41637, 64553760

#
```

18. Verify your new partition table with the prtvtoc command.

```
# prtvtoc /dev/rdsck/c0t1d0s2
* /dev/rdsck/c0t1d0s2 partition map
*
* Dimensions:
*      512 bytes/sector
*      255 sectors/track
*      16 tracks/cylinder
*      4080 sectors/cylinder
*      57461 cylinders
*      57459 accessible cylinders
*
* Flags:
*      1: unmountable
*      10: read-only
*
*                               First       Sector      Last
* Partition  Tag  Flags    Sector      Count     Sector Mount Directory
*          0   2   00           0  56626320  56626319
*          1   3   01  56626320  56626320 113252639
*          2   5   01           0 234432720 234432719
*          3   0   00 113252640  56626320 169878959
*          4   0   00 169878960  64553760 234432719
#
```

19. Create a directory called /vtoc.

```
# mkdir /vtoc
```

20. Use the prtvtoc command to print the partition table that you just created and complete the following steps:

- Save its output to a file in the /vtoc directory.
- Name the file so that it corresponds with the disk you are examining.
- Use the cat command to verify that valid information exists in the file that you create.

```
# prtvtoc /dev/rdsck/c0t1d0s2 > /vtoc/c0t1d0
# cat /vtoc/c0t1d0
* /dev/rdsck/c0t1d0s2 partition map
*
* Dimensions:
*      512 bytes/sector
*      255 sectors/track
*      16 tracks/cylinder
*     4080 sectors/cylinder
*    57461 cylinders
*   57459 accessible cylinders
*
* Flags:
*   1: unmountable
*  10: read-only
*
*          First       Sector      Last
* Partition Tag  Flags    Sector      Count      Sector Mount Directory
      0      2    00           0  56626320  56626319
      1      3    01  56626320  56626320 113252639
      2      5    01           0 234432720 234432719
      3      0    00 113252640  56626320 169878959
      4      0    00 169878960  64553760 234432719
#
```

21. Use the following dd command to destroy the disk label and complete the following steps:

- Be certain to specify the correct disk device name for the of= argument.
- Enter all other arguments exactly as listed.

```
# dd if=/dev/zero of=/dev/rdsck/c0t1d0s2 bs=512 count=1
1+0 records in
1+0 records out
#
```

22. Attempt to read the label from the same disk by using the prtvtoc command.

```
# prtvtoc /dev/rdsck/c0t1d0s2
```

What happens?

Different disk types present different results. SCSI disks might report messages that indicate that the disk label is unreadable. For example:

```
prtvtoc: /dev/rdsck/c1t3d0s2: Unable to read Disk
geometry errno = 0x16
```

IDE disks might report a partition table where only slice 2, or slice 0 and 2 remain defined. For example:

*			First	Sector	Last		
*	Partition	Tag	Flags	Sector	Count	Sector	Mount Directory
	2	5	01	0	17801280	17801279	

Example from a Sun Blade 1500:

*			First	Sector	Last		
*	Partition	Tag	Flags	Sector	Count	Sector	Mount Directory
	0	0	00	0	234432720	234432719	
	2	5	01	0	234432720	234432719	

#

23. Complete one of the following steps:

- If the prtvtoc command reported an “Unable to read Disk geometry” message, use the format utility to place a default label on the disk for which you destroyed the label earlier. For example:

Note – The device in this example, c1t3d0, is a disk of the type that would match the condition described above.

```
# format
```

```
Searching for disks...done
```

```
c1t3d0: configured with capacity of 8.43GB
```

AVAILABLE DISK SELECTIONS:

0. c0t0d0 <ST38410A cyl 16706 alt 2 hd 16 sec 63>
/pci@1f,0/pci@1,1/ide@3/dad@0,0
1. c1t3d0 <SUN9.0G cyl 4924 alt 2 hd 27 sec 133>
/pci@1f,0/pci@1/scsi@1/sd@3,0

Specify disk (enter its number): **1**

selecting c1t3d0

Exercise Solutions - Managing Local Disk Devices

```
[disk formatted]
Disk not labeled. Label it now? y
```

```
(format menu)
format> q
#
# prtvtoc /dev/rdsck/c1t3d0s2
(output omitted)
```

- If the `prtvtoc` command reports that only slice 2 exists, or slice 0 and 2 exist on the disk, proceed to step 24.
24. Use the `fmthard` command to write to the disk the label information you saved earlier.

```
# fmthard -s /vtoc/c0t1d0 /dev/rdsck/c0t1d0s2
fmthard: New volume table of contents now in place.
#
```

25. Attempt to read the label from the same disk.

```
# prtvtoc /dev/rdsck/c0t1d0s2
* /dev/rdsck/c0t1d0s2 partition map
*
* Dimensions:
*   512 bytes/sector
*   255 sectors/track
*   16 tracks/cylinder
*   4080 sectors/cylinder
*   57461 cylinders
*   57459 accessible cylinders
*
* Flags:
*   1: unmountable
*   10: read-only
*
*
*                               First      Sector      Last
* Partition  Tag  Flags    Sector     Count    Sector  Mount Directory
0          2    00        0  56626320  56626319
1          3    01  56626320  56626320 113252639
2          5    01        0 234432720 234432719
3          0    00 113252640  56626320 169878959
4          0    00 169878960  64553760 234432719
#
```

Was this successful?

This command should successfully read the disk label.

Task 4 - Working With Disks and Partitions on x86/x64

Complete the following steps:

1. Log in as the root user and complete the following steps:
 - a. Open a terminal window.
 - b. Run the **format** utility.

```
# format
```

Searching for disks...done

AVAILABLE DISK SELECTIONS:

0. c1d0 <DEFAULT cyl 39887 alt 2 hd 64 sec 63>
 /pci@0,0/pci-ide@7/ide@0/cmdk@0,0
1. c2d0 <DEFAULT cyl 10008 alt 2 hd 255 sec 63>
 /pci@0,0/pci-ide@7/ide@1/cmdk@0,0

Specify disk (enter its number):

2. Complete the following steps:
 - a. Record the list of disks presented by the **format** utility. For example, c1d0 and c2d0.
 - b. Press **Control-D** to exit the **format** utility.

Specify disk (enter its number): ^D

```
#
```

3. Use the **prtvtoc** command to list the VTOC for each of the disks you found in the previous step, and complete the following steps:
 - a. If the **prtvtoc** command displays an error similar to this example, you have identified the unused disk:

```
# prtvtoc /dev/rdsk/c2d0s2
```

prtvtoc: /dev/rdsk/c2d0s2: Unable to read Disk geometry errno = 0x5

```
#
```

If this error displays, record the name of this disk and skip steps b and c.

- b. If the **prtvtoc** command displays a partition table for both disks, examine the Mount Directory field in the command output.
- c. Record the name of a disk that has no mount directory listed. For your classroom environment, this is an unused disk.



Note – Systems in production can use unmounted slices to store production data. Do not assume that unmounted slices are unused on systems in production.

```
# prtvtoc /dev/rdsk/c2d0s2
* /dev/rdsk/c2d0s2 partition map
*
* Dimensions:
*      512 bytes/sector
*      63 sectors/track
*      255 tracks/cylinder
*      16065 sectors/cylinder
*      10010 cylinders
*      10008 accessible cylinders
*
* Flags:
*      1: unmountable
*      10: read-only
*
*
*          First        Sector      Last
* Partition Tag  Flags    Sector   Count   Sector Mount Directory
*      0      11  00      48195  819315  867509
*      1       3  01     867510 1060290 1927799
*      2       5  01      0 160778520 160778519
*      7       8  00 1927800 158850720 160778519
*      8       1  01      0      16065      16064
*      9       9  00      16065      32130      48194
#

```

Unused disk: *The disk you identify will depend on your specific system.*

4. Run the format utility again, and select the unused disk from the list of disks presented.
 - a. If the disk has no fdisk partition table defined, format asks you to select a drive type. Select the DEFAULT entry. The FORMAT menu displays.

```
# format
Searching for disks...done
```

AVAILABLE DISK SELECTIONS:

- 0. c1d0 <DEFAULT cyl 10002 alt 2 hd 255 sec 63>
/pci@0,0/pci-ide@7/ide@0/cmdk@0,0
- 1. c2d0 <drive type unknown>
/pci@0,0/pci-ide@7/ide@1/cmdk@0,0

Specify disk (enter its number): **1**

AVAILABLE DRIVE TYPES:

- 0. DEFAULT
- 1. other

Specify disk type (enter its number): **0**

selecting c2d0

No current partition list

No defect list found

[disk formatted, no defect list found]

FORMAT MENU:

- | | |
|------|-----------------|
| disk | - select a disk |
|------|-----------------|

(output omitted)

format>

- b. If the disk has a defined fdisk partition table, the FORMAT menu displays:

format

Searching for disks...done

AVAILABLE DISK SELECTIONS:

- 0. c1d0 <DEFAULT cyl 39887 alt 2 hd 64 sec 63>
/pci@0,0/pci-ide@7/ide@0/cmdk@0,0
- 1. c2d0 <DEFAULT cyl 10008 alt 2 hd 255 sec 63>
/pci@0,0/pci-ide@7/ide@1/cmdk@0,0

Specify disk (enter its number): **1**

selecting c2d0

Controller working list found

[disk formatted, defect list found]

FORMAT MENU:

- | | |
|-----------|-------------------------------------|
| disk | - select a disk |
| type | - select (define) a disk type |
| partition | - select (define) a partition table |
| current | - describe the current disk |
| format | - format and analyze the disk |

Exercise Solutions - Managing Local Disk Devices

```
fdisk      - run the fdisk program
repair     - repair a defective sector
show       - translate a disk address
label      - write label to the disk
analyze    - surface analysis
defect     - defect list management
backup     - search for backup labels
verify     - read and display labels
save       - save new disk/partition definitions
volname    - set 8-character volume name
!<cmd>   - execute <cmd>, then return
quit
format>
```

5. Display the fdisk menu and complete the following steps to create one Solaris fdisk partition that uses the whole disk:
 - a. If the disk has no defined fdisk partition table, the fdisk menu asks if you want to accept the default Solaris partition. Enter n to decline, and then edit the partition table.

format> **fdisk**

No fdisk table exists. The default partition for the disk is:

a 100% "SOLARIS System" partition

Type "y" to accept the default partition, otherwise type "n" to edit the partition table. **n**

Total disk size is 10011 cylinders
Cylinder size is 16065 (512 byte) blocks

Partition	Status	Type	Start	End	Length	Cylinders	%
=====	=====	=====	====	==	=====	=====	==

WARNING: no partitions are defined!

SELECT ONE OF THE FOLLOWING:

1. Create a partition
2. Specify the active partition
3. Delete a partition
4. Change between Solaris and Solaris2 Partition IDs
5. Exit (update disk configuration and exit)
6. Cancel (exit without updating disk configuration)

Enter Selection:

- b. If the disk has a defined fdisk partition table, the table displays when you enter the fdisk menu.

format> **fdisk**

Total disk size is 10011 cylinders
 Cylinder size is 16065 (512 byte) blocks

Partition	Status	Type	Cylinders				%
			Start	End	Length		
1		Solaris2	1	10010	10010		100

SELECT ONE OF THE FOLLOWING:

1. Create a partition
2. Specify the active partition
3. Delete a partition
4. Change between Solaris and Solaris2 Partition IDs
5. Exit (update disk configuration and exit)
6. Cancel (exit without updating disk configuration)

Enter Selection:

- c. Delete all fdisk partitions that exist. For example:

Enter Selection: **3**

Specify the partition number to delete (or enter 0 to exit): **1**

Are you sure you want to delete partition 1? This will make all files and programs in this partition inaccessible (type "y" or "n").

y

Total disk size is 10011 cylinders
 Cylinder size is 16065 (512 byte) blocks

Partition	Status	Type	Cylinders				%
			Start	End	Length		

WARNING: no partitions are defined!

SELECT ONE OF THE FOLLOWING:

1. Create a partition
2. Specify the active partition
3. Delete a partition
4. Change between Solaris and Solaris2 Partition IDs
5. Exit (update disk configuration and exit)
6. Cancel (exit without updating disk configuration)

Exercise Solutions - Managing Local Disk Devices

Enter Selection:

Partition 1 has been deleted.

6. Use the fdisk menu to create one Solaris fdisk partition that uses the whole disk.

Enter Selection: **1**

- a. Select the SOLARIS2 partition type

Select the partition type to create: **1**

1=SOLARIS2	2=UNIX	3=PCIXOS	4=Other
5=DOS12	6=DOS16	7=DOSEXT	8=DOSBIG
9=DOS16LBA	A=x86 Boot	B=Diagnostic	C=FAT32
D=FAT32LBA	E=DOSEXTLBA	F=EFI	0=Exit?

- b. Enter 100 to specify using 100% of the disk

Specify the percentage of disk to use for this partition
(or type "c" to specify the size in cylinders). **100**

- c. Enter n to prevent this partition from being the active partition

Should this become the active partition? If yes, it will be activated each time the computer is reset or turned on.

Please type "y" or "n". **n**

Total disk size is 10011 cylinders
Cylinder size is 16065 (512 byte) blocks

Partition	Status	Type	Cylinders				%
			Start	End	Length		
1		Solaris2	1	10010	10010		100

SELECT ONE OF THE FOLLOWING:

1. Create a partition
2. Specify the active partition
3. Delete a partition
4. Change between Solaris and Solaris2 Partition IDs
5. Exit (update disk configuration and exit)
6. Cancel (exit without updating disk configuration)

Enter Selection:

- d. Select item 5 from the fdisk menu to update the disk configuration, and exit the fdisk menu.

Enter Selection: **5**

format>

7. Display the partition menu and complete the following steps:
 - a. Print the current partition table.
 - b. Record the number of gigabytes (Gbytes) assigned to Slice 2.

format> **part**

PARTITION MENU:

```

0      - change '0' partition
1      - change '1' partition
2      - change '2' partition
3      - change '3' partition
4      - change '4' partition
5      - change '5' partition
6      - change '6' partition
7      - change '7' partition
select - select a predefined table
modify - modify a predefined partition table
name   - name the current table
print   - display the current table
label   - write partition map and label to the disk
!<cmd> - execute <cmd>, then return
quit
  
```

partition> **print**

Current partition table (original):

Total disk cylinders available: 10008 + 2 (reserved cylinders)

Part	Tag	Flag	Cylinders	Size	Blocks
0	reserved	wm	3 - 53	400.06MB	(51/0/0) 819315
1	swap	wu	54 - 119	517.72MB	(66/0/0) 1060290
2	backup	wu	0 - 10007	76.67GB	(10008/0/0) 160778520
3	unassigned	wu	0	0	(0/0/0) 0
4	unassigned	wu	0	0	(0/0/0) 0
5	unassigned	wu	0	0	(0/0/0) 0
6	unassigned	wu	0	0	(0/0/0) 0
7	home	wm	120 - 10007	75.75GB	(9888/0/0) 158850720
8	boot	wu	0 - 0	7.84MB	(1/0/0) 16065
9	alternates	wm	1 - 2	15.69MB	(2/0/0) 32130

partition>

Note – Various slices may already be defined. This task replaces all slices that already exist.



Gbytes: Your entry will depend on the disks present in your system.

8. Divide the number of Gbytes by 4 and complete the following steps:
 - a. Use the result as the number of Gbytes to assign as disk space to four slices.
 - b. Round *down* to the next whole Gbyte if the result includes a fraction.

For example, if the slice 2 is 76.67 Gbytes, $76.67 / 4 = 19.16$ Gbytes. In this case use 19 Gbytes as the size for the four slices.

Gbytes/4: Your entry will depend on your system.

9. Display the PARTITION menu again and complete the following steps:
 - a. Select Slice 0.
 - b. Accept the defaults for tags and flags.
 - c. Use the default starting cylinder. Typically this is cylinder 3.
 - d. Enter the resulting number of Gbytes from the previous step for the slice size.
 - e. Print the partition table again to verify the change.

```
partition> 0
Part      Tag    Flag    Cylinders          Size            Blocks
     0  reserved   wm        3 -      53    400.06MB    (51/0/0)
819315
```

Enter partition id tag[reserved]: <Return>

Enter partition permission flags[wm]: <Return>

Enter new starting cyl[3]: 3

Enter partition size[819315b, 51c, 53e, 400.06mb, 0.39gb]: 19gb

partition> print

Current partition table (unnamed):

Total disk cylinders available: 10008 + 2 (reserved cylinders)

Part	Tag	Flag	Cylinders	Size	Blocks
0	reserved	wm	3 - 2483	19.01GB	(2481/0/0) 39857265
1	swap	wu	54 - 119	517.72MB	(66/0/0) 1060290
2	backup	wu	0 - 10007	76.67GB	(10008/0/0) 160778520
3	unassigned	wu	0	0	(0/0/0) 0
4	unassigned	wu	0	0	(0/0/0) 0
5	unassigned	wu	0	0	(0/0/0) 0
6	unassigned	wu	0	0	(0/0/0) 0
7	home	wm	120 - 10007	75.75GB	(9888/0/0) 158850720
8	boot	wu	0 - 0	7.84MB	(1/0/0) 16065
9	alternates	wm	1 - 2	15.69MB	(2/0/0) 32130

partition>

10. Attempt to select slice 8 and 9 for modification. What happens?

partition> **8**

'8' is not expected.

partition> **9**

'9' is not expected.

partition>

The format utility prevents you from modifying these two pre-defined slices.

11. Set the sizes of Slices 1, 3, and 4 so that they are the same as Slice 0.

You should begin each successive slice on the cylinder that follows the ending cylinder of the previous slice. For example:

partition> **1**

Part	Tag	Flag	Cylinders	Size	Blocks
1	swap	wu	54 - 119	517.72MB	(66/0/0) 1060290

Enter partition id tag[swap] : <Return>

Enter partition permission flags[wu] : <Return>

Enter new starting cyl[54] : **2484**

Enter partition size[1060290b, 66c, 2549e, 517.72mb, 0.51gb] : **19gb**

partition> **print**

Current partition table (unnamed) :

Total disk cylinders available: 10008 + 2 (reserved cylinders)

Exercise Solutions - Managing Local Disk Devices

Part	Tag	Flag	Cylinders	Size	Blocks
0	reserved	wm	3 - 2483	19.01GB	(2481/0/0) 39857265
1	swap	wu	2484 - 4964	19.01GB	(2481/0/0) 39857265
2	backup	wu	0 - 10007	76.67GB	(10008/0/0) 160778520
3	unassigned	wu	0	0	(0/0/0) 0
4	unassigned	wu	0	0	(0/0/0) 0
5	unassigned	wu	0	0	(0/0/0) 0
6	unassigned	wu	0	0	(0/0/0) 0
7	home	wm	120 - 10007	75.75GB	(9888/0/0) 158850720
8	boot	wu	0 - 0	7.84MB	(1/0/0) 16065
9	alternates	wm	1 - 2	15.69MB	(2/0/0) 32130

partition>

12. Set Slices 5, 6, and 7 to start at cylinder 0, and then assign them 0 Mbytes.

partition> 7

Part	Tag	Flag	Cylinders	Size	Blocks
7	home	wm	120 - 10007	75.75GB	(9888/0/0)

Enter partition id tag[home]: <Return>
Enter partition permission flags[wm] : <Return>
Enter new starting cyl[120] : 0
Enter partition size[158850720b, 9888c, 9887e, 77563.83mb, 75.75gb] : 0
partition>

13. Complete the following steps:

- a. Print the partition table.

partition> print

Current partition table (unnamed) :

Total disk cylinders available: 10008 + 2 (reserved cylinders)

Part	Tag	Flag	Cylinders	Size	Blocks
0	reserved	wm	3 - 2483	19.01GB	(2481/0/0) 39857265
1	swap	wu	2484 - 4964	19.01GB	(2481/0/0) 39857265
2	backup	wu	0 - 10007	76.67GB	(10008/0/0) 160778520
3	unassigned	wu	4965 - 7445	19.01GB	(2481/0/0) 39857265
4	unassigned	wu	7446 - 9926	19.01GB	(2481/0/0) 39857265
5	unassigned	wu	0	0	(0/0/0) 0
6	unassigned	wu	0	0	(0/0/0) 0
7	unassigned	wm	0	0	(0/0/0) 0
8	boot	wu	0 - 0	7.84MB	(1/0/0) 16065
9	alternates	wm	1 - 2	15.69MB	(2/0/0) 32130

partition>

Except for slice 2, is there any overlap of ending and beginning cylinders for any of the slices listed?

No.

b. Proceed to the following steps to introduce this problem.

14. Complete the following steps:

a. Add 1 to the number Gbytes/4 value listed in Step 6.

(Gbytes/4) + 1: Your entry will depend on your system.

b. Change Slice 0 so that it uses the new size.

```
partition> 0
Part      Tag    Flag     Cylinders          Size        Blocks
  0  reserved   wm       3 -  2483      19.01GB    (2481/0/0)
39857265
```

Enter partition id tag[reserved]: <Return>

Enter partition permission flags[wm]: <Return>

Enter new starting cyl[3]: 3

Enter partition size[39857265b, 2481c, 2483e, 19461.55mb, 19.01gb]: 20gb

partition> print

Current partition table (unnamed):

Total disk cylinders available: 10008 + 2 (reserved cylinders)

Part	Tag	Flag	Cylinders	Size	Blocks
0	reserved	wm	3 - 2613	20.00GB	(2611/0/0) 41945715
1	swap	wu	2484 - 4964	19.01GB	(2481/0/0) 39857265
2	backup	wu	0 - 10007	76.67GB	(10008/0/0) 160778520
3	unassigned	wu	4965 - 7445	19.01GB	(2481/0/0) 39857265
4	unassigned	wu	7446 - 9926	19.01GB	(2481/0/0) 39857265
5	unassigned	wu	0	0	(0/0/0) 0
6	unassigned	wu	0	0	(0/0/0) 0
7	unassigned	wm	0	0	(0/0/0) 0
8	boot	wu	0 - 0	7.84MB	(1/0/0) 16065
9	alternates	wm	1 - 2	15.69MB	(2/0/0) 32130

partition>

The Cylinders column in the partition table should now indicate that Slice 0 ends after Slice 1 begins.

15. Complete the following steps:

- a. Use the modify command from the partition menu to attempt to fix this problem.
- b. Select Item 0 to modify the current partition table.

```
partition> ?
```

Expecting one of the following: (abbreviations ok) :

```
0      - change '0' partition
1      - change '1' partition
2      - change '2' partition
3      - change '3' partition
4      - change '4' partition
5      - change '5' partition
6      - change '6' partition
7      - change '7' partition
select - select a predefined table
modify - modify a predefined partition table
name   - name the current table
print   - display the current table
label   - write partition map and label to the disk
!<cmd> - execute <cmd>, then return
quit
```

```
partition> modify
```

Select partitioning base:

0. Current partition table (unnamed)
1. All Free Hog

```
Choose base (enter number) [0]? 0
```

Warning: Overlapping partition (1) in table.

Warning: Fix, or select a different partition table.

```
partition>
```

Which warnings display?

Warning: Overlapping partition (1) in table.

Warning: Fix, or select a different partition table.

16. Modify the partition table and select Item 1 to use the All Free Hog option.

partition> **modify**

Select partitioning base:

0. Current partition table (unnamed)
1. All Free Hog

Choose base (enter number) [0]? **1**

Part	Tag	Flag	Cylinders	Size	Blocks
0	root	wm	0	0	(0/0/0) 0
1	swap	wu	0	0	(0/0/0) 0
2	backup	wu	0 - 10007	76.67GB	(10008/0/0) 160778520
3	unassigned	wm	0	0	(0/0/0) 0
4	unassigned	wm	0	0	(0/0/0) 0
5	unassigned	wm	0	0	(0/0/0) 0
6	usr	wm	0	0	(0/0/0) 0
7	unassigned	wm	0	0	(0/0/0) 0
8	boot	wu	0 - 0	7.84MB	(1/0/0) 16065
9	alternates	wm	1 - 2	15.69MB	(2/0/0) 32130

Do you wish to continue creating a new partition table based on above table[yes]?

The partition table displays.

17. Observe the Cylinders and Size columns, and then verify that they are all zero except for slice 2, 8, and possibly 9.

18. Respond to the prompts to continue the process and complete the following steps:

- Select Slice 4 as the All Free Hog partition.
- Use the size listed in Step 8 for Slices 0, 1, and 3.
- Enter Return to set the other slices to size 0.

table based on above table[yes]? **y**

Free Hog partition[6]? **4**

```
Enter size of partition '0' [0b, 0c, 0.00mb, 0.00gb]: 19gb
Enter size of partition '1' [0b, 0c, 0.00mb, 0.00gb]: 19gb
Enter size of partition '3' [0b, 0c, 0.00mb, 0.00gb]: 19gb
Enter size of partition '5' [0b, 0c, 0.00mb, 0.00gb]: <Return>
Enter size of partition '6' [0b, 0c, 0.00mb, 0.00gb]: <Return>
Enter size of partition '7' [0b, 0c, 0.00mb, 0.00gb]: <Return>
```

Part	Tag	Flag	Cylinders	Size	Blocks
0	root	wm	3 - 2483	19.01GB (2481/0/0)	39857265
1	swap	wu	2484 - 4964	19.01GB (2481/0/0)	39857265
2	backup	wu	0 - 10007	76.67GB (10008/0/0)	160778520
3	unassigned	wm	4965 - 7445	19.01GB (2481/0/0)	39857265
4	unassigned	wm	7446 - 10007	19.63GB (2562/0/0)	41158530
5	unassigned	wm	0	0 (0/0/0)	0
6	usr	wm	0	0 (0/0/0)	0
7	unassigned	wm	0	0 (0/0/0)	0
8	boot	wu	0 - 0	7.84MB (1/0/0)	16065
9	alternates	wm	1 - 2	15.69MB (2/0/0)	32130

Okay to make this the current partition table [yes]?

At the end of this process, you should have three slices of equal size, where Slice 4 takes up any extra room if it exists.

19. Name the partition table "MYDISKpartition", and then label the disk.

Okay to make this the current partition table[yes]? **y**
Enter table name (remember quotes): "**MYDISKpartition**"

Ready to label disk, continue? **y**

partition> **q**

FORMAT MENU:

disk	- select a disk
type	- select (define) a disk type
partition	- select (define) a partition table
current	- describe the current disk
format	- format and analyze the disk
fdisk	- run the fdisk program
repair	- repair a defective sector
show	- translate a disk address
label	- write label to the disk
analyze	- surface analysis
defect	- defect list management
backup	- search for backup labels
verify	- read and display labels
save	- save new disk/partition definitions
volname	- set 8-character volume name
!<cmd>	- execute <cmd>, then return
quit	
format>	

20. Save your new partition table to the /etc/format.dat file and complete the following steps:

- a. Carefully read the message that the format utility displays, and then enter the correct file name.
- b. Quit the format utility when you have finished.
- c. Use the tail command to view the contents of the /etc/format.dat file.

- d. Your new partition information is appended to the file.

```
format> save
```

Saving new disk and partition definitions

Enter file name["./format.dat"]: **/etc/format.dat**

```
format> quit
```

```
# tail /etc/format.dat
```

```
: ctrlr = ata : ncyl = 10008 : acyl = 2 : pcyl = 10010 \
: nhead = 255 : nsect = 63 : rpm = 3600
```

```
partition = "MYDISKpartition" \
```

```
: disk = "DEFAULT" : ctrlr = ata \
: 0 = reserved, wm, 3, 39857265 : 1 = 2484, 39857265 \
: 2 = 0, 160778520 : 3 = unassigned, wu, 4965, 39857265 \
: 4 = unassigned, wu, 7446, 41158530 : 8 = 0, 16065 \
: 9 = 1, 32130
```

```
#
```

21. Verify your new partition table with the prtvtoc command.

```
# prtvtoc /dev/rdsck/c2d0s2
```

* /dev/rdsck/c2d0s2 partition map

*

* Dimensions:

- * 512 bytes/sector
- * 63 sectors/track
- * 255 tracks/cylinder
- * 16065 sectors/cylinder
- * 10010 cylinders
- * 10008 accessible cylinders

*

* Flags:

- * 1: unmountable
- * 10: read-only

*

* Partition	Tag	Flags	First	Sector	Last	Mount	Directory
			Sector	Count	Sector		
0	11	00	48195	39857265	39905459		
1	3	01	39905460	39857265	79762724		
2	5	01		0	160778520	160778519	
3	0	01	79762725	39857265	119619989		
4	0	01	119619990	41158530	160778519		
8	1	01		0	16065	16064	
9	9	00	16065	32130	48194		

```
#
```

22. Create two directories, one called /vtoc and another called /fdisk.

```
# mkdir /vtoc
# mkdir /fdisk
```

23. Use the prtvtoc command to print the partition table that you just created and complete the following steps:

- Save its output to a file in the /vtoc directory.
- Name the file so that it corresponds with the disk you are examining.
- Use the cat command to verify that valid information exists in the file that you create.

```
# prtvtoc /dev/rdsk/c2d0s2 > /vtoc/c2d0
# cat /vtoc/c2d0
* /dev/rdsk/c2d0s2 partition map
*
* Dimensions:
*      512 bytes/sector
*      63 sectors/track
*      255 tracks/cylinder
*     16065 sectors/cylinder
*    10010 cylinders
*   10008 accessible cylinders
*
* Flags:
*   1: unmountable
*  10: read-only
*
*          First       Sector      Last
* Partition Tag  Flags      Sector      Count      Sector  Mount Directory
*            0        11    00        48195  39857265  39905459
*            1         3    01      39905460  39857265  79762724
*            2         5    01                  0 160778520  160778519
*            3         0    01      79762725  39857265  119619989
*            4         0    01    119619990  41158530  160778519
*            8         1    01                  0      16065      16064
*            9         9    00      16065      32130      48194
*
```

24. Verify the fdisk partition table for your spare disk with the fdisk command.

```
# fdisk -W - /dev/rdsk/c2d0p0
*
* /dev/rdsk/c2d0p0 default fdisk table
* Dimensions:
*      512 bytes/sector
```

Exercise Solutions - Managing Local Disk Devices

* 63 sectors/track
* 255 tracks/cylinder
* 10011 cylinders
*
* systid:
* 1: DOSOS12
* 2: PCIXOS
* 4: DOSOS16
* 5: EXTDOS
* 6: DOSBIG
* 7: FDISK_IFS
* 8: FDISK_AIXBOOT
* 9: FDISK_AIXDATA
* 10: FDISK_0S2BOOT
* 11: FDISK_WINDOWS
* 12: FDISK_EXT_WIN
* 14: FDISK_FAT95
* 15: FDISK_EXTLBA
* 18: DIAGPART
* 65: FDISK_LINUX
* 82: FDISK_CPM
* 86: DOSDATA
* 98: OTHEROS
* 99: UNIXOS
* 101: FDISK_NOVELL3
* 119: FDISK_QNX4
* 120: FDISK_QNX42
* 121: FDISK_QNX43
* 130: SUNIXOS
* 131: FDISK_LINUXNAT
* 134: FDISK_NTFSVOL1
* 135: FDISK_NTFSVOL2
* 165: FDISK_BSD
* 167: FDISK_NEXTSTEP
* 183: FDISK_BSDIFS
* 184: FDISK_BSDISWAP
* 190: X86BOOT
* 191: SUNIXOS2
* 238: EFI_PMBR
* 239: EFI_FS
*
* Id Act Bhead Bsect Bcyl Ehead Esect Ecyl Rsect Numsect
 191 0 0 1 1 254 63 1023 16065
160810650
#

25. Use the fdisk command to print the fdisk partition table and complete the following steps:
- Save its output to a file in the /fdisk directory.
 - Name the file so that it corresponds with the disk you are examining.
 - Use the cat command to verify that valid information exists in the file that you create.

```
# fdisk -W /fdisk/c2d0p0 /dev/rdsck/c2d0p0
# cat /fdisk/c2d0p0

* /dev/rdsck/c2d0p0 default fdisk table
* Dimensions:
*   512 bytes/sector
*   63 sectors/track
*   255 tracks/cylinder
*   10011 cylinders
*
(output omitted)
```

26. Use the following dd command to destroy the x86/x64 VTOC and complete the following steps:
- Be certain to specify the correct disk device name for the of= argument.
 - Enter all other arguments exactly as listed.

```
# dd if=/dev/zero of=/dev/rdsck/c2d0s2 bs=512 count=2
2+0 records in
2+0 records out
#
#
```

27. Attempt to read the label from the same disk by using the prtvtoc command. What happens?

```
# prtvtoc /dev/rdsck/c2d0s2
* /dev/rdsck/c2d0s2 partition map
*
* Dimensions:
*   512 bytes/sector
*   63 sectors/track
*   255 tracks/cylinder
*   16065 sectors/cylinder
*   10010 cylinders
*   10008 accessible cylinders
*
* Flags:
*   1: unmountable
```

Exercise Solutions - Managing Local Disk Devices

```
* 10: read-only
*
* Unallocated space:
*      First    Sector    Last
*      Sector    Count    Sector
*      48195 160730325 160778519
*
*                                         First    Sector    Last
* Partition Tag  Flags      Sector    Count    Sector  Mount Directory
*      2      5   01          0 160778520 160778519
*      8      1   01          0 16065      16064
*      9      9   00          16065     32130      48194
#
```

Typically on x86/x64 systems, a default partition table displays. In this example from an Ultra 20, only slices 2, 8 and 9 remain defined. All other partitions have been destroyed.

28. Run the **format** utility, and perform the following steps:

- a. Select your spare disk.

```
# format
```

```
Searching for disks...done
```

AVAILABLE DISK SELECTIONS:

0. c1d0 <DEFAULT cyl 39887 alt 2 hd 64 sec 63>
/pci@0,0/pci-ide@7/ide@0/cmdk@0,0
1. c2d0 <DEFAULT cyl 10008 alt 2 hd 255 sec 63>
/pci@0,0/pci-ide@7/ide@1/cmdk@0,0

Specify disk (enter its number) : **1**

selecting c2d0

Controller working list found

[disk formatted, defect list found]

FORMAT MENU:

- | | |
|-----------|-------------------------------------|
| disk | - select a disk |
| type | - select (define) a disk type |
| partition | - select (define) a partition table |
| current | - describe the current disk |
| format | - format and analyze the disk |
| fdisk | - run the fdisk program |
| repair | - repair a defective sector |
| show | - translate a disk address |
| label | - write label to the disk |
| analyze | - surface analysis |

defect - defect list management
 backup - search for backup labels
 verify - read and display labels
 save - save new disk/partition definitions
 volname - set 8-character volume name
 !<cmd> - execute <cmd>, then return
 quit

format>

- b. Display the fdisk menu.

format> **fdisk**

Total disk size is 10011 cylinders
 Cylinder size is 16065 (512 byte) blocks

Partition	Status	Type	Cylinders				%
			Start	End	Length	====	
1		Solaris2	1	10010	10010		100

SELECT ONE OF THE FOLLOWING:

1. Create a partition
2. Specify the active partition
3. Delete a partition
4. Change between Solaris and Solaris2 Partition IDs
5. Exit (update disk configuration and exit)
6. Cancel (exit without updating disk configuration)

Enter Selection:

- c. Delete the defined Solaris fdisk partition

Enter Selection: **3**

Specify the partition number to delete (or enter 0 to exit): **1**

Are you sure you want to delete partition 1? This will make all files and programs in this partition inaccessible (type "y" or "n").

y

- d. Save the fdisk partition table, exit the fdisk menu, and quit the format utility.

```
Total disk size is 10011 cylinders
Cylinder size is 16065 (512 byte) blocks
```

Partition	Status	Type	Start	End	Length	%
=====	=====	=====	====	==	=====	==

WARNING: no partitions are defined!

SELECT ONE OF THE FOLLOWING:

1. Create a partition
2. Specify the active partition
3. Delete a partition
4. Change between Solaris and Solaris2 Partition IDs
5. Exit (update disk configuration and exit)
6. Cancel (exit without updating disk configuration)

Enter Selection: **5**

Solaris fdisk partition not found

No fdisk solaris partition found

format> **q**

#

29. Attempt to use the fmthard command to write to the disk the x86/x64 VTOC information you saved earlier. What happens?

```
# fmthard -s /vtoc/c2d0 /dev/rdsck/c2d0s2
/dev/rdsck/c2d0s2: Cannot get disk geometry
#
```

The attempt fails because the fmthard command can not find a slice 2 defined on the disk.

30. Use the fdisk command to replace the fdisk partition table with the fdisk information you saved earlier.

```
# fdisk -F /fdisk/c2d0p0 /dev/rdsck/c2d0s0
#
```

31. Use the fmthard command to write to the disk the label information you saved earlier.

```
# fmthard -s /vtoc/c2d0 /dev/rdsck/c2d0s2
fmthard: New volume table of contents now in place.
#
```

32. Attempt to read the label from the same disk.

```
# prtvtoc /dev/rdsk/c2d0s2
* /dev/rdsk/c2d0s2 partition map
*
* Dimensions:
*      512 bytes/sector
*      63 sectors/track
*     255 tracks/cylinder
*    16065 sectors/cylinder
*   10010 cylinders
*  10008 accessible cylinders
*
* Flags:
*   1: unmountable
*  10: read-only
*
*          First       Sector      Last
* Partition Tag  Flags    Sector      Count    Sector Mount Directory
*            0        00      48195  39857265  39905459
*            1        03      39905460 39857265  79762724
*            2        05      01           0 160778520 160778519
*            3        00      01      79762725 39857265 119619989
*            4        00      01  119619990 41158530 160778519
*            8        01      01           0      16065      16064
*            9        09      00      16065      32130      48194
#
```

Was this successful?

This command should successfully read the disk label.

Notes:

Module 3

Interface Configurations

Objectives

Upon completion of this module, you should be able to:

- Control and monitor network interfaces
- Configure internet protocol version 4 (IPv4) interfaces at boot time

Controlling and Monitoring Network Interfaces

Network commands, such as `ifconfig`, `ping`, and `snoop`, control and monitor the functionality of network interfaces.

Displaying the MAC Address

The media access control (MAC) address is your computer's unique hardware address on a local area network (LAN). The MAC address is also the Ethernet address on an Ethernet LAN. When you are connected to a LAN, an address resolution table maps your computer's physical MAC address to an Internet Protocol (IP) address on the LAN. Two ways to display the MAC address or the Ethernet address are:

- Use the `ifconfig -a` command
- Use the boot programmable read-only memory (PROM) banner command

Note – The MAC address is displayed only if run as the root user.



ifconfig -a

```
lo0: flags=1000849<UP,LOOPBACK,RUNNING,MULTICAST,IPv4,VIRTUAL> mtu 8232
index 1 inet 127.0.0.1 netmask ff000000
nge0: flags=1000843<UP,BROADCAST,RUNNING,MULTICAST,IPv4> mtu 1500 index 2
    inet 192.168.30.41 netmask ffffff00 broadcast 192.168.30.255
    ether 8:0:20:93:c9:af
```

The MAC address is listed as `8:0:20:93:c9:af` in this example.

You can also retrieve the MAC address from a SPARC-based system that has not yet been booted by running the `banner` command at the `ok` prompt.

ok **banner**

```
Sun Ultra 5/10 UPA/PCI (UltraSPARC-III 300MHz), Keyboard Present
OpenBoot 3.31 256 MB (60ns) memory installed, Serial #9685423.
Ethernet address 8:0:20:93:c9:af, Host ID: 8093c9af.
```

Displaying the IP Address

The `ifconfig -a` command displays the current configuration for the network interfaces.

With the `-a` option, the `ifconfig` command displays the current configuration for all network interfaces in the system.

```
# ifconfig -a
lo0: flags=1000849<UP,LOOPBACK,RUNNING,MULTICAST,IPv4,VIRTUAL> mtu 8232
index 1 inet 127.0.0.1 netmask ff000000
nge0: flags=1000843<UP,BROADCAST,RUNNING,MULTICAST,IPv4> mtu 1500 index 2
        inet 192.168.30.41 netmask ffffff00 broadcast 192.168.30.255
              ether 8:0:20:93:c9:af
```

The previous example shows that the loopback interface (`lo0`) is up, running, and configured with an IP address of `127.0.0.1`. The `nge0` interface is up, running, and configured with an IP address of `192.168.30.41`.

Marking an Ethernet Interface as Down

When an Ethernet interface is marked as down, it means that it cannot communicate. You can use the `ifconfig` command to mark an Ethernet interface as up or down. For example, to mark the `nge0` interface as down, perform the commands:

```
# ifconfig nge0 down
# ifconfig -a
lo0: flags=1000849<UP,LOOPBACK,RUNNING,MULTICAST,IPv4,VIRTUAL> mtu 8232
index 1 inet 127.0.0.1 netmask ff000000
nge0: flags=1000842<BROADCAST,RUNNING,MULTICAST,IPv4> mtu 1500 index 2
        inet 192.168.30.41 netmask ffffff00 broadcast 192.168.30.255
              ether 8:0:20:93:c9:af
```



Note – The UP flag for `nge0` is no longer present. When an interface is flagged as UP, it is ready to communicate.

The following example shows that when you mark an interface as up, the UP status appears in the flags field of the `ifconfig` command output:

```
# ifconfig nge0 up
# ifconfig -a
```

```
lo0: flags=1000849<UP,LOOPBACK,RUNNING,MULTICAST,IPv4,VIRTUAL> mtu 8232
index 1 inet 127.0.0.1 netmask ff000000
nge0: flags=1000843<UP,BROADCAST,RUNNING,MULTICAST,IPv4> mtu 1500 index 2
    inet 192.168.30.41 netmask ffffff00 broadcast 192.168.30.255
    ether 8:0:20:93:c9:af
```

Sending ICMP ECHO_REQUEST Packets

To determine if you can contact another system over the network, enter the ping command:

```
# ping sys41
sys41 is alive
```

The previous response indicates the host name sys41 is alive. A response of no answer from sys41 indicates that you cannot contact host sys41. This implies a problem with host sys41, or a problem with the network.

For the ping command to succeed, the following conditions must be satisfied on both systems:

- The interface must be plumbed.
- The interface must be configured.
- The interface must be up.
- The interface must be physically connected.
- The interface must have valid routes configured.

For more information on ifconfig and plumbed, see the ifconfig man page.

Note – You can track the route a packet follows using traceroute. This utility traces the route that an IP packet follows to another internet host.



Capturing and Inspecting Network Packets

Use the snoop utility to capture and inspect network packets to determine what kind of data is transferred between systems. Use the snoop utility to see what happens when one system uses the ping command to communicate with another system. To view network traffic between two specific systems, perform the command:

```
# snoop sys41 sys42
sys41 -> sys42 ICMP Echo request (ID: 615 Sequence number: 0)
sys42 -> sys41 ICMP Echo reply (ID: 615 Sequence number: 0)
```

Use the **-a** option to enable audible clicks, which notify you of any network traffic. Although noisy, the clicks are useful when troubleshooting.

The following example shows how to turn on audible clicks for all network traffic related to a Dynamic Host Configuration Protocol (DHCP) boot:

```
# snoop -a dhcp
```

Additional snoop options include:

snoop	Summary output
snoop -o <i>filename</i>	Saves captured packets in <i>filename</i> as they are captured
snoop -i <i>filename</i>	Displays packets previously captured in <i>filename</i>
snoop -i <i>filename</i> -v	Filters output to verbose mode
snoop -i <i>filename</i> -V	Filters output to summary verbose mode
snoop -i <i>filename</i> -x <i>offset</i>	Displays packet data in hexadecimal and ASCII format.
snoop -d <i>device</i>	Receives packets from a network interface specified by <i>device</i>

Note – Press Control-C to stop the snoop utility.



Configuring IPv4 Interfaces at Boot Time

This section describes the files and scripts involved with configuring IPv4 network interfaces.

Introducing IPv4 Interface Files

You can get a basic understanding of network interfaces within the Solaris OS by learning the function of a few files and services. The services and files are the following:

- The svc:/network/physical:default service
- The /etc/hostname.*xxn* file
- The /etc/inet/hosts file
- The /etc/inet/ipnodes file

The svc:/network/physical:default Service

The svc:/network/physical:default service calls the /lib/svc/method/net-physical method script. It is one of the startup scripts that runs each time you boot the system. This script uses the ifconfig utility to configure each interface with an IP address and other required network information. The script searches for files called hostname.*xxn* in the /etc directory, where *xx* is an interface type and *n* is the instance of the interface. For every file named /etc/hostname.*xxn*, the script uses the ifconfig command with the plumb option to make the kernel ready to talk to this type of interface. The script then configures the named interface using other options to the ifconfig command. The /etc/hostname.nge0 file is an example of an interface configuration file.

Note – In Solaris 8 and 9 OS, the etc/rcS.d/S30network.sh file is used to perform the same function. Before Solaris 8 OS, the /etc/rcS.d/S30rootusr.sh file was used.



The `/etc/hostname.xxn` File

The `/etc/hostname.xxn` file contains an entry that configures a corresponding interface. The variable component of the file name is replaced by an interface type and a number that differentiates between multiple interfaces of the same type configured in the system. Table 3-1 shows examples of file entries for Ethernet interfaces commonly found in Sun systems.

Table 3-1 The `/etc/hostname.xxn` File Entries and Corresponding Interfaces

Entry	Interface
<code>/etc/hostname.e1000g0</code>	First e1000g (Intel PRO/1000 Gigabit family device driver) Ethernet interface in the system
<code>/etc/hostname.e1000g1</code>	Second e1000g Ethernet interface in the system
<code>/etc/hostname.bge0</code>	First bge (Broadcom Gigabit Ethernet device driver) Ethernet interface in the system
<code>/etc/hostname.bge1</code>	Second bge Ethernet interface in the system
<code>/etc/hostname.ce0</code>	First ce (Cassini Gigabit-Ethernet device driver) Ethernet interface in the system
<code>/etc/hostname.qfe0</code>	First qfe (Quad Fast-Ethernet device driver) Ethernet interface in the system
<code>/etc/hostname.hme0</code>	First hme (Fast-Ethernet device driver) Ethernet interface in the system
<code>/etc/hostname.eri0</code>	First eri (eri Fast-Ethernet device driver) Ethernet interface in the system
<code>/etc/hostname.nge0</code>	First nge (Nvidia Gigabit Ethernet driver) Ethernet interface in the system

The codes for the interface types are product codes. These codes originate from varying sources. For example, the `e1000g` code is for the Intel PRO/1000 Gigabit and the `bge` code is an abbreviation for Broadcom Gigabit Ethernet.

The /etc/hostname.nge0 file contains either the host name or the IP address of the system that contains the nge0 interface. The host name contained in the file must exist in the /etc/inet/hosts file so that it can be resolved to an IP address at system boot time. You can edit the /etc/hostname.nge0 file to contain either the host name or the IP address from the /etc/inet/hosts file.

```
# cat /etc/hostname.nge0
sys41
or
# cat /etc/hostname.nge0
192.168.30.41
```

The /etc/inet/hosts File

The /etc/inet/hosts file is a local database that associates the IP addresses of hosts with their names. You can use the /etc/inet/hosts file with, or instead of, other hosts databases, including the Domain Name System (DNS), the Network Information Service (NIS) hosts map, and the Network Information Service Plus (NIS+) hosts table. Programs use library interfaces to access information in the /etc/inet/hosts file.

The /etc/inet/hosts file contains at least the loopback and host information. The file has one entry for each IP address of each host. If a host has more than one IP address, this file will have one entry for each address, on separate lines. The format of each line is:

IP-address official-host-name [aliases] . . .

Items are separated by any number of spaces or tab characters. The first item on a line is the host's IP address. The second entry is the host's official name. Subsequent entries on the same line are alternative names for the same machine, or nicknames. Nicknames are optional.

```
# cat /etc/inet/hosts
.
<output truncated>
.
127.0.0.1 localhost
.
<output truncated>
.
192.168.30.41 sys41          loghost      #connection to nge0 interface
.
<output truncated>
```



Note – The /etc/inet/hosts file is the official (system V release 4) SVr4 name of the hosts file. The symbolic link /etc/inet/hosts exists for Berkeley Software Distribution (BSD) compatibility.

The /etc/iphnodes File

The ipnodes file is a symbolic link to the /etc/inet/hosts file. It associates the names of nodes with their internet protocol (IP) addresses. The ipnodes file can be used in conjunction with, or instead of, other ipnodes databases, including the domain name system (DNS), the NIS ipnodes map, and LDAP.

The ipnodes file has one entry for each IP address of each node, and can contain either IPv4 or IPv6 addresses.

If a node has more than one IP address, it will have one entry for each, on consecutive lines. The format of each line is:

IP-address official-node-name nicknames...

Items are separated spaces or tab characters. The first item on a line is the host IP address. The second entry is the host official name. Subsequent entries on the same line are alternative names for the same machine, or nicknames. Nicknames are optional.

```
# cat /etc/inet/iphnodes
#
# Internet host table
#
::1      localhost
127.0.0.1      localhost
192.168.30.41sys41      loghost
```

IP addresses can be defined in the ipnodes file or in the hosts file. The ipnodes file will be searched first, then the hosts file.

Changing the System Host Name

The host name of a system is contained in four files on the system. You must modify all of these files, and perform a reboot, to successfully change a system's host name. The files that contain the host name of a system are:

- The /etc/nodename file
- The /etc/hostname.*xxn* file
- The /etc/inet/hosts file

Note – /etc/inet/nodes was used for IPv6 addresses but now it is a link to /etc/inet/hosts.

Note – If crash dump is enabled on a system, the system name must be changed under /var/crash. Older versions of Solaris also had the hostname in files located under /etc/net/tic*/*.

Use the `uname -s` command to change the system hostname on a temporarily.

Editing the /etc/nodename File

Each Solaris OS has a canonical name, which is the official name used when referring to a system. By convention, the system name is the same as the host name associated with the IP address of the primary network interface; for example, `hostname.nge0`.

The following example shows a system's /etc/nodename file:

```
# cat /etc/nodename
sys41
```

Change the canonical name by editing the /etc/nodename file, and rebooting the system.

If the machine network configuration is managed remotely and delivered by the DHCP or remote procedure call (RPC) bootparams protocols, the /etc/nodename file is not used. The file is not used because the remote service delivers the canonical name.

The /etc/hostname.*xxn* File

You do not have to edit the /etc/inet/ipnodes file. It is the /etc/hosts file.

Editing the /etc/inet/hosts File

Network addresses are written in the conventional decimal-dot notation.

Host names are text strings up to 24 characters. Alphabetic characters, numbers, the (-) sign, and a (.) are allowed in the host name. Periods are only allowed when they serve to delimit components of domain style names. Spaces are not allowed in the host name. The first character must be an alphabetic character. The last character must not be a (-) or a (.).

No distinction is made between uppercase and lowercase characters, unless the NIS naming service is used. Uppercase characters in names have been known to cause problems with NIS.

A (#) indicates the beginning of a comment. After a comment character, all characters, up to the end of the line, are not interpreted.

The sys-unconfig Command

Use the /usr/sbin/sys-unconfig command to undo a system configuration. Use the /usr/sbin/sys-unconfig command to restore a system configuration to an unconfigured state. It will be ready to be reconfigured.

The sys-unconfig command does the following:

- Saves the current /etc/inet/hosts file information in the /etc/inet/hosts.saved file.
- If the current /etc/vfstab file contains Network File System (NFS) mount entries, it saves the /etc/vfstab file to the /etc/vfstab.orig file.
- Restores the default /etc/inet/hosts file.
- Removes the default host name in the /etc/hostname.*xxn* files for all configured interfaces.
- Removes the default domain name in the /etc/defaultdomain file.
- Restores the time zone to PST8PDT in the /etc/TIMEZONE file.
- Resets naming services to local files.

Configuring IPv4 Interfaces at Boot Time

- Removes the /etc/inet/netmasks file.
- Removes the /etc/defaultrouter file.
- Removes the password set for the root user in the /etc/shadow file.
- Removes the /etc/.rootkey file for NIS+.
- Executes all system configuration applications. These applications are defined by prior executions of a sysidconfig -a command.
- Removes the /etc/resolv.conf file for DNS clients.
- Disables Lightweight Directory Access Protocol (LDAP) by removing:
 - The /var/ldap/ldap_client_cache file
 - The /var/ldap/ldap_client_file file
 - The /var/ldap/ldap_client_cred file
 - The /var/ldap/cachemgr.log file
- Regenerates keys for the Secure Shell Daemon (sshd).

When the sys-unconfig command is finished, it performs a system shutdown. The sys-unconfig command is a potentially dangerous utility and can only be run by the root user.

When you restart the system, a configuration script prompts you to configure the system information. The sys-unconfig command is not available on diskless clients.

Exercise: The Solaris OS Network Commands

In this exercise, you use basic network-related commands.

Preparation

To prepare for this exercise, perform the following tasks:

- Check that you have two systems listed in each /etc/inet/hosts file on each system.
- Work with a partner for this exercise, and perform all steps on both systems, unless noted otherwise.

Task

Complete the following steps using the ifconfig utility, the ping command, and the snoop utility.



Note – Be sure to work closely with your partner during the lab to ensure you are both working on the same steps.

1. On both systems, log in as the root user, and open a terminal window. Using the ifconfig -a command, display basic configuration information about your network interfaces.

For your primary interface (for example nge0), what does the ifconfig command report for the following attributes? Enter your values into Table 3-2.

Table 3-2 Primary Interface Values

Attribute	Value
IP address	
Ethernet address	
Interface up/down	

2. On both systems, open a new terminal window. In the new window, enter the snoop command to display the network traffic between your two systems only.

Exercise: The Solaris OS Network Commands

3. Use the ping command to verify that your system can contact the network interface on your partner's system.
 4. Observe the output from the snoop command. Which protocol does the ping command use?
-

Does the snoop output contain requests and replies (yes or no)?

Requests: Replies:

5. On one system, use the ifconfig command to mark its primary interface as down and then again to display its configuration information. Warn your lab partner that the system's interface will go down.

Has anything changed in the information that the ifconfig command reports?

6. On the system whose interface remains up, attempt to use the ping command to contact the system whose interface is down.

What does the ping command report?

7. Observe the output from the snoop utility on both systems. How does the snoop output differ from the ping command output before and after you marked the interface as down?
-

How many requests does the ping command send by default?

Does the target system see the ping command requests? If so, how are these requests handled?

8. On the system whose interface is down, use the ifconfig command to mark its primary interface as up. Check that the change took place.
9. On the system whose interface remained up, use the ping command to contact the other system.

What does the ping command report?

Does the snoop utility report a reply from the target host?

Exercise Summary



Discussion – Take a few minutes to discuss what experiences, issues, or discoveries you had during the lab exercise.

- Experiences
- Interpretations
- Conclusions
- Applications

Exercise Solutions: The Solaris OS Network Commands

In this exercise, you use basic network-related commands.

Preparation

To prepare for this exercise, perform the following tasks:

- Check that you have two systems listed in each /etc/inet/hosts file on each system.
- Work with a partner for this exercise, and perform all steps on both systems, unless noted otherwise.



Task and Solutions

This section describes the tasks for you to perform, and lists the solutions. Complete the following steps using the `ifconfig` utility, the `ping` command, and the `snoop` utility.

Note – Be sure to work closely with your partner during the lab to ensure you are both working on the same steps.

1. On both systems, log in as the `root` user, and open a terminal window. Using the `ifconfig -a` command, display basic configuration information about your network interfaces.

For your primary interface (for example `ng0`), what does the `ifconfig` command report for the following attributes? Enter your values into Table 3-3.

Table 3-3 Primary Interface Values

Attribute	Value
IP address	<i>It varies according to the system in use.</i>
Ethernet address	<i>It varies according to the system in use.</i>
Interface up/down	<i>The interface should be UP.</i>

2. On both systems, open a new terminal window. In the new window, enter the `snoop` command to display the network traffic between your two systems only.

`snoop host1 host2`

3. Use the `ping` command to verify that your system can contact the network interface on your partner's system.

`ping host`

4. Observe the output from the `snoop` command. Which protocol does the `ping` command use?

ICMP

Does the `snoop` output contain requests and replies (yes or no)?

Requests: Yes Replies: Yes

5. On one system, use the ifconfig command to mark its primary interface as down and then again to display its configuration information. Warn your lab partner that the system's interface will go down.

```
# ifconfig nge0 down  
# ifconfig nge0
```

Has anything changed in the information that the ifconfig command reports?

The ifconfig command no longer lists the interface as UP.

6. On the system whose interface remains up, attempt to use the ping command to contact the system whose interface is down.

What does the ping command report?

After a time-out period, the ping command reports no answer from host.

7. Observe the output from the snoop utility on both systems. How does the snoop output differ from the ping command output before and after you marked the interface as down?

The snoop utility only shows the ping command requests—no replies.

How many requests does the ping command send by default?

Twenty

Does the target system see the ping command requests? If so, how are these requests handled?

Yes it does, but it does not send a reply.

8. On the system whose interface is down, use the ifconfig command to mark its primary interface as up. Check that the change took place.

```
# ifconfig nge0 up  
# ifconfig nge0
```

9. On the system whose interface remained up, use the ping command to contact the other system.

What does the ping command report?

The host is alive.

Does the snoop utility report a reply from the target host?

Yes.

Module 4

Managing the Solaris OS UFS

Objectives

Upon completion of this module, you should be able to:

- Identify disk-based, distributed, and pseudo file systems in the Solaris OS
- Describe the Solaris OS Unix File System (UFS)
- Create a new UFS file system
- Check a file system using the `fsck` command
- Resolve file system inconsistencies
- Monitor file system use

Introducing Solaris OS File Systems

A file system is a collection of files and directories that make up a structured set of information. The Solaris OS supports three different types of file systems:

- Disk-based file systems
- Distributed file systems
- Pseudo file systems

Disk-Based File Systems

Disk-based file systems are found on hard disks, CD-ROMs, diskettes, and DVDs. The following are examples of disk-based file systems:

- ufs – The UNIX file system in the Solaris OS is based on the Berkeley fast file system. Enhancements in the Solaris 10 OS allow the ufs to grow to multiple terabytes in size.
- hsfs – The High Sierra file system is a special-purpose file system developed for use on CD-ROM media.
- pcfs – The PC file system is a UNIX implementation of the disk operating system (DOS) file allocation table (FAT32) file system. The pcfs file system allows the Solaris OS to access PC-DOS formatted file systems. Users can use UNIX commands for direct read and write access to PC-DOS files.
- udfs – The Universal Disk Format file system is used for optical storage targeted at DVD and CD-ROM media. The UDF file system allows universal data exchange and supports read and write operations.
- ZFS - A file system created by Sun that features simplified administration, pooled storage, checksums and self-healing data, snapshots, clones, and unparalleled scalability.

Distributed File Systems

Distributed file systems provide network access to file system resources.

- NFS – The network file system allows users to share files among many types of systems on the network. The NFS file system makes part of a file system on one system appear as though it were part of the local directory tree.

Pseudo File Systems

Pseudo file systems are memory based. These file systems provide for better system performance, in addition to providing access to kernel information and facilities. Pseudo file systems include:

- `tmpfs` – The temporary file system stores files in memory, which avoids the overhead of writing to a disk-based file system. The `tmpfs` file system is created and destroyed every time the system is rebooted.
- `swapfs` – The swap file system is used by the kernel to manage swap space on disks.
- `fdfs` – The file descriptor file system provides explicit names for opening files by using file descriptors (for example, `/dev/fd/0`, `/dev/fd/1`, `/dev/fd/2`) in the `/dev/fd` directory.
- `procfs` – The process file system contains a list of active processes in the `/proc` directory. The processes are listed by process number. Information in this directory is used by commands, such as the `ps` command.
- `lofs` - The loopback file system provides access to existing files using alternate pathnames.
- `mntfs` – The mount file system provides read-only information from the kernel about locally mounted file systems.
- `objfs` – The kernel object file system. This file system is used by the kernel to store details relating to the modules currently loaded by the kernel. The object file system is used for the `/system/object` directory.



- `devfs` – The device file system is used to manage the namespace of all devices on the system. This file system is used for the `/devices` directory.
- `ctfs` – The contract file system is associated with the `/system/contract` directory. This is used by the Service Management Facility to track the processes which compose a service, so that a failure in a part of a multi-process service can be identified as a failure of that service.

Note – The `mount` command with the `-p` or `-v` option displays the file system types currently in use.

Creating a New ufs File System

This section describes the ufs file system in the Solaris OS.

Viewing the Solaris OS ufs File System

The user views the ufs file system differently than the operating system does in the Solaris OS. To a user, a file system appears as a collection of files and directories used to store and organize data for access by the system and its users. To the operating system, a file system is a collection of control structures and data blocks that occupy the space defined by a partition, which allow for data storage and management.

The Solaris OS stores data in a logical file hierarchy often consisting of several file systems. This file hierarchy is referred to as the Solaris directory hierarchy.

Creating a New ufs File System

Figure 4-1 shows the Solaris OS directory hierarchy beginning with the / (root) directory.

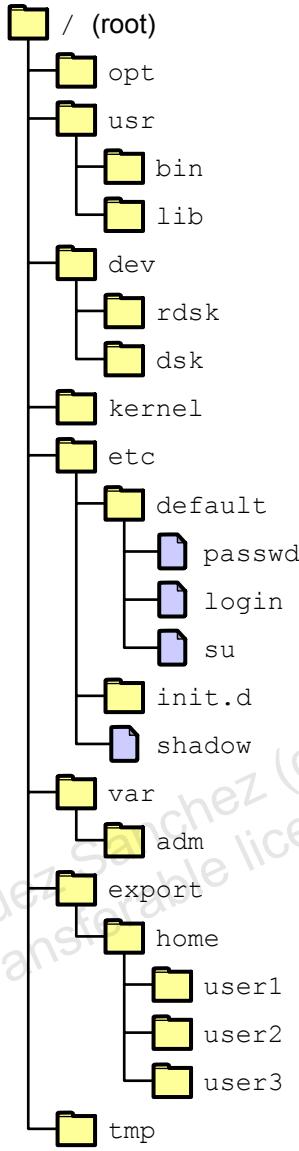


Figure 4-1 Solaris OS Directory Hierarchy

Note – Figure 4-1 is not a complete representation of a Solaris OS directory hierarchy.



A ufs file system must be created on disk before it can be used in the Solaris OS.

You can create ufs file systems on different types of storage resources, including disk slices found on a single disk, and slices on virtual storage objects, such as volumes created by software volume managers.

This description focuses on using ufs file systems on disk slices. Solaris systems typically use a number of ufs file systems to store the directories and files that make up the Solaris OS. The Solaris directory hierarchy may consist of data found on a number of separate ufs file systems.

Figure 4-2 shows ufs file systems that are located on different disk slices.

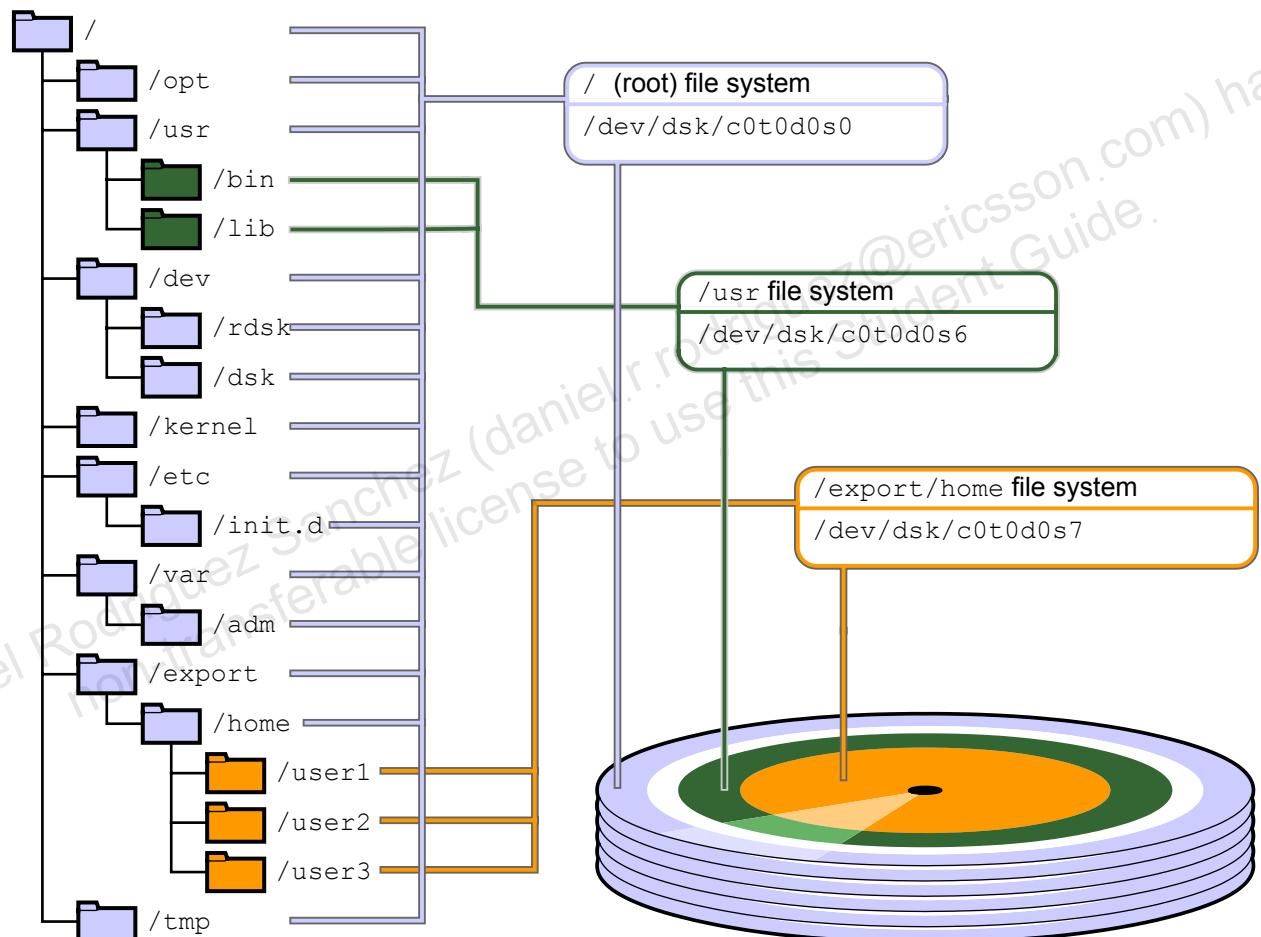


Figure 4-2 Solaris ufs File Systems Residing on Disk Slices

Partition Tables and Reserved Space

In order to create a ufs file system, slices on the disk must be defined in a partition table within a disk label. The space used to hold a ufs file system is provided by a disk slice.

- Disk labels record the locations and boundaries of disk slices.
- The label that contains these slice definitions is also known as the volume table of contents, or VTOC.
- On SPARC systems, the VTOC is located in sector 0 of the disk.
- On x86/x64 systems, the VTOC is located in the second sector of the Solaris fdisk partition. On x86/x64 systems, slice 8 occupies the first cylinder of the Solaris fdisk partition, and protects the x86/x64 VTOC.

The VTOC itself is not a file system structure, but the slices it defines are necessary for a ufs file system to be created on a disk.

Solaris OS ufs file systems do not use the following sectors of the slices in which they reside:

- Sector 0 - This sector is reserved for the VTOC on SPARC systems. If a ufs file system on a SPARC system occupies a slice that starts on the first disk cylinder, this sector would contain the SPARC VTOC. Every ufs file system on SPARC and x86/x64 systems skips sector 0.
- Sectors 1 through 15 - These sectors are reserved for the bootstrap program (bootblk) used by SPARC systems. Only a slice that contains a file system from which a SPARC system boots would contain a bootstrap program in these sectors. Every ufs file system on SPARC and x86/x64 systems skips sectors 1 through 15.

Solaris OS ufs file system structures begin at sector 16 of the slice that holds the file system. This is true of ufs file systems on SPARC and x86/x64 systems.

ufs File System Structures

Solaris ufs file systems use the structures described in this section.

Primary Superblock

The primary superblock resides in the 16 disk sectors (Sectors 16–31 relative to the first sector in the slice) that follow the space reserved for a SPARC boot block. The superblock is a table of information that describes the file system, including:

- The number of data blocks
- The number of cylinder groups
- The size of a data block and fragment
- A description of the hardware, derived from the label
- The name of the mount point
- File system state flag: clean, stable, active, logging, or unknown

Backup Superblocks

When the file system is created, a backup copy of the superblock is created beginning at sector 32, and an additional backup is created at the beginning of each cylinder group. This replication protects the critical data in the superblock against catastrophic loss.

Cylinder Groups

Each file system is divided into cylinder groups with a minimum default size of 16 cylinders per group. Cylinder groups improve disk access.

The file system constantly optimizes disk performance by attempting to place a file's data into a single cylinder group, which reduces the distance a head has to travel to access the file's data. The file system stores large files across several cylinder groups, if needed.

Cylinder Group Blocks

The cylinder group block is a table in each cylinder group that describes the cylinder group, including:

- The number of inodes
- The number of data blocks in the cylinder group
- The number of directories
- Free blocks, free inodes, and free fragments in the cylinder group
- The free block map
- The used inode map

Figure 4-3 shows a series of cylinder groups in a ufs file system. The sectors reserved for the SPARC disk label (VTOC) and SPARC bootblk program are skipped when you create a ufs file system on both SPARC and x86/64 systems.

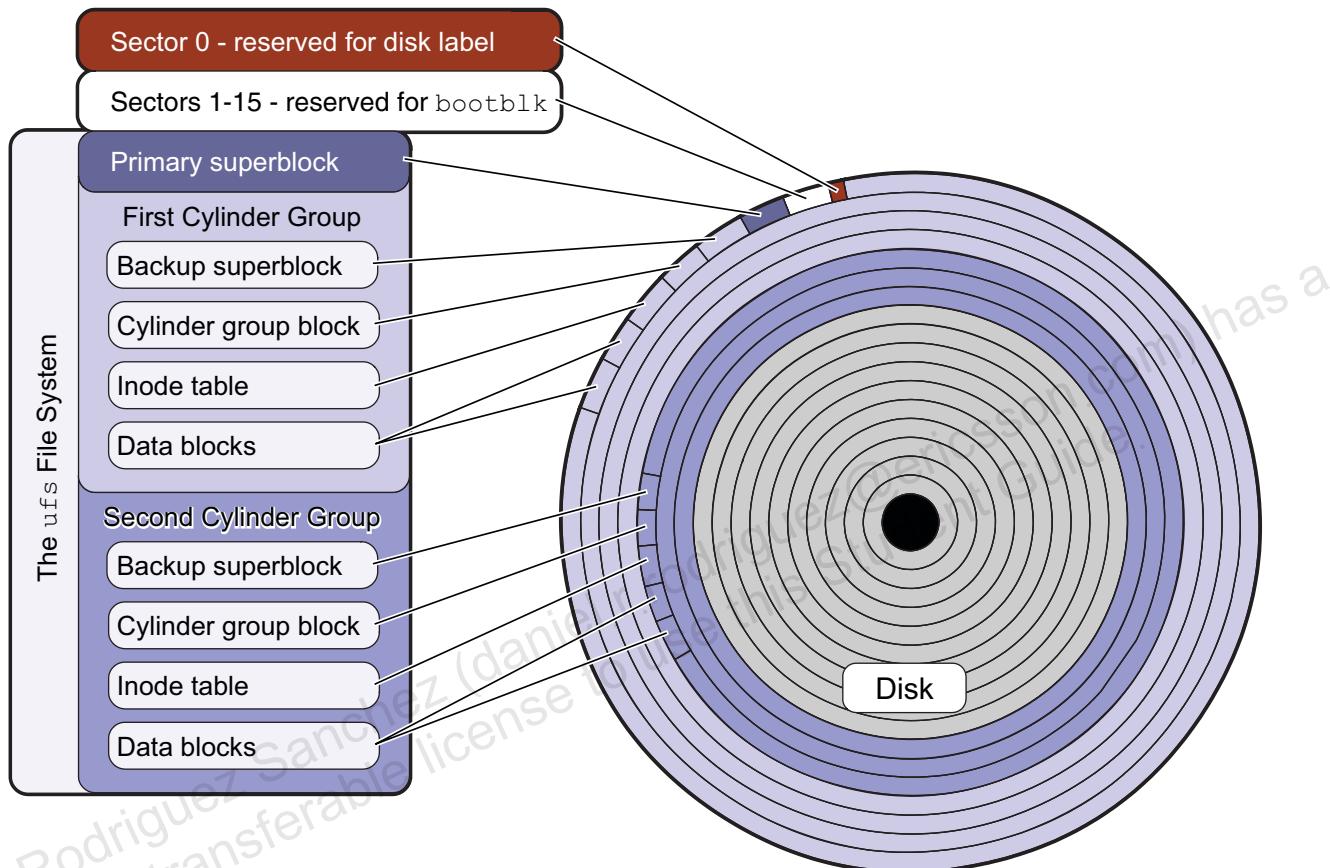


Figure 4-3 Solaris OS ufs File System Structure

The ufs Inode

An inode contains the following information about a file:

- The type of file and the access modes
- The user identification (UID) and group identification (GID) numbers of the file's owner and group
- The size of the file
- The link count
- The time the file was last accessed and modified and the inode changed
- The total number of data blocks used by or allocated to the file
- Two types of pointers: direct pointers and indirect pointers

Creating a New ufs File System

Figure 4-4 shows some of the information contained in an inode.

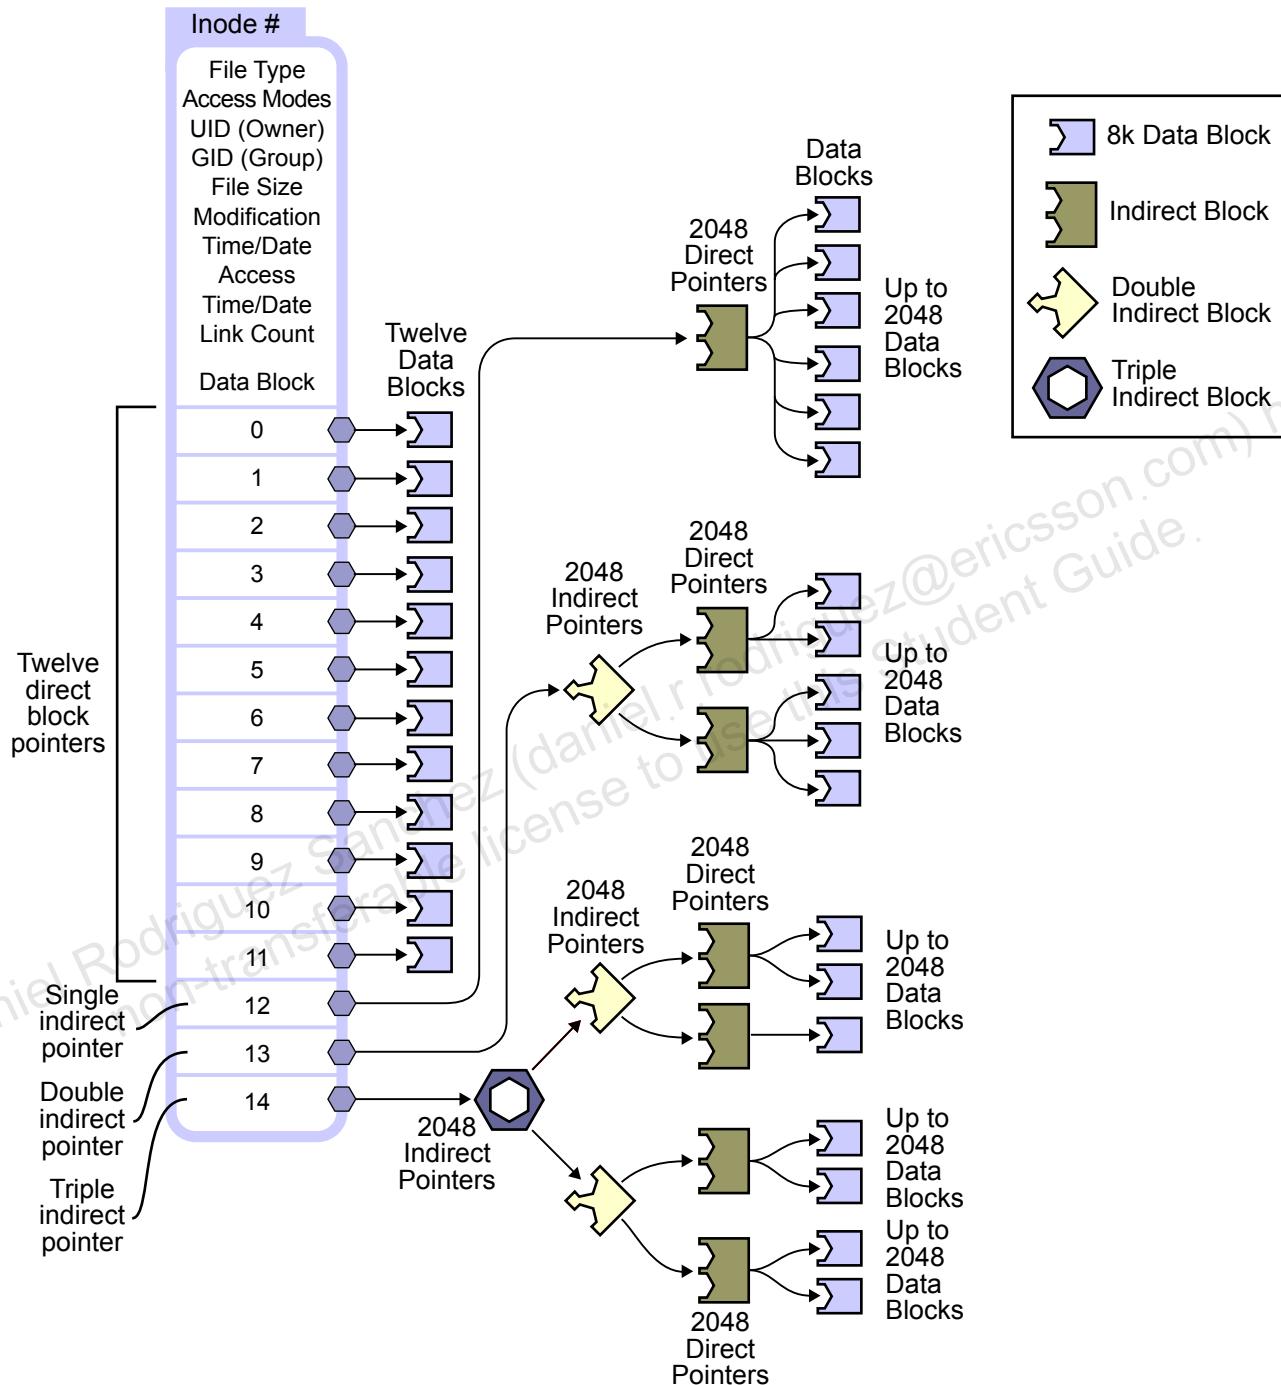


Figure 4-4 Structure of a ufs Inode



Note – To view some of the information contained in a file or directory inode, use the `ls -l` command. To view the inode number assigned to the file or directory, use the `ls -i` command.

Direct Pointers

Inside the inode there are 12 direct pointers, which contain addresses for the file's first 12 data blocks. The 12 direct pointers can each reference 8-Kbyte data blocks for a file that is up to 96 Kbytes.

Indirect Pointers

The three types of indirect pointers within an inode are:

- Single indirect pointer – Refers to a file system block that contains pointers to data blocks. This file system block contains 2048 additional addresses of 8-Kbyte data blocks, which can point to an additional 16 Mbytes of data.
- Double indirect pointer – Refers to a file system block that contains single indirect pointers. Each indirect pointer refers to a file system block that contains the data block pointers. Double indirect pointers point to an additional 32 Gbytes of data.
- Triple indirect pointer – Can reference up to an additional 64 Tbytes of data.

Data Blocks

The remaining space allocated to the ufs file system holds data blocks. Data blocks are allocated, by default, in 8-Kbyte logical block sizes. The blocks are further divided into 1-Kbyte fragments. For a regular file, the data blocks contain the contents of the file. For a directory, the data blocks contain entries that associate the inode numbers and the file names of the files and directories contained in that directory.

Within a file system, those blocks that are currently not being used as files, directories, indirect address blocks, or storage blocks are marked as free in the cylinder group map. This map also keeps track of fragments to prevent disk performance from degrading.

Fragmentation

Fragmentation is the method used by the ufs file system to allocate disk space efficiently. Files less than 96 Kbytes in size are stored using fragmentation.

By default, data blocks can be divided into eight fragments of 1024 bytes each. Fragments store files and pieces of files smaller than 8192 bytes. For files larger than 96 Kbytes, fragments are never allocated and full blocks are exclusively used.

Figure 4-5 shows a fragment in a data block.

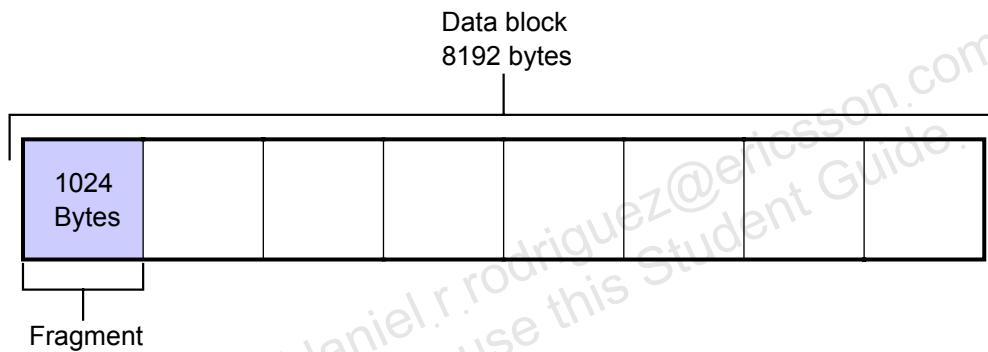


Figure 4-5 Divided Data Block

If a file contained in a fragment grows and requires more space, it is allocated one or more additional fragments in the same data block.

Figure 4-6 shows the contents of two different files stored in fragments in the same data block.

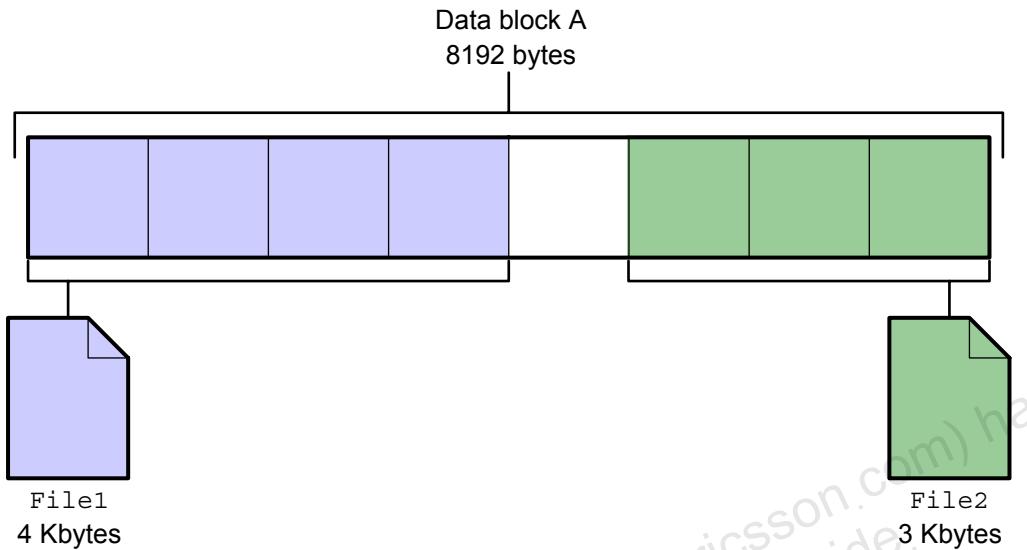


Figure 4-6 Two Files Stored in One Data Block

For example, if File1 requires more space than is currently available in the shared data block, the entire contents of that expanding file are moved by the ufs file system into a free data block. This requirement by the ufs file system assures that all of a file's fragments are contained in a whole data block. The ufs file system does not allow fragments of the same file to be stored in two different data blocks.

Using the newfs Command

To use the disk to store directories or files, a file system must be created on every disk partition. As the `root` user, you can construct a ufs file system on a disk slice by using the `newfs` command.

The `newfs` command is an easy-to-use front-end to the `mkfs` command, which you use to create file systems. The `newfs` command is located in the `/usr/sbin` directory.



Caution – Creating a new file system is destructive. The `mkfs` and `newfs` commands re-create the file system structure (new superblock, cylinder group blocks, and inode tables). Access to all existing data on the filesystem will be lost.

Creating a New ufs File System

To create a ufs file system, by using the newfs command, perform the following steps:

1. As the root user, create a file system on a slice of a newly partitioned disk by entering the command:
newfs /dev/rdsk/c2d0s0
2. The newfs command asks for confirmation before continuing. Verify that the correct disk slice on the correct disk is selected. To proceed, type y, to terminate the process, type n.

```
newfs: construct a new file system /dev/rdsk/c2d0s0: (y/n)? y
```

The newfs command displays information about the new file system being created.

```
Warning: 5008 sector(s) in last cylinder unallocated  
/dev/rdsk/c2d0s0:      39857264 sectors in 6488 cylinders of 48 tracks,  
128 sectors  
          19461.6MB in 406 cyl groups (16 c/g, 48.00MB/g, 5824 i/g)  
super-block backups (for fsck -F ufs -o b=#) at:  
    32, 98464, 196896, 295328, 393760, 492192, 590624, 689056, 787488,  
885920,  
Initializing cylinder groups:  
.....  
super-block backups for last 10 cylinder groups at:  
 38929952, 39028384, 39126816, 39225248, 39323680, 39422112, 39520544,  
39618976, 39717408, 39815840  
#
```

Output printed by the newfs command describes the basic disk geometry and the ufs file system created in this slice. The beginning sector locations of backup superblocks also displays.

The newfs -i command is used to specify the density of the number of bytes per inode in the file system. To create more inodes, a specify a smaller number of bytes per inode.

Note – This process also creates a lost+found directory for the ufs file system, which is a directory that is used by the file system check and repair (fsck) utility.

3. Repeat Steps 1 and 2 for every slice on any newly partitioned disk that needs to contain a file system.



The newfs command reserves between 1 and 10 percent of the file system space, depending on the size of the file system, for maintenance. This free space, referred to as minfree, specifies the amount of space on the slice that is reserved or held back from regular users. You can use the newfs -m %free command to preset the percentage of free space when you create a new file system.

To show the value of minfree for a file system, use the fstyp command with the -v option.

```
# fstyp -v /dev/dsk/c2d0s0 | grep minfree
minfree 1%      maxbpg 2048      optim   time
# newfs -m 2 /dev/dsk/c2d0s0
newfs: construct a new file system /dev/rdsck/c2d0s0: (y/n)? y
(output omitted)
# fstyp -v /dev/dsk/c2d0s0 | grep minfree
minfree 2%      maxbpg 2048      optim   time
#
```

The following command shows the minfree value for the file system on the c0t0d0s0 device.

```
# fstyp -v /dev/rdsck/c0t0d0s0 | head
ufs
magic 11954    format dynamic time     Fri Oct 22 10:09:11 2004
sblkno 16       cblkno 24      iblkno 32      dblkno 456
sbsize 5120    cgsizs 5120    cgoffset 72      cgmask 0xffffffff
ncg 110        size 3147511 blocks 3099093
bsize 8192     shift 13      mask 0xfffffe000
fsize 1024     shift 10      mask 0xfffffc00
frag 8         shift 3       fsbtodb 1
minfree 2%   maxbpg 2048    optim   time
maxcontig 128  rotdelay 0ms   rps     120
```

To change the minimum percentage value of free space on an existing file system, you can use the tunefs -m %free command.

The following command changes the minimum percentage of free space on the /dev/rdsck/c0t0d0s0 device to 1 percent.

```
# tuneefs -m 1 /dev/rdsck/c0t0d0s0
minimum percentage of free space changes from 2% to 1%
#
```

Checking the File System Using the `fsck` Command

A file system can become damaged if it is corrupted from a power failure, a software error in the kernel, a hardware failure, or an improper shutdown of the system. The file system check program, `fsck`, checks the data consistency of a file system and attempts to correct or repair any inconsistencies or damage found.



Caution – Never run the `fsck` command on a mounted file system. This could leave the file system in an unusable state. It could also delete data. The `/` (root), `/usr`, and `/var` file systems should have the `fsck` command run on them while in single-user mode.

Every time you boot a system, the operating system determines which file systems the `fsck` command should check. The `fsck` command checks and repairs any problems encountered in file systems before they are mounted.



Note – The status of a file system's state flag determines whether the file system needs to be scanned by the `fsck` command. When the state flag is "clean," "stable," or "logging," file system scans are not run.

Data Inconsistencies Checked by the `fsck` Command

The `fsck` command makes several passes through a file system. During each pass, the `fsck` command checks for several types of file system inconsistencies.

Superblock Consistency

The file system superblock is checked for inconsistencies involving such parameters as file system size, free block count, and free inode count.

Cylinder Group Block Consistency

The `fsck` command checks any unallocated data blocks claimed by inodes, the unallocated data block count, and the unallocated inode count.

Inode Consistency

The `fsck` command checks for the allocation state of inodes, as well as the type, the link count, duplicate blocks (blocks already claimed by another inode), bad blocks, the inode size, and the block count for each inode. Any unreferenced inode with a nonzero link count is linked to the file system's `lost+found` directory.

Data Block Consistency

The `fsck` command cannot check ordinary data blocks, but it can check directory data blocks. In directory data blocks, the `fsck` command checks for inodes that point to unallocated blocks, unallocated blocks tagged as in use, allocated blocks tagged as free (incorrect inodes for `.` and `..`) and directories not connected to the file system. These directories are linked back to the file system in its `lost+found` directory.

The `lost+found` Directory

The `fsck` command puts files and directories that are allocated but unreferenced in the `lost+found` directory located in that file system. The inode number of each file is assigned as the file name. If the `lost+found` directory does not exist, the `fsck` command creates it. If not enough space exists in the `lost+found` directory, the `fsck` command increases the directory's size.

Non-Interactive Mode

During a normal system boot, the `fsck` command operates in noninteractive mode, which is often referred to as preen, or silent mode. In this mode, the `fsck` command addresses only minor inconsistency problems that can be corrected. If a more serious inconsistency is found and a decision has to be made, the `fsck` program terminates and requests the root password to enter single-user mode. Execute the `fsck` command in interactive mode to continue.

Interactive Mode

In interactive mode, the `fsck` command lists each problem it encounters, followed by a suggested corrective action in the form of a question that requires a yes or no response.

The following example shows how the `fsck` command displays a message that asks if you want to correct the block count.

```
# fsck /dev/rdsck/c0t0d0s7
** /dev/rdsck/c0t0d0s7
** Last Mounted on /export/home
** Phase 1 - Check Blocks and Sizes
INCORRECT BLOCK COUNT I=743 (5 should be 2)
CORRECT?
```

If you respond with yes, the `fsck` command applies the corrective action and moves on. If you respond with no, the `fsck` command repeats the message about the original problem and suggests corrective action. It does not fix the inconsistency until you respond yes.

The following examples demonstrate how you as the system's root user can run the `fsck` command to check the integrity of file systems.

- To check a single unmounted file system, perform the command:

```
# fsck /dev/rdsck/c0t0d0s7
```

This is the only way to check a file system that has not been entered in the `/etc/vfstab` file.

- To check a file system using the mount point directory name as listed in the `/etc/vfstab` file, perform the command:

```
# fsck /export/home
```

In the following example, the `fsck` command checks and repairs the file system with the force (`f`) and preen (`p`) options.

```
# fsck -o f,p /dev/rdsck/c0t0d0s7
/dev/rdsck/c0t0d0s7: 77 files, 9621 used, 46089 free
/dev/rdsck/c0t0d0s7: (4 frags, 57 blocks, 0.0% fragmentation)
```

The `f` option of the `fsck` command forces a file system check, regardless of the state of the file system's superblock state flag.

The `p` option checks and fixes the file system noninteractively (preen). The program exits immediately if a problem requiring intervention is found.

Resolving File System Inconsistencies

If problems are located in a file system, you are alerted by the fsck utility. Some of the more common file system errors that require interactive intervention are:

- Allocated unreferenced file
- Inconsistent link count
- Free block count corruption
- Superblock corruption

Reconnecting an Allocated Unreferenced File

If the fsck command discovers an inode that is allocated but unreferenced or not linked in any directory, the command sends a message that asks you if you want to reconnect the inode.

```
** Phase 3a - Check Connectivity
UNREF FILE I=788 OWNER=root MODE=100644
SIZE=19994 MTIME=Oct 18 10:49 2004
RECONNECT? y
```

A yes response causes the fsck command to save the file to the lost+found directory. The fsck command references the inode number.

To determine the type of file moved to the lost+found directory by the fsck command, perform the following steps:

1. List the contents of the file system's lost+found directory.

```
# ls /export/home/lost+found
#788
```

2. Determine the file type by using the file command.

```
# file /export/home/lost+found/#788
/export/home/lost+found/#788: ascii text
```

3. To view the contents of an ASCII text file, use the more or cat command. To view the contents of a binary file, use the strings command. If the file is associated with an application, such as a word processing document, use the application to view the contents of the file.

```
# cat /export/home/lost+found/#788
```

4. If the file is intact and you know where it belongs, you can copy the file back to its original location in the file system.

```
# cp /export/home/lost+found/#788 /export/home/user1/report
```

Adjusting a Link Counter

If the fsck program discovers that the value of a directory inode link counter and the actual number of directory links are inconsistent, the command displays a message that asks you if you want to adjust the counter.

```
** Phase 4 - Check Reference Counts
LINK COUNT DIR I=2 OWNER=root MODE=40755
SIZE=512 MTIME=Oct 18 15:59 2004 COUNT 4 SHOULD BE 3
ADJUST? y
```

In the example, a **y** (yes) response causes the fsck command to correct the directory inode link counter from 4 to 3.

During this phase, you might also be asked to clear or remove a link.

```
BAD/DUP type I=200 OWNER=root MODE=40755
SIZE=512 MTIME=Mar 14 08:03 2004
CLEAR? y
```

Salvaging the Free List

If the fsck utility discovers that the unallocated block count and the free block number listed in the superblock are inconsistent, the fsck command displays a message that asks if you want to salvage the free block count by rectifying it with the unallocated block count.

```
** Phase 5 - Check Cyl groups
CG 0: BAD MAGIC NUMBER
FREE BLK COUNT(S) WRONG IN SUPERBLK
SALVAGE? y
```

In the example, a **y** (yes) response causes the fsck command to update the information in the file system superblock.

Using Backup Superblocks

Superblock corruption can cause a file system to be unmountable. A file system is unusable when the message such as "Can't mount *file_system_name*" or "*device_name* is not this fstype" appears. For example:

```
Can't mount /dev/dsk/c0t0d0s7
```

```
mount: /dev/dsk/c2d0s0 is not this fstype
```

This message can appear during a system boot or when you are manually mounting the file system.

If the fsck command fails because of a corrupted superblock, you see a series of messages that suggest searching for alternate superblocks using different methods. For example:

```
# fsck /dev/rdsk/c2d0s3
** /dev/rdsk/c2d0s3
BAD SUPERBLOCK AT BLOCK 16: MAGIC NUMBER WRONG

LOOK FOR ALTERNATE SUPERBLOCKS WITH MKFS? n

LOOK FOR ALTERNATE SUPERBLOCKS WITH NEWFS? n

SEARCH FOR ALTERNATE SUPERBLOCKS FAILED.

USE GENERIC SUPERBLOCK FROM MKFS? n

USE GENERIC SUPERBLOCK FROM NEWFS? n

SEARCH FOR ALTERNATE SUPERBLOCKS FAILED. YOU MUST USE THE -o b OPTION
TO FSCK TO SPECIFY THE LOCATION OF A VALID ALTERNATE SUPERBLOCK TO
SUPPLY NEEDED INFORMATION; SEE fsck(1M).
#
```

In this example the n responses are meant to illustrate where fsck would look for backup superblocks.

You can try any of these methods to locate a backup superblock. For example:

```
# fsck /dev/rdsk/c2d0s3
** /dev/rdsk/c2d0s3
BAD SUPERBLOCK AT BLOCK 16: MAGIC NUMBER WRONG

LOOK FOR ALTERNATE SUPERBLOCKS WITH MKFS? y

FOUND ALTERNATE SUPERBLOCK 32 WITH MKFS

USE ALTERNATE SUPERBLOCK? y

FOUND ALTERNATE SUPERBLOCK AT 32 USING MKFS
If filesystem was created with manually-specified geometry, using
auto-discovered superblock may result in irrecoverable damage to
filesystem and user data.

CANCEL FILESYSTEM CHECK? y

Please verify that the indicated block contains a proper
superblock for the filesystem (see fsdb(1M)).
```

#

Note the warning issued by fsck. If you manually specified the file system geometry when you created it, this automatic discovery process may not provide a valid superblock, and could damage the file system if used. Use these automatic discover methods only on file systems that were created using default parameters.

To use a backup superblock that you manually select, use the fsck command with the -o option and with the b flag followed by a backup superblock number. Every file system has an alternative backup superblock at block number 32, which can be used with the fsck command to repair the primary superblock.

The following command uses a backup superblock.

```
# fsck -o b=32 /dev/rdsck/c2d0s0
Alternate super block location: 32.
** /dev/rdsck/c2d0s0
** Last Mounted on
** Phase 1 - Check Blocks and Sizes
** Phase 2 - Check Pathnames
** Phase 3a - Check Connectivity
** Phase 3b - Verify Shadows/ACLs
** Phase 4 - Check Reference Counts
** Phase 5 - Check Cylinder Groups

UPDATE STANDARD SUPERBLOCK? y

2 files, 9 used, 19626536 free (8 frags, 2453316 blocks, 0.0%
fragmentation)

***** FILE SYSTEM WAS MODIFIED *****
#
```

The `fsck` utility compares the information in the backup superblock with the actual file system and attempts to rebuild the primary superblock. However, if the first backup superblock is part of the file system that was damaged, it might be unusable. If it is unusable, select another backup superblock to continue the `fsck` command.

To list the locations of all the alternative backup superblocks in the file system, run the `newfs -N` command.



Caution – This method works if the underlying file system was built using the `newfs` default parameters. If the file system was not built with these defaults, execute the `newfs -N` command, using the same parameters originally used, to generate identical superblock locations.

Resolving File System Inconsistencies

Use the **-N** option to view the file system parameters that you could use to create a new file system without actually creating the file system. A portion of the output lists the locations of alternative backup superblocks that can be used with the `fsck -o b=# command`.

```
# newfs -N /dev/rdsk/c2d0s0
Warning: 5008 sector(s) in last cylinder unallocated
/dev/rdsk/c2d0s0:      39857264 sectors in 6488 cylinders of 48 tracks,
128 sectors
      19461.6MB in 406 cyl groups (16 c/g, 48.00MB/g, 5824 i/g)
super-block backups (for fsck -F ufs -o b=#) at:
  32, 98464, 196896, 295328, 393760, 492192, 590624, 689056, 787488,
885920,
Initializing cylinder groups:
.....
super-block backups for last 10 cylinder groups at:
  38929952, 39028384, 39126816, 39225248, 39323680, 39422112, 39520544,
  39618976, 39717408, 39815840
#
```

If you use the `newfs` command with the **-T** option, the structure created allows the file system to grow to a multi-terabyte file system. You can view the file system parameters using this option without actually creating the file system.

 **Note** – Using the `newfs -T` option on smaller file systems that will contain many small files can cause a shortage of inodes.

```
# newfs -N -T /dev/rdsk/c2d0s0
Warning: 14176 sector(s) in last cylinder unallocated
/dev/rdsk/c2d0s0:      39857264 sectors in 2637 cylinders of 240 tracks,
63 sectors
      19461.6MB in 48 cyl groups (56 c/g, 413.44MB/g, 448 i/g)
super-block backups (for fsck -F ufs -o b=#) at:
  32, 846816, 1693600, 2540384, 3387168, 4233952, 5080736, 5927520,
6774304,
  7621088,
  32177824, 33024608, 33871392, 34718176, 35564960, 36411744, 37258528,
  38105312, 38952096, 39798880
#
```

You can use any other alternative superblock number in the list with the fsck command.

```
# fsck -o b=590624 /dev/rdsk/c2d0s0
Alternate super block location: 590624.
** /dev/rdsk/c2d0s0
** Last Mounted on
** Phase 1 - Check Blocks and Sizes
** Phase 2 - Check Pathnames
** Phase 3a - Check Connectivity
** Phase 3b - Verify Shadows/ACLs
** Phase 4 - Check Reference Counts
** Phase 5 - Check Cylinder Groups

UPDATE STANDARD SUPERBLOCK? y

2 files, 9 used, 19626536 free (8 frags, 2453316 blocks, 0.0%
fragmentation)

***** FILE SYSTEM WAS MODIFIED *****
#
```

Monitoring File System Use

An important activity of a system administrator is to monitor file system use on a regular basis. There are three useful commands available for this task:

- `df` – Displays the number of free disk blocks
- `du` – Summarizes disk use
- `quot` – Summarizes file system ownership

Using the `df` Command

Use the `df` command to display the amount of disk space used in file systems. This command lists the amount of used and available space and the amount of the file system's total capacity being used.

The format for the `df` command is:

```
df -option resource
```

Table 4-1 lists some of the more common options used with the `df` command.

Table 4-1 Partial Listing of Options for the `df` Command

Option	Description
<code>-a</code>	Reports on all file systems, including those with entries in the <code>/etc/mnttab</code> file for which the ignore option is set
<code>-b</code>	Prints the total number of Kbytes free
<code>-e</code>	Prints only the number of files free
<code>-k</code>	Displays disk allocation in Kbytes
<code>-h</code>	Acts like the <code>-k</code> option, except that sizes are in a more readable format, for example, 14K, 234M, 2.7G, or 3.0T
<code>-l</code>	Reports on local file systems only
<code>-F FSType</code>	Specifies the file system type on which to operate. This is intended for use on unmounted file systems.

To display the capacity of file systems, perform the command:

```
# df -k
Filesystem      kbytes   used   avail capacity  Mounted on
/dev/dsk/c0t0d0s0    5042262 156789 4835051      4%       /
/devices          0        0     0        0%       /devices
ctfs             0        0     0        0%       /system/contract
proc             0        0     0        0%       /proc
mnttab           0        0     0        0%       /etc/mnttab
swap             1149584  1224 1148360      1%       /etc/svc/volatile
objfs            0        0     0        0%       /system/object
/dev/dsk/c0t0d0s6    10085284 5132912 4851520      52%      /usr
/platform/sun4u-us3/lib/libc_psr/libc_psr_hwcap1.so.1
                    5042262 156789 4835051      4%       /platform/sun4u-
us3/lib/libc_psr.so.1
/platform/sun4u-us3/lib/sparcv9/libc_psr/libc_psr_hwcap1.so.1
                    5042262 156789 4835051      4%       /platform/sun4u-
us3/lib/sparcv9/libc_psr.so.1
fd                0        0     0        0%       /dev/fd
/dev/dsk/c0t0d0s3    493487  81718 362421      19%      /var
swap             1148392   32 1148360      1%       /tmp
swap             1148416   56 1148360      1%       /var/run
/dev/dsk/c0t0d0s5    481551  5562 427834      2%       /opt
/dev/dsk/c0t0d0s7    98771928 65750 97718459      1%      /export/home
```

The same file system displayed with the **-h** option would appear in human-readable format.

```
# df -h
Filesystem      size   used   avail capacity  Mounted on
/dev/dsk/c0t0d0s0    4.8G  153M  4.6G      4%       /
/devices          0K    0K    0K      0%       /devices
ctfs             0K    0K    0K      0%       /system/contract
proc             0K    0K    0K      0%       /proc
mnttab           0K    0K    0K      0%       /etc/mnttab
swap             1.1G  1.2M  1.1G      1%       /etc/svc/volatile
objfs            0K    0K    0K      0%       /system/object
/dev/dsk/c0t0d0s6    9.6G  4.9G  4.6G      52%      /usr
/platform/sun4u-us3/lib/libc_psr/libc_psr_hwcap1.so.1
                    4.8G  153M  4.6G      4%       /platform/sun4u-
us3/lib/libc_psr.so.1
/platform/sun4u-us3/lib/sparcv9/libc_psr/libc_psr_hwcap1.so.1
                    4.8G  153M  4.6G      4%       /platform/sun4u-
us3/lib/sparcv9/libc_psr.so.1
fd                0K    0K    0K      0%       /dev/fd
/dev/dsk/c0t0d0s3    482M  80M  354M      19%      /var
swap             1.1G  32K  1.1G      1%       /tmp
swap             1.1G  56K  1.1G      1%       /var/run
/dev/dsk/c0t0d0s5    470M  5.4M  418M      2%       /opt
/dev/dsk/c0t0d0s7    94G   64M  93G      1%      /export/home
```

Table 4-2 defines the fields displayed by the `df -k` command.

Table 4-2 Fields for the `df -k` Command

Field	Definition
Filesystem	The mounted file system
kbytes	The size of the file system in Kbytes (1024 bytes)
used	The number of Kbytes used
avail	The number of Kbytes available
capacity	The percentage of file system capacity used
Mounted on	The mount point

The amount of space that is reported as used and `avail` is typically less than the amount of total space in the file system. A fraction of space, from 1 to 10 percent, is reserved in each file system as the `minfree` value.

When all of the reported space on the file system is in use, the file system capacity is displayed as 100 percent. Regular users receive the message “File System Full” and cannot continue working. The reserved space is available to the root user, who can then delete or back up files in the file system.

The `df -k` command can be used with the device as the resource to show available space on the device:

```
# df -k /dev/dsk/c0t1d0s6
Filesystem          kbytes   used   avail capacity  Mounted on
/dev/dsk/c0t1d0s6     17153338      9  16810225      0%#
#
```

Note – This command does not work on a partition without a file system, but does work on a partition with an unmounted file system.



Using the du Command

Use the du command to display the number of disk blocks used by directories and files. Each disk block consists of 512 bytes. The format for the du command is:

du -options directory

Table 4-3 describes the options for the du command.

Table 4-3 Options for the du Command

Option	Description
-k	Displays disk use in Kbytes.
-s	Displays only the summary in 512-byte blocks. Using the s and k options together shows the summary in Kbytes.
-a	Displays the number of blocks used by all files in addition to directories within the specified directory hierarchy.

To display disk usage in kilobytes, perform the command:

```
# cd /opt
# du -k
3          ./SUNWits/Graphics-sw/xil/lib
4          ./SUNWits/Graphics-sw/xil
5          ./SUNWits/Graphics-sw
6          ./SUNWits
15         ./SUNWmlib/lib/sparcv8
15         ./SUNWmlib/lib/sparcv8plus
15         ./SUNWmlib/lib/sparcv8plus+vis
15         ./SUNWmlib/lib/sparcv8plus+vis2
15         ./SUNWmlib/lib/sparcv9
15         ./SUNWmlib/lib/sparcv9+vis
15         ./SUNWmlib/lib/sparcv9+vis2
120        ./SUNWmlib/lib
24          ./SUNWmlib/include
146        ./SUNWmlib
376        ./SUNWrvtvc/bin
10          ./SUNWrvtvc/examples/rtvc_capture_movie
24          ./SUNWrvtvc/examples/rtvc_display
68          ./SUNWrvtvc/examples/rtvc_video_conference
25          ./SUNWrvtvc/examples/test
128        ./SUNWrvtvc/examples
7           ./SUNWrvtvc/man/man1
```

Monitoring File System Use

```
19      ./SUNWrtvc/man/man3
28      ./SUNWrtvc/man
533     ./SUNWrtvc
686     .
```

To display disk use in human readable form, perform the command:

```
# du -h /opt |more
3K   /opt/SUNWits/Graphics-sw/xil/lib
4K   /opt/SUNWits/Graphics-sw/xil
5K   /opt/SUNWits/Graphics-sw
6K   /opt/SUNWits
24K  /opt/SUNWmlib/include
15K  /opt/SUNWmlib/lib/sparcv8
15K  /opt/SUNWmlib/lib/sparcv8plus
15K  /opt/SUNWmlib/lib/sparcv8plus+vis
15K  /opt/SUNWmlib/lib/sparcv8plus+vis2
15K  /opt/SUNWmlib/lib/sparcv9
15K  /opt/SUNWmlib/lib/sparcv9+vis
15K  /opt/SUNWmlib/lib/sparcv9+vis2
120K /opt/SUNWmlib/lib
146K /opt/SUNWmlib
376K /opt/SUNWrtvc/bin
10K  /opt/SUNWrtvc/examples/rtvc_capture_movie
24K  /opt/SUNWrtvc/examples/rtvc_display
68K  /opt/SUNWrtvc/examples/rtvc_video_conference
25K  /opt/SUNWrtvc/examples/test
128K /opt/SUNWrtvc/examples
7K   /opt/SUNWrtvc/man/man1
19K  /opt/SUNWrtvc/man/man3
```

To display disk use including files, perform the command:

```
# du -ak /opt
1     /opt/SUNWits/Graphics-sw/xil/lib/libxil.so
1     /opt/SUNWits/Graphics-sw/xil/lib/libxil.so.1
3     /opt/SUNWits/Graphics-sw/xil/lib
4     /opt/SUNWits/Graphics-sw/xil
(output removed for brevity)
19    /opt/SUNWrtvc/man/man3
1     /opt/SUNWrtvc/man/windex
28    /opt/SUNWrtvc/man
533   /opt/SUNWrtvc
686   /opt
```

To display a disk use summary, perform the command:

```
# du -sk /opt
686      /opt
```

Using the quot Command

Use the quot command to display how much disk space, in kilobytes, is being used by users. The format for the quot command is:

```
quot -options filesystem
```

Table 4-4 describes the options for the quot command.

Table 4-4 Options for the quot Command

Option	Description
-a	Reports on all mounted file systems
-f	Includes the number of files

To display disk space being used by users on all mounted file systems, perform the command:

```
# quot -af
/dev/rdsck/c0t0d0s0 (/):
112410  5246  root
      31     12  uucp
      11     11  lp
        1      1  adm
/dev/rdsck/c0t0d0s6 (/usr):
2313692 102415  root
    806     15  uucp
    11      4  bin
      1      1  adm
...
...
```

The columns represent kilobytes used, number of files, and owner, respectively.

Monitoring File System Use

To display a count of the number of files and space owned by each user for a specific file system, enter the following:

```
# quot -f /dev/dsk/c0t0d0s7
/dev/rdsk/c0t0d0s7 (/export/home) :
      9      2    root
     15     35    sue
     51     51   paul
     23     25   jeff
```

Exercise: Creating and Maintaining ufs File Systems

In this exercise, you complete the following tasks:

- Create ufs file systems
- Calculate and adjust minfree values
- Destroy the superblock on an unused file system and repair it using an alternative superblock

Preparation

This exercise requires an unused disk, divided into four slices. Slices 0, 1, and 3 must be equal in size, roughly one quarter of the entire disk each, and Slice 4 takes up the remaining space on the disk. If it is necessary to partition this disk, this exercise requires an understanding of how to use the `format` utility.

This exercise applies to both SPARC and x86/x64 systems. Keep in mind that these different systems typically use different disk device names. For example:

- The spare disk of a Sun Blade 1500 (SPARC) is typically `c0t1d0`.
- The spare disk of an Ultra 20 (x86/x64) is typically `c2d0`.

Throughout this exercise, be certain to use the correct device names for your particular system.

Task

Complete the following steps:

1. Log in as the root user and complete the following steps:
 - a. Open a terminal window.
 - b. Change directory to `/dev/rdsk`.
2. Complete the following steps to identify a spare disk:
 - a. To find a spare disk, use the `ls` command to display a list of possible disks and the `prtvtoc` command to display the VTOC for each disk you find.

Exercise: Creating and Maintaining ufs File Systems

- b. Examine the Partition list, as well as the Mount Directory field that the prtvtoc command displays.

On systems in the classroom, disks that are not in use have no mount directory listed.

- c. Record the name of the unused disk.

Unused disk:

Note – This procedure works for the classroom environment. A disk that does not show mounted slices in the Mount Directory field of the prtvtoc output is not necessarily unused.

3. If the spare disk is not divided into four slices as described in the Preparation section of this exercise, complete the following steps:
- On x86/x64 systems, use the fdisk menu in the format utility to create a single Solaris fdisk partition that uses the entire disk. Make sure the partition is not marked as *active*.
 - Use the partition menu in the format utility to create the required slices.
 - Make slices 0, 1, and 3 exactly the same size (approximately 25 percent of the total disk space each), and use slice 4 for the remainder of the available space. Using the All Free Hog method and choosing slice 4 as the free hog slice makes this easy.
On SPARC systems, you can also use the Solaris Management Console to partition the disk.
 - Save your changes and quit the format utility when you are finished.



Example of a correct SPARC partition table:

Part	Tag	Flag	Cylinders	Size	Blocks
0	root	wm	0 - 13878	27.00GB	(13879/0/0) 56626320
1	swap	wu	13879 - 27757	27.00GB	(13879/0/0) 56626320
2	backup	wu	0 - 57458	111.79GB	(57459/0/0) 234432720
3	unassigned	wm	27758 - 41636	27.00GB	(13879/0/0) 56626320
4	unassigned	wm	41637 - 57458	30.78GB	(15822/0/0) 64553760
5	unassigned	wm	0	0	(0/0/0) 0
6	usr	wm	0	0	(0/0/0) 0
7	unassigned	wm	0	0	(0/0/0) 0

Example of a correct x86/x64 partition table:

Part	Tag	Flag	Cylinders	Size	Blocks
0	root	wm	3 - 2483	19.01GB	(2481/0/0) 39857265
1	swap	wu	2484 - 4964	19.01GB	(2481/0/0) 39857265
2	backup	wu	0 - 10007	76.67GB	(10008/0/0) 160778520
3	unassigned	wm	4965 - 7445	19.01GB	(2481/0/0) 39857265
4	unassigned	wm	7446 - 10007	19.63GB	(2562/0/0) 41158530
5	unassigned	wm	0	0	(0/0/0) 0
6	usr	wm	0	0	(0/0/0) 0
7	unassigned	wm	0	0	(0/0/0) 0
8	boot	wu	0 - 0	7.84MB	(1/0/0) 16065
9	alternates	wm	1 - 2	15.69MB	(2/0/0) 32130

4. Use the newfs command without options to create a new file system on Slice 0 on the spare disk and complete the following steps:
- Observe how quickly the newfs command creates cylinder groups on this slice.
 - Record the number of cylinder groups, the number of cylinders per group, and the number of inodes per group.

Cylinder groups: _____

Cylinders per group: _____

Inodes per group: _____

5. Use the newfs command to create a new file system on slice 1 on the spare disk, and complete the following steps:
- Use the -i option to create one inode per 16,384 bytes of data space.
 - Observe how quickly the newfs command creates cylinder groups on this slice.

Exercise: Creating and Maintaining ufs File Systems

- c. Record the number of cylinder groups, the number of cylinders per group, and the number of inodes per group.

Cylinder groups: _____

Cylinders per group: _____

Inodes per group: _____

According to the statistics you have gathered, how do the file systems on slices 0 and 1 differ?

6. Complete the following steps:

- Use the `df` command to display statistics for the file systems on slices 0 and 1 that you used in the previous steps, for example:
- Record the values listed in the kbytes, used, and avail columns.

Which file system has the larger amount of available data space and why?

7. Complete the following steps:

- For each file system, add the used and avail values, and compare the sum to the kbytes value.

Expressed as a percentage, how much larger is the kbytes value than the sum of used and avail? This percentage should approximately match the minfree value.

- Use the `fstyp -v /dev/rdsck/c#t#d#s# | grep minfree` command to verify your result.

8. Use the `tunefs -m X /dev/rdsck/c#t#d#s0` command, where `X` is the minfree value, to change the minfree value for the file system on slice 0 of the spare disk.

- If the current minfree value is greater than 5 percent, reduce it by 3 percent.
- If it is less than or equal to 5 percent, add 3 percent.

What message does the `tunefs` command display?

9. Complete the following steps:
 - a. Use the `df -k` command to display the space utilization for the file system on slice 0 of your spare disk. Observe what that the `minfree` value has changed in the output.
 - b. Record the values listed in the `kbytes`, `used`, and `avail` columns.

Which value has changed from the information you gathered in Step 6?

10. Create new file systems on slices 3 and 4 of your spare disk.
11. Run the `fsck` command interactively to check the new file system that you created on slice 3 of the spare disk.
Did the `fsck` command report errors?

12. Use the `dd` command to destroy the main superblock of the file system on slice 3. For example:
`# dd if=/dev/zero of=/dev/rdsck/c0t1d0s3 count=32 bs=512`
13. Run the `fsck` command interactively to check the new file system. Respond `n` to the prompts presented by `fsck`.
Did the `fsck` command report errors? If so, what corrective action does the `fsck` command suggest?

14. Run the `fsck` command, and specify an alternative superblock to use on slice 3.
Block 32 is always one of the alternatives available.
15. Run the `fsck` command again to verify that the file system was repaired.

Exercise Summary



Discussion – Take a few minutes to discuss what experiences, issues, or discoveries you had during the lab exercises.

- Experiences
- Interpretations
- Conclusions
- Applications

Exercise Solutions: Creating and Maintaining ufs File Systems

This section contains solutions to the exercises.

Task

Review the following solutions:

1. Log in as the root user and complete the following steps:
 - a. Open a terminal window.
 - b. Change directory to /dev/rdsks.
`# cd /dev/rdsks`
 2. Complete the following steps to identify a spare disk:
 - a. To find a spare disk, use the ls command to display a list of possible disks and the prtvtoc command to display the VTOC for each disk you find.
 - b. Examine the Partition list, as well as the Mount Directory field that the prtvtoc command displays.
On systems in the classroom, disks that are not in use have no mount directory listed.
 - c. Record the name of the unused disk.
`# ls *s2`
`# prtvtoc /dev/rdsks/c0t1d0s0`
- Unused disk:



Note – This procedure works for the classroom environment. A disk that does not show mounted slices in the Mount Directory field of the prtvtoc output is not necessarily unused.

3. If the spare disk is not divided into four slices as described in the Preparation section of this exercise, complete the following steps:
 - a. On x86/x64 systems, use the fdisk menu in the format utility to create a single Solaris fdisk partition that uses the entire disk. Make sure the partition is not marked as *active*.

- b. Use the partition menu in the format utility to create the required slices.
- c. Make slices 0, 1, and 3 exactly the same size (approximately 25 percent of the total disk space each), and use slice 4 for the remainder of the available space. Using the All Free Hog method and choosing slice 4 as the free hog slice makes this easy.

On SPARC systems, you can also use the Solaris Management Console to partition the disk.

- d. Save your changes and quit the format utility when you are finished.

Example of a correct SPARC partition table:

Part	Tag	Flag	Cylinders	Size	Blocks
0	root	wm	0 - 13878	27.00GB	(13879/0/0) 56626320
1	swap	wu	13879 - 27757	27.00GB	(13879/0/0) 56626320
2	backup	wu	0 - 57458	111.79GB	(57459/0/0) 234432720
3	unassigned	wm	27758 - 41636	27.00GB	(13879/0/0) 56626320
4	unassigned	wm	41637 - 57458	30.78GB	(15822/0/0) 64553760
5	unassigned	wm	0	0	(0/0/0) 0
6	usr	wm	0	0	(0/0/0) 0
7	unassigned	wm	0	0	(0/0/0) 0

Example of a correct x86/x64 partition table:

Part	Tag	Flag	Cylinders	Size	Blocks
0	root	wm	3 - 2483	19.01GB	(2481/0/0) 39857265
1	swap	wu	2484 - 4964	19.01GB	(2481/0/0) 39857265
2	backup	wu	0 - 10007	76.67GB	(10008/0/0) 160778520
3	unassigned	wm	4965 - 7445	19.01GB	(2481/0/0) 39857265
4	unassigned	wm	7446 - 10007	19.63GB	(2562/0/0) 41158530
5	unassigned	wm	0	0	(0/0/0) 0
6	usr	wm	0	0	(0/0/0) 0
7	unassigned	wm	0	0	(0/0/0) 0
8	boot	wu	0 - 0	7.84MB	(1/0/0) 16065
9	alternates	wm	1 - 2	15.69MB	(2/0/0) 32130

4. Use the newfs command without options to create a new file system on Slice 0 on the spare disk and complete the following steps:

- Observe how quickly the newfs command creates cylinder groups on this slice.
- Record the number of cylinder groups, the number of cylinders per group, and the number of inodes per group.

```
# newfs /dev/rdsck/c0t1d0s0
```

Cylinder groups:

Cylinders per group:

Inodes per group:

5. Use the newfs command to create a new file system on slice 1 on the spare disk, and complete the following steps:

- Use the -i option to create one inode per 16,384 bytes of data space.
- Observe how quickly the newfs command creates cylinder groups on this slice.
- Record the number of cylinder groups, the number of cylinders per group, and the number of inodes per group.

```
# newfs -i 16384 /dev/rdsck/c0t1d0s1
```

Cylinder groups:

Cylinders per group:

Inodes per group:

According to the statistics you have gathered, how do the file systems on slices 0 and 1 differ?

The number of inodes per group is less on file system 1 than on file system 0. The number of cylinder groups and cylinders per group have not changed.

6. Complete the following steps:

- Use the **df** command to display statistics for the file systems on slices 0 and 1 that you used in the previous steps, for example:

```
# df -k /dev/dsk/c0t1d0s0
Filesystem          kbytes   used   avail capacity  Mounted on
/dev/dsk/c0t1d0s0    27883846      9  27604953      0%
# df -k /dev/dsk/c0t1d0s1
Filesystem          kbytes   used   avail capacity  Mounted on
/dev/dsk/c0t1d0s1    28091566      9  27810576      0%
```

- Record the values listed in the kbytes, used, and avail columns.

Which file system has the larger amount of available data space and why?

File system 1 has the larger amount of available data space because it holds fewer inode records.

7. Complete the following steps:

- For each file system, add the used and avail values, and compare the sum to the kbytes value.

Expressed as a percentage, how much larger is the kbytes value than the sum of used and avail? This percentage should approximately match the minfree value.

- Use the **fstyp -v /dev/rdsks/c#t#d#s# | grep minfree** command to verify your result.

To calculate the percentage difference between the sum of used and avail and the kbytes value, complete the following steps:

- Add the values listed as used and avail, for example:

$$9 + 27604953 = 27604962$$

- Divide the sum of used and avail by the kbytes value, for example:

$$27604962 / 27883846 = 0.989998$$

- Multiply the result of Step b by 100, for example:

$$0.989998 * 100 = 98.9998$$

- Subtract the result of Step c from 100, for example:

$$100 - 98.9998 = 1.0002$$

- Round the result of Step d to the nearest whole number, for example:

$$1.0002 = 1\text{percent}$$

8. Use the `tunefs -m X /dev/rdsk/c#t#d#s0` command, where *X* is the minfree value, to change the minfree value for the file system on slice 0 of the spare disk.

- If the current minfree value is greater than 5 percent, reduce it by 3 percent.
- If it is less than or equal to 5 percent, add 3 percent, for example:

```
# tunefs -m 4 /dev/rdsk/c0t1d0s0
```

What message does the `tunefs` command display?

minimum percentage of free space changes from (*old value*) to (*new value*)

9. Complete the following steps:

- a. Use the `df -k` command to display the space utilization for the file system on slice 0 of your spare disk. Observe that the minfree value has changed in the output.
- b. Record the values listed in the kbytes, used, and avail columns, for example:

```
# df -k /dev/dsk/c0t1d0s0
Filesystem      kbytes   used   avail capacity  Mounted on
/dev/dsk/c0t1d0s0    27883846      9  26768439      0%
```

Which value has changed from the information you gathered in Step 6?

The avail column changes, but not the kbytes or used columns.

10. Create new file systems on slices 3 and 4 of your spare disk, for example:

```
# newfs /dev/rdsk/c0t1d0s3
# newfs /dev/rdsk/c0t1d0s4
```

11. Run the `fsck` command interactively to check the new file system that you created on slice 3 of the spare disk.

```
# fsck /dev/rdsk/c0t1d0s3
```

Did the `fsck` command report errors?

No.

12. Use the `dd` command to destroy the main superblock of the file system on slice 3.

```
# dd if=/dev/zero of=/dev/rdsk/c0t1d0s3 count=32 bs=512
```

13. Run the fsck command interactively to check the new file system.
Respond n to the prompts presented by fsck.

```
# fsck /dev/rdsk/c0t1d0s3
```

Did the fsck command report errors? If so, what corrective action does the fsck command suggest?

The fsck command indicates that the magic number in the superblock is wrong and suggests looking for alternate superblocks in different ways, and repairing it by using one of them, for example:

```
# fsck /dev/rdsk/c0t1d0s3
```

```
** /dev/rdsk/c0t1d0s3
```

```
BAD SUPERBLOCK AT BLOCK 16: MAGIC NUMBER WRONG
```

```
LOOK FOR ALTERNATE SUPERBLOCKS WITH MKFS? n
```

```
LOOK FOR ALTERNATE SUPERBLOCKS WITH NEWFS? n
```

```
SEARCH FOR ALTERNATE SUPERBLOCKS FAILED.
```

```
USE GENERIC SUPERBLOCK FROM MKFS? n
```

```
USE GENERIC SUPERBLOCK FROM NEWFS? n
```

```
SEARCH FOR ALTERNATE SUPERBLOCKS FAILED. YOU MUST USE THE -o b OPTION  
TO FSCK TO SPECIFY THE LOCATION OF A VALID ALTERNATE SUPERBLOCK TO  
SUPPLY NEEDED INFORMATION; SEE fsck(1M).
```

```
#
```

14. Run the fsck command, and specify an alternative superblock to use on slice 3.

Block 32 is always one of the alternatives available.

```
# fsck -o b=32 /dev/rdsck/c0t1d0s3
Alternate super block location: 32.
** /dev/rdsck/c0t1d0s3
** Last Mounted on
** Phase 1 - Check Blocks and Sizes
** Phase 2 - Check Pathnames
** Phase 3a - Check Connectivity
** Phase 3b - Verify Shadows/ACLs
** Phase 4 - Check Reference Counts
** Phase 5 - Check Cylinder Groups
```

UPDATE STANDARD SUPERBLOCK? **y**

2 files, 9 used, 27883837 free (13 frags, 3485478 blocks, 0.0%
fragmentation)

```
***** FILE SYSTEM WAS MODIFIED *****
#
```

15. Run the fsck command again to verify that the file system was repaired.

```
# fsck /dev/rdsck/c0t1d0s3
```

*This time the fsck command output does not report that the file system
was modified.*

Notes:

Module 5

ZFS File System Introduction

Objectives

Upon completion of this module, you should be able to:

- Describe the Solaris ZFS file system
- Understand ZFS terminology
- Create new ZFS pools and file systems
- Modify ZFS file system properties
- Destroy ZFS pools and file systems
- Work with ZFS snapshots and clones

What Is ZFS?

ZFS is a file system that:

- Integrates filesystem and volume management to simplify administration
- Supports high storage capacities and avoids arbitrary limits
- Supports unlimited snapshots and clones
- Checksums all data
- Is open source, and available in Solaris 10, OpenSolaris, MacOS X and FreeBSD.

ZFS Pooled Storage

ZFS uses storage pools to manage physical storage. Historically, file systems were constructed on top of a single physical device. To address multiple devices and provide for data redundancy, the concept of a volume manager was introduced to provide the image of a single device so that file systems would not have to be modified to take advantage of multiple devices.

ZFS eliminates the volume management altogether. Instead of forcing you to create virtualized volumes, ZFS aggregates devices into a storage pool. The storage pool describes the physical characteristics of the storage (device layout, data redundancy, and so on) and acts as an arbitrary data store from which file systems can be created.

File systems are no longer constrained to individual devices, allowing them to share space with all file systems in the pool. You don't need to predetermine the file system size, as file systems grow automatically within the space allocated to the storage pool.

Transactional Semantics

ZFS is a transactional file system, which means that the file system state is always consistent on disk. With a transactional file system, data is managed using copy on write semantics.

Data is never overwritten, and any sequence of operations is either entirely committed or entirely ignored. This mechanism means that the file system can never be corrupted through accidental loss of power or a system crash. So, no need for a fsck equivalent exists. While the most recently written pieces of data might be lost, the file system itself will always be consistent.

Checksums and Self-Healing Data

With ZFS, all data and metadata is checksummed using a user-selectable algorithm. ZFS checksums are stored in a way such that these failure modes are detected and can be recovered from gracefully. All checksumming and data recovery is done at the file system layer, and is transparent to applications.

In addition, ZFS provides for self-healing data. ZFS supports storage pools with varying levels of data redundancy, including mirroring and a variation on RAID-5. When a bad data block is detected, ZFS fetches the correct data from another replicated copy, and repairs the bad data, replacing it with the good copy.

Scalability

The ZFS file system was designed to be the most scalable file system, ever. The file system is 128-bit, allowing for 256 quadrillion zettabytes of storage. All metadata is allocated dynamically, so you don't need to pre-allocate inodes or otherwise limit file system scalability when you first create it. All the algorithms were written with scalability in mind. Directories can have up to 2^{48} (256 trillion) entries, and no limit exists on the number of file systems or number of files that can be contained within a file system.

Snapshots

A snapshot is a read-only copy of a file system or volume. You can create snapshots. Initially, snapshots consume no additional space within the pool.

Administration

The ZFS file system provides simplified administration. Through the use of hierarchical file system layout, property inheritance, and automanagement of mount points and NFS share semantics, ZFS makes it easy to create and manage file systems without needing multiple commands or editing configuration files. You can set quotas or reservations, turn compression on or off, or manage mount points for numerous file systems with a single command.

Devices can be examined or repaired without having to understand a separate set of volume manager commands. You can take an unlimited number of instantaneous snapshots of file systems. You can backup and restore individual file systems.

ZFS Terminology

Term	Definition
checksum	A 256-bit data hash in a file-system block. A checksum entity can range from the simple and fast fletcher2 (the default) to cryptographically strong hashes like SHA256.
clone	A file system whose initial contents are identical to the contents of a snapshot.
dataset	A generic name for the following ZFS entities: clones, file systems, snapshots, or volumes. Each dataset is identified by a unique name in the ZFS namespace. Identify datasets with the following format: <i>pool/path[@snapshot]</i>
file system	A ZFS dataset of type filesystem that is mounted within the standard system namespace and behaves like other file systems.
mirror	A virtual device that stores identical data copies on two or more disks. If any disk in a mirror fails, any other disk in that mirror can provide the same data.
pool	Identifies the storage pool name that contains the dataset. A logical group of devices describing the layout and physical characteristics of the available storage. Dataset space is allocated from a pool.
RAID-Z	A virtual device that stores data and parity on multiple disks, similar to RAID-5. For more information about RAID-Z, see “RAID-Z Storage Pool Configuration” on page 5-12.
resilvering	The process of transferring data from one device to another device. For example, if a mirror component is replaced or taken offline, the data from the up-to-date mirror component is copied to the newly restored mirror component. This process is called <i>mirror resynchronization</i> in traditional volume management products.
snapshot	A read-only image of a file system or volume at a given point in time.
virtual device	A logical device in a pool, which can be a physical device, a file, or a collection of devices.
volume	A dataset used to emulate a physical device. For example, you can create a ZFS volume as a swap device.

Creating a Basic ZFS File System

The ZFS file system reduces the commands you need to create a usable file system. When you create a new pool, a new ZFS file system is created and mounted automatically.

The following example illustrates how to create a storage pool named tank and a ZFS name tank in one command. Assume that the whole disk /dev/dsk/c1t0d0 is available for use.

```
# zpool create tank c1t0d0
```

The new ZFS file system, tank, can use as much of the disk space on c1t0d0 as needed, and is automatically mounted at /tank.

```
# mkfile 100m /tank/foo
# df -h /tank
Filesystem           size   used  avail capacity
Mounted on
tank                  80G   200M   80G      1%
/tank
```

Within a pool, you may want to create additional file systems. File systems provide points of administration that allow you to manage different sets of data within the same pool. In most cases, you will want to create and organize a hierarchy of file systems that matches your organizational needs

ZFS Storage Pool Components

This section provides detailed information about the following storage pool components:

- Disks
- Files
- Virtual devices
- Cache devices

Using Disks in a ZFS Storage Pool

The most basic element of a storage pool is a piece of physical storage. Physical storage can be any block device of at least 128 Mbytes. Typically, this device is a hard drive in the /dev/dsk directory.

A storage device can be a whole disk (c1t0d0) or an individual slice (c0t0d0s7). Sun recommends that you use an entire disk so that you don't have to specially format the disk. The disk ZFS formats the disk using an EFI label to contain a single large slice.

ZFS applies an EFI label when you create a storage pool with whole disks. Disks can be labeled with a traditional Solaris VTOC label when you create a storage pool with a disk slice.

Use a slice only under the following conditions:

- The device name is non-standard.
- A single disk is shared between ZFS and another file system, such as the UFS.
- A disk is used as a swap or a dump device.

Specify a disk by using a full path like /dev/dsk/c1t0d0 or a shorthand name that consists of the device name within the /dev/dsk directory, such as c1t0d0. For example, the following are valid disk names:

- c1t0d0
- /dev/dsk/c1t0d0
- c0t0d6s2
- /dev/foo/disk

Using whole physical disks is the simplest way to create ZFS storage pools. ZFS configurations become progressively more complex, from management, reliability, and performance perspectives, when you build pools from disk slices, LUNs in hardware RAID arrays, or volumes presented by software-based volume managers. The following considerations might help you determine how to configure ZFS with other hardware or software storage solutions:

- If you construct ZFS configurations on top of LUNs from hardware RAID arrays, you need to understand the relationship between ZFS redundancy features and the redundancy features offered by the array. Certain configurations might provide adequate redundancy and performance, but other configurations might not.
- You can construct logical devices for the ZFS using volumes presented by software-based volume managers, such as Solaris™ Volume Manager (SVM) or Veritas Volume Manager (VxVM). However, these configurations are not recommended. While the ZFS functions properly on such devices, less-than-optimal performance could result.

Using Files in a ZFS Storage Pool

The ZFS file system allows you to use UFS files as virtual devices in your storage pool. This feature enables testing and simple experimentation. It is not for production use. This is because any file use relies on the underlying file system for consistency. If you create a ZFS pool backed by files on a UFS file system, then you are relying on the UFS to guarantee correctness and synchronous semantics.

Virtual Devices in a Storage Pool

Each storage pool consists of one or more virtual devices. A virtual device is an internal representation of the storage pool that describes the layout of physical storage and its fault characteristics. As such, a virtual device represents the disk devices or files that are used to create the storage pool.

Two top-level virtual devices provide data redundancy: mirror and RAID-Z virtual devices. These virtual devices consist of disks, disk slices, or files.

Disk, disk slices, or files that are used in pools outside of mirrors and RAID-Z virtual devices, function as top-level virtual devices.

Storage pools typically contain multiple top-level virtual devices. ZFS file dynamically stripes data among all of the top-level virtual devices in a pool.

Using Cache Devices in a ZFS Storage Pool

In the Solaris 10 10/09 release, you can create a ZFS storage pool and specify cache devices, which are used to cache storage pool data. Cache devices provide an additional layer of caching between main memory and disk memory. Using cache devices provide the greatest performance improvement for random read-workloads of mostly static content. One or more cache devices can be specified when the pool is created. For example:

```
# zpool create pool mirror c0t2d0 c0t4d0 cache c0t0d0
# zpool status pool
  pool: pool
    state: ONLINE
      scrub: none requested
      config:

          NAME        STATE   READ WRITE CKSUM
          pool        ONLINE    0     0     0
            mirror    ONLINE    0     0     0
              c0t2d0  ONLINE    0     0     0
              c0t4d0  ONLINE    0     0     0
            cache
              c0t0d0  ONLINE    0     0     0
      errors: No known data errors
```

Cache devices can be added, or removed from a storage pool after the pool is created.



Note – For information about determining whether using cache devices is appropriate for your environment, see the *Solaris ZFS Administration Guide*.

ZFS Storage Pool Replication Features

Mirrored Storage Pool Configuration

A mirrored storage pool configuration requires at least two disks, preferably on separate controllers. Many disks can be used in a mirrored configuration. In addition, you can create more than one mirror in each pool. Conceptually, a simple mirrored configuration would look like this:

```
mirror c1t0d0 c2t0d0
```

Conceptually, a more complex mirrored configuration would look like this:

```
mirror c1t0d0 c2t0d0 c3t0d0 mirror c4t0d0 c5t0d0 c6t0d0
```

Each `mirror` argument defines a separate top-level virtual device. The first example defines one two-way mirror device. The second example defines two, three-way mirror devices.

RAID-Z Storage Pool Configuration

In addition to a mirrored storage pool configuration, the ZFS file system provides a RAID-Z configuration. RAID-Z is similar to RAID-5.

In RAID-Z, ZFS file uses variable-width RAID stripes so that all writes are full-stripe writes. This design is only possible because ZFS integrates file system and device management in such a way that the file system's metadata has enough information about the underlying data replication model to handle variable-width RAID stripes.

A RAID-Z configuration with N disks of size X with P parity disks can hold approximately $(N-P)*X$ bytes and can withstand one device failing before data integrity is compromised. You need at least two disks for a single-parity RAID-Z configuration and at least three disks for a double-parity RAID-Z configuration.

For example, if you have three disks in a single-parity RAID-Z configuration, parity data occupies space equal to one of the three disks. Otherwise, no special hardware is required to create a RAID-Z configuration.

Conceptually, a RAID-Z configuration with three disks looks like:

```
raidz c1t0d0 c2t0d0 c3t0d0
```

A more complex conceptual RAID-Z configuration looks like:

```
raidz c1t0d0 c2t0d0 c3t0d0 c4t0d0 c5t0d0 c6t0d0 c7t0d0  
raidz c8t0d0 c9t0d0 c10t0d0 c11t0d0 c12t0d0 c13t0d0  
c14t0d0
```

If you are creating a RAID-Z configuration with many disks, as in above example, a RAID-Z configuration with 14 disks is better split into two 7-disk groupings. RAID-Z configurations with single-digit groupings of disks should perform better.

Each `raidz` argument defines a separate top-level virtual device. The first example defines a single RAID-Z device that consists of three disks. The second example defines two RAID-Z devices, where each device consists of seven disks.

Self-Healing Data in a Replicated Configuration

ZFS file enables self-healing data in a mirrored or RAID-Z configuration. When a bad data block is detected, ZFS fetches the correct data from another replicated copy, and repairs the bad data by replacing it with the good copy.

Dynamic Striping in a Storage Pool

For each virtual device that is added to a pool, ZFS file dynamically stripes data across all available devices. The decision about where to place data is done at write time, so no fixed width stripes are created at allocation time.

When virtual devices are added to a pool, ZFS gradually allocates data to the new device in order to maintain performance and space allocation policies. Each virtual device can also be a mirror or a RAID-Z device that contains other disk devices or files. This configuration allows for flexibility in controlling the fault characteristics of your pool. For example, you could create the following configurations out of 4 disks:

- Four disks using dynamic striping
- One four-way RAID-Z configuration
- Two two-way mirrors using dynamic striping

The ZFS file system supports combining different virtual devices within the same pool, but this isn't recommended. For example, you can create a pool with a two-way mirror and a three-way RAID-Z configuration. But, your fault tolerance will be as good as your worst virtual device, RAID-Z in this case. Sun recommends that you use top-level virtual devices of the same type with the same replication level in each device.

Creating and Destroying ZFS Storage Pools

Creating and destroying pools is fast and easy. However, be cautious when doing so. Although checks are performed to prevent using devices known to be in use in a new pool, the ZFS file system cannot always know when a device is already in use. Also, use `zpool destroy` with caution. Using this command can have significant consequences.

Creating a ZFS Storage Pool

Use the `zpool create` command to create a storage pool. This command takes a pool name and any number of virtual devices as arguments. The pool name must conform to the ZFS component naming conventions as described in the *Solaris ZFS Administration Guide*.



Caution – Do not add a disk that is currently configured as a Sun Cluster quorum device to a ZFS storage pool. After the disk is added to a storage pool, then the disk can be configured as a quorum device.

Creating a Basic Storage Pool

The following command creates a new pool named `tank` that consists of the disks `c1t0d0` and `c1t1d0`:

```
# zpool create tank c1t0d0 c1t1d0
```

These whole disks are found in the `/dev/dsk` directory and are labelled appropriately by ZFS to contain a single, large slice. Data is dynamically striped across both disks.

Creating a Mirrored Storage Pool

To create a mirrored pool, use the `mirror` keyword, followed by any number of storage devices that will comprise the mirror. Multiple mirrors can be specified by repeating the `mirror` keyword on the command line. The following command creates a pool with two, two-way mirrors:

```
# zpool create tank mirror c1d0 c2d0 mirror c3d0 c4d0
```

The second `mirror` keyword indicates that a new top-level virtual device is being specified. Data is dynamically striped across both mirrors, with data being replicated between each disk appropriately.

Creating a Single-Parity RAID-Z Storage Pool

Creating a single-parity RAID-Z pool is identical to creating a mirrored pool, except that the `raidz` keyword is used instead of `mirror`. The following example shows how to create a pool with a single RAID-Z device that consists of five disks:

```
# zpool create tank raidz c1t0d0 c2t0d0 c3t0d0 c4t0d0 /dev/dsk/c5t0d0
```

This example demonstrates that disks can be specified by using their full paths. The `/dev/dsk/c5t0d0` device is identical to the `c5t0d0` device.

A similar configuration could be created with disk slices. For example:

```
# zpool create tank raidz c1t0d0s0 c2t0d0s0 c3t0d0s0 c4t0d0s0 c5t0d0s0
```

However, the disks must be preformatted to have an appropriately sized slice zero.

Destroying ZFS Storage Pools

Use the Pools `zpool destroy` command to destroy a pool. This command destroys the pool even if it contains mounted datasets.

```
# zpool destroy tank
```

Caution – Be careful when you destroy a pool. Make sure you are destroying the right pool and you have copies of your data. If you accidentally destroy the wrong pool, you can attempt to recover the pool.



Querying ZFS Storage Pool Status

Use the `zpool` command to request pool status. This section describes how to display basic use information and pool HEALTH status.

Getting Basic ZFS Storage Pool Information

Use the `zpool list` command to display basic pool information. With no arguments, the `zpool list` command displays all the fields for all pools on the system. Specify the pool name to gather statistics for that pool. For example:

```
# zpool list
NAME          SIZE    USED   AVAIL    CAP  HEALTH      ALTROOT
tank          80.0G  22.3G  47.7G  28%  ONLINE      -
dozer         1.2T   384G   816G  32%  ONLINE      -
# zpool list tank
NAME          SIZE    USED   AVAIL    CAP  HEALTH      ALTROOT
tank          80.0G  22.3G  47.7G  28%  ONLINE      -
```

This output displays the following information:

NAME	Pool name
SIZE	Pool size (equal to the sum of the size of all top-level virtual devices)
USED	Space allocated by all datasets and internal metadata
AVAILABLE	Unallocated space in the pool.
CAPACITY (CAP)	Space used (expressed as a percentage of total space)
HEALTH	Current pool health status
ALTROOT	Alternate root of the pool, if any.

Double Parity RAID-Z

With Solaris 10 11/06, a redundant RAID-Z configuration can have either single- or double-parity, which means that one or two device failures can be sustained respectively, without any data loss. Specify the `raidz2` keyword for a double-parity RAID-Z configuration. Specify the `raidz` or `raidz1` keyword for a single-parity RAID-Z configuration. For more information, see the `zpool(1M)` man page.

ZFS Storage Pool HEALTH Status

ZFS provides an integrated method of examining pool and device HEALTH status. Use the `zpool status` command to display pool HEALTH status. In addition to the `zpool status` command, the fault manager daemon (`fmd`) displays potential pool and device failures on the system console and in the `/var/adm/messages` file. A device can be in one of the following states:

ONLINE	Some transient errors may still occur, the virtual device is working normally.
DEGRADED	The virtual device experienced a failure but still functions. This state is most common when a mirror or RAID-Z device has lost one or more constituent devices. The fault tolerance of the pool might be compromised, as a subsequent fault in another device might be unrecoverable.
FAULTED	The virtual device is completely inaccessible. This status typically indicates total failure of the device, such that the ZFS file system is incapable of sending or receiving data from it. If a top-level virtual device is in this state, then the pool is completely inaccessible.
OFFLINE	The virtual device has been explicitly taken offline by the administrator.
UNAVAILABLE	The device or virtual device cannot be opened. In some cases, pools with UNAVAILABLE devices appear in DEGRADED mode. If a top-level virtual device is unavailable, then nothing in the pool can be accessed.

Determine the HEALTH of a pool from the status of all its top-level virtual devices. If all the virtual devices are ONLINE, then the pool is also ONLINE. If any of the virtual devices is DEGRADED or UNAVAILABLE, then the pool is also DEGRADED. If a top-level virtual device is FAULTED or OFFLINE, then the pool is also FAULTED. A FAULTED pool is completely inaccessible.

Getting Basic Storage Pool HEALTH Status

To request a quick overview of pool HEALTH, use the `zpool status` command:

```
# zpool status -x
all pools are healthy
```

Specify a pool name to examine a specific pool. Investigate potential problems in any pool that isn't in the ONLINE state.

Getting Detailed Storage Pool HEALTH Status

Use the `-v` option to request a detailed HEALTH summary. The following output displays a complete description of why the pool is in its current state. It includes a readable description of the problem and a link to a knowledge article for more information. Each knowledge article provides up-to-date information on the best way to recover from your current problem. Using the detailed configuration information, you should be able to determine which device is damaged and how to repair the pool.

```
# zpool status -v tank
pool: tank
state: DEGRADED
status: One or more devices could not be opened. Sufficient replicas
exist for the pool to continue functioning in a degraded state.
action: Attach the missing device and online it using 'zpool online'.
see: http://www.sun.com/msg/ZFS-8000-2Q
```

```
scrub: none requested
config:
```

NAME	STATE	READ	WRITE	CKSUM
tank	DEGRADED	0	0	0
mirror	DEGRADED	0	0	0
c1t0d0	FAULTED	0	0	0 cannot open
c1t1d0	ONLINE	0	0	0

```
errors: No known data errors
```

ZFS Log Device Recovery

In the Solaris 10 10/09 release, ZFS identifies intent-log failures in the `zpool status` command. ZFS describes how to recover from an intent-log failure.

For example, if the system shuts down abruptly before synchronous write operations are committed to a pool with a separate log device, you will see intent-log related error messages in the `zpool status` output.

Note – For information about resolving log device failures, see the *Solaris ZFS Administration Guide*.



Creating and Destroying ZFS File Systems

ZFS is a lightweight POSIX file system that is built on top of a storage pool. You can dynamically create and destroy file systems without allocating or formatting underlying space. Because file systems are lightweight and the central point of administration in ZFS, you will probably create many.

Use the `zfs` commands to administer ZFS. The `zfs` command includes subcommands that perform specific operations on file systems. Use the `zfs create` and `zfs destroy` command to create and destroy file systems.

Note – The term *dataset* is used in this module as a generic term to refer to a filesystem, snapshot, clone, or volume.



Creating a ZFS File System

Use the `zfs create` command to create a ZFS file system. The `create` subcommand takes a single argument: the name of the file system to create. The file system name is specified as a path name starting from the name of the pool:

```
pool-name/[filesystem-name/]filesystem-name
```

The pool name and initial file system names in the path identify the location in the hierarchy where the new file system will be created. All the intermediate file system names must already exist in the pool. The last name in the path identifies the name of the file system to be created.

In the following example, a file system named `bonwick` is created in the `tank/home` file system.

```
# zfs create tank/home/bonwick
```

The ZFS file system automatically mounts the newly created file system if it is created successfully. By default, file systems are mounted as `/dataset`, using the path provided for the file system name in the `create` subcommand. In this example, the newly created `bonwick` file system is at `/tank/home/bonwick`.

Destroying a ZFS File System

Use the `zfs destroy` command to destroy a ZFS file system. The destroyed file system is automatically unmounted and unshared.

In the following example, the `tabriz` file system is destroyed.

```
# zfs destroy tank/home/tabriz
```

Caution – No confirmation prompt appears with the `destroy` subcommand. Use it with extreme caution.



If the file system to be destroyed is busy and therefore can't be unmounted, the `zfs destroy` command fails. To destroy an active file system, use the `-f` option. Use this option with caution as it can unmount, unshare, and destroy active file systems, causing unexpected application behavior.

```
# zfs destroy tank/home/ahrens  
cannot unmount 'tank/home/ahrens': Device busy  
# zfs destroy -f tank/home/ahrens
```

The `zfs destroy` command also fails if a file system has children. To recursively destroy a file system and all its descendants, use the `-r` option. Note that a recursive destroy also destroys snapshots so use this option with caution.

```
# zfs destroy tank/ws  
cannot destroy 'tank/ws': filesystem has children  
use '-r' to destroy the following datasets:  
tank/ws/billm  
tank/ws/bonwick  
tank/ws/maybe  
  
# zfs destroy -r tank/ws
```

If the file system to be destroyed has indirect dependents, even the recursive destroy command described above fails. To force the destruction of all dependents, including cloned file systems outside the target hierarchy, the **-R** option must be used. Use extreme caution with this option.

```
# zfs destroy -r tank/home/schrock
cannot destroy 'tank/home/schrock': filesystem has
dependent clones
use '-R' to destroy the following datasets:
tank/clones/schrock-clone

# zfs destroy -R tank/home/schrock
```

Renaming a ZFS File System

Use `zfs rename` to rename a ZFS file system. Using the `rename` subcommand, you can perform the following operations:

- Change the name of a file system
- Relocate the file system to a new location within the ZFS hierarchy
- Change the file system name and relocate it with the ZFS hierarchy

The following example uses the `rename` subcommand to rename a file system. This example renames the `kustarz` file system to `kustarz_old`.

```
# zfs rename tank/home/kustarz tank/home/kustarz_old
```

The following example shows how to use `zfs rename` to relocate a file system. In this example, the `maybebe` file system is relocated from `tank/home` to `tank/ws`. When you relocate a file system through `rename`, the new location must be within the same pool and it must have enough space to hold this new file system. If the new location does not have enough space, possibly because it has reached its quota, the `rename` will fail.

```
# zfs rename tank/home/maybebe tank/ws/maybebe
```

The `rename` operation attempts an unmount/remount sequence for the file system and any descendant file systems. The `rename` fails if the operation is unable to unmount an active file system. If this problem occurs, you will need to force unmount the file system.

ZFS Properties

Properties are the main mechanism that you use to control the behavior of file systems, volumes, snapshots, and clones. Unless stated otherwise, the properties defined in the section apply to all the dataset types.

Properties are either read-only statistics or settable properties. Most settable properties are also inheritable. An inheritable property is a property that, when set on a parent, is propagated down to all of its descendants.

All inheritable properties have an associated source. The source indicates how a property was obtained. A property source can have the following values:

default	A default value means that the property setting was not inherited or set locally. This source is a result of no ancestor having the property as source local.
local	A local source indicates that the property was explicitly set on the dataset by using the <code>zfs set</code> command.
inherited from dataset- name	A value of inherited from <i>dataset-name</i> means that the property was inherited from the named ancestor.
temporary	A temporary value means that the property value was set by using the <code>zfs mount -o</code> option and is valid for the mount lifetime.
- (none)	Read-only property. Its value is generated by ZFS.

Read-Only ZFS Properties

You can retrieve but not set read-only properties. Read-only properties are not inherited. Some properties are specific to a particular type of dataset. In such cases, the particular dataset type is mentioned in the description in the previous table. ZFS read-only properties are:

- available
- compressratio
- creation

- mounted
- origin
- referenced
- type
- used
- usedbychildren
- usedbydataset
- usedbyreservation
- usedbysnapshots

Settable ZFS Properties

You can retrieve and set “settable” properties. Use `zfs set` to set properties. With the exception of quotas and reservations, settable properties are inherited.

Some settable properties are specific to a dataset type. In such cases, the dataset type is mentioned in the description in the previous table. If not specifically mentioned, a property applies to all dataset types: file systems, volumes, clones, and snapshots.

The following table explains settable properties. For more information see the `zfs(1M)` man page and the *Solaris ZFS System Administration Guide*.

Table 6–1 ZFS Native Property Descriptions

Property Name	Type	Default Value	Description
aclinherit	String	secure	Controls how ACL entries are inherited when files and directories are created. The values are <code>discard</code> , <code>noallow</code> , <code>secure</code> , and <code>passthrough</code> . For a description of these values, see ACL Property Modes .
aclmode	String	groupmask	Controls how an ACL entry is modified during a <code>chmod</code> operation. The values are <code>discard</code> , <code>groupmask</code> , and <code>passthrough</code> . For a description of these values, see ACL Property Modes .

Property Name	Type	Default Value	Description
atime	Boolean	on	Controls whether the access time for files is updated when they are read. Turning this property off avoids producing write traffic when reading files and can result in significant performance gains, though it might confuse mailers and other similar utilities.
canmount	Boolean	on	Controls whether the given file system can be mounted with the <code>zfs mount</code> command. This property can be set on any file system and the property isn't inheritable. When this property is set to off, a mountpoint can be inherited to descendent file systems, but the file system is never mounted.
checksum	String	on	When the <code>noauto</code> option is set, a dataset can only be mounted and unmounted explicitly. The dataset is not mounted automatically when the dataset is created or imported, nor is it mounted by the <code>zfs mount-a</code> command or unmounted by the <code>zfs umount-a</code> command.
compression	String	off	Controls the checksum used to verify data integrity. The default value is on, which automatically selects an appropriate algorithm, currently fletcher2. The values are on, off, fletcher2, fletcher4, and sha256. A value of off disables integrity checking on user data. A value of off is not recommended.

Property Name	Type	Default Value	Description
copies	Number	1	Sets the number of copies of user data per file system. Available values are 1, 2 or 3. These copies are in addition to any pool-level redundancy. Space used by multiple copies of user data is charged to the corresponding file and dataset and counts against quotas and reservations. In addition, the used property is updated when multiple copies are enabled. Consider setting this property when the file system is created because changing this property on an existing file system only affects newly written data.
devices	Boolean	on	Controls the ability to open device files in the file system.
exec	Boolean	on	Controls whether programs within this file system are allowed to be executed. Also, when set to off, mmap(2) calls with PROT_EXEC are disallowed.
groupquota	Number	none	Sets a quota on the amount of space consumed by files that are owned by a particular group.
mountpoint	String	NA	Controls the mount point used for this file system. When the mountpoint property is changed for a file system, the file system and any children that inherit the mount point are unmounted. If the new value is legacy, then they remain unmounted. Otherwise, they are automatically remounted in the new location if the property was previously legacy or none, or if they were mounted before the property was changed. In addition, any shared file systems are unshared and shared in the new location.

Property Name	Type	Default Value	Description
primarycache	String	NA	Controls the primary cache (ARC) on an existing ZFS file system or when the file system is created. The default, all caches both user data and metadata. The other values are none where neither user data or metadata is cached, or metadata where only metadata is cached.
quota	Number (or none)	none	Limits the amount of space a dataset and its descendants can consume. This property enforces a hard limit on the amount of space used, including all space consumed by descendants, including file systems and snapshots. Setting a quota on a descendent of a dataset that already has a quota does not override the ancestor's quota, but rather imposes an additional limit. Quotas cannot be set on volumes, as the volsize property acts as an implicit quota.
readonly	Boolean	off	Controls whether this dataset can be modified. When set to on, no modifications can be made to the dataset. This property can also be referred to by its shortened column name, rdonly.
recordsize	Number	128K	Specifies a suggested block size for files in the file system. This property can also be referred to by its shortened column name, recsize.
refquota	Number (or none)	none	Sets the amount of space that a dataset can consume. This property enforces a hard limit on the amount of space used. This hard limit does not include space used by descendants, such as snapshots and clones.

Property Name	Type	Default Value	Description
refreservation	Number (or none)	none	Sets the minimum amount of space that is guaranteed to a dataset, not including descendants, such as snapshots and clones. When the amount of space that is used is below this value, the dataset is treated as if it were taking up the amount of space specified by refreservation. The refreservation reservation is accounted for in the parent datasets' space used, and counts against the parent datasets' quotas and reservations.
			If refreservation is set, a snapshot is only allowed if enough free pool space is available outside of this reservation to accommodate the current number of referenced bytes in the dataset.
reservation	Number (or none)	none	This property can be referred to by its shortened column name, <code>refreserv</code> .
			The minimum amount of space guaranteed to a dataset and its descendants. When the amount of space used is below this value, the dataset is treated as if it were using the amount of space specified by its reservation. Reservations are accounted for in the parent datasets' space used, and count against the parent datasets' quotas and reservations. This property can also be referred to by its shortened column name, <code>reserv</code> .
secondarycache	String	NA	Controls the secondary cache (L2ARC) on an existing ZFS file system or when the file system is created. The default, <code>all</code> caches both user data and metadata. The other values are <code>none</code> where neither user data or metadata is cached, or <code>metadata</code> where only metadata is cached.
setuid	Boolean	on	Controls whether the setuid bit is honored in the file system.

Property Name	Type	Default Value	Description
sharenfs	String	off	Controls whether the file system is available over NFS, and what options are used. If set to on, the <code>zfs share</code> command is invoked with no options. Otherwise, the <code>zfs share</code> command is invoked with options equivalent to the contents of this property. If set to off, the file system is managed by using the legacy share and unshare commands and the <code>dfstab</code> file.
snapdir	String	hidden	Controls whether the <code>.zfs</code> directory is hidden or visible in the root of the file system.
userquota	Number	none	Sets a quota on the amount of space consumed by files that are owned by a particular user.
volsize	Number	NA	For volumes, specifies the logical size of the volume.
volblocksize	Number	8 Kbytes	For volumes, specifies the block size of the volume. The block size cannot be changed once the volume has been written, so set the block size at volume creation time. The default block size for volumes is 8 Kbytes. Any power of 2 from 512 bytes to 128 Kbytes is valid. This property can also be referred to by its shortened column name, <code>volblock</code> .
zoned	Boolean	NA	Indicates whether this dataset has been added to a non-global zone. If this property is set, then the mount point is not honored in the global zone, and ZFS cannot mount such a file system when requested. When a zone is first installed, this property is set for any added file systems.

ZFS Property Enhancements

The Solaris 10 10/09 release provides several new ZFS property enhancements for the read-only ZFS properties and settable properties.

Using Space (Usage) Accounting Properties

In the Solaris 10 10/09 release you can use the space accounting properties to identify space usage for clones, file systems, and volumes, but not snapshots. The space accounting properties are as follows:

- `usedbychildren` – Identifies the amount of space that is used by children of this dataset, which would be freed if all the dataset's children were destroyed. The property abbreviation is `usedchild`.
- `usedbydataset` – Identifies the amount of space that is used by this dataset itself, which would be freed if the dataset was destroyed, after first destroying any snapshots and removing any refreservation. The property abbreviation is `useddds`.
- `usedbyrefreservation` – Identifies the amount of space that is used by a refreservation set on this dataset, which would be freed if the refreservation was removed. The property abbreviation is `usedrefreserv`.
- `usedbysnapshots` – Identifies the amount of space that is consumed by snapshots of this dataset. In particular, it is the amount of space that would be freed if all of this dataset's snapshots were destroyed. Note that this is not simply the sum of the snapshots' used properties, because space can be shared by multiple snapshots. The property abbreviation is `usesnap`.

These new properties break down the value of the used property into the various elements that consume space. In particular, the value of the used property breaks down as follows:

`used property = usedbychildren + usedbydataset + usedbyrefreservation + usedbysnapshots`

You can view these properties by using the `zfs list -o space` command. For example:

# <code>zfs list -o space</code>	NAME	AVAIL	USED	USED SNAP	USED DDS	USED REFRESERV	USED CHILD
	pool	33.2G	7.2K	0	21K	0	51K
	rpool	27.0G	6.27G	20.5K	97K	0	6.27G
	rpool/ROOT	27.0G	4.73G	0	21K	0	4.73G

rpool/ROOT/zfsBE	27.0G	4.73G	97.5M	4.63G	0	0
rpool/dump	27.0G	1.00G	16K	1.00G	0	0
rpool/export	27.0G	60K	16K	23K	0	21K
rpool/export/home	27.0G	21K	0	21K	0	0
rpool/swap	27.5G	553M	0	41.5M	512M	0

Setting ZFS File System Properties at Pool Creation Time

You can set ZFS file system properties at pool creation time. In the following example, compression is enabled on the ZFS file system that is created when the pool is created.

```
# zpool create -o compression=on pool mirror c0t1d0 c0t2d0
```

Setting Cache Properties on a ZFS File System

There are two new ZFS file system properties that control the primary cache and secondary cache.

- primarycache
- secondarycache

The values for setting cache properties are as follows:

- all – default both user data and metadata are cached
- none – neither user data or metadata is cached
- metadata – only metadata is cached

You can set these properties on an existing ZFS file system, or when the file system is created:

```
# zfs set primarycache=metadata tank/datab
```

```
# zfs create -o primarycache=metadata tank/newdatab
```

Some database environments might benefit from not caching user data. You will have to determine if setting cache properties is appropriate for your environment.

Setting ZFS User and Group Quotas

In previous Solaris OS releases, you could apply quotas and reservations to ZFS file systems to manage and reserve space. In the Solaris 10 10/09 release, you can now set a quota on the amount of space consumed by files that are owned by a particular user or group.

You might consider setting user and group quotas in any environment that has a large number of users or groups.

You can set user or group quotas by using the `zfs userspace` and `zfs groupspace` properties as follows:

```
# zfs set userquota@user1=5G tank/data  
# zfs set groupquota@staff=10G tank/staff/admins
```

You can display a user's or group's current quota setting as follows:

```
# zfs get userquota@user1 tank/data  
NAME PROPERTY VALUE SOURCE  
tank/data userquota@user1 5G local  
  
# zfs get groupquota@staff tank/staff/admins  
NAME PROPERTY VALUE SOURCE  
tank/staff/admins groupquota@staff 10G local
```

To display general quota information, type the following:

```
# zfs userspace tank/data  
TYPE NAME USED QUOTA  
POSIX User root 3K none  
POSIX User user1 0 5G  
  
# zfs groupspace tank/staff/admins  
TYPE NAME USED QUOTA  
POSIX Group root 3K none  
POSIX Group staff 0 10G
```

Querying ZFS Information

The `zfs list` command provides an extensible mechanism for viewing and querying dataset information. Both basic and complex queries are explained in this section.

Listing Basic ZFS Information

List basic dataset information by using the `zfs list` command with no options. This command displays system dataset names including their used, available, referenced, and mountpoint properties. For example:

```
# zfs list
NAME          USED  AVAIL  REFER  MOUNTPOINT
pool           84.0K 33.5G   -      /pool
pool/clone     0     33.5G  8.50K  /pool/clone
pool/test      8K    33.5G   8K    /test
pool/home      17.5K 33.5G  9.00K  /pool/home
pool/home/marks 8.50K 33.5G  8.50K  /pool/home/marks
pool/home/marks@snap 0     -      8.50K  /pool/home/marks@snap
```

Use `zfs list` to display specific datasets by providing the dataset name on the command line. Use the `-r` option to display dataset descendants. The following example shows how to display `tank/home/chua` and its descendant datasets. For additional information about the `zfs list` command, see `zfs(1M)`.

```
# zfs list -r tank/home/chua
tank/home/chua           26.0K  4.81G  10.0K  /tank/home/chua
tank/home/chua/projects  16K    4.81G  9.0K   /tank/home/chua/projects
tank/home/chua/projects/fs1 8K    4.81G   8K
/tank/home/chua/projects/fs1
tank/home/chua/projects/fs2 8K    4.81G   8K
/tank/home/chua/projects/fs2
```

ZFS Snapshots Omitted from `zfs list` Output

In the Solaris 10 10/09 release, snapshots are omitted from the `zfs list` output. The `listsnap` pool property controls whether snapshot information is displayed by the `zfs list` command. If you use the `zfs list -t snapshots` command, snapshot information is displayed. The default value is `off`, which means snapshot information is not displayed by default.

Sending and Receiving ZFS Data

The `zfs send` command creates a stream representation of a snapshot that is written to standard output. By default, a full stream is generated. You can redirect the output to a file or to a different system. The `zfs receive` command creates a snapshot whose contents are specified in the stream that is provided on standard input. If a full stream is received, a new file system is created. You can send ZFS snapshot data and receive ZFS snapshot data and file systems with these commands. For more information on sending and receiving ZFS data, see the *Solaris ZFS System Administration Guide*.

Creating and Destroying ZFS Snapshots

Use the `zfs snapshot` command to create snapshots. Enter the snapshot name for the `zfs snapshot` command argument as follows. You must follow snapshot naming conventions. See the Solaris *ZFS System Administration Guide* to find the naming conventions.

```
filesystem@snapname
volume@snapname
```

In the following, a snapshot of `tank/home/ahrens` named `friday` is created.

```
# zfs snapshot tank/home/ahrens@friday
```

Use the `-r` option to create snapshots for all descendent file systems.

```
# zfs snapshot -r tank/home@now
# zfs list -t snapshot
NAME          USED  AVAIL  REFER  MOUNTPOINT
tank/home@now      0     -  29.5K  -
tank/home/ahrens@now  0     -   2.15M  -
tank/home/anne@now   0     -   1.89M  -
tank/home/bob@now    0     -   1.89M  -
tank/home/cindys@now  0     -   2.15M  -
```

Snapshots have no modifiable properties. Dataset properties can't be applied to a snapshot.

```
# zfs set compression=on tank/home/ahrens@tuesday
```

cannot set compression property for 'tank/home/ahrens@tuesday':
snapshotproperties cannot be modified

Use `zfs destroy` to destroy snapshots.

```
# zfs destroy tank/home/ahrens@friday
```

A dataset can't be destroyed if snapshots of the dataset exist (see example below). If clones were created from a snapshot, you must destroy them before you can destroy the snapshot.

```
# zfs destroy tank/home/ahrens
```

cannot destroy 'tank/home/ahrens': filesystem has children

use '-r' to destroy the following datasets:

`tank/home/ahrens@tuesday`

`tank/home/ahrens@wednesday`

`tank/home/ahrens@thursday`

Exercise: ZFS File System Introduction

In this exercise, you complete the following tasks:

- Create and destroy ZFS storage pools
- Create, rename, and destroy ZFS file systems
- Work with ZFS properties

Preparation

This exercise requires a minimum of one unused spare disk. Tasks 1 and 2 provide steps to create storage pools that consist of 1, 2, or 3 disks, if they are available. You must evaluate the list of disk devices on your system to determine the correct choices to make in tasks 1 and 2. The path names of your disk devices might differ from those described in these exercises. Use the device path names appropriate to your system.

Note – If only one spare disk is available on your system, this exercise asks you to use a single disk to create a storage pool. This is not a recommended configuration. Single-disk pools offer no data redundancy. Failure of the disk supporting the pool causes the pool and its data to become inaccessible.



Exercise Summary

This exercise includes tasks that illustrate the main ZFS features and cover the following concepts:

- Working with ZFS storage pools - creating and destroying a basic ZFS storage pool using spare disks on your system.
- Creating, renaming, and destroying ZFS file systems - creating a basic storage pool, and manipulating a ZFS file within that pool.
- Working with ZFS properties - setting and listing ZFS properties, principles of ZFS property inheritance.
- Working with ZFS snapshots and clones - creating and destroying ZFS snapshots and clones, and dependencies that exist among them.
- Removing the ZFS configuration - removing the objects you create during the exercise.

Task 1 – Creating and Destroying ZFS Storage Pools

Complete the following steps:

1. Use the `zpool` command to display the list of ZFS pools. Verify that no ZFS pool currently exists.
2. Use the `zfs` command to display the list of ZFS file systems. Verify that no ZFS file system currently exists.
3. Use the `df` command to display the list of UFS file systems currently in use on your system. Use this list to determine which disk or disks on your system are currently in use.



Note – The `df` command with the `-F` option allows you to list file systems of a specific type.

4. Use the `format` command to list what disks are attached to your system. Compare this list with the output of step 3 to determine what disks are currently not in use on your system. Use `<Ctrl-D>` to exit your `format` session.



Note – You may find that your system uses different sized disks. For this exercise, although it is not the best practice, use the `zpool` command with the `-f` option to override errors that this problem causes.



Note – When you try to create a pool, your system will probably display errors that indicate UFSs exist on your spare disk. If it does, verify that the spare disk is not in use, then use the `zpool` command with the `-f` option to override the errors and create the pool.

5. If you have three unused disks available, use the `zpool` command to create a storage pool called `firstpool` that consists of one, three-disk RAID-Z virtual device.
6. If you have only two unused disks available, use the `zpool` command to create a storage pool called `firstpool` that consists of one mirror virtual device.
7. If you have only one unused disk available, use the `zpool` command to create a storage pool called `firstpool` that uses the single disk as a non-redundant, top-level virtual device.
8. Use the `zpool` command to display the list of ZFS pools.

9. Use the `zpool` command to display the basic health status of your pool.
10. Use the `zpool` command to display the detailed health status of your pool. What kind of top-level virtual device does the `zpool` command report, and what devices does it consist of? What is the state of the top-level virtual device, and of the pool itself?
11. Use the `zfs` command to list ZFS file systems, and display the space utilization for `firstpool`. Does the space available match what you would expect for a pool of the type you created?
12. Use the `zpool` command to display the name of your pool, and only its space used and space available statistics.

Note – For a RAID-Z pool, the `zpool list` command reports the sum of all disk space in the AVAIL column. This differs from the value reported by `zfs list`. For a RAID-Z pool, the `zfs list` command reduces the value reported in the AVAIL column by the amount of parity data that the pool requires.

13. Use the `zfs` command to display the list of currently mounted ZFS file systems. Display a long listing for the `/firstpool` directory.
14. Use the `zpool` command to destroy `firstpool`. Verify that `firstpool` no longer exists.
15. Verify that the `firstpool` file system and the `/firstpool` mount point no longer exist.

Task 2 – Create, Rename, and Destroy ZFS File Systems

Complete the following steps:

1. Use the `zpool` command to display the list of ZFS pools. Verify that no ZFS pool currently exists.
2. If you have three unused disks available, use the `zpool` command to create a storage pool called `mypool` that consists of one, three-disk RAID-Z virtual device.
3. If you have only two unused disks available, use the `zpool` command to create a storage pool called `mypool` that consists of one mirror virtual device.



4. If you have only one unused disk available, use the `zpool` command to create a storage pool called `mypool` that uses the single disk as a non-redundant, top-level virtual device.



Note – Example command output throughout the remainder of this exercise assumes that a non-redundant pool is in use. This is not a recommended pool configuration.

5. Use the `zfs` command to list ZFS file systems, and display the space use for `/mypool`.
6. Attempt to create a file system named `mypool/fs1/fs2`. What happens and why?
7. Create a directory called `/mypool/fs1`. Again attempt to create a file system named `mypool/fs1/fs2`. What happens and why?

The attempt fails because the intermediate file system fs1 does not exist. ZFS only creates new file systems immediately below a parent file system.
8. Create a new empty file named `/mypool/fs1/file1`.
9. Attempt to create a new file system named `mypool/fs1`. Use the `zfs list` and `zfs mount` commands to list the ZFS file systems that currently exist and that are currently mounted. What happens, and why?
10. Destroy the `mypool/fs1` file system.
11. Remove `file1` from the `/mypool/fs1` directory.
12. Again attempt to create a new file system named `mypool/fs1`. What happens?
13. Verify that the new file system exists and that it is mounted.
14. Rename `mypool/fs1` to `mypool/fs2`.
15. Verify that `mypool/fs2` exists and that it is mounted.
16. List `/mypool`. Does the `fs1` directory exist?
17. Destroy the `mypool/fs2` file system.
18. Verify that `mypool/fs2` is not mounted, and that it no longer exists.
19. List `/mypool`. Does the `fs2` directory exist within it?



Task 3 – Working with ZFS Properties

1. Use the `zfs` command to create a new ZFS in `mypool` called `cmp`. List the space use for `/mypool/cmp`.
2. Use the `tar` command to create a partial archive of the `/usr/lib` directory, and save the archive as `/mypool/cmp/archive.tar`. Use the `-k` option to limit the size of this archive to 400 Mbytes. You will use this file to demonstrate how the `compression` property functions in ZFS file systems.

Note – The `tar` command will ask you to press RETURN when it has created an archive of the size (in kilobytes) you specified. Enter Control-C in response to this prompt. Doing so prevents `tar` from overwriting `archive.tar` with the next set of files that exceed the limit you specified.

3. Use the `zfs get` command to display all of the properties for the `mypool/cmp` file system. Do the properties use the same source? If so, which source, and why?
4. Use `ls` with the `-lh` options to list the size of `/mypool/cmp/archive.tar`. Note the size displayed.
5. Use the `zfs get` command to display the `mypool/cmp` used, compression, and compressratio property settings. Verify the following:
 - That the compression is off
 - That the compression ratio is 1.00x
 - That the space used matches the `archive.tar` size.
6. Set the compression property for `mypool/cmp` to on.
7. Use the `zfs get` command to display the current settings of the used, compression and compressratio properties for `mypool/cmp`. Verify that compression is on. Have the compression ratio or space-used values changed? If not, why? What source is now listed for the compression property?
8. Change the directory to `/mypool/cmp`. Copy `archive.tar` to `archive1.tar`.
9. Use `ls -lh` and `du -h` to display the sizes of `archive.tar` and `archive1.tar`. What file sizes do these commands report?



Note – The size of the compressed archive and compression ratio values may vary, depending on the archive content.

10. Use the `zfs get` command to display the current settings of the `used`, `compression` and `compressratio` properties for `mypool/cmp`. How have the `compressratio` and `used` values changed from the last time you checked them, and why?
11. Set the `compression` property for `mypool/cmp` to `off`.
12. Copy `archive.tar` to `archive2.tar`.
13. Use `ls -lh` and `du -h` to display the sizes of `archive.tar` and `archive1.tar`. What file sizes do these commands report?
14. Use the `zfs get` command to display the current settings of the `used`, `compression` and `compressratio` properties for `mypool/cmp`. How have the `compressratio` and `used` values changed from the last time you checked them, and why? What source is listed for the `compression` property?
15. Use the `zfs inherit` command to clear the `compression` property for `mypool/cmp`. Verify that the change was made. What source is now listed for the `compression` property?
16. Create two new file systems named `mypool/cmp/a` and `mypool/cmp/a/b`.
17. Use the `zfs get` command to recursively display the `compression` property for `mypool/cmp`. Do all three file systems list the same value and source?
18. Use the `zfs set` command to set the `compression` property to `on` for `mypool/cmp`.
19. Use the `zfs get` command to recursively display the `compression` property for `mypool/cmp`. Do all three file systems list the same value and source?
20. Use the `zfs set` command to set the `compression` property to `off` for `mypool/cmp/a`.
21. Use the `zfs get` command to recursively display the `compression` property for `mypool/cmp`. What has changed in the reported values and sources?
22. Use the `zfs inherit` command to clear the `compression` property for `mypool/cmp`.

Exercise: ZFS File System Introduction

23. Use the `zfs get` command to recursively display the compression property for `mypool/cmp`. What has changed in the reported values and sources?
24. Change directory to root (`/`). Use the `zfs destroy` command to destroy `mypool/cmp` and its dependent file systems. Verify that the file systems no longer exist.

Exercise Summary



Discussion – Take a few minutes to discuss what experiences, issues, or discoveries you had during the lab exercises.

- Experiences
- Interpretations
- Conclusions
- Applications

Exercise Solution: ZFS File System Introduction

In this exercise, you complete the following tasks:

- Create and destroy ZFS storage pools
- Create, rename, and destroy ZFS file system
- Work with ZFS properties

Preparation

This exercise requires a minimum of one unused spare disk. Tasks 1 and 2 provide steps to create storage pools that consist of one, two, or three disks, if they are available. You must evaluate the list of disk devices on your system to determine the correct choices to make in tasks 1 and 2. The path names of your disk devices might differ from those described in these exercise. Use the device path names appropriate to your system.



Note – If only one spare disk is available on your system, this exercise asks you to use a single disk to create a storage pool. This is not a recommended configuration. Single-disk pools offer no data redundancy. Failure of the disk supporting the pool causes the pool and its data to become inaccessible.

Exercise Summary

This exercise includes tasks that illustrate the main features of the ZFS file system. The tasks in this exercise illustrate the following concepts:

- Working with ZFS storage pools - creating and destroying a basic ZFS storage pool using spare disks on your system.
- Creating, renaming, and destroying ZFS file systems - creating a basic storage pool, and manipulating ZFS file systems within that pool.
- Working with ZFS properties - setting and listing ZFS properties, and demonstrates principles of ZFS property inheritance.
- Working with ZFS snapshots and clones - creating and destroying ZFS snapshots and clones, and dependencies.
- Removing the ZFS Configuration - removing all of the objects you created during the exercise.

Task 1 – Creating and Destroying ZFS Storage Pools

Complete the following steps:

1. Use the **zpool** command to display the list of ZFS pools. Verify that no ZFS pool currently exists.

```
# zpool list  
no pools available  
#
```

2. Use the **zfs** command to display the list of ZFS file systems. Verify that no ZFS exists.

```
# zfs list  
no datasets available  
#
```

3. Use the **df** command to display the list of UFS file systems currently in use on your system. Use this list to determine which disk or disks on your system are currently in use.

Note – The **df** command with the **-F** option allows you to list file systems of a specific type.



```
# df -F ufs -h  
Filesystem      size   used  avail capacity  Mounted on  
/dev/dsk/c1t0d0s0    9.9G  3.8G   6.0G    39%      /  
/dev/dsk/c1t0d0s7    9.9G  1.9G   7.9G    20%  /export/home  
#
```

4. Use the format command to list what disks are attached to your system. Compare this list with the output of step 3 to determine what disks are currently not in use on your system. Use <Ctrl-D> to exit your format session.

```
# format  
Searching for disks...done
```

AVAILABLE DISK SELECTIONS:

0. c1t0d0 <SUN36G cyl 24620 alt 2 hd 27 sec 107>
/pci@1f,0/pci@1/scsi@8/sd@0,0
1. c1t1d0 <FUJITSU-MAP3367N SUN36G-0401-33.92GB>
/pci@1f,0/pci@1/scsi@8/sd@1,0
2. c2t0d0 <SEAGATE-ST373307LSUN72G-0507-68.37GB>
/pci@1f,0/pci@1/scsi@8,1/sd@0,0
3. c2t1d0 <SEAGATE-ST373307LSUN72G-0507-68.37GB>
/pci@1f,0/pci@1/scsi@8,1/sd@1,0
4. c2t2d0 <SEAGATE-ST373307LSUN72G-0507-68.37GB>
/pci@1f,0/pci@1/scsi@8,1/sd@2,0

Specify disk (enter its number) : ^D
#

Note – Your system may use different sized disks. For this exercise, it's not the best practice. Use the zpool command with the -f option to override errors that this problem causes.

Note – When you try to create a pool, your system will probably display errors that indicate UFS file systems exist on your spare disk. If it does, verify that the spare disk is not in use, then use the zpool command with the -f option to override the errors and create the pool.

5. If you have three unused disks available, use the zpool command to create a storage pool called firstpool that consists of one, three-disk RAID-Z virtual device.

```
# zpool create firstpool raidz c2t0d0 c2t1d0 c2t2d0
```

6. If you have only two unused disks available, use the zpool command to create a storage pool called firstpool that consists of one mirror virtual device.

```
# zpool create firstpool mirror c2t0d0 c2t1d0
```



7. If you have only one unused disk available, use the zpool command to create a storage pool called `firstpool` that uses the single disk as a non-redundant, top-level virtual device.

```
# zpool create firstpool c2t0d0
```

8. Use the zpool command to display the list of ZFS pools.

RAID-Z pool result:

```
# zpool list
```

NAME	SIZE	USED	AVAIL	CAP	HEALTH	ALTROOT
firstpool	204G	141K	204G	0%	ONLINE	-
#						

Mirrored pool result:

```
# zpool list
```

NAME	SIZE	USED	AVAIL	CAP	HEALTH	ALTROOT
firstpool	68G	80K	68.0G	0%	ONLINE	-
#						

Non-redundant pool result:

```
# zpool list
```

NAME	SIZE	USED	AVAIL	CAP	HEALTH	ALTROOT
firstpool	68G	79K	68.0G	0%	ONLINE	-
#						

9. Use the zpool command to display the basic health status of your pool.

```
# zpool status -x
```

all pools are healthy

```
#
```

10. Use the zpool command to display the detailed health status of your pool. What kind of top-level virtual device does the zpool command report, and what devices does it consist of? What is the state of the top-level virtual device, and of the pool itself?

RAID-Z pool result:

```
# zpool status -v
```

pool: firstpool

state: ONLINE

scrub: none requested

config:

NAME	STATE	READ	WRITE	CKSUM
firstpool	ONLINE	0	0	0
raidz1	ONLINE	0	0	0
c2t0d0	ONLINE	0	0	0
c2t1d0	ONLINE	0	0	0

```
c2t2d0  ONLINE      0      0      0
```

errors: No known data errors
#

Mirrored pool result:

```
# zpool status -v
  pool: firstpool
    state: ONLINE
    scrub: none requested
  config:
```

NAME	STATE	READ	WRITE	CKSUM
firstpool	ONLINE	0	0	0
mirror	ONLINE	0	0	0
c2t0d0	ONLINE	0	0	0
c2t1d0	ONLINE	0	0	0

errors: No known data errors
#

Non-redundant pool result:

```
# zpool status -v
  pool: firstpool
    state: ONLINE
    scrub: none requested
  config:
```

NAME	STATE	READ	WRITE	CKSUM
firstpool	ONLINE	0	0	0
c2t0d0	ONLINE	0	0	0

errors: No known data errors
#

The pool *firstpool* should contain one top-level virtual device that matches the type you specified when you created the pool; a *raidz1*, a *mirror*, or in the non-redundant case, the single disk you specify serves as the only top-level device. The *raidz1* device should consist of three disks. The *mirror* device should consist of two disks. Both the top-level virtual device and the pool should be in the *ONLINE* state.

11. Use the *zfs* command to list ZFS files, and display the space utilization for *firstpool*. Does the space available match what you would expect for a pool of the type you created?

RAID-Z pool result:

```
# zfs list
NAME          USED  AVAIL  REFER  MOUNTPOINT
```

```
firstpool          89.9K  134G  32.6K  /firstpool
#
```

Mirrored pool result:

```
# zfs list
NAME          USED  AVAIL  REFER  MOUNTPOINT
firstpool    77K   66.9G  24.5K  /firstpool
#
```

Non-redundant pool result:

```
# zfs list
NAME          USED  AVAIL  REFER  MOUNTPOINT
firstpool    76K   66.9G  24.5K  /firstpool
#
```

The space available should be roughly equal to twice the size of the smallest disk, if you created a pool that consists of one, three-disk RAID-Z device. This accounts for the space required to store parity data in this particular pool. The space available should roughly match the size of the smallest disk, if you created a pool with a single, two-disk mirror. The space available should roughly match the size of the disk you specified, if you created a non-redundant pool.

12. Use the zpool command to display the name of your pool, and only its space used and space available statistics.

RAID-Z pool result:

```
# zpool list -o name,used,available
NAME          USED  AVAIL
firstpool    141K  204G
#
```

Mirrored pool result:

```
# zpool list -o name,used,available
NAME          USED  AVAIL
firstpool    80K   68.0G
#
```

Non-redundant pool result:

```
# zpool list -o name,used,available
NAME          USED  AVAIL
firstpool    79K   68.0G
#
```



Note – For a RAID-Z pool, the `zpool list` command reports the sum of all disk space in the AVAIL column. This differs from the value reported by `zfs list`. For a RAID-Z pool, the `zfs list` command reduces the value reported in the AVAIL column by the amount of parity data that the pool requires.

13. Use the `zfs` command to display the list of currently mounted ZFS file systems. Display a long listing for the `/firstpool` directory.

```
# zfs mount
firstpool                               /firstpool
# ls -ld /firstpool
drwxr-xr-x  2 root      sys            2 Dec 12 19:16
/firstpool
#
```

14. Use the `zpool` command to destroy `firstpool`. Verify that `firstpool` no longer exists.

```
# zpool destroy firstpool
# zpool list
no pools available
#
```

15. Verify that the `firstpool` file system and the `/firstpool` mount point no longer exist.

```
# zfs list
no datasets available
# ls -ld /firstpool
/firstpool: No such file or directory
#
```

Task 2 – Creating, Renaming, and Destroying ZFS File Systems

Complete the following steps:

1. Use the `zpool` command to display the list of ZFS pools. Verify that no ZFS pool currently exists.

```
# zpool list
no pools available
#
```

2. If you have three unused disks available, use the `zpool` command to create a storage pool called `mypool` that consists of one, three-disk RAID-Z virtual device.



- ```
zpool create mypool raidz c2t0d0 c2t1d0 c2t2d0
```
3. If you have only two unused disks available, use the zpool command to create a storage pool called mypool that consists of one mirror virtual device.

```
zpool create mypool mirror c2t0d0 c2t1d0
```

  4. If you have only one unused disk available, use the zpool command to create a storage pool called mypool that uses the single disk as a non-redundant, top-level virtual device.

```
zpool create mypool c2t0d0
```

**Note** – Example command output throughout the remainder of this exercise assumes that a non-redundant pool is in use. This isn't a recommended pool configuration.

5. Use the zfs command to list ZFS file systems and display /mypool space use.

```
zfs list
```

| NAME   | USED | AVAIL | REFER | MOUNTPOINT |
|--------|------|-------|-------|------------|
| mypool | 76K  | 66.9G | 24.5K | /mypool    |
| #      |      |       |       |            |

6. Attempt to create a file system named mypool/fs1/fs2. What happens and why?

```
zfs create mypool/fs1/fs2
```

cannot create 'mypool/fs1/fs2': parent does not exist  
#

*The attempt fails because the intermediate file system fs1 does not exist.*

7. Create a directory called /mypool/fs1. Again attempt to create a file system named mypool/fs1/fs2. What happens and why?

```
mkdir /mypool/fs1
```

```
zfs create mypool/fs1/fs2
```

cannot create 'mypool/fs1/fs2': parent does not exist  
#

*The attempt fails because the intermediate file system fs1 does not exist. The ZFS only creates new file systems immediately below a parent file system.*

8. Create a new empty file named /mypool/fs1/file1.

```
touch /mypool/fs1/file1
```

#

9. Attempt to create a new file system named mypool/fs1. Use the zfs list and zfs mount commands to list the ZFS file systems that currently exist and that are currently mounted. What happens, and why?

```
zfs create mypool/fs1
cannot mount '/mypool/fs1': directory is not empty
filesystem successfully created, but not mounted
zfs list
NAME USED AVAIL REFER MOUNTPOINT
mypool 108K 66.9G 27.5K /mypool
mypool/fs1 24.5K 66.9G 24.5K /mypool/fs1
zfs mount
mypool
#
```

*The new file system exists, but isn't mounted. The empty file prevented the /mypool/fs1 file system from mounting.*

10. Destroy the mypool/fs1 file system.

```
zfs destroy mypool/fs1
#
```

11. Remove file1 from the /mypool/fs1 directory.

```
rm /mypool/fs1/file1
#
```

12. Again attempt to create a new file system named mypool/fs1. What happens?

```
zfs create mypool/fs1
#
```

*The mypool/fs1 file system is created successfully.*

13. Verify that the new file system exists and that it is mounted.

```
zfs list
NAME USED AVAIL REFER MOUNTPOINT
mypool 108K 66.9G 26.5K /mypool
mypool/fs1 24.5K 66.9G 24.5K /mypool/fs1
zfs mount
mypool
#
```

14. Rename mypool/fs1 to mypool/fs2.

```
zfs rename mypool/fs1 mypool/fs2
```

15. Verify that mypool/fs2 exists and that it is mounted.

```
zfs list
NAME USED AVAIL REFER MOUNTPOINT
mypool 108K 66.9G 26.5K /mypool
mypool/fs2 24.5K 66.9G 24.5K /mypool/fs2
zfs mount
mypool /mypool
mypool/fs2 /mypool/fs2
#
```

16. List /mypool. Does the fs1 directory exist?

```
ls /mypool
fs2
#
```

*No, the fs1 directory no longer exists.*

17. Destroy the mypool/fs2 file system.

```
zfs destroy mypool/fs2
#
```

18. Verify that mypool/fs2 isn't mounted and no longer exists.

```
zfs mount
mypool /mypool
zfs list
NAME USED AVAIL REFER MOUNTPOINT
mypool 76K 66.9G 24.5K /mypool
#
```

19. List /mypool. Does the fs2 directory exist?

```
ls /mypool
#
#
```

*No, the /mypool directory is empty.*

## Task 3 – Working With ZFS Properties

1. Use the zfs command to create a new ZFS in mypool called cmp. List the space use for /mypool/cmp.

```
zfs create mypool/cmp
zfs list
NAME USED AVAIL REFER MOUNTPOINT
mypool 104K 66.9G 24.5K /mypool
mypool/cmp 24.5K 66.9G 24.5K /mypool/cmp
#
```

2. Use the tar command to create a partial archive of the /usr/lib directory, and save the archive as /mypool/cmp/archive.tar. Use the -k option to limit the size of this archive to 400 Mbytes. You will use this file to demonstrate how the compression property functions in ZFS file systems.

```
tar cfk /mypool/cmp/archive.tar 409600 /usr/lib
tar: please insert new volume, then press RETURN. (Enter Control-C)
#
```



**Note** – The tar command will ask you to press RETURN when it has created an archive of the size (in kilobytes) you specified. Enter Control-C in response to this prompt. Doing so prevents tar from overwriting archive.tar with the next set of files that exceed the limit you specified.

3. Use the zfs get command to display all of the properties for the mypool/cmp file system. Do the properties all use the same source? If so, which source is it, and why?

```
zfs get all mypool/cmp
```

| NAME       | PROPERTY      | VALUE                | SOURCE  |
|------------|---------------|----------------------|---------|
| mypool/cmp | type          | filesystem           | -       |
| mypool/cmp | creation      | Mon Oct 5 12:57 2009 | -       |
| mypool/cmp | used          | 400M                 | -       |
| mypool/cmp | available     | 72.9G                | -       |
| mypool/cmp | referenced    | 400M                 | -       |
| mypool/cmp | compressratio | 1.00x                | -       |
| mypool/cmp | mounted       | yes                  | -       |
| mypool/cmp | quota         | none                 | default |
| mypool/cmp | reservation   | none                 | default |
| mypool/cmp | recordsize    | 128K                 | default |
| mypool/cmp | mountpoint    | /mypool/cmp          | default |
| mypool/cmp | sharenfs      | off                  | default |
| mypool/cmp | checksum      | on                   | default |
| mypool/cmp | compression   | off                  | default |
| mypool/cmp | atime         | on                   | default |
| mypool/cmp | devices       | on                   | default |
| mypool/cmp | exec          | on                   | default |
| mypool/cmp | setuid        | on                   | default |
| mypool/cmp | readonly      | off                  | default |
| mypool/cmp | zoned         | off                  | default |
| mypool/cmp | snapdir       | hidden               | default |
| mypool/cmp | aclmode       | groupmask            | default |
| mypool/cmp | aclinherit    | restricted           | default |
| mypool/cmp | canmount      | on                   | default |
| mypool/cmp | shareiscsi    | off                  | default |

```

mypool/cmp xattr on default
mypool/cmp copies 1 default
mypool/cmp version 4 -
mypool/cmp utf8only off -
mypool/cmp normalization none -
mypool/cmp casesensitivity sensitive -
mypool/cmp vscan off default
mypool/cmp nbmand off default
mypool/cmp sharesmb off default
mypool/cmp refquota none default
mypool/cmp refreservation none default
mypool/cmp primarycache all default
mypool/cmp secondarycache all default
mypool/cmp usedbysnapshots 0 -
mypool/cmp usedbydataset 400M -
mypool/cmp usedbychildren 0 -
mypool/cmp usedbyrefreservation 0 -
#

```

*All of the settable properties have the same source: default. This is true because none of the properties have been set locally or inherited from a locally-set property.*

4. Use ls with the -lh options to list the size of /mypool/cmp/archive.tar. Make note of the size displayed.

```

ls -lh /mypool/cmp/archive.tar
-rw-r--r-- 1 root root 400M Dec 18 17:18
/mypool/cmp/archive.tar
#

```

5. Use the zfs get command to display the current settings of the used, compression and compressratio properties for mypool/cmp. Verify that compression is off, the compression ratio is 1.00x, and that the space used matches the size of archive.tar.

```

zfs get used,compression,compressratio mypool/cmp
NAME PROPERTY VALUE SOURCE
mypool/cmp used 400M -
mypool/cmp compression off default
mypool/cmp compressratio 1.00x -
#

```

6. Set the compression property for mypool/cmp to on.

```

zfs set compression=on mypool/cmp
#

```

7. Use `zfs get` to display the current settings of the used, compression and compressratio properties for `mypool/cmp`. Verify that compression is now on. Have the compression ratio or space used values changed? If not, why not? What source is listed for the compression property?

```
zfs get used,compression,compressratio mypool/cmp
```

| NAME       | PROPERTY      | VALUE | SOURCE |
|------------|---------------|-------|--------|
| mypool/cmp | used          | 400M  | -      |
| mypool/cmp | compression   | on    | local  |
| mypool/cmp | compressratio | 1.00x | -      |
| #          |               |       |        |

*No, the compression ratio remains 1.00, and the space used remains 400Mbytes. Turning on compression for a file system does not compress files that already exist within the file system. The compression property lists the local source.*

8. Change directory to `/mypool/cmp`. Copy `archive.tar` to `archive1.tar`.

```
cd /mypool/cmp
cp archive.tar archive1.tar
```

9. Use `ls -lh` and `du -h` to display the sizes of `archive.tar` and `archive1.tar`. What file sizes do these commands report?

```
ls -lh
total 1348943
-rw-r--r-- 1 root root 400M Dec 18 17:18 archive.tar
-rw-r--r-- 1 root root 400M Dec 18 17:38 archive1.tar
du -h *
400M archive.tar
234M archive1.tar
#
```

*The `ls -lh` command reports that both files are 400 Mbytes in size. The `du -h` command reports that `archive.tar` uses 400 Mbytes of space, and `archive1.tar` uses about 234 Mbytes.*

---

**Note** – The size of the compressed archive and compression ratio values may vary, depending on the archive content.



10. Use the `zfs get` command to display the current settings of the used, compression and compressratio properties for `mypool/cmp`. How have the compressratio and used values changed from the last time you checked them, and why?

```
zfs get used,compression,compressratio mypool/cmp
```

| NAME       | PROPERTY      | VALUE | SOURCE |
|------------|---------------|-------|--------|
| mypool/cmp | used          | 659M  | -      |
| mypool/cmp | compression   | on    | local  |
| mypool/cmp | compressratio | 1.24x | -      |
| #          |               |       |        |

The used property now lists about 659 Mbytes, and the compressratio property now indicates about 1.24x compression. The space used matches the sum of the file sizes, and the compression ratio indicates the compression achieved for all files in the file system.

11. Set the compression property for mypool/cmp to off.

```
zfs set compression=off mypool/cmp
#
```

12. Copy archive.tar to archive2.tar.

```
cp archive.tar archive2.tar
#
```

13. Use `ls -lh` and `du -h` to display the sizes of `archive.tar` and `archive1.tar`. What file sizes do these commands report?

```
ls -lh
total 2168760
-rw-r--r-- 1 root root 400M Dec 18 17:18 archive.tar
-rw-r--r-- 1 root root 400M Dec 18 17:38 archive1.tar
-rw-r--r-- 1 root root 400M Dec 18 18:56 archive2.tar
du -h *
400M archive.tar
234M archive1.tar
334M archive2.tar
```

*The `ls -lh` command reports that all three files are 400 Mbytes in size. The `du -h` command reports that `archive.tar` uses 400 Mbytes, `archive1.tar` uses 234 Mbytes of space, and `archive2.tar` uses 334 Mbytes.*

14. Use `zfs get` to display the current settings of the `used`, `compression` and `compressratio` properties for `mypool/cmp`. How have the `compressratio` and `used` values changed from the last time you checked them, and why? What source is listed for the `compression` property?

```
zfs get used,compression,compressratio mypool/cmp
NAME PROPERTY VALUE SOURCE
mypool/cmp used 1.01G -
mypool/cmp compression off local
mypool/cmp compressratio 1.15x -
#
```

*The `used` property now lists about 1.01 Gbytes, and the `compressratio` property now indicates about 1.15x compression. The space used matches the sum of the file sizes. The compression ratio is smaller because a smaller proportion of the data in the file system is compressed. The `compression` property still lists the `local` source.*

15. Use `zfs inherit` to clear the `compression` property for `mypool/cmp`. Verify that the change was made. What source is listed for the `compression` property?

```
zfs inherit compression mypool/cmp
zfs get compression mypool/cmp
NAME PROPERTY VALUE SOURCE
mypool/cmp compression off default
#
```

16. Create two new file systems named `mypool/cmp/a` and `mypool/cmp/a/b`.

```
zfs create mypool/cmp/a
zfs create mypool/cmp/a/b
```

17. Use `zfs get` to recursively display the compression property for `mypool/cmp`. Do all three file systems list the same value and source?

```
zfs get -r compression mypool/cmp
```

| NAME                        | PROPERTY                 | VALUE            | SOURCE               |
|-----------------------------|--------------------------|------------------|----------------------|
| <code>mypool/cmp</code>     | <code>compression</code> | <code>off</code> | <code>default</code> |
| <code>mypool/cmp/a</code>   | <code>compression</code> | <code>off</code> | <code>default</code> |
| <code>mypool/cmp/a/b</code> | <code>compression</code> | <code>off</code> | <code>default</code> |
| #                           |                          |                  |                      |

*Yes, all three file systems have compression turned off, and they list the default source for the compression property.*

18. Use the `zfs set` command to set the compression property to on for `mypool/cmp`.

```
zfs set compression=on mypool/cmp
#
```

19. Use the `zfs get` command to recursively display the compression property for `mypool/cmp`. Do all three file systems list the same value and source?

```
zfs get -r compression mypool/cmp
```

| NAME                        | PROPERTY                 | VALUE           | SOURCE                      |
|-----------------------------|--------------------------|-----------------|-----------------------------|
| <code>mypool/cmp</code>     | <code>compression</code> | <code>on</code> | <code>local</code>          |
| <code>mypool/cmp/a</code>   | <code>compression</code> | <code>on</code> | <code>inherited from</code> |
| <code>mypool/cmp</code>     | <code>compression</code> | <code>on</code> | <code>inherited from</code> |
| <code>mypool/cmp/a/b</code> | <code>compression</code> | <code>on</code> | <code>inherited from</code> |
| <code>mypool/cmp</code>     |                          |                 |                             |

*All three file systems have compression turned on. The `mypool/cmp` file system lists the local source for the compression property, `mypool/cmp/a` and `mypool/cmp/a/b` indicate that they inherited their compression values from `mypool/cmp`.*

20. Use the `zfs set` command to set the compression property to off for `mypool/cmp/a`.

```
zfs set compression=off mypool/cmp/a
```

21. Use the `zfs get` command to recursively display the compression property for `mypool/cmp`. What has changed in the reported values and sources?

```
zfs get -r compression mypool/cmp
```

| NAME                        | PROPERTY                 | VALUE            | SOURCE                      |
|-----------------------------|--------------------------|------------------|-----------------------------|
| <code>mypool/cmp</code>     | <code>compression</code> | <code>on</code>  | <code>local</code>          |
| <code>mypool/cmp/a</code>   | <code>compression</code> | <code>off</code> | <code>local</code>          |
| <code>mypool/cmp/a/b</code> | <code>compression</code> | <code>off</code> | <code>inherited from</code> |
| <code>mypool/cmp/a</code>   |                          |                  |                             |

The mypool/cmp/a and mypool/cmp/a/b file systems now have compression turned off. The mypool/cmp/a file system lists the local source for the compression property, and mypool/cmp/a/b now inherits its compression value from mypool/cmp/a.

22. Use the zfs inherit command to clear the compression property for mypool/cmp.

```
zfs inherit compression mypool/cmp
```

23. Use the zfs get command to recursively display the compression property for mypool/cmp. What has changed in the reported values and sources?

```
zfs get -r compression mypool/cmp
NAME PROPERTY VALUE
mypool/cmp compression off
mypool/cmp/a compression off
mypool/cmp/a/b compression off
mypool/cmp/a
#
```

All three file systems have compression turned off. The mypool/cmp, mypool/cmp/a, and mypool/cmp/a/b file systems list the default, local, and inherited from sources, respectively. The mypool/cmp/a/b file system inherits its compression value from mypool/cmp/a.

24. Change directory to root (/). Use the zfs destroy command to destroy mypool/cmp and its dependent file systems. Verify that the file systems no longer exist.

```
cd /
zfs destroy -r mypool/cmp
zfs list
NAME USED AVAIL REFER MOUNTPOINT
mypool 103K 66.9G 25.5K /mypool
#
```

25. Use the zpool command to destroy mypool file system.

```
zpool destroy mypool
#
```

26. Verify that the mypool file system no longer exist.

```
zpool list
no pools available
zfs list
no datasets available
#
```

---

## Module 6

---

# Performing Mounts and Unmounts

---

## Objectives

Upon completion of this module, you should be able to:

- Identify mounting basics
- Perform mounts
- Mount ZFS file systems
- Perform unmounts
- Repair important files if boot fails
- Access a mounted diskette, CD-ROM, or DVD
- Restrict access to a mounted diskette, CD-ROM, or DVD
- Access a diskette, CD-ROM or DVD without volume management (vol&d)

## Working With Mounting Basics

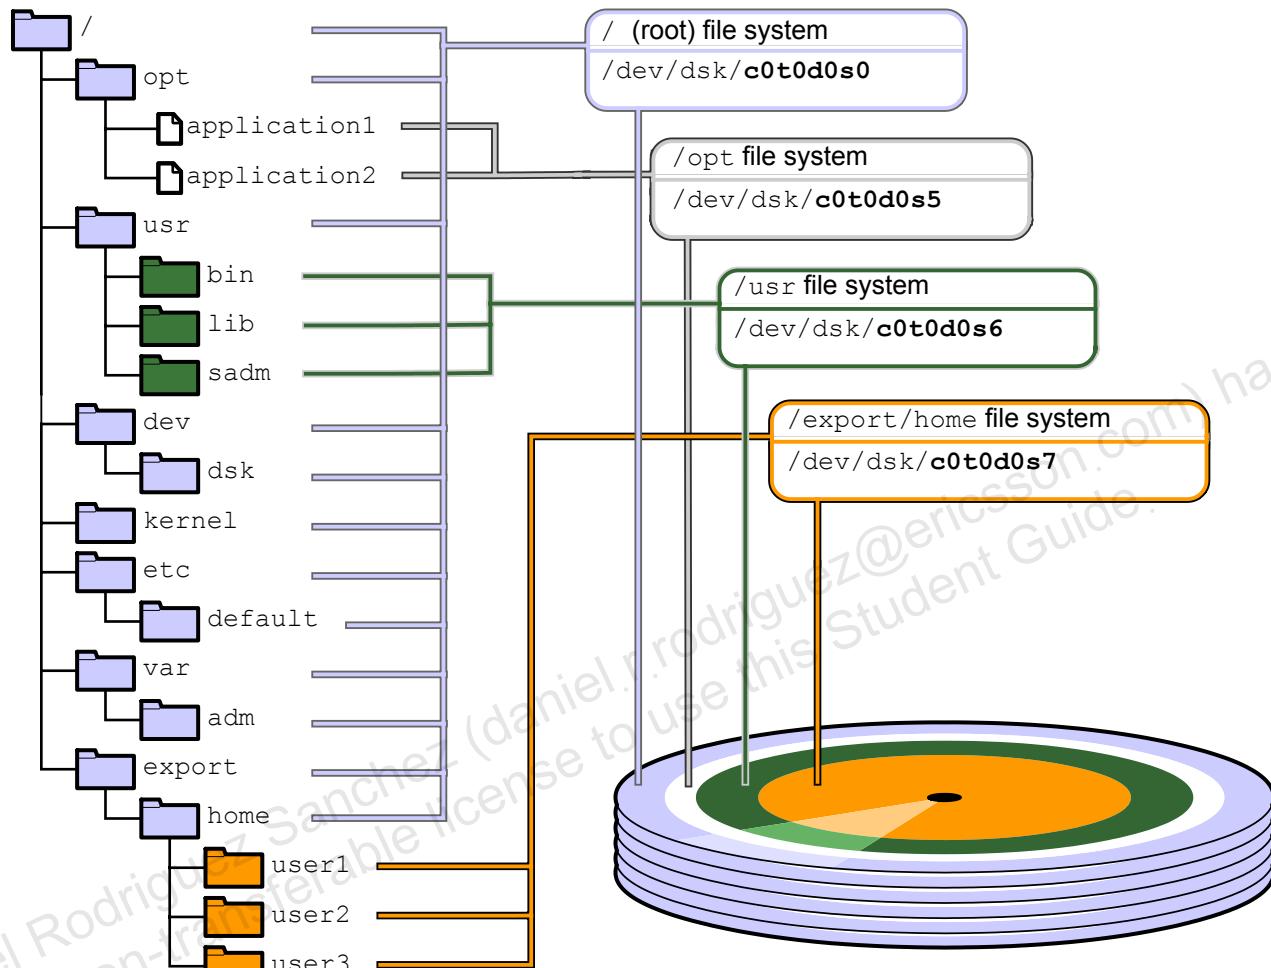
In the Solaris OS, use the mounting process to attach individual file systems to their mount points on the directory hierarchy. This action makes a file system accessible to the system and to the users.

Use the unmounting process to detach a file system from its mount point in the directory hierarchy. This action makes a file system unavailable to the system or users.

After you create a UNIX<sup>R</sup> file system (UFS) with the newfs command, you must attach it to the Solaris OS directory hierarchy at a mount point. A mount point is a directory that is the point of connection for a file system. File systems are commonly referred to by the names of their mount points, for example, the / (root) file system or the /usr file system.

This module focuses on methods used to mount and unmount UFSs.

Figure 6-1 shows how the directory hierarchy spans from one file system to the next.



**Figure 6-1** File Systems and Mount Points

File systems do not contain their own mount point directories.

## Determining Which File Systems Are Currently Mounted

Determine which file systems are currently mounted by using the mount or df command.

The df command displays the amount of disk space occupied by mounted or unmounted file systems and, depending on the options used, displays both locally mounted and virtual file system information.

The mount command, which is located in the /usr/sbin directory, maintains a table of currently mounted file systems in the /etc/mnttab file system table. When the mount command is used without arguments, it lists all of the mounted file systems in the /etc/mnttab file system table. When used with only a partial argument list, the mount command searches the /etc/vfstab file for an entry that supplies the missing arguments.

---

**Note** – Options to the mount command are discussed later in this module.

---

**Note** – The symbolic link /usr/sbin/mount points to the /sbin/mount command.

---

## Mounting a File System Automatically

The Solaris OS provides several methods for automating file system mounts.

The Solaris OS creates a default /etc/vfstab file during software installation, based on your selections. You can edit the /etc/vfstab file whenever file system entries need to be added or modified.

---

**Note** – The automounter can mount network file systems on demand.

---





## Introducing the Virtual File System Table: /etc/vfstab

The /etc/vfstab file lists all the file systems to be automatically mounted at system boot time, with the exception of the /etc/mnttab and /var/run file systems.

The file format includes seven fields per line entry. By default, a tab separates each field, but any white space can be used for separators. The dash (-) character is used as a placeholder for fields when text arguments are not appropriate. Commented lines begin with the (#) symbol.

---

**Note –** Because the default is to use tabs to separate the fields in the /etc/vfstab file, the fields often do not line up under their respective headings. This can lead to some confusion when you are viewing this file in a terminal window.

---

An example /etc/vfstab file from a Sun Blade 1500:

```
cat /etc/vfstab
#device device mount FS fsck mount
mount
#to mount to fsck point type pass at boot
options
#
fd - /dev/fd fd - no -
/proc - /proc proc - no -
/dev/dsk/c0t0d0s1 - - swap - no -
/dev/dsk/c0t0d0s0 /dev/rdsk/c0t0d0s0 / ufs 1
no - -
/dev/dsk/c0t0d0s6 /dev/rdsk/c0t0d0s6 /usr ufs 1
no - -
/dev/dsk/c0t0d0s3 /dev/rdsk/c0t0d0s3 /var ufs 1
no - -
/dev/dsk/c0t0d0s7 /dev/rdsk/c0t0d0s7 /export/home ufs 2
yes - -
/dev/dsk/c0t0d0s5 /dev/rdsk/c0t0d0s5 /opt ufs 2
yes - -
/devices - /devices devfs - no -
ctfs - /system/contract ctfss - no -
objfs - /system/object objfs - no -
swap - /tmp tmpfs - yes -
#
```

## Working With Mounting Basics

An example /etc/vfstab file from an Ultra 20 system:

```
cat /etc/vfstab
#device device mount FS fsck mount
mount
#to mount to fsck point type pass at boot
options
#
fd - /dev/fd fd - no -
/proc - /proc proc - no -
/dev/dsk/c1d0s1 - - swap - no -
/dev/dsk/c1d0s0 /dev/rdsk/c1d0s0 / ufs 1 no -
/dev/dsk/c1d0s6 /dev/rdsk/c1d0s6 /usr ufs 1 no -
/dev/dsk/c1d0s3 /dev/rdsk/c1d0s3 /var ufs 1 no -
/dev/dsk/c1d0s7 /dev/rdsk/c1d0s7 /export/home ufs 2 -
yes - -
/dev/dsk/c1d0s5 /dev/rdsk/c1d0s5 /opt ufs 2 yes -
/devices - /devices devfs - no -
ctfs - /system/contract ctfss - no -
objfs - /system/object objfs - no -
swap - /tmp tmpfs - yes -
#
```

To add a new entry to the /etc/vfstab file, you need the following information:

|                 |                                                                                                                                                                      |
|-----------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| device to mount | The device to be mounted. For example, a local ufs file system /dev/dsk/c#t#d#s#, or a pseudo file system /proc.                                                     |
| device to fsck  | The raw or <i>character</i> device checked by the file system check program (fsck) if applicable. Pseudo and distributed file systems have a dash (-) in this field. |
| mount point     | The name of the directory that serves as the attach mount point in the Solaris OS directory hierarchy.                                                               |
| FS type         | The type of file system to be mounted.                                                                                                                               |

## fsck pass

The pass number used by the `fsck` command to decide whether to check a file system. When the field contains a (-), the file system is not checked. When the field contains a zero, UFS file systems are not checked, however, non-UFS file systems are checked. When the field contains a value greater than zero, the file system is always checked. All file systems with a value of 1 in this field are checked one at a time in the order they appear in the `vfstab` file. When the `fsck` command is run on multiple UFSs that have `fsck` pass values greater than 1 and the `preen` option (`-o p`) is used, the `fsck` command automatically checks the file systems on different disks in parallel to maximize efficiency. Otherwise, the value of the pass number does not have any effect.

## mount at boot

Enter yes to enable the `mountall` command to mount the file systems at boot time. Enter no to prevent a file system mount at boot time.



**Note** – For / (root), /usr, and /var (if it is a separate file system) file systems, the mount at boot field value is specified as no. The kernel mounts these file systems as part of the boot sequence before the `mountall` command is run. SMF uses scripts in the `/lib/svc/method` directory whose names begin with `fs-` to mount all required file systems.

## mount options

A comma-separated list of options passed to the `mount` command. A dash (-) indicates the use of default mount options.

## Introducing the /etc/mnttab File System Table

The file `/etc/mnttab` is an `mntfs` file system that provides read-only access to the table of mounted file systems for the current host.

Each time a file system is mounted, the `mount` command adds an entry to this table. Whenever a file system is unmounted, its entry is removed from this table. These are the fields in the `/etc/mnttab` file system table:

|                  |                                                                                                                                                                                    |
|------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Device Name      | The name of the device that is mounted at the mount point. This block device is where the file system is physically located.                                                       |
| Mount Point      | The mount point or directory name where the file system is attached (for example, <code>/usr</code> , <code>/opt</code> ).                                                         |
| File System Type | The file system type of the mounted file system.                                                                                                                                   |
| Mount Options    | The list of mount options in effect for the file system. This includes a <code>dev=number</code> which reflects the minor device number of the mounted file system in hexadecimal. |
| Time             | The time at which the file system was mounted.                                                                                                                                     |

The `/var/run` file system is a `tmpfs` mounted file system in the Solaris OS. It is the repository for temporary operating system files that are not needed across system reboots in this Solaris OS release. It is mounted as a pseudo file system rather than a disk-based file system.

The `/var/run` directory requires no administration. For security reasons, it is owned by the `root` user.

The `/tmp` directory continues to be a `tmpfs` mounted file system in the Solaris OS. It is the repository for temporary user and application files that are not needed across system reboots. It is a pseudo file system rather than a disk-based file system.

The following examples show two ways to display currently mounted file systems.

```
cat /etc/mnttab
/dev/dsk/c0t0d0s0 / ufs
rw,intr,largefiles,logging,xattr,onerror=panic,dev=2200008 1170808975
/devices /devices devfs dev=53c0000 1170808960
ctfs /system/contract ctf dev=5400001 1170808960
proc /proc proc dev=5440000 1170808960
mnttab /etc/mnttab mntfs dev=5480001 1170808960
swap /etc/svc/volatile tmpfs xattr,dev=54c0001 1170808960
objfs /system/object objfs dev=5500001 1170808960
/dev/dsk/c0t0d0s6 /usr ufs
rw,intr,largefiles,logging,xattr,onerror=panic,dev=220000e 1170808975
/platform/sun4u-us3/lib/libc_psr/libc_psr_hwcap1.so.1 /platform/sun4u-
us3/lib/libc_psr.so.1 lofs dev=2200008 1170808971
/platform/sun4u-us3/lib/sparcv9/libc_psr/libc_psr_hwcap1.so.1 /platform/sun4u-
us3/lib/sparcv9/libc_psr.so.1 lofs dev=2200008 1170808972
fd /dev/fd fd rw,dev=56c0001 1170808975
/dev/dsk/c0t0d0s3 /var ufs
rw,intr,largefiles,logging,xattr,onerror=panic,dev=220000b 1170808989
swap /tmp tmpfs xattr,dev=54c0002 1170808989
swap /var/run tmpfs xattr,dev=54c0003 1170808989
/dev/dsk/c0t0d0s5 /opt ufs
rw,intr,largefiles,logging,xattr,onerror=panic,dev=220000d 1170808994
/dev/dsk/c0t0d0s7 /export/home ufs
rw,intr,largefiles,logging,xattr,onerror=panic,dev=220000f 1170808994
-hosts /net autofs nosuid,indirect,ignore,nobrowse,dev=5780001
1170809014
auto_home /home autofs indirect,ignore,nobrowse,dev=5780002
1170809014
sys-05:vold(pid517) /vol nfs ignore,noquota,dev=5740001
1170809022
#
```

## Working With Mounting Basics

---

```
mount
/ on /dev/dsk/c0t0d0s0
read/write/setuid/devices/intr/largefiles/logging/xattr/onerror=panic/dev=220000
8 on Tue Feb 6 17:42:55 2007
/devices on /devices read/write/setuid/devices/dev=53c0000 on Tue Feb 6
17:42:40 2007
/system/contract on ctfs read/write/setuid/devices/dev=5400001 on Tue Feb 6
17:42:40 2007
/proc on proc read/write/setuid/devices/dev=5440000 on Tue Feb 6 17:42:40 2007
/etc/mnttab on mnttab read/write/setuid/devices/dev=5480001 on Tue Feb 6
17:42:40 2007
/etc/svc/volatile on swap read/write/setuid/devices/xattr/dev=54c0001 on Tue Feb
6 17:42:40 2007
/system/object on objfs read/write/setuid/devices/dev=5500001 on Tue Feb 6
17:42:40 2007
/usr on /dev/dsk/c0t0d0s6
read/write/setuid/devices/intr/largefiles/logging/xattr/onerror=panic/dev=220000
e on Tue Feb 6 17:42:55 2007
/platform/sun4u-us3/lib/libc_psr.so.1 on /platform/sun4u-
us3/lib/libc_psr/libc_psr_hwcap1.so.1 read/write/setuid/devices/dev=2200008 on
Tue Feb 6 17:42:51 2007
/platform/sun4u-us3/lib/sparcv9/libc_psr.so.1 on /platform/sun4u-
us3/lib/sparcv9/libc_psr/libc_psr_hwcap1.so.1
read/write/setuid/devices/dev=2200008 on Tue Feb 6 17:42:52 2007
/dev/fd on fd read/write/setuid/devices/dev=56c0001 on Tue Feb 6 17:42:55 2007
/var on /dev/dsk/c0t0d0s3
read/write/setuid/devices/intr/largefiles/logging/xattr/onerror=panic/dev=220000
b on Tue Feb 6 17:43:09 2007
/tmp on swap read/write/setuid/devices/xattr/dev=54c0002 on Tue Feb 6 17:43:09
2007
/var/run on swap read/write/setuid/devices/xattr/dev=54c0003 on Tue Feb 6
17:43:09 2007
/opt on /dev/dsk/c0t0d0s5
read/write/setuid/devices/intr/largefiles/logging/xattr/onerror=panic/dev=220000
d on Tue Feb 6 17:43:14 2007
/export/home on /dev/dsk/c0t0d0s7
read/write/setuid/devices/intr/largefiles/logging/xattr/onerror=panic/dev=220000
f on Tue Feb 6 17:43:14 2007
#
```

# Performing Mounts

Mount file systems manually by running the `mount` command, or the system can automatically mount file systems at boot time after consulting the `/etc/vfstab` file.

## Mounting a Local File System Manually

The `mount` command not only lists which file systems are currently mounted, it also provides you with a method for mounting file systems.

### Default Behavior of the `mount` Command

To mount a local file system manually, you need to know the name of the device where the file system resides and its mount point path name. Perform the command:

```
mount /dev/dsk/c0t0d0s7 /export/home
```

In this example, the default action mounts the file system with the following options: `read/write`, `setuid`, `intr`, `logging`, `largefiles`, `xattr`, and `onerror`.

The following list explains the default options for the `mount` command.

|                          |                                                                                                                                       |
|--------------------------|---------------------------------------------------------------------------------------------------------------------------------------|
| <code>read/write</code>  | Indicates whether reads and writes are allowed on the file system.                                                                    |
| <code>setuid</code>      | Permits the execution of <code>setuid</code> programs in the file system.                                                             |
| <code>intr/nointr</code> | Allows and forbids keyboard interrupts to kill a process that is waiting for an operation on a locked file system.                    |
| <code>logging</code>     | Indicates that logging is enabled for the <code>ufs</code> file system. This is the default for the Solaris 10 OS.                    |
| <code>largefiles</code>  | Allows for the creation of files larger than 2 Gbytes. A file system mounted with this option can contain files larger than 2 Gbytes. |
| <code>xattr</code>       | Supports extended attributes not found in standard UNIX filesystems. See <code>man fsattr(5)</code> .                                 |



**Note** – Due to file system overhead, the largest file that can be created in a ufs file system is approximately 1 Tbyte. The data capacity of a 1 Tbyte file system is approximately 1 Tbyte minus 0.5% overhead and the recommended 1% free space.

---

**onerror=action**      Specifies the action that the ufs file system should take to recover from an internal inconsistency on a file system. An action can be specified as:

**panic** — Causes a forced system shutdown. This is the default.

**lock** — Applies a file system lock to the file system.

**umount** — Forcibly unmounts the file system.

The /etc/vfstab file provides you with another important feature. Because the /etc/vfstab file contains the mapping between the mount point and the actual device name, the root user can manually mount a file system specifying only the mount point on the command line.

```
mount /export/home
```

## Using the `mount` Command Options

When you are using mount options on the command line, remember that the options are preceded by the `-o` flag. When you are using multiple options, enter them as a comma-separated list following the `-o` flag.

```
mount -o option,option,... device_name mount_point
```

---

**Note** – Mount options are described in detail in the `mount_ufs` man page. This man page describes mount command options specific to the ufs file system.

---



Some options used to mount local file systems include: `ro`, `nosetuid`, `noatime`, `nolargefiles`, and `nologging`.

- `ro` – Mounts the file system as read-only.

The following is an example using this option on the command line:

```
mount -o ro /dev/dsk/c0t0d0s7 /export/home
```

- `nosuid` – Prohibits the execution of setuid programs in the file system. This does not restrict the creation of setuid programs.

The following example shows the use of multiple options on the command line:

```
mount -o ro,nosuid /dev/dsk/c0t0d0s7 /export/home
```

- `noatime` – Suppresses the time-last-accessed modification on inodes, which reduces disk activity on a file system where access times are not important. Specifying this option generally improves file access times and boosts overall performance, for example:

```
mount -o noatime /dev/dsk/c0t0d0s7 /export/home
```

- `nolargefiles` – Prevents a file system that contains one or more “large files” from being mounted, for example:

```
mount -o nolargefiles /dev/dsk/c0t0d0s7 /export/home
```

Use of the `nolargefiles` option fails if the file system to be mounted contains a large file or did contain a large file at one time.

If the file system currently contains a large file and the root user needs to mount it with this option, then the large file must be located and moved or removed from the file system. Then you must execute the `fsck` command manually to update the superblock information.

The mount also fails if the file system at one time contained a large file, even though it was moved or removed. You must execute the `fsck` command to clear the old information and allow the file system to be mounted.

## Mounting All File Systems Manually

The `/etc/vfstab` file is read by the `/usr/sbin/mountall` command during the system boot sequence and mounts all file systems that have a `yes` in the `mount at boot` field.

The root user can use the `mountall` command to mount manually every file system in the `/etc/vfstab` file that has a `yes` in the `mount at boot` field, for example:

```
mountall
```

To mount only the local file systems listed in the `/etc/vfstab` file, execute:

```
mountall -l
```

During the boot sequence, the `fsck` utility checks each local file system in the `/etc/vfstab` file that has a device to `fsck` entry and an `fsck` pass number greater than 0. The utility determines if the file system is in a usable state to be safely mounted.

If the `fsck` utility determines that the file system is in an unusable state (for example, corrupted), the `fsck` utility will only perform non-interactive repairs before attempting to mount the file system. If the repairs require interaction, the boot process will take the system to single-user mode to allow you to run `fsck` manually.

The system attempts to mount any local file systems that have a `-` (dash) or `0` (zero) entry in the `fsck` `pass` field without checking the file system itself.

## Mounting a New File System

To add a new disk to the system, prepare the disk to hold a file system, and mount the file system, perform these general steps:

1. Set up the disk hardware, which might include setting address switches and connecting cables.
  2. Perform a reconfiguration boot or run the `devfsadm` utility to add support for the new disk.
  3. Use the `format` utility to define one or more slices. On x86/x64 systems, create a Solaris `fdisk` partition, and then create one or more slices within the Solaris partition.
  4. Create a new file system on one slice by using the `newfs` command.
  5. Create a mount point for the file system by using the `mkdir` command to create a new directory in the `/` (root) file system.
- ```
# mkdir /data
```
6. Mount the new file system manually by using the `mount` command.
- ```
mount /dev/dsk/c1t3d0s7 /data
```
7. Use the `mount` command to determine if the file system is mounted.

```
mount
```

```
...
```

```
/data on /dev/dsk/c1t3d0s7
read/write/setuid/devices/intr/largefiles/logging/xattr/onerror=panic/dev
=800027 on Sun Oct 24 11:55:34 2004
```

8. Edit the `/etc/vfstab` file to add a line entry for the new file system.

```
vi /etc/vfstab
fd - /dev/fd fd - no -
/proc - /proc proc - no -
/dev/dsk/c0t0d0s1 - - swap - no -
/dev/dsk/c0t0d0s0 /dev/rdsk/c0t0d0s0 / ufs 1 no -
/dev/dsk/c0t0d0s6 /dev/rdsk/c0t0d0s6 /usr ufs 1 no -
/dev/dsk/c0t0d0s3 /dev/rdsk/c0t0d0s3 /var ufs 1 no -
/dev/dsk/c0t0d0s7 /dev/rdsk/c0t0d0s7 /export/home ufs 2 yes -
/devices - /devices devfs - no -
ctfs - /system/contract ctfss - no -
objfs - /system/object objfs - no -
swap - /tmp tmpfs - yes -
/dev/dsk/c1t3d0s7 /dev/rdsk/c1t3d0s7 /data ufs 2 yes -
#
```

The file system automatically mounts whenever the system boots.

## Mounting Different Types of File Systems

Different file system types have unique properties that affect how the mount command functions.

By default, the mount command assumes it is mounting a ufs-type file system. However, when you are mounting a different type of file system, you might have to specify its type on the command line.

You use the `-F` option with the mount command to specify the type of file system mounted. The file system type must be determinable from the `/etc/vfstab`, `/etc/default/fs`, or `/etc/dfs/fstypes` files.

### Determining a File System's Type

Because the mount command needs the file system type to function properly, the file system type must be explicitly specified or determined by searching the following files:

- The `/etc/vfstab` file for the FS type field
- The `/etc/default/fs` file for a local file system type
- The `/etc/dfs/fstypes` file for a remote file system type

If the file system's type has not been explicitly specified on the command line using the mount `-F FStype` option, the mount command examines the `/etc/vfstab` file to determine the file system's type. The mount command makes this determination by using the file system's block device name, raw device name, or mount point directory name.

If the mount command cannot determine the file system's type by searching the `/etc/vfstab` file, the mount command uses the default file system type specified in either the `/etc/default/fs` file or the `/etc/dfs/fstypes` file, depending on whether the file system is local or remote.

The default local file system type is specified in the /etc/default/fs file by the line entry LOCAL=*fstype*.

```
LOCAL=ufs
```

The first line entry in the /etc/dfs/fstypes file determines the default remote file system type.

```
nfs NFS Utilities
autofs AUTOFS Utilities
cachefs CACHEFS Utilities
```

## Using the `fstyp` Command

You can also use the `fstyp` command with the raw device name of the disk slice to determine a file system's type.

```
fstyp /dev/rdsk/c0t0d0s7
ufs
```

## Specifying a ZFS File System Type

To mount a ZFS file system using the `mount` command, as the `root` user, perform the command:

```
mount -F zfs rzpool/users/jf /home/jf
```

In this example, the file system type is `zfs`, the file system is named `rzpool/users/jf` and resides within the ZFS pool called `rzpool`. The mount point `/home/jf` is an existing directory in the Solaris OS.



**Note** – To use the `mount` command to mount a ZFS file system, you must set the `mountpoint` property for the ZFS file system to the `legacy` value. By default, ZFS file systems mount automatically when you create them. See the `zfs(1M)` command man page for more information.

## Specifying an hsfs File System Type

To mount a file system that resides on a CD-ROM when the Volume Management (vold) services are stopped, as the `root` user, perform the command:

```
mount -F hsfs -o ro /dev/dsk/c0t6d0s0 /cdrom
```

In this example, the file system type is `hsfs`, the file system resides on disk slice `/dev/dsk/c0t6d0s0`, and the mount point `/cdrom` is an existing directory in the Solaris OS.

## Specifying a pcfs File System Type

To mount a file system that resides on a diskette when the Volume Management (vold) services are stopped, perform the commands:

```
mkdir /pcfs
mount -F pcfs /dev/diskette /pcfs
```

In this example, the file system type is `pcfs`. This file system resides on the device `/dev/diskette`, and the mount point is `/pcfs`.

# Managing ZFS Mount Points

By default, all ZFS file systems are mounted by ZFS at boot by using the SMF svc://system/filesystem/local service. File systems are mounted under `/path`, where `path` is the name of the file system.

You can override the default mount point by setting the `mountpoint` property to a specific path by using the `zfs set` command. ZFS automatically creates this mount point, if needed, and automatically mounts this file system when the `zfs mount` –a command is invoked, without requiring you to edit the `/etc/vfstab` file.

The `mountpoint` property is inherited. For example, if `pool/home` has its `mountpoint` set to `/export/stuff`, then `pool/home/user` inherits `/export/stuff/user` for its `mountpoint` property.

The `mountpoint` property can be set to `none` to prevent the file system from being mounted.

If desired, file systems can also be explicitly managed through legacy mount interfaces by setting the `mountpoint` property to `legacy` by using `zfs set`. Doing so prevents ZFS from automatically mounting and managing this file system. Legacy tools including the `mount` and `umount` commands, and the `/etc/vfstab` file must be used instead.

When changing mount point management strategies, the following behaviors apply:

- Automatic mount point behavior
- Legacy mount point behavior

## Automatic Mount Points

- When changing from `legacy` or `none`, ZFS automatically mounts the file system.
- If ZFS is currently managing the file system but it is currently unmounted, and the `mountpoint` property is changed, the file system remains unmounted.

You can also set the default mount point for the root dataset at pool creation time by using `zpool create` with the `-m` option.

Any dataset whose mountpoint property is not legacy is managed by ZFS. In the following example, a dataset is created whose mount point is automatically managed by ZFS.

```
zfs create pool/filesystem
zfs get mountpoint pool/filesystem
NAME PROPERTY VALUE SOURCE
pool/filesystem mountpoint /pool/filesystem default
zfs get mounted pool/filesystem
NAME PROPERTY VALUE SOURCE
pool/filesystem mounted yes -
```

You can also explicitly set the mountpoint property as shown in the following example:

```
zfs set mountpoint=/mnt pool/filesystem
zfs get mountpoint pool/filesystem
NAME PROPERTY VALUE SOURCE
pool/filesystem mountpoint /mnt local
zfs get mounted pool/filesystem
NAME PROPERTY VALUE SOURCE
pool/filesystem mounted yes -
```

When the mountpoint property is changed, the file system is automatically unmounted from the old mount point and remounted to the new mount point. Mount point directories are created as needed. If ZFS is unable to unmount a file system due to it being active, an error is reported and a forced manual unmount is necessary.

## Legacy Mount Points

Manage ZFS file systems with legacy tools by setting the mountpoint property to legacy. Manage the legacy file systems through the mount and umount commands and the /etc/vfstab file. ZFS does not automatically mount legacy file systems on boot, and the zfs mount and zfs umount commands do not operate on datasets of this type. The following examples show how to set up and manage a ZFS dataset in legacy mode:

```
zfs set mountpoint=legacy tank/home/eschrock
mount -F zfs tank/home/eschrock /mnt
```

If you set up separate ZFS /usr or /var file systems, you must indicate that they are legacy file systems. You must mount them by creating entries in the /etc/vfstab file. Otherwise, the system/filesystem/local service enters maintenance mode when the system boots.

To automatically mount a legacy file system on boot, you must add an entry to the /etc/vfstab file. The following example shows what the entry in the /etc/vfstab file might look like:

| #device<br>mount     | device  | mount | FS  | fsck   | mount        |
|----------------------|---------|-------|-----|--------|--------------|
| #to mount<br>options | to fsck | point |     | type   | pass at boot |
| #                    |         |       |     |        |              |
| tank/home/eschrock - | /mnt    |       | zfs | - yes- |              |

The device to fsck and fsck pass entries are set to -. This syntax is used because the fsck command is not applicable to ZFS file systems.

## Mounting a ZFS File System

ZFS automatically mounts file systems when file systems are created or when the system boots. Use the zfs mount command only when you are changing mount options or explicitly mounting or unmounting file systems.

The zfs mount command with no arguments shows all currently mounted file systems that are managed by ZFS. Legacy managed mount points are not displayed. For example:

```
zfs mount
tank /tank
tank/home /tank/home
tank/home/bonwick /tank/home/bonwick
tank/ws /tank/ws
```

Use the -a option to mount all ZFS managed file systems. Legacy managed file systems are not mounted. For example:

```
zfs mount -a
```

## Managing ZFS Mount Points

---

By default, ZFS doesn't allow mounting on top of a non-empty directory. To force a mount on top of a non-empty directory, use the -O option. For example:

```
zfs mount tank/home/lalt
cannot mount '/export/home/lalt': directory is not empty
use legacy mountpoint to allow this behavior, or use the -O flag
zfs mount -O tank/home/lalt
```

Use legacy tools to manage legacy mount points. If you use ZFS tools, you will get errors. For example:

```
zfs mount pool/home/billm
cannot mount 'pool/home/billm': legacy mountpoint
use mount(1M) to mount this filesystem
mount -F zfs tank/home/billm
```

# Performing Unmounts

A file system is commonly unmounted if it needs to be checked and repaired by the `fsck` command, or if it needs to be backed up completely.

## Unmounting a File System

Some file system administration tasks cannot be performed on mounted file systems.

To unmount a file system to prepare it for system maintenance, use the `umount` command. Unmounting a file system by using the `umount` command detaches it from the file system mount point and deletes its entry from the `/etc/mnttab` file system table.

---

**Note** – Notify users before unmounting a file system that they are currently accessing.

---



To unmount a file system manually by using the directory mount point, perform the command:

```
umount /export/home
```

To unmount a file system manually by using the logical disk device name, perform the command:

```
umount /dev/dsk/c0t0d0s7
```

## Unmounting All File Systems

The `/etc/mnttab` file system table is read by the `/usr/sbin/umountall` command during the system shutdown sequence or when `umountall` is invoked from the command line. The `umountall` unmounts all file systems specified in the `vfstab` file except `/ (root)`, `/usr`, `/proc`, `/dev/fd`, `/var`, `/var/run`, and `/tmp`.

Run the `umountall` command as the `root` user when you want to unmount manually all the file systems listed in `/etc/mnttab`, for example:

```
umountall
```

## Performing Unmounts

---

To unmount only the local file systems listed in /etc/mnttab, perform the command:

```
umountall -1
```

To verify that a file system or a number of file systems have been unmounted, invoke the mount command and check the output.

## Unmounting a Busy File System

Any file system that is busy is not available for unmounting. Both the umount and umountall commands display the error message:

```
umount: file_system_name busy
```

A file system is considered to be busy if one of the following conditions exists:

- A program is accessing a file or directory in the file system
- A user is accessing a directory or file in the file system
- A program has a file open in the file system
- The file system is being shared

Aside from eliminating the conditions above, there are two methods to make a file system available for unmounting if it is busy:

- fuser command – Lists all of the processes that are accessing the file system and kills them if necessary
- umount -f command – Forces the unmount of a file system

---

**Note** – The fuser command displays the process IDs of all processes currently using the specified file system. Each process ID is followed by a letter code. These letter codes are described in the man page for this command.



## Using the fuser Command

To stop all processes that are currently accessing a file system, follow these steps:

1. As the root user, list all of the processes that are accessing the file system. Use the following command to identify which processes need to be terminated.

```
fuser -cu mount_point
```

This command displays the name of the file system and the user login name for each process currently active in the file system.

2. Kill all processes accessing the file system.

```
fuser -ck mount_point
```

A SIGKILL message is sent to each process that is using the file system.

3. Verify that there are no processes accessing the file system.

```
fuser -c mount_point
```

4. Unmount the file system.

```
umount mount_point
```

## Using the umount -f Command

As the root user, you can use the umount command with the -f (force) option to unmount a file system even if it is busy. The following is the format for this command:

```
umount -f mount_point
```

This command unmounts the file system even if it contains open files. A forced unmount can result in loss of data and in zombie processes that are left running on the system. However, it is particularly useful for unmounting a shared file system if the remote file server is nonfunctional.

## Unmounting a ZFS File System

Unmount ZFS file systems with the `zfs unmount` subcommand. The `unmount` command can take either the mount point or the file system name as arguments.

In the following example, a file system is unmounted by file system name:

```
zfs unmount tank/home/tabriz
```

In the following example, the file system is unmounted by mount point:

```
zfs unmount /export/home/tabriz
```

The `unmount` command fails if the file system is active or busy. To forcibly unmount a file system, you can use the `-f` option. Be cautious when forcibly unmounting a file system, if its contents are actively being used. Unpredictable application behavior can result.

```
zfs unmount tank/home/eschrock
cannot unmount '/export/home/eschrock': Device busy
zfs unmount -f tank/home/eschrock
```

For backward compatibility, use the legacy `umount` command to unmount a ZFS file system. For example:

```
umount /export/home/bob
```

## Repairing Important Files if Boot Fails on a SPARC System

The following procedure describes how to boot from a Solaris OS DVD to edit a misconfigured /etc/vfstab file.

1. Insert the Solaris 10 OS Software DVD into the CD-ROM drive.
2. Execute a single-user boot from the DVD.

```
ok boot cdrom -s
Boot device: /pci@1e,600000/ide@d/cdrom@2,0:f File and args -s
<<Output omitted>>
Booting to milestone "milestone/single-user:default"
Configuring devices.
Using RPC Bootparams for network configuration information.
Attempting to configure interface bge0...
Skipped interface bge0
Requesting System Maintenance Mode
SINGLE USER MODE
#
```



**Note** – Performing a single-user boot operation from this software CD-ROM creates an *in-memory* copy of the /root file system, which supports your ability to perform administrative tasks.

3. Use the fsck command on the faulty / (root) partition to check and repair any potential problems in the file system and make the device writable.
 

```
fsck /dev/rdsk/c0t0d0s0
```
4. If the fsck command is successful, mount the / (root) file system on the /a directory to gain access to the file system on disk.
 

```
mount /dev/dsk/c0t0d0s0 /a
```
5. Set and export the TERM variable, which enables the vi editor to work properly.
 

```
TERM=sun
export TERM
```
6. Edit the /etc/vfstab file, and correct any problems. Then exit the file.
 

```
vi /a/etc/vfstab
:wq!
```

7. Unmount the file system.

```
cd /
umount /a
```

8. Reboot the system.

```
init 6
```

## Repairing Files if Boot Fails on x86/x64 Systems

The following procedure describes how to boot from the Solaris OS DVD to edit a misconfigured /etc/vfstab file.

1. Power on the system and insert the Solaris 10 OS Software DVD into the DVD-ROM drive.

**Note** – Your system needs to be configured to boot from the DVD. This might need to be set in your BIOS settings before booting.

2. The system detects the DVD, loads the grand unified boot (GRUB) loader, and presents the GRUB boot menu. Three boot choices are displayed:

Solaris

Solaris Serial Console ttya

Solaris Serial Console ttvb (for lx50, v60, and v65x)

The first boot choice (Solaris) is the default, and it boots the system from a DVD. With the Solaris choice the system automatically boots after 60 seconds. Press Return to boot immediately from this source.

3. The system presents the following list of installation methods and possible tasks to perform. Press the space bar to pause at this menu.

1.Solaris Interactive (default)

2.Custom Jump-start

3.Solaris Interactive Text (Desktop session)

4.Solaris Interactive Text (Console session)

5.Apply driver updates

6.Single user shell



4. Enter 6 to select the single user shell. The system searches for installed Solaris OS instances, and asks if you want to mount the root file system on the /a mount point. For example:

Searching for installed OS instances...

Solaris 10 10/09 s10x\_u8wos\_08a X86 was found on /dev/dsk/c1d0s0.  
Do you wish to have it mounted read-write on /a? [y,n,?]

5. Enter n to choose not to mount the file system.

Do you wish to have it mounted read-write on /a? [y,n,?] **n**

Starting shell.

#

6. Use the fsck command on the faulty / (root) partition to check and repair any potential problems in the file system and make the device writable.

**# fsck /dev/rdsck/c1d0s0**

7. If the fsck command is successful, mount the / (root) file system on the /a directory to gain access to the file system.

**# mount /dev/dsk/c1d0s0 /a**

8. Set and export the TERM variable, which enables the vi editor to work properly.

**# TERM=sun**

**# export TERM**

9. Edit the /etc/vfstab file, and correct any problems. Then exit the file.

**# vi /a/etc/vfstab**

**:wq!**

10. Unmount the file system.

**# cd /**

**# umount /a**

11. Eject the DVD.

12. Reboot the system.

**# init 6**

## Accessing Mounted Diskettes, CD-ROMs, or DVDs

To provide access to file systems on diskettes and CD-ROMs, the Solaris OS provides the volume management service standard interface.



---

**Note** – The Solaris 10 OS includes support for additional removable media such as DVDs, Jaz drives, and Zip drives. For more information on using these devices, see the resources available on the Solaris 10 Documentation CD or visit <http://docs.sun.com> to access online documentation.

---

The volume management service provides two major benefits:

- It automatically mounts removable media for both the root user and non-root users.
- It can give other systems on the network automatic access to any removable media currently inserted in the local system.

The volume management service is controlled by the /usr/sbin/vold daemon. With a default install, this service automatically manages diskettes and CD-ROMs for regular users.

The volume management service features automatic DVD and CD-ROM detection. It doesn't detect the presence of a diskette that has been inserted in the drive until the volcheck command is run. This command instructs the vold daemon to check the diskette drive for any inserted media. The volume management service can mount ufs, pcfs, hsfs, and udfs file systems.

## Using the Volume Management Service

To make working with diskettes, DVDs, and CD-ROMs simple, each device is easy to mount and mounts at an easy-to-remember location.

If the `vold` daemon detects that the mounted device contains a file system, then the device is mounted at the directory location.

Table 6-1 lists the directory locations of mounted devices that contain file systems.

**Table 6-1** Directory Locations

| Media Device              | Access File Systems On |
|---------------------------|------------------------|
| First diskette drive      | /floppy/floppy0        |
| First CD-ROM or DVD drive | /cdrom/cdrom0          |
| First Jaz drive           | /rmdisk/jaz0           |
| First Zip drive           | /rmdrive/zip0          |
| First PCMCIA card         | /pcmem0                |

If the `vold` daemon detects that the mounted device does not contain a file system, the device is accessible through a path.

Table 6-2 lists the paths for mounted devices that don't contain file systems.

**Table 6-2** Paths for Accessing Devices

| Media Device              | Access Raw Device On     |
|---------------------------|--------------------------|
| First diskette drive      | /vol/dev/aliases/floppy0 |
| First CD-ROM or DVD drive | /vol/dev/aliases/cdrom0  |
| First Jaz drive           | /vol/dev/aliases/jaz0    |
| First Zip drive           | /vol/dev/aliases/zip0    |
| First PCMCIA card         | /vol/dev/aliases/pcmem0  |

When the Volume Management service is running on the system, a regular user can easily access a diskette or CD-ROM by following these basic steps:

1. Insert the media.
2. For diskettes only, enter the `volcheck` command.
3. Use the `cd` command to change to the directory of the mounted volume.
4. Work with files on the media.
5. Use the `cd` command to leave the directory structure of the mounted volume.
6. Eject the media.

Table 6-3 shows the configuration files used by the Volume Management service.

**Table 6-3** Volume Management Service Configuration Files

| File              | Description                                                                                                                                                                                                      |
|-------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| /etc/vold.conf    | The vold configuration file. This file defines items, such as what action should be taken when media is inserted or ejected, which devices are managed by vold, and which file system types are unsafe to eject. |
| /etc/rmmount.conf | The rmmount command configuration file. The rmmount command is a removable media mounter that is executed by the vold daemon whenever a CD-ROM or diskette is inserted.                                          |

## Restricting Access to Mounted Diskettes, CD-ROMs, or DVDs

To restrict regular users from accessing diskettes or CD-ROMs on the system, you can, as the root user, disable the Volume Management service.

### Stopping Volume Management (vold)

To stop Volume Management service from running on a system, as the root user, perform the command:

```
svcadm disable svc:/system/filesystem/volfs:default
```

To restart the Volume Management service, as the root user, perform the command:

```
svcadm enable svc:/system/filesystem/volfs:default
```

Legacy syntax to start and stop the Volume Management service is provided in Solaris 10. The /etc/init.d/volmgt script calls the svcadm command with the arguments required to stop and start the vold daemon. You can use the following legacy commands to stop and start the Volume Management service:

```
/etc/init.d/volmgt stop
/etc/init.d/volmgt start
```

## Troubleshooting Volume Management Service

If a CD-ROM fails to eject from the drive, as the `root` user, attempt to stop Volume Management service. If this is unsuccessful, stop the `vold` daemon.

```
svcadm disable svc:/system/filesystem/volfs:default
```

or as a last resort:

```
pkill -9 vold
```

Push the button on the system to eject the CD-ROM. The CD-ROM tray ejects. Remove the CD-ROM, and leave the tray out. Then enable the Volume Management service.

```
svcadm enable svc:/system/filesystem/volfs:default
```

Wait a few seconds, and then push the CD-ROM tray back into the drive.

# Accessing a Diskette, CD-ROM, or DVD Without VMS

When the `vold` daemon is not running, only the `root` user can mount and access a diskette or CD-ROM. Follow these steps:

1. Insert the media device.
2. Become the `root` user.
3. Create a mount point, if necessary.
4. Determine the file system type.
5. Mount the device by using the mount options listed in the following sections.
6. Work with files on the media device.
7. Unmount the media device.
8. Eject the media device.
9. Exit the `root` session.

## Using the `mount` Command

To mount a file system that resides on a DVD or CD-ROM when the Volume Management service is stopped, as the `root` user, perform the command:

```
mount -F hsfs -o ro /dev/dsk/c0t2d0s2 /cdrom
```

In this example, the file system type is `hsfs`, the file system resides on CD-ROM slice `/dev/dsk/c0t2d0s2`, and the mount point `/cdrom` is an existing directory in the Solaris OS.




---

**Note** – Different systems use different names for their DVD or CD-ROM devices. The example above is from a Sun Blade 1500.

---

To mount a file system that resides on a diskette when the Volume Management service is stopped, as the `root` user, perform the command:

```
mkdir /pcfs
mount -F pcfs /dev/diskette /pcfs
```

In this example, the file system type is `pcfs`. This file system resides on the `/dev/diskette` device, and the mount point used is `/pcfs`.

# Exercise: Mounting File Systems

In this exercise, you complete the following tasks:

- Create mount points
- Mount file systems
- Specify mount options

## Preparation

This exercise requires a spare disk that contains 4 unmounted UFSs on slices 0, 1, 3, and 4 of the spare disk on your system. If you have to partition this disk for this exercise, you must understand how to use the format utility. If you need to make file systems on the slices for this exercise, you must understand how to use the newfs command.

This exercise applies to both SPARC and x86/x64 systems. These systems typically have different disk device names. For example:

- The spare disk of a Sun Blade 1500 (SPARC) is typically c0t1d0.
- The spare disk of an Ultra 20 (x86/x64) is typically c2d0.

Throughout this exercise, use the correct device names for your system.

## Task

Complete the following steps:

1. Log in as the root user and open a terminal window.
2. Use the prtvtoc command to determine if the required slices exist on your spare disk.
3. Use the fstyp command to determine if a ufs file system exists on the required slices on your spare disk.
4. If the spare disk is not divided into 4 slices as described in the “Preparation” section of this exercise, complete the following steps:
  - a. On x86/x64 systems, use the fdisk menu in the format utility to create define a single Solaris fdisk partition that uses the entire disk. Make sure the partition is not marked as *active*.

- Unauthorized reproduction or distribution prohibited. Copyright© 2011, Oracle and/or its affiliates.
- b. Use the partition menu in the format utility to create the required slices.
  - c. Make slices 0, 1, and 3 exactly the same size (approximately 25 percent of the total disk space each), and use slice 4 for the remainder of the available space. Using the All Free Hog method and choosing slice 4 as the free hog slice makes this easy.

On SPARC systems, you can also use the Solaris Management Console to partition the disk.

- d. Save your changes and quit the format utility when you are finished.

Example of a correct SPARC partition table:

| Part | Tag        | Flag | Cylinders     | Size     | Blocks                |
|------|------------|------|---------------|----------|-----------------------|
| 0    | root       | wm   | 0 - 13878     | 27.00GB  | (13879/0/0) 56626320  |
| 1    | swap       | wu   | 13879 - 27757 | 27.00GB  | (13879/0/0) 56626320  |
| 2    | backup     | wu   | 0 - 57458     | 111.79GB | (57459/0/0) 234432720 |
| 3    | unassigned | wm   | 27758 - 41636 | 27.00GB  | (13879/0/0) 56626320  |
| 4    | unassigned | wm   | 41637 - 57458 | 30.78GB  | (15822/0/0) 64553760  |
| 5    | unassigned | wm   | 0             | 0        | (0/0/0) 0             |
| 6    | usr        | wm   | 0             | 0        | (0/0/0) 0             |
| 7    | unassigned | wm   | 0             | 0        | (0/0/0) 0             |

Example of a correct x86/x64 partition table:

| Part | Tag        | Flag | Cylinders    | Size    | Blocks                |
|------|------------|------|--------------|---------|-----------------------|
| 0    | root       | wm   | 3 - 2483     | 19.01GB | (2481/0/0) 39857265   |
| 1    | swap       | wu   | 2484 - 4964  | 19.01GB | (2481/0/0) 39857265   |
| 2    | backup     | wu   | 0 - 10007    | 76.67GB | (10008/0/0) 160778520 |
| 3    | unassigned | wm   | 4965 - 7445  | 19.01GB | (2481/0/0) 39857265   |
| 4    | unassigned | wm   | 7446 - 10007 | 19.63GB | (2562/0/0) 41158530   |
| 5    | unassigned | wm   | 0            | 0       | (0/0/0) 0             |
| 6    | usr        | wm   | 0            | 0       | (0/0/0) 0             |
| 7    | unassigned | wm   | 0            | 0       | (0/0/0) 0             |
| 8    | boot       | wu   | 0 - 0        | 7.84MB  | (1/0/0) 16065         |
| 9    | alternates | wm   | 1 - 2        | 15.69MB | (2/0/0) 32130         |

- 5. If the 4 slices on the spare disk do not have UFSs on them as described in the “Preparation” section of this exercise, create the required file systems.
- 6. Use the mount command to list the file systems that are currently mounted on your system.

What are the default mount options applied to the / (root) file system?

---



---

## Exercise: Mounting File Systems

---

7. Create a directory named /morespace to use as a mount point.
8. Complete the following steps:
  - a. Mount the file system on slice 4 of your spare disk to the /morespace directory.
  - b. Record the default mount options that were applied to this mount.

---

---

9. Change to the /morespace directory, and then create a new file that has one line of content.
10. Display a long listing for this file, and then record the time value it reports.

This time value represents when the file was last modified.

---

11. Add the -u option to the ls command to show when the file was last accessed, and then record the time value it reports.

This time value is updated whenever you read the file.

---

12. Wait one minute or more and complete the following steps:
    - a. Use the cat command to display the file.
    - b. Again, check and record the access time.
- The access time should differ from the access time indicated in the previous step.

---

13. Change to the / (root) directory and complete the following steps:
  - a. Unmount the /morespace file system.
  - b. Mount the same file system to the /morespace directory, but add the option that prevents update of access time values.
  - c. Verify that the noatime option is now set for the /morespace mount.

14. Return to the /morespace file system and complete the following steps:

- Use the cat command to display your test file.
- Again, check and record the access time.

The access time should match the access time that existed before you mounted the /morespace file system with the noatime option.

- 
15. Use the vi editor to add a line to the /etc/vfstab file to make the mount for the /morespace file system happen when you boot the system.
  16. Reboot your system, and complete the following steps:
    - Log in as the root user.
    - Open a terminal window.
    - Use the mount command to verify that the /morespace file system is mounted, and the noatime option is set.
  17. Create a directory called /dir0 and mount the file system that resides on slice 0 of your spare disk as /dir0.
  18. Create a directory called /dir0/dir1 and mount the file system that resides on slice 1 of your spare disk as /dir0/dir1.
  19. Open a second terminal window and change the directory in this second window to /dir0/dir1.
  20. In your original terminal window, attempt to unmount the file system mounted below the /dir0/dir1 directory.

What message is displayed? Does the file system unmount?

---



**Note –** To discover why you could not unmount the file system, use the fuser -cu /dir0/dir1 command. The fuser command should show the process ID of the shell currently using the filesystem.

21. In your original terminal window, add the -f option to the umount command, and attempt to unmount the file system mounted below the /dir0/dir1 directory.

Which message displays? Does the file system unmount?

---

## Exercise: Mounting File Systems

---



22. In the second terminal window, complete the following steps:
  - a. Attempt to determine your current working directory.  
Which message displays?
  - b. Change the directory to / (root), and then verify that the `pwd` command works.

---

**Note** – Different shells display different messages. Users of the BASH shell see the correct directory.

---

23. Unmount the file systems mounted below `/dir0` and `/morespace`.
24. Use the `vi` editor to remove the line from `/etc/vfstab` that mounts a file system below `/morespace`.

## Exercise Summary



**Discussion** – Take a few minutes to discuss what experiences, issues, or discoveries you had during the lab exercises.

- Experiences
- Interpretations
- Conclusions
- Applications

# Exercise Solutions: Mounting File Systems

This section contains solutions to the exercise.

## Task

Complete the following steps:

1. Log in as the root user and open a terminal window.
2. Use the prtvtoc command to determine if the required slices exist on your spare disk. For example:  

```
prtvtoc /dev/rdsk/c2d0s2
```
3. Use the fstyp command to determine if a ufs file system exists on the required slices on your spare disk. For example:  

```
fstyp /dev/dsk/c2d0s0
ufs
#
```
4. If the spare disk is not divided into four slices as described in the Preparation section of this exercise, complete the following steps:
  - a. On x86/x64 systems, use the fdisk menu in the format utility to define a single Solaris fdisk partition that uses the entire disk. Make sure the partition is not marked as *active*.
  - b. Use the partition menu in the format utility to create the required slices.
  - c. Make slices 0, 1, and 3 exactly the same size (approximately 25 percent of the total disk space each), and use slice 4 for the remainder of the available space. Using the All Free Hog method and choosing slice 4 as the free hog slice makes this easy.

On SPARC systems, you can also use the Solaris Management Console to partition the disk.

- d. Save your changes and quit the format utility when you are finished.

Example of a correct SPARC partition table:

| Part | Tag        | Flag | Cylinders     | Size     | Blocks                |
|------|------------|------|---------------|----------|-----------------------|
| 0    | root       | wm   | 0 - 13878     | 27.00GB  | (13879/0/0) 56626320  |
| 1    | swap       | wu   | 13879 - 27757 | 27.00GB  | (13879/0/0) 56626320  |
| 2    | backup     | wu   | 0 - 57458     | 111.79GB | (57459/0/0) 234432720 |
| 3    | unassigned | wm   | 27758 - 41636 | 27.00GB  | (13879/0/0) 56626320  |
| 4    | unassigned | wm   | 41637 - 57458 | 30.78GB  | (15822/0/0) 64553760  |
| 5    | unassigned | wm   | 0             | 0        | (0/0/0) 0             |
| 6    | usr        | wm   | 0             | 0        | (0/0/0) 0             |
| 7    | unassigned | wm   | 0             | 0        | (0/0/0) 0             |

Example of a correct x86/x64 partition table:

| Part | Tag        | Flag | Cylinders    | Size    | Blocks                |
|------|------------|------|--------------|---------|-----------------------|
| 0    | root       | wm   | 3 - 2483     | 19.01GB | (2481/0/0) 39857265   |
| 1    | swap       | wu   | 2484 - 4964  | 19.01GB | (2481/0/0) 39857265   |
| 2    | backup     | wu   | 0 - 10007    | 76.67GB | (10008/0/0) 160778520 |
| 3    | unassigned | wm   | 4965 - 7445  | 19.01GB | (2481/0/0) 39857265   |
| 4    | unassigned | wm   | 7446 - 10007 | 19.63GB | (2562/0/0) 41158530   |
| 5    | unassigned | wm   | 0            | 0       | (0/0/0) 0             |
| 6    | usr        | wm   | 0            | 0       | (0/0/0) 0             |
| 7    | unassigned | wm   | 0            | 0       | (0/0/0) 0             |
| 8    | boot       | wu   | 0 - 0        | 7.84MB  | (1/0/0) 16065         |
| 9    | alternates | wm   | 1 - 2        | 15.69MB | (2/0/0) 32130         |

5. If the four slices on the spare disk do not have ufs file systems on them as described in the Preparation section of this exercise, create the required file systems.

# **newfs /dev/rdsk/c2d0s0**

6. Use the mount command to list the file systems that are currently mounted on your system.

What are the default mount options applied to the / (root) file system?

# **mount**

```
/ on /dev/dsk/c1d0s0
read/write/setuid/devices/intr/largefiles/logging/xattr/onerror=panic/dev
=1980000 on Tue Feb 6 16:54:32 2007
(output omitted)
```

*The mount options should match those in this example, and your dev-number depends on the devices in use on your system.*

7. Create a directory named /morespace to use as a mount point.

# **mkdir /morespace**

8. Complete the following steps:
- Mount the file system on slice 4 of your spare disk to the /morespace directory.
  - Record the default mount options that were applied to this mount.

```
mount /dev/dsk/c2d0s4 /morespace
mount
(output omitted)
/morespace on /dev/dsk/c2d0s4
read/write/setuid/devices/intr/largefiles/logging/xattr/onerror=panic/dev
=1980044 on Sun Feb 11 17:20:24 2007
```

- Change to the /morespace directory, and then create a new file that has one line of content.

```
cd /morespace
echo "Some Text" > testfile
#
```

- Display a long listing for this file, and then record the time value it reports.

This time value represents when the file was last modified.

```
ls -l
total 18
drwx----- 2 root root 8192 Feb 8 17:42 lost+found
-rw-r--r-- 1 root root 10 Feb 11 17:21 testfile
#
```

- Add the -u option to the ls command to show when the file was last accessed, and then record the time value it reports.

This time value is updated whenever you read the file.

```
ls -lu testfile
-rw-r--r-- 1 root root 10 Feb 11 17:21 testfile
#
```

12. Wait one minute or more and complete the following steps:

- Use the cat command to display the file.
- Again, check and record the access time.

The access time should differ from the access time indicated in the previous step.

```
cat testfile
Some Text
ls -lu testfile
-rw-r--r-- 1 root root 10 Feb 11 17:23 testfile
#
```

13. Change to the / (root) directory and complete the following steps:

- Unmount the /morespace file system.
- Mount the same file system to the /morespace directory, but add the option that prevents update of access time values.
- Verify that the noatime option is now set for the /morespace mount.

```
cd /
umount /morespace
mount -o noatime /dev/dsk/c2d0s4 /morespace
mount
(output omitted)
/morespace on /dev/dsk/c2d0s4
read/write/setuid/devices/intr/largefiles/logging/xattr/noatime/onerror=panic/dev=1980044 on Sun Feb 11 17:24:50 2007
#
```

14. Return to the /morespace file system and complete the following steps:

- Use the cat command to display your test file.
- Again, check and record the access time.

The access time should match the access time that existed before you mounted the /morespace file system with the noatime option.

```
cd /morespace
cat testfile
Some Text
ls -lu testfile
-rw-r--r-- 1 root root 10 Feb 11 17:23 testfile
#
```

15. Use the vi editor to add a line to the /etc/vfstab file to make the mount for the /morespace file system happen when you boot the system.

```
/dev/dsk/c2d0s4 /dev/rdsck/c2d0s4 /morespace ufs 1 yes noatime
```

16. Reboot your system, and complete the following steps:
- Log in as the root user.
  - Open a terminal window.
  - Use the mount command to verify that the /morespace file system is mounted, and the noatime option is set.

```
init 6
(reboot messages & login prompts display)
mount
(output omitted)
/morespace on /dev/dsk/c2d0s4
read/write/setuid/devices/intr/largefiles/logging/xattr/noatime/onerror=panic/dev=1980044 on Sun Feb 11 17:32:12 2007
```

17. Create a directory called /dir0 and mount the file system that resides on slice 0 of your spare disk as /dir0.

```
mkdir /dir0
mount /dev/dsk/c2d0s0 /dir0
#
```

18. Create a directory called /dir0/dir1 and mount the file system that resides on slice 1 of your spare disk as /dir0/dir1.

```
mkdir /dir0/dir1
mount /dev/dsk/c2d0s1 /dir0/dir1
```

19. Open a second terminal window and change the directory in this second window to /dir0/dir1.

```
cd /dir0/dir1
```

20. In your original terminal window, attempt to unmount the file system mounted below the /dir0/dir1 directory.

What message is displayed? Does the file system unmount?

```
umount /dev/dsk/c2d0s1
umount: /dir0/dir1 busy
mount
(output omitted)
/dir0/dir1 on /dev/dsk/c2d0s1
read/write/setuid/devices/intr/largefiles/logging/xattr/onerror=panic/dev=1980041 on Sun Feb 11 17:46:38 2007
```



*The file system does not unmount.*

---

**Note** – To discover why you could not unmount the file system, use the fuser -cu /dir0/dir1 command. The fuser command should show the process ID of the shell currently using the filesystem.

---

21. In your original terminal window, add the -f option to the umount command, and attempt to unmount the file system mounted below the /dir0/dir1 directory.

Which message displays? Does the file system unmount?

```
umount -f /dir0/dir1
```

No message displays. The file system unmounts.

22. In the second terminal window, complete the following steps:

- a. Attempt to determine your current working directory.

Which message displays?

- b. Change the directory to / (root), and then verify that the pwd command works.

```
pwd
```

cannot determine current directory

```
#
```

---

**Note** – Different shells display different messages. Users of the BASH shell see the correct directory.

---

```
cd /
```

```
pwd
```

```
/
```

23. Unmount the file systems mounted below /dir0 and /morespace.

```
umount /dir0
```

```
umount /morespace
```

24. Use the vi editor to remove the line from /etc/vfstab that mounts a file system below /morespace.

## Notes:

---

## Module 7

---

# Performing Solaris 10 OS Package Administration

---

## Objectives

Upon completion of this module, you should be able to:

- Describe the fundamentals of package administration
- Administer packages using the command-line interface

# Introducing the Fundamentals of Package Administration

Software package administration adds software to systems and removes software from systems. Sun and its third-party vendors deliver software products to users in software packages.

## Software Packages

The term package refers to the method of distributing software products and installing them in systems. In its simplest form, a package is a collection of files and directories.



**Note** – The Solaris OS software installation process installs all the required software packages automatically, based on the software group configuration choice.

Software packages contain:

- Files that describe the package and the amount of disk space required for installation
- Compressed software files to be installed on the system
- Optional scripts that run when the package is added or removed

## The /var/sadm/install/contents File

The /var/sadm/install/contents file is a complete record of all the software packages installed on the local system disk. It references every file and directory belonging to every software package and shows the configuration of each product installed. To list the contents of the /var/sadm/install/contents file, perform the command:

```
more /var/sadm/install/contents
/bin=./usr/bin s none SUNWcsr
/dev d none 0755 root sys SUNWcsr SUNWcsd
/dev/allkmem=../devices/pseudo/mm@0:allkmem s none SUNWcsd
/dev/arp=../devices/pseudo/arp@0:arp s none SUNWcsd
/dev/conslog=../devices/pseudo/log@0:conslog s none SUNWcsd
/dev/console=../devices/pseudo/cn@0:console s none SUNWcsd
/dev/dsk d none 0755 root sys SUNWcsd
(output omitted)
```

The pkgadd command updates the /var/sadm/install/contents file each time new packages are installed.

The pkgrm command uses the /var/sadm/install/contents file to determine where the files for a software package are located on the system. When a package is removed from the system, the pkgrm command updates the /var/sadm/install/contents file.

To determine if a particular file was installed on the system disk and to find the directory in which it is located, use the pkgchk command with either the full or partial path name of the command you want to report on. For example, to verify that the showrev command is installed on the system disk, perform the command:

```
pkgchk -l -P showrev
Pathname: /usr/bin/showrev
Type: regular file
Expected mode: 0755
Expected owner: root
Expected group: sys
Expected file size (bytes): 29992
Expected sum(1) of contents: 59728
Expected last modification: Nov 08 23:02:51 2006
Referenced by the following packages:
 SUNWadmc
Current status: installed

Pathname: /usr/share/man/man1m/showrev.1m
Type: regular file
Expected mode: 0644
Expected owner: root
Expected group: root
Expected file size (bytes): 3835
Expected sum(1) of contents: 64380
Expected last modification: Sep 15 18:42:55 2006
Referenced by the following packages:
 SUNWman
Current status: installed
```

## Package Formats

Solaris OS packages can be in one of two formats:

- File system (or Directory) format
- Data stream format

Packages delivered in file system format consist of multiple files and directories. Packages delivered in data stream format consist of a single file only.

### File System Format

An example of a package (SUNWzsh) in file system format:

```
ls -ld SUNWzsh
drwxr-xr-x 5 root staff 2048 Nov 14 20:54 SUNWzsh
cd SUNWzsh
ls -l
total 3280
total 127
drwxr-xr-x 2 root staff 2048 Nov 14 20:54 archive
drwxr-xr-x 2 root staff 2048 Nov 14 20:54 install
-rw-r--r-- 1 root staff 455 Nov 14 15:56 pkginfo
-rw-r--r-- 1 root staff 58248 Nov 14 15:56 pkgmap
drwxr-xr-x 3 root staff 2048 Nov 14 20:54 reloc
#
```

The package consists of a directory that matches the package name, and other files and directories including the `pkginfo` and `pkgmap` files.

## Data Stream Format

An example of a package in data stream format:

```
ls -l SUNWzsh.pkg
-rw-r--r-- 1 root root 1092608 Feb 15 14:56 SUNWzsh.pkg
file SUNWzsh.pkg
SUNWzsh.pkg: package datastream
head SUNWzsh.pkg
PaCkAgE DaTaStReAm
SUNWzsh 1 2110
end of header
NAME=Z shell (zsh)
ARCH=sparc
VERSION=11.10.0,REV=2005.01.08.05.16
SUNW_PRODNAME=SunOS
SUNW_PRODVERS=5.10/SunOS Development
SUNW_PKGTYPE=usr
MAXINST=1000
#
```

Packages downloaded from the Internet are often in data stream format.

# Administering Packages From the Command Line

From the command line, you can translate, add, remove, check the state of, and display information about packages.

The command-line tools are used for translating packages, viewing software, adding software, and removing software from a workstation after the Solaris OS software is installed on a system include:

|          |                                                |
|----------|------------------------------------------------|
| pkgtrans | Translates packages from one format to another |
| pkgadd   | Installs software packages to the system       |
| pkgrm    | Removes a package from the system              |
| pkginfo  | Displays software package information          |
| pkgchk   | Checks package installation state              |

## Translating Package Formats

Use the pkgtrans command to translate a package from file system format to data stream format, or from data stream format to file system format. The command syntax for the pkgtrans command is:

```
pkgtrans source_device destination_device [package_name ...]
```

For example, to translate a package from file system format in /var/tmp to data stream format, use:

```
pkgtrans /var/tmp /tmp/SUNWrsc.pkg SUNWrsc
Transferring <SUNWrsc> package instance
```

The first argument above specifies the source directory where the file system format package is stored. The second argument specifies the destination package data stream file. The third argument is the package to translate that exists in the source directory specified by the first argument.

If a package name is not given, the pkgtrans command provides a list of all packages in the source directory, and prompts the user for the packages to translate.

---

**Note** – Students need to insert the Solaris 10 OS Software DVD to demonstrate the steps described in this module.

---

## Displaying Installed Software Packages Information

Use the `pkginfo` command to display information about the software packages installed on a local system disk. The `/var/sadm/pkg` directory maintains a record of installed packages. For example, to display information about software packages installed on a local system disk:

```
pkginfo | more
(output omitted)
```

|                  |                 |                                                                |
|------------------|-----------------|----------------------------------------------------------------|
| system           | SUNWcti2        | Netra ct I2C and System Drivers                                |
| ALE              | SUNWctltk       | Simplified Chinese ToolTalk Runtime Package                    |
| system           | SUNWctlu        | Print utilities for CTL locales                                |
| CTL              | SUNWctpls       | Portable layout services for Complex Text Layout support       |
| ALE              | SUNWcttf        | Simplified Chinese (EUC) True Type Fonts                       |
| ALE              | SUNWcttfe       | Simplified Chinese (EUC) True Type Fonts Extension             |
| application      | SUNWcudc        | Simplified Chinese User Defined Character tool for Solaris CDE |
| ALE              | SUNWcufnt       | Simplified Chinese (UTF-8) X Windows Platform required Fonts   |
| Package          |                 |                                                                |
| system           | SUNWcupdatemgru | Simplified Chinese Update Manager GUI                          |
| system           | SUNWcvc         | Network Console                                                |
| system           | SUNWcvcr        | Network Console daemon                                         |
| (output omitted) |                 |                                                                |

The column on the left displays the package category, such as application, system, complex text layout (CTL), or alternate language environment (ALE). A CTL language is any language that stores text differently than it is displayed. An ALE is an alternate language, different from the basic Solaris OS languages.

The center column displays the software package name. If it begins with SUNW, it is a Sun Microsystems product. Otherwise, it represents a third-party package.

The column on the right displays a brief description of the software product.

### Displaying Information for All Packages

To display all the available information about the software packages, use the `pkginfo` command with the `-l` option.

For example, to view additional information about each software package installed on the local systems hard drive, perform the command:

```
pkginfo -l | more
(output omitted)
```

## Displaying Information for a Specific Package

To display the information for a specific software package, specify its name on the command line. For example, to view the information for the SUNWman software package, enter:

```
pkginfo -l SUNWman
```

```
PKGINST: SUNWman
 NAME: On-Line Manual Pages
CATEGORY: system
 ARCH: sparc
VERSION: 43.0,REV=75.0
BASEDIR: /usr
VENDOR: Sun Microsystems, Inc.
DESC: System Reference Manual Pages
PSTAMP: 2004.12.21.15.30
INSTDATE: Jan 11 2007 13:07
HOTLINE: Please contact your local service provider
STATUS: completely installed
FILES: 12402 installed pathnames
 21 shared pathnames
 190 directories
 67327 blocks used (approx)
```

#

The last line identifies the size of the package. The number of blocks used defines how much space is needed on the disk to install the package.



---

**Note –** A block is a 512-byte disk block.

---

To determine how many packages are currently installed on disk, perform the command:

```
pkginfo | wc -l
1720
```

## Displaying Information for Software Packages

To view information about packages that are located on a SPARC platform-based Solaris 10 OS Software DVD, perform the command:

```
pkginfo -d /cdrom/cdrom0/Solaris_10/Product | more
```

To view information about packages that are located on an x86/x64 Solaris 10 OS Software DVD, perform the command:

```
pkginfo -d /cdrom/cdrom0/Solaris_10/Product | more
```

To view information about packages that are located on an x86/x64 Solaris 10 OS Software 1 CD-ROM, perform the command:

```
pkginfo -d /cdrom/cdrom0/Solaris_10/Product | more
```

The software groups located on a Solaris 10 OS Software 1 CD-ROM are Reduced Networking Core System Support and Core System Support.

To view information about packages that are located on any of the Solaris 10 Software CD-ROMs other than the Software 1 CD-ROM, perform the command:

```
pkginfo -d /cdrom/cdrom0/Solaris_10/Product | more
```

The software groups located on Solaris 10 OS Software 2, 3, and 4 CD-ROMs are the End User System Support, Developer System Support, Entire Distribution, and Entire Distribution Plus OEM Support software groups.

## Adding a Software Package

When you add a software package, the pkgadd command copies the files from the installation media to the local system's disk and executes scripts to uncompress files. By default, the pkgadd command requires confirmation during the package add process.

For example, on a SPARC system, to transfer the SUNWvts software package from a DVD and install it on the system, perform the commands:

```
cd /cdrom/cdrom0/Solaris_10/ExtraValue/CoBundled/SunVTS_6.3/Packages
pkgadd -d . SUNWvts
```

```
Processing package instance <SUNWvts> from
</cdrom/sol_10_1106_sparc/s0/Solaris_10/ExtraValue/CoBundled/SunVTS_6.3/P
ackages>
```

```
SunVTS Framework(sparc) 6.3,REV=2006.10.19.18.34
Copyright 2006 Sun Microsystems, Inc. All rights reserved.
Use is subject to license terms.
Using </opt> as the package base directory.
Processing package information.
```

## Administering Packages From the Command Line

---

```
Processing system information.
Verifying package dependencies.
Verifying disk space requirements.
Checking for conflicts with packages already installed.
Checking for setuid/setgid programs.
```

This package contains scripts which will be executed with super-user permission during the process of installing this package.

Do you want to continue with the installation of <SUNWvts> [y,n,?] **y**

Installing SunVTS Framework as <SUNWvts>

```
Installing part 1 of 1.
8094 blocks
```

Installation of <SUNWvts> was successful.

```
#
```

On an x86/x64 system, you would find the SUNWvts package in the following directory on the Solaris OS DVD:

/cdrom/cdrom0/Solaris\_10/ExtraValue/CoBundled/SunVTS\_6.3/Packages

To install all packages in a data stream format package, perform the command:

```
pkgadd -d /tmp/SUNWzsh.pkg all
```

Processing package instance <SUNWzsh> from </tmp/SUNWzsh.pkg>

Z shell (zsh) (sparc) 11.10.0,REV=2005.01.08.05.16

(output omitted)

Using </> as the package base directory.

```
Processing package information.
Processing system information.
6 package pathnames are already properly installed.
Verifying package dependencies.
Verifying disk space requirements.
Checking for conflicts with packages already installed.
Checking for setuid/setgid programs.
```

This package contains scripts which will be executed with super-user permission during the process of installing this package.

Do you want to continue with the installation of <SUNWzsh> [y,n,?] **y**

Installing Z shell (zsh) as <SUNWzsh>

```
Installing part 1 of 1.
5481 blocks
```

Installation of <SUNWzsh> was successful.  
#

Packages in data stream format can also be added from a web server. The following example commands require a web server running on the instructor system in the classroom, and the SPARC and x86/x64 versions of the SUNWzsh package stored in the specified locations.



---

**Note** – Before you try these examples, use the pkgrm SUNWzsh command to remove the SUNWzsh package that is installed on your system.

---

```
pkgadd -d http://instructor/packages/sparc/SUNWzsh.pkg all
```

```
pkgadd -d http://instructor/packages/x86/SUNWzsh.pkg all
```

SPARC and x86/x64 systems display installation messages similar to the following:

```
Downloading...
.....25%.....50%.....75%.....100%
Download Complete
```

Processing package instance <SUNWzsh> from  
<http://instructor/packages/SUNWzsh.pkg>

Z shell (zsh) (sparc) 11.10.0,REV=2005.01.08.05.16  
(output omitted)

Using </> as the package base directory.

```
Processing package information.
Processing system information.
3 package pathnames are already properly installed.
Verifying package dependencies.
Verifying disk space requirements.
Checking for conflicts with packages already installed.
```

## Administering Packages From the Command Line

---

The following files are already installed on the system and are being used by another package:

```
/usr <attribute change only>
/usr/sfw <attribute change only>
/usr/sfw/lib <attribute change only>
```

Do you want to install these conflicting files [y,n,?,q] **y**  
## Checking for setuid/setgid programs.

This package contains scripts which will be executed with super-user permission during the process of installing this package.

Do you want to continue with the installation of <SUNWzsh> [y,n,?] **y**

Installing Z shell (zsh) as <SUNWzsh>

```
Installing part 1 of 1.
5481 blocks
```

Installation of <SUNWzsh> was successful.

```
#
```

## Checking a Package Installation

The **pkgchk** command checks to determine if a package has been completely installed on the system. The **pkgchk** command also checks the path name, the file size and checksum, and the file attributes of a package. If the **pkgchk** command does not display a message, it indicates the package was installed successfully and that no changes have been made to any files or directories in the package.

The following example checks the contents and attributes of the SUNWladm software package currently installed on the system.

```
pkgchk SUNWladm
#
```

To list the files contained in a software package, use the **-v** option.

For example, to list the files in the SUNWladm software package, perform the command:

```
pkgchk -v SUNWladm
/usr
/usr/sadm
/usr/sadm/lib
/usr/sadm/lib/localeadm
/usr/sadm/lib/localeadm/Locale_config_S10.txt
/usr/sadm/lib/localeadm/admin
/usr/sadm/lib/localeadm/core_packages.txt
/usr/sadm/lib/localeadm/gen_configfile.pl
/usr/sadm/lib/localeadm/retrievecd.pl
/usr/sadm/lib/localeadm/retrievecombined.pl
/usr/sbin
/usr/sbin/localeadm
#
```

To determine if the contents or attributes of a file have changed from when it was installed, use the **-p** option.

For example, to check the **/etc/shadow** file, perform the command:

```
pkgchk -p /etc/shadow
ERROR: /etc/shadow
modtime <11/09/06 09:51:50 AM> expected <12/27/06 05:44:53 PM> actual
file size <296> expected <344> actual
file cksum <20180> expected <23909> actual
```

The differences in modtime, file size, and checksum indicate that the original /etc/shadow file has changed in size. For this particular file, this happened because of the initial Solaris OS software installation.

The -l option with the pkgchk command lists information about selected files that make up a package.

For example, to list information about the /usr/bin/showrev file, perform the command:

```
pkgchk -l -p /usr/bin/showrev
Pathname: /usr/bin/showrev
Type: regular file
Expected mode: 0755
Expected owner: root
Expected group: sys
Expected file size (bytes): 29992
Expected sum(1) of contents: 59728
Expected last modification: Nov 08 23:02:51 2006
Referenced by the following packages:
 SUNWadmc
Current status: installed
```

If the -p option is used, the full path must be typed for the pkgchk command to return information about the file. If the -P option is used, a partial path name can be supplied.

For example, the pkgchk command returns no information if the /usr/bin/ path is omitted from the previous example.

```
pkgchk -l -p showrev
#
```

## Removing a Software Package

The pkgrm command removes a software package from the system and deletes all of the files associated with that package, unless other packages share those files.

By default, the pkgrm command requires confirmation to continue removing a package and issues a message to warn about possible package dependencies. If package dependencies do exist, the command again requires confirmation to continue with the package removal process.

The following command removes the SUNWapchr software package from the system.



**Caution** – Be cautious of the dependency warnings you receive when removing a package. The system allows you to remove these packages even though they may be required by a different package.

```
pkgrm SUNWapchr
```

The following package is currently installed:

```
SUNWapchr Apache Web Server (root)
(sparc) 11.10.0,REV=2005.01.08.05.16
```

Do you want to remove this package? [y,n,?,q] **y**

```
Removing installed package instance <SUNWapchr>
Verifying package <SUNWapchr> dependencies in global zone
WARNING:
```

```
The <SUNWipplr> package depends on the package
currently being removed.
```

WARNING:

```
The <SUNWapchu> package depends on the package
currently being removed.
```

WARNING:

```
The <SUNWapchd> package depends on the package
currently being removed.
```

WARNING:

```
The <SUNWserweb> package depends on the package
currently being removed.
```

Dependency checking failed.

Do you want to continue with the removal of this package [y,n,?,q] **y**

```
Processing package information.
Removing pathnames in class <initd>
/etc/rcS.d/K16apache
/etc/rc3.d/S50apache
```

(output omitted)

```
/etc/apache <shared pathname not removed>
/etc <shared pathname not removed>
Updating system information.
```

Removal of <SUNWapchr> was successful.



**Note** – A file shared by two or more packages displays the message  
*filename <shared pathname not removed>*.

## Adding Packages by Using a Spool Directory

For convenience, you can copy frequently installed software packages from the Solaris 10 Software DVD or Solaris 10 Software CD-ROMs to a spool directory on the system.

The default installation directory for packages that have been spooled, but not installed, is `/var/spool/pkg`. The `pkgadd` command, by default, looks in the `/var/spool/pkg` directory for any packages specified on the command line.

Copying packages from the CD-ROM or DVD into a spool directory is not the same as installing the packages on disk.

On a SPARC system, to copy a package from the Solaris 10 OS Software DVD into the `/var/spool/pkg` directory, perform the command:

```
pkgadd -d /cdrom/cdrom0/Solaris_10/Product -s spool SUNWauda
Transferring <SUNWauda> package instance
```

On an x86/x64 system, to copy a package from the Solaris 10 OS Software DVD into the `/var/spool/pkg` directory, perform the command:

```
pkgadd -d /cdrom/cdrom0/Solaris_10/Product -s spool SUNWauda
Transferring <SUNWauda> package instance
```

The `-s` option with the keyword `spool` copies the package into the `/var/spool/pkg` directory by default.

To verify that the package exists in the `spool` directory, perform the command:

```
ls -al /var/spool/pkg
total 6
drwxrwxrwt 3 root bin 512 Feb 15 22:19 .
drwxr-xr-x 12 root bin 512 Feb 6 17:26 ..
drwxr-xr-x 5 root root 512 Feb 15 22:19 SUNWauda
```

To add the package from the spool area, perform the following:

```
pkgadd SUNWauda
(output omitted)
```

To remove software packages from a spool directory, use the `pkgrm` command with the `-s` option.

```
pkgrm -s spool SUNWauda
```

The following package is currently spooled:

```
SUNWauda Audio Applications
(sparc) 11.10.0,REV=2005.01.21.15.53
```

Do you want to remove this package? [y,n,?,q] y

Removing spooled package instance <SUNWauda>

If alternative spooling directories exist, specify the directory to use by substituting the directory path for the `spool` keyword.

For example, to select the `/export/pkg` directory, perform the commands:

```
pkgadd -d /cdrom/cdrom0/Solaris_10/Product -s /export/pkg SUNWauda
pkgrm -s /export/pkg SUNWauda
```

## Streaming One or More Packages

Packages can be individually or collectively packaged into a data stream file format. The data stream file can then be made available as a shared network file or from a web page.

### Worked Example

To create a data streamed package, perform the following commands:

On a SPARC system:

```
cd /cdrom/cdrom0/s0/Solaris_10
```

On an x86/x64 system:

```
cd /cdrom/cdrom0/Solaris_10
```

On both SPARC and x86/x64 systems:

```
pkgtrans -s Product /var/tmp/stream.pkg SUNWzlib SUNWftpr SUNWftpu
Transferring <SUNWzlib> package instance
Transferring <SUNWftpr> package instance
Transferring <SUNWftpu> package instance

file /var/tmp/stream.pkg
/var/tmp/stream.pkg: package datastream

head -5 /var/tmp/stream.pkg
PaCkAgE DaTaStReAm
SUNWzlib 1 186
SUNWftpr 1 60
SUNWftpu 1 300
end of header

pkgadd -d /var/tmp/stream.pkg
```

The following packages are available:

- 1 SUNWftpr      FTP Server, (Root)  
(sparc) 11.10.0,REV=2005.01.21.15.53
- 2 SUNWftpu      FTP Server, (Usr)  
(sparc) 11.10.0,REV=2005.01.21.15.53
- 3 SUNWzlib      The Zip compression library  
(sparc) 11.10.0,REV=2005.01.08.05.16

Select package(s) you wish to process (or 'all' to process all packages). (default: all) [?,??,q]: **q**

## Reviewing Package Administration

This section details the package administration tasks.

Table 7-1 summarizes the commands used for package administration.

**Table 7-1** Package Administration Commands

| Command Name | Description                                                                                                                                             |
|--------------|---------------------------------------------------------------------------------------------------------------------------------------------------------|
| pkginfo      | Lists packages installed on the system or available on distribution media                                                                               |
| pkgadd       | Installs packages                                                                                                                                       |
| pkgrm        | Removes packages                                                                                                                                        |
| pkgchk       | Verifies the attributes of the path names that belong to packages                                                                                       |
| pkgtrans     | Translates a package from one format to another, such as a file system format to a datastream, or one file system format to another file system format. |

Table 7-2 summarizes the files and directories used in package administration.

**Table 7-2** Package Administration Files and Directories

| File or Directory                                    | Description                                                           |
|------------------------------------------------------|-----------------------------------------------------------------------|
| /var/sadm/install/contents                           | A software package map of the entire system                           |
| /opt/ <i>pkgname</i>                                 | The preferred location for the installation of unbundled packages     |
| /opt/ <i>pkgname</i> /bin or /opt/bin                | The preferred location for the executable files of unbundled packages |
| /var/opt/ <i>pkgname</i> or /etc/opt/ <i>pkgname</i> | The preferred location for log files of unbundled packages            |

# Exercise 1: Adding a Software Package

In this exercise, you use the package-related commands to add software packages to the Solaris 10 OS.

## Preparation

This exercise requires the packages found in the /opt/ses/lab/packages directory on student systems. If the packages are located on a server, your instructor will provide its location and information about how to retrieve them.

Refer to the lecture notes as necessary to perform the tasks listed.



**Note** – Use the packages in the /opt/ses/lab/packages directory to understand how to use package administration commands. The packages used for this exercise are for demonstration only.

---

## Task

Complete the following steps:

1. List the packages found in the /opt/ses/lab/packages directory.
2. Display a long-format listing of the information for the packages.
3. Add the packages from the /opt/ses/lab/packages directory. Respond **y** to questions asked by the pkgadd command.
4. Use pkginfo to list information about the installed packages.
5. Use pkgchk to verify that the packages were installed successfully.

## Exercise Summary



**Discussion** – Take a few minutes to discuss what experiences, issues, or discoveries you had during the lab exercises.

- Experiences
- Interpretations
- Conclusions
- Applications

# Exercise 1 Solutions: Adding a Software Package

In this exercise, you use the package-related commands to add a software packages to the Solaris 10 OS.

## Preparation

This exercise requires the packages found in the /opt/ses/lab/packages directory on student systems. If the packages are located on a server, your instructor will provide its location and information about how to retrieve them.



**Note** – Use the packages in the /opt/ses/lab/packages directory to understand how to use package administration commands. The packages used for this exercise are for demonstration only.

## Task

Complete the following steps:

1. List the packages found in the /opt/ses/lab/packages directory.

```
ls /opt/ses/lab/packages
SUNWoptdir SUNWusrdir
```

2. Display a long-format listing of the information for the packages.

```
pkginfo -d /opt/ses/lab/packages -l SUNWoptdir
pkginfo -d /opt/ses/lab/packages -l SUNWusrdir
```

3. Add the packages from the /opt/ses/lab/packages directory.

Respond **y** to questions asked by the pkgadd command.

```
pkgadd -d /opt/ses/lab/packages SUNWoptdir
pkgadd -d /opt/ses/lab/packages SUNWusrdir
```

4. Use pkginfo to list information about the installed packages.

```
pkginfo SUNWoptdir
pkginfo SUNWusrdir
```

5. Use pkgchk to verify that the packages were installed successfully.

```
pkgchk -v SUNWoptdir
pkgchk -v SUNWusrdir
```

# Exercise 2: Manipulating Software Packages

In this exercise, you use package-related commands to remove, install, and spool packages.

## Preparation

This exercise requires the Solaris 10 OS Software DVD appropriate for your system architecture. Your instructor will provide these items.

Refer to the lecture notes as necessary to perform the tasks listed.



**Note** – If you are using CD-ROM media, the SUNWman and SUNWdoc packages are located on the CD 4 of 4 and CD 2 of 4 Solaris 10 OS Software CD-ROMs, respectively.

## Task

Complete the following steps:

1. Insert the Solaris 10 Software SPARC or x86/x64 DVD into the drive.
2. Use the pkginfo command to search for packages currently on your system that are related to manuals.

Which packages were listed?

---

---

---

---



**Note** – The list that displays depends on selections that were made during the installation process.

## Exercise 2: Manipulating Software Packages

---

3. Display a long-format listing of the information for the SUNWman package installed on your system.

What is listed for the following:

- The status of this package?  
\_\_\_\_\_
- The install date of this package?  
\_\_\_\_\_
- The number of files used by this package?  
\_\_\_\_\_
- The number of blocks used by this package?  
\_\_\_\_\_

**Note** – These values may differ among different releases of Solaris 10. Values will be the same for SPARC and x86/x64 systems for any given release.

4. Display a long-format listing of the information for the SUNWman package on the Solaris 10 OS Software DVD.

Obtain the same information as in the previous step.

- The status of this package?  
\_\_\_\_\_
- The install date of this package?  
\_\_\_\_\_
- The number of files used by this package?  
\_\_\_\_\_
- The number of blocks used by this package?  
\_\_\_\_\_

5. Remove the SUNWman package from your system and complete the following steps:

- a. Respond **y** to questions asked by the `pkgrm` command.
- b. Verify that SUNWman has been removed by using the `pkginfo` command, and trying to access the manual pages.



6. Reinstall the SUNWman package from the Solaris 10 OS Software DVD and complete the following steps:
  - a. Respond **y** to questions asked by the `pkgadd` command.
  - b. Use `pkginfo` to list information about the SUNWman package.
  - c. Verify that the manual pages work.
7. Check if the SUNWdoc package is installed on your system, and complete the following steps:
  - a. Remove the SUNWdoc package from your system.
  - b. Respond **y** to questions asked by the `pkgrm` command.

Are there any package dependencies related to removing this package?

---

  8. Use the `pkgadd` command to spool the SUNWdoc package into the default spool area.
  9. Use the `pkginfo` command (with the appropriate options) to verify the presence of the SUNWdoc package in the default spool area.
  10. Install the SUNWdoc package and complete the following steps:
    - a. Observe the messages.
    - b. Respond **y** to questions asked by the `pkgadd` command.
    - c. Verify that the package was installed from the `/var/spool/pkg` directory.
  11. Remove the SUNWdoc package from the default spool area.
  12. Verify that the SUNWdoc package no longer exists in the spool area and that it is installed on your system.

## Exercise Summary



**Discussion** – Take a few minutes to discuss what experiences, issues, or discoveries you had during the lab exercises.

- Experiences
- Interpretations
- Conclusions
- Applications

# Exercise 2 Solutions: Manipulating Software Packages

In this exercise, you use package-related commands to remove, install, and spool packages.

## Preparation

This exercise requires the Solaris 10 OS Software DVD appropriate for your system architecture. Your instructor will provide these items.

Refer to the lecture notes as necessary to perform the tasks listed.



**Note** – If you are using CD-ROM media, the SUNWman and SUNWdoc packages are located on the CD 4 of 4 and CD 2 of 4 Solaris 10 OS Software CD-ROMs, respectively.

## Task

Complete the following steps:

1. Insert the Solaris 10 Software SPARC or x86/x64 DVD into the drive.
2. Use the pkginfo command to search for packages currently on your system that are related to manuals.

```
pkginfo | grep -i manual
```

Which packages were listed?

|                     |                 |                             |
|---------------------|-----------------|-----------------------------|
| system              | SUNWcrman       | Encryption Kit On-Line      |
| <b>Manual Pages</b> |                 |                             |
| system              | SUNWman         | On-Line Manual Pages        |
| system              | SUNWmfman       | CDE Motif Manuals           |
| system              | SUNWopenssl-man | OpenSSL Manual Pages        |
| system              | SUNWperl584man  | Perl 5.8.4 Reference Manual |
| <b>Pages</b>        |                 |                             |
| system              | SUNWpl15m       | Perl 5.6.1 Reference Manual |
| <b>Pages</b>        |                 |                             |
| system              | SUNWt1tkm       | ToolTalk manual pages       |



**Note** – The list that displays depends on selections that were made during the installation process.

---

3. Display a long-format listing of the information for the SUNWman package installed on your system.

# **pkginfo -l SUNWman**

What is listed for the following:

- The status of this package?
- The install date of this package?
- The number of files used by this package?
- The number of blocks used by this package?

*Status:* Completely installed

*Install date:* Should match the date and time when you installed Solaris OS on your system

*Number of files:* 12402 installed path names, 21 shared directories, 190 directories

*Number of blocks:* 67327

**Note** – These values may differ among different releases of Solaris 10. Values will be the same for SPARC and x86/x64 systems for any given release. The values listed are examples only.

---



4. Display a long-format listing of the information for the SUNWman package on the Solaris 10 OS Software DVD.

Obtain the same information as in the previous step.

For example:

```
pkginfo -d /cdrom/cdrom0/Solaris_10/Product -l SUNWman
```

*Status:* Spooled

*Install date:* No install date indicated

*Number of files:* 12406 spooled path names, 190 directories, 4 package information files

*Number of blocks:* 67327

5. Remove the SUNWman package from your system and complete the following steps:

- Respond **y** to questions asked by the **pkgrm** command.
- Verify that SUNWman has been removed by using the **pkginfo** command, and trying to access the manual pages.

```
pkgrm SUNWman
<output omitted>
pkginfo SUNWman
ERROR: information for "SUNWman" was not found
man ls
No manual entry for ls.
#
```

## Exercise 2 Solutions: Manipulating Software Packages

---

6. Reinstall the SUNWman package from the Solaris 10 OS Software DVD and complete the following steps:

- Respond **y** to questions asked by the pkgadd command.
- Use pkginfo to list information about the SUNWman package.
- Verify that the manual pages work.

For example:

```
pkgadd -d /cdrom/cdrom0/Solaris_10/Product SUNWman
(output omitted)
pkginfo SUNWman
system SUNWman On-Line Manual Pages
man ls
<output omitted>
```

*The manual page for ls displays.*

7. Check if the SUNWdoc package is installed on your system, and complete the following steps:
- Remove the SUNWdoc package from your system.
  - Respond **y** to questions asked by the pkgrm command.

```
pkginfo SUNWdoc
system SUNWdoc Documentation Tools
pkgrm SUNWdoc
```

The following package is currently installed:

```
SUNWdoc Documentation Tools
(sparc) 11.10.0,REV=2005.01.21.15.53
```

Do you want to remove this package? [y,n,?,q] **y**

```
Removing installed package instance <SUNWdoc>
```

This package contains scripts which will be executed with super-user permission during the process of removing this package.

Do you want to continue with the removal of this package [y,n,?,q] **y**

```
Verifying package <SUNWdoc> dependencies in global zone
```

WARNING:

The <SUNWuium> package depends on the package currently being removed.

WARNING:

The <SUNWman> package depends on the package currently being removed.

WARNING:

The <SUNWpmowm> package depends on the package

currently being removed.

**WARNING:**

The <SUNWopenssl-man> package depends on the package currently being removed.

**WARNING:**

The <SUNWperl584man> package depends on the package currently being removed.

**WARNING:**

The <SUNWpl15m> package depends on the package currently being removed.

**WARNING:**

The <SUNWvtsmn> package depends on the package currently being removed.

Dependency checking failed.

Do you want to continue with the removal of this package [y,n,?,q] **y**  
## Processing package information.

## Executing preremove script.

## Removing pathnames in class <none>

/usr/share/man <shared pathname not removed>

/usr/share/lib/tmac/vgrind

(output omitted)

/usr/bin/apropos

/usr/bin/addbib

/usr/bin <shared pathname not removed>

/usr <shared pathname not removed>

## Updating system information.

Removal of <SUNWdoc> was successful.

Are there any package dependencies related to removing this package?

*Yes, there are. There are seven other packages dependent on the SUNWdoc package.*

## Exercise 2 Solutions: Manipulating Software Packages

---

8. Use the pkgadd command to spool the SUNWdoc package into the default spool area.

For example:

```
pkgadd -d /cdrom/cdrom0/Solaris_10/Product -s spool SUNWdoc
```

9. Use the pkginfo command (with the appropriate options) to verify the presence of the SUNWdoc package in the default spool area.

```
pkginfo -d spool SUNWdoc
```

```
system SUNWdoc Documentation Tools
```

```
pkginfo -d /var/spool/pkg -l SUNWdoc
```

```
PKGINST: SUNWdoc
```

```
(output omitted)
```

10. Install the SUNWdoc package and complete the following steps:
  - a. Observe the messages.
  - b. Respond **y** to questions asked by the pkgadd command.
  - c. Verify that the package was installed from the /var/spool/pkg directory.

```
pkgadd SUNWdoc
```

```
Processing package instance <SUNWdoc> from </var/spool/pkg>
```

```
(output omitted)
```

11. Remove the SUNWdoc package from the default spool area.

```
pkgrm -s spool SUNWdoc
```

The following package is currently spooled:

```
SUNWdoc Documentation Tools
(sparc) 11.10.0,REV=2005.01.21.15.53
```

Do you want to remove this package? [y,n,?,q] **y**

Removing spooled package instance <SUNWdoc>

```
#
```

12. Verify that the SUNWdoc package no longer exists in the spool area and that it is installed on your system.

```
pkginfo -d spool SUNWdoc
ERROR: information for "SUNWdoc" was not found
pkginfo -l SUNWdoc
PKGINST: SUNWdoc
(output omitted)
```

13. Eject the Solaris 10 OS Software DVD.

## Notes:

---

## Module 8

---

# Managing Software Patches on the Solaris 10 OS

---

## Objectives

Upon completion of this module, you should be able to:

- Describe patch administration fundamentals
- Describe patch administration best practices
- Obtain patches using SunSolve
- Install and remove patches
- Install recommended patch clusters
- Implement patch management using the `smpatch` command

# Patch Administration Preparation

Patch administration involves installing or removing Solaris OS patches from a running Solaris OS. Please read the patch README file before installing any patch.

For information on patching with Live Upgrade, see the “System Installation” module.

## Introducing Solaris OS Patches

A patch contains a collection of files and directories. This collection replaces existing files and directories that prevent proper software execution. Some patches contain product enhancements. Patches are also described as *software updates*.

Solaris OS patch types include:

- Standard patches – Patches that fix specific problems with the Solaris OS and other Sun hardware and software products.
- Recommended patches – Sun Customer Service designates operating system patches that are of universal interest or reflect security concerns to be “recommended” and “security” patches, respectively.
- Firmware and PROM patches.
- Patch clusters – A group of standard, recommended, security, or Y2K patches that have been bundled into a single archive for easy downloading and installation.

A patch is distributed as a directory that is named using a unique number. The number assigned to a patch includes the patch base code first, a hyphen, and a number that represents the patch revision number.

For example, a patch directory named 105050-01, indicates that 105050 is the patch number and 01 is the revision number.

The Solaris 10 OS unsigned patches are delivered in zip format, for example, 105050-01.zip.

Signed patches are delivered as Java Archive format files, for example, 105050-01.jar.

A signed patch is a patch that is signed with a valid digital signature. A signed patch offers greater security than an unsigned patch. The digital signature of the patch can be verified before the patch is applied to your system. A valid digital signature ensures that the signed patch has not been modified since the signature was applied.

To add a patch with a digital signature to your system, you must set-up a keystore. A keystore is a repository of certificates and keys that is queried when you apply a signed patch.

For more information about signed patches and setting up keystores, see the *Managing Solaris Patches and Updates (Overview)* section of the *Solaris 10 System Administration Guide: Basic Administration* found on <http://docs.sun.com/app/docs/doc/817-1985/6mhm8o611>.



**Note** – Not all patches available from Sun Microsystems must be installed. Review the README documents for each patch, and then decide if you should install the patch.

## Best Practices

Following are some of the patching related best practices:

- Always install the latest patch and package utility patches first.
- Upgrade to the latest Solaris 10 Update release during your next major maintenance effort. Each Solaris Update is intensely tested and provides a quality baseline.
- Between major maintenance efforts, stay current with the Sun Alert patch cluster contents.
- Use Solaris Live Upgrade to patch or upgrade an inactive boot environment.
  - Live Upgrade enables you to avoid much of the risk and downtime associated with patching the live boot environment.
  - Live Upgrade provides a simple roll-back mechanism.

## Accessing Patches and Patch Information

Use SunSolve to access software updates (patches). SunSolve provides direct access to a variety of support resources, including patches and updates, the Sun Knowledge Database, support forums, the Sun System Handbook, support cases, Sun Alerts, and Sun Security Resources.

### Discontinued Patch Access Methods

Formerly, Sun customers were able to gain access to patches through the sunsolve.sun.com anonymous ftp server. In September of 2006, Sun ceased providing access to this anonymous ftp server for the purpose of downloading patches. These changes to the methods used to access patches are described in two infodoc articles, 82023 and 83061, available using these two URLs:

<http://sunsolve.sun.com/search/document.do?assetkey=1-9-82023-1>  
<http://sunsolve.sun.com/search/document.do?assetkey=1-9-83061-1>

Access to patches through PatchPro Interactive has also been discontinued.

### Requirements for Patch Access

Currently, customers are required to have a Sun Online account to obtain patches. This requires creating a Sun Online Account, and accepting the new click through Software Agreement.

You already have a Sun Online Account if you have an account on MySun, Sun Store, SunSolve, the Online Support Center, or Sun Download Center. Use the same user name and password during registration and to log in to SunSolve web site.

A Sun Service Plan is not required to access security patches and device drivers; however, you do need to have a Sun Online account and accept the Software License Agreement applicable to all Sun patches to have access to such patches.

Restricted patches require a Sun Service Plan or Solaris Subscription to gain access. Restricted patches include recommended patch clusters. Service plan descriptions are available using this URL:

<http://www.sun.com/service/serviceplans/index.jsp>

Access to Solaris 8 and 9 patches (except for security patches and device drivers) will also require a Solaris Subscription or Sun System Service Plan to gain access.

## Accessing Patches Using SunSolve

To access patches through the World Wide Web, use the following URLs:

<http://sunsolve.sun.com> – United States

<http://sunsolve.sun.com.au> – Australia

<http://sunsolve.sun.fr> – France

<http://sunsolve.sun.de> – Germany

<http://sunsolve.sun.co.jp> – Japan

<http://sunsolve.sun.se> – Sweden

<http://sunsolve.sun.ch> – Switzerland

<http://sunsolve.sun.co.uk> – United Kingdom

The comprehensive set of patches and patch information is available to contract customers when they log in through these sites.

You can register to create a Sun Online Account by clicking the Register link found on any of these sites, reading and accepting the Software License Agreement, and providing the information requested on the SunSolve Online Registration page. Registering to create a Sun Online Account is all that is required to obtain access to security patches and driver updates.

You can provide your Sun Service Plan number at the time that you register to create a Sun Online Account. If you have a Sun Online account without a registered Service Plan, after you log in, you can use the Edit Account link to submit a request to update your account with your Service Plan contract number.

## Accessing Patch Documents

Before you install patches on your system, you should review the related patch documents available on the SunSolve site. When reviewing patch documentation, start with the Patch Report document. This report is divided into categories that include information about all patches for a specific Solaris OS release. The Patch Report helps you identify patches to choose to install.

Individual patches and patch clusters have associated README files that provide specific information about patches and patch dependencies. Review these README files to learn more about what patches and patch clusters require, and what they provide. These documents are available to you through SunSolve whether or not you log in to the site.

### Accessing a Patch Report

On the SunSolve site, you can access Patch Reports from the Patches and Updates page, under the section titled Patch Documents and Articles.

Click on --Select to view report-- located under Patch Reports, and make your selection for the Solaris OS release and system architecture that you wish to view.

-- Select to view report --

Solaris 10  
Solaris 10\_x86  
Solaris 9  
Solaris 9\_x86  
Solaris 8  
Solaris 7  
Solaris 2.6  
Solaris 2.5.1  
Solaris 8\_x86  
Solaris 7\_x86  
Solaris 2.6\_x86  
Solaris 2.5.1\_x86

-----

Archived Patch Reports  
Archived Patch Reports  
2.5  
2.5\_x86  
2.4  
2.4\_x86  
1.1.2

The Patch Report describe different categories of patches that are available. For example, the Solaris 10 Patch Report Update from October 1, 2009, lists the following patches in the Quick Reference Section:

New Patches Released Since Last Report:

---

124581-01	commcli 7.0-1.00: core patch
139620-01	CDE 1.6: Dthelp patch
142288-01	SunOS 5.10: su patch
142371-01	Sun Studio 12.1 Update 1: Patch for dbx
142702-01	Hardware/PROM: Sun Fire V480 / V490 Flash PROM Update
142703-01	Hardware/PROM: Sun Ultra 45 Flash PROM Update
142704-01	Hardware/PROM: Sun Ultra 25 Flash PROM Update
142705-01	Hardware/PROM: Sun Fire V125 Flash PROM Update
142706-01	Hardware/PROM: Sun Fire V880 / V890 Flash PROM Update
142840-01	Logical Domains (LDom) Manager 1.2 ldmd patch

Updated Revs Released Since Last Report:

---

118245-21	comm_dssetup 6.4-5.05: core patch
119169-32	Sun Java System App Server Enterprise Ed 8.1 2005Q1 Solaris
Patch28 :	File based patch
119173-32	Sun Java System App Server Platform Ed 8.1 2005Q1 Solaris
Patch28 :	File based patch
119963-18	SunOS 5.10: Shared library patch for C++
120094-27	X11 6.6.2: xscreensaver patch
120460-16	GNOME 2.6.0: Gnome libs Patch
121657-40	Calendar Server SunOS 5.9 5.10: Core patch (output omitted)

---

**Note** – Patch Reports are published at the beginning and middle of each month.

---



## Accessing Information about Specific Patches

The patch numbers listed in the online Patch Report are links to the pages that describe their respective patches. For example, to view information for the su patch listed above, you would click the 142288-01 link.

Access the same information by using the PatchFinder link located at the top of the Patches and Updates page. You click on the PatchFinder link to open the Software Update Finder page, enter the number of the patch you wish to find, and click on the Search button.



---

**Note** – Omitting the patch revision number causes the Software Update Finder to display information for the latest revision of the patch you specify. Including the revision number causes the Software Update Finder to display information for that specific patch and revision.

---

## Latest Patch Utility Patches

You may also find patches using the Latest Patch Utility Patches pulldown menu located directly under the PatchFinder link on the Patches and Updates page to access information about patches.

Click on --Select SunOS and Patch Number-- from the pulldown menu and you make your selection for the Solaris OS release, system architecture, and patch you wish to view.

-- Select SunOS and Patch Number --

SunOS 5.10: 138253  
SunOS 5.10: 128342  
SunOS 5.10: 127884  
SunOS 5.10: 125555  
SunOS 5.10: 122034  
SunOS 5.10: 121296  
SunOS 5.10: 119317  
SunOS 5.10: 119254

SunOS 5.10x86: 138254  
SunOS 5.10x86: 128343  
SunOS 5.10x86: 127885  
SunOS 5.10x86: 125556  
SunOS 5.10x86: 122035  
SunOS 5.10x86: 121297  
SunOS 5.10x86: 119318  
SunOS 5.10x86: 119255  
(output omitted)

Selection of the latest patch utility patches is dependent on the Solaris OS release and system architecture that you wish to view. For example if you have a Solaris 10 SPARC system, you will want to select patch 119254. If you have a Solaris 10 x86 system, you will want to choose patch 119255.



**Note** – You will need to download and apply any other patches listed for Solaris 10. All the patches shown are required to avoid issues that can impact correct patch application.

For example if you have a Solaris 10 SPARC system you can click on patch 119254. The README file for the patch utilities patch is displayed. Read through this file and give special attention to any installation requirements, required patches or dependencies, patch installation instructions, and special install instructions.

After you've read through the README file, you are ready to download the patch.

To download the patch, simply click on the HTTP link to the right of Download Patch.

Download Patch (1663600 bytes) : [HTTP](#)

Notice that you have the option of downloading a signed version of the patch.

Download Signed Patch (1664612 bytes) : [HTTP](#)

A signed patch is one that has a digital signature applied to it. A patch that has its digital signature verified has not been modified since the signature was applied. The digital signature of a signed patch is verified after the patch is downloaded to your system.

After you click on the HTTP link, the download process begins.

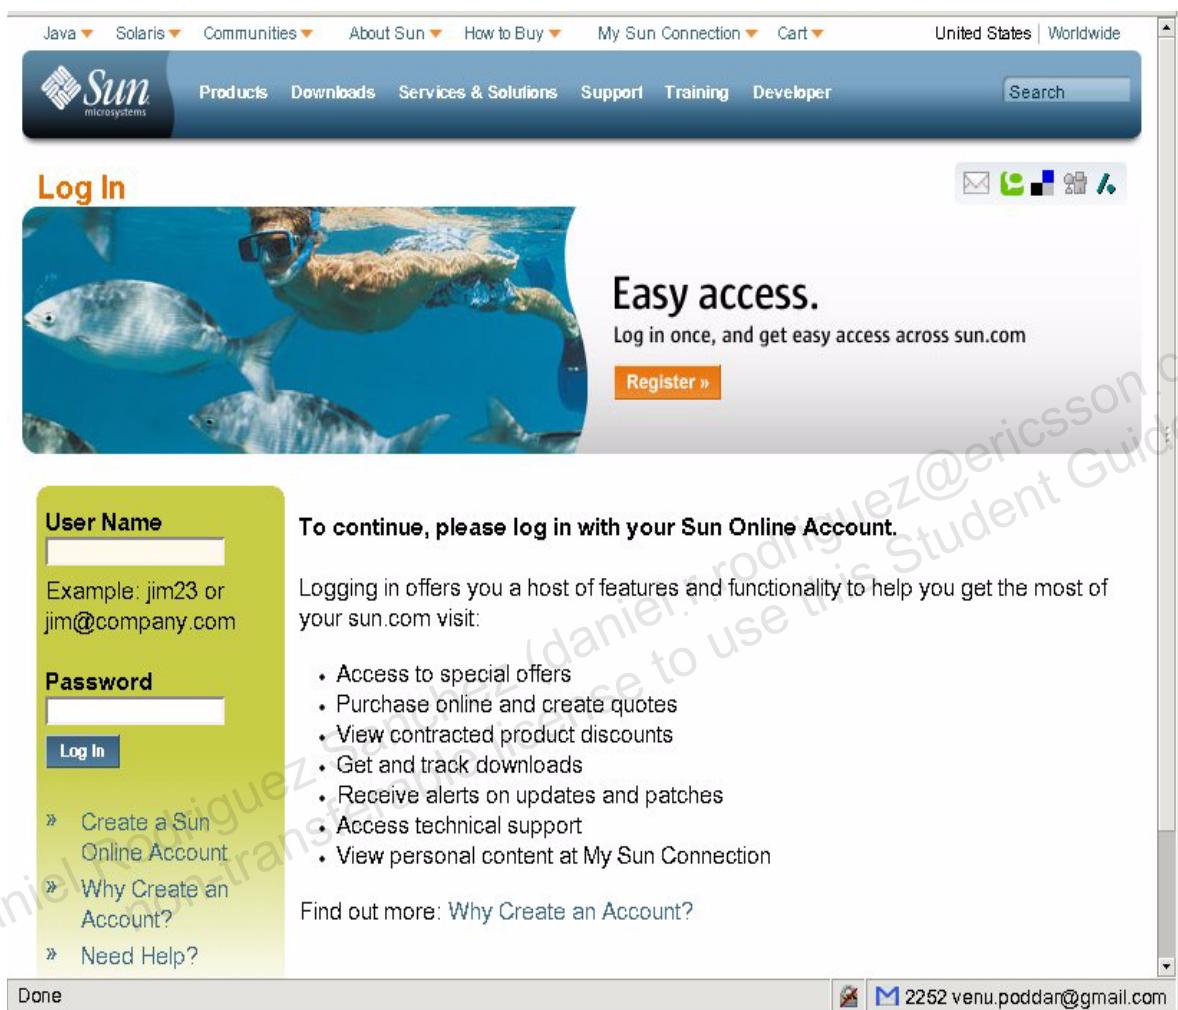
You'll need to repeat this process for each of the patches listed for Solaris 10 in the Latest Patch Utility Patches area.



**Note** – Just so you are aware, the Patch Utilities Patch is also contained in the Recommended and Sun Alert patch clusters; however, you have to have a valid support contract to download either of these patch clusters, located on the Patch Cluster and Patch Bundle Downloads Page.

## Downloading Patches Using SunSolve

To download any patch from SunSolve you must login using a Sun Online account. Figure 8-1 shows the Sun Online Account login screen.



**Figure 8-1** The Sun Online Account Login Screen



## Patch Cluster and Patch Bundle Downloads

Patch Clusters and Patch Bundles are popular collections of patches providing specific functionality. They contain an install script for ease of installation. A Sun support contract which covers Solaris is required to access or download a Patch Cluster or Patch Bundle containing Solaris OS patches, and they may only be installed and used on systems covered by an active Sun support contract for Solaris.

---

**Note** – The Patch Clusters and Patch Bundle Downloads page replaces both the Recommended Patch Cluster and Recommended and Security Patches pages.

---

Once you have logged into SunSolve and the Patches and Updates page appears, scroll down to the Downloads section and locate the Patch Cluster and Patch Bundle Downloads link. Click on this link to access this page.

The Patch Cluster and Patch Bundle Downloads page contains all the patch clusters.

The Solaris Patch Clusters include:

- Recommended Patch Clusters

The Solaris Recommended Patch Clusters contain Solaris Operating System patches that include the latest revision of the patch and package utility patches to help ensure correct patching operations. The latest revision of any Solaris OS patch which address Sun Alert issues. These are patches which fix Security, Data Corruption, or System Availability issues. And any patch that is required to correctly install the above patches.

- Sun Alert Patch Clusters

The Solaris Sun Alert Patch Clusters contain Solaris Operating System patches that include the latest revision of the patch and package utility patches to help ensure correct patching operations. The minimum revision of any Solaris Operating System patch which address Sun Alert issues. These are patches which fix Security, Data Corruption, or System Availability issues. And any patch that is required to correctly install the above patches. The clusters are updated whenever a patch meeting the above inclusion criteria is released.

- Solaris Update Patch Bundles

These Solaris Patch Bundles contain the equivalent set of patches that are pre-applied to the corresponding Solaris Update image. They are provided for customers whose Change Control procedures make it difficult to upgrade but OK to patch. Applying these Patch Bundles is not equivalent to installing or upgrading to the corresponding Solaris Update as these patch bundles don't typically contain any new packages contained in the corresponding Solaris Update.



**Note** – After applying one of these Patch Bundles, you must apply the Sun Alert or Recommended Patch Cluster to obtain all available Solaris OS patches which address Security, Data Corruption, and System Availability issues released since these Patch Bundles were created.

---

- Solaris Live Upgrade Bundles

The Live Upgrade Zones Starter Patch Bundles are provided specifically for use on systems with non-global zones that are running Solaris 10 8/07 (Update 4) or an earlier Solaris 10 Update.

### Accessing Sun Alert Patch Clusters

On the Patch Cluster and Patch Bundle Download page, scroll down to the Solaris Patch Clusters and click on the Sun Alert Patch Clusters link to view the available patch clusters.

Patch Cluster Name	Cluster Size	View README	Download Cluster
Solaris 10 SPARC Sun Alert Patch Cluster	1103 MB	README	Download
Solaris 10 x86 Sun Alert Patch Cluster	838 MB	README	Download
Solaris 9 SPARC Sun Alert Patch Cluster	507 MB	README	Download
Solaris 9 x86 Sun Alert Patch Cluster	357 MB	README	Download
Solaris 8 Sun Alert Patch Cluster	277 MB	README	Download
Solaris 8 x86 Sun Alert Patch Cluster	158 MB	README	Download



**Note** – The **Release Date** field for each of the patch clusters listed above has been omitted from this example.

---

If you have a SPARC system, select the Solaris 10 SPARC Sun Alert Patch Cluster. After you select the cluster, you have the option to view the README file or to download the patch cluster. Click on README and review the file before you download the patch cluster.

The patch cluster README files have a different look from the individual patch README files. You will want to take notice of what patches are included in the cluster, the Important Notes and Warnings, and the install instructions.

When you have finished reading the patch cluster's README file, you will return to the Patch Cluster and Patch Bundle Download Page. At this point you can download the Solaris 10 SPARC Sun Alert Patch Cluster.

---

**Note** – In applying Sun Alert Patch Clusters be aware that any `reconfigimmediate` and `rebootimmediate` patches require a reboot of the system before normal operations can be resumed. If you use Solaris Live Upgrade to install a patch cluster, you will not need to reboot the system.

---

## Checking Patch Levels

Before installing operating system patches, determine which patches are already installed on your system.

The showrev command and the patchadd command provide useful information about currently installed patches. The patchadd -p command can only be run as root. A normal user can use showrev -p. The command output is the same for patchadd -p and showrev -p. The patchadd command sometimes takes longer to display information.

```
showrev -p
```

```
Patch: 118344-14 Obsoletes: 122397-01 Requires: Incompatibles:
Packages: SUNWcsu, SUNWcsl, SUNWckr, SUNWhea, SUNWarc, SUNWfmd
Patch: 118368-04 Obsoletes: Requires: Incompatibles: Packages: SUNWcsu
```

```
. . .
```

```
patchadd -p
```

```
Patch: 118344-14 Obsoletes: 122397-01 Requires: Incompatibles: Packages:
SUNWcsu SUNWcsl SUNWckr SUNWhea SUNWarc SUNWfmd
Patch: 118368-04 Obsoletes: Requires: Incompatibles: Packages: SUNWcsu
```

```
. . .
```

The /var/sadm/patch directory includes the following information:

- Historical information about patches currently installed on a system
- A listing of patches you can remove using the patchrm command

```
ls /var/sadm/patch
```

```
107558-05 107594-04 107630-01 107663-01 107683-01 107696-01
107817-01 107582-01 107612-06 107640-03
```




---

**Caution** – Deleting files from the /var/sadm directory to make more space is a Solution Center call generator. The only way to correct the problems that occur is to restore the deleted files from backup tapes or to reload the Solaris OS.

---




---

**Note** – Allocated enough space for the /var file system. There must be sufficient space for the /var/sadm directory to grow as new software packages and patches are installed on the system.

---

## Preparing Patches for Installation

When patches are downloaded to the local system, you must place the patches in a temporary directory to prepare them for installation. Use the /var/tmp directory.

One of the common reasons for patch installation failure is directory permission or ownership problems. The /var/tmp directory is open to all and eliminates any of these types of problems.

The Solaris 7, Solaris 8, Solaris 9, and Solaris 10 OS unsigned patches are in zip format, for example, 105050-01.zip.

Use the unzip command to unpack the patch files.

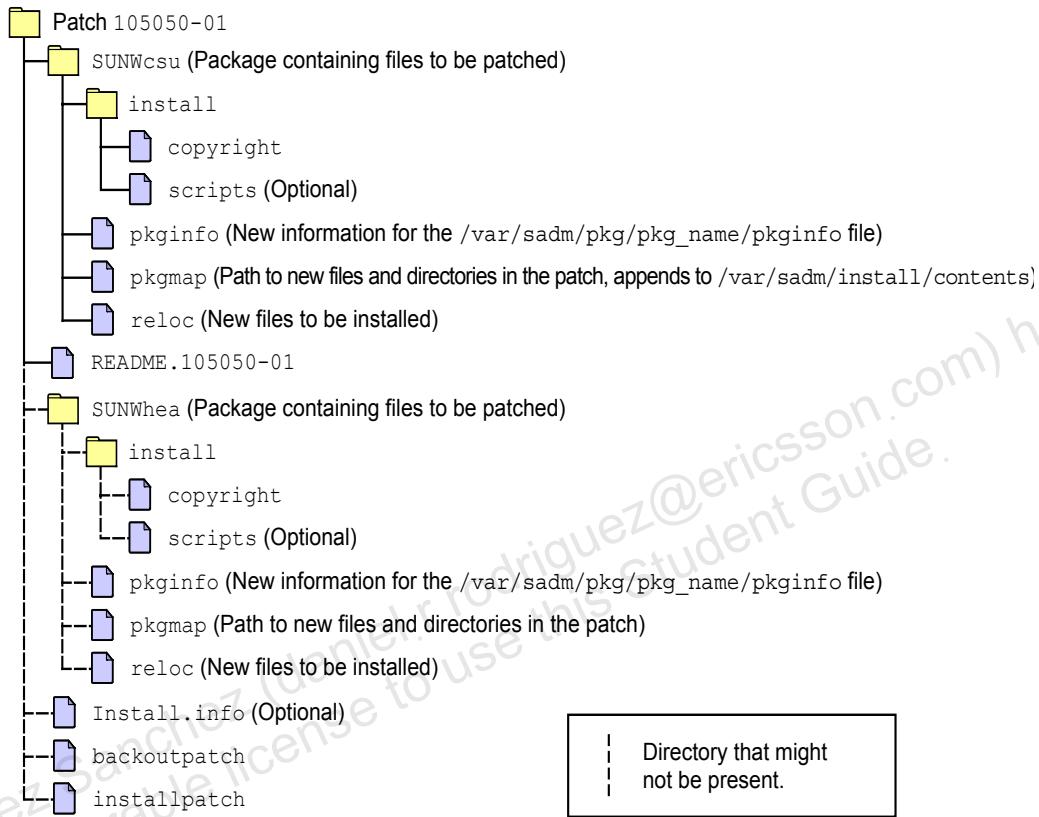
```
/usr/bin/unzip 124998-01.zip
 creating: 124998-01/
 inflating: 124998-01/.diPatch
 inflating: 124998-01/patchinfo
 creating: 124998-01/SUNWcsu/
 inflating: 124998-01/SUNWcsu/pkgmap
 inflating: 124998-01/SUNWcsu/pkginfo
 creating: 124998-01/SUNWcsu/install/
 inflating: 124998-01/SUNWcsu/install/checkinstall
 inflating: 124998-01/SUNWcsu/install/copyright
 inflating: 124998-01/SUNWcsu/install/i.none
 inflating: 124998-01/SUNWcsu/install/patch_checkinstall
 inflating: 124998-01/SUNWcsu/install/patch_postinstall
 inflating: 124998-01/SUNWcsu/install/postinstall
 inflating: 124998-01/SUNWcsu/install/preinstall
 creating: 124998-01/SUNWcsu/reloc/
 creating: 124998-01/SUNWcsu/reloc/usr/
 creating: 124998-01/SUNWcsu/reloc/usr/bin/
 inflating: 124998-01/SUNWcsu/reloc/usr/bin/tip
 inflating: 124998-01/README.124998-01
 inflating: 124998-01/LEGAL_LICENSE.TXT
#
```

Earlier versions of the Solaris OS used patches in compressed tar archive format, for example, 101010-01.tar.z. For these patches, use the zcat command to uncompress the tar archive, and the tar command to extract the patch files from the tar archive.

```
/usr/bin/zcat 101010-01.tar.z | tar xvf -
```

## Patch Contents

Figure 8-2 shows the contents of a patch directory after it is extracted from the **.zip** file.



**Figure 8-2** An Extracted Patch Directory

# Installing and Removing Patches

The commands used to manage patches are:

- **patchadd** – Installs uncompressed patches to the Solaris OS
- **patchrm** – Removes patches installed on the Solaris OS
- **install\_cluster** – Installs the patches delivered in patch clusters
- **smpatch** – Manages the update process on a single system

## Installing a Patch

When you install a patch, the **patchadd** command calls the **pkgadd** command to install the package files that the patch delivers.

The following example shows installing a patch using the **patchadd** command. This example assumes that the patch to be installed exists in the **/var/tmp** directory and has been unzipped for installation.

---

**Note** – The patches are delivered as compressed tar files.

---

```
cd /var/tmp
patchadd 124998-01
Validating patches...
```

Loading patches installed on the system...

Done!

Loading patches requested to install.

Done!

Checking patches that you specified for installation.

Done!

Approved patches will be installed in this order:

124998-01

Checking installed patches...

Verifying sufficient filesystem capacity (dry run method) ...

## Installing and Removing Patches

## Installing patch packages...

Patch 124998-01 has been successfully installed.  
See /var/sadm/patch/124998-01/log for details

Patch packages installed:

SUNWcsu

#



**Caution** – Patches can be added using the `-d (nosave)` option to save space. When this option is used, `patchadd` does not save copies of the files being updated or replaced. Patches added in this way cannot be backed out.

## Removing a Patch

When you remove a patch, the `patchrm` command restores all files that were modified or replaced by that patch, unless:

- The patch was installed with the patchadd -d option (which instructs the patchadd command not to save copies of files being updated or replaced)
  - The patch is required by another patch
  - The patch has been obsoleted by a later patch

The `patchrm` command calls the `pkgadd` utility to restore packages that were saved during the initial patch installation.

The following example shows how to remove a patch using the `patchrm` command.

```
patchrm 124998-01
Validating patches...
```

Loading patches installed on the system...

Done!

Checking patches that you specified for removal.

Done!

Approved patches will be removed in this order:

```
124998-01
```

```
Checking installed patches...
```

```
Backing out patch 124998-01...
```

```
Patch 124998-01 has been backed out.
```

```
#
```

## Installing Patch Clusters

Patch clusters provide a selected set of patches for a designated Solaris OS level, conveniently wrapped for one-step installation. Patch clusters usually consist of a set of recommended and security patches.

You should not install cluster patches on systems with limited disk space. Patch clusters often consume substantial amounts of disk space. For example, the Solaris 10\_x86 Recommended Patch Cluster, made available February 13, 2007, is approximately 176 Mbytes in its zipped form, and 464 Mbytes in its unzipped form.

By default, the cluster installation procedure saves the base objects being patched. Prior to installing the patches, the cluster installation script first determines if enough system disk space is available in the /var/sadm/pkg directory to save the base packages and terminates if not enough space is available.

---

**Caution** – You can override the save feature by using the -nosave option when you are executing the cluster installation script. If you use the -nosave option, you cannot back out these patches if the need arises.

---



You can remove individual patches that were installed by the patch cluster using the patchrm command. README files for the patches installed by a patch cluster are located in the patch directories that the installation process creates below the /var/sadm/patch directory.

It is generally recommended to install patch clusters with the system operating in single user mode.

To install a patch cluster, perform the following steps:

1. Be sure the patch cluster has been unzipped, or uncompressed and extracted if the cluster was received as a `.tar.Z` file. For example, on a Sparc system:  
`# unzip 10_Recommended.zip`  
On an x86/x64 system:  
`# unzip 10_x86_Recommended.zip`
2. Decide which method to use to install the cluster—using the recommended default save option, or the `-nosave` option.
3. Change to the directory that contains the patch cluster. Read the `CLUSTER_README` file, which contains information about the bundled set of patches, including:
  - Cluster description
  - Patches included
  - Important notes and warnings
  - Save and backout options
  - Special install instructions
  - Special patch circumstances
  - Any notices and other recommendations

---

**Caution** – Recent Solaris 10 patch clusters exceed the capability of the `unzip` utility shipped with Solaris 10 versions prior to Solaris 10 10/08. Install the latest `unzip` patch before you extract files.

---

For example:

```
cd 10_x86_Recommended
more CLUSTER_README
CLUSTER_README
```

NAME: Solaris 10\_x86 Recommended Patch Cluster  
DATE: Feb/13/07

```
#####
```

This patch cluster is intended to provide a selected set of patches for the designated Solaris release level. This is a bundled set of patches conveniently wrapped for one-step installation. Only install this cluster on the appropriate Solaris system. Carefully read all important notes and install instructions provided in this README file before

installing the cluster. A cluster grouping does not necessarily imply that additional compatibility testing has occurred since the individual patches were released.

WARNING!! IT IS HIGHLY RECOMMENDED that the installation of this patch cluster be performed in single-user mode (Run Level S).

```
#####
#
```

#### CLUSTER DESCRIPTION

---

These Solaris Recommended patches are considered the most important and highly recommended patches that avoid the most critical system, user, or security related bugs which have been reported and fixed to date. In

4. Run the `install_cluster` script.

```
./install_cluster
```

Patch cluster install script for Solaris 10\_x86 Recommended Patch Cluster

\*WARNING\* SYSTEMS WITH LIMITED DISK SPACE SHOULD \*NOT\* INSTALL PATCHES: With or without using the save option, the patch installation process will still require some amount of disk space for installation and administrative tasks in the /, /usr, /var, or /opt partitions where patches are typically installed. The exact amount of space will depend on the machine's architecture, software packages already installed, and the difference in the patched objects size. To be safe, it is not recommended that a patch cluster be installed on a system with less than 4 MBytes of available space in each of these partitions. Running out of disk space during installation may result in only partially loaded patches. Check and be sure adequate disk space is available before continuing.

Are you ready to continue with install? [y/n]: **y**

(output omitted)

The following patches were able to be installed:

```
120720-02
122035-05
119255-34
123840-04
```

## Installing and Removing Patches

---

(output omitted)

```
ERROR: The following patches were not able to be installed:
121297-01
120901-03
118873-04
121334-04
```

(output omitted)

For more installation messages refer to the installation logfile:

```
/var/sadm/install_data/Solaris_10_x86_Recommended_Patch_Cluster_log
```

Use '/usr/bin/showrev -p' to verify installed patch-ids.

Refer to individual patch README files for more patch detail.

Rebooting the system is usually necessary after installation.

```
#
```

5. Read individual patch README files to determine if any additional steps are required to fully install any individual patch.
6. Check the patch cluster log file if more detail is needed.

Reviewing the log provides information about why the patches listed above were not able to be installed:

```
more
/var/sadm/install_data/Solaris_10_x86_Recommended_Patch_Cluster_log
*** Install Solaris 10_x86 Recommended Patch Cluster begins Sun Feb 25
22:01:19 MST 2007 ***
*** PATCHDIR = /patches/10_x86_Recommended ***

Installing 120720-02...
Validating patches...

Loading patches installed on the system...
```

Done!

Loading patches requested to install.

Done!

The following requested patches have packages not installed on the system. Package SUNWgzipS from directory SUNWgzipS in patch 120720-02 is not installed on the system. Changes for package SUNWgzipS will not be applied to the system.

Checking patches that you specified for installation.

Done!

Approved patches will be installed in this order:

```
120720-02
(output omitted)
#
```

7. Reboot the system for all patches to take effect.



**Note** – For recent Solaris 10 patch clusters, you must use an option with the `install_cluster` command. Refer to the README to find out which option you must use.

## Patch Administration From the Command Line (CLI)

An update (also known as a patch) contains a collection of files and directories. This collection replaces existing files and directories that prevent proper execution of the software. Some updates contain product enhancements.

A Solaris OS update types include:

- Standard updates – Updates that fix specific problems with the Solaris OS and other Sun hardware and software products.
- Recommended patches – Solaris OS updates that fix problems that might occur on a large percentage of systems.
- Update clusters – A group of standard, recommended, or security updates that have been bundled into a single archive for easy downloading and installation.

---

**Note** – In previous versions of the Solaris OS, maintenance updates were also available. These were sets of patches that had been tested together and packaged for one-step installation. Maintenance updates were available to service contract customers. Maintenance updates are now replaced by the Solaris OS distributions of the Solaris Express Program. Such updates to the Solaris OS are free for download and are available on a monthly schedule. See the following URL for details about the Solaris Express program:

<http://www.sun.com/software/solaris/solaris-express/>

---

An update is distributed as a directory that is identified by a unique number. The number assigned to an update includes the update base code first, a hyphen, and a number that represents the update revision number. For example, an update directory named 105050-01, indicates that 105050 is the base code and 01 is the revision number.

---

**Note** – Not all updates from Sun Microsystems must be installed. Only install the recommended and security updates, and those required to fix problems specific to your site.

---



## Using the `smpatch` Command

The `smpatch` command (and its sub-commands) are the preferred commands to use for update/patch management using the CLI.

Starting with the Solaris 9 OS, the `smpatch` command is available in two modes, local mode and remote mode:

- Local mode can only be run on the local system in single-user or multiuser mode.
- Remote mode can be used to perform tasks on remote systems, using the `-n system_name` option with `smpatch` commands.

By default, `smpatch` runs in local mode. In local mode none of the authentication options or options that refer to remote systems are available.

In releases earlier than Solaris 10, if you specify any of the remote or authentication options (except for `-L`), remote mode is used.

Remote mode is currently not supported in the Solaris 10 OS.

In Solaris 8 only local mode `smpatch` is available.

## Phases for Applying Updates

The full sequence for applying an update involves these phases or steps:

1. Analyzing your system.
2. Downloading the necessary updates.
3. Applying the updates.

You can exercise as much control of the phases as needed:

- The `smpatch update` command will perform all three functions using one command.

This command requires multiuser mode and will not apply an update that has interactive as one of its property install types. The application of updates will be governed by the update policy.

- The `smpatch analyze` and `smpatch update` commands will perform all three functions using two commands.

## Using the smpatch Command

---

If you want to first analyze your system and then download and apply them in a single subsequent step, first use the `smpatch analyze` command followed by the `smpatch update` command.

The `smpatch analyze` command requires multiuser mode.

The `smpatch update` command will also download any prerequisite patches.

- The `smpatch analyze`, `smpatch download`, and `smpatch add` commands will perform all three functions using three commands.

If you want to analyze your system, download the updates and add them to your system in three separate steps, first use the `smpatch analyze` command followed by the `smpatch download` command followed by the `smpatch add` command.

The `smpatch add` command can be used in single user mode or multiuser mode. The `smpatch add` command will not consult the update policy.

## Applying an Update In Three Steps

Using three commands gives more control and flexibility when you apply a patch.

1. Assume that you want to have the latest update(s) for the North America locales. The following command will analyze your local system and determine the appropriate available updates for it. It will not download or apply them. The command will write the list to the file `plist`. Look in the `plist` file for updates involving locales.

```
smpatch analyze > plist
vi plist
...
119397-06 SunOS 5.10: patch for North America region locales issues
...
```



The `patchadd -p` command shows what updates were applied to the system. Use it to verify that the locales update you found in the plist file is not already on the system:

```
patchadd -p | grep 119397
#
```

---

**Note** – You can still use `showrev -p` to accomplish the same thing, and it executes more quickly.

---

2. The following command will download (but not apply) the new update for the locales issue:

```
smpatch download -i 119397-06
119397-06 has been validated.
```

The update was downloaded to the download area and validated. By default, this directory is `/var/sadm/spool`. If it was changed from the default, query the system with `smpatch get` to find the new location. The default is still in effect:

```
smpatch get | grep download
patchpro.download.directory - /var/sadm/spool
```

The following commands show the update has been downloaded as a \*.jar file:

```
cd /var/sadm/spool; ls
119397-06.jar
...
```

3. Apply or install this update using the following `smpatch add` command:

```
smpatch add -i 119397-06
add patch 119397-06
Patch 119397-06 has been successfully installed.
```

## Using the smpatch Command

Verify that the patch is installed on your system using this command:

```
patchadd -p | grep 119397-06
Patch: 119397-06 Obsoletes: Requires: 121734-01 Incompatibles: Packages:
SUNWnameos SUNWnamdt SUNWnamow
```

A subsequent analysis of this system will no longer show this update as appropriate:

```
smpatch analyze | grep 119397-06
#
```

An update is easily removed (backed out). The following command does this for the update just applied:

```
smpatch remove -i 119397-06
remove patch 119397-06
Transition old-style patching.
Patch 119397-06 has been backed out.
```

An analysis now shows that this update is once again appropriate and available for this system:

```
smpatch analyze | grep 119397-06
119397-06 SunOS 5.10: patch for North America region locales issues
```

## Applying an Update In One Step

Use the smpatch update to analyze your system, download and apply the update in one step. For example this awk/nawk patch can be applied as follows:

```
smpatch update -i 118815-05
118815-05 has been validated.
Installing patches from /var/sadm/spool...
118815-05 has been applied.
/var/sadm/spool/patchpro_dnld_2007.03.16@12:36:36:MST.txt has been moved
to
/var/sadm/spool/patchproSequester/patchpro_dnld_2007.03.16@12:36:36:MST.t
xt
```

# Configuring the Patch Management Environment

The `smpatch get`, `smpatch set` and `smpatch unset` commands are used to configure the patch management environment:

- `smpatch get` displays the current settings for environment parameters.
- `smpatch set` changes values for environment parameters. The new values are not validated in anyway so verify the intended changes.
- `smpatch unset` enables the default values for environment parameters.

## Examples for Configuring Patch Management Environment

Use the following command to display the current environment parameter values:

```
smpatch get
patchpro.backout.directory - ""
patchpro.baseline.directory - /var/sadm/spool
patchpro.download.directory - /var/sadm/spool
patchpro.install.types - " "
rebootafter:reconfigafter:standard
patchpro.patch.source http://192.168.201.1:3816/solaris/
https://getupdates1.sun.com/solaris/
patchpro.patchset - current
patchpro.proxy.host - ""
patchpro.proxy.passwd **** ****
patchpro.proxy.port - 8080
patchpro.proxy.user - "
```




---

**Note** – The `smpatch(1M)` man page contains a detailed description of the environment parameters.

---

The following `smpatch set` and `get` commands will set a new value for the update source. (This is typically what you would do to direct your local client to a new update proxy server.)

```
smpatch set
patchpro.patch.source=http://newproxy.apex.com:3816/solaris/
smpatch get
```

## Configuring the Patch Management Environment

```
patchpro.backout.directory - ""
patchpro.baseline.directory - /var/sadm/spool
patchpro.download.directory - /var/sadm/spool
patchpro.install.types - ""
rebootafter:reconfigafter:standard
patchpro.patch.source http://newproxy.apex.com:3816/solaris/
https://getupdates1.sun.com/solaris/
patchpro.patchset - current
patchpro.proxy.host - ""
patchpro.proxy.passwd **** *****
patchpro.proxy.port - 8080
patchpro.proxy.user - ""
```

You can also set the source of updates to a local or remote directory as the following examples illustrate:

```
smpatch set patchpro.patch.source=file:/net/sys-04/export/updates
smpatch set patchpro.patch.source=file:/local/updates
smpatch set patchpro.patch.source=file:/cdrom/cdrom0
```

The following command sets the patchpro.patch.source parameter back to the default value:

```
smpatch unset patchpro.patch.source
smpatch get
patchpro.backout.directory - ""
patchpro.baseline.directory - /var/sadm/spool
patchpro.download.directory - /var/sadm/spool
patchpro.install.types - rebootafter:reconfigafter:standard
patchpro.patch.source - ""
https://getupdates1.sun.com/solaris/
patchpro.patchset - current
patchpro.proxy.host - ""
patchpro.proxy.passwd **** *****
patchpro.proxy.port - 8080
patchpro.proxy.user - ""
```

You can configure an update set which defines a subset of updates that commands will work with. For example, the following commands will result in an analysis only on recommended updates:

```
smpatch set patchpro.patchset=recommended
smpatch analyze
```

## Using the Update Policy for Applying Updates

The `patchpro.install.types` property defines the update policy in effect for the update management environment.

When you apply patches using the `smpatch update` command the update policy is consulted before an update is actually applied.

The following are the types of updates that are applied to the system:

- Standard updates that are applied immediately and require no system restart
- Updates that require a system restart

If you use the `smpatch update` command to update your system, you get the benefit of the guidelines established by update/patch developers in how best to apply the update. However, you can customize the policy for applying updates using the `patchpro.install.types` parameter.

Table 8-1 shows the `patchpro.install.types` parameter values and also describes the value and if it is part of the default update policy.

**Table 8-1** Install Type Parameter Values

<b>patchpro. install. types value</b>	<b>Description</b>
standard	A default. Can be applied in multiuser mode and visible immediately unless the application being updated is running while the update is applied. In this case, the effects of the update are visible after the affected application is restarted.
rebootafter	A default. Effects not visible until system reboot.
reconfigafter	A default. Effects not visible until a reconfiguration reboot ( <code>boot -r</code> ). See the <code>boot(1M)</code> man page.
rebootimmediate	System becomes unstable (unpredictable behavior or possible data loss) until system reboot.
reconfigimmediate	System becomes unstable (unpredictable behavior or possible data loss) until system reconfiguration reboot ( <code>boot -r</code> ).

**Table 8-1** Install Type Parameter Values (Continued)

<b>patchpro. install. types value</b>	<b>Description</b>
singleuser	Do not apply this update in multiuser mode. You must apply this update on a quiet system with no network traffic and with extremely restricted I/O activity.
interactive	Only downloaded to your system and must be applied manually according to the instructions in the update's README file.
clientroot   clientuser	Only applies to diskless clients and indicates which portion of the client's filesystem is affected by the update (clientroot indicates / and clientuser indicates /usr)

The default value for this parameter is shown with this **smpatch get** command:

```
smpatch get patchpro.install.types
patchpro.install.types - rebootafter:reconfigafter:standard
```

### Per Update Policy Value

Each update has properties associated with it. The PATCH\_PROPERTIES values are the install types for the update. You can learn these values with the following command sequence if the update is downloaded in the spool area of your system:

```
cd /var/sadm/spool
jar xvf 119578-15.jar 119578-15/patchinfo
inflated: 119578-15/patchinfo
cat 119578-15/patchinfo
PATCHINFOVERSION="1.0"
PATCHID=119578-15
PATCH_CORRECTS='BaseOS.SolarisCore-5.10 BaseOS.SolarisFaultMgmt-5.10'
PATCH_ARCH='sparc'
PATCH_OS='SunOS'
PATCH_OSRELEASE='5.10'
PATCH_PROPERTIES='rebootafter'
PATCH_OBSOLETES="119330-01 119331-01 119559-01 119576-01 120635-02"
```

The patchinfo file is included in the collection of files in an update. It contains more metadata than just the PATCH\_PROPERTIES value.

## Example of Using the Update Policy

It is good practice to always use the update policy when adding patches, otherwise you might miss an important patch property such as immediate reboot or a prerequisite patch.

### Not Using the smpatch update Command

Analyze your system and learn if any updates involving wanboot are appropriate and available:

```
smpatch analyze | grep wanboot
119681-06 SunOS 5.10: wanboot patch
```

Determine if any prior versions of the wanboot update are already on the system:

```
patchadd -p | grep 119681
```

Patch: 119681-05 Obsoletes: Requires: Incompatibles: Packages: SUNWcakr

Only the earlier 05 version of this update is already installed.

Download, but do not apply, the newer wanboot update:

```
smpatch download -i 119681-06
119681-06 has been validated.
```

Apply the patch using the smpatch add command:

```
smpatch add -i 119681-06
add patch 119681-06
...
Validating patches...
Loading patches installed on the system...
Done!
Loading patches requested to install.
Done!
Checking patches that you specified for installation.
Done!
Approved patches will be installed in this order:
119681-06
```

Patch 119681-06 has been successfully installed.

Verify that the patch is installed on your system:

```
patchadd -p | grep 119681
Patch: 119681-05 Obsoletes: Requires: Incompatibles: Packages: SUNWcakr
Patch: 119681-06 Obsoletes: Requires: Incompatibles: Packages: SUNWcakr
```

A subsequent analysis of this system will no longer show this update as appropriate:

```
smpatch analyze | grep 119681-06
#
```

Especially when you use the smpatch add command it is always a good practice to read information about the update. Go to the download spool area and see what information there is about this update:

```
cd /var/sadm/spool; ls
119681-06.jar
cache
patchpro_dnld_2006.02.13@10:10:29:MST.txt
cat *.txt
This patch bundle was generated by PatchPro.
```

Please refer to the README file within each patch for installation instructions. To properly patch your system, the following patches should be installed in the listed order:

1) 119681-06!!! IMMEDIATE REBOOT!!!

The \*.txt and other README files often contain important information. In this case the warning to immediately reboot implies that the PATCH\_PROPERTIES value for install type is either reconfigimmediate or rebootimmediate.

The following command sequence will display the install type value for this update:

```
cd /var/sadm/spool
jar xvf 119681-06.jar 119681-06/patchinfo
inflated: 119681-06/patchinfo
grep PROP 119681-06/patchinfo
PATCH_PROPERTIES='reconfigimmediate'
```

A reconfiguration reboot (`boot -r`) should be done on this system to render it stable again.

---

**Note** – The `/var/adm/messages` file identifies problems that are found when applying a patch to a system.

---



## Using the `smpatch update` Command

The `smpatch update` command will analyze your system, download the update and apply it in one step. It also provides safeguards that are not available with `smpatch add` because it consults the update policy.

The `smpatch update` command also is knowledgeable about update dependencies and applies any dependencies for the updates you specify.

```
smpatch update -i 119681-06
```

119681-06 has been validated.

Installing patches from `/var/sadm/spool...`

NOTICE: Patch 119681-06 cannot be installed until the next system shutdown.

`/var/sadm/spool/patchpro_dnld_2006.02.15@06:02:43:MST.txt` has been moved to

`/var/sadm/spool/patchproSequester/patchpro_dnld_2006.02.15@06:02:43:MST.txt`

`/var/sadm/spool/patchpro_dnld_2006.02.15@06:09:14:MST.txt` has been moved to

`/var/sadm/spool/patchproSequester/patchpro_dnld_2006.02.15@06:09:14:MST.txt`

ID's of the updates that are disallowed by installation policy have been written to file

`/var/sadm/spool/disallowed_patch_list`

One or more updates that you installed requires a system shutdown to activate it. To initiate the system shutdown, you must use one of the following commands:

- o Power down the system - `init 0` or `shutdown -i 0`
- o Drop to the firmware prompt - `init 5` or `shutdown -i 5`
- o Restart the system - `init 6` or `shutdown -i 6`

Recall that `smpatch add` command informs you about the required reboot in the \*.txt in the download spool area. `smpatch update`, on the other hand, displayed this to standard out, creates a `disallowed_patch_list` and gave instructions about the reboot.

```
cat /var/sadm/spool/disallowed_patch_list
119681-06
```

Part of the `smpatch update` command applies the updates. Updates that cannot be applied for some reason are listed in the `disallowed_patch_list`. Typically you attend to updates listed in this file manually.

Verify that the only version of this update installed on the system is the prior version (05):

```
patchadd -p | grep 119681
Patch: 119681-05 Obsoletes: Requires: Incompatibles: Packages: SUNWcakr
```

A subsequent analysis of the system still shows that this patch is available and still appropriate for this system. It is in the spooled area awaiting installation and a system reboot.

```
smpatch analyze | grep wanboot
119681-06 SunOS 5.10: wanboot patch
```

Manually add the patch:

```
smpatch add -i 119681-06
Validating patches...
Loading patches installed on the system...
Done!
Loading patches requested to install.
Architecture for package SUNWcakr from directory SUNWcakr.v in patch
119681-06 differs from the package installed on the system.
Architecture for package SUNWcakr from directory SUNWcakr.us in patch
119681-06 differs from the package installed on the system.
Done!
Checking patches that you specified for installation.
Done!
Approved patches will be installed in this order:
119681-06
Patch 119681-06 has been successfully installed.
```

Verify that it is installed:

```
patchadd -p | grep 119681
Patch: 119681-05 Obsoletes: Requires: Incompatibles: Packages: SUNWcakr
Patch: 119681-06 Obsoletes: Requires: Incompatibles: Packages: SUNWcakr
```

Analyze the system to show that it is no longer appropriate and available:

```
smpatch analyze | grep wanboot
#
```

Since this update is marked as reconfigimmediate, reboot the system with the -r option:

```
ok boot -r
...
```

Remove this update's entry in the `disallowed_patch_list` file so you know you are finished administering this update:

```
cat /var/sadm/spool/disallowed_patch_list
#
```

## Working With Multiple Updates

Many of the `smpatch` sub-commands can be applied to multiple updates. Following are some examples:

Multiple instances of the `-i` option are permitted if you just have a few updates to apply:

```
smpatch update -i 118927-02 -i 118822-15 -i 119681-06
```

A list of update IDs can be listing in a file, one per line, and referenced using the `-x idlist=` option:

```
smpatch update -x idlist=/var/sadm/spool/disallowed_patch_list
```




---

**Note** – If you specify particular patches by using the `-i` or `-x idlist=` options, the list is augmented with patches on which they depend before the update occurs.

The following example shows how to create a list of patches that you actually want to apply from the larger list available and appropriate. It also resolves the dependencies for the updates you want to apply.

Generate the full list of updates available and appropriate for your system:

```
smpatch analyze > my.list
```

Edit my.list and remove the ones you are not interested in:

```
vi my.list
...
```

Analyze just the ones that are left and resolve dependencies:

```
smpatch analyze -x idlist=my.list > /tmp/justdothese.list
```

Apply the updates:

```
smpatch update -x -idlist=/tmp/justdothese.list
```

## Authorization and Authentication

The root user certainly can manage updates. You can designate non root users for these tasks by having them assume a role that includes the Software Installation profile or the solaris.admin.patchmgr.\* authorization.

The following command shows a smpatch get command requiring assumption of the role (-r) called patchman:

```
sys-01> smpatch get -r patchman
```

# Exercise: Maintaining Patches

In this exercise, you use commands to prepare patches for installation, evaluate patch descriptions, and install and remove patches.

## Preparation

This exercise requires the set of patches found in `/opt/ses/lab/patches` directory on student systems. If the patches are located on a server, your instructor will provide its location and information about how to retrieve them.

**Note** – Use the patch archives in the `/opt/ses/lab/patches` directory to understand how to use patch administration commands.



## Task

Complete the following steps

1. Copy the patch archives found in `/opt/ses/lab/patches` directory to the `/var/tmp` directory.

**Note** – If the patch archives are located on a server, follow the instructions your instructor provides to copy them to the `/var/tmp` directory on your system.



2. Change directory to `/var/tmp`.
3. Choose two patches that are named using consecutive numbers, and use the `unzip` command to extract the patches from their `.zip` archives.
4. Use the `uname -r -p` command to display your system's OS release level, and architecture (processor) type.
5. Display the `README` files for the two patches you extracted, and determine which of them is the correct patch for your system architecture. Also verify that the patch you choose is for the Solaris release level in use on your system.

## Exercise: Maintaining Patches

---

6. Use the patchadd command to determine if the patch you identified in the previous step is currently installed on your system.
  7. Verify that the /var/sadm/patch directory is empty.
  8. Add the patch.
  9. Verify that the patch is now installed.
- 
- 

10. Examine the patch installation log file.
11. Remove the patch you just installed, and verify that the patch is no longer installed.

## Exercise Summary



**Discussion** – Take a few minutes to discuss what experiences, issues, or discoveries you had during the lab exercises.

- Experiences
- Interpretations
- Conclusions
- Applications

# Exercise Solutions: Maintaining Patches

This section provides task steps and their solutions.

## Task

Complete the following steps:

1. Copy the patch archives found in /opt/ses/lab/patches to the /var/tmp directory.

**Note** – Use the patch archives in the /opt/ses/lab/patches directory to understand how to use patch administration commands. Use the patch that is appropriate for your system architecture. The patches for this exercise for demonstration only.

```
cd /opt/ses/lab/patches
ls
123456-01.zip 654321-01.zip
cp *.zip /var/tmp
#
```

**Note** – If the patch archives are located on a server, follow the instructions your instructor provides to copy them to the /var/tmp directory on your system.

2. Change directory to /var/tmp.

```
cd /var/tmp
```

3. Extract the patches from their archives. For example:

```
unzip 123456-01.zip
Archive: 123456-01.zip
creating: 123456-01/
...
(output omitted)
unzip 654321-01.zip
Archive: 654321-01.zip
creating: 654321-01/
...
(output omitted)
```

4. Use the `uname -r -p` command to display your system OS release level, and architecture (processor) type.

```
uname -r -p
5.10 sparc
#
```

5. Display the `README` files for the two patches you extracted, and determine which of them is the correct patch for your system architecture. Also verify that the patch you choose is for the Solaris release level in use on your system. For example:

```
more 123456-01/README*
more 654321-01/README*
```

6. Use the `patchadd` command to determine if the patch you identified in the previous step is currently installed on your system. For example:

```
patchadd -p | grep 123456
patchadd -p | grep 654321
```

7. Verify that the `/var/sadm/patch` directory is empty.

```
ls /var/sadm/patch
#
```

8. Add the patch. For example:

```
patchadd 123456-01
patchadd 654321-01
```

9. Verify that the patch is now installed.

What package does the patch affect?

For example:

```
patchadd -p | grep 123456-01
patchadd -p | grep 654321-01
```

10. Examine the patch installation log file. This file is located in the `/var/sadm/patch` directory. For example:

```
cd /var/sadm/patch/123456-01
more log
cd /var/sadm/patch/654321-01
more log
```

11. Remove the patch you just installed, and verify that the patch is no longer installed.

```
cd
patchrm 123456-01
patchrm 654321-01
```

## Notes:

## Module 9

# Using Boot PROM Commands

## Objectives

Upon completion of this module, you should be able to:

- Identify boot programmable read-only memory (PROM) fundamentals
- Use basic boot PROM commands
- Identify the system boot device
- Create and remove custom device aliases
- View and change non-volatile random access memory (NVRAM) parameters from the shell
- Interrupt an unresponsive system

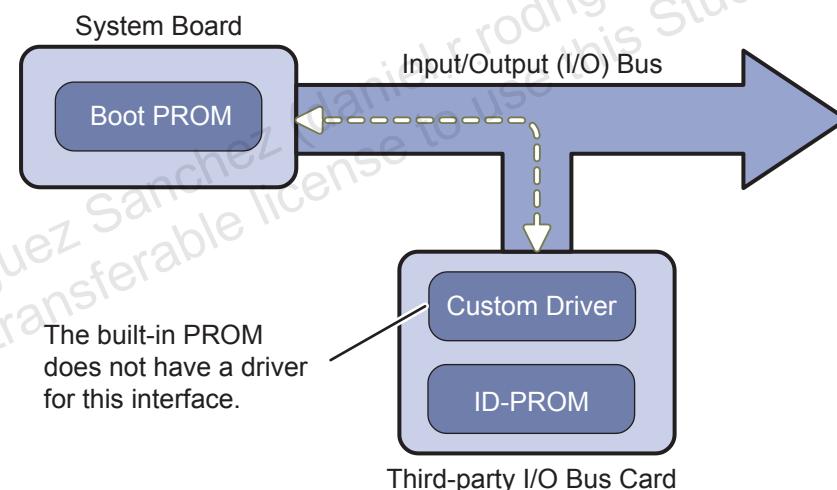
# Introducing Boot PROM Fundamentals

All Sun systems have resident boot PROM firmware that provides basic hardware testing and initialization prior to booting. The boot PROM also enables you to boot from a wide range of devices. In addition, there is a user interface that provides several important functions.

The Sun boot PROM has access to a standard set of generic device drivers. The system needs these drivers to access and control the buses and the boot device to boot the system properly.

All versions of the OpenBoot™ architecture allow a third-party board to identify itself and load its own plug-in device driver. Each device identifies its type and furnishes its plug-in device driver when requested by the OpenBoot PROM during the system hardware configuration phase of the boot process.

Figure 9-1 shows the identification process.



**Figure 9-1** Third-Party Device Identification Process

## Goal of the OpenBoot™ Architecture Standard

The overall goal of the Institute of Electrical and Electronics Engineers (IEEE) standard #1275 for the OpenBoot architecture is to provide the capabilities to do the following:

- Test and initialize system hardware
- Determine the system hardware configuration
- Boot the operating system (OS)
- Provide an interactive interface for configuration, testing, and debugging
- Enable the use of third-party devices

## Boot PROM

Each Sun SPARC system has a boot PROM chip. This 1-Mbyte chip is typically located on the same board as the central processing unit (CPU). Boot PROM chips are usually found in a pluggable socket on older systems. As of the 3.x PROM, they are permanently soldered to the main system board.

The Ultra workstations use a reprogrammable boot PROM called a flash PROM (FPROM). The FPROM allows you to load new boot program data into the PROM by using software, instead of having to replace the chip.

Desktop systems have a write-protect jumper that must be moved before you can write to the PROM. You have to move the jumper because the default position is write-protect. Refer to the *Sun Flash PROM Guide for Workstations and Workgroup Servers - Standalone Version* part number 802-3233-27, for the jumper location on your system.



**Caution** – Many systems have the jumper under an installed frame buffer or other removable card. Be careful when removing or replacing this card.

The main functions of the boot PROM are to test the system hardware and to boot the operating system. The boot PROM firmware is referred to as the *monitor* program.



The boot PROM firmware controls the operation of the system before the operating system has been booted and the kernel is available. The boot PROM also provides the user with a user interface and firmware utility commands, known as the FORTH command set. Commands include the boot commands, diagnostics commands, and commands to modify the default configuration.

**Note** – The boot PROM does not work with the Solaris OS file systems or files. It handles mainly hardware devices. The OS works with and is dependent on firmware, but firmware is independent of the OS.

To determine which revision of OpenBoot PROM is running on the system, you can use the **prtdiag** or **prtconf** command. For example:

```
/usr/platform/`uname -m`/sbin/prtdiag -v

(output omitted)
System PROM revisions:

OBP 4.16.4 2004/12/18 05:21 Sun Blade 1500 (Silver)
OBDIAG 4.16.4 2004/12/18 05:21

or
prtconf -v
OBP 4.16.4 2004/12/18 05:21
#
```

## System Configuration Information

Another important element in each Sun system is the system configuration information. The system configuration information includes the following:

- The Ethernet or MAC address, such as 8:0:20:5d:6f:9e
- The system host ID, such as 805d6f9e
- User-configurable parameters which have been modified from the default settings

The user-configurable parameters are known as NVRAM variables, or EEPROM parameters. They allow an administrator to control things such as the default boot device, the level of Power-on self-test (POST), and so on.

Depending on the system, one of three different components store the system configuration information:

- NVRAM chip
- Serial Electronically Erasable Programmable Read Only Memory (SEEPROM) chip
- System Configuration Card (SCC)

### NVRAM Chip

Older systems contain a removable NVRAM chip, normally located on the main system board. In addition to the system configuration information, the NVRAM chip contains an integrated lithium battery which provides battery backup for the configuration information and also provides the system's time-of-day (TOD) function.

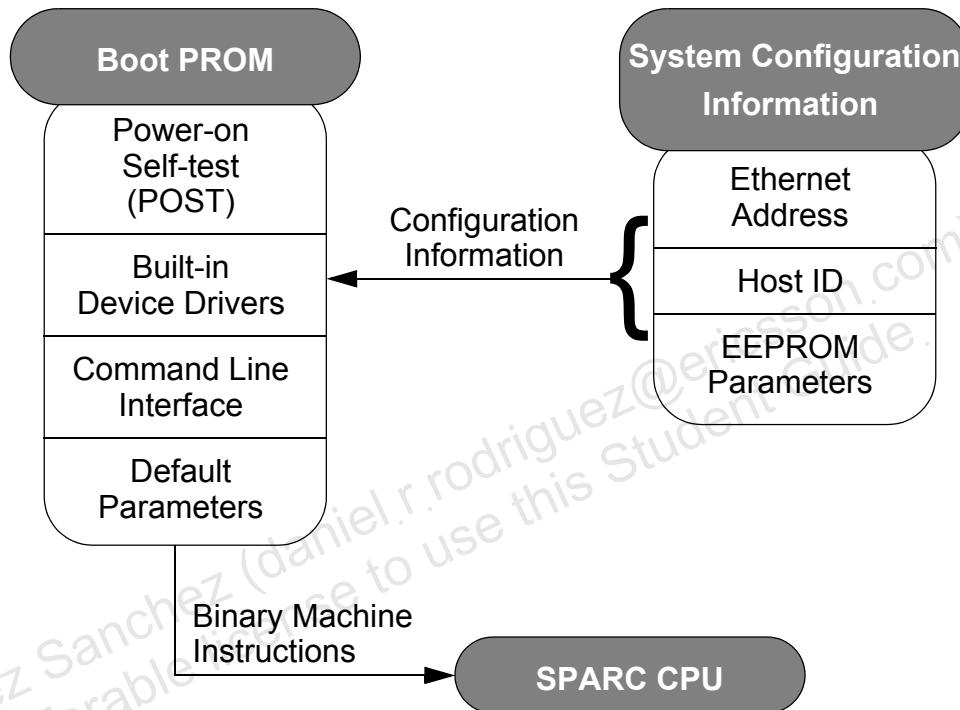
### SEEPROM Chip

Most newer systems contain a non-removable SEEPROM chip, normally located on the main system board. SEEPROM chips do not require a battery to maintain the system configuration information.

## System Configuration Card

Some newer systems contain a removable System Configuration Card which holds the system configuration information. It is inserted into the System Configuration Card Reader.

Figure 9-2 shows the basic elements of the Boot PROM and NVRAM



**Figure 9-2** Basic Elements of the Boot PROM and NVRAM POST

When a system's power is turned on, a low-level POST is initiated. This low-level POST code is stored in the boot PROM and is designed to test the most basic functions of the system hardware.

At the successful completion of the low-level POST phase, the boot PROM firmware takes control and performs the following initialization sequence:

- Probes the memory and then the CPU
- Probes bus devices, interprets their drivers, and builds a device tree
- Installs the console

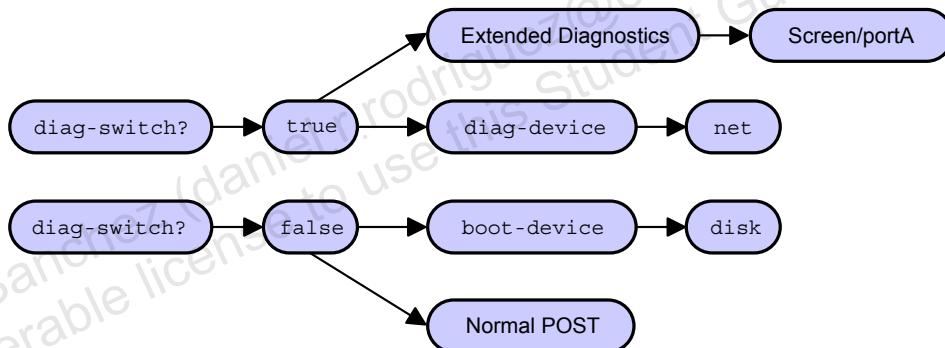
After the boot PROM initializes the system, the banner displays on the console. The system checks parameters stored in the boot PROM and NVRAM to determine if and how to boot the operating system.

One of the first tests that POST runs is to check to determine if a keyboard is connected to the system and if a Stop-key option is present.

The Stop-key is located on the left side of the keyboard. To enable various diagnostic modes, hold down the Stop-key simultaneously with another key. The Stop-key sequences have an effect on the OpenBoot PROM and define how POST runs when the system power is on. The following is a list of the Stop-key sequences:

- Stop-D key sequence – Hold down the Stop and D keys simultaneously while system power is turned on, and the firmware automatically switches to diagnostic mode. This mode runs more extensive POST diagnostics on the system hardware. The OpenBoot PROM variable diag-switch? is set to true.

See Figure 9-3 to show the effect of the variable diag-switch?.



**Figure 9-3** Post Diagnostics

**Note** – The Stop-D key sequence is not available on a serial port terminal.



- Stop-N key sequence – Hold down the Stop and N keys simultaneously while the system power is turned on to set the NVRAM parameters to the default values. When you see the light emitting diodes (LEDs) on the keyboard begin to flash, you can release the keys, and the system should continue to boot.

Incorrect NVRAM settings can cause system boot failure. For example, during a flash PROM download, if a power failure occurs, some of the contents of the NVRAM can become unusable.

If the system does not boot and you suspect that the NVRAM parameters are set incorrectly, the parameters can easily be changed to the default values.



**Caution** – Where possible, capture non-default NVRAM values before using the Stop-N key sequence.

---

- Stop-A key sequence – Hold down the Stop and A keys simultaneously to interrupt any program that is running at the time these keys are pressed and to put the system into the command entry mode for the OpenBoot PROM. The system presents an *ok* prompt for the user, which signifies it is ready to accept OpenBoot PROM commands.



**Caution** – The Stop-A key sequence, as a method for getting to the *ok* prompt, is not recommended unless there is no alternative. The Stop-A key sequence can cause Solaris OS file system corruption which can be difficult to repair.

---

## Disabling the Abort Sequence

You may want to disable the abort key sequence on a system to prevent possible corruption of a file system or to provide tighter security. To disable the abort key sequence do the following. When you complete these steps, the system allows the Stop-A key sequence only during the boot process.

1. Edit the `/etc/default/kbd` file.

Inside the file, the statement `KEYBOARD_ABORT=disable` is commented out.

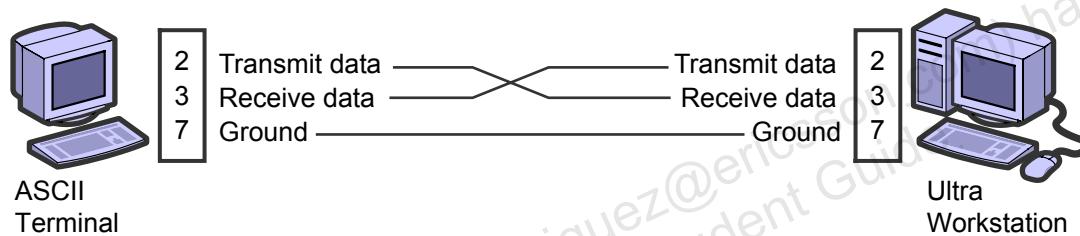
2. Remove the comment character from the beginning of the line.
3. Save the file
4. Execute the command `kbd -i`.

You can also configure the system to change the keyboard abort sequence to an alternate keystroke. Review the `kbd` command man page for more information.

## Displaying POST to the Serial Port

Attach a terminal to the serial port of a system to capture more information from the POST output. When the power is on, POST looks for a keyboard. If there is no keyboard present, POST diverts system output to serial port A. POST runs more extensive tests when the system is in diagnostic mode with the PROM parameter diag-switch? set to true. Attach the correct type of null modem cable for your system type to serial port A.

Some systems require a serial port cable, as shown in Figure 9-4.



**Figure 9-4** Serial Port Connection to a Sun Workstation

The following example is the POST output from a Sun Blade™ 1000:

```
@(#) 4.0 Version 29 created 2000/07/12 16:46
Clearing TLBs Done
Reset: 0000.0000.0000.0010 SPOR
Loading Configuration
Membase: 0000.0000.0000.0000
MemSize: 0000.0000.2000.0000
Init CPU arrays Done
Init E$ tags Done
Setup TLB Done
MMUs ON
Block Scrubbing Done
Copy Done
PC = 0000.07ff.f000.3138
Decompressing Done
Size = 0000.0000.0006.e3b0
ttya initialized
Start Reason: Soft Reset
System Reset: (SPOR)
Probing gptwo at 0,0 SUNW,UltraSPARC-III (750 MHz @ 5:1, 8 MB)
 memory-controller
Probing gptwo at 1,0 Nothing there
Probing gptwo at 8,0 pci pci upa ppm
```

## Introducing Boot PROM Fundamentals

---

Loading Support Packages: kbd-translator

Loading onboard drivers: ebus flashprom bbc ppm i2c dimm-fru dimm-fru  
dimm-fru dimm-fru dimm-fru dimm-fru dimm-fru nvram idprom  
i2c cpu-fru temperature fan-control card-reader motherboard-fru

Memory Configuration:

Segment @ Base: 0 Size: 512 MB (2-Way)

Probing /upa@8,480000 Device 0,0 Nothing there

Probing /upa@8,480000 Device 1,0 Nothing there

Probing /pci@8,600000 Device 4 SUNW,qlc fp disk

Probing /pci@8,600000 Device 1 Nothing there

Probing /pci@8,700000 Device 5 network firewire usb

dev-descrip

next-add

node made

Probing /pci@8,700000 Device 6 scsi disk tape scsi disk tape

Probing /pci@8,700000 Device 1 Nothing there

Probing /pci@8,700000 Device 2 Nothing there

(UltraSPARC-III) , Keyboard Present

OpenBoot 4.0, 512 MB memory installed, Serial #12134217.

Ethernet address 8:0:20:b9:27:49, Host ID: 80b92749.

# Using Basic Boot PROM Commands

The boot PROM monitor provides a user interface for invoking OpenBoot commands.



**Caution** – Modifying OPB contents can seriously damage your system. Use caution when modifying settings.



**Note** – The `ok` prompt indicates that the Solaris OS is currently not running.

Table 9-1 shows typical commands entered at the `ok` prompt.

**Table 9-1** Typical Commands Used at the `ok` Prompt

Command	Description
<code>banner</code>	Displays the power-on banner
<code>boot</code>	Boots the system
<code>help</code>	Lists the main help categories
<code>printenv</code>	Displays all parameter current and default values
<code>setenv</code>	Sets the specified NVRAM parameter to some value
<code>reset-all</code>	Resets the entire system; similar to a power cycle
<code>set-defaults</code>	Resets all parameter values to the factory defaults
<code>sifting text</code>	Displays the FORTH commands containing <code>text</code>
<code>.registers</code>	Displays the contents of the registers
<code>probe-scsi</code>	Identifies the devices on the internal Small Computer System Interface (SCSI) bus
<code>probe-scsi-all</code>	Identifies the devices on all SCSI buses
<code>probe-ide</code>	Identifies devices on the internal integrated device electronics (IDE) bus
<code>probe-fcal-all</code>	Identifies devices on all Fibre Channel loops

**Table 9-1** Typical Commands Used at the `ok` Prompt (Continued)

<code>show-devs</code>	Displays the entire device tree
<code>devalias</code>	Identifies the current boot device alias for the system
<code>nvalias</code>	Creates a new device alias name
<code>nvunalias</code>	Removes a device alias name
<code>show-disks</code>	Displays and allows a selection of device paths for the disks to be used for <code>nvalias</code>
<code>sync</code>	Manually attempts to flush memory and synchronize file systems
<code>test</code>	Runs self-tests on specified devices

## Identifying the System Boot PROM Version

The `banner` command lists system information, for example:

- model name
- boot PROM version number (for example, `1.x`, `2.x`, `3.x`, `4.x`, or `5.x`)
- memory
- Ethernet address
- host ID

The following example shows output from the `banner` command.

```
ok banner
```

```
Sun Blade 1500 (Silver), Keyboard Present
Copyright 1998-2004 Sun Microsystems, Inc. All rights reserved.
OpenBoot 4.16.4, 1024 MB memory installed, Serial #64512991.
Ethernet address 0:3:ba:d8:63:df, Host ID: 83d863df.
```

## Booting the System

Use the boot command to boot the Solaris OS from the ok prompt. This command has several options available for booting the system in different situations. The format for the boot command is:

```
boot device_name -options
```

Enter the boot command at the ok prompt to boot the system to multiuser mode automatically.

```
ok boot
```

The following list describes some of the options for the boot command:

- **-s** – Boots the system to a single-user mode and asks the user for the root password.  
ok **boot -s**
- **cdrom -s** – Boots the system to single user mode from a CD-ROM or a DVD.  
ok **boot cdrom -s**
- **-a** – Boots the system interactively. Use this option if an alternative file must be executed during boot. The boot program asks for the following information. Enter the Return key to accept the values presented.

```
ok boot -a
```

(output omitted)

Rebooting with command: boot -a

Boot device: /pci@1e,600000/ide@d/disk@0,0 File and args: -a

Enter filename [kernel/sparcv9/unix]: <**Return**>

Enter default directory for modules [/platform/SUNW,Sun-Blade-1500/kernel /platform/sun4u/kernel /kernel /usr/kernel]:<**Return**>

Name of system file [etc/system]:<**Return**>

SunOS Release 5.10 Version Generic\_118833-33 64-bit

Copyright 1983-2006 Sun Microsystems, Inc. All rights reserved.

Use is subject to license terms.

root filesystem type [ufs]:<**Return**>

Enter physical name of root device

[/pci@1e,600000/ide@d/disk@0,0:a]:<**Return**>

(output omitted)

- **-r** – Performs a reconfiguration boot. Use this option to find a newly attached device and to create new device entries in the /devices and /dev directories. It also updates the /etc/path\_to\_inst file.

ok **boot -r**

- **-v** – Boots the system while displaying more detailed device information to the console. Use this option to troubleshoot problems during the boot process. You can use this option with other options.

ok **boot -v**

ok **boot -rv**

ok **boot -sv**

## Accessing More Detailed Information

Use the **help** command to obtain help on the main categories in the OpenBoot firmware. The following is an example of the help output from a Sun Blade 1500 system that is running OpenBoot PROM version 4.16.4:

```
ok help
```

Enter 'help command-name' or 'help category-name' for more help  
(Use ONLY the first word of a category description)

Examples: help select -or- help line

Main categories are:

Breakpoints (debugging)

Repeated loops

Defining new commands

Numeric output

Radix (number base conversions)

Arithmetic

Memory access

Line editor

System and boot configuration parameters

Select I/O devices

eject devices

Power on reset

Diag (diagnostic routines)

Resume execution

File download and boot

nvramrc (making new commands permanent)

ok

The help command listing provides keywords you can use to view further details. For example, to view information for one of the main categories listed in the preceding example, perform one of the following commands:

```
ok help boot
```

```
ok help nvramrc
```

```
ok help diag
```

```
ok help line
```

## Listing NVRAM Parameters

Use the `printenv` command to list all the NVRAM parameters. If the parameter can be modified, the `printenv` command displays its default setting and current setting. The following example shows output from the `printenv` command.

Variable Name	Value	Default Value
asr-policy	normal	normal
test-args		
diag-passes	1	1
local-mac-address?	true	true
fcode-debug?	false	false
scsi-initiator-id	7	7
oem-logo		No default
oem-logo?	false	false
oem-banner		No default
oem-banner?	false	false
ansi-terminal?	true	true
screen-#columns	80	80
screen-#rows	34	34
ttyb-rts-dtr-off	false	false
ttyb-ignore-cd	true	true
ttya-rts-dtr-off	false	false
ttya-ignore-cd	true	true
ttyb-mode	9600,8,n,1,-	9600,8,n,1,-
ttya-mode	9600,8,n,1,-	9600,8,n,1,-
output-device	ttya	screen
input-device	ttya	keyboard
auto-boot-on-error?	false	false
error-reset-recovery	sync	sync
load-base	16384	16384
auto-boot?	true	true
boot-command	boot	boot
diag-file		
diag-device	net	net
boot-file		
boot-device	disk net	disk net
use-nvramrc?	false	false
nvramrc		
security-mode	none	No default
security-password		No default
security-#badlogins	0	No default
verbosity	min	min

diag-trigger	none	none
service-mode?	false	false
diag-script	normal	normal
diag-level	max	max
diag-switch?	false	false
ok		

Use the **printenv** command to display a single parameter and its values. For example, to display only the **boot-device** parameter, perform the following command:

```
ok printenv boot-device
boot-device = disk net
```

The possible values of the **boot-device** parameter include **disk**, **net**, and **cdrom**.

---

**Note** – Some OpenBoot PROM parameters, such as **auto-boot?**, end in a question mark. The question mark is part of the parameter name. If an OpenBoot PROM parameter ends in a question mark, the parameter value is typically either **true** or **false**.



## Changing NVRAM Parameters

Use the **setenv** command to change the current values assigned to NVRAM parameters.

If the **auto-boot?** parameter is set to **true**, the system boots automatically. If it is set to **false**, the system stops at the **ok** prompt.

The following example changes the **auto-boot?** parameter from its default setting of **true** to the value of **false**.

```
ok printenv auto-boot?
auto-boot? = true
ok
ok setenv auto-boot? false
auto-boot? = false
```

The **reset-all** command halts the system, clears all buffers and registers, and performs a software simulated power-off/power-on of the system.

```
ok reset-all
Resetting ...
```

**Note** – The **reset-all** command, combined with the **auto-boot? = false** setting clears system registers, which is required on a system with a PROM 3.x or higher, before you can use the **probe** command or perform other tests.



## Restoring Default NVRAM Parameters

Use the **set-defaults** command to reset all NVRAM parameters to their default values. It affects only parameters that have assigned default values.

```
ok set-defaults
Setting NVRAM parameters to default values.
```

To reset a specific parameter to its default value, use the **set-default** command followed by the parameter name.

```
ok set-default parameter-name
```

For example, to reset the **diag-level** parameter, perform the command:

```
ok set-default diag-level
```

## Displaying Devices Connected to the System

Use the `probe` command to identify peripheral devices, such as disks, tape devices, or CD-ROMs, currently connected to a system.

Use the `sifting` command to identify available probe commands (see the following example) and to find OpenBoot PROM commands when you don't know the command syntax. Commonly-used probe commands include `probe-scsi`, `probe-scsi-all`, and `probe-ide`.

**ok sifting probe**

In vocabulary	forth		
(f0060930) probe-all	(f005fe98) probe-ide	(f005fc88) probe-scsi-all	
(f005023c) probe-all	(f00501cc) probe-io	(f004a7c4) probe-root-	
slot			
(f00342e0) probe	(f00341a0) probe-virtual	(f002fc10) xprobe	
(f002fbe8) lprobe	(f002fbcc0) wprobe	(f002fb98) cprobe	
(output omitted)			



**Caution** – The following warning message might be displayed if you invoke the `probe` commands on Sun systems that contain a boot PROM that is version 3.x and above.

This command may hang the system if a Stop-A or halt command has been executed. Please type `reset-all` to reset the system before executing this command.

Do you wish to continue? (y/n) **n**

## Using Basic Boot PROM Commands

If any portion of the Solaris OS was loaded into memory when the system was aborted or shut down, the probe commands can hang the system. Look at OBP registers to determine if a **reset-all** command is necessary. Non-zero values in the registers listed indicate the possibility that the system may hang if you use a probe command. For example:

```
ok .registers
 Normal Alternate MMU Vector
0: 0 0 0 0
1: 1047cf0 ffff686c0 f00447d4 798
2: 8 f00000000 0 4
3: 8 3 0 180c660
4: 30005c8ac00 3 3ff 180c000
5: 9 fe000000000f fff6c000 10
6: 10 100769c 800000003ff640b6 180c9f0
7: 300022f09a0 4058 2 20
%PC f00475c8 %nPC f00475cc
%TBA f0000000 %CCR 0 XCC:nzvc ICC:nzvc
ok
```

To avoid hanging your system, perform the commands:

```
ok setenv auto-boot? false
ok reset-all
```

Use the **.registers** command again, and verify that all values are set to zero. For example:

```
ok .registers
 Normal Alternate MMU Vector
0: 0 0 0 0
1: 0 0 0 0
2: 0 0 0 0
3: 0 0 0 0
4: 0 0 0 0
5: 0 0 0 0
6: 0 0 0 0
7: 0 0 0 0
%PC 0 %nPC 0
%TBA 0 %CCR 0 XCC:nzvc ICC:nzvc
okc
```

With registers in this state, probe commands will function properly.

## The **probe-scsi** Command

The **probe-scsi** command identifies peripheral devices attached to the primary on-board SCSI controller. The **probe-scsi** command identifies such peripheral devices as disks, tape drives, or CD-ROMs by their target addresses. This example shows the two internal SCSI disks of a Sun Fire V120 system:

```
ok probe-scsi
Target 0
 Unit 0 Disk FUJITSU MAP3367N SUN36G 0401
Target 1
 Unit 0 Disk FUJITSU MAP3367N SUN36G 0401
ok
```

Not all systems support the **probe-scsi** command. For example, the **probe-scsi** command is absent from SunBlade 1500 systems:

```
ok probe-scsi
probe-scsi ?
ok
```

## The **probe-scsi-all** Command

The **probe-scsi-all** command identifies the peripheral devices that are attached to the on-board SCSI controller and all peripheral devices attached to additional SBUS or PCI SCSI controllers.

The following example from a Sun Fire V120 system lists three disks in an array attached to the external SCSI initiator identified by the `/pci@1f,0/pci@1/scsi@8,1` path, and the two internal SCSI disks attached to the internal SCSI initiator identified by the `/pci@1f,0/pci@1/scsi@8` path. Both of these SCSI initiators are found on the system board of the V120, but the **probe-scsi-all** command is required to show the initiator used to connect external devices.

```
ok probe-scsi-all
/pci@1f,0/pci@1/scsi@8,1
Target 0
 Unit 0 Disk SEAGATE ST373307LSUN72G 0507
Target 1
 Unit 0 Disk SEAGATE ST373307LSUN72G 0507
Target 2
 Unit 0 Disk SEAGATE ST373307LSUN72G 0507
/pci@1f,0/pci@1/scsi@8
```

## Using Basic Boot PROM Commands

---

```
Target 0
 Unit 0 Disk FUJITSU MAP3367N SUN36G 0401
Target 1
 Unit 0 Disk FUJITSU MAP3367N SUN36G 0401

ok
```

The Target and Unit values identify the *target* and *disk* (logical unit number, or LUN) values used in the physical device paths for these disks. Each disk attached to the same SCSI initiator will use a unique target value. The disk or LUN value used in the physical device paths for these disks is 0. This is true for disks that use embedded target controllers.

### The **probe-ide** Command

The **probe-ide** command identifies disks and CD-ROMs that are attached to the on-board IDE controller. This command displays the device numbers of internal devices. This output from a SunBlade 1500 system lists two internal disks and one DVD-ROM drive.

```
ok probe-ide
 Device 0 (Primary Master)
 ATA Model: ST3120026A

 Device 1 (Primary Slave)
 ATA Model: ST3120026A

 Device 2 (Secondary Master)
 Removable ATAPI Model: TOSHIBA ODD-DVD SD-
R1512

 Device 3 (Secondary Slave)
 Not Present

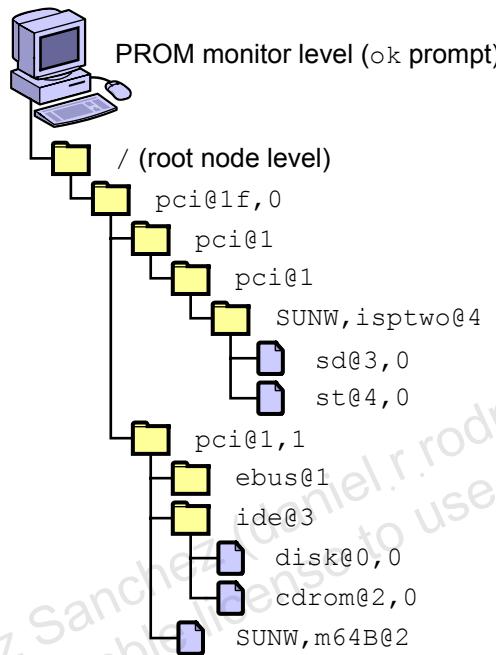
ok
```

On the SunBlade 1500, the Device values of the disks, in this case, 0 and 1, reflect the *target* values used in the physical device paths for these disks. The *disk* (logical unit number, or LUN) value used in the physical device paths for these disks is 0.

# Identifying the System Boot Device

Sun hardware uses the concept of a device tree to organize devices that are attached to the system.

Figure 9-5 shows the organizational structure of a device tree for an Ultra 5 or an Ultra 10 workstation.



**Figure 9-5** Partial Device Tree for an Ultra 5 or Ultra 10 Workstation



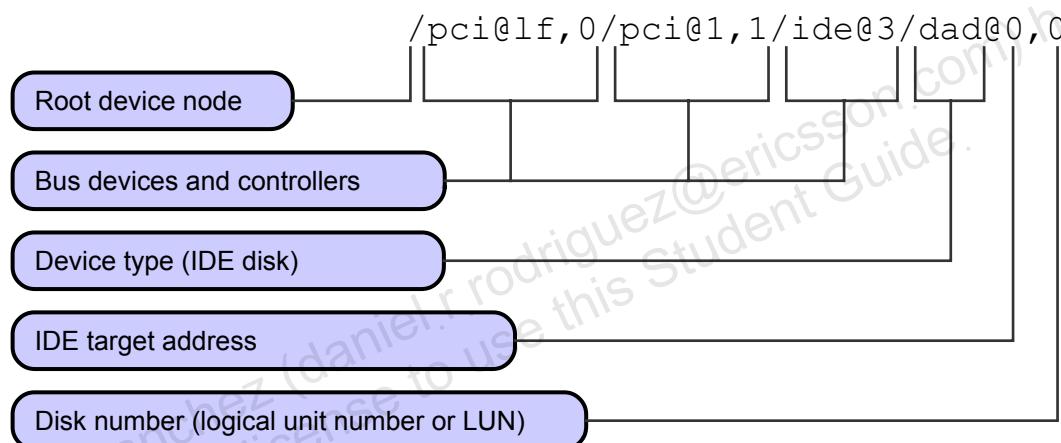
**Note** – In Figure 9-5, some license was taken in naming these directories to simplify the illustration.

The OpenBoot firmware builds the device tree from information gathered at the POST. This device tree is loaded into memory and is used by the kernel during the boot process to identify all configured devices. The top of the device tree is the root device node. Following the root device node is a bus nexus node. Connected to a bus nexus node is a leaf node, typically a controller for an attached device.

In Figure 9-5, the `disk@0,0` device is the IDE device for the hard disk drive, and the `cdrom@2,0` device is the IDE device for the CD-ROM drive. Both are attached to the IDE controller `ide@3`. Similarly, the `sd@3,0` device is the SCSI disk device and the `st@4,0` device is the SCSI tape device. Both are attached to the PCI-based SCSI controller `SUNW,isptwo@4`.

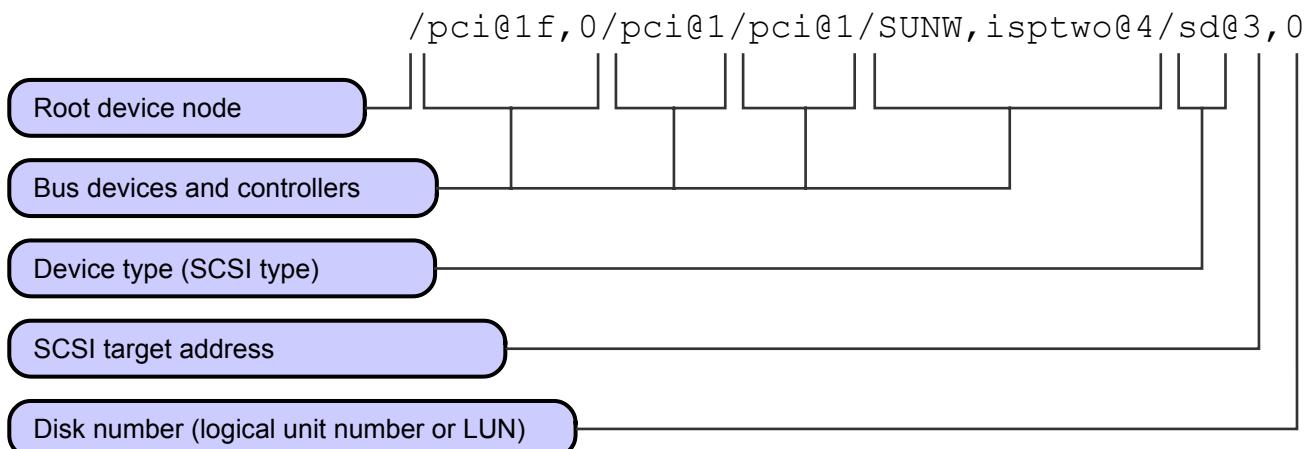
The paths built in the device tree by the OpenBoot firmware vary depending on the type of system and its device configuration.

Figure 9-6 shows a sample disk device path on an Ultra workstation with a PCI bus.



**Figure 9-6** Disk Device Path – Ultra Workstation With a PCI IDE Bus

Figure 9-7 shows a sample disk device path on an Ultra workstation with a PCI-SCSI bus.



**Figure 9-7** Disk Device Path – Ultra Workstation With a PCI-SCSI Bus

## The show-devs Command

Use the show-devs command to view the entire device tree. The following example shows partial output from the show-devs command on an Ultra 5 or Ultra 10 system:

```
ok show-devs
/SUNW,UltraSPARC-III@0,0
/pci@1f,0
/virtual-memory
/memory@0,10000000
/pci@1f,0/pci@1
/pci@1f,0/pci@1,1
/pci@1f,0/pci@1,1/ide@3
/pci@1f,0/pci@1,1/SUNW,m64B@2
/pci@1f,0/pci@1,1/network@1,1
/pci@1f,0/pci@1,1/ebus@1
/pci@1f,0/pci@1,1/ide@3/cdrom
/pci@1f,0/pci@1,1/ide@3/disk
/pci@1f,0/pci@1,1/ebus@1/SUNW,CS4231@14,200000
/pci@1f,0/pci@1,1/ebus@1/flashprom@10,0
/pci@1f,0/pci@1,1/ebus@1/eeprom@14,0
/pci@1f,0/pci@1/pci@1
/pci@1f,0/pci@1/pci@1/SUNW,isptwo@4
(output omitted)
ok
```

The following example shows partial output from the show-devs command on a SunBlade 1500 system:

```
ok show-devs
/i2c@1f,464000
/pci@1f,700000
/ppm@1e,0
/pci@1e,600000
/memory-controller@0,0
/SUNW,UltraSPARC-III@0,0
/virtual-memory
/memory@m0,0
/aliases
/options
/openprom
/chosen
/packages
/i2c@1f,464000/idprom@0,ae
/i2c@1f,464000/nvram@0,ae
```

## Identifying the System Boot Device

---

```
/pci@1f,700000/SUNW,XVR-100@3
/pci@1f,700000/network@2
/pci@1e,600000/pci@3
/pci@1e,600000/pci@2
/pci@1e,600000/ide@d
/pci@1e,600000/usb@b
/pci@1e,600000/usb@a
/pci@1e,600000/sound@8
/pci@1e,600000/pmu@6
/pci@1e,600000/isa@7
/pci@1e,600000/pci@3/SUNW,isptwo@4
/pci@1e,600000/pci@3/SUNW,hme@0,1
/pci@1e,600000/pci@3/pci108e,1000@0
/pci@1e,600000/pci@3/SUNW,isptwo@4/st
/pci@1e,600000/pci@3/SUNW,isptwo@4/sd
(output omitted)
ok
```

---

**Note** – In addition to the show-devs command, use the following additional OpenBoot PROM commands to view specific device information: show-ttys, show-displays, show-nets, show-disks, and show-tapes.

---



## The devalias Command

Use the devalias command to display the list of defined device aliases on your system. Device aliases provide short names used to identify longer physical device paths. This example shows output from the devalias command on an Ultra 5 or Ultra 10 system:

```
ok devalias
screen /pci@1f,0/pci@1,1/SUNW,m64B@2
net /pci@1f,0/pci@1,1/network@1,1
cdrom /pci@1f,0/pci@1,1/ide@3/cdrom@2,0:f
disk /pci@1f,0/pci@1,1/ide@3/disk@0,0
disk3 /pci@1f,0/pci@1,1/ide@3/disk@3,0
disk2 /pci@1f,0/pci@1,1/ide@3/disk@2,0
disk1 /pci@1f,0/pci@1,1/ide@3/disk@1,0
disk0 /pci@1f,0/pci@1,1/ide@3/disk@0,0
ide /pci@1f,0/pci@1,1/ide@3
floppy /pci@1f,0/pci@1,1/ebus@1/fdthree
ttyb /pci@1f,0/pci@1,1/ebus@1/se:b
ttya /pci@1f,0/pci@1,1/ebus@1/se:a
keyboard!
/pci@1f,0/pci@1,1/ebus@1/su@14,3083f8:forcemode
keyboard
/pci@1f,0/pci@1,1/ebus@1/su@14,3083f8
mouse /pci@1f,0/pci@1,1/ebus@1/su@14,3062f8
name aliases
```

This example shows output from the devalias command on a SunBlade 1500 system:

```
ok devalias
screen /pci@1f,700000/SUNW,XVR-100@3
mouse /pci@1e,600000/usb@b/mouse@2
keyboard /pci@1e,600000/usb@b/keyboard@1
net /pci@1f,700000/network@2
disk /pci@1e,600000/ide@d/disk@0,0
cdrom /pci@1e,600000/ide@d/cdrom@2,0:f
ide /pci@1e,600000/ide@d
ttyb /pci@1e,600000/isa@7/serial@0,2e8
ttya /pci@1e,600000/isa@7/serial@0,3f8
name aliases
ok
```

The left side of the devalias output lists the device alias names, and the right side of the output lists the physical address of each device.

Predefined device aliases are built into the OpenBoot PROM firmware, and they are easier to remember and use than the physical device addresses. The disk device alias typically identifies the default boot device for the system. Some systems list more predefined device aliases than others.

## The boot-device Parameter

The value of the boot-device parameter specifies the device from which the system will boot by default. The default value of the boot-device parameter is disk net. Both disk and net are device aliases.

Use the printenv command to view the current setting of the boot-device parameter. For example:

```
ok printenv boot-device
boot-device = disk net
ok
```

To determine what devices the disk and net device aliases represent, use the devalias command. For example:

```
ok devalias disk
disk /pci@1e,600000/ide@d/disk@0,0
ok devalias net
net /pci@1f,700000/network@2
ok
```

To boot the system from the default device, use the boot command:

```
ok boot
```

With the boot-device parameter set to these two values, the system attempts to boot from the disk device first, and attempts to boot from the net device only if it cannot boot from disk.

Knowing what physical device it is that the disk device alias represents is useful when you use the show-disks command to list disk devices that are attached to your system.

# Creating and Removing Custom Device Aliases

A portion of the NVRAM called NVRAMRC contains registers to hold custom parameters and is also reserved for storing new device alias names. External devices do not, by default, have built-in device aliases associated with them.

The NVRAMRC is affected by the commands `nvalias` and `nvunalias`, and the `use-nvramrc?` parameter.

## The `nvalias` Command

You use the `nvalias` command to create a new device alias name to access a newly attached external device. The command format is:

```
nvalias aliasname device_path
```

The effect of the `nvalias` command is to store the following command line in the NVRAMRC:

```
devalias aliasname device_path
```

The following example procedure shows how to use the `nvunalias` command to create a user-defined device alias, and how to boot the system using the new device alias.

**Note** – A shortcut that the `show-disks` command provides enables you to select a device path and use the Control-Y keys to copy the device path onto the command line.



## Creating and Removing Custom Device Aliases

The following example uses the **devalias** and **show-disks** commands to select the device path for the disk. It then uses the **nvalias** command to create a new device alias called **mydisk**.

1. Verify that the **boot-device** parameter is set to **disk net**.

```
ok printenv boot-device
boot-device = disk net
ok
```

2. Use the **devalias** command to display the device path that the **disk** device alias represents:

```
ok devalias disk
disk /pci@1e,600000/ide@d/disk@0,0
ok
```

The **disk@0,0** device identifies target 0, and disk or LUN 0, attached to the **/pci@1e,600000/ide@d** controller.

3. Use the **show-disks** command to display the list of disk types available on your system.

```
ok show-disks
a) /pci@1e,600000/pci@3/SUNW,isptwo@4/sd
b) /pci@1e,600000/ide@d/cdrom
c) /pci@1e,600000/ide@d/disk
q) NO SELECTION
Enter Selection, q to quit:
```

4. Enter the letter that selects the disk type that matches the current boot disk.

Enter Selection, q to quit: **c**

**/pci@1e,600000/ide@d/disk** has been selected.

Type **^Y** ( Control-Y ) to insert it in the command line.

e.g. ok **nvalias mydev ^Y**

for creating devalias **mydev** for **/pci@1e,600000/ide@d/disk**

ok

5. Use the **nvalias** command to define the new **mydisk** alias, and enter Control-Y to paste the device path you selected from the **show-disks** command into the command line.

```
ok nvalias mydisk ^y
```

The device path inserts into the command line:

```
ok nvalias mydisk /pci@1e,600000/ide@d/disk
```

6. Complete the disk device path by adding the target and disk (LUN) numbers so the device path matches the device that the disk device alias represents, and press Return.

```
ok nvalias mydisk /pci@1e,600000/ide@d/disk@0,0
ok
```

7. Use the devalias command to verify that the new device alias is defined.

```
ok devalias mydisk
mydisk /pci@1e,600000/ide@d/disk@0,0
ok
```

8. Use the printenv command to display the contents of the NVRAMRC registers. The devalias command used at boot time to establish the mydisk device alias should display.

```
ok printenv nvramrc
nvramrc = devalias mydisk /pci@1e,600000/ide@d/disk@0,0
ok
```

9. Set the boot-device parameter to the new value, in this case mydisk, and boot the system.

```
ok setenv boot-device mydisk
boot-device = mydisk
ok boot
```

When you create a device alias and use the partial path that the show-disks command provides, you must enter the correct target and disk (LUN) values to complete the device path.

A device path that ends with the disk (LUN) value implies that slice 0 on that disk is the boot slice.

To specify that a slice other than slice 0 is the boot slice, you must enter a letter that identifies the correct slice. For this purpose, slices 0 through 7 are represented by the letters a through h.

For example, the device path that ends with disk@0,0 is interpreted to mean disk@0,0:a, where :a specifies slice 0.

If slice 3 was the boot slice on this disk, you would enter disk@0,0:d to specify the correct device path. The complete nvalias command would look like this:

```
ok nvalias mydisk /pci@1e,600000/ide@d/disk@0,0:d
```

## The nvunalias Command

Use the nvunalias command to remove an alias name.

To remove a custom device alias name, use the following command format:

```
ok nvunalias aliasname
```



---

**Note** – The nvunalias command is the single exception to the rule that changes to NVRAM occur immediately and do not require a reset-all command.

---

In this example procedure shows how to use the nvunalias command to remove a user-defined device alias.

1. Use the devalias command to verify that the mydisk device alias is defined.

```
ok devalias mydisk
mydisk /pci@1e,600000/ide@d/disk@0,0
ok
```

2. Use the printenv command to verify that the NVRAMRC registers contain command used at boot time to establish the mydisk device alias.

```
ok printenv nramrc
nramrc = devalias mydisk /pci@1e,600000/ide@d/disk@0,0
ok
```

3. Use the nvunalias command to remove the devalias command for mydisk from the NVRAMRC registers.

```
ok nvunalias mydisk
ok
```

4. Use the printenv command to verify that the NVRAMRC registers have been cleared.

```
ok printenv nramrc
nramrc =
ok
```

5. Use the **devalias** command to verify that the **mydisk** device alias is still defined. The **nvunalias** command removes the **devalias** command for **mydisk** from the NVRAMRC registers, but does not remove the **mydisk** alias that was defined when the system booted.

```
ok devalias mydisk
```

```
mydisk
```

```
/pci@1e,600000/ide@d/disk@0,0
```

```
ok
```

6. Use the **set-default** command to set the **boot-device** parameter back to its default value. Verify the change.

```
ok set-default boot-device
```

```
ok printenv boot-device
```

```
boot-device = disk net
```

```
ok
```

7. Run the **reset-all** command, and then verify that the **mydisk** alias is no longer defined. This procedure assumes that the **auto-boot?** parameter is set to **false**.

```
ok reset-all
```

Sun Blade 1500 (Silver), Keyboard Present

Copyright 1998-2004 Sun Microsystems, Inc. All rights reserved.

OpenBoot 4.16.4, 1024 MB memory installed, Serial #64512991.

Ethernet address 0:3:ba:d8:63:df, Host ID: 83d863df.

```
ok devalias mydisk
```

```
mydisk : no such alias
```

```
ok
```

# Viewing and Changing NVRAM Parameters From the OS

Use the `/usr/sbin/eeprom` command to view and to change the NVRAM parameters while the Solaris OS is running.

## Using the `eeprom` Command

Be aware of the following guidelines when using the `eeprom` command:

- Only the root user can change the value of a parameter.
- You must enclose parameters with a trailing question mark in single quotation marks (single quotes) when the command is executed in the C shell.

The following examples use the `eeprom` command to view and change NVRAM parameters.

- To list all of the parameters with their current values, perform the command:

```
eeprom
asr-policy=normal
test-args: data not available.
diag-passes=1
local-mac-address?=true
fcode-debug?=false
scsi-initiator-id=7
oem-logo: data not available.
oem-logo?=false
(output omitted)
```

- To list a single parameter and its value, in this case, the `boot-device` parameter, perform the command:

```
eeprom boot-device
boot-device=disk net
#
```

- To change the value of the default boot device to `disk2`, perform the command:

```
eeprom boot-device=disk2
eeprom boot-device
boot-device=disk2
#
```



**Note** – The predefined device alias `disk2` is present on Ultra 5 and Ultra 10 systems, and not present on SunBlade 1500 systems.

---

- To change the value of the `auto-boot?` parameter, perform the command:

```
eeprom auto-boot?=true
eeprom auto-boot?
auto-boot?=true
#
```

## Interrupting an Unresponsive System

When a system freezes or stops responding to the keyboard, you might have to interrupt it. When you interrupt the system, all active processes stop immediately, and the processor services the OpenBoot PROM exclusively. It does not allow you to flush memory or to synchronize file systems.

## Aborting an Unresponsive System

To abort or interrupt an unresponsive system:

1. Attempt a remote login on the unresponsive system to locate and kill the offending process.
2. Attempt to reboot the unresponsive system gracefully.
3. Hold down the Stop-A key sequence on the keyboard of the unresponsive system. The system is placed at the **ok** prompt.

---

**Note** – If an ASCII terminal is being used as the system console, use the Break sequence keys.

---

4. Manually synchronize the file systems by using the OpenBoot PROM sync command.

**ok sync**

This command causes the system to panic, synchronize the file systems, perform a crash dump of memory, and then reboot the system.



# Exercise: Using Boot PROM Commands

In this exercise, you use the OpenBoot™ PROM and Solaris OS commands to set and view OBP parameters, create a new device alias, and boot the system using the new alias.

## Preparation

Refer to the lecture notes, as necessary, to complete the following steps and answer the questions listed.

This exercise lists examples from two systems, a V120 server and a Sun Blade 1500. These provide examples of systems that use SCSI and IDE internal disks.

## Task

Complete the following steps:

1. If the Solaris OS is currently running, log in as the `root` user and halt your system.
2. When the `ok` prompt displays, use the `help` command to display the list of help topics.
3. Use the `help` command to display information about the `boot` command.

What does the `help` command list for `boot`?

---

---

4. Use the `banner` command to obtain the following information:

OpenBoot PROM revision: \_\_\_\_\_

Megabytes of installed memory: \_\_\_\_\_

System type: \_\_\_\_\_

NVRAM serial number: \_\_\_\_\_

Ethernet address: \_\_\_\_\_

Host ID: \_\_\_\_\_

## Exercise: Using Boot PROM Commands

---

5. Use the printenv command to display the list of OpenBoot PROM parameters and record the current values for the following parameters:

output-device  
input-device  
auto-boot?  
boot-device

6. Prevent the system from booting automatically after using the reset-all command by setting the auto-boot? parameter to false.
7. Use the reset-all command to verify that the new auto-boot? value is in effect.

The system should remain at the ok prompt after the reset-all command completes.

8. Use the probe-scsi, probe-scsi-all and probe-ide commands to display the list of devices attached to your system's SCSI and IDE busses.

Not all of these commands are present on all systems.

What are the main differences that you see in the information that these commands display?

---

---

---

9. Record the target number and device type (disk, tape, or CD-ROM) of all the devices shown by the probe-scsi, probe-scsi-all and probe-ide commands.

10. Verify that your default boot-device is set to disk net.

11. Use the devalias command to display the full device path for the disk alias and record the path name reported:

---

---

12. Use the show-disks command to select the device path that relates to the disk recorded in Step 11 and complete the following steps:

- a. Use the nvalias command to create a new device alias called mydisk.
- b. Set the mydisk alias to the path and disk name you recorded in Step 11.

Remember to use the Control-Y key sequence to paste the disk path into your nvalias command. You must manually complete the path to specify the disk you want to use.

13. Verify that the new alias is correctly set.
14. Use the printenv command to display the contents of the nvramrc file.

What command does the nvramrc file contain that creates the mydisk alias?

---

---

15. Use the printenv command to display the setting of the use-nvramrc? parameter.

What is the current setting of the use-nvramrc? parameter?

---

---

16. Boot your system using the mydisk alias.
17. Log in as the root user on your system and open a new terminal window.
18. Use the eeprom command to list all NVRAM parameters.
19. Use the eeprom command to list the setting of the boot-device parameter.
20. Use the eeprom command to set the boot-device parameter to the alias mydisk.
21. Bring your system to run level 0.
22. Verify that the eeprom command set the boot-device parameter to the alias mydisk.
23. Set the boot-device parameter to its default value and verify the setting.
24. Use the nvunalias command to remove the alias mydisk.
25. Verify that the mydisk alias is no longer in the nvramrc file.
26. Use the devalias command to see if the mydisk alias has been removed from the list of device aliases.

Has it?

---

## Exercise: Using Boot PROM Commands

---

27. Run the `reset-all` command, and then check again if the `mydisk` alias has been removed from the list of device aliases.

(If your system reboots, interrupt the reboot with a Stop-A key sequence.)

Has it?

---

28. Set the OpenBoot PROM parameters back to their default values and boot the system from the default device.

29. Log in as the `root` user.

## Exercise Summary



**Discussion** – Take a few minutes to discuss what experiences, issues, or discoveries you had during the lab exercises.

- Experiences
- Interpretations
- Conclusions
- Applications

# Exercise Solutions: Using Boot PROM Commands

This section contains solutions to the exercise.

## Task

Complete the following steps:

1. If the Solaris OS is currently running, log in as the root user and halt your system.  
# **init 0**
2. When the ok prompt displays, use the help command to display the list of help topics.

V120 example:

```
ok help
Enter 'help command-name' or 'help category-name' for more help
(Use ONLY the first word of a category description)
Examples: help select -or- help line
```

Main categories are:

```
Breakpoints (debugging)
Repeated loops
Defining new commands
Numeric output
Radix (number base conversions)
Arithmetic
Memory access
Line editor
System and boot configuration parameters
Select I/O devices
Floppy eject
Power on reset
Diag (diagnostic routines)
Resume execution
File download and boot
nvramrc (making new commands permanent)
ok
```

Sun Blade 1500 example:

```
ok help
Enter 'help command-name' or 'help category-name' for more help
(Use ONLY the first word of a category description)
Examples: help select -or- help line
```

Main categories are:

Breakpoints (debugging)  
 Repeated loops  
 Defining new commands  
 Numeric output  
 Radix (number base conversions)  
 Arithmetic  
 Memory access  
 Line editor  
 System and boot configuration parameters  
 Select I/O devices  
 eject devices  
 Power on reset  
 Diag (diagnostic routines)  
 Resume execution  
 File download and boot  
 nvramrc (making new commands permanent)  
 ok

3. Use the help command to display information about the boot command.

Sun Blade 1500 example:

```

ok help boot
boot <specifier> (--) boot kernel (default) or other file
Examples:
 boot - boot kernel from default device.
 Factory default is to boot
 from DISK if present, otherwise from NET.
 boot net - boot kernel from network
 boot cdrom - boot kernel from CD-ROM
 boot disk1:h - boot from disk1 partition h
 boot tape - boot default file from tape
 boot disk myunix -as - boot myunix from disk with flags "-as"
dload <filename> (addr --) debug load of file over network at
address
Examples:
 4000 dload /export/root/foo/test
 ?go - if executable program, execute it
 or if Forth program, compile it
ok

```

What does the help command list for boot?

*The help command lists examples of boot command syntax used to boot the system from various devices.*

4. Use the banner command to obtain the following information:

OpenBoot PROM revision:

## Exercise Solutions: Using Boot PROM Commands

Megabytes of installed memory:

System type:

NVRAM serial number:

Ethernet address:

Host ID:

*Each system presents its own unique information.*

V120 example:

ok **banner**

Sun Fire V120 (UltraSPARC-IIe 648MHz), No Keyboard  
OpenBoot 4.0, 2048 MB memory installed, Serial #57165033.  
Ethernet address 0:3:ba:68:44:e9, Host ID: 836844e9.

ok

Sun Blade 1500 example:

ok **banner**

Sun Blade 1500 (Silver), Keyboard Present  
Copyright 1998-2004 Sun Microsystems, Inc. All rights reserved.  
OpenBoot 4.16.4, 1024 MB memory installed, Serial #64512991.  
Ethernet address 0:3:ba:d8:63:df, Host ID: 83d863df.

ok

5. Use the **printenv** command to display the list of OpenBoot PROM parameters and record the current values for the following parameters:

output-device - screen  
input-device - keyboard  
auto-boot? - true  
boot-device - disk net

ok **printenv**

Variable Name	Value	Default Value
ras-shutdown-enabled?	false	false
shutdown-temp	75	75
warning-temp	70	70
env-monitor	disabled	disabled
diag-passes	1	1
(output omitted)		

6. Prevent the system from booting automatically after using the reset-all command by setting the auto-boot? parameter to false.

```
ok setenv auto-boot? false
auto-boot? = false
ok
```

7. Use the reset-all command to verify that the new auto-boot? value is in effect.

The system should remain at the ok prompt after the reset-all command completes.

```
ok reset-all
Resetting
LOM event: +14d+1h39m5s host reset
g ...
?

Sun Fire V120 (UltraSPARC-IIe 648MHz), No Keyboard
OpenBoot 4.0, 2048 MB memory installed, Serial #57165033.
Ethernet address 0:3:ba:68:44:e9, Host ID: 836844e9.
```

ok

8. Use the probe-scsi, probe-scsi-all and probe-ide commands to display the list of devices attached to your system's SCSI and IDE busses. Not all of these commands are present on all systems.

V120 example:

```
ok probe-scsi
Target 0
 Unit 0 Disk FUJITSU MAP3367N SUN36G 0401
Target 1
 Unit 0 Disk FUJITSU MAP3367N SUN36G 0401
ok probe-scsi-all
/pci@1f,0/pci@1/scsi@8,1

/pci@1f,0/pci@1/scsi@8
Target 0
 Unit 0 Disk FUJITSU MAP3367N SUN36G 0401
Target 1
 Unit 0 Disk FUJITSU MAP3367N SUN36G 0401

ok probe-ide
Device 0 (Primary Master)
Removable ATAPI Model: DV-28E-C
```

## Exercise Solutions: Using Boot PROM Commands

Device 1 ( Primary Slave )  
Not Present

Device 2 ( Secondary Master )  
Not Present

Device 3 ( Secondary Slave )  
Not Present

ok

Sun Blade 1500 example:

ok **probe-scsi**

probe-scsi ?

ok **probe-scsi-all**

/pci@1e,600000/pci@3/SUNW,isptwo@4

Target 4

Unit 0 Removable Tape HP C1537A L706

ok **probe-ide**

Device 0 ( Primary Master )  
ATA Model: ST3120026A

Device 1 ( Primary Slave )  
ATA Model: ST3120026A

Device 2 ( Secondary Master )  
Removable ATAPI Model: TOSHIBA ODD-DVD SD-R1512

Device 3 ( Secondary Slave )  
Not Present

ok

What are the main differences that you see in the information that these commands display?

*The probe-scsi command only lists devices on the built-in SCSI chain and does not list the full device paths. The probe-scsi-all command lists all devices on all SCSI chains and their full device paths. The probe-ide command reports the list of IDE devices attached to the system.*

9. Record the target number and device type (disk, tape, or CD-ROM) of all the devices shown by the probe-scsi, probe-scsi-all and probe-ide commands.

*Each system presents its own unique information.*

- Verify that your default boot-device is set to disk net.

```
ok printenv boot-device
boot-device = disk net
ok
```

- Use the devalias command to display the full device path for the disk alias.

```
ok devalias disk
```

Record the path name reported:

*This differs from system to system. On an Ultra™ 5 workstation, the alias is defined as follows:*

```
/pci@1f,0/pci@1,1/ide@3/disk@0,0
```

*On a Sun Fire™ V120 server, the alias is defined as follows:*

```
/pci@1f,0/pci@1/scsi@8/disk@0,0
```

*On a Sun Blade™ 1500 system, the alias is defined as follows:*

```
/pci@1e,600000/ide@d/disk@0,0
```

- Use the show-disks command to select the device path that relates to the disk recorded in Step 11 and complete the following steps:
  - Use the nvalias command to create a new device alias called mydisk.
  - Set the mydisk alias to the path and disk name you recorded in Step 11.

Remember to use the Control-Y key sequence to paste the disk path into your nvalias command. You must manually complete the path to specify the disk you want to use.

V120 example:

```
ok show-disks
a) /pci@1f,0/pci@1/scsi@8,1/disk
b) /pci@1f,0/pci@1/scsi@8/disk
c) /pci@1f,0/pci@1,1/ide@d/cdrom
d) /pci@1f,0/pci@1,1/ide@d/disk
e) /pci@1f,0/pci@1,1/ebus@c/SUNW,lomh@14,200000
q) NO SELECTION
Enter Selection, q to quit: b
/pci@1f,0/pci@1/scsi@8/disk has been selected.
Type ^Y (Control-Y) to insert it in the command line.
e.g. ok nvalias mydev ^Y
for creating devalias mydev for
```

## Exercise Solutions: Using Boot PROM Commands

```
/pci@1f,0/pci@1/scsi@8/disk
ok nvalias mydisk /pci@1f,0/pci@1/scsi@8/disk@0,0
ok
```

Sun Blade 1500 example:

```
ok show-disks
a) /pci@1e,600000/pci@3/SUNW,isptwo@4/sd
b) /pci@1e,600000/ide@d/cdrom
c) /pci@1e,600000/ide@d/disk
q) NO SELECTION
Enter Selection, q to quit: c
/pci@1e,600000/ide@d/disk has been selected.
Type ^Y (Control-Y) to insert it in the command line.
e.g. ok nvalias mydev ^Y
 for creating devalias mydev for
/pci@1e,600000/ide@d/disk
ok nvalias mydisk /pci@1e,600000/ide@d/disk@0,0
ok
```

13. Verify that the new alias is correctly set.

V120 example:

```
ok devalias mydisk
mydisk
/pci@1f,0/pci@1/scsi@8/disk@0,0
ok
```

Sun Blade 1500 example:

```
ok devalias mydisk
mydisk /pci@1e,600000/ide@d/disk@0,0
ok
```

14. Use the printenv command to display the contents of the nvramrc file.

ok printenv nvramrc

What command does the nvramrc file contain that creates the mydisk alias?

*Systems differ according to the disk devices they use. An Ultra 5 workstation would report the following:*

```
nvramrc = devalias mydisk /pci@1f,0/pci@1,1/ide@3/disk@0,0
```

*A Sun Fire V120 server would report the following:*

```
nvramrc = devalias mydisk /pci@1f,0/pci@1/scsi@8/disk@0,0
```

*A SunBlade 1500 system would report the following:*

```
nvramrc = devalias mydisk /pci@1e,600000/ide@d/disk@0,0
```

15. Use the **printenv** command to display the setting of the **use-nvramrc?** parameter.

```
ok printenv use-nvramrc?
use-nvramrc? = true
ok
```

What is the current setting of the **use-nvramrc?** parameter?

true

16. Boot your system using the **mydisk** alias.

```
ok boot mydisk
(output omitted)
```

17. Log in as the **root** user on your system and open a new terminal window.

18. Use the **eeprom** command to list all NVRAM parameters.

```
eeprom
ras-shutdown-enabled?=false
shutdown-temp=75
warning-temp=70
env-monitor=disabled
diag-passes=1
diag-continue?=0
diag-targets=0
(output omitted)
```

19. Use the **eeprom** command to list the setting of the **boot-device** parameter.

```
eeprom boot-device
boot-device=disk net
#
```

20. Use the **eeprom** command to set the **boot-device** parameter to the alias **mydisk**.

```
eeprom boot-device=mydisk
```

21. Bring your system to run level 0.

```
init 0
```

22. Verify that the **eeprom** command set the **boot-device** parameter to the alias **mydisk**.

```
ok printenv boot-device
boot-device = mydisk
ok
```

23. Set the **boot-device** parameter to its default value and verify the setting.

```
ok set-default boot-device
ok printenv boot-device
boot-device = disk net
ok
```

24. Use the nvunalias command to remove the alias mydisk.

```
ok nvunalias mydisk
```

25. Verify that the mydisk alias is no longer in the nvramrc file.

```
ok printenv nvramrc
nvramrc =
ok
```

26. Use the devalias command to see if the mydisk alias has been removed from the list of device aliases.

```
ok devalias mydisk
mydisk
/pci@1f,0/pci@1/scsi@8/disk@0,0
```

ok

Has it?

No

27. Run the reset-all command, and then check again if the mydisk alias has been removed from the list of device aliases.

```
ok reset-all
```

(If your system reboots, interrupt the reboot with a Stop-A key sequence.)

```
ok devalias mydisk
mydisk : no such alias
ok
```

Has it?

Yes

28. Set the OpenBoot PROM parameters back to their default values and boot the system from the default device.

```
ok set-defaults
```

```
ok printenv
```

Variable Name	Value	Default Value
ras-shutdown-enabled?	false	false
shutdown-temp	75	75
warning-temp	70	70
env-monitor	disabled	disabled
diag-passes	1	1
(output omitted)		
ok <b>reset-all</b>		

29. Log in as the root user.

## Notes:

---

## Module 10

---

# Using the Grand Unified Bootloader

---

## Objectives

Upon completion of this module, you should be able to:

- Discuss the purpose of the Grand Unified Bootloader (GRUB)
- Describe GRUB terminology and basic functions
- Modify x86 system boot behavior in the Solaris OS
- Use the `eeprom` command
- Use the `kernel` command
- Use the `findroot` command
- Describe and manage GRUB boot archives
- Boot a system in the GRUB-based boot environment
- Interrupt an unresponsive system

# Introduction to GRUB

Starting with the Solaris 10 1/06 release, the open source GNU GRUB is available on x86-based systems that are running the Solaris OS. GRUB is the boot loader that is responsible for loading a boot archive into a system's memory. The boot archive contains the kernel modules and configuration files that are required to boot the system.

To support GRUB, x86-based systems must have at least 256 Mbytes of main memory.

## GRUB Terminology

The following GRUB terms are used throughout this module, and are helpful in understanding GRUB concepts.

- Boot archive

A boot archive is a collection of critical files that is used to boot the Solaris OS. These files are needed during system startup before the root (/) file system is mounted. Two boot archives are maintained on a system:

- The boot archive that is used to boot the Solaris OS on a system. This boot archive is sometimes called the primary boot archive.
- The boot archive that is used for recovery when the primary boot archive is damaged. This boot archive starts the system without mounting the root (/) file system. On the GRUB menu, this boot archive is called Solaris failsafe. The failsafe archive's essential purpose is to boot the system and regenerate the primary boot archive, which is usually used to boot the system.

- Boot loader

The boot loader is the first software program that runs after you turn on a system. This program begins the booting process.

- GRUB edit menu

A boot menu that is a submenu of the GRUB main menu. GRUB commands are displayed on this menu. These commands can be edited to change boot behavior.

- GRUB main menu  
A boot menu that lists the operating systems that are installed on a system. From this menu, you can easily boot an operating system without modifying the BIOS or fdisk partition settings.
- menu.lst file  
A file that lists all the operating systems that are installed on a system. The contents of this file dictate the list of operating systems that is displayed on the GRUB menu. From the GRUB menu, you can easily boot an operating system without modifying the BIOS or fdisk partition settings.
- Miniroot  
A minimal, bootable root (/) file system that resides on the Solaris installation media. A miniroot consists of the Solaris software that is required to install and upgrade systems. On x86-based systems, the miniroot is copied to the system to be used as the failsafe boot archive.

## GRUB-Based Booting

The boot loader is the first software program that runs after you turn on an x86-based system. This program begins the boot process. After an x86-based system is powered on, the Basic Input/Output System (BIOS) initializes the CPU, the memory, and the platform hardware. When the initialization phase has completed, the BIOS loads the boot loader from the configured boot device, and then gives control of the system to the boot loader.

GRUB implements a simple menu interface that includes boot options that are predefined in a configuration file. GRUB also has a command-line interface that is accessible from the menu interface for performing various boot commands. The GRUB implementation on x86-based systems in the Solaris OS is compliant with the Multiboot Specification, which is described in detail at <http://www.gnu.org/software/grub/grub.html>.

Because the Solaris kernel is fully compliant with the Multiboot Specification, you can boot x86-based systems using GRUB. With GRUB, you can more easily boot the various operating systems that are installed on your system.

For example, on one x86-based system, you could individually boot the following operating systems:

- Solaris OS
- Linux
- Microsoft Windows

When you boot an x86-based system, the GRUB menu displays. This menu provides a list of boot entries from which to choose. A boot entry is an OS instance that is installed on your system.

The GRUB main menu is based on a configuration file that you can modify.

- If you install or upgrade the Solaris OS, the GRUB menu is automatically updated. The Solaris OS is displayed as a new boot entry.
- If you install an operating system other than the Solaris OS, you must modify the /boot/grub/menu.1st configuration file to include the new OS instance. Adding the new OS instance enables the new boot entry to appear in the GRUB menu the next time that you boot the system.

## How GRUB-Based Booting Works

After GRUB gains control of the system, the GRUB menu displays on the console.

In the GRUB main menu, you can do the following:

- Select a boot entry.
- Modify a boot entry using the GRUB edit menu.
- Load an OS kernel from the command line manually.

The following is a sample of a typical GRUB main menu. This menu lists the boot entries to choose from. If you do not choose an entry, the system boots the default OS. The default OS is usually the first entry that is displayed in the GRUB main menu.

```
GNU GRUB version 0.95 (631K lower / 2095488K upper memory)
+-----+
| Solaris 10 10/09 s10x_u8wos_08a x86
| Solaris failsafe
| Diagnostic Partition
|
+-----+
```

Use the ^ and v keys to select which entry is highlighted.  
Press enter to boot the selected OS, 'e' to edit the commands before booting, or 'c' for a command-line.

Use the arrow keys to highlight a boot entry, then press Enter to boot the system using the selected entry.

If you want to edit a boot entry using the GRUB edit menu, use the arrow keys to highlight a boot entry in the GRUB main menu. Then, type e to display the GRUB edit menu. If you want to use the command-line interface, use the arrow keys to select a boot entry in the GRUB main menu, then type c.

When the system boots, GRUB loads the primary boot archive and the multiboot program. The primary boot archive, named /platform/i86pc/boot\_archive, is a ram image of the file system that contains the Solaris kernel modules and data.

GRUB puts this information in memory, without any interpretations. After the operation has completed, GRUB gives control of the system to the multiboot program. GRUB itself then becomes inactive, and system memory is reclaimed.

By reading the boot archive and passing boot-related information to the kernel, the multiboot program is responsible for assembling core kernel modules into memory. The Multiboot Specification specifies the boot-related information.

## GRUB Device Naming Conventions

GRUB uses device naming conventions that are slightly different from previous Solaris OS versions. Understanding the GRUB device naming conventions can assist you in correctly specifying drive and partition information when you configure GRUB on your system.

The following table describes the GRUB device naming conventions.

**Table 10-1** Naming Conventions for GRUB Devices

Device Name	Description
(fd0), (fd1)	First diskette, second diskette
(nd)	Network device
(hd0, 0), (hd0, 1)	First and second fdisk partition of first bios disk
(hd0, 0, a), (hd0, 0, b)	Solaris/BSD slice 0 and 1 (a and b) on the first fdisk partition on the first bios disk

---

**Note** – All GRUB device names must be enclosed in parentheses. Partition numbers are counted from 0 (zero), not from 1.

---



## GRUB Functional Components

GRUB consists of the following functional components:

- stage1 – Is installed on the first sector of the Solaris fdisk partition.
- stage2 – Is installed in a reserved area in the Solaris fdisk partition. stage2 is the core image of GRUB.
- menu.1st – Is a file that is typically located in the /boot/grub directory. The GRUB stage2 file reads this file.

You cannot use the dd command to write stage1 and stage2 to disk. stage1 must be able to receive information about the location of stage2 on the disk. Use the installgrub command, which is the supported method for installing GRUB boot blocks.



## The GRUB Menu

By default, the Solaris menu file resides in `/boot/grub/menu.1st`. The contents of this file dictate what is displayed in the GRUB menu when you boot the system. The GRUB menu file contains entries for all the OS instances that are installed on your system, as well as other important boot directives.

**Note** – In some situations, the GRUB menu.1st file resides elsewhere. For example, on a system that uses Solaris Live Upgrade, the GRUB menu.1st file might exist on a boot environment that is not the currently running boot environment. Or, if you have upgraded a system with an x86-based boot partition, the menu.1st file might reside in the `/stubboot` directory.

Only the active GRUB menu.1st file is used to boot the system. In order to modify the GRUB menu that is displayed when you boot the system, the active GRUB menu.1st file must be modified. Changing any other GRUB menu.1st file has no effect on the menu that is displayed when you boot the system.

To determine the location of the active GRUB menu.1st file, use the `bootadm` command. The `list-menu` subcommand displays the location of the active GRUB menu. For example:

```
bootadm list-menu
```

```
The location for the active GRUB menu is: /boot/grub/menu.1st
default 0
timeout 10
0 Solaris 10 11/06 s10x_u3wos_10 X86
1 Solaris failsafe
#
```

You might need to customize the menu.1st file for one of the following reasons:

- To add new OS entries
- To add GRUB console redirection information
- To change an aspect of system boot behavior that is not controlled using the eeprom or kernel commands

The instructions for the syntax to use to edit the file can be found in the /boot/grub/menu.1st file comments. Because these changes are not directly related to the Solaris OS, you cannot make them using the eeprom command.

When you install the Solaris OS, two GRUB menu entries are installed on the system by default. The first entry is the Solaris OS entry. The second entry is the miniroot (failsafe) archive, which is intended to be used for system recovery.

The GRUB menu entries are installed and then updated automatically as part of the Solaris software installation and upgrade process. The OS directly manages these entries and they should not be manually edited.

### Modifying menu.1st to Prevent Automatic System Boot

The timeout value in the menu.1st file determines if the system will boot automatically, and if so, how long the GRUB main menu displays before the boot process proceeds.

To prevent the system from booting automatically, and cause the GRUB main menu to remain displayed until you provide a response, edit the menu.1st file and set the timeout value to -1. For example:

```
timeout -1
```

A positive timeout value specifies the timeout in seconds before the system boots from the default source. The comments in the /boot/grub/menu.1st file describe this parameter.

Although you can use the eeprom command to set the auto-boot? parameter to false in the /boot/solaris/bootenv.rc file, this has no effect on the x86 system boot process.

# Modifying x86 System Boot Behavior in the Solaris OS

This section describes the methods you use to modify boot behavior in the Solaris OS on x86 systems. The primary ways you can modify boot behavior are:

- Use the `eeprom` command to assign a different value to a standard set of properties. These values, which are the equivalent to the SPARC OpenBoot PROM NVRAM variables, are stored in the `/boot/solaris/bootenv.rc` file. Changes that are made to the Solaris boot behavior using the `eeprom` command persist over each system reboot.
- Use the `kernel` command, which is available in the GRUB menu when the system boots. Boot entries listed in the GRUB main menu provide access to a list of commands used to boot the system, including the `kernel` command. Editing `kernel` command in the GRUB menu allows you change the system boot behavior for a single boot cycle.



**Note** – Changes that you make with the `kernel` command override options that you set using the `eeprom` command. These changes remain in effect until the next time you boot the system.

You can also modify the Solaris boot behavior in the GRUB menu.1st file. Modifying the Solaris boot behavior this way is similar to using the `eeprom` command, but isn't recommended because the Solaris OS entries in the menu.1st file could be modified during a software upgrade. In this instance, any changes that you made to the file are lost. In contrast, changes you make using the `eeprom` command are added to the `/boot/solaris/bootenv.rc` file. This file is handled properly during an upgrade.

## Setting Solaris Boot Parameters Using `eeprom`

Display or set Solaris boot parameters on x86 systems using the `eeprom` command. The parameters that you set with the `eeprom` command are stored in the `/boot/solaris/bootenv.rc` file. Changes that you make using the `eeprom` command persist across system reboots. However, you can override `eeprom` settings if you boot the system by specifying options and arguments to the `kernel` command in the GRUB menu.

Use the `eeprom` command to display the current values stored in the `bootenv.rc` file. For example:

```
kbd-type=US-English
ata-dma-enabled=1
atapi-cd-dma-enabled=0
ttyb-rts-dtr-off=false
ttyb-ignore-cd=true
ttya-rts-dtr-off=false
ttya-ignore-cd=true
ttyb-mode=9600,8,n,1,-
ttya-mode=9600,8,n,1,-
lba-access-ok=1
prealloc-chunk-size=0x2000
bootpath=/pci@0,0/pci-ide@5/ide@0/cmdk@0,0:a
console=text
bootpath=/pci@0,0/pci-ide@5/ide@0/cmdk@0,0:a
```

For more information about the `eeprom` command and parameters, see the `eeprom(1M)` man page.

---

**Note** – Not all systems support all `eeprom` parameters. Defaults vary depending on the system configuration. Use the following procedure to display and set boot parameters using the `eeprom` command.



1. Become superuser or assume an equivalent role.
2. To check if a parameter is listed in the `/boot/solaris/bootenv.rc` file and view its current value, use the `eeprom` command and specify the parameter name. For example, if the parameter is not listed in the `bootenv.rc` file, the `eeprom` command reports that data is not available for it:

```
eeprom boot-file
boot-file: data not available.
```

If the parameter is listed in the `bootenv.rc` file, but is set to a null value, the `eeprom` command reports the parameter name, and no value:

```
eeprom boot-file
boot-file=
```

3. To set or change a parameter, enter the **eeprom** command with the appropriate arguments. For example, to cause the system to boot using the 64-bit UNIX kernel found in the `/platform/i86pc/kernel/amd64` directory, you would use the following command:

```
eeprom boot-file=kernel/amd64/unix
```

4. Use the **eeprom** command to verify that the new parameter has been set.

```
eeprom boot-file
boot-file=kernel/amd64/unix
```

5. You can use the **grep** command to view the entry that the **eeprom** command placed in the `bootenv.rc` file.

```
grep boot-file /boot/solaris/bootenv.rc
setprop boot-file 'kernel/amd64/unix'
```

6. To remove a setting for a parameter, use double quotes to specify a null parameter value. For example, setting the `boot-file` parameter to a null value on an x86 system causes the system to use its default autodetect boot behavior:

```
eeprom boot-file=""
eeprom boot-file
boot-file=
```

## Identifying the Boot Device

Use the following **eeprom** command to identify the boot device the system uses:

```
eeprom bootpath
bootpath=/pci@0,0/pci-ide@5/ide@0/cmdk@0,0:a
#
```

## Modifying Boot Behavior Using the kernel Command

You can modify the Solaris boot behavior by editing kernel commands in the GRUB menu when the system first boots.

When you first power on an x86 system, the GRUB main menu displays. For example:

```
GNU GRUB version 0.95 (631K lower / 2095488K upper memory)
+-----+
| Solaris 10 10/09 s10x_u8wos_08a X86
| Solaris failsafe
| Diagnostic Partition
|
+-----+
```

Use the ^ and v keys to select which entry is highlighted.

Press enter to boot the selected OS, 'e' to edit the commands before booting, or 'c' for a command-line.

Pressing the space bar or any key other than e interrupts the boot process and causes the GRUB main menu to remain displayed. To display the command list you want to edit, highlight the desired boot entry in the GRUB main menu, and enter e.

Typing e when the GRUB main menu displays interrupts the boot process, and displays the command list for the default boot entry. For example:

```
GNU GRUB version 0.95 (631K lower / 2095488K upper memory)
+-----+
| kernel /platform/i86pc/multiboot
| module /platform/i86pc/boot_archive
|
+-----+
```

Each line in this menu is an individual command for the boot entry that you chose in the GRUB main menu. To modify the boot behavior for the boot entry you have selected, edit the kernel command for that boot entry.

To edit a kernel command, highlight the command in the command list, and enter e. The GRUB edit menu displays the kernel command you selected. For example:

```
[Minimal BASH-like line editing is supported. For the first word, TAB lists possible command completions. Anywhere else TAB lists the possible completions of a device/filename. ESC at any time exits.]
```

```
grub edit> kernel /platform/i86pc/multiboot
```

Typically you would edit a kernel command to add an option that affects the boot process. For example, to cause the system to boot to single-user mode, you would add the -s option. For example:

```
grub edit> kernel /platform/i86pc/multiboot -s
```

Enter Return when the kernel command reflects the options you want to use. The command list displays. For example:

```
GNU GRUB version 0.95 (631K lower / 2095488K upper memory)
+-----+
| kernel /platform/i86pc/multiboot -s
| module /platform/i86pc/boot_archive
|
+-----+
```

To boot the system using the new kernel command, enter b.

## Using kernel Command Options and Arguments

Table 10-2 describes some of the kernel command options and arguments you can use to modify the boot behavior of x86 systems. See the kernel(1M) man page for more information.

**Table 10-2** kernel Command Options and Arguments

Option or Argument	Description
-a	Asks the user for configuration information, such as where to find the system file, where to mount root, and even override the name of the kernel itself. Default responses will be contained in square brackets.

**Table 10-2** kernel Command Options and Arguments (Continued)

Option or Argument	Description
<b>-m</b> <i>smf_options</i>	<p>The <b>-m</b> option accepts the <b>quiet</b> and <b>verbose</b> options to control the type and amount of messages that <b>smf(5)</b> displays during boot, and the <b>debug</b> and <b>milestone=milestone_name</b> options to determine the services which are used to boot the system.</p> <p>For the <b>milestone=milestone_name</b> option, valid milestones are <b>none</b>, <b>single-user</b>, <b>multi-user</b>, <b>multi-user-server</b>, and <b>all</b>.</p>
<b>-r</b>	<p>Reconfiguration boot. The system will probe all attached hardware devices and configure the logical namespace in <b>/dev</b>.</p>
<b>-s</b>	<p>Boots only to init level 's'. See <b>init(1M)</b>.</p>
<i>kernel_name</i>	<p>You can specify an alternate kernel to boot by supplying its name as a kernel command argument. The <b>kernel/unix</b> argument specifies the 32-bit kernel. The <b>kernel/amd64/unix</b> argument specifies the 64-bit kernel. Both of these kernels reside below the <b>/platform/i86pc</b> directory structure.</p>

### kernel Command Examples

Adding the **-m** option to the kernel command and specifying the **single-user** milestone boots the system to single-user mode. For example:

```
grub edit> kernel /platform/i86pc/multiboot -m milestone=single-user
```

Adding the **-m** option to the kernel command and specifying the **verbose** argument boots the system to the default milestone and displays verbose messages from **smf**. For example:

```
grub edit> kernel /platform/i86pc/multiboot -m verbose
```

Adding the **kernel/unix** argument to the kernel command boots the system to the default milestone using the 32-bit kernel. For example:

```
grub edit> kernel /platform/i86pc/multiboot kernel/unix
```

Adding the kernel/amd64/unix argument and the -s option to the kernel command boots the system using the 64-bit kernel to the single-user milestone. For example:

```
grub edit> kernel /platform/i86pc/multiboot kernel/amd64/unix -s
```



**Note** – Specify the kernel/amd64/unix argument only on systems capable of booting a 64-bit kernel.

### Verifying Which Kernel the System is Using

If you used an eeprom or kernel command to specify a kernel for the system to boot, you may want to check if the system is using the kernel you specified.

You can use the prtconf command to print the system configuration, and the grep command to search for the path in which the kernel you specified resides. For example:

```
prtconf -v | grep /platform/i86pc/kernel
 value='/platform/i86pc/kernel/amd64/unix'
#
```

The value displayed comes from the system property that lists the name of the kernel. For example:

```
name='whoami' type=string items=1
 value='/platform/i86pc/kernel/amd64/unix'
```

## The `findroot` Command for x86

With the Solaris 10 5/09 release, all Solaris installation methods, including Solaris Live Upgrade, use the `findroot` command for specifying which disk slice to boot on an x86 based system. Previously, the root command, `root (hd0.0.a)`, was used to explicitly specify which disk slice to boot. This information is located in the `menu.1st` file that is used by GRUB. The most common form of the entry in the `menu.1st` file is:

```
findroot (rootfs0,0,a)
kernel$ /platform/i86pc/kernel/$ISADIR/unix
modules$ /platform/i86pc/$ISADIR/boot_archive
```

The `findroot` command discovers the targeted disk, irrespective of the boot device. In addition, the `findroot` command provides enhanced support for booting systems with ZFS and UFS roots. In addition to the `findroot` command, the name of a signature file on the slice, (`<mysign>, 0, a`) is provided. The boot signature is located in the `/boot/grub/bootsign` directory on the system. The name of the signature file varies, depending on the installation method that is used.

For more information, see “Implementation of the `findroot` Command” in the *System Administration Guide: Basic Administration* guide.

# GRUB Boot Archives

GRUB uses two kinds of boot archives:

- Failsafe boot archive
- Primary boot archive

A failsafe boot archive has the following benefits and characteristics:

- Is self-sufficient
- Can boot on its own
- Is created by default during installation of the OS
- Requires no maintenance

By default, the failsafe boot archive is /boot/x86.miniroot-safe.

A primary boot archive shadows a root (/) file system. This boot archive contains all of the kernel modules, driver.conf files, plus a few configuration files. These files are located in the /etc directory. The kernel reads the files in the boot archive before the root (/) file system is mounted. After the root (/) file system is mounted, the kernel discards the boot archive from memory. Then, file I/O is performed against the root device.

By default, the primary boot archive is /platform/i86pc/boot\_archive. In the list of commands represented by a GRUB boot entry, the boot archive is listed as the argument to the module command. For example the module command for the Solaris 10 11/06 s10x\_u3wos\_10 X86 GRUB boot entry lists the primary boot archive:

```
GNU GRUB version 0.95 (631K lower / 2095488K upper memory)
+-----+
| root (hd0,1,a)
| kernel /platform/i86pc/multiboot
| module /platform/i86pc/boot_archive
|
+-----+
```

The module command for the Solaris Failsafe boot entry in the GRUB main menu lists the failsafe boot archive. For example:

```
GNU GRUB version 0.95 (631K lower / 2095488K upper memory)
```

```
+-----+
| root (hd0,1,a)
| kernel /boot/multiboot
| module /boot/x86.miniroot-safe
|
+-----+
```

You can use the `bootadm list-archive` command to list the contents of the primary boot archive. For example:

```
bootadm list-archive
etc/rtc_config
etc/system
etc/name_to_major
etc/driver_aliases
etc/name_to_sysnum
etc/dacf.conf
etc/driver_classes
etc/path_to_inst
etc/mach
etc/devices/devid_cache
etc/devices/mdi_scsi_vhci_cache
etc/devices/mdi_ib_cache
kernel
platform/i86pc/biosint
platform/i86pc/kernel
boot/solaris.xpm
boot/solaris/bootenv.rc
boot/solaris/devicedb/master
boot/acpi/tables
#
```

## Managing the Primary Boot Archive

If any of the files in the operating system that are listed in the archive change, the boot archive must be rebuilt. For modifications to take effect, the rebuild of the archive must take place before the next system reboot.

When the system shuts down, the svc:/system/boot-archive:default service rebuilds the primary boot archive if necessary.

To rebuild the boot archive manually, use the `bootadm` command. The `update-archive` argument causes `bootadm` to rebuild the primary boot archive, and the `-v` option lists the stale files found. For example:

```
bootadm update-archive -v
 changed /etc/path_to_inst
updating /platform/i86pc/boot_archive...this may take a minute
#
```

You can also rebuild the boot archive by booting the system using the failsafe archive which is an available option in the GRUB main menu. During the failsafe boot procedure, when prompted by the system, type **y** to rebuild the primary boot archive.

The `bootadm` command accepts these options:

- `-f` – Forces an update of the boot archive
- `-R` – Allows you to specify an alternate root where the boot archive is located
- `-n` – In an `update-archive` operation, archive content is checked, but not updated

See the `bootadm(1M)` man page for more information.

# Booting a System in the GRUB-Based Boot Environment

This section describes performing a GRUB based disk boot on an x86-based system.

Some of the procedures in this section ask you to use the Reset button to restart the system. If your system does not have a Reset button, use the power switch to restart the system. You might be able to press Ctrl-Alt-Del to interrupt system operation, depending upon the state of the system.

## Booting a System to Default Run Level (Multiuser)

Use this procedure to boot a system that is currently at run level 0 to the default run level.

1. If the system displays Press any key to reboot prompt, press any key to reboot the system.

You can also use the Reset button at this prompt. If the system is shut down, turn the system on with the power switch.

When the boot sequence begins, the GRUB menu displays.

2. To boot the system to the default run level, press Enter to boot the default OS instance.

If you do not choose an entry within 10 seconds, the system automatically boots to the default run level.

3. Login as root and verify that the system has booted to run level 3.

```
who -r
. run-level 3 Mar 1 11:32 3 0 S
#
```

---

**Note** – The who command displays run level 3 because it is not milestone aware.

---

## How to Boot a System to Run Level S (Single-User Level)

Use this procedure to boot a system that is at run level 0 to run level S. Single-user level is used for performing system maintenance.

1. If the system displays the Press any key to reboot prompt, press any key to reboot the system.

You can also use the Reset button at this prompt. If the system is shut down, turn the system on with the power switch.

When the boot sequence begins, the GRUB menu displays.

2. To boot the system to run level S, type e when the GRUB main menu displays.
3. Use the arrow keys to choose the kernel /platform/i86pc/multiboot boot entry.

**Note** – If you cannot use the arrow keys, use the ^ key to scroll up and the v key to scroll down.

4. Type e to edit the kernel /platform/i86pc/multiboot boot entry.

The GRUB edit menu is displayed, enabling you to add options and arguments to the kernel command.

5. Enter a space and -s at the end of the kernel /platform/i86pc/multiboot line. Then, press Enter. For example:

```
grub edit> kernel /platform/i86pc/multiboot -s <Enter>
```

After you press Enter, you are returned to the GRUB main menu. For example:

```
root (hd0,1,a)
kernel /platform/i86pc/multiboot -s
module /platform/i86pc/boot_archive
```

6. To boot the system to the single-user level, type b. The system boots to the single-user level.
7. Type the root password, if prompted. For example:

Root password for system maintenance (control-d to bypass): **password**

8. The superuser prompt displays.

```
Mar 11 16:00:00 su: 'su root' succeeded for root on /dev/console
Sun Microsystems Inc. SunOS 5.10 Generic January 2005
```



#

9. Verify that the system is at run level S.

```
who -r
 . run-level S Mar 11 16:08 S 0 0
#
```

10. Perform the system maintenance task that required the run level change to S.
11. After you complete the system maintenance task, type Control-D to bring the system to the multiuser state.

## How to Boot a System Interactively

Use this procedure to boot a system when you need to specify alternate kernel module directories or an alternate /etc/system file.

1. If the system displays the Press any key to reboot prompt, press any key to reboot the system.

You can also use the Reset button at this prompt. If the system is shut down, turn the system on with the power switch.

When the boot sequence begins, the GRUB menu displays.

2. Type e to access the GRUB edit menu.
3. Use the arrow keys to select the kernel /platform/i86pc/multiboot line.
4. Type e to edit the line.

The GRUB edit menu is displayed, enabling you to add options to the kernel command.

5. Enter a space and -a at the end of the kernel /platform/i86pc/multiboot line. Then, press Enter. For example:  
`grub edit> kernel /platform/i86pc/multiboot -a <Enter>`

The GRUB main menu displays. For example:

```
root (hd0,1,a)
kernel /platform/i86pc/multiboot -a
module /platform/i86pc/boot_archive
```

6. To boot the system interactively, type b. A series of system prompts are displayed. If you want to accept the default, press Enter.
7. Provide the names of alternate directories for kernel modules, or accept the default.

Enter default directory for modules [/platform/i86pc/kernel /kernel /usr/kernel] : **Return**

8. Provide the name of an alternate /etc/system file or accept the default. To help repair a damaged /etc/system file, enter /dev/null as the alternate for /etc/system. For more information see the system(4) man page.

Name of system file [etc/system] : **Return**

The system boots to the default milestone.

## Stopping an x86 System for Recovery

Stop the system using one of the following commands, if possible:

- If the keyboard and mouse are functional, become superuser. Then, type init 0 to stop the system. After the Press any key to reboot prompt appears, press any key to reboot the system.
- If the keyboard and mouse are functional, become superuser. then, type init 6 to reboot the system.

If the system does not respond to any input from the mouse or the keyboard, press the Reset key, if it exists, to reboot the system. Or, you can use the power switch to reboot the system.

## Booting the Failsafe Archive for Recovery

Follow this procedure to boot the Solaris failsafe archive, so that you can update the primary boot archive or repair a critical system resource.



**Note** – Starting with the Solaris 10 6/06 release, the GRUB failsafe interaction changed. Previously, when you booted the failsafe archive, you were prompted to update the boot archives whether the boot archives were out-of-date or not. Now, you are prompted to update the boot archives only if the system detects any out-of-date boot archives. For more information, see the `grub(5)` man page.

1. Use the procedures described the previous section to stop your system and boot it again, or if the system displays the Press any key to reboot prompt, press any key to reboot the system.

When the boot sequence begins, the GRUB main menu displays. For example:

```
GNU GRUB version 0.95 (631K lower / 2095488K upper memory)
```

```
+-----+
| Solaris 10 11/06 s10x_u3wos_10 x86
| Solaris failsafe
| Diagnostic Partition
+-----+
```

Use the ^ and v keys to select which entry is highlighted.

Press enter to boot the selected OS, 'e' to edit the commands before booting, or 'c' for a command-line.

2. When the GRUB main menu displays, press the Space bar to interrupt the boot process.
3. Use the arrow keys to select the Solaris failsafe entry.

4. Type **b** or press Return to boot the failsafe archive.

The system searches for installed OS instances. If the system detects any out-of-date boot archives, a message similar to the following displays:

Searching for installed OS instances...

An out of sync boot archive was detected on /dev/dsk/c1d0s0.

The boot archive is a cache of files used during boot and should be kept in sync to ensure proper system operation.

Do you wish to automatically update this boot archive? [y,n,?]

5. Type **y** to automatically update the out-of-date boot archive. If multiple out-of-date boot archives are found, the system prompts you to update each boot archive, one at a time. Type **y** and press Return to update each boot archive.
6. After the boot archive has been updated successfully, the system searches again for all installed OS instances. If only one OS instance is found, a message similar to the following displays:

Searching for installed OS instances...

Solaris 10 11/06 s10x\_u3wos\_10 X86 was found on /dev/dsk/c1d0s0.

Do you wish to have it mounted read-write on /a? [y,n,?]

If the system finds multiple OS instances, it presents prompts that allow you to select the OS instance to mount on /a.

7. Type **y** to mount the OS instance /a.
8. Perform the tasks required to correct the problem encountered. For example, to forcibly update a corrupt boot archive, type:

```
bootadm update-archive -f -R /a
```

Creating ram disk on /a

updating /a/platform/i86pc/boot\_archive...this may take a minute

#

The -f option forces the update, and the -R option allows you to specify an alternate root, in this case /a.

9. When you have finished correcting the problem, change directory to root (/) and unmount the device mounted on /a.

```
cd /
umount /a
```

10. Reboot the system.

```
init 6
```

## Interrupting an Unresponsive System

When a system freezes or stops responding to the keyboard, you might have to interrupt it. When you interrupt the system, all active processes stop immediately. It does not allow you to flush memory or to synchronize file systems.

## Aborting an Unresponsive System

To abort or interrupt an unresponsive system:

- Attempt a remote login on the unresponsive system to locate and kill the offending process.
- Attempt to reboot the unresponsive system gracefully.
- Hold down the Ctrl+Alt+Del key sequence on the keyboard of the unresponsive system. The system reboots.
- Press the reset button.
- Power off the system and power it back on.
- Some Sun systems allow the system to be reset remotely using the LOM software.

# Exercise: Using GRUB

In this exercise, you change entries in the GRUB menu to boot the system in different ways, and use the `eeprom` command to set boot parameters.

## Preparation

Refer to the lecture notes as necessary to perform the following tasks and answer the questions listed.

This exercise requires a system that is capable of using a 64-bit kernel in order to complete all lab steps. If you are using a system that is only 32-bit capable, you will only be able to run a subset of the lab steps.

## Task

Complete the following steps:

1. Login as `root` and open a terminal window.
2. Use the `init` command to shut down the system to run level 0.
3. When the `Press any key to reboot` prompt displays, press any key to start the boot process.
4. When the GRUB main menu displays, press the Space bar to interrupt the boot process. The GRUB main menu remains displayed.
5. Select the `Solaris 10 10/09 s10x_u8wos_08a x86` entry, and press `e` to display the list of commands it represents.
  - What is the pathname of the boot archive this boot entry uses?

---

- Is a specific kernel listed in the `kernel` command line?

---

6. Press the `Esc` key to return to the GRUB main menu.
7. Select the `Solaris failsafe` entry, and press `e` to display the list of commands it represents. What is the pathname of the boot archive this boot entry uses?

---
8. Press the `Esc` key to return to the GRUB main menu.

## Exercise: Using GRUB

---

- Unauthorized reproduction or distribution prohibited. Copyright© 2011, Oracle and/or its affiliates.
9. Select the Solaris 10 10/09 s10x\_u8wos\_08a X86 entry, and press Return to boot the system to the default milestone.
  10. Login as root and open a terminal window.
  11. Use the eeprom command to display the list of boot parameters defined in the bootenv.rc file. Is the boot-file parameter currently defined?
- 
12. Use the prtconf command to display the name of the kernel that the system is currently using. If the kernel listed is a 64-bit kernel, for example /platform/i86pc/kernel/amd64/unix, proceed to the next step. If the kernel listed is a 32-bit kernel, for example, /platform/i86pc/kernel/unix, skip to step 20.

**Note** – You cannot use a 64-bit kernel to boot a 32-bit system.

---



13. Use the eeprom command to set the boot-file parameter so the system will boot using the 32-bit kernel. Verify that the parameter is now set.
  14. Use the init command to reboot your system. When the boot process completes, login as root and open a terminal window.
  15. Use the prtconf command to display the name of the kernel that the system is now using.
- 
16. Use the init command to reboot your system. When the GRUB main menu displays, press e to interrupt the boot process and edit the default boot entry. Does the kernel command list a specific kernel to boot?
- 
17. Select the kernel command from the GRUB menu, and e to enter the GRUB editor.
  18. Edit the kernel command to add the name of the 64-bit kernel to the end of the command line.
  19. Verify the kernel command specifies the 64-bit kernel, and enter b to boot the system.
  20. Login as root and open a terminal window.

21. Use the `prtconf` command to display the name of the kernel that the system is currently using. Is the system using the 32-bit kernel you specified with the `eeprom` command, or the 64-bit kernel you specified using the `kernel` command?

---

22. Use the `eeprom` command to set the `boot-file` parameter to a null value, and verify the change.
23. Use the `init` command to reboot your system. When the GRUB main menu displays, press the Space bar to interrupt the boot process.
24. Select the Solaris `failsafe` boot option and boot the system. Allow the Solaris instance that the system finds to mount on `/a`.
25. Change directory to `/a/platform/i86pc`, and create a directory named `backup`.
26. Move the `boot_archive` file into the `backup` directory and verify that it no longer exists in `/a/platform/i86pc`.
27. Use the `bootadm` command to create a new boot archive in the `/a/platform/i86pc` directory. Verify that it exists.
28. Change directory to root (`/`), unmount `/a`, and reboot the system.

## Exercise Summary



**Discussion** – Take a few minutes to discuss what experiences, issues, or discoveries you had during the lab exercises.

- Experiences
- Interpretations
- Conclusions
- Applications

# Exercise Solutions: Using GRUB

This section contains solutions to the exercise.

## Task

Complete the following steps:

1. Login as root and open a terminal window.
2. Use the init command to shut down the system to run level 0.  
`# init 0`
3. When the Press any key to reboot prompt displays, press any key to start the boot process.
4. When the GRUB main menu displays, press the Space bar to interrupt the boot process. The GRUB main menu remains displayed. For example:

```
GNU GRUB version 0.95 (572K lower / 3380808K upper memory)
+-----+
| Solaris 10 10/09 s10x_u8wos_08a X86
| Solaris failsafe
| Diagnostic Partition
+-----+
```

Use the ^ and v keys to select which entry is highlighted.

Press enter to boot the selected OS, 'e' to edit the commands before booting, or 'c' for a command-line.

---

**Note** – The Diagnostic Partition item in the GRUB menu is not present unless the system had an unrecoverable startup failure.

---

5. Select the Solaris 10 10/09 s10x\_u8wos\_08a X86 entry, and press e to display the list of commands it represents.
  - What is the pathname of the boot archive this boot entry uses?
  - Is a specific kernel listed in the kernel command line?

GNU GRUB version 0.95 (572K lower / 3380808K upper memory)

```
+-----+
| findroot (rootfs0,0,a)
| kernel /platform/i86pc/multiboot
| module /platform/i86pc/boot_archive
+-----+
```

*The Solaris 10 10/09 s10x\_u8wos\_08a X86 boot entry uses the /platform/i86pc/boot\_archive boot archive. The kernel command does not list a specific kernel to load.*

6. Press the Esc key to return to the GRUB main menu.
7. Select the Solaris failsafe entry, and press e to display the list of commands it represents. What is the pathname of the boot archive this boot entry uses?

GNU GRUB version 0.95 (572K lower / 3380808K upper memory)

```
+-----+
| findroot (rootfs0,0,a)
| kernel /boot/multiboot kernel/unix -s
| module /boot/x86.miniroot-safe
+-----+
```

*The Solaris failsafe boot entry uses the /boot/x86.miniroot-safe boot archive.*

8. Press the Esc key to return to the GRUB main menu.
9. Select the Solaris 10 10/09 s10x\_u8wos\_08a X86 entry, and press Return to boot the system to the default milestone.
10. Login as root and open a terminal window.
11. Use the eeprom command to display the list of boot parameters defined in the bootenv.rc file. Is the boot-file parameter currently defined?

```
eeprom
kbd-type=US-English
ata-dma-enabled=1
atapi-cd-dma-enabled=0
ttyb-rts-dtr-off=false
ttyb-ignore-cd=true
```

```
ttya-rts-dtr-off=false
ttya-ignore-cd=true
ttyb-mode=9600,8,n,1,-
ttya-mode=9600,8,n,1,-
lba-access-ok=1
prealloc-chunk-size=0x2000
bootpath=/pci@0,0/pci-ide@5/ide@0/cmdk@0,0:a
console=text
#
No, the boot-file parameter is not defined.
```

12. Use the **prtconf** command to display the name of the kernel that the system is currently using. If the kernel listed is a 64-bit kernel, for example /platform/i86pc/kernel/amd64/unix, proceed to the next step. If the kernel listed is a 32-bit kernel, for example, /platform/i86pc/kernel/unix, skip to step 20.

```
prtconf -v | grep /platform/i86pc/kernel
 value='/platform/i86pc/kernel/amd64/unix'
#
```

---

**Note –** You cannot use a 64-bit kernel to boot a 32-bit system.

---



13. Use the **eeprom** command to set the **boot-file** parameter so the system will boot using the 32-bit kernel. Verify that the parameter is now set.

```
eeprom boot-file=kernel/unix
eeprom boot-file
boot-file=kernel/unix
#
```

14. Use the **init** command to reboot your system. When the boot process completes, login as root and open a terminal window.

```
init 6
```

15. Use the **prtconf** command to display the name of the kernel that the system is now using.

```
prtconf -v | grep /platform/i86pc/kernel
 value='/platform/i86pc/kernel/unix'
#
```

16. Use the `init` command to reboot your system. When the GRUB main menu displays, press `e` to interrupt the boot process and edit the default boot entry. Does the kernel command list a specific kernel to boot?

```
init 6
```

(output omitted)

```
GNU GRUB version 0.95 (572K lower / 3380808K upper memory)
```

```
+-----+
| findroot (rootfs0,0,a)
| kernel /platform/i86pc/multiboot
| module /platform/i86pc/boot_archive
+-----+
```

*No, the kernel command still lists no specific kernel.*

17. Select the `kernel` command from the GRUB menu, and press `e` to enter the GRUB editor.

```
grub edit> kernel /platform/i86pc/multiboot
```

18. Edit the `kernel` command to add the name of the 64-bit kernel to the end of the command line.

```
grub edit> kernel /platform/i86pc/multiboot kernel1/amd64/unix
```

19. Verify the `kernel` command specifies the 64-bit kernel, and enter `b` to boot the system.

20. Login as root and open a terminal window.

21. Use the `prtconf` command to display the name of the kernel that the system is currently using. Is the system using the 32-bit kernel you specified with the `eeprom` command, or the 64-bit kernel you specified using the `kernel` command?

```
prtconf -v | grep /platform/i86pc/kernel
```

```
value='/platform/i86pc/kernel/amd64/unix'
```

*The system is using the 64-bit kernel you specified using the `kernel` command, overriding the entry in `bootenv.rc` file created by the `eeprom` command.*

22. Use the `eeprom` command to set the `boot-file` parameter to a null value, and verify the change.

```
eeprom boot-file=""
eeprom boot-file
boot-file=
#
```

23. Use the `init` command to reboot your system. When the GRUB main menu displays, press the Space bar to interrupt the boot process.

```
init 6
```

(output omitted)

```
GNU GRUB version 0.95 (572K lower / 3380808K upper memory)
```

```
+-----+
| Solaris 10 5/09 s10x_u7wos_08 X86
| Solaris failsafe
| Diagnostic Partition
|
```

Use the ^ and v keys to select which entry is highlighted.

Press enter to boot the selected OS, 'e' to edit the commands before booting, or 'c' for a command-line.

24. Select the Solaris failsafe boot option and boot the system. Allow the Solaris instance that the system finds to mount on /a.

Solaris 10 5/09 s10x\_u7wos\_08 X86 was found on /dev/dsk/c1d0s0.

Do you wish to have it mounted read-write on /a? [y,n,?] **y**

25. Change directory to /a/platform/i86pc, and create a directory named backup.

```
cd /a/platform/i86pc
mkdir backup
#
```

26. Move the boot\_archive file into the backup directory and verify that it no longer exists in /a/platform/i86pc.

```
mv boot_archive backup
ls
backup biosint boot kernel multiboot ucode
#
```

27. Use the `bootadm` command to create a new boot archive in the `/a/platform/i86pc` directory. Verify that it exists.

```
bootadm update-archive -f -R /a
```

```
Creating boot_archive for /a
updating /a/platform/i86pc/boot_archive
```

```
ls boot_archive
```

```
boot_archive
```

```
#
```

28. Change directory to root (/), unmount /a, and reboot the system.

```
cd /
```

```
umount /a
```

```
init 6
```

---

## Module 11

---

# Performing Legacy Boot and Shutdown Procedures

---

## Objectives

Upon completion of this module, you should be able to:

- Describe the Solaris SPARC boot design
- Identify boot process phases
- Identify run level fundamentals
- Control boot processes
- Perform system shutdown procedures

# Solaris SPARC Boot Design Overview

The Solaris x86 boot design feature was introduced in the Solaris 10 1/06 release. With the Solaris 10 10/08 release, the Solaris SPARC bootstrap was redesigned to be in line with the Solaris x86 boot design. The improved Solaris boot design provides support for Solaris boot archives, which was earlier available only on x86 platforms, on the SPARC platform. With this simplified boot design, other non-UFS file systems, such as ZFS are supported as root file systems. Although the implementation of the SPARC boot has changed, boot tasks remain the same.

## New Design

The Solaris boot design architecture for SPARC is similar to the x86 architecture. The difference is how the boot device and fields are selected at boot time. The SPARC platform continues to use the open boot PROM (OBP) and its commands for selections, while the x86 platform uses the BIOS and the GRUB menu. On the SPARC and x86 platforms, there is one primary boot archive and one fail-safe boot archive. The following are the four phases of the new boot process:

1. Boot loader phase: Loads the Solaris root file system archive from the media to memory.
2. Booter phase: Reads the ramdisk directly and executes it. This is the only portion of the boot process that requires knowledge of the root file system format.
3. Ramdisk phase: Extracts the kernel image from the boot archive and executes it. Neither the booter nor the kernel need to know about the format of the boot archive.
4. Kernel phase: Extracts the rest of the primary modules from the boot archive, initializes itself, mounts the real root file system, and discards the boot archive.

For example, when a system powers on, the firmware loads the boot block, which in turn, retrieves a boot loader. The boot loader gathers a root file system and places it in a ramdisk. The ramdisk loads and executes a secondary loader, which assembles core kernel module from the ramdisk and transfers control to the Solaris kernel. The kernel probes input and output (I/O devices and then mounts the root file system.

## Features

The SPARC platform bootadm command functions the same as on the x86 platform. It handles the boot archive update and verification. Use it to boot archives:

```
bootadm update-archive
```

The boot archive service is controlled with the svcadm command and managed by the SMF:

```
svcs:/system/boot-archive:default
```

Use the `-p` option to specify the platform or machine hardware class on a client system where the client platform differs from the server platform. You must use the `-p` option with the `-R` option to obtain an alternate root path. For example:

```
boot -p platform -R [altroot]
```

where the platform specified can be `i86pc`, `sun4u`, or `sun4v`.

The boot loader and the OS kernel are separated in the new boot design. The loader can evolve independently from the Solaris kernel, making it a replaceable component. The boot loader is disassociated with the root file system and does not need to comprehend it. The boot loader accepts any file system to be the Solaris root without changing the boot code.

Disk boot and network boot from disk are unified, making the LAN and WAN identical from the Solaris perspective. The kernel always initializes from in-memory text and data regardless from where they are loaded.

## System Requirements

The following are the SPARC and x86 boot design minimum system requirements, dependencies, and limitations:

- Memory
  - Text-based installation requires a minimum of 384 Mbytes.
  - GUI based installation requires a minimum of 768 Mbytes.
  - SPARC and x86 platform memory requirements are the same.
- Use the `root_archive` command to pack and unpack the miniroot. Only SPARC based systems that support the new boot architecture have the ability to pack and unpack a compressed miniroot.
- The Solaris 10 `root_archive` tool isn't compatible with Solaris release tool versions. Perform ramdisk manipulation only on systems running the same Solaris and archive release.

## Legacy Boot and Shutdown

Prior to Solaris 10, Sun systems used a series of scripts to start and stop processes typically associated with run levels. Each run level had an associated script located in the /sbin directory, with some scripts hard-linked to each other. The init daemon was responsible for starting and stopping system services. These structures are still present and are referred to through out this module as legacy boot processes. With Solaris 10, you can use the Service Management Facility (SMF) to start and stop processes typically associated with run levels. The SMF is covered in the next module of this course. This module discusses the features and structures of the legacy boot architecture.

## Identifying Boot Process Phases

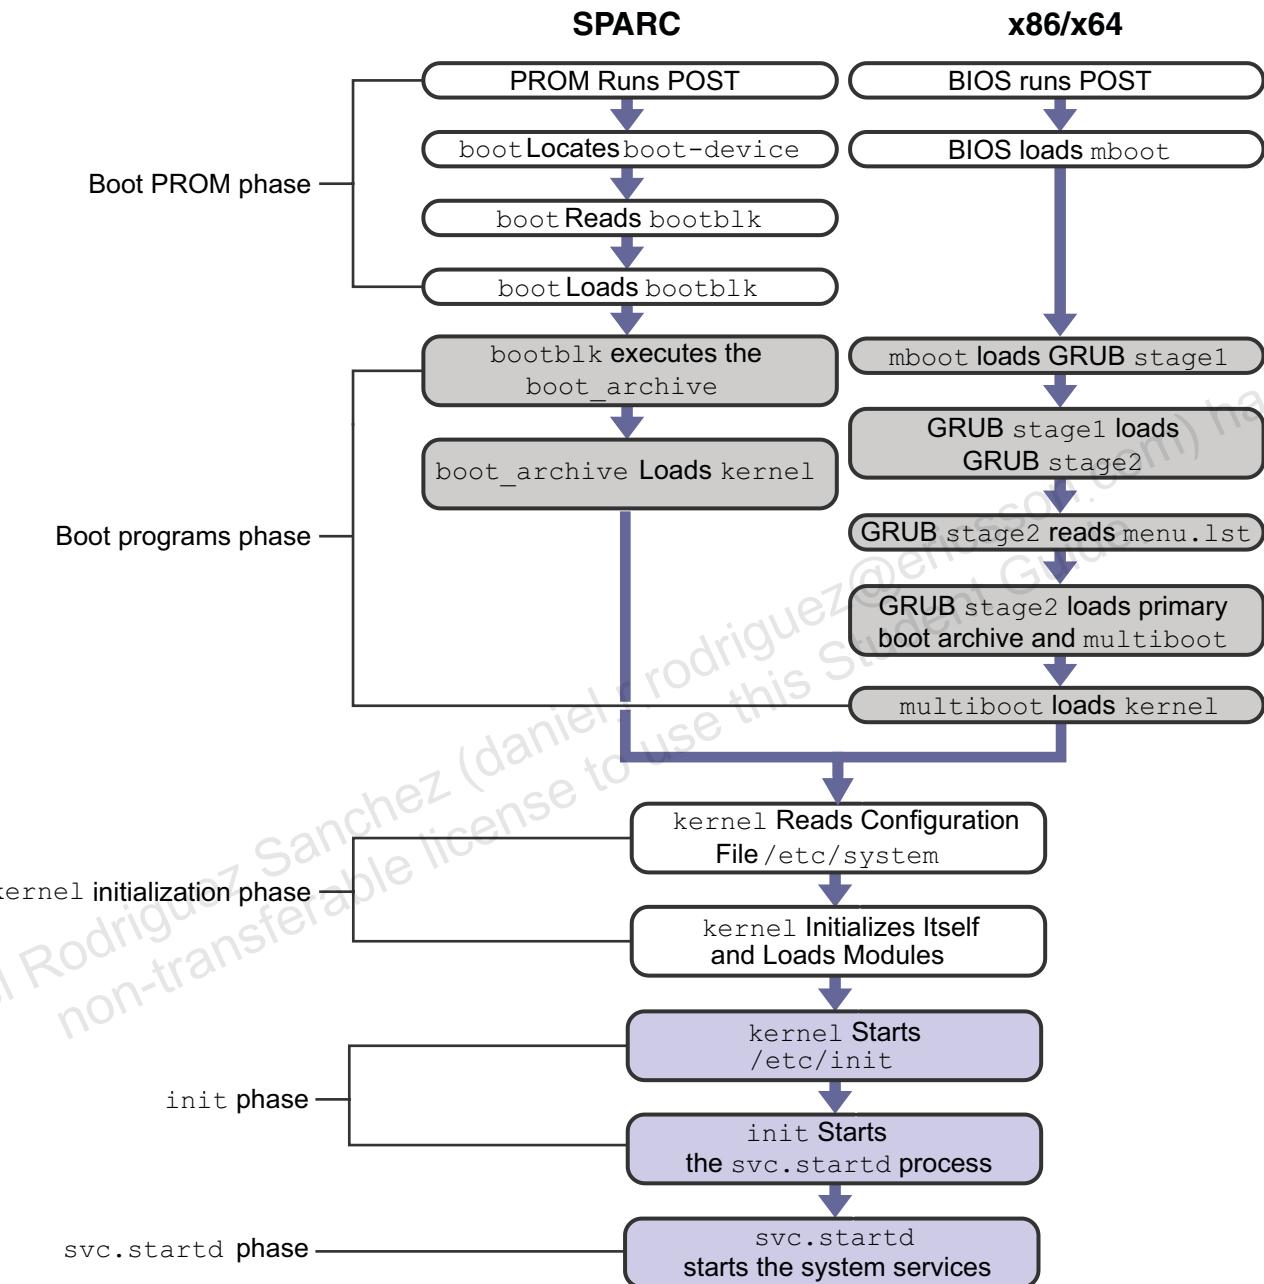
In general, when a system is powered on, the PROM monitor on SPARC systems, and the BIOS monitor on x86/x64 systems, runs a POST procedure that checks the hardware and memory on the system. If no errors are found, and the `auto-boot?` parameter is set to `true` on SPARC systems, or the `timeout` parameter is set to a positive value in the `/boot/grub/menu.1st` file on x86/x64 systems, the system automatically begins the boot process.

The entire boot process is described by five distinct phases:

- The boot PROM phase
- The boot programs phase
- The kernel initialization phase
- The `init` phase
- The `svc.startd` phase

The boot PROM and boot programs phases differ significantly between SPARC and x86/x64 systems, but subsequent phases on these platforms are generally the same and use similar sets of files.

Figure 11-1 shows the phases of the boot process.



**Figure 11-1** Phases of the Boot Process



## SPARC Boot PROM Phase

---

**Note** – This section describes the first part of the boot sequence on systems with a single board only. Multi-board systems run a sequence that identifies and sets up their hardware before the boot PROM is involved.

---

The boot PROM performs the following steps during the first part of the boot sequence on a SPARC-based system:

- The PROM runs the POST.

The boot PROM firmware runs the POST to verify the system hardware and memory. After successful completion of the diagnostics, the boot sequence begins.
- The PROM displays the system identification banner.

The model type, processor type and speed, keyboard status, PROM revision number, amount of installed random access memory (RAM), NVRAM serial number, Ethernet address, and host ID are displayed.
- The boot PROM determines the boot device by reading the PROM parameter boot-device.
- The boot PROM reads the disk label located at Sector 0 on the default boot device.
- The boot PROM finds the boot program from the default boot device programmed into the PROM.

The boot PROM program reads the system primary boot program called bootblk (located at sectors 1 through 15) that contains a UNIX file system (UFS) reader. (The bootblk program is placed on the disk by the installboot command during system installation.)

The boot command loads the bootblk program from its location on the boot device into memory.

## SPARC Boot Programs Phase

**Note** – Solaris releases that support the SPARC Newboot feature might require SPARC systems with at least 512 Mbytes of memory.

The following describes the boot programs phase on a SPARC-based system:

- The bootblk program loads the secondary boot program, ufsboot, from the boot device into memory.

The path to ufsboot is recorded in the bootblk program, which is installed by the installboot command.

- The ufsboot program locates and loads the appropriate two-part kernel.

The core of the kernel is two pieces of static code called genunix and unix, where genunix is the platform-independent generic kernel file and unix is the platform-specific kernel file.

When ufsboot loads these two files into memory, they are combined to form the running kernel.

On a SPARC system running in 64-bit mode, the two-part kernel is located in the directory:

```
/platform/`uname -m'/kernel/sparcv9
```

**Note** – Solaris 10 for SPARC only runs on 64-bit systems.



**Note** – To determine the platform name (for example, the system hardware class), type the uname -m command. For example, when you type this command on an Ultra 10 workstation, the console displays sun4u.

## x86/x64 Boot PROM Phase

The BIOS ROM performs the following steps during the first part of the boot sequence on an x86/x64-based system:

- After a PC-compatible machine is turned on, the system firmware in the BIOS ROM executes a power-on self test (POST), runs BIOS extensions in peripheral board ROMs, and invokes software interrupt INT 19h, Bootstrap.
- The INT 19h handler typically performs the standard PC-compatible boot, which consists of trying to read the first physical sector from the first diskette drive, or, if that fails, from the first hard disk. The processor then jumps to the first byte of the sector image in memory.
- The first sector on a hard disk on an x86/x64 system contains the master boot block, which contains the master boot (mboot) program and the FDISK table, named for the PC program that maintains it.

## x86/x64 Boot Programs Phase

The following describes the boot programs phase on an x86/x64-based system:

- The master boot finds the active partition in the FDISK table, loads its first sector (GRUB stage1), and jumps to its first byte in memory.
- If GRUB stage1 is installed on the master boot block (see the -m option of `installgrub(1M)`), then stage2 is loaded directly from the Solaris FDISK partition regardless of the active partition.
- An x86 FDISK partition for the Solaris software begins with a one-cylinder boot slice, which contains GRUB stage1 in the first sector, the standard Solaris disk label and volume table of contents (VTOC) in the second and third sectors, and GRUB stage2 in the fiftieth and subsequent sectors. The area from sector 4 to 49 might contain boot blocks for older versions of Solaris.
- When the FDISK partition for the Solaris software is the active partition, the master boot program (mboot) reads the partition boot program in the first sector of the active partition (GRUB stage1) into memory and jumps to it.
- GRUB stage1 in turn reads the GRUB stage2 program into memory and jumps to it. The GRUB stage2 program contains code that allows it to navigate the UFS structure on the root file system.

- The GRUB stage2 program locates the GRUB menu configuration file /boot/grub/menu.1st and displays the GRUB main menu.
- When the GRUB menu displays, the user can choose to boot an operating system on a different partition, a different disk, or possibly from the network.
- Either selected by the user or by default timeout, GRUB executes commands from /boot/grub/menu.1st to load a pre-constructed boot archive containing a kernel program and data. The default primary boot archive for Solaris 10 Update 3 is /platform/i86pc/boot\_archive.
- GRUB loads a program called multiboot, which implements the kernel side of the Multiboot Specification.
- The multiboot program assembles the core kernel modules from the boot\_archive and starts the operating system, links in the necessary modules from the boot archive and mounts the root filesystem on the real root device.
- At this point, the kernel regains storage I/O, mounts additional file systems (see vfstab(4)), and starts various operating system services (see smf(5)).

## The kernel Initialization Phase

The following describes the kernel initialization phase:

- The kernel reads its configuration file, called /etc/system
- The kernel initializes itself and begins loading modules.

The kernel uses the ufsboot command to load the files. When it has loaded enough modules to mount the / (root) file system, it unmaps the ufsboot program and continues.

- The kernel starts the /etc/init daemon.

---

**Note** – The /etc/init file is a symbolic link to /sbin/init.

---



The SunOS™ kernel is a small static core, consisting of genunix and unix and many dynamically loadable kernel modules.

Modules can consist of device drivers, binary files to support file systems, and streams, as well as other module types used for specific tasks within the system.

## Identifying Boot Process Phases

---

The modules that make up the kernel typically reside in the directories /kernel and /usr/kernel. Platform-dependent modules reside in the /platform/'uname -m'/kernel and /platform/'uname -i'/kernel directories.

On x86/x64 systems, 32-bit and 64-bit modules share parent directory structures in which subdirectories named amd64 store 64-bit modules. For example, 32-bit device drivers are found in /kernel/drv, and 64-bit device drivers are found in /kernel/drv/amd64:

```
file /kernel/drv/cmdk /kernel/drv/amd64/cmdk
/kernel/drv/cmdk: ELF 32-bit LSB relocatable 80386 Version 1
/kernel/drv/amd64/cmdk: ELF 64-bit LSB relocatable AMD64 Version 1
#
```

The following describes the types of module subdirectories contained in the /kernel, /usr/kernel, /platform/'uname -m'/kernel, or /platform/'uname -i'/kernel directories:

- drv/sparcv9, drv, drv/amd64 – Device drivers
- exec/sparcv9, exec, exec/amd64 – Executable file formats
- fs/sparcv9, fs, fs/amd64 – File system types, for example, ufs, nfs, and proc
- misc/sparcv9, misc, misc/amd64 – Miscellaneous modules, for example, usb
- sched/sparcv9, sched, sched/amd64 – Scheduling classes (process execution scheduling)
- strmod/sparcv9, strmod, strmod/amd64 – Streams modules (generalized connection between users and device drivers)
- sys/sparcv9, sys, sys/amd64 – System calls (defined interfaces for applications to use)

On SPARC systems, the /kernel/drv/sparcv9 directory contains all of the device drivers that are used for system boot. The /usr/kernel/drv/sparcv9 directory is used for all other device drivers.

---

**Note** – The sparcv9 CPU is the type of SPARC CPU that supports 64-bit processing.

---

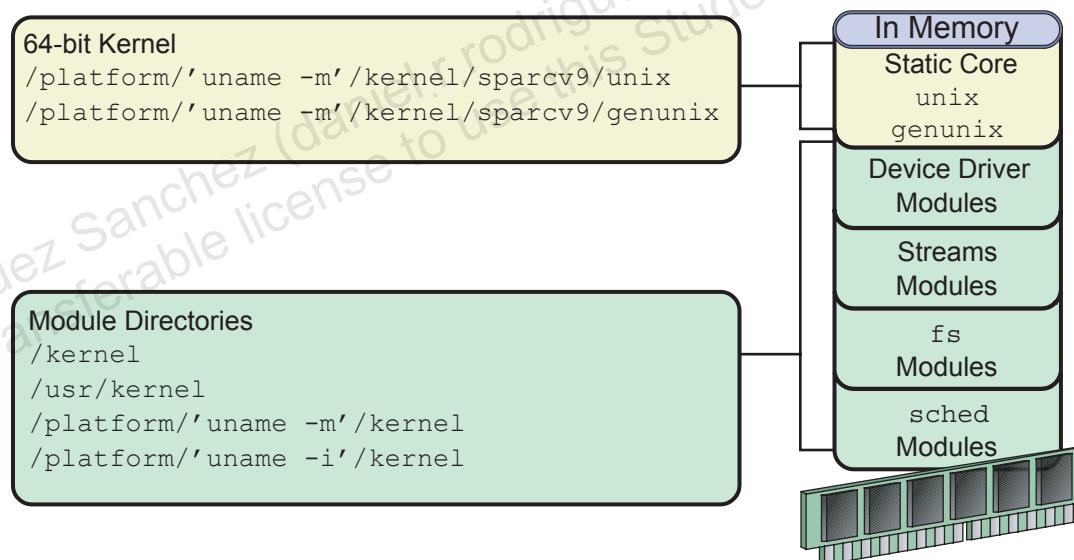


Modules are loaded automatically as needed either at boot time or on demand, if requested by an application. When a module is no longer in use, it might be unloaded on the basis that the memory it uses is needed for another task.

After the boot process is complete, device drivers are loaded when devices, such as tape devices, are accessed. This process is called autoconfiguration because some kernel driver modules are loaded automatically when needed.

Upon initial or reconfiguration boot, the system does a self-test and checks for all devices that are attached.

The advantage of this dynamic kernel arrangement is that the overall size of the kernel is smaller, which makes more efficient use of memory and allows for simpler modification and tuning. Figure 11-2 shows this arrangement on a SPARC-based system.



**Figure 11-2** Dynamic kernel Arrangement on a SPARC system

On x86/x64 systems the 32-bit kernel static core is composed of these files:

- /platform/i86pc/kernel/unix
- /kernel/genunix

On x86/x64 systems the 64-bit kernel static core is composed of these files:

- /platform/i86pc/kernel/amd64/unix
- /kernel/amd64/genunix



## The /etc/system File and Kernel Configuration

**Caution** – The Solaris OS builds the kernel based upon the system size (memory, CPUs, etc). The performance of the default kernel that is built can handle most system activities. Make modifications with extreme caution.

---

The /etc/system file is the control file for modifying which modules and parameters are to be loaded by the kernel at boot time. By default, all lines in this file are commented out.

Modifying the kernel's behavior (or configuration) requires editing the /etc/system file. Altering this file allows you to modify the kernel's treatment of loadable modules as well as to modify kernel parameters for some performance tuning.

The ufsboot program contains a list of default loadable kernel modules that are loaded at boot time. However, you can override this list by modifying the /etc/system file to control which modules, as well as which parameters, are loaded.

All changes to this file take effect after a reboot.

The /etc/system file can explicitly control:

- The search path for default kernel modules to be loaded at boot time
- The root file system type and device
- The modules that are excluded from loading automatically at boot time
- The modules to be forcibly loaded at boot time, rather than at first access
- The new values to override the default kernel parameter values

---

**Note** – Command lines must be 80 characters or less in length, and comment lines must begin with an asterisk (\*) and end with a newline character.

---



The /etc/system file is divided into five distinct sections:

- moddir:  
Sets the search path for default loadable kernel modules. You can list together multiple directories to search, delimited either by blank spaces or colons. If the module is not found in the first directory, the second directory is searched, and so on.
- root device and root file system configuration:  
Sets the root file system type to the listed value. The default is rootfs:ufs.  
rootdev:/sbus@1,f8000000/esp@0,800000/sd@3,0:a
- exclude:  
Does not allow the loadable kernel modules to be loaded during kernel initialization, for example:  
exclude: sys/shmsys
- forceload:  
Forces the kernel modules to be loaded during kernel initialization, for example:  
forceload: drv/vx  
The default action is to load a kernel module automatically when its services are first accessed during runtime by a user or an application.
- set:  
Changes kernel parameters to modify the operation of the system, for example:  
set maxusers=40



## Editing the /etc/system File

**Caution** – Before you edit the /etc/system file, you should make a backup copy. If you enter incorrect values in this file, the system might not be able to boot.

---

The following example shows how to copy the original /etc/system file to a backup file and then edit the /etc/system file.

```
cp /etc/system /etc/system.orig
vi /etc/system
```

If a boot process fails because of an unusable /etc/system file, on a SPARC system, issue the interactive boot command: boot -a. When the system prompts you to enter the name of the system file, type in the name of your backup system file, or, alternatively, enter /dev/null for a null configuration file.

```
ok boot -a
Enter filename [kernel/sparcv9/unix]: <Return>
Enter default directory for modules [/platform...]: <Return>
Name of system file [etc/system]: etc/system.orig - or - /dev/null
root filesystem type [ufs]: <Return>
Enter physical name of root device [/...]: <Return>
(further boot messages omitted)
```

Solaris x86/x64-based systems also provide an interactive boot procedure. In the GRUB menu, adding the -a option to the kernel command used to boot the system initiates an interactive boot procedure similar to that found on SPARC systems.

In the command list for the boot option you want to use, add the -a option as in the following example, then boot the system using the modified command.

```
grub edit> kernel /platform/i86pc/multiboot -a
```

---

**Note** – Some x86/x64 systems may require the -ar options to initiate an interactive boot procedure.

---

You can also boot x86/x64 systems using the failsafe boot archive. Doing so allows you to mount the root file system, and make corrections to the /etc/system file as necessary.

To boot the system using the failsafe archive, select the Solaris failsafe entry from the GRUB main menu, and boot the system. The root file system mounts automatically below /a, so you would edit /a/etc/system to correct the file.

## The init Phase

The next to the last phase of the boot process is the init phase. During this phase, the init daemon starts the svc.startd daemon that is responsible for starting and stopping services as requested. The /sbin/init phase uses information stored in the /etc/inittab file.

### The /etc/inittab File

Each line in the /etc/inittab file contains the following four fields:

*id:rstate:action:process*

Table 11-1 describes the fields in an inittab entry.

**Table 11-1** Fields in the inittab File

Field	Description
id	Two character identifier for the entry
rstate	Run levels to which this entry applies
action	Defines how the process listed should be run  For a description of the action keywords see man inittab
process	Defines the command to execute

The following example shows the default inittab file installed with the Solaris 10 OS. The lines of output are described after the example:

```
ap::sysinit:/sbin/autopush -f /etc/iu.ap
sp::sysinit:/sbin/soconfig -f /etc/sock2path
smf::sysinit:/lib/svc/bin/svc.startd>/dev/msglog 2<>/dev/msglog
</dev/console
p3:s1234:powerfail:/usr/sbin/shutdown -y -i5 -g0 >/dev/msglog
2<>/dev/msglog
```



**Note** – Message output from rc scripts is directed to the logfile for the appropriate milestone. Prior to the Solaris 8 OS, all of these messages were written to the /dev/console file. The /dev/msglog file is used for message output collection from system startup or background applications.

Table 11-2 shows an explanation for each action keyword.

**Table 11-2** The action Field Keywords

Keyword	Explanation
sysinit	Executes the process before the init process tries to access the console, for example, the console login prompt. The init process waits for completion of the process before it continues to read the inittab file.
powerfail	Executes the process only if the init process receives a power fail signal.

Information about additional action keywords is available in the inittab man page. The following describes each of the lines in the /etc/inittab file in order:

1. Initializes the STREAMS modules used for communication services.
2. Configures the socket transport providers for network connections.
3. Initializes the svc.startd daemon for SMF.
4. Describes a power fail shutdown.

# Legacy Run Level Fundamentals

A run level is a system state, represented by a digit or letter, that defines the services and resources that are currently available to users.

Table 11-3 shows the eight run levels found in the Solaris OS.

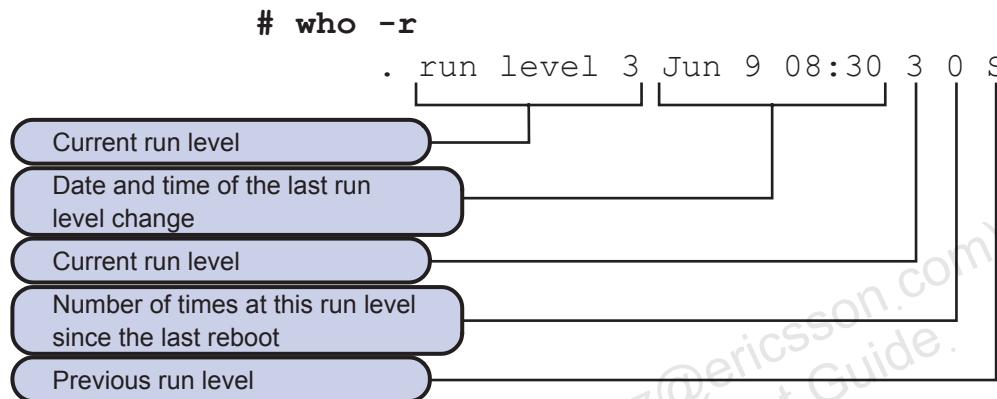
**Table 11-3** Solaris OS Run Levels

Run Level	Milestone	Function
0		System is running the PROM monitor.
s or S	single-user	Solaris OS single-user mode with critical file systems mounted and accessible.
1		The system is running in a single-user administrative state with access to all available file systems.
2	multi-user	The system is supporting multiuser operations. Multiple users can access the system. All system daemons are running except for the Network File System (NFS) server and some other network resource server related daemons.
3	multi-user-server	The system is supporting multiuser operations and has NFS resource sharing and other network resource servers available.
4		This level is currently not implemented.
5		A transitional run level in which the Solaris OS is shut down and the system is powered off.
6		A transitional run level in which the Solaris OS is shut down and the system reboots to the default run level.

## Determining the Current Run Level of a System

To determine the current run level of a system, use the `who -r` command.

Figure 11-3 shows output from the command.



**Figure 11-3** The System's Current Run Level

## Changing Run Levels

Run levels are sometimes referred to as `init` states because you can use the `init` command to transition between run levels. The `init` command passes the required run level to `svc.startd`.

Use the `init` command to manually initiate run-level transitions. Change run levels with the `shutdown`, `halt`, `reboot`, and `poweroff` commands. Use the `svcadm` command to change system milestones.

The `svcadm` command doesn't change run levels. It changes system milestones.

---

**Note** – Prior to Solaris 10, the `init` daemon started and stopped system services. With Solaris 10, the `svc.startd` daemon starts and stops system services. The `init` daemon initializes stream modules, configures socket transport providers, sets up the system for a correct response to a power fail shutdown, and starts the `svc.startd` daemon.



# Controlling Legacy Boot Processes

The Solaris OS provides a series of legacy scripts to stop and start processes typically associated with run levels or milestones.

## The /sbin Directory

Each run level has an associated script located in the /sbin directory, with some scripts hard-linked to each other.

The scripts are executed by the svc.startd daemon to set up variables, test conditions, and make calls to other scripts that start and stop processes for that run level.

The rc0, rc5, and rc6 scripts are hard-linked to each other. Notice that each script is assigned the same inode number. The following is an example of the hard links:

```
ls -li /sbin/rc*
2317 -rwxr--r-- 3 root sys 1983 Dec 22 18:06 rc0
2318 -rwxr--r-- 1 root sys 2242 Dec 22 18:06 rc1
2319 -rwxr--r-- 1 root sys 2536 Dec 22 18:06 rc2
2320 -rwxr--r-- 1 root sys 2567 Dec 22 18:06 rc3
2317 -rwxr--r-- 3 root sys 1983 Dec 22 18:06 rc5
2317 -rwxr--r-- 3 root sys 1983 Dec 22 18:06 rc6
2321 -rwxr--r-- 1 root sys 5125 Dec 22 18:06 rcs
```

The Solaris OS provides the same series of rc scripts in the /etc directory for backward compatibility. These scripts are symbolic link files to the rc scripts in the /sbin directory. The following example shows this connection:

```
ls -l /etc/rc?
ls -l rc?
lrwxrwxrwx 1 root root 11 Oct 12 17:15 rc0 -> ../sbin/rc0
lrwxrwxrwx 1 root root 11 Oct 12 17:15 rc1 -> ../sbin/rc1
lrwxrwxrwx 1 root root 11 Oct 12 17:15 rc2 -> ../sbin/rc2
lrwxrwxrwx 1 root root 11 Oct 12 17:15 rc3 -> ../sbin/rc3
lrwxrwxrwx 1 root root 11 Oct 12 17:15 rc5 -> ../sbin/rc5
lrwxrwxrwx 1 root root 11 Oct 12 17:15 rc6 -> ../sbin/rc6
lrwxrwxrwx 1 root root 11 Oct 12 17:15 rcs -> ../sbin/rcs
rcm:
total 2
drwxr-xr-x 2 root sys 512 Oct 12 17:18 scripts
```

## Controlling Legacy Boot Processes

Table 11-4 summarizes the function performed by each /sbin scripts. These scripts execute legacy rc scripts.

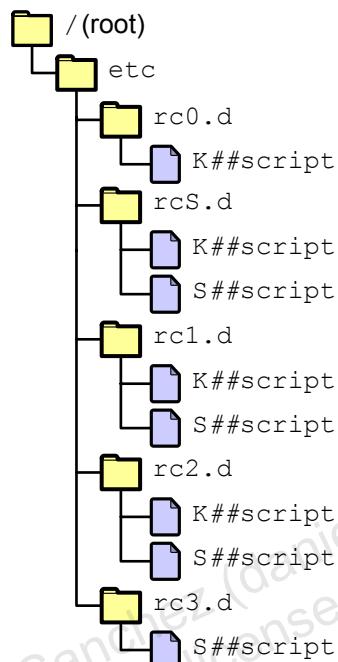
**Table 11-4** Legacy Run Control (rc) Scripts and Functions

Script	Function
/sbin/rc0	Runs the /etc/rc0.d/K* scripts and the /etc/rc0.d/S* scripts to stop system services and daemons. Start scripts should only perform fast system cleanup functions.
/sbin/rc1	Runs the /etc/rc1.d/S* scripts to perform the following tasks: <ul style="list-style-type: none"><li>• Stops system services and daemons</li><li>• Terminates certain running application processes</li><li>• Unmounts all remote file systems</li></ul>
/sbin/rc2	Runs the /etc/rc2.d/K* scripts and the /etc/rc2.d/S* scripts to start certain application daemons.
/sbin/rc3	Runs the /etc/rc3.d/K* scripts and the /etc/rc3.d/S* scripts to start certain application daemons. The K scripts are not normally present in the /etc/rc3.d directory, although if they were present, they would be run.
/sbin/rc5 /sbin/rc6	Runs the /etc/rc0.d/K* scripts and then the /etc/rc0.d/S* scripts to perform the following tasks: <ul style="list-style-type: none"><li>• Stops system services and daemons</li><li>• Starts scripts that should only perform fast system cleanup functions</li></ul>
/sbin/rcS	Runs the /etc/rcS.d scripts to bring up the system to run level S and establishes a minimal network.

## The /etc/rc#.d Directories

The /etc/rc#.d directories contain scripts that start and stop system processes for that run level.

Figure 11-4 shows an example of /etc/rc#.d directories.



**Figure 11-4** /etc/rc#.d directories

The /etc/rc#.d directories, for example /etc/rc2.d, contain scripts to start and stop processes for run level 2. The following output shows a partial script list.

```
ls -l /etc/rc2.d
total 130
-rwxr--r-- 6 root sys 324 Dec 10 11:26 K03samba
-rwxr--r-- 5 root sys 824 May 27 2004 K05appserv
(some output omitted)
-rwxr--r-- 5 root sys 2329 Dec 11 08:52 S94ncalodg
-rwxr--r-- 2 root sys 733 Dec 11 08:54 S98deallocate
-rwxr--r-- 5 root sys 1023 Dec 11 08:53 S99audit
-rwxr--r-- 5 root sys 2804 Dec 7 20:52 S99dtlogin
```

## Start Run Control Scripts

The `/etc/rc#.d` start scripts are run in the sort order shown by `ls`. The files beginning with S are run to start a system process. These scripts are called by the appropriate `/sbin/rc#` and this script passes the argument `start` to them if their names do not end in `.sh`. There are no arguments passed to `.sh` scripts. These files have names in the form of `S##name-of-script`. For example, the `S99dtlogin` script starts the graphical login process.

## Stop Run Control Scripts

The `/etc/rc#.d` stop scripts (also referred to as the kill scripts) are always run in the sort order shown by the `ls` command. The files that begin with K are run to stop or kill a system process. These scripts are called by the appropriate `/sbin/rc#`, and this script passes the argument "stop" to them if their names do not end in `.sh`.

These files have names in the form of:

`K##name-of-script`

For example, the `K03samba` script stops the Samba server.

---

**Note** – File names that begin with a lowercase k or s are ignored, and are not executed. To disable a script, rename it with the appropriate lowercase letter.

---



## The /etc/init.d Directory

Run control scripts are located in the /etc/init.d directory.

The run control script /etc/init.d/samba is hard-linked to the corresponding run control script /etc/rc3.d/S90samba, as shown by the ls commands:

```
cd /etc/init.d
ls -i samba
4715 samba

cd /etc/rc3.d
ls -i S90samba
4715 samba
```

You can stop a process or start a process without changing the system run level.

For example, to stop and restart the samba file and print sharing service, run the following command with a start or stop argument:

```
/etc/init.d/samba start
/etc/init.d/samba stop
```

### Adding Scripts to Start and Stop Legacy Services

To add run control scripts to start and stop a service not managed by SMF, create the script in the /etc/init.d directory and create links in the appropriate /etc/rc#.d directory for the run level in which the service is to be started and stopped.

The following procedure describes how to add a run control script:

1. Create the script in the /etc/init.d directory.

```
vi /etc/init.d/filename
chmod 744 /etc/init.d/filename
chgrp sys /etc/init.d/filename
```

2. Create links to the appropriate /etc/init.d directory.

```
cd /etc/init.d
ln filename /etc/rc#.d/S##filename
ln filename /etc/rc#.d/K##filename
```

For instance, you might link a file in /etc/init.d called myapplication to a file called /etc/rc3.d/S99myapplication, with a corresponding stop script called /etc/rc3.d/K01myapplication.

3. Use the ls command to verify that the script has links in the appropriate directories.

```
ls -li /etc/init.d/filename
ls -li /etc/rc#.d/S##filename
ls -li /etc/rc#.d/K##filename
```

4. Test the *filename* by performing the following commands:

```
/etc/init.d/filename start
```

# Performing System Shutdown Procedures

Shut down the Solaris OS as follows:

- To perform administration or maintenance
- If you are anticipating a power outage
- If you need to move the system to a new location.

The Solaris OS requires a clean and orderly shutdown. A “clean and orderly” shutdown does the following:

- Stops processes
- Writes data in memory to disks
- Unmounts file systems.

The work you need to do while the system is shut down will determine how you shut the system down and which command you use to do so. The following describes the system shut-down types:

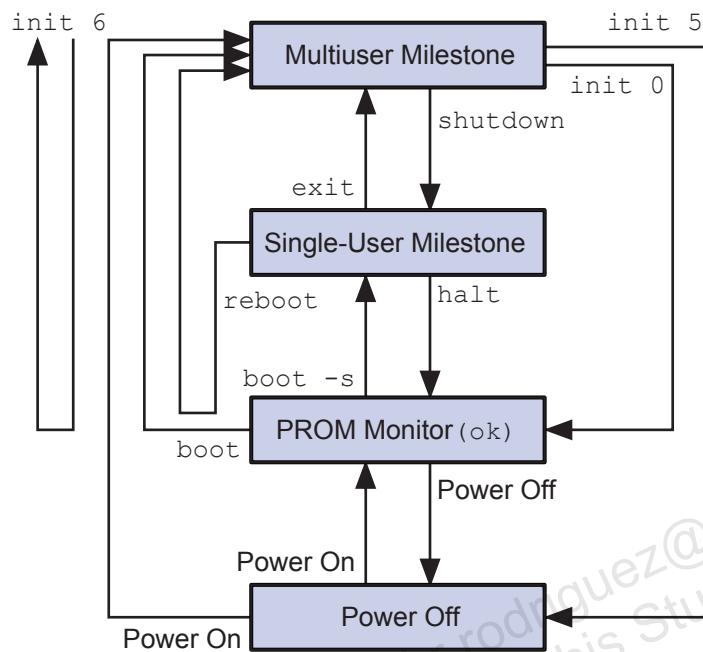
- Shut down the system to single-user mode
- Shut down the system to stop the Solaris OS, and display the ok or Press any key to reboot prompt
- Shut down the system and turn off power
- Shut down the system and automatically reboot to multiuser mode

The commands available to you, as the root user, for performing system shutdowns include:

- `/sbin/init` (using run levels S, 0, 1, 5, or 6)  
The `init` command accepts more arguments than listed here. Those arguments are not listed here because they are not part of system shutdown procedures.
- `/usr/sbin/shutdown` (using run levels S, 0, 1, 5, or 6)



Figure 11-5 shows the run-level transitions that occur during system boot or shut down on a SPARC-based system.



**Figure 11-5** Run-Level Transitions

On x86/x64-based systems, init 0, halt, and the shutdown command with the -i 0 option, shut down the Solaris OS, and leave the system at the Press any key to reboot prompt, rather than the ok prompt. Pressing any key causes a reset of the x86/x64 system. Without user intervention, x86/x64 systems boot to the default milestone when the boot process is initiated in this manner.

The init 6, init 5, and reboot commands on x86/x64 systems function as they do on SPARC-based systems.

Powering-on x86/x64 systems initiates a boot process to the default run level from the default source, and as such is equivalent to the boot command on SPARC systems. Similarly, selecting a boot option from the GRUB main menu and typing b boots the system using the source you selected.

Editing the kernel command in the command list for a GRUB boot option allows you to specify the -s option and boot an x86/x64 system to the single-user milestone.

## The /usr/sbin/init Command

Use the `init` command to shut down, power off, or reboot a system in a clean and orderly manner. The `init` command informs the `svc.startd` daemon of the runlevel change. The `init` command doesn't warn logged-in users that the system is being shut down, and there is no grace period.

Command	Execution
# <code>init 1</code>	Shuts down the system to single-user mode. Use run level S or 1
# <code>init 0</code>	Shuts down the system, stops the Solaris OS, and displays the ok or Press any key to reboot prompt.
# <code>init 5</code>	Shuts down the system and turns the system power off
# <code>init 6</code>	Shut down the system and reboots to multiuser mode

## The /usr/sbin/shutdown Command

The `shutdown` command is a script that invokes the `init` daemon to shut down, power off, or reboot the system. It executes the `rc0` kill scripts to shut down processes and applications gracefully. But unlike the `init` command, the `shutdown` command does the following:

- Notifies all logged-in users that the system is being shut down
- Delays the shutdown for 60 seconds by default
- Enables you to include an optional descriptive message to inform your users of what will transpire

The command format for the `shutdown` command is:

```
shutdown -y -g grace-period -i init-state
optional message
```

The `-y` option pre-answers the final shutdown confirmation question so that the command runs without your intervention.

## Performing System Shutdown Procedures

---



The `-g grace-period` allows you to change the number of seconds from the 60-second default.

The `-i init-state` specifies the run level that the system is to attain. By default, system state S is used.

---

**Note** – If the shutdown command displays the error message: “shutdown: ‘i’ – unknown flag,” it indicates that the shell has located and executed the /usr/ucb/shutdown command. Re-issue the command using its full path (for example, /usr/sbin/shutdown), or set the PATH variable to ensure /usr/sbin comes before /usr/ucb.

---

To shut down the system to single-user mode, enter the shutdown command without options.

```
shutdown
```

To shut down the system to stop the Solaris OS, and display the ok or Press any key to reboot prompt, enter:

```
shutdown -i0
```

To shut down the system and turn off its power automatically, enter:

```
shutdown -i5
```

To shut down the system and reboot to multiuser mode, enter:

```
shutdown -i6
```

You can use the `-i` option with other command options. For example, to shut down the system and reboot to multiuser mode, answer yes to the questions presented, provide a grace period of two minutes, and provide a message to the users by entering:

```
shutdown -y -g120 -i6 "The system is being rebooted"
```

## “Ungraceful” Shutdown Commands

The following commands perform an immediate system shutdown. They do not execute the rc0 kill scripts. They do not notify logged-in users, and there is no grace period.

```
halt
poweroff
reboot
```



**Caution** – These commands should be used with extreme caution and only when there is no alternative.

### Setting the Default Boot-Time Milestone

Use the svcadm command to control the milestone the svc.startd daemon meets on boot. The default milestone, if one is not specified, is “all” which is an abstract milestone where all system services are started.

To ensure that the svc.startd daemon meets the requirements of the multi-user-server milestone on the next reboot, use the following command before rebooting:

```
svcadm -v milestone -d multi-user-server:default
```

Valid options for default boot level using the svcadm command include the following:

- all
- none
- svc:/milestone/single-user:default
- svc:/milestone/multi-user:default
- svc:/milestone/multi-user-server:default

## Notes:

---

## Module 12

---

# Service Management Facility

---

## Objectives

Upon completion of this module, you should be able to:

- Describe Service Management Facility (SMF) features
- Compare run levels and SMF milestones
- Use SMF administrative commands

# Introduction to SMF

Prior to Solaris 10, Sun systems used a series of scripts to stop and start processes associated with run levels. Each run level had an associated script located in the /sbin directory, with some scripts hard-linked to each other. The init daemon was responsible for starting and stopping system services. These structures are still present. See the “Performing Legacy Boot and Shutdown Procedures” module for more information.

The SMF was introduced in Solaris 10. It delivers a unified Solaris service configuration infrastructure that accurately models Solaris services and their interaction with Solaris and other services. SMF starts services in parallel, according to dependencies, which allows the system to boot faster, and reduces dependency conflicts.

The SMF provides a centralized configuration structure for managing system services and the interaction of a service with other services. The SMF includes the following:

- A mechanism to establish and formalize dependency relationships between services.
- Information on procedures to start, stop, and restart services.
- A centralized repository for information on startup behavior and service status.
- A structured mechanism for Fault Management of system services.
- Detailed information about misconfigured services such as an explanation of why a service is not running.
- Individual log files for each service.

## Features

The SMF infrastructure includes the following:

- Service configuration repository
- Process re-starter
- Administrative CLI utilities
- Supporting kernel functionality

This infrastructure allows Solaris services to express:

- Restart requirements for the presence of prerequisite services and system facilities (such as networking)
- Requirements for identity and privileges for various tasks
- Configuration setting per service instance

Solaris services are modeled by describing them in terms of an SMF schema and associated service methods. For existing services converted to SMF services, compatibility or conversion of legacy configuration files is handled on a service-by-service basis. Once service descriptions are bootstrapped into SMF, instances of such services can be created, started, stopped, and status collected by the infrastructure. This saves time and system administration effort.

## Service Definition

A *service* provides a resource or list of capabilities to applications and other services, both local and remote. A service isn't a running process, such as a web server. A service can be the software state of a device, such as a configured network device, or a mounted file system.

A system can have more than one occurrence of a service running. For example, a system can have more than one configured network interface, or more than one mounted file system.

## Service Identifiers

In the SMF, each service instance has a *service identifier*. The service identifier is in the form of a fault management resource identifier (FMRI). The FMRI indicates the service type or category, name, and instance. The service categories include the following:

- application
- device
- legacy
- milestone
- network
- platform
- site
- system

An example of a service instance FMRI is:

```
svc:/system/filesystem/root:default
```

Where:

- The prefix `svc` indicates that this service is managed by SMF
- The category of the service is `system`
- The service itself is a `filesystem`
- The instance of the service is the `root` file system
- The word `default` identifies the first, in this case only, instance of the service

Another example of an FMRI for a service is:

```
lrc:/etc/rc3_d/S90samba
```

Where:

- The prefix, `legacy run control (lrc)`, indicates that this service isn't currently managed by the SMF
- The pathname `/etc/rc3_d` refers to the directory `/etc/rc3.d` where there is a script used to manage this service
- The name of the script is `S90samba`

## Listing Service Information

List service instance names and the service state of with the svcs command.

#	svcs	STATE	STIME	FMRI
	legacy_run	Feb_10		lrc:/etc/rc2_d/S10lu
	legacy_run	Feb_10		lrc:/etc/rc2_d/S20sysetup
	legacy_run	Feb_10		lrc:/etc/rc2_d/S90wbem
	legacy_run	Feb_10		lrc:/etc/rc2_d/S99dtlogin
	legacy_run	Feb_10		lrc:/etc/rc3_d/S81volmgt
	(output omitted)			
	online	Feb_10		svc:/system/system-log:default
	online	Feb_10		svc:/system/fmd:default
	online	Feb_10		svc:/system/console-login:default
	online	Feb_10		svc:/network/smtp:sendmail
	online	Feb_10		svc:/milestone/multi-user:default
	online	Feb_10		svc:/milestone/multi-user-server:default
	online	Feb_10		svc:/system/zones:default
	offline	Feb_10		svc:/application/print/ipp-listener:default
	offline	Feb_10		svc:/application/print/rfc1179:default
	maintenance	10:24:15		svc:/network/rpc/spray:default

## Service States

Use the svcs command to list service identifiers and the service instance state. You can enable or disable a service. Service states include the following:

- online  
The service instance is enabled and has successfully started.
- offline  
The service instance is enabled, but the service is not yet running or available to run.
- disabled  
The service instance is not enabled and is not running.
- legacy\_run  
The legacy service is not managed by SMF, but the service can be observed. This state is only used by legacy services.

- uninitialized

This state is the initial state for all services before their configuration has been read.

- maintenance

The service instance has encountered an error that must be resolved by the administrator.

- degraded

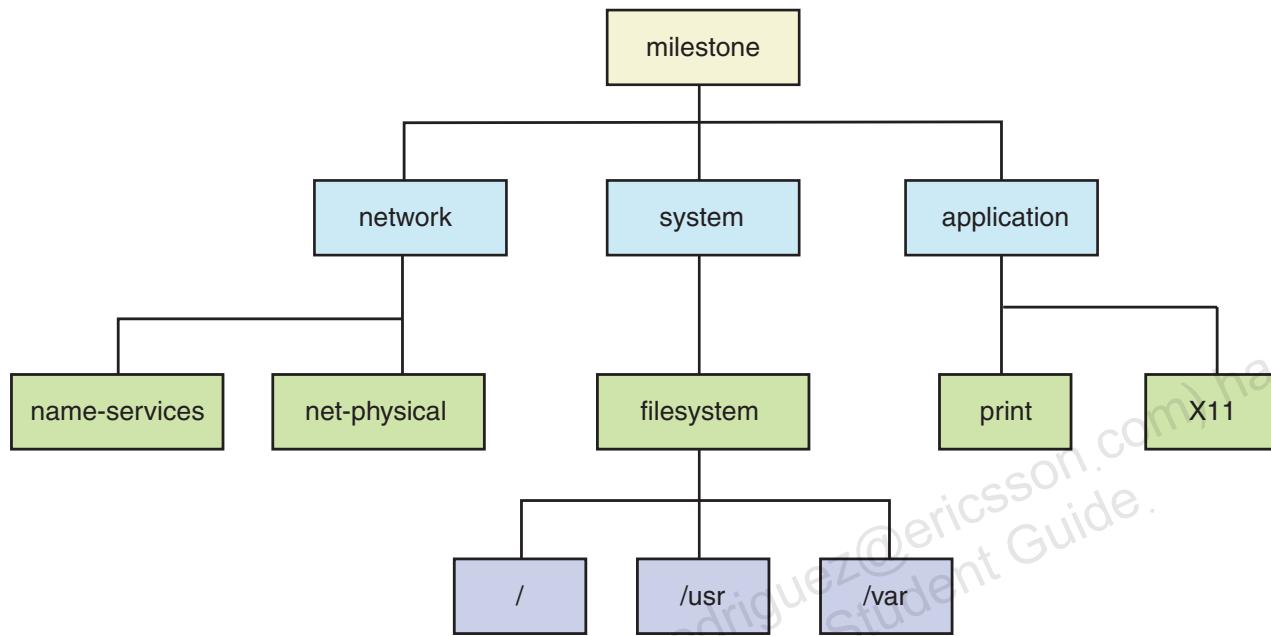
The service instance is enabled, but is running at a limited capacity.

## Milestones

A milestone service includes a defined set of other services. Regard a milestone as a system state to reach. A *milestone system state* requires a defined set of services to be running. These services depend on other services being available. Hence, there is a hierarchy of dependency relationships. The SMF manages this core feature. There are seven milestones.

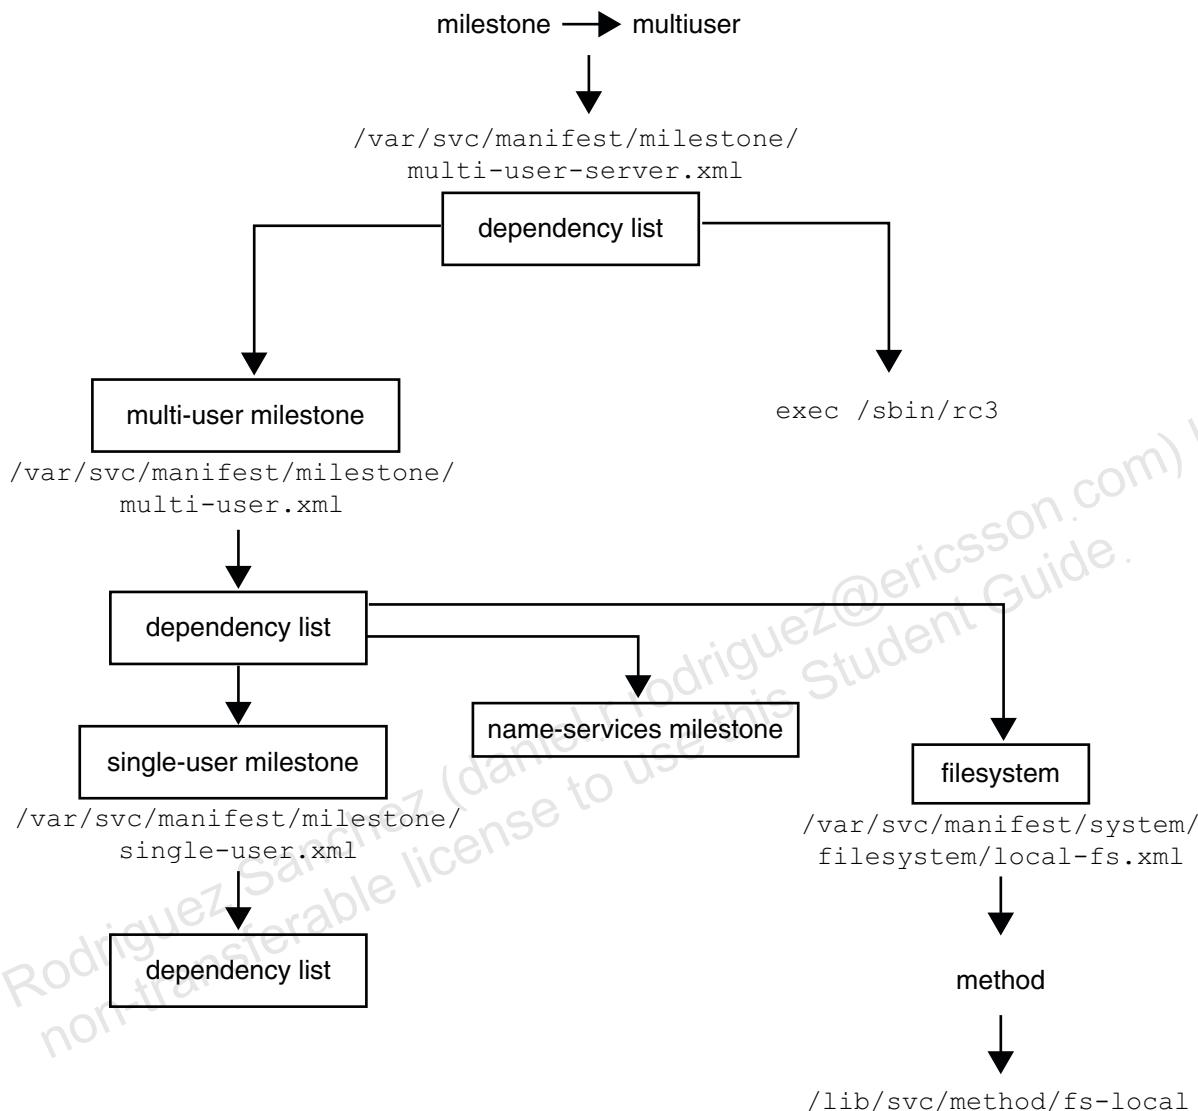
- single-user
- multi-user
- multi-user-server
- network
- name-services
- sysconfig
- devices

Figure 12-1 shows the relationship between a milestone and services.



**Figure 12-1** SMF Milestone and Services

Figure 12-2 shows an example of the dependency relationships.



**Figure 12-2** SMF Dependency Relationships

## The svc.startd Daemon

The svc.startd daemon maintains system services and ensures that the system boots to the appropriate milestone. If you don't specify a milestone at boot up, svc.startd boots to the built-in milestone "all" which includes all system services. Use the following milestones at boot time:

- none
- single-user
- multi-user
- multi-user-server
- all

To boot a system to a specific milestone, the -m option is passed to the boot command from OBP.

```
ok> boot -m milestone=single-user
```

The svc.startd daemon could also be called the master restarter daemon because it is responsible for ensuring the correct running, starting, and restarting of system services. The svc.startd daemon can obtain information about services from the repository. This init process did this before.

The svc.startd daemon delegates responsibility for services to other delegated restarter daemons, for example, the inetd daemon.

The svc.startd daemon starts the appropriate processes for the achieved run level. The svc.startd daemon uses information in the repository to determine the required milestone and then starts to process the manifests located in the /var/svc/manifest directory.

/var/svc/manifest/milestone contains XML files which describe the dependencies for this milestone. A milestone includes multiple SMF services.

The following files are in the /var/svc/manifest/milestone directory:

- single-user.xml
- multi-user.xml
- multi-user-server.xml
- network.xml
- name-services.xml
- sysconfig.xml

These.xml files may refer to other directories that contain commands to run, such as:

- /sbin/rc2
- /lib/svc/method/fs-local

These.xml files may refer to other.xml files in subdirectories below /var/svc/manifest that contain commands to run, such as:

- /sbin/rc2
- /lib/svc/method/fs-local

## The inetd Restarter

inetd is the delegated restarter for SMF. It manages service states in response to administrative requests, system failures, and service failures. When appropriate, inetd listens for network service requests.

With SMF, services are no longer managed by editing the `inetd.conf` configuration file. Instead, you use `inetconv` to convert the configuration file content into SMF format services, then you manage these services using `inetadm` and `svcadm`.

### Setting inetd Managed Services Backlog Queue Size

Server-side resource demands can fill up a backlog queue. In Solaris 10 5/08, the `SMF connection_backlog` tunable property was introduced to enable you to manage an `inetd` managed services backlog queue size. The `connection_backlog` queue size default is 10, which can easily accommodate client requests. Modify the `connection_backlog` property by using the `inetadm` command.

To list the properties, type:

```
#inetadm -l fmri/pattern
```

To change the value for a specific service, type:

```
#inetadm -m fmri/pattern connection_backlog=new value
```

To change the value globally, type:

```
#inetadm -M connection_backlog=newvalue
```

## Boot Archive Service

The boot archive service is controlled with the `svcadm` command and managed by SMF:

```
svc:/system/boot-archive:default
```

## The Service Configuration Repository

The repository database stores information about the state of each service instance. It also stores configuration information about the services and system. The repository is distributed among local memory and local disk-based files. The disk-based database is `/etc/svc/repository.db`.

The repository is managed by the `svc.configd` daemon. The `svc.configd` daemon backs up the repository before applying any changes issued by the SMF commands and utilities. These backup copies of the repository ensure that fallback is possible.

A corrupt repository prevents the system from booting. A corrupt repository can be repaired by booting the system to single-user, and running the `/lib/svc/bin/restore_repository` command.

## The Service Repository Database

A database is saved in the /etc/svc directory that contains details of the available services and their settings. Use the /lib/svc/bin/restore\_repository utility to repair or restore a corrupt repository.

To see how the repository database is used, perform the following steps:

1. **cd /lib/svc/bin**
2. **./restore\_repository**

---

**Note** – This example of the restore\_repository script output is for the FCS Solaris 10 release, newer releases of Solaris 10 have different output from this script.

---

### Repository Restore utility

See <http://sun.com/msg/SMF-8000-MY> for more information on the use of this script to restore backup copies of the smf(5) repository.

If there are any problems which need human intervention, this script will give instructions and then exit back to your shell.

Note that upon full completion of this script, the system will be rebooted using reboot(1M), which will interrupt any active services.

The following backups of /etc/svc/repository.db exist, from oldest to newest:

```
manifest_import-20050221_112255
manifest_import-20050221_144358
boot-20050223_100423
boot-20050223_211258
boot-20050224_095929
boot-20050225_134532
```

The backups are named based on their type and the time what they were taken.

Backups beginning with "boot" are made before the first change is made to the repository after system boot. Backups beginning with "manifest\_import"

are made after svc:/system/manifest-import:default finishes its processing.

The time of backup is given in YYYYMMDD\_HHMMSS format.

Please enter one of:

- 1) boot, for the most recent post-boot backup
- 2) manifest\_import, for the most recent manifest\_import backup.
- 3) a specific backup repository from the above list
- 4) -seed-, the initial starting repository. (All customizations will be lost.)
- 5) -quit-, to cancel.

Enter response [boot]: manifest\_import-20050221\_144358

After confirmation, the following steps will be taken:

svc.startd(1M) and svc.configd(1M) will be quiesced, if running.

/etc/svc/repository.db

-- renamed --> /etc/svc/repository.db\_old\_20050225\_163816

/etc/svc/repository-manifest\_import-20050221\_144358

-- copied --> /etc/svc/repository.db

and the system will be rebooted with reboot(1M).

Proceed [yes/no] ? **no**

Exiting...

#

In the above example, the alternate repository.db file naming convention is YYYYMMDD\_number, therefore the repository.db\_old\_20050225\_163816 file was saved on February 25, 2005.

# Stopping and Starting Services Using SMF Commands

## Listing Services With the svcs Command

The svcs command is used to monitor SMF services. It is useful for examining the status of services, and for following the dependency relationship between services.

```
svcs
STATE STIME FMRI
legacy_run 13:45:11 lrc:/etc/rcS_d/S29wrsmcfg
legacy_run 13:45:37 lrc:/etc/rc2_d/S10lu
legacy_run 13:45:38 lrc:/etc/rc2_d/S20sysetup
legacy_run 13:45:38 lrc:/etc/rc2_d/S401lc2
legacy_run 13:45:38 lrc:/etc/rc2_d/S42ncakmod
legacy_run 13:45:39 lrc:/etc/rc2_d/S47pppd
(output omitted)
online 13:45:36 svc:/network/smtp:sendmail
online 13:45:38 svc:/network/ssh:default
online 13:45:38 svc:/system/fmd:default
online 13:45:38 svc:/application/print/server:default
online 13:45:39 svc:/application/print/rfc1179:default
online 13:45:41 svc:/application/print/ipp-listener:default
online 13:45:45 svc:/milestone/multi-user:default
online 13:45:53 svc:/milestone/multi-user-server:default
online 13:45:54 svc:/system/zones:default
online 8:46:25 svc:/system/filesystem/local:default
online 8:46:26 svc:/network/inetd:default
online 8:46:32 svc:/network/rpc/meta:tcp
online 8:46:32 svc:/system/mdmonitor:default
online 8:46:38 svc:/milestone/multi-user:default
online 13:14:35 svc:/network/telnet:default
maintenance 8:46:21 svc:/network/rpc/keyserv:default
```

The svcs command can also be used to examine the status of a specific service instance. For example:

```
svcs svc:/system/console-login:default
STATE STIME FMRI
online 14:38:27 svc:/system/console-login:default
#
```

Use the svcs command with the -d and the -D options to examine the dependency relationships of services. The -d option shows what other services the named service depends on. The -D option shows what other services depend on the named service.

The following example shows what the service dependencies are for the filesystem/local:default service instance:

```
svcs -d svc:/system/filesystem/local:default
STATE STIME FMRI
online Mar_16 svc:/system/filesystem/minimal:default
online Mar_16 svc:/milestone/single-user:default
#
#
```

The following command is a example of changing the default milestone from milestone all to milestone multi-user-server:

```
svcs -d milestone/multi-user:default
STATE STIME FMRI
disabled Mar_16 svc:/application/print/server:default
disabled Mar_16 svc:/network/ntp:default
disabled Mar_16 svc:/system/auditd:default
disabled Mar_16 svc:/system/mdmonitor:default
disabled Mar_16 svc:/system/rcap:default
online Mar_16 svc:/milestone/name-services:default
online Mar_16 svc:/system/rmtmpfiles:default
online Mar_16 svc:/system/name-service-cache:default
online Mar_16 svc:/system/power:default
online Mar_16 svc:/milestone/single-user:default
online Mar_16 svc:/system/filesystem/local:default
online Mar_16 svc:/system/cron:default
online Mar_16 svc:/network/rpc/bind:default
online Mar_16 svc:/platform/i86pc/kdmconfig:default
online Mar_16 svc:/milestone/sysconfig:default
online Mar_16 svc:/system/utmp:default
online Mar_16 svc:/network/inetd:default
online Mar_16 svc:/network/nfs/client:default
online Mar_16 svc:/system/filesystem/autofs:default
online Mar_16 svc:/system/system-log:default
online Mar_16 svc:/system/system-log:default
online Mar_16 svc:/network/smtp:sendmail
#
#
```

The following example shows what other services depend on the system/filesystem/local service:

```
svcs -D svc:/system/filesystem/local
STATE STIME FMRI
disabled Mar_16 svc:/application/management/webmin:default
disabled Mar_16 svc:/application/print/server:default
disabled Mar_16 svc:/application/gdm2-login:default
disabled Mar_16 svc:/network/apocd/udp:default
disabled Mar_16 svc:/system/auditd:default
disabled Mar_16 svc:/system/pools/dynamic:default
disabled Mar_16 svc:/network/inetd-upgrade:default
disabled Mar_16 svc:/network/nfs/server:default
online Mar_16 svc:/system/cron:default
online Mar_16 svc:/application/opengl/ogl-select:default
online Mar_16 svc:/network/nfs/status:default
online Mar_16 svc:/system/sac:default
online Mar_16 svc:/network/inetd:default
online Mar_16 svc:/network/rpc/cde-ttddbserver:tcp
online Mar_16 svc:/network/rpc/cde-calendar-manager:default
online Mar_16 svc:/system/filesystem/volfs:default
online Mar_16 svc:/system/filesystem/autofs:default
online Mar_16 svc:/system/system-log:default
online Mar_16 svc:/system/dumpadm:default
online Mar_16 svc:/network/ssh:default
online Mar_16 svc:/application/font/fc-cache:default
online Mar_16 svc:/network/smtp:sendmail
online Mar_16 svc:/milestone/multi-user:default
online Mar_16 svc:/application/management/common-agent-
container-1:default
#
```

## Changing Service States Using the svcadm Command

Use the svcadm command to change the service state. For example, to verify the status of the cron service:

```
pgrep -fl cron
 180 /usr/sbin/cron
svcs cron
STATE STIME FMRI
online Mar_16 svc:/system/cron:default
svcadm -v disable system/cron:default
svc:/system/cron:default disabled.
svcs cron
STATE STIME FMRI
disabled 16:21:52 svc:/system/cron:default
pgrep -fl cron
svcadm -v enable system/cron:default
svc:/system/cron:default enabled.
svcs cron
STATE STIME FMRI
online 16:22:28 svc:/system/cron:default
pgrep -fl cron
 5502 /usr/sbin/cron
#
```

Use the svcadm -v disable command to disable the service. With the svcadm -v disable command, the service will remain disabled until you enable it from the command line. Disable the service temporarily, until the next reboot, by using the -t option.

```
svcadm -v disable -t system/cron:default
svc:/system/cron:default temporarily disabled.
```

## Using svcs to Determine Why Services are Not Running

Use the svcs command to troubleshoot why services are not running. In the following example, the cron service was temporarily disabled by the administrator.

```
svcs -x cron
svc:/system/cron:default (clock daemon (cron))
 State: disabled since Fri Feb 25 15:05:47 2005
 Reason: Temporarily disabled by an administrator.
 See: http://sun.com/msg/SMF-8000-1S
 See: cron(1M)
 See: crontab(1)
 See: /var/svc/log/system-cron:default.log
 Impact: This service is not running.
```

To find out more about a service, see the /var/svc/log/system-cron:default.log log file, and <http://sun.com/msg/SMF-8000-1S>.

## Creating New Service Scripts

You can create new scripts to start and stop additional processes or services to customize a system. For example, to eliminate the requirement for a data base server manual start, you can create a script to start the database server automatically after the appropriate network services start. You could then create another script to terminate this service and shut down the database server before the network services are stopped.

You must incorporate the new service into the SMF. Doing this can be complex. The required steps are as follows:

- Determine the service start and stop process.
- Establish a service name and category.
- Determine if the service runs multiple instances.
- Identify any dependency relationships between this service and any other services.
- If a script is required to start and stop the process, create the script and place it in a local directory such as /usr/local/svc/method.

- Create a service manifest file for your service. This file describes the service and any dependency relationships. Service manifests are pulled into the repository either by using the svccfg command or at boot time.
- Incorporate the script into the SMF using the svccfg utility.

Here's an example:

```
vi /usr/local/svc/method/newservice
#!/sbin/sh
#
Copyright 2004 Sun Microsystems, Inc. All rights reserved.
Use is subject to license terms.
#
ident "@(#)newservice 1.14 04/08/30 SMI"

case "$1" in
 'start')
 /usr/bin/newservice &
 ;;

 'stop')
 /usr/bin/pkill -x -u 0 newservice
 ;;
*)
 echo "Usage: $0 { start | stop }"
 ;;
esac
exit 0

chmod 544 /usr/local/svc/method/newservice

cd /var/svc/manifest/site
vi newservice.xml
<?xml version="1.0"?>
<!DOCTYPE service_bundle SYSTEM
"/usr/share/lib/xml/dtd/service_bundle.dtd.1">
<!--
Copyright 2004 Sun Microsystems, Inc. All rights reserved.
Use is subject to license terms.

ident "@(#)newservice.xml" 1.2 04/09/13 SMI"
-->

<service_bundle type='manifest' name='OPTnew:newservice'>
```

## Introduction to SMF

---

```
<service
 name='site/newservice'
 type='service'
 version='1'>

 <single_instance/>
 <dependency
 name='usr'
 type='service'
 grouping='require_all'
 restart_on='none'>
 <service_fmri value='svc:/system/filesystem/local' />
 </dependency>

 <dependent
 name='newservice'
 grouping='require_all'
 restart_on='none'>
 <service_fmri value='svc:/milestone/multi-user' />
 </dependent>

 <exec_method
 type='method'
 name='start'
 exec='/usr/local/svc/method/newservice start'
 timeout_seconds='30' />

 <exec_method
 type='method'
 name='stop'
 exec='/usr/local/svc/method/newservice stop'
 timeout_seconds='30' />

 <property_group name='startd' type='framework'>
 <propval name='duration' type='astring' value='transient' />
 </property_group>

 <instance name='default' enabled='true' />

 <stability value='Unstable' />

 <template>
 <common_name>
 <loctext xml:lang='C'>
```

```

 New service
 </loctext>
 </common_name>
</template>
</service>

</service_bundle>
```

The following describes the entries in the file:

- Standard header.

```
<?xml version="1.0"?>
<!DOCTYPE service_bundle SYSTEM
 "/usr/share/lib/xml/dtd/service_bundle.dtd.1">
```

- Comment section.

```
<!--
Copyright 2004 Sun Microsystems, Inc. All rights reserved.
Use is subject to license terms.
```

```
ident "@(#)newservice.xml 1.2 04/09/13 SMI"
-->
```

- The name of the service. The type (manifest) indicates a simple service rather than a milestone, the package providing the service, and the service name.

```
<service_bundle type='manifest' name='OPTnew:newservice'>
```

- Service category, type, name, and version.

```
<service
 name='site/newservice'
 type='service'
 version='1'>
```

- Whether multiple instances of the service will run.

```
<single_instance/>
```

- The service model to use. The entry shows that the service will be started by svc.startd. Transient services are started once and not restarted.

```
<property_group name='startd' type='framework'>
 <propval name='duration' type='astring' value='transient'
/>
</property_group>
```

- How the service is started and stopped.

```
<exec_method
 type='method'
 name='start'
 exec='/usr/local/svc/method/newservice start'
 timeout_seconds='30' />
```

```
<exec_method
 type='method'
 name='stop'
 exec='/usr/local/svc/method/newservice stop'
 timeout_seconds='30' />
```

- Define any dependencies for this service. The first entry states that the newservice requires the filesystem/local service.

```
<dependency
 name='usr'
 type='service'
 grouping='require_all'
 restart_on='none'>
 <service_fmri value='svc:/system/filesystem/local' />
</dependency>
```

- The second entry makes sure that your service is associated with the multi-user milestone and that the multi-user milestone requires this service.

```
<dependent
 name='newservice'
 grouping='require_all'
 restart_on='none'>
 <service_fmri value='svc:/milestone/multi-user' />
</dependent>
```

- Creating the instance.

```
<instance name='default' enabled='true' />
```

```
<stability value='Unstable' />
```

- Creating information to describe the service.

```
<template>
 <common_name>
 <loctext xml:lang='C'>
 New service
 </loctext>
 </common_name>
</template>
```

The new service (newservice) now needs to be imported into SMF.

This is done by running the svccfg utility:

```
svccfg import /var svc/manifest/site/newservice.xml
```

After the service has been imported into SMF it should be visible using the svcs command.

```
svcs newservice
STATE STIME FMRI
online 8:43:45 svc:/site/newservice:default
#
```

It should also be possible to manipulate the service using svcadm.

```
svcadm -v disable site/newservice
site/newservice disabled.
svcs newservice
STATE STIME FMRI
disabled 9:11:38 svc:/site/newservice:default
svcadm -v enable site/newservice
site/newservice enabled.
svcs newservice
STATE STIME FMRI
online 9:11:54 svc:/site/newservice:default
#
```

The multiuser milestone has a dependency on the newservice. This dependency was created by the dependent entry in the newservice xml file. You can observe that the newservice has started before the multiuser milestone completed.

```
svcs -d milestone/multi-user:default
STATE STIME FMRI
disabled 8:43:16 svc:/platform/sun4u/sf880drd:default
online 8:43:16 svc:/milestone/name-services:default
online 8:43:33 svc:/system/rmtpfiles:default
online 8:43:42 svc:/network/rpc/bind:default
online 8:43:46 svc:/milestone/single-user:default
online 8:43:46 svc:/system/utmp:default
online 8:43:47 svc:/system/system-log:default
online 8:43:47 svc:/system/system-log:default
online 8:43:49 svc:/system/filesystem/local:default
online 8:44:01 svc:/system/mdmonitor:default
online 9:11:54 svc:/site/newservice:default
#
```

## Manipulating Legacy Services Not Managed by SMF

If the FMRI prefix for a service is lrc, then that service is not currently managed by SMF. To start and stop a service without changing run levels, you must manually run the script associated with that service. The script to run a legacy service (non-SMF managed) is in /etc/init.d.

To identify and start or stop a legacy run service, perform the following steps:

1. Use the svcs command to display the list of legacy run services and identify the service you want to start or stop.

```
svcs | grep legacy
legacy_run Apr_05 lrc:/etc/rcS_d/S50sk98sol
legacy_run Apr_05 lrc:/etc/rc2_d/S10lu
legacy_run Apr_05 lrc:/etc/rc2_d/S20syssetup
legacy_run Apr_05 lrc:/etc/rc2_d/S40llc2
legacy_run Apr_05 lrc:/etc/rc2_d/S42ncakmod
legacy_run Apr_05 lrc:/etc/rc2_d/S47pppd
legacy_run Apr_05 lrc:/etc/rc2_d/S70uucp
legacy_run Apr_05 lrc:/etc/rc2_d/S72autoinstall
legacy_run Apr_05 lrc:/etc/rc2_d/S73cachefs_daemon
legacy_run Apr_05 lrc:/etc/rc2_d/S81dodatadm_udapl
legacy_run Apr_05 lrc:/etc/rc2_d/S89PRESERVE
```

```

legacy_run Apr_05 lrc:/etc/rc2_d/S94ncalodg
legacy_run Apr_05 lrc:/etc/rc2_d/S98deallocate
legacy_run Apr_05 lrc:/etc/rc3_d/S16boot_server
legacy_run Apr_05 lrc:/etc/rc3_d/S50apache
legacy_run Apr_05 lrc:/etc/rc3_d/S52imq
legacy_run Apr_05 lrc:/etc/rc3_d/S80mipagent
legacy_run Apr_05 lrc:/etc/rc3_d/S84appserv
legacy_run Apr_05 lrc:/etc/rc3_d/S90samba
#

```

- Verify that the script associated with the service you want to start or stop exists in the /etc/init.d directory.

```

ls /etc/init.d/apache
/etc/init.d/apache
#

```

- To start the service, run the script in /etc/init.d and specify the start argument.

```

/etc/init.d/apache start
#

```

**Note** – Some legacy services require configuration steps before you run these scripts and specify the start argument. Without a complete configuration, the processes that provide these services may not start.



- To stop the service, run the script in /etc/init.d and specify the stop argument.

```

/etc/init.d/apache stop
#

```

## Secure By Default

Traditionally, Solaris systems provided many network services by default. This open approach was convenient, but it also made it easy for remote attackers to exploit any vulnerabilities that existed in the network services software. The Secure by Default (SBD) enhancement reduces vulnerabilities by disabling as many network services as possible while still leaving a useful system. The only network-listening service left enabled in the hardened configuration is SSH.

### Installation and Secure by Default

Solaris Express (since 7/06) installs hardened by default. Initial installations of Solaris 10 OS update releases present a Secure by Default option during the installation process.

### Custom Configurations

SBD uses the SMF to control network services. (Some legacy services were converted to SMF control in support of SBD.) Some services are disabled completely, while others are configured using SMF properties to accept input only from clients on the local system. Starting from this hardened baseline, you use SMF to enable services for network listening using the svcadm(1M) and svccfg(1M) commands.

Use the netservices(1M) command to disable network services. For example, use the following command to re-establish the hardened state after enabling individual services:

```
netservices limited
```

---

**Note** – The existing /var/svc/profile/generic\_limited\_net.xml SMF profile was extended to implement this hardened state. The netservices script may become obsolete once SMF supports setting service properties in its profiles.

---

You can enable default services, as they were in previous Solaris releases, by running the following command:

```
netservices open
```



## Limited Networking Profile

Table 12-1 includes SMF service settings implemented for SBD plus the settings for existing services. It contains more than networking services (for completeness). The SSH service is enabled. New service properties were added to the first five listed services to implement the `local_only` configuration.

**Table 12-1** Limited Networking Profile Details

Service	FMRI	Property	Action Taken
rpcbind	svc:/network/rpc/bind	config/local_only	limit to local connections
syslogd	svc:/system/system-log	config/log_from_remote	limit to local connections
sendmail	svc:/network/smtp:sendmail	config/local_only	limit to local connections
smcwebserver	svc:/system/webconsole:console	options/tcp_listen	limit to local connections
wbem	svc:/application/management/wbem	options/tcp_listen	limit to local connections
X server	svc:/application/x11/x11-server	options/tcp_listen	limit to local connections
X font server	svc:/application/x11/xfs		disabled
dtlogin	svc:/application/graphical-login/cde-login	dtlogin/args	limit to local connections
ToolTalk	svc:/network/rpc/cde-ttdbserver:tcp	proto=ticotsord	limit to local connections
dtcm	svc:/network/rpc/cde-calendar-manager	proto=ticlts	limit to local connections
BSD print	svc:/application/print/rfc1179:default	bind_addr=localhost	limit to local connections
Internet print protocol	svc:/application/print/ipp-listener:default		disabled

**Table 12-1** Limited Networking Profile Details (Continued)

Service	FMRI	Property	Action Taken
SVM remote metaset	svc:/network/rpc/meta		disabled
SVM remote mediator	svc:/network/rpc/metamed		disabled
SVM remote multihost disk	svc:/network/rpc/metamh		disabled
SVM communication	svc:/network/rpc/mdcomm		disabled
rstatd	svc:/network/rpc/rstat:default		disabled
rusersd	svc:/network/rpc/rusers:default		disabled
telnetd	svc:/network/telnet:default		disabled
statd	svc:/network/nfs/status		disabled
lockd	svc:/network/nfs/nlockmgr		disabled
NFS client	svc:/network/nfs/client		disabled
NFS server	svc:/network/nfs/server		disabled
rquotad	svc:/network/nfs/rquota		disabled
NFS v4 callback daemon	svc:/network/nfs/cbd		disabled
NFS id mapping	svc:/network/nfs/mapid		disabled
ftpd	svc:/network/ftp:default		disabled
fingerd	svc:/network/finger:default		disabled
rlogind	svc:/network/login:rlogin		disabled
rshd	svc:/network/shell:default		disabled
Secure Shell	svc:/network/ssh:default		enabled

## Opening Selected Services

To establish a secure or hardened system, use the netservices limited command to lock down the system, and then selectively open the services you need.

Examine the symbolic link for /var/svc/profile to determine the current system security configuration. For example, the following command indicates that the system is open:

```
sys11# cd /var/svc/profile
sys11# ls -l generic.xml
lrwxrwxrwx 1 root root 18 Jan 3 10:48 generic.xml ->
./generic_open.xml
```

The following netservices command sets the system in secure mode again:

```
sys11# netservices limited
restarting syslogd
Jan 4 14:54:56 sys11 syslogd: going down on signal 15
restarting sendmail
restarting wbem
dtlogin needs to be restarted. Restart now? [Y] Y
restarting dtlogin
```

The symlink verifies the restricted configuration:

```
sys11# ls -l generic.xml
lrwxrwxrwx 1 root root 18 Jan 3 10:48 generic.xml ->
./generic_limited_net.xml
```

Services like telnet are now disabled as this captured session indicates:

```
instructor# telnet sys11ext
Trying 192.168.30.31...
telnet: Unable to connect to remote host: Connection refused
```

The following svcs command verifies that telnet is disabled:

```
sys11# svcs telnet
STATE STIME FMRI
disabled 14:54:53 svc:/network/telnet:default
```

Even for a limited system, SSH is still enabled as this session shows:

```
instructor# ssh root@sys11ext
The authenticity of host 'sys11ext (192.168.30.31)' can't be established.
RSA key fingerprint is c8:8b:35:00:65:0d:fd:94:96:eb:f3:95:27:87:7b:cf.
Are you sure you want to continue connecting (yes/no)? yes
Warning: Permanently added 'sys11ext,192.168.30.31' (RSA) to the list of
known hosts.
Password: password
Last login: Wed Jan 3 16:12:52 2007
Sun Microsystems Inc. SunOS 5.10 Generic January 2005
Welcome to SA300-S10_B on sys11
sys11#
```

The svcs command also verifies that sshd is accepting connections:

```
sys11# svcs svc:/network/ssh
STATE STIME FMRI
online Jan_03 svc:/network/ssh:default
```

## Removing Restrictions

After you verify that a system is closed, use the svcadm(1M) command to completely enable selected services. For example, the following command enables telnet in an otherwise hardened system. You can also enable services that support local mode.

```
svcadm enable telnet
svcs telnet
STATE STIME FMRI
online 10:09:58 svc:/network/telnet:default
```

As indicated, some services are enabled in local\_only mode as a result of executing the netservices limited command. For services that support this restrictive configuration, use the svccfg(1M) command to remove the local\_only restriction and completely open the service. For example, the following commands enable logging from remote systems for syslogd and make the change immediate:

```
svccfg -s system-log setprop config/log_from_remote =
true
svcadm refresh system-log
svcadm restart system-log
```

Use the `svcprop` command to query service property settings. The following example shows how to verify that `syslogd` will not allow logging from remote hosts:

```
sys11# svcprop system-log | grep config
...
config/log_from_remote boolean true
...
```

In the following example, the `sendmail` service is configured in `local_only` mode after the `netservices limited` command is executed.

```
sys11# svcprop sendmail | grep config
...
config/local_only boolean true
...
```

You can use the `svccfg` and `svcadm` commands to remove the restriction and make the change immediate:

```
svccfg -s sendmail setprop config/local_only = false
svcadm refresh sendmail
svcadm restart sendmail
```

Use the `svcprop` command to verify the change:

```
svcprop sendmail | grep config
...
config/local_only boolean false
...
```

There are two configuration options:

1. Enabling or disabling a service
2. Setting a `local_only` property

For example, if the `system-log` service was configured `local_only` but disabled. Use `svcadm enable` to enable it for local use only. Similarly, if you set the `local_only` property using the `svccfg` command on a disabled service, nothing will happen until you enable it with the `svcadm` command.

This example shows how you use `svccfg` to open the `xserver` service:

```
svccfg -s x11-server setprop options/tcp_listen = true
```

In some cases, you may see status information that conflicts with the service state. For example, the results from entering the following netstat command indicate that rpcbind will accept network connections when it will not. This is because the SMF is controlling the service with a local\_only property value of true.

```
netstat -f inet -P tcp -a

TCP: IPv4
 Local Address Remote Address Swind Send-Q Rwind Recv-Q
State

...
..sunrpc *.* 0 0 49152 0
LISTEN
...
..ssh *.* 0 0 49152 0
LISTEN
...
```

Most services implement the local\_only mode by not binding to system IP addresses. The rpcbind command binds to a system IP address and discards any packets that are received from remote systems. This is necessary because RPC clients on a local system expect to be able to contact rpcbind using the system IP address.

Use the svcprop command to verify the correct security configuration.

The following example illustrates that while rpcbind listens for incoming connections, it refuses connection requests coming from remote systems.

Check the status of the rpcbind service on the local host:

```
sys-05# svcprop bind |grep config/local
config/local_only boolean false
sys-05# netstat -a |grep rpc
 *.sunrpc Idle
 *.sunrpc *.* 0 0 49152 0
LISTEN
sys-05#
```

From a remote host, attempt a connection using an RPC service, in this example, rusers:

```
sys-06# rusers -l sys-05
root sys-05:pts/1 Apr 8 19:35 1
(192.168.201.1)
sys-06#
```

The connection succeeds. Set the `local_only` property for `rpcbind` on the local system to `true`, and verify that the status of `rpcbind` has not changed:

```
sys-05# svccfg -s bind setprop config/local_only = true
sys-05# svcadm refresh bind
sys-05# svcprop bind |grep config/local
config/local_only boolean true
sys-05# netstat -a |grep rpc
 *.sunrpc Idle
 *.sunrpc *.* 0 0 49152 0
LISTEN
sys-05#
```

Then test the remote connection with the same RPC service. The connection fails.

```
sys-06# rusers -l sys-05
sys-05: RPC: Rpcbind failure
sys-06#
```

# Exercise: Managing Services With SMF

In this exercise, you practice adding, using, and removing services in the Solaris Management Facility. You also demonstrate the function of the /etc/system file.

## Preparation

This exercise requires the banner-rc, banner-smf, and banner-smf.xml files, found in the /opt/ses/lab/smf directory. Your instructor will indicate if the files are located elsewhere. Refer to the lecture notes, as necessary, to perform the tasks listed.

## Task

Complete the following steps:

1. Change to the directory /opt/ses/lab/smf, and check that the files banner-rc, banner-smf, and banner-smf.xml all have execute permission.
2. If necessary, use chmod to add execute permissions to these files.
3. Create the /usr/local/svc/method directory.
4. Copy the banner-smf file to the /usr/local/svc/method directory.
5. Copy the banner-smf.xml file to the /var/svc/manifest/site directory.
6. Verify that the banner-smf service is not yet part of SMF.
7. Import the new banner-smf service into the SMF repository.
8. Verify that the banner-smf service is now part of SMF.
9. Open a console window to view the output of the banner-smf service.
10. Use the svcadm command to disable the banner-smf service and view the output in the console window.
11. Use the svcadm command to enable the banner-smf service and view the output in the console window.
12. Shut down the system to run-state 0.

- Unauthorized reproduction or distribution prohibited. Copyright© 2011, Oracle and/or its affiliates.
13. If you are using a SPARC system, boot the system with the `-m` verbose option.  
Does the startup message from `banner-smf` display?
  14. If you are using an x86/x64-based system, press any key to initiate the boot process. Then:
    - a. Enter `e` when the GRUB main menu displays.
    - b. Select the kernel command line and enter `e` to use the GRUB editor to modify the line.
    - c. Add the `-m` verbose option to the end of the kernel command line and press Return.
    - d. Enter `b` to boot the system.Does the startup message from `banner-smf` display?
  15. Log in as the `root` user and complete the following steps:
    - a. Open a terminal window.
    - b. Open a separate console window.
    - c. Remove the `banner-smf` service from the system.
  16. Copy the `banner-rc` script from the `/opt/ses/lab/smf` directory to the `/etc/init.d` directory.
  17. Change directory to `/etc/init.d` and verify that the `banner-rc` script runs with both the `start` and `stop` arguments.
  18. Change directory to `/etc/rc2.d` and create a hard link called `S22banner` that points to the same data as the `/etc/init.d/banner-rc` file.
  19. Change directory to `/etc/rcS.d` and create a hard link called `K99banner` that points to the same data as the `/etc/init.d/banner-rc` file.
  20. Reboot the system and watch for the output of the script you just installed.  
Does the startup message from `S22banner` display?.
  21. Log in as the `root` user and complete the following steps:
    - a. Open a terminal window.
    - b. Use the `init` command to change to run level S.Does the shutdown message from `K99banner` appear?

## Exercise: Managing Services With SMF

---

- Unauthorized reproduction or distribution prohibited. Copyright© 2011, Oracle and/or its affiliates.
22. Enter the password for the root user when it is requested and change to run level 3 by typing **Control-D**.  
Does the startup message from S22banner display?
  23. Log in as the root user and complete the following steps:
    - a. Open a terminal window.
    - b. Change to the directory /etc.
  24. Make a backup copy of the /etc/system file and name the backup file system.orig.
  25. If your system uses a SCSI tape device, complete the following steps:
    - a. Log in as the root user, and open a terminal window.
    - b. Use the prtconf command to list instances of the st driver currently loaded.  
How many instances are reported?

---
    - c. Edit the /etc/system file so that it includes the following line:  
forceload: drv/st
    - d. Reboot the system.
    - e. Log in as root and open a terminal window.
    - f. Again, list instances of the st driver currently loaded.  
How many instances are reported?

---
  26. Edit the /etc/system file so that it excludes the main disk driver for your system.

---

**Note** – On x86/x64 systems you must exclude both the 32-bit and 64-bit version of the driver.

---



27. Shut down the system to run level 0, and then attempt to boot it again.  
What happened?

28. If you are using a SPARC-based system, use the boot -a command to boot the system and complete the following steps:
- Supply the name of your backup file called etc/system.orig.

**Note** – There is *not* a slash before the etc directory name.



- Press Return to accept the default values for all other boot parameters.
29. If you are using a SPARC-based system, log in as the root user and complete the following steps:
- Open a terminal window.
  - Copy the /etc/system.orig file to the /etc/system file.
  - Reboot the system.
30. If you are using an x86/x64-based system, complete the following steps:
- When the GRUB main menu displays, press the space bar to cancel the timeout and automatic boot.
  - Select the Solaris failsafe boot option, and boot the system.
  - When prompted, choose to mount the Solaris instance below /a.
  - Change directory to /a/etc, and copy system.orig to system.
  - Change directory to root (/), and reboot the system.

## Exercise Summary



**Discussion** – Take a few minutes to discuss what experiences, issues, or discoveries you had during the lab exercises.

- Experiences
- Interpretations
- Conclusions
- Applications

# Exercise Solutions: Managing Services With SMF

In this exercise, you practice adding, using, and removing services in the Solaris Management Facility. You also demonstrate the function of the /etc/system file.

## Preparation

This exercise requires the banner-rc, banner-smf, and banner-smf.xml files, found in the /opt/ses/lab/smf directory. Your instructor will indicate if the files are located elsewhere. Refer to the lecture notes, as necessary, to perform the tasks listed.

## Task

Complete the following steps:

1. Change to the directory /opt/ses/lab/smf, and check that the files banner-rc, banner-smf, and banner-smf.xml all have execute permission.

```
cd /opt/ses/lab/smf
ls -l banner*
-rwxr-xr-x 1 root root 347 Feb 28 14:26 banner-rc
-rwxr-xr-x 1 root root 314 Feb 28 14:29 banner-smf
-rwxr-xr-x 1 root root 989 Feb 28 14:30 banner-smf.xml
#
```

2. If necessary, use chmod to add execute permissions to these files.

```
chmod a+x banner*
```

3. Create the /usr/local/svc/method directory.

```
mkdir -p /usr/local/svc/method
```

4. Copy the banner-smf file to the /usr/local/svc/method directory.

```
cp banner-smf /usr/local/svc/method
```

5. Copy the banner-smf.xml file to the /var/svc/manifest/site directory.

```
cp banner-smf.xml /var/svc/manifest/site
```

6. Verify that the banner-smf service is not yet part of SMF.

```
svcs site/banner-smf
svcs: Pattern 'site/banner-smf' doesn't match any instances
STATE STIME FMRI
#
```

7. Import the new banner-smf service into the SMF repository.

```
svccfg import /var svc/manifest/site/banner-smf.xml
#
```

8. Verify that the banner-smf service is now part of SMF.

```
svcs site/banner-smf
STATE STIME FMRI
online 9:11:54 svc:/site/banner-smf:default
#
```

9. Open a console window to view the output of the banner-smf service.

```
dtterm -C &
```

Or:

```
xterm -C &
```

10. Use the svcadm command to disable the banner-smf service and view the output in the console window.

```
svcadm -v disable site/banner-smf
```

svc:/site/banner-smf:default disabled.

```
svcs site/banner-smf
```

STATE STIME FMRI

disabled 9:11:54 svc:/site/banner-smf:default

11. Use the svcadm command to enable the banner-smf service and view the output in the console window.

```
svcadm -v enable site/banner-smf
```

svc:/site/banner-smf:default enabled.

```
svcs site/banner-smf
```

STATE STIME FMRI

online 9:11:54 svc:/site/banner-smf:default

12. Shut down the system to run-state 0.

```
init 0
```

SPARC-based systems display the ok prompt, x86/x64-based systems display the Press any key to reboot prompt.

13. If you are using a SPARC system, boot the system with the **-m verbose** option. Use the following command:

```
ok boot -m verbose
```

Does the startup message from banner-smf display?

*Yes.*

14. If you are using an x86/x64-based system, press any key to initiate the boot process. Then:

- Enter **e** when the GRUB main menu displays.
- Select the kernel command line and enter **e** to use the GRUB editor to modify the line.
- Add the **-m verbose** option to the end of the kernel command line and press Return. For example:

```
grub edit> kernel /platform/i86pc/multiboot -m verbose <Return>
```

- Enter **b** to boot the system.

Does the startup message from banner-smf display?

*Yes.*

15. Log in as the root user and complete the following steps:

- Open a terminal window.
- Open a separate console window.

```
dtterm -C &
```

Or:

```
xterm -C &
```

- Remove the banner-smf service from the system.

```
svcadm -v disable site/banner-smf
svc:/site/banner-smf:default disabled
svccfg delete -f svc:/site/banner-smf:default
svcs site/banner-smf
svcs: Pattern 'site/banner-smf' doesn't match any instances
STATE STIME FMRI
rm /var/svc/manifest/site/banner-smf.xml
```

16. Copy the banner-rc script from the /opt/ses/lab/smf directory to the /etc/init.d directory.

```
cp /opt/ses/lab/smf/banner-rc /etc/init.d
```

17. Change directory to /etc/init.d and verify that the banner-rc script runs with both the start and stop arguments.

```
cd /etc/init.d
./banner-rc start
./banner-rc stop
```

18. Change directory to /etc/rc2.d and create a hard link called S22banner that points to the same data as the /etc/init.d/banner-rc file.

```
cd /etc/rc2.d
ln /etc/init.d/banner-rc S22banner
```

19. Change directory to /etc/rcS.d and create a hard link called K99banner that points to the same data as the /etc/init.d/banner-rc file.

```
cd /etc/rcS.d
ln /etc/init.d/banner-rc K99banner
```

20. Reboot the system and watch for the output of the script you just installed.

```
init 6
```

Does the startup message from S22banner display?

Yes.

21. Log in as the root user and complete the following steps:

a. Open a terminal window.

b. Use the init command to change to run level S.

```
init s
```

Does the shutdown message from K99banner appear?

Yes.

22. Enter the password for the root user when it is requested and change to run level 3 by typing **Control-D**.

```
Control-d
```

Does the startup message from S22banner display?

Yes.

23. Log in as the root user and complete the following steps:

a. Open a terminal window.

b. Change to the directory /etc.

```
cd /etc
```

24. Make a backup copy of the /etc/system file and name the backup file system.orig.

```
cp system system.orig
```

25. If your system uses a SCSI tape device, complete the following steps:

- Log in as the root user, and open a terminal window.
- Use the prtconf command to list instances of the st driver currently loaded.

```
prtconf | grep "st, instance"
```

How many instances are reported?

*None*

- Edit the /etc/system file so that it includes the following line:  
forceload: drv/st
- Reboot the system.

```
init 6
```

- Log in as root and open a terminal window.
- Again, list instances of the st driver currently loaded.

```
prtconf | grep "st, instance"
```

How many instances are reported?

*The number varies depending on how many SCSI controllers are present. You should see Instances 0 through 6 for a system with one controller.*

26. Edit the /etc/system file so that it excludes the main disk driver for your system.

- On SPARC systems using SCSI disks, add the following line:

*exclude: drv/sd*

- On SPARC systems using IDE disks, add the following line:

*exclude: drv/dad*

- On x86/x64 systems using SATA disks, add the following two lines:

*exclude: drv/cmdk*

*exclude: drv/amd64/cmdk*

If you are uncertain of the type of disk or disk driver used for your boot disk, use the following procedure to determine the driver in use:

- a. Use the `df` command to display information for the root (/) file system.

```
df -h /
Filesystem size used avail capacity Mounted on
/dev/dsk/c1d0s0 4.8G 228M 4.5G 5% /
#
```

- b. Use the `ls -lL` command to display physical device file information for the root device.

```
ls -lL /dev/dsk/c1d0s0
brw-r---- 1 root sys 102, 0 Mar 1 11:31 /dev/dsk/c1d0s0
#
```

- c. Search in `/etc/name_to_major` for the major device number listed in the previous command. This identifies the driver in use, in this example, `cmdk`. Specify this driver in the `exclude` directive that you create in `/etc/system`.

```
grep 102 /etc/name_to_major
cmdk 102
#
```

27. Shut down the system to run level 0, and then attempt to boot it again.

```
shutdown -y -i0 -g0
(shutdown messages)
```

On a SPARC system:

ok **boot**

On an x86/x64 system:

Press any key to reboot <**press any key**>

What happened?

*The system is unable to boot. Excluding this driver prevents you from using the boot disk so long as you use the same /etc/system file. You must boot using the -a option to be able to supply an alternative file for the /etc/system file, or boot from the failsafe boot archive on x86/x64 systems to correct the problem.*



28. If you are using a SPARC-based system, use the `boot -a` command to boot the system and complete the following steps:
- Supply the name of your backup file called `etc/system.orig`.

---

**Note** – There is *not* a slash before the `etc` directory name.

---

- Press Return to accept the default values for all other boot parameters. For example:

ok **boot -a**

```
Name of system file [/etc/system]: etc/system.orig
root filesystem type [ufs]: <Return>
Enter physical name of root device [/...]: <Return>
```

29. If you are using a SPARC-based system, log in as the `root` user and complete the following steps:

- Open a terminal window.
- Copy the `/etc/system.orig` file to the `/etc/system` file.
- Reboot the system.

```
cd /etc
cp system.orig system
init 6
```

30. If you are using an x86/x64-based system, complete the following steps:

- When the GRUB main menu displays, press the space bar to cancel the timeout and automatic boot.
- Select the Solaris failsafe boot option, and boot the system.
- When prompted, choose to mount the Solaris instance below `/a`.
- Change directory to `/a/etc`, and copy `system.orig` to `system`.

```
cd /a/etc
cp system.orig system
#
```

- Change directory to root (`/`), and reboot the system.

```
cd /
init 6
(output omitted)
```

## Notes:

## Module 13

---

# Performing User Administration

---

## Objectives

Upon completion of this module, you should be able to:

- Describe user administration fundamentals
- Manage user accounts
- Manage initialization files

# Introducing User Administration

An important system administration task is setting up user accounts for each user who requires system access. Each user needs a unique account name, a user identification (UID) number, a home directory, and a login shell. You also have to determine which groups a user may access.

## Main Components of a User Account

The following is a list of the main components of a user account:

- User name – A unique name that a user enters to log in to a system. The user name is also called the login name.
- Password – A combination of up to 256 letters, numbers, or special characters that a user enters with the login name to gain access to a system.

---

**Note** – To enable 256 character passwords, the encryption policy in /etc/security/policy.conf needs to be changed to either md5 or blowfish. The line that reads: CRYPT\_DEFAULT=\_unix\_ needs to be changed to CRYPT\_DEFAULT=2a (blowfish). The values 1, 2a, and md5 are explained in crypt.conf.

---

- UID number – A user account's unique numerical identification within the system.
- Group identification (GID) number – A unique numerical identification of the group to which the user belongs.

---

**Note** – You can add a user to predefined groups listed in the /etc/group file.

---

- Comment – Information that identifies the user. A comment generally contains the full name of the user and optional information, such as a phone number or a location.
- User's home directory – A directory into which the user is placed after login. The directory is provided to the user to store and create files.
- User's login shell – The user's work environment is set up by the initialization files that are defined by the user's login shell.

## System Files That Store User Account Information

The Solaris 10 OS stores user account and group entry information in the following system files:

- /etc/passwd
- /etc/shadow
- /etc/group

Authorized system users have login account entries in the /etc/passwd file.

The /etc/shadow file is a separate file that contains the encrypted passwords. To further control user passwords, you can enforce password aging. This information is also maintained in the /etc/shadow file.

The /etc/group file defines the default system group entries. You use this file to create new group entries or modify existing group entries on the system.

### The /etc/passwd File

Due to the critical nature of the /etc/passwd file, you should refrain from editing this file directly. Instead, you should use the Solaris™ Management Console or command-line tools to maintain the file.

The following is an example of an /etc/passwd file that contains the default system account entries.

```
root:x:0:0:Super-User:/sbin/sh
daemon:x:1:1:::
bin:x:2:2::/usr/bin:
sys:x:3:3:::
adm:x:4:4:Admin:/var/adm:
lp:x:71:8:Line Printer Admin:/usr/spool/lp:
uucp:x:5:5:uucp Admin:/usr/lib/uucp:
nuucp:x:9:9:uucp Admin:/var/spool/uucppublic:/usr/lib/uucp/uucico
smmsp:x:25:25:SendMail Message Submission Program:/:
listen:x:37:4:Network Admin:/usr/net/nls:
gdm:x:50:50:GDM Reserved UID:/:
webservd:x:80:80:WebServer Reserved UID:/:
nobody:x:60001:60001:NFS Anonymous Access User:/:
noaccess:x:60002:60002>No Access User:/:
nobody4:x:65534:65534:SunOS 4.x NFS Anonymous Access User:/:
```

Each entry in the /etc/passwd file contains seven fields. A colon separates each field. The following is the format for an entry:

*loginID:x:UID:GID:comment:home\_directory:login\_shell*

Figure 13-1 defines the requirements for each of the seven fields.

**Figure 13-1** Fields in the /etc/passwd File

Field	Description
<i>loginID</i>	<p>Represents the user's login name. It should be unique to each user. The field should contain a string of no more than eight letters (A-Z, a-z) and numbers (0-9). The first character should be a letter, and at least one character should be lowercase.</p> <p><b>Note</b> – Even though some programs allow a maximum of 32 characters, as well as user names that contain periods (.), underscores (_), and hyphens (-), this practice is not recommended and might cause problems with other programs.</p>
<i>x</i>	Represents a placeholder for the user's encrypted password, which is kept in the /etc/shadow file.
<i>UID</i>	<p>Contains the UID number used by the system to identify the user. UID numbers for users range from 100 to 60000. Values 0 through 99 are reserved for system accounts. UID number 60001 is reserved for the nobody account. UID number 60002 is reserved for the noaccess account. While duplicate UID numbers are allowed, they should be avoided unless absolutely required by a program.</p> <p><b>Note</b> – The maximum value for a UID is 2147483647. However, the UIDs over 60000 do not have full utility and are incompatible with some Solaris OS features. Avoid using UIDs over 60000 so as to be compatible with earlier versions of the operating system.</p>

**Figure 13-1** Fields in the /etc/passwd File (Continued)

Field	Description
<i>GID</i>	Contains the GID number used by the system to identify the user's primary group. GID numbers for users range from 100 to 60000. (Those between 0 and 99 are reserved for system accounts.)
<i>comment</i>	Typically contains the user's full name.
<i>home_directory</i>	Contains the full path name to the user's home directory.
<i>login_shell</i>	Defines the user's login shell. There are six possible login shells in the Solaris OS: the Bourne shell, the Korn shell, the C shell, the Z shell, the BASH shell, and the TC shell.

Table 13-1 shows the default system account data for entries in the /etc/passwd file.

**Table 13-1** Default System Account Entries

User Name	User ID	Description
root	0	The root account that has access to the entire system. It has almost no restrictions and overrides all other logins, protections, and permissions.
daemon	1	The system daemon account that is associated with routine system tasks.
bin	2	The administrative daemon account that is associated with running system binary files.
sys	3	The administrative daemon account that is associated with system logging or updating files in temporary directories.
adm	4	The administrative daemon account that is associated with system logging.
lp	71	The line printer (lp) daemon account.
uucp	5	The daemon account associated with UNIX®-to-UNIX Copy Protocol (UUCP) functions.
nuucp	6	The account used by remote systems to log in to the host and start file transfers using uucp.

**Table 13-1** Default System Account Entries (Continued)

User Name	User ID	Description
smmssp	25	The sendmail message submission daemon account.
listen	37	The network listener daemon account.
gdm	50	Gnome Display Manager daemon.
webservd	80	Account reserved for WebServer access.
nobody	60001	The anonymous user account that is assigned by a Network File System (NFS) server when an unauthorized root user makes a request. The nobody user account is assigned to software processes that do not need any special permissions.
noaccess	60002	The account assigned to a user or a process that needs access to a system through some application instead of through a system login procedure.
nobody4	65534	The anonymous user account that is the SunOS™ 4.X software version of the nobody account

**Note** – The nobody account secures NFS resources. When a user is logged in as root on an NFS client and attempts to access a remote file resource, the UID number changes from 0 to the UID of nobody (60001)



### The /etc/shadow File

Due to the critical nature of /etc/shadow, don't edit it directly. Use the command-line tools. Only a root user can read the /etc/shadow file.

The following is an example /etc/shadow file that contains initial system account entries.

```
root:rJrdhjNWQQHoY:6445::::::
daemon:NP:6445::::::
bin:NP:6445::::::
sys:NP:6445::::::
adm:NP:6445::::::
lp:NP:6445::::::
uucp:NP:6445::::::
nuucp:NP:6445::::::
smmsp:NP:6445::::::
```

```

listen:*LK*::::::
gdm:*LK*::::::
webservd:*LK*::::::
nobody:*LK*:6445:::::
noaccess:*LK*:6445:::::
nobody4:*LK*:6445:::::

```

Each entry in the /etc/shadow file contains nine fields. A colon separates each field.

Following is the format of an entry:

*loginID:password:lastchg:min:max:warn:inactive:expire:flag*

Table 13-2 defines the requirements for each of the eight fields.

**Table 13-2** Fields in the /etc/shadow File

Field	Description
<i>loginID</i>	The user's login name.
<i>password</i>	A 13-character encrypted password. The string *LK* indicates a locked account, and the string NP indicates no valid password. Passwords must be constructed to meet the following requirements:  Each password must be at least six characters and contain at least two alphabetic characters and at least one numeric or special character. It cannot be the same as the login ID or the reverse of the login ID.
<i>lastchg</i>	The number of days between January 1, 1970, and the last password modification date.
<i>min</i>	The minimum number of days required between password changes.
<i>max</i>	The maximum number of days the password is valid before the user is prompted to enter a new password at login.
<i>warn</i>	The number of days the user is warned before the password expires.
<i>inactive</i>	The number of inactive days allowed for the user before the user's account is locked.

**Table 13-2** Fields in the /etc/shadow File (Continued)

<i>expire</i>	The date (given as number of days since January 1, 1970) when the user account expires. After the date is exceeded, the user can no longer log in.
<i>flag</i>	To track failed logins. The count is in low order four bits. The remainder is reserved for future use, set to zero.

### The /etc/group File

Each user belongs to a group that is referred to as the user's primary group. The GID number, located in the user's account entry within the /etc/passwd file, specifies the user's primary group.

Each user can also belong to up to 15 additional groups, known as secondary groups. In the /etc/group file, you can add users to group entries, thus establishing the user's secondary group affiliations.

---

**Note** – Line length limits of each group entry differ depending on the source. The maximum length for a line in the /etc/group file is 512 characters. The maximum length for a line in an NIS group is 1024 characters. The maximum length for a line in an NIS+ group is 4096 characters.

---

The following is an example of the default entries in an /etc/group file:

```
root::0:
other::1:root
bin::2:root,daemon
sys::3:root,bin,adm
adm::4:root,daemon
uucp::5:root
mail::6:root
tty::7:root,adm
lp::8:root,adm
nuucp::9:root
staff::10:
daemon::12:root
sysadmin::14:
smmsp::25:
gdm::50:
webservd::80:
```

```
nobody::60001:
noaccess::60002:
nogroup::65534::
```

Each line entry in the /etc/group file contains four fields. A colon character separates each field. The following is the format for an entry:

*groupname:group-password:GID:username-list*

Table 13-3 defines the requirements for each of the four fields.

**Table 13-3** Fields in the /etc/group File

Field	Description
<i>groupname</i>	Contains the name assigned to the group. Group names contain up to a maximum of eight characters.
<i>group-password</i>	Usually contains an empty field or an asterisk. This is a relic of earlier versions of UNIX.  <b>Caution</b> – A group-password is a security hole because it might allow an unauthorized user who is not a member of the group but who knows the group password, to enter the group.
	<b>Note</b> – The newgrp command changes a user's primary group association within the shell environment from which it is executed. If this new, active group has a password and the user is not a listed member in that group, the user must enter the password before the newgrp command can continue.
<i>GID</i>	Contains the group's GID number. It is unique on the local system and should be unique across the organization. Numbers 0 to 99, 60001, 60002 and 65534 are reserved for system group entries. User-defined groups range from 100 to 60000.

**Table 13-3** Fields in the /etc/group File (Continued)

Field	Description
<i>username-list</i>	<p>Contains a comma-separated list of user names that represent the user's secondary group memberships. By default, each user can belong to a maximum of 15 secondary groups.</p> <hr/> <p><b>Note</b> – The maximum number of groups is set by the kernel parameter called <code>ngroups_max</code>. You can set this parameter in the /etc/system file to allow for a maximum of 32 groups. Not all applications will be able to reference group memberships greater than 16. NFS is a notable example.</p>

### The /etc/default/passwd File

Set values for the following parameters in the /etc/default/passwd file to control properties for all users' passwords on the system:

- **MAXWEEKS** – Sets the maximum time period (in weeks) that the password is valid.
- **MINWEEKS** – Sets the minimum time period before the password can be changed.
- **PASSLENGTH** – Sets the minimum number of characters for a password. Valid entries are 6, 7, and 8.
- **WARNWEEKS** – Sets the time period prior to a password's expiration to warn the user that the password will expire.

---

**Note** – The `WARNWEEKS` value does not exist by default in the /etc/default/passwd file, but it can be added.

---



The password aging parameters `MAXWEEKS`, `MINWEEKS`, and `WARNWEEKS` are default values. If set in the /etc/shadow file, the parameters in that file override those in the /etc/default/passwd file for individual users.

The Solaris 10 OS release introduces a number of new controls for password management. These controls are configured by setting values in the /etc/default/passwd file. These controls are commented out by default.

- NAMECHECK=NO – Sets the password controls to verify that the user is not using the login name as a component of the password.
- HISTORY=0 – Forces the passwd program to log up to 26 changes to the user's password. This prevents the user from reusing the same password for 26 changes. If the HISTORY value is set to another number other than zero (0), and then set back to zero, it causes the password log for a user to be removed on the next password change.
- DICTIONLIST= – Causes the passwd program to perform dictionary word lookups from comma-separated dictionary files.
- DICTIONDBDIR=/var/passwd – The location of the dictionary where the generated dictionary databases reside. This directory must be created manually.

**Note** – To pre-build the dictionary database, refer to the man page for mkpwdict(1M).



Complexity of the password can be controlled using the following parameters:

```
#MINDIFF=3
#MINALPHA=2
#MINNONALPHA=1
#MINUPPER=0
#MINLOWER=0
#MAXREPEATS=0
#MINSPECIAL=0
#MINDIGIT=0
#WHITESPACE=YES
```

By default, all of the above parameters are commented out.

**Note** – By forcing greater complexity of password structure, you may inadvertently cause the users to write down their passwords as they may be too difficult for the user to remember. When setting a password change policy, you must not underestimate the problems that too much complexity may cause.



## Password Management

The Solaris 10 OS has new security enhancements. The pam\_unix\_auth module implements account locking for local users. Account locking is enabled by the LOCK\_AFTER\_RETRIES tunable parameter in /etc/security/policy.conf and the lock\_after-retries key in /etc/user\_attr.

The LOCK\_AFTER\_RETRIES=YES|NO parameter specifies whether a local account is locked after the number of failed login attempts for a user is equal to, or exceeds the allowed number of retries. The number of retries is defined by RETRIES in /etc/default/login.



---

**Note** – These files are discussed in greater detail in:  
*SA-202-S10; System Administration for the Solaris 10 Operating System, Part 2.*

---

The passwd command has two new options, -N and -u. The -N option creates a password entry for a non-login account. This option is useful for accounts that should not be logged in to, but must run cron jobs. The -u option unlocks a previously locked account. The passwd -N *username* command sets the password field in /etc/shadow to NP which is an unmatchable password. This effectively disables the account from logging in.

For more information, see the passwd(1) man page.

The following example shows how to prevent a user from reusing too many previous passwords.

```
vi /etc/default/passwd
(output omitted)
```

Locate the line called #HISTORY=0, and remove the comment from the beginning of the line. Modify the number to 3, so the line shows as HISTORY=3. Write and quit the file. As a regular user, log in and attempt to change your password a number of times, using different passwords and then one of the previous passwords.

```
telnet localhost
Trying 127.0.0.1...
Connected to localhost.
Escape character is '^]'.
login: testuser
```

```
 Password: 123pass
$ passwd
passwd: Changing password for testuser
Enter existing login password: 123pass
New Password: pass123
Re-enter new Password: pass123
passwd: password successfully changed for testuser
$ passwd
passwd: Changing password for testuser
Enter existing login password: pass123
New Password: 123pass
passwd: Password in history list.
Please try again
New Password: newpas1
Re-enter new Password: newpas1
passwd: password successfully changed for testuser
$
```

By uncommenting the HISTORY= line in the /etc/default/passwd file, prior password history is checked. By changing the value to 3, the number of prior password changes to keep and check when a user changes passwords is set to three.

## Setting Password Aging on a User Account

If you are using NIS+ or the /etc files to store user account information, you can set up password aging on a user's password. Starting in the Solaris 9 12/02 release, password aging is also supported in the LDAP directory service.

Password aging enables you to force users to change their passwords periodically or to prevent a user from changing a password before a specified interval. If you want to prevent an intruder from gaining undetected access to the system by using an old and inactive account, you can also set a password expiration date when the account becomes disabled. You can set password aging attributes with the passwd command.

# Managing User Accounts

Each of the following sections present two sets of command-line tools for managing user accounts: the command-line tools used in the Solaris OS versions prior to the Solaris 10 OS, and the new set of command-line tools developed for the Solaris 10 OS.

## Introducing Command-Line Tools

The Solaris OS provides these command-line tools, defined as follows:

- `useradd` – Adds a new user account on the local system
- `usermod` – Modifies a user's account on the local system
- `userdel` – Deletes a user's account from the local system
- `groupadd` – Adds a new group entry to the system
- `groupmod` – Modifies a group entry on the system
- `groupdel` – Deletes a group entry from the system

In addition to these standard command-line tools, the Solaris 9 and 10 OS has a set of command-line tools that accomplish the same tasks. They are the `smuser` and `smgroup` commands.

The `smuser` command enables you to manage one or more users on the system with the following set of subcommands:

- `add` – Adds a new user account
- `modify` – Modifies a user's account
- `delete` – Deletes a user's account
- `list` – Lists one or more user entries

The `smuser` and `smgroup` commands interact with naming services, can use autohome functionality, and are better suited for remote management.



**Note** – The `smuser` and `smgroup` commands are the command-line interface equivalent to the Solaris Management Console range of operation, and allow you to perform Solaris Management Console actions in scripts. Therefore, the `smuser` and `smgroup` commands have numerous subcommands and options designed to function across domains and multiple systems. This module describes only the basic commands.

The smgroup command enables you to manage one or more groups on the system with the following set of subcommands:

- add – Adds a new group entry
- modify – Modifies a group entry
- delete – Deletes a group entry
- list – Lists one or more group entries

Any subcommand to add, modify, list, or delete users with the smuser and smgroup commands requires authentication with the Solaris Management Console server and requires the initialization of the Solaris Management Console. For example, the following is the command format for the smuser command:

```
/usr/sadm/bin/smuser subcommand [auth_args] -- [subcommand_args]
```

The authorization arguments are all optional. However, if you do not specify the authorization argument, the system might prompt you for additional information, such as a password for authentication purposes.

The -- option separates the subcommand-specific options from the authorization arguments. The -- option must be entered even if an authorization argument is not specified because it must precede the subcommand arguments.

The subcommand arguments are quite numerous. For a complete listing of the subcommands, refer to the smuser man page. It is important to note that descriptions and other arguments that contain white space must be enclosed in double quotation marks.

## Creating a User Account

Use the useradd or smuser add command to add new user accounts to the local system. These commands add an entry for a new user into the /etc/passwd and /etc/shadow files.

These commands also automatically copy all the initialization files from the /etc/skel directory to the user's new home directory.

## The useradd Command Format and Options

The following is the command format for the useradd command:

```
useradd [-u uid][-g gid][-G gid [,gid,...]]
[-d dir][-m][-s shell][-c comment] loginname
```

Table 13-4 shows the options for the useradd command.

**Table 13-4** Options for the useradd Command

Option	Definition
<code>-u uid</code>	Sets the UID number for the new user
<code>-g gid</code>	Defines the new user's primary group
<code>-G gid</code>	Defines the new user's secondary group memberships
<code>-d dir</code>	Defines the full path name for the user's home directory
<code>-m</code>	Creates the user's home directory if it does not already exist
<code>-s shell</code>	Defines the full path name for the shell program of the user's login shell
<code>-c comment</code>	Specifies any comment, such as the user's full name and location
<code>loginname</code>	Defines the user's login name for the user account
<code>-D</code>	Displays the defaults that are applied to the useradd command

The following example uses the useradd command to create an account for a user named newuser1. It assigns 100 as the UID number, adds the user to the group other, creates a home directory in the /export/home directory, and sets /bin/ksh as the login shell for the user account.

```
useradd -u 100 -g other -d /export/home/newuser1 -m -s /bin/ksh -c
"Regular User Account" newuser1
64 blocks
#
```

The useradd command has a preset range of default values. These values can be displayed using the useradd -D command. When this command has been used for the first time, the useradd command generates a file called /var/sadm/defadduser that contains the default values. If the contents of this file are amended, the new contents become the default values for the next time the useradd command is used.

```
ls -l /usr/sadm/defadduser
/usr/sadm/defadduser: No such file or directory

useradd -D
group=other,1 project=default,3 basedir=/home
skel=/etc/skel shell=/bin/sh inactive=0
expire= auths= profiles= roles= limitpriv=
defaultpriv= lock_after_retries=
ls -l /usr/sadm/defadduser
-rw-r--r-- 1 root root 286 Oct 17 09:04
/usr/sadm/defadduser
cat /usr/sadm/defadduser
Default values for useradd. Changed Sun Oct 17 09:04:27 2004
defgroup=1
defgname=other
defparent=/home
defskel=/etc/skel
defshell=/bin/sh
definact=0
defexpire=
defauthorization=
defrole=
defprofile=
defproj=3
defprojname=default
deflimitpriv=
defdefaultpriv=
deflock_after_retries=
```

User accounts are locked by default when added with the useradd command. This can be verified by viewing the contents of the /etc/shadow file:

```
grep 'newuser1' /etc/shadow
newuser1:*LK*:12708::::::
```

By convention, a user's login name is also the user's home directory name.

Use the `passwd` command to create a password for the new account.

```
passwd newuser1
New Password: 123pass
Re-enter new Password: 123pass
passwd: password successfully changed for newuser1
```

This password setting can be verified by viewing the contents of the `/etc/shadow` file:

```
grep 'newuser1' /etc/shadow
newuser1:M0/jo1fmSbYio:12708:::::::
```

## The `smuser add` Command Format and Options

The following is the command format for the `smuser add` command:

```
smuser add [auth_args] -- [subcommand_args]
```

Table 13-5 shows some of the most common subcommand arguments for the `smuser add` command.

**Table 13-5** Subcommand Arguments for the `smuser add` Command

Subcommand Argument	Definition
<code>-c comment</code>	A short description of the login, typically the user's name. This string can be up to 256 characters.
<code>-d directory</code>	Specifies the home directory of the new user and is limited to 1024 characters.
<code>-g group</code>	Specifies the new user's primary group membership.
Subcommand Argument	Definition
<code>-G group</code>	Specifies the user's secondary group membership.
<code>-n login</code>	Specifies the user's login name.
<code>-s shell</code>	Specifies the full path name of the user's login shell.

**Table 13-5** Subcommand Arguments for the smuser add Command  
(Continued)

<b>-u uid</b>	Specifies the user ID of the user you want to add. If you do not specify this option, the system assigns the next available unique UID greater than 100.
<b>-x autohome=Y/N</b>	Sets the home directory to automount if set to Y.

The following example uses the smuser add command to create an account for a user named newuser2. It designates the login name as newuser2, assigns the UID number 500, adds the user to the group other, creates a home directory in the /export/home directory, and sets /bin/ksh as the login shell for the user account.



**Note** – The -x autohome=N option to the smuser command adds the user without automounting the user's home directory. See the man page for automount for more information.

```
/usr/sadm/bin/smuser add -- -n newuser2 -u 500 -g other -d
/export/home/newuser2 -c "Regular User Account 2" -s /bin/ksh -x
autohome=N
```

Authenticating as user: root

```
Type /? for help, pressing <enter> accepts the default denoted by []
Please enter a string value for: password :: Enter The root Password
Loading Tool: com.sun.admin.usermgr.cli.user.UserMgrCli from sys-02
Login to sys-02 as user root was successful.
Download of com.sun.admin.usermgr.cli.user.UserMgrCli from sys-02
was successful.
```

Users are added without a password by default with the smuser command. This can be verified by viewing the appropriate entry in the /etc/shadow file:

```
grep 'newuser2' /etc/shadow
newuser2::12708::::::
```

Use the passwd command to create a new password for the user.

```
passwd newuser2
New Password: 123pass
Re-enter new Password: 123pass
passwd: password successfully changed for newuser2
```

Confirm that the password change has been applied by viewing the entry for that user in the /etc/shadow file:

```
grep 'newuser2' /etc/shadow
newuser2:FSMOsxncoc6yI:12708:::::::
```

## Modifying a User Account

Use the usermod or smuser modify command to modify a user's login account on the system.

### The usermod Command Format and Options

The following is the command format for the usermod command:

```
usermod [-u uid [-o]] [-g gid] [-G gid [, gid . . .]]
[-d dir] [-m] [-s shell] [-c comment]
[-l newlogname] loginname
```

In general, the options for the usermod command function the same as those for the useradd command.

Table 13-6 shows the key options to the usermod command.

**Table 13-6** Key Options for the usermod Command

Option	Definition
-o	Allows a UID to be duplicated.
-m	Moves the user's home directory to the new location specified with the -d option.
-l newlogname	Changes a user's login name for the specified user account.
-f inactive	Sets the number of inactive days that are allowed on a user account. If the account is not logged in to for the specified number of days, it is locked.
-e expire	Sets an expiration date on the user account. Specifies the date (mm/dd/yy) on which a user can no longer log in and access the account. After that date, the account is locked.
loginname	Identifies the user's login name for the current user account.

The following example changes the login name and home directory for newuser1 to usera.

```
usermod -m -d /export/home/usera -l usera newuser1
```

### The smuser modify Command Format and Options

The following is the command format for the smuser modify command:

```
smuser modify [auth_args] -- [subcommand_args]
```

In general, the options for the smuser modify command function the same as for the smuser add command. Refer to the smuser(1M) man page for additional options.

Table 13-7 shows the options for the smuser modify command.

**Table 13-7** Options for the smuser modify Command

Option	Definition
-n <i>login</i>	Specifies the user's login name
-N <i>login</i>	Specifies the user's new login name

The following example changes the login name and home directory for newuser2 to userb.

```
/usr/sadm/bin/smuser modify -- -n newuser2 -N userb -d
/export/home/userb
Authenticating as user: root
```

```
Type /? for help, pressing <enter> accepts the default denoted by []
Please enter a string value for: password :: Enter_The_root_Password
Loading Tool: com.sun.admin.usermgr.cli.user.UserMgrCli from sys-02
Login to sys-02 as user root was successful.
Download of com.sun.admin.usermgr.cli.user.UserMgrCli from sys-02 was
successful.
```

## Deleting a User Account

Use the `userdel` command or `smuser delete` command to delete a user's login account from the system.

The following is the command format for the `userdel` command:

```
userdel -r login
```

The `userdel` command also removes the user's home directory and all of its contents if you request it to do so. Use the `-r` option to remove the user's home directory from the local file system. This directory must exist.

The following example removes the login account for a user named `usera`.

```
userdel usera
```

To request that both the user's account and home directory be removed from the system at the same time, perform the command:

```
userdel -r usera
```

---

**Note** – This command does not remove all files owned by the user, just the home directory. The system administrator should run a `find` command to locate all files owned by the user to be backed up or removed.



## The `smuser delete` Command Format and Options

The following is the command format for the `smuser delete` command:

```
smuser delete [auth_args] -- [subcommand_args]
```

The following example removes the `userb` account from the system:

```
/usr/sadm/bin/smuser delete -- -n userb
```

Authenticating as user: root

Type `/?` for help, pressing `<enter>` accepts the default denoted by `[ ]`

Please enter a string value for: password :: **Enter\_The\_root\_Password**

Loading Tool: com.sun.admin.usermgr.cli.user.UserMgrCli from sys-02

Login to sys-02 as user root was successful.

Download of com.sun.admin.usermgr.cli.user.UserMgrCli from sys-02 was successful.



**Note** – Unlike the userdel command, the smuser delete command has no -r equivalent option for deleting the home directory. The user's home directory must be deleted manually.

## Creating a Group Entry

As the root user, you create new group entries on the local system by using the groupadd or smgroup add command. These commands add an entry for the new group into the /etc/group file. Like the smuser command, the smgroup add command uses the same subcommands and authentication arguments derived from the Solaris Management Console.

### The groupadd Command Format and Options

The following is the command format for the groupadd command:

```
groupadd [-g gid [-o]] groupname
```

Table 13-8 shows the options for the groupadd command.

**Table 13-8** Options for the groupadd Command

Option	Description
-g gid	Assigns the GID number for the new group
-o	Allows the GID number to be duplicated

The following example uses the groupadd command to create the new group classroom on the local system:

```
groupadd -g 301 classroom
UX: groupadd: classroom name too long.
```



**Note** – A warning message will be written if the string exceeds MAXGLEN, which is usually set at eight characters in the /usr/include/userdefs file.

This example will work without changing the defaults:

```
groupadd -g 301 class1
```

## The smgroup add Command Format and Options

The following is the command format for the smgroup add command:

```
/usr/sadm/bin/smgroup subcommand [auth_args] -- [subcommand_args]
```

Table 13-9 shows the options for the smgroup add command.

**Table 13-9** Options for the smgroup add Command

Option	Description
<code>-g gid</code>	Specifies the GID number for the new group
<code>-m group_member</code>	Specifies the new members to add to the group
<code>-n group_name</code>	Specifies the name of the new group

The following example uses the smgroup add command to create a new group called workgroup with a GID of 123, and to add usera to the group:

```
/usr/sadm/bin/smgroup add -- -n workgroup -g 123 -m usera
```

Authenticating as user: root

```
Type /? for help, pressing <enter> accepts the default denoted by []
Please enter a string value for: password :: Enter_The_root_Password
Loading Tool: com.sun.admin.usermgr.cli.group.UserMgrGroupCli from sys-02
Login to sys-02 as user root was successful.
Download of com.sun.admin.usermgr.cli.group.UserMgrGroupCli from sys-02
was successful.
```

## Modifying a Group Entry

You can use the following commands to modify a group entry:

- The groupmod command
- The smgroup modify command

## The groupmod Command Format and Options

The following is the command format for the groupmod command:

```
groupmod [-g gid [-o]] [-n name] groupname
```

Table 13-10 defines the options for the groupmod command:

**Table 13-10 Options for the groupmod Command**

Options	Description
<code>-g gid</code>	Specifies the new GID number for the group
<code>-o</code>	Allows the GID number to be duplicated
<code>-n name</code>	Specifies the new name for the group

The following example changes the class1 account group GID number to 400:

```
groupmod -g 400 class1
```

### The smgroup modify Command Format and Options

The following is the command format for the smgroup modify command:

```
/usr/sadm/bin/smgroup subcommand [auth_args] -- [subcommand_args]
```

Table 13-11 shows the options for the smgroup modify command.

**Table 13-11 Options for the smgroup modify Command**

Option	Description
<code>-n name</code>	Specifies the name of the group you want to modify
<code>-m new_member</code>	Specifies the new members to add to the group
<code>-N new_group</code>	Specifies the new group name

The following example changes the group workgroup to schoolgroup:

```
/usr/sadm/bin/smgroup modify -- -n workgroup -N schoolgroup
Authenticating as user: root
```

```
Type /? for help, pressing <enter> accepts the default denoted by []
Please enter a string value for: password :: Enter_The_root_Password
Loading Tool: com.sun.admin.usermgr.cli.group.UserMgrGroupCli from sys-02
Login to sys-02 as user root was successful.
Download of com.sun.admin.usermgr.cli.group.UserMgrGroupCli from sys-02
was successful.
```

## Deleting a Group Entry

Use the `groupdel` or `smgroup delete` commands to delete a group entry from the `/etc/group` file on the system.

### The `groupdel` Command Format

The following is the command format for the `groupdel` command:

```
groupdel groupname
```

The following example removes the group entry `class1` from the local system:

```
groupdel class1
```

### The `smgroup delete` Command Format and Options

The following is the command format for the `smgroup delete` command:

```
/usr/sadm/bin/smgroup subcommand [auth_args] -- [subcommand_args]
```

Use the `-n group_name` option with the `smgroup delete` command to specify the name of the group you want to delete.

The following example deletes the group entry `schoolgroup` from the local system:

```
/usr/sadm/bin/smgroup delete -- -n schoolgroup
```

```
Loading Tool: com.sun.admin.usermgr.cli.group.UserMgrGroupCli from sys-02
Login to sys-02 as user root was successful.
```

```
Download of com.sun.admin.usermgr.cli.group.UserMgrGroupCli from sys-02
was successful.
```

## Troubleshooting Login Issues

Some of the most common problems you might encounter as a system administrator are user login problems. There are two categories of login problems: login problems when the user logs in at the command line and login problems when the user logs in from the Common Desktop Environment (CDE) or the GNOME desktop.

The CDE uses more configuration files, so there are more potential problems associated with logging in from the CDE. When you troubleshoot a login problem, first determine whether you can log in from the command line. Attempt to log in from another system by using either the telnet command or the rlogin command, or click Options from the CDE login panel and select Command Line Login. If you can log in successfully at the command line, then the problem is with the CDE configuration files. If you cannot log in at the command line, then the problem is more serious and involves key configuration files.

### Login Problems at the Command Line

Table 13-12 presents an overview of common login problems that occur when the user logs in at the command line.

**Table 13-12** Login Problems at the Command Line

Login Problem	Description
Login incorrect	This message occurs when there are problems with the login information. The most common cause of an incorrect login message is a mistyped password. Make sure the that correct password is being used, and then attempt to enter it again. Remember that passwords are case-sensitive, so you cannot interchange uppercase letters and lowercase letters. In the same way, the letter "o" is not interchangeable with the numeral "0" nor is the letter "l" interchangeable with the numeral "1."
Permission denied	This message occurs when there are login, password, or NIS+ security problems. Most often, an administrator has locked the user's password or the user's account has been terminated.

**Table 13-12** Login Problems at the Command Line (Continued)

Login Problem	Description
Password will not work at lockscreen	A common error is to have the Caps Lock key on, which causes all letters to be uppercase. This does not work if the password contains lowercase letters.
No shell	This message occurs when the user's shell does not exist, is typed incorrectly, or is wrong in the /etc/passwd file.
No directory! Logging in with home=/	This message occurs when the user cannot access the home directory for one of the following reasons: An entry in the /etc/passwd file is incorrect, or the home directory has been removed or is missing, or the home directory exists on a mount point that is currently unavailable.
Choose a new password (followed by the New password: prompt)	This message occurs the first time a user logs in and chooses an initial password to access the account.
Couldn't fork a process!	This message occurs then the server could not fork a child process during login. The most common cause of this message is that the system has reached its maximum number of processes. You can either kill some unneeded processes (if you are already logged into that system as root) or increase the number of processes your system can handle.

### Login Problems in the CDE

Problems associated with logging into the CDE range from a user being unable to login (and returning to the CDE login screen), to the custom environment not loading properly. In general, the system does not return error messages to the user from the CDE. The following is a list of files and directories that provide troubleshooting information about the CDE:

- /usr/dt/bin/Xsession

This file is the configuration script for the login manager. This file should not be edited. The first user-specific file that the Xsession script calls is the \$HOME/.dtprofile file.

- `$HOME/.dtprofile`

By default, the file does not contain much content, except for examples. It contains a few echo statements for session logging purposes, and the `DTSOURCEPROFILE` variable is set. But it also contains information about how it might be edited. The user can edit this file to add user-specific environment variables.

- `DTSOURCEPROFILE=true`

This line allows the user's `$HOME/.login` file (for csh users) or the `$HOME/.profile` (for other shell users) to be sourced as part of the startup process.

Sometimes a `.login` or `.profile` file contains problem commands that cause the shell to crash. If the `.dtprofile` file is set to source a `.login` or `.profile` file that has problem commands, desktop startup might fail.

Consequently, no desktop appears. Instead, the system redisplays the Solaris OS CDE login screen. Startup errors from the `.login` or `.profile` file are usually noted in the `$HOME/.dt/startlog` file. Use a Failsafe login Session or a command-line login to debug problem commands in the `.login` or `.profile` files.

- `$HOME/.dt/sessions`

This directory structure contains files and directories that configure the display of the user's custom desktop and determine the applications that start when the user logs in. Look for recent changes to files and for changes to the directory structure. For example, examine the `home` directory and the `home.old` directory or a `current` directory and the `current.old` directory. Compare the changes. The changes could provide information on a new application or on changes in the saved desktop that cause the user's login to fail.

- `$HOME/.dt`

Upon removing the entire `.dt` directory structure, log out, and log back in again for the system to rebuild a default `.dt` file structure. This action allows the user to get back into the system if the problem with the CDE files cannot be resolved.

Table 13-13 shows the locations of and information found in error logs for the CDE.

**Table 13-13 CDE Error Log Locations**

Location	Error Log
/var/dt/Xerrors	The Solaris OS CDE login window system errors that occur prior to user login
\$HOME/.dt/startlog	The Solaris OS CDE errors that occur during the startup of the Xsession script, while processing the .dtprofile, .login, or .profile file
\$HOME/.dt/errorlog.old \$HOME/.dt/errorlog.older	The Solaris OS CDE errors that occur after the Xsession script start up
\$HOME/.dt/sessionlogs	Directory of session logs for Session Manager and Window Manager errors

You may now perform the Adding User and Group Accounts exercise.

## Login Problem in the GNOME Desktop

There may be a problem with the GNOME Desktop. A menu item was missing from the GNOME 2.0 Desktop and it may not appear in the login screen.

GNOME dtlogin resource files are currently only installed for the C locale. To make the GNOME login option appear in other locales, you must copy this resource file to the correct directory for your locale. Use the following command to copy the dtlogin file. Future releases will include localized versions of this file.

```
cp /usr/dt/config/C/Xresources.d/Xresources.Sun-gnome-2.0* /usr/dt/config/{your locale}/Xresources.d Managing Initialization Files
```

# Managing Initialization Files

The environment maintained by the shell includes variables that are defined by the `login` program, the system initialization files, and the user initialization files.

When users log in to the system, their login shells look for and execute two different types of initialization files. The first type controls the system-wide environment. The second type controls the user's environment. The six shells available in the Solaris 10 OS provide basic features and a set of variables which the `root` user or a regular user can set in the initialization files to customize the shell environment.

The shells support two types of variables:

- Environment variables – Variables that provide information about the user's environment to every shell program that is started.
- Local variables – Variables that affect only the current shell. Any subshell started would not have knowledge of these variables.

## Introducing System-Wide Initialization Files

As the system administrator, you maintain the system-wide initialization files. These files provide an environment for the entire community of users who log in to the system. The Solaris OS provides the system initialization files. They reside in the `/etc` directory.

The `/etc/profile` file and the `/etc/.login` file are the two main system initialization files.

The Bourne, Korn, and BASH login shells look for and execute the system initialization file `/etc/profile` during login.

The C login shell looks for and executes the system initialization file `/etc/.login` during the login process.

---

**Note** – The default files `/etc/profile` and `/etc/.login` check disk usage quotas, print the message of the day from the `/etc/motd` file, and check for mail. None of the messages are printed to the screen if the `.hushlogin` file exists in the user's home directory.

---



## Introducing User Initialization Files

As the system administrator, you set up the user initialization files that are placed in each user account's home directory when the user is created.

The primary purpose of the user initialization files is to define the characteristics of a user's work environment, such as the command-line prompt, the environment variables, and the windowing environment.

Only the owners of the files or the root user can change or customize the content of these files.

Table 13-14 shows the initialization files necessary for each primary shell available in the Solaris 10 OS.

**Table 13-14** Initialization Files for the Primary Shells

Shells	System-Wide Initialization Files	Primary User Initialization Files Read at Login	User Initialization Files Read When a New Shell Is Started	Shell Path Name
Bourne	/etc/profile	\$HOME/.profile		/bin/sh
Korn	/etc/profile	\$HOME/.profile \$HOME/.kshrc	\$HOME/.kshrc	/bin/ksh
C	/etc/.login	\$HOME/.cshrc \$HOME/.login	\$HOME/.cshrc	/bin/csh

For additional information about the Z, BASH, and TC shells available in the Solaris 10 OS, refer to the online manual pages.



**Note** – By default, the root user's login shell is the Bourne shell, and the shell entry in the /etc/passwd file appears as /sbin/sh.

When a user logs in to the system, the system invokes the user's login shell program. The shell program looks for its initialization files in a specific order, executes the commands contained in each file, and displays the shell prompt on the user's screen.

## Customizing the User's Work Environment

The Solaris OS provides a set of initialization file templates. The /etc/skel directory contains the initialization file templates. Table 13-15 shows the default initialization file templates and the user initialization files for the Bourne, Korn, and C shells.

**Table 13-15** Default User Initialization Files

Shell	Initialization File Templates	User Initialization Files
Bourne	/etc/skel/local.profile	\$HOME/.profile
Korn	/etc/skel/local.profile	\$HOME/.profile
C	/etc/skel/local.cshrc /etc/skel/local.login	\$HOME/.cshrc \$HOME/.login

**Note** – The useradd command copies files from the /etc/skel directory to the \$HOME directory. The smuser command copies files from the /etc/skel directory to the \$HOME directory and renames them to the appropriate file names.

The root user can customize these templates to create a standard set of user initialization files. A standard set of user initialization files provides a common work environment for each user. When the root user creates new user accounts, some or all of these initialization files are automatically copied to each new user's home directory.

Users can then edit their initialization files to further customize their environments for each shell.

Table 13-16 shows some of the variables available for customizing a user's shell environment.

**Table 13-16** Login Variables

Variable Name	Set By	Description
LOGNAME	Login	Defines the user's login name.
HOME	Login	Sets the path to the user's home directory. It is the default argument for the cd command.
SHELL	Login	Sets the path to the default shell.
PATH	Login	Sets the default path that the shell searches to find commands.
MAIL	Login	Sets the path to the user's mailbox.
TERM	Login	Defines the terminal.
LPDEST	Not set by default	Sets the user's default printer.
PWD	Shell	Defines the current working directory.
PS1	Shell	Defines the shell prompt for the Bourne or Korn shell.
prompt	Shell	Defines the shell prompt for the C shell.



**Note** – For complete information on all variables used by the default shells, see the following man pages: sh(1), ksh(1), csh(1), zsh(1), bash(1), and tcsh(1).

A user can change the values of the predefined variables and specify additional variables.

## Managing Initialization Files

---

Table 13-17 shows how to set environment variables in the user initialization files of the Bourne, Korn, and C shells.

**Table 13-17 Setting Environment Variables**

Shell	User's Initialization File
Bourne or Korn	<i>VARIABLE=value ; export VARIABLE</i> For example: <code>PS1="\$HOSTNAME "; export PS1</code>
C	<code>setenv variable value</code> For example: <code>setenv LPDEST laserprinter</code>

You may now perform the Modifying Initialization Files exercise.

# Exercise: Adding User and Group Accounts

In this exercise, you use the Solaris Management Console, and the smuser, smgroup, usermod, userdel, groupadd, and groupdel commands, to create, modify, and delete multiple user accounts and group entries.

## Preparation

This exercise assumes that no user accounts exist except for the default accounts created during a default Solaris OS installation. Use the logins command to display the list of user accounts. Look for accounts named using the names listed in Table 13-19. Use the userdel -r command to remove any of these accounts that you find.

Refer to the lecture notes, as necessary, to perform the tasks listed. Refer to Table 13-18 and Table 13-19 as needed.

**Table 13-18** Group Specifications

Group Name	GID Number
class1	101
class2	102

**Table 13-19** User Specifications

User Name	Password	Shell	UID	Primary Group	Secondary Group
user3	123pass	Korn	1003	10	class1
user4	123pass	C	1004	10	class1
user5	123pass	Bourne	1005	10	
locked1	Select Account is Locked	Bourne	2001	10	
cleared1	Select User must set password at next login	Bourne	2002	10	



**Note** – Some of the commands displayed in this section are quite long and will wrap to the next line. You should consider all of the bold typeface commands that follow a command line prompt to be all one line.

---

## Task 1 – Disabling the Solaris OS Registration Window

Complete the following steps:

1. Log in as the root user and open a terminal window (or use the **su** command to change your identity to root).
2. Change to the /etc/default directory.
3. In the /etc/default directory, use the **vi** editor to create a file called **solregis**.
4. In the **solregis** file, create a line that reads:  
**DISABLE=1**

**Note** – The character “1” is the number one.

---

5. Save this file, and then exit the editor.



## Task 2 – Adding Group Entries

Complete the following steps:



**Note** – Refer to Table 13-18 for details while adding groups with the tools used in this exercise.

---

1. As the root user, open a terminal window.
2. Use the **groupadd** and **smgroup** commands to add the groups **class1** and **class2**, respectively.  
The **smgroup** command will request a password for authentication. You should respond to the prompts as required.



## Task 3 – Adding User Accounts

Complete the following steps:

**Note** – Refer to Table 13-19 for details while adding users with the tools used in this exercise.

1. If the /export/home directory does not exist, create it.
2. Use the useradd command to add a user named user3, and then set a password for user3.
3. Use the smuser command to add a user named user4 and complete the following steps:
  - a. Respond to smuser authentication prompts, as required.

**Note** – Bug ID 6478291 for Solaris 10 Update 3 identifies an issue with the smuser and smrole commands. If smuser reports this error:

Authentication failed to the Solaris Management Console server, and you have specified the correct authentication information, use useradd instead to add the user4 account.

- b. As an alternative, use useradd to create the user4 account.
  - c. Set a password for user4.
4. Launch the Solaris Management Console by typing **smc&** on the command line.

After the Solaris Management Console displays, complete the following steps to create a user template and add user accounts that do not use automounted home directories:

- a. Select the This Computer toolbox in the Open Toolbox panel that displays when you start SMC, and click the Open button.
- b. Select This Computer, then System Configuration, then Users, and then User Templates to open the User Templates tool.  
You should respond to authentication prompts as required.
- c. From the Menu Bar, select Action, and then select Add User Template.

The Add User Template window appears, containing blank fields for a template name and description.

- d. Enter the name SolarisUser in the User Template Name field and Class Template for the Description field.

## Exercise: Adding User and Group Accounts

---

- e. Click the Home Directory Tab and complete the following steps:
  1. Uncheck the Automatically Mount Home Directory check box.
  2. Enter the name of your system in the Home Directory Server field.
- f. Click OK to create your template.
5. Complete the following steps:
  - a. Click User Accounts, and add the user5 account by selecting Action, then Add User, and then selecting From Template on the menu bar.  
The Add User From Template window appears.
  - b. Enter user5 in the User Name field and enter 1005 in the User ID Number field.
  - c. Click the User Must Use button, and then enter 123pass in both password fields for password information.
  - d. Click OK, and confirm your choice as requested.
6. From the Solaris Management Console, complete the following steps:
  - a. Use the SolarisUser template to add the users cleared1 and locked1.
  - b. Use the UID numbers as listed in Table 13-21.
  - c. Leave the User Account Is Locked button selected for both users.
7. Double-click the cleared1 user and complete the following steps:
  - a. Select the Password tab.
  - b. Select the password option User Must Set Password At Next Login.
  - c. Click OK.
8. Double-click the locked1 user and complete the following steps:
  - a. Select the General tab.
  - b. Under the Account Availability section, verify that the Account is Locked button is selected.
9. Exit the Solaris Management Console.

## Task 4 – Examining Configuration Files

Complete the following steps:

1. Examine the contents of the /etc/passwd file.

What are the full path names of the shells used by user3, user4, and user5?

user3

user4

user5

2. Examine the contents of the /etc/shadow file.

What text is found in the password field for the users locked1 and cleared1?

locked1

cleared1

You used the same password for user3 through user5.

Are the password strings the same in the /etc/shadow file?

- 
3. Examine the contents of the /etc/group file, and then verify that user3 and user4 are both listed as secondary members of the class1 group.

Are they?

- 
4. Log out of the window manager, and then attempt to log in as locked1.

Are you able to log in?

- 
5. Complete the following steps:

- a. Attempt to log in as cleared1.

What happens?

## Exercise: Adding User and Group Accounts

- b. Attempt to use the password abcdefg.

What are the system requirements for the password? You must not press Return when you are asked for an initial password.

---

---

- c. Use the password abc123, and then log in as cleared1 after you establish a password to verify that the login works.
- d. Log out, and then log in as the root user. Open a terminal window.

## Task 5 – Establishing Password Aging

Complete the following steps:

1. Start the Solaris Management Console and complete the following steps:
  - a. Go back into the User Accounts tool.
  - b. Respond to authentication prompts as required.
  - c. Select user5 from the list of users.
2. Complete the following steps:
  - a. Change the password options information for user5 so that it matches the following information.
  - b. Click OK when you are finished, and exit the Solaris Management Console.

User Must Keep For: 1 (one day)

Before Change Alert User: 2 (two days)

User Must Change Within: 2 (two days)

Expires If Not Used For: 1 (one day)

3. Log out of your root login session, and attempt to log in as user5. What happens?

---

**Note** – If the password expiry warning does not display when you log in, you can test the password aging feature by using the su user5 command. Doing so will cause the message Your password will expire in 2 days to display in the terminal window where you run the su user5 command.

---

4. Complete the login as user5 and complete the following steps:
- Open a terminal window.
  - Attempt to change the password for user5.
- What happens?

- 
- 
5. Log out, and then log in again as the root user. Open a terminal window.

## Task 6 – Modifying User Accounts and Group Entries

Complete the following steps:

- Use the groupadd command to create a new group called class3 that uses GID number 103.
- Complete the following steps:
  - Use the usermod command to change the following information for the locked1 user:
    - The login name to user6
    - The UID to 3001
    - The home directory to user6
  - Verify that the changes you request are recorded in the /etc/passwd file and that the directory was moved.

## Exercise: Adding User and Group Accounts

---

3. Use the smuser modify command to change the login shell of user5 to /bin/ksh and complete the following steps:
  - a. Respond to authentication prompts as required.



**Note** – Bug ID 6478291 for Solaris 10 Update 3 identifies an issue with the smuser and smrole commands. If smuser reports this error: Authentication failed to the Solaris Management Console server, and you have specified the correct authentication information, use usermod instead to modify the user5 account.

---

- b. As an alternative, use usermod to modify the user5 account.
- c. Verify that the changes you request are recorded in the /etc/passwd file.
4. Complete the following steps:
  - a. Use the userdel command to delete the user account cleared1 and the related home directory.
  - b. Verify that the /export/home/cleared1 directory no longer exists.
5. Use the smgroup command to change the group name of class1 to group1 and complete the following steps:
  - a. Respond to authentication prompts as required.
  - b. Verify the change in the /etc/group file.
6. Complete the following steps:
  - a. Use the groupdel command to remove the group entry class2.
  - b. Verify that class2 is no longer in the /etc/group file.

## Exercise Summary



**Discussion** – Take a few minutes to discuss what experiences, issues, or discoveries you had during the lab exercises.

- Experiences
- Interpretations
- Conclusions
- Applications

# Exercise Solutions: Adding User and Group Accounts

This section contains solutions to the exercises.

Refer to the lecture notes as necessary to perform the tasks listed. Refer to Table 13-20 and Table 13-21 as needed.

**Table 13-20** Group Specifications – Solution Section

Group Name	GID Number
class1	101
class2	102

**Table 13-21** User Specifications – Solution Section

User Name	Password	Shell	UID	Primary Group	Secondary Group
user3	123pass	Korn	1003	10	class1
user4	123pass	C	1004	10	class1
user5	123pass	Bourne	1005	10	
locked1	Select Account is Locked	Bourne	2001	10	
cleared1	Select User must set password at next login	Bourne	2002	10	

**Note** – Some of the commands displayed in this section are quite long and will wrap to the next line. You should consider all of the bold typeface commands that follow a command line prompt to be all one line.



## Task 1 – Disabling the Solaris OS Registration Window

Complete the following steps:

1. Log in as the root user and open a terminal window (or use the su command to change your identity to root).
2. Change to the /etc/default directory.
3. In the /etc/default directory, use the vi editor to create a file called solregis.  
**# vi solregis**
4. In the solregis file, create a line that reads:  
**DISABLE=1**

---

**Note** – The character “1” is the number one.

5. Save this file, and then exit the editor.



## Task 2 – Adding Group Entries

Complete the following steps:




---

**Note** – Refer to Table 13-20 for details while adding groups with the tools used in this exercise.

1. As the root user, open a terminal window.
2. Use the groupadd and smgroup commands to add the groups class1 and class2, respectively.

The smgroup command will request a password for authentication. You should respond to the prompts as required.

```
groupadd -g 101 class1
/usr/sadm/bin/smgroup add -- -n class2 -g 102
Authenticating as user: root
```

Type /? for help, pressing <enter> accepts the default denoted by [ ]  
Please enter a string value for: password :: **cangetin**

Loading Tool: com.sun.admin.usermgr.cli.group.UserMgrGroupCli from host2  
Login to host2 as user root was successful.

Download of com.sun.admin.usermgr.cli.group.UserMgrGroupCli from host2 was successful.

#

## Task 3 – Adding User Accounts

Complete the following steps:



**Note** – Refer to Table 13-21 for details while adding users with the tools used in this exercise.

1. If the /export/home directory does not exist, create it.

```
ls /export/home
/export/home: No such file or directory
mkdir /export/home
#
```

2. Use the useradd command to add a user named user3, and then set a password for user3.

```
useradd -u 1003 -g 10 -G class1 -d /export/home/user3 -m -s /bin/ksh
user3
64 blocks
passwd user3
New Password: 123pass
Re-enter new Password: 123pass
passwd: password successfully changed for user3
```

3. Use the smuser command to add a user named user4 and complete the following steps:

- a. Respond to smuser authentication prompts, as required.

```
/usr/sadm/bin/smuser add -- -n user4 -u 1004 -g 10 -G class1 -d
/export/home/user4 -s /bin/csh -x autohome=N
Authenticating as user: root
```

Type /? for help, pressing <enter> accepts the default denoted by [ ]  
Please enter a string value for: password :: **cangetin**

Loading Tool: com.sun.admin.usermgr.cli.user.UserMgrCli from host2

Login to host2 as user root was successful.

Download of com.sun.admin.usermgr.cli.user.UserMgrCli from host2 was successful.

#

- b. As an alternative, use useradd to create the user4 account.

```
useradd -u 1004 -g 10 -G class1 -d /export/home/user4 -m -s /bin/csh
user4
64 blocks
#
```

- c. Set a password for user4.

```
passwd user4
```

```
New Password: 123pass
```

```
Re-enter new Password: 123pass
```

```
passwd: password successfully changed for user4
```

4. Launch the Solaris Management Console by typing **smc&** on the command line.

After the Solaris Management Console displays, complete the following steps to create a user template and add user accounts that do not use automounted home directories:

- a. Select This Computer, then System Configuration, then Users, and then User Templates to open the User Templates tool.

You should respond to authentication prompts as required.

- b. From the Menu Bar, select Action, and then select Add User Template.

The Add User Template window appears, containing blank fields for a template name and description.

- c. Enter the name SolarisUser in the User Template Name field and Class Template for the Description field.

- d. Click the Home Directory Tab and complete the following steps:

1. Uncheck the Automatically Mount Home Directory check box.

2. Enter the name of your system in the Home Directory Server field.

- e. Click OK to create your template.

5. Complete the following steps:

- a. Click User Accounts, and add the user5 account by selecting Action, then Add User, and then selecting From Template on the menu bar.

The Add User From Template window appears.

- b. Enter user5 in the User Name field and enter 1005 in the User ID Number field.

## Exercise Solutions: Adding User and Group Accounts

---

- 
- 
- 
- 
- 
- c. Click the User Must Use button, and then enter 123pass in both password fields for password information.
- d. Click OK, and confirm your choice as requested.
6. From the Solaris Management Console, complete the following steps:
  - a. Use the SolarisUser template to add the users cleared1 and locked1.
  - b. Use the UID numbers as listed in Table 13-21.
  - c. Leave the User Account Is Locked button selected for both users.
7. Double-click the cleared1 user and complete the following steps:
  - a. Select the Password tab.
  - b. Select the password option User Must Set Password At Next Login.
  - c. Click OK.
8. Double-click the locked1 user and complete the following steps:
  - a. Select the General tab.
  - b. Under the Account Availability section, verify that the Account is Locked button is selected.
9. Exit the Solaris Management Console.

## Task 4 – Examining Configuration Files

Complete the following steps:

1. Examine the contents of the /etc/passwd file.

What are the full path names of the shells used by user3, user4, and user5?

user3	/bin/ksh
user4	/bin/csh
user5	/bin/sh

2. Examine the contents of the /etc/shadow file.

What text is found in the password field for the users locked1 and cleared1?

locked1	*LK*
cleared1	no value

You used the same password for user3 through user5.

Are the password strings the same in the /etc/shadow file?

No.

3. Examine the contents of the /etc/group file, and then verify that user3 and user4 are both listed as secondary members of the class1 group.

Are they?

*The names user3 and user4 should be listed in the last field for the class1 group.*

4. Log out of the window manager, and then attempt to log in as locked1.

Are you able to log in?

*No, you get a message that says login incorrect, no matter what you use as a password.*

5. Complete the following steps:
- Attempt to log in as cleared1.  
What happens?
  - Attempt to use the password abcdefg.  
What are the system requirements for the password? You must not press Return when you are asked for an initial password.  
*You must choose an initial password for this user and then log in again. The first six characters must contain at least two alphabetic characters and at least one numeric or special character.*
  - Use the password abc123, and then log in as cleared1 after you establish a password to verify that the login works.
  - Log out, and then log in as the root user. Open a terminal window.

## Task 5 – Establishing Password Aging

Complete the following steps:

- Start the Solaris Management Console and complete the following steps:
  - Go back into the User Accounts tool.
  - Respond to authentication prompts as required.
  - Select user5 from the list of users.
- Complete the following steps:
  - Change the password options information for user5 so that it matches the following information.

User Must Keep For: 1 (one day)

Before Change Alert User: 2 (two days)

User Must Change Within: 2 (two days)

Expires If Not Used For: 1 (one day)

- Click OK when you are finished, and exit the Solaris Management Console.

3. Log out of your root login session, and then attempt to log in as user5.

What happens?

*You can log in directly as user5, and a warning that indicates that your password expires in two days displays.*




---

**Note** – If the password expiry warning does not display when you log in, you can test the password aging feature by using the `su user5` command. Doing so will cause the message `Your password will expire in 2 days` to display in the terminal window where you run the `su user5` command.

---

4. Complete the login as user5 and complete the following steps:

a. Open a terminal window.

b. Attempt to change the password for user5.

What happens?

*When you try to change your password, the following error message displays:*

`passwd: Sorry: less than 1 days since the last change.  
Permission denied`

5. Log out, and then log in again as the root user. Open a terminal window.

## Task 6 – Modifying User Accounts and Group Entries

Complete the following steps:

1. Use the `groupadd` command to create a new group called `class3` that uses GID number 103.

`# groupadd -g 103 class3`

2. Complete the following steps:

a. Use the `usermod` command to change the following information for the `locked1` user:

- The login name to `user6`
- The UID to 3001
- The home directory to `/user6`

- b. Verify that the changes you request are recorded in the /etc/passwd file and that the directory was moved.

```
ls -d /export/home/locked1
/export/home/locked1
usermod -u 3001 -d /export/home/user6 -m -l user6 locked1
48 blocks
ls -d /export/home/locked1
/export/home/locked1: No such file or directory
ls -d /export/home/user6
/export/home/user6
grep user6 /etc/passwd
user6:x:3001:10:::/export/home/user6:/bin/sh
#
```

*The /etc/passwd file should reflect the new UID number and user name. The directory under /export/home should be renamed from locked1 to user6.*

3. Use the smuser modify command to change the login shell of user5 to /bin/ksh and complete the following steps:
- Respond to authentication prompts as required.

```
grep user5 /etc/passwd
user5:x:1005:10:::/export/home/user5:/bin/sh
/usr/sadm/bin/smuser modify -- -n user5 -s /bin/ksh
Authenticating as user: root
```

Type /? for help, pressing <enter> accepts the default denoted by [ ]  
Please enter a string value for: password :: **cangetin**  
Loading Tool: com.sun.admin.usermgr.cli.user.UserMgrCli from host2  
Login to host2 as user root was successful.  
Download of com.sun.admin.usermgr.cli.user.UserMgrCli from host2 was successful.

- b. As an alternative, use usermod to modify the user5 account.

```
usermod -s /bin/ksh user5
#
```

- c. Verify that the changes you request are recorded in the /etc/passwd file.

```
grep user5 /etc/passwd
user5:x:1005:10::/export/home/user5:/bin/ksh
#
```

*The /etc/passwd file should show that the shell for user5 is /bin/ksh.*

4. Complete the following steps:

- Use the userdel command to delete the user account cleared1 and the related home directory.
- Verify that the /export/home/cleared1 directory no longer exists.

```
userdel -r cleared1
ls /export/home/cleared1
/export/home/cleared1: No such file or directory
#
```

*The /export/home/cleared1 directory should no longer exist.*

5. Use the smgroup command to change the group name of class1 to group1 and complete the following steps:
- Respond to authentication prompts as required.
  - Verify the change in the /etc/group file.

```
grep class1 /etc/group
class1::101:user3,user4
/usr/sadm/bin/smgroup modify -- -n class1 -N group1
Authenticating as user: root
```

Type /? for help, pressing <enter> accepts the default denoted by [ ]  
Please enter a string value for: password :: cangetin

Loading Tool: com.sun.admin.usermgr.cli.group.UserMgrGroupCli from host2  
Login to host2 as user root was successful.

Download of com.sun.admin.usermgr.cli.group.UserMgrGroupCli from host2  
was successful.

```
grep class1 /etc/group
grep group1 /etc/group
group1::101:user3,user4
#
```

6. Complete the following steps:

- a. Use the groupdel command to remove the group entry class2.
- b. Verify that class2 is no longer in the /etc/group file.

```
groupdel class2
grep class2 /etc/group
#
```

*The group class2 should not exist.*

---

## Module 14

---

# Controlling System Processes

---

## Objectives

Upon completion of this module, you should be able to:

- View system processes
- Kill frozen processes
- Schedule an automatic one-time execution of a command
- Schedule an automatic recurring execution of a command

# Viewing System Processes

A process is any program running on a system. All processes are assigned a unique process identification (PID) number, which is used by the kernel to track and manage the process. The PID numbers are used by the root and regular users to identify and control their processes.

## Using the prstat Command

The prstat command examines and displays information about active processes on the system.

This command enables you to view information by specific processes, user identification (UID) numbers, central processing unit (CPU) IDs, or processor sets. By default, the prstat command displays information about all processes sorted by CPU usage. To use the prstat command, perform the command:

```
prstat
 PID USERNAME SIZE RSS STATE PRI NICE TIME CPU PROCESS/NLWP
 1641 root 4864K 4520K cpu0 59 0 0:00:00 0.5% prstat/1
 1635 root 1504K 1168K sleep 59 0 0:00:00 0.3% ksh/1
 9 root 6096K 4072K sleep 59 0 0:00:29 0.1% svc.configd/11
 566 root 82M 30M sleep 29 10 0:00:36 0.1% java/14
 1633 root 2232K 1520K sleep 59 0 0:00:00 0.1% in.rlogind/1
 531 root 8200K 2928K sleep 59 0 0:00:12 0.1% dtgreet/1
 474 root 21M 7168K sleep 59 0 0:00:11 0.1% Xsun/1
 236 root 4768K 2184K sleep 59 0 0:00:03 0.0% inetd/4
 86 root 3504K 1848K sleep 59 0 0:00:01 0.0% nscd/24
 7 root 5544K 1744K sleep 59 0 0:00:06 0.0% svc.startd/12
 154 root 2280K 824K sleep 59 0 0:00:01 0.0% in.routed/1
 509 root 6888K 2592K sleep 59 0 0:00:02 0.0% httpd/1
 240 root 5888K 1256K sleep 59 0 0:00:01 0.0% sendmail/1
 145 root 2944K 816K sleep 59 0 0:00:01 0.0% httpd/1
 347 daemon 2608K 776K sleep 59 0 0:00:00 0.0% nfsmapid/3
 206 root 1288K 600K sleep 59 0 0:00:00 0.0% utmpd/1
 344 daemon 2272K 1248K sleep 60 -20 0:00:00 0.0% nfsd/2
 241 smmsp 5792K 960K sleep 59 0 0:00:00 0.0% sendmail/1
 107 root 2584K 784K sleep 59 0 0:00:00 0.0% syseventd/14
 123 root 3064K 880K sleep 59 0 0:00:00 0.0% picld/4
 146 lp 2976K 448K sleep 59 0 0:00:00 0.0% httpd/1
Total: 53 processes, 171 lwps, load averages: 0.02, 0.04, 0.07
```

#

To quit the prstat command, type **q**.

Table 14-1 shows the column headings and their meanings in a prstat report.

**Table 14-1** prstat Report Column Headings

Default Column Heading	Description
PID	PID process number
USERNAME	Login name or UID of the owner of the process
SIZE	Total virtual memory size of the process
RSS	Resident process set size in kilobytes, megabytes, or gigabytes
STATE	Process state <ul style="list-style-type: none"> <li>• <b>cpu</b> – process is running on CPU.</li> <li>• <b>sleep</b> – process is waiting for event to complete.</li> <li>• <b>run</b> – process is in run queue</li> <li>• <b>zombie</b> – process terminated, and the parent is not waiting</li> <li>• <b>stop</b> – process stopped</li> </ul>
PRI	Process priority
NICE	Value used in priority computation
TIME	Cumulative process execution time
CPU	Percentage of recent CPU time used by the process.
PROCESS/NLWP	Name of process/number of lightweight processes (LWPs) in the process.



**Note** – The kernel and many applications are now multithreaded. A thread is a logical sequence of program instructions written to accomplish a particular task. Each application thread is independently scheduled to run on an LWP, which functions as a virtual CPU. LWPs in turn, are attached to kernel threads, which are scheduled to run on actual CPUs.



**Note** – Use the priocntl(1) command to assign processes to a priority class and to manage process priorities. The nice(1) command is only supported for backward compatibility to previous Solaris OS releases. The priocntl command provides more flexibility in managing processes.

Table 14-2 shows the options for the prstat command.

**Table 14-2** Options for the prstat Command

Option	Description
<code>-a</code>	Displays separate reports about processes and users at the same time.
<code>-c</code>	Continuously prints new reports below previous reports.
<code>-n nproc</code>	Restricts the number of output lines.
<code>-p pidlist</code>	Reports only on processes that have a PID in the given list.
<code>-s key</code>	Sorts output lines by <i>key</i> in descending order. The five possible keys include: <i>cpu</i> , <i>time</i> , <i>size</i> , <i>rss</i> , and <i>pri</i> . You can use only one key at a time.
<code>-S key</code>	Sorts output lines by <i>key</i> in ascending order.
<code>-t</code>	Reports total usage summary for each user.
<code>-u euidlist</code>	Reports only processes that have an effective user ID (EUID) in the given list.
<code>-U uidlist</code>	Reports only processes that have a real UID in the given list.

# Killing Frozen Processes

Use the kill command or the pkill command to send a signal to one or more running processes. You would typically use these commands to kill an unwanted process.

## Using the kill and pkill Commands

You use the kill or pkill commands to terminate one or more processes.

The format for the kill command is:

```
kill -signal PID
```

To show all of the available signals used with the kill command:

```
kill -l
```

The format for the pkill command is:

```
pkill -signal Process
```

Before you can terminate a process, you must know its name or PID. Use either the ps or pgrep command to locate the PID for the process.

The following examples uses the pgrep command to locate the PID for the mail processes.

```
pgrep -l mail
 241 sendmail
 240 sendmail
#
pkill sendmail
```

The following examples use the ps and pkill to locate and end sendmail.

```
ps -e | grep sendmail
 241 ? 0:00 sendmail
 240 ? 0:02 sendmail
kill 241
```

To terminate more than one process at the same time, use the following syntax:

```
kill -signal PID PID PID PID
pkill signal process process
```

Use the `kill` command without a signal on the command line to send the default Signal 15 to the process. This signal usually causes the process to terminate.

Table 14-3 shows some signals and names.

**Table 14-3** Process Signal Numbers and Names

Signal Number	Signal Name	Event	Default Action
1	SIGHUP	Hangup	Exit
2	SIGINT	Interrupt	Exit
9	SIGKILL	Kill	Exit
15	SIGTERM	Terminate	Exit

- 1, SIGHUP – A hangup signal to cause a telephone line or terminal connection to be dropped. For certain daemons, such as `inetd` and `in.named`, a hangup signal will cause the daemon to reread its configuration file.
- 2, SIGINT – An interrupt signal from your keyboard—usually from a Control-C key combination.
- 9, SIGKILL – A signal to kill a process. A process cannot ignore this signal.
- 15, SIGTERM – A signal to terminate a process in an orderly manner. Some processes ignore this signal.

Use `kill -l`, or see the `signal` man page, for a complete list of signals that the `kill` command can send.

```
man -s3head signal
```

Some processes can be written to ignore Signal 15. Processes that do not respond to a Signal 15 can be terminated by force by using Signal 9 with the `kill` or `pkill` commands. You use the following syntax:

```
kill -9 PID
pkill -9 process
```



**Caution** – Use the kill -9 or pkill -9 command as a last resort to terminate a process. Using the -9 signal on a process that controls a database application or a program that updates files can be disastrous. The process is terminated instantly with no opportunity to perform an orderly shutdown.

## Performing a Remote Login

When a workstation is not responding to your keyboard or mouse input, the CDE might be frozen. In such cases, you may be able to remotely access your workstation by using the rlogin command or by using the telnet command from another system.

### Killing the Process for a Frozen Login

After you are connected remotely to your system, invoke the pkill command to terminate the corrupted session.

In the following examples, the rlogin command is used to log in to sys42, from which you can issue a pkill or a kill command.

```
rlogin sys-02
Password:
Last login: Sun Oct 24 13:44:51 from sys-01
Sun Microsystems Inc. SunOS 5.10 s10_68 Sep. 20, 2004
pkill -9 Xsun
or
ps -e | grep Xsun
 442 ? 0:01 Xsun
kill -9 442
```

# Scheduling an Automatic One-Time Execution of a Command

Use the `at` command to automatically execute a job only once at a specified time.

## Using the `at` Command

The format for the `at` command is:

```
at -m -q queuename time date
at -r job
at -l
```

Table 14-4 shows the options you can use to instruct the cron process on how to execute an `at` job.

**Table 14-4** Options for the `at` Command

Option	Description
<code>-m</code>	Sends mail to the user after the job has finished
<code>-r job</code>	Removes a scheduled <code>at</code> job from the queue
<code>-q queuename</code>	Specifies a specific queue
<code>time</code>	Specifies a time for the command to execute
<code>-l</code>	Reports all jobs scheduled for the invoking user
<code>date</code>	Specifies an optional date for the command to execute, which is either a month name followed by a day number or a day of the week

For example, to create an at job to run at 9:00 p.m. to locate and verify the file type of core files from the /export/home directory, perform the command:

```
at 9:00 pm
at> find /export/home -name core -exec file {} \; >> /var/tmp/corelog
at> <EOT>
commands will be executed using /sbin/sh
job 1098648000.a at Mon Oct 25 21:00:00 2004
```




---

**Note** – The at command job numbers show the job's execution time as reflected in the amount of time in seconds since 12:00 A.M. January 1, 1970, which is referred to as the UNIX® epoch date.

---

To display information about the execution times of jobs, perform the command:

```
at -l 1098648000.a
1098648000.a Mon Oct 25 21:00:00 2004
```

To display the jobs queued to run at specified times by chronological order of execution, perform the command:

```
atq
Rank Execution Date Owner Job Queue Job Name
1st Oct 25, 2004 21:00 root 1098648000.a a stdin
```

To view all the at jobs currently scheduled in the queue, perform the command:

```
ls -l /var/spool/cron/atjobs
total 4
-r-Sr--r-- 1 root root 1044 Oct 25 13:48 1098648000.a
```

You can also use the at command to remove a job from the at queue.

For example, to remove job 1098648000.a from the at queue, perform the command:

```
at -r 1098648000.a
atq
Rank Execution Date Owner Job Queue Job Name
```

## Controlling Access to the at Command

As the root user, you control who has access to the at command with the at.deny and at.allow files.

### The /etc/cron.d/at.deny File

By default, the Solaris OS includes the /etc/cron.d/at.deny file. This file identifies users who are prohibited from using the at command. The file format is one user name per line. The file initially contains:

```
daemon
bin
nuucp
listen
nobody
noaccess
```

A user who is denied access to the at command receives the following message when attempting to use the command:

```
at: you are not authorized to use at. Sorry.
```

If only the /etc/cron.d/at.deny file exists but is empty, then all logged-in users can access the at command.

### The /etc/cron.d/at.allow File

The /etc/cron.d/at.allow file does not exist by default, so all users (except those listed in the /etc/cron.d/at.deny file) can create at jobs. By creating the /etc/cron.d/at.allow file, you create a list of only those users who are allowed to execute at commands. The /etc/cron.d/at.allow file consists of user names, one per line.

The interaction between the at.allow and the at.deny files follows these rules:

- If the at.allow file exists, only the users listed in this file can execute at commands. This rule also applies to the root user.
- If the at.allow file does not exist, all users, except for users listed in the at.deny file, can execute at commands.
- If neither file exists, only the root user can use the at command.
- If a user is listed in both files, the user is allowed to use at.

# Scheduling an Automatic Recurring Execution of a Command

You can use the cron facility to schedule regularly recurring commands. Users can submit a command to the cron facility by modifying their crontab file.

All crontab files are maintained in the /var/spool/cron/crontabs directory and are stored as the login name of the user that created the cron job.

The cron daemon is responsible for scheduling and running these jobs.

---

**Note** – The clock daemon, cron, starts at system boot and runs continuously in the background.

---



## Introducing the crontab File Format

A crontab file consists of lines of six fields each. The fields are separated by spaces or tabs. The first five fields provide the date and time the command is to be scheduled. The last field is the full path to the command.

---

**Note** – If the command field contains a percent (%) character, then all subsequent characters are passed to the command as standard input.

---



## Scheduling an Automatic Recurring Execution of a Command

These first five fields are separated by spaces and indicate when the command will be executed. See Figure 14-1.

10 3 \* \* 0 /usr/sbin/log/adm

The *minute* field can hold values between 0 and 59.

The *hour* field can hold values between 0 and 23.

The *day-of-month* field can hold values between 1 and 31.

The *month* field can hold values between 1 and 12, January to December.

The *day-of-week* field can hold values between 0 and 6. Sunday is 0.

The *command* field contains the full path name to the command to be run by the `cron` utility.

**Figure 14-1** First Five Fields in a crontab File

The first five fields follow the format rules shown in Table 14-5.

**Table 14-5** Rules for the crontab Fields

Value	Rule	Example
<i>n</i>	Matches if field value is <i>n</i>	As shown in the preceding figure for hour or minute, a 3 or 10
<i>n,p,q</i>	Matches if field value is <i>n</i> , <i>p</i> , or <i>q</i>	Every 10 minutes would be represented by 0,10,20,30,40,50
<i>n-p</i>	Matches if field has values between <i>n</i> and <i>p</i> inclusive	The hours between 1:00 a.m. and 4:00 a.m. would be represented by 1-4
*	Matches all legal values	As in the preceding example for the month, representing every month.

## Using the crontab Command

The crontab command enables the user to view, edit, or remove a crontab file.

### Viewing a crontab File

To view the contents of the root crontab file, run the crontab -l command as the root user.

```
crontab -l
#ident "@(#)root 1.21 04/03/23 SMI"
#
The root crontab should be used to perform accounting data collection.
#
#
10 3 * * * /usr/sbin/logadm
15 3 * * 0 /usr/lib/fs/nfs/nfsfind
30 3 * * * [-x /usr/lib/gss/gsscared_clean] &&
/usr/lib/gss/gsscared_clean
#10 3 * * * /usr/lib/krb5/kprop_script __slave_kdcs__
```

This is the same command that users run to view the contents of their own crontab file.

As the root user, you can view the contents of any regular user's crontab file by performing the command:

```
crontab -l username
```

### Editing a crontab File



---

**Caution** – If you accidentally enter the crontab command on the command line without an option (-l, -e, -r), press the interrupt keys Control-C to exit. Do not press Control-D, this action overwrites the existing crontab file with an empty file.

---

To create or edit a crontab file, follow these steps:

1. Check that the EDITOR variable is set to the editor you want to use. This instructs the cron utility which editor to use to open the file.  

```
EDITOR=vi
export EDITOR
```
2. Run the following crontab command to open your crontab file, and add the appropriate entry.

```
crontab -e
30 17 * * 5 /usr/bin/banner "Time to go!" > /dev/console
:wq
```



---

**Note** – If the users do not redirect the standard output and standard error of their commands in the crontab file, any generated output or errors are mailed electronically to the user.

---

### Removing a crontab File

The correct way to remove a crontab file is to invoke the command:

```
crontab -r username
```

Typical users can remove only their own crontab file. The root user can delete any user's crontab file.

## Controlling Access to the crontab Command

You can control access to the crontab command with two files in the /etc/cron.d directory—the cron.deny file and the cron.allow file.

These files permit only specified users to perform crontab tasks, such as creating, editing, displaying, or removing their own crontab files.

### The /etc/cron.d/cron.deny File

The Solaris OS provides a default cron.deny file. The file consists of a list of user names, one per line, of the users who are not allowed to use cron. The following is an example of the contents of a cron.deny file:

```
daemon
bin
nuucp
listen
nobody
noaccess
```

### The /etc/cron.d/cron.allow File

The /etc/cron.d/cron.allow file does not exist by default, so all users (except those listed in the cron.deny file) can access their crontab file. By creating a cron.allow file, you can list only those users who can access crontab commands.

The file consists of a list of user names, one per line.

The interaction between the cron.allow and the cron.deny files follows these rules:

- If the cron.allow file exists, only the users listed in this file can create, edit, display, or remove crontab files.
- If the cron.allow file does not exist, all users, except for users listed in the cron.deny file, can create, edit, display, or remove crontab files.

## Scheduling an Automatic Recurring Execution of a Command

---

- If neither file exists, only the root user can run the crontab command.

---

**Note** – The Solaris OS provides commands that check authorizations. The batch job-related commands, such as at, atq, batch, and crontab are such commands

---

# Exercise: Using Process Control

In this exercise, you complete the following tasks:

- Use the Process Manager Tool and the `prstat` command to monitor and kill processes.
- Create an `at` job and create an entry in a `crontab` file.

## Preparation

Refer to the lecture notes as necessary to perform the tasks listed.

## Task

Complete the following steps:

1. Log in as the `root` user, and open a terminal window. Press `Ctrl+Alt+Del` to start the Process Manager or invoke the appropriate command from the command line.
2. In the Process Tool display, sort the listing according to `CPU%`, and change the sample time to five seconds.
3. Open a second terminal window, and run the `prstat` command.
4. Position the Process Tool and the window in which the `prstat` command is running so that you can observe both simultaneously. In an available window, run the `find` command to list all files on your system. Observe how the Process Tool and the `prstat` command display statistics for the `find` command.

What is the maximum percentage of recent CPU time used by the `find` command as it executes?

- 
5. Open a third terminal window, and run the `ps` command to determine the PID of the shell. Record the PID you find.

---

  6. In the Process Tool, locate and select the shell process you identified in the previous step. Select the Show Ancestry option from the Process menu in the Process Tool. What is the name and PID of the first process listed?

---

## Exercise: Using Process Control

---

7. Close the Show Ancestry window. Again, select the shell process you identified in Step 4. From the Process menu in the Process Tool, select the Kill option. What happens?
- 
8. In the Process Tool, use the Find function to locate the prstat process. Select the Signal option from the Process menu. In the Signal fill-in field, enter the TERM signal, and click OK. What happens to the prstat process? Close the Process Tool when you are finished.
- 
9. Identify the device associated with your current terminal by using the `tty` command, and display the current time of day.
- 
10. Submit an `at` job that echoes Test Complete to your current window. Have the job run about five minutes from the current time, and submit it to the queue called `x`.
11. Display the `at` job in the queue.
12. Open a new window and set and export the `EDITOR` environment variable to use the `vi` editor to edit `crontab` files.
13. Use the `crontab` command to view the current `crontab` file for the root user.
14. When is the `logadm` process scheduled to run?
- 
15. Use the `crontab` command to edit the `crontab` file for the root user. Add an entry that sends the message It works! to your current window five minutes from now. For example, if the current time is 10:25, make an entry in your `crontab` file for the 30th minute of the same hour.

Save the file, and quit the `vi` edit session. In about five minutes, you should see the result in your window.

## Exercise Summary



**Discussion** – Take a few minutes to discuss what experiences, issues, or discoveries you had during the lab exercises.

- Experiences
- Interpretations
- Conclusions
- Applications

# Exercise Solutions: Using Process Control

This section contains solutions to the exercise.

## Task

Complete the following steps:

1. Log in as the root user, and open a terminal window. Press Ctrl+Alt+Del to start the Process Manager or invoke the appropriate command from the command line.

**# /usr/dt/bin/sdtprocess &**

In the Process Tool display, sort the listing according to CPU%, and change the sample time to five seconds.

2. Open a second terminal window, and run the prstat command.

**# prstat**

3. Position the Process Tool and the window in which the prstat command is running so that you can observe both simultaneously. In an available window, run the find command to list all files on your system. Observe how the Process Tool and the prstat command display statistics for the find command.

**# find /**

What is the maximum percentage of recent CPU time used by the find command as it executes?

*This varies according to your system configuration. Some systems might display values in the 20-percent range.*

4. Open a third terminal window, and run the ps command to determine the PID of the shell. Record the PID you find.

**# ps**

*Your value appears here.*

5. In the Process Tool, locate and select the shell process you identified in the previous step. Select the Show Ancestry option from the Process menu in the Process Tool. What is the name and PID of the first process listed?

*On systems running JDS, the first process should be gnome-terminal.*

6. Close the Show Ancestry window. Again, select the shell process you identified in Step 4. From the Process menu in the Process Tool, select the Kill option. What happens?

*The process stops, and the window no longer appears.*

7. In the Process Tool, use the Find function to locate the prstat process. Select the Signal option from the Process menu. In the Signal fill-in field, enter the TERM signal, and click OK. What happens to the prstat process? Close the Process Tool when you are finished.

*The prstat process terminates, and the prompt appears in the window in which it ran.*

8. Identify the device associated with your current terminal by using the tty command, and display the current time of day.

```
tty
/dev/pts/2
#
date
Tue Mar 27 18:07:52 MDT 2007
#
```

9. Submit an at job that echoes Test Complete to your current window. Have the job run about five minutes from the current time, and submit it to the queue called x.

```
at -q x 18:13
at> echo "Test Complete" > /dev/pts/# (# is from the tty command)
at> <Control-D>
commands will be executed using /bin/ksh
job 1175040780.x at Tue Mar 27 18:13:00 2007
#
```

10. Display the at job in the queue.

```
atq
atq
Rank Execution Date Owner Job Queue Job Name
 1st Mar 27, 2007 18:13 root 1175040780.x x stdin
#
```

11. Open a new window and set and export the EDITOR environment variable to use the vi editor to edit crontab files.

If you are using the Bourne or Korn shell, perform the command:

```
EDITOR=vi
export EDITOR
```

If you are using the C shell, perform the command:

```
setenv EDITOR vi
```

12. Use the crontab command to view the current crontab file for the root user.

```
crontab -l
```

## Exercise Solutions: Using Process Control

---

```
#ident "@(#)" root 1.21 04/03/23 SMI"
#
The root crontab should be used to perform accounting data collection.
#
#
10 3 * * * /usr/sbin/logadm
15 3 * * 0 /usr/lib/fs/nfs/nfsfind
(output omitted)
```

13. When is the `logadm` process scheduled to run?

*Ten minutes after 3:00 a.m. on all days*

14. Use the `crontab` command to edit the crontab file for the root user. Add an entry that sends the message `It works!` to your current window five minutes from now. For example, if the current time is 10:25, make an entry in your `crontab` file for the 30th minute of the same hour.

```
tty
/dev/pts/2
date
Tue Mar 27 18:16:02 MDT 2007
crontab -e
```

*Add the following line, but substitute the correct time and terminal device:*

```
30 10 * * * /usr/bin/echo "It works!" > /dev/pts/2
```

*Save the file, and quit the vi edit session. In about five minutes, you should see the result in your window.*

---

## Module 15

---

# Solaris™ 10 Operating System Installation Requirements

---

## Objectives

Upon completing this module, you should be able to:

- Identify hardware requirements for Solaris 10 OS installation
- Identify the fundamentals of the Solaris 10 OS installations
- Identify the Solaris 10 OS software components
- Define guidelines to install the Solaris 10 OS from a DVD
- Describe the Secure by Default installation enhancement

# Identifying the Hardware Requirements for Solaris 10 OS Installation

A Solaris 10 OS installation requires the following:

For SPARC platform systems:

- 512 Mbytes of memory is recommended, 256 Mbytes minimum
- At least 5 Gbytes of disk space
- Access to a bootable DVD drive or an installation server

For x86/x64 platform systems:

- 120-MHz or faster processor is recommended. Hardware floating-point support is required.
- 512 Mbytes of memory is recommended, 256 Mbytes minimum.
- At least 5 Gbytes of disk space.
- Access to a bootable DVD drive or an installation server.
- If you are installing on a system that is not produced by Sun Microsystems, Inc., a system listed in the Hardware Compatibility List (HCL) found at: <http://www.sun.com/bigadmin/hcl>

## Two-Terabyte Disk Support

Starting with the Solaris 10 10/09 release, you can install and boot the Solaris OS from a disk that is up to 2 Tbytes in size. In previous Solaris releases, you could not install and boot the Solaris OS from a disk that was greater than 1 Tbyte in size.

When installing and booting the Solaris OS on a 2 Tbyte disk it must be connected to a system that runs a 64-bit kernel.

When booting an x86 system with a target disk greater than 1 Tbyte, a minimum of 1 Gbyte of memory is required.

# Identifying the Fundamentals of Solaris 10 OS Installations

This section introduces the various methods used to install or upgrade the Solaris OS.

## Solaris 10 OS Installation Options

There are two methods available to install the Solaris 10 OS on your system; `install-solaris` and Flash Archive installation. Procedures that use `install-solaris` install groups of software packages onto your system. Flash Archive installations use flash archives to install references configurations of the Solaris OS.

The following installation procedures use `install-solaris`:

- Solaris installation Graphical User Interface (GUI)
- Solaris installation Command Line Interpreter (CLI)
- Solaris Custom JumpStart™ software (JumpStart) installation

The following installation procedures use Flash archives:

- Solaris Flash Archive installation
- Solaris WAN boot installation

This module provides an overview of the Solaris OS installation concepts.

### GUI Installation

The default installation method uses a GUI if the host has sufficient memory and a display capable of displaying GUI screens. Otherwise, installation proceeds using the CLI. To run the installation GUI program requires a minimum of 768 Mbytes of memory.

These memory requirements relate to installation processes, and may represent values that are insufficient to support specific workloads. Memory requirements for systems should be determined according to anticipated production workloads.

On SPARC systems, initiate the GUI-based installation from DVD by issuing the following command:

```
ok boot cdrom
```

On x86/x64 systems, the GUI-based installation begins automatically when you insert the Solaris OS software DVD, power-on the system, and accept the default choices from the first two menus that display.

### CLI Installation

Hosts that do not have a graphical screen or that have insufficient memory cannot run the GUI installation, and consequently use the CLI. You can specify that you want to use the CLI when you start the installation process. Text based installation requires a minimum of 384 Mbytes of memory. On SPARC systems, using the `nowin` option when you boot from DVD initiates the CLI. For example:

```
ok boot cdrom - nowin
```

Windowing x86/x64 systems present sets of menus when you boot from DVD. Choosing the Solaris Interactive Text (Console session) option causes the installation software to use the CLI on the system console. Choosing the Solaris Interactive Text (Desktop session) option causes the installation software to use the CLI in a terminal window within the Java Desktop System (JDS).

When using the CLI, installation options display in a menu-driven format. Use the space bar to select options, and the F2 key (or the equivalent escape key sequence), to accept selected options.

### Custom JumpStart Installation

The Solaris JumpStart procedure installs Solaris OS software on a system through use of a user-defined profile found on a JumpStart server. You can customize profiles for different system types.

A JumpStart installation provides a hands-off installation across the network, and is based on a centrally-configured server. or set of servers. Custom JumpStart configurations allow you to incorporate shell scripts into the installation process, and customize software beyond simply installing the Solaris OS. Custom JumpStart shell scripts may perform both pre-installation and post-installation tasks. On a SPARC system, use the following command to initiate a custom JumpStart installation:

```
ok boot net - install
```

On a x86/x64 systems, you initiate a custom JumpStart installation by holding the F12 key during the initial power-on sequence. This causes the system to initiate the Intel Preboot Installation Environment (PXE) installation process, that requires a server to respond to DHCP and other requests made by the x86/x64 system. No boot media is required on the client system.

### Flash Archive Installation

The Solaris Flash Archive Installation enables you to install many systems based on a reference configuration that you create on a master system. After you have installed and configured the master system, you then create a flash archive from that system. You associate the flash archive with a JumpStart profile to allow JumpStart clients to use the flash archive as their installation source. You can create many flash archives to capture many master system configurations.

Solaris Flash Archive installations can offer speed advantages over custom JumpStart installations. Custom JumpStart installations install each Solaris OS package individually. Solaris Flash Archive installations use a single archive source, and avoid individual package installations.

You initiate Solaris Flash Archive Installations in the same way that you initiate custom JumpStart installations. The server configuration determines which type of installation occurs.

### Solaris WAN Boot

The WAN boot installation method enables you to boot and install software over a wide area network (WAN) using HTTP. The WAN boot installation method enables you to transmit an encrypted Solaris Flash archive over a public network to a remote SPARC-based client. The WAN boot programs then install the client system using custom JumpStart files and the Flash archive.

To protect the integrity of the installation, you can use private keys to authenticate and encrypt data. You can also transmit your installation data and files over a secure HTTP connection by configuring your systems to use digital certificates.

## Solaris 10 OS Upgrade Options

The two Solaris 10 OS Upgrade options are as follows:

- Standard Upgrade
- Solaris Live Upgrade

### Standard Upgrade

A standard upgrade merges the new version of the Solaris OS with the existing files on the system's disk. The methods available for a standard upgrade are Solaris GUI installation, the CLI installation, and the custom JumpStart procedure.

A standard upgrade saves many of the modifications that were made to the OS with the previous version of the Solaris OS. Because the Solaris OS is unavailable to users during the standard upgrade, the standard upgrade results in longer periods of downtime.

### Solaris Live Upgrade

Solaris Live Upgrade is the default method for upgrading your system in the recent releases of Solaris 10 OS. The Solaris Live Upgrade software upgrades a duplicate boot environment while the active boot environment is still running. This method eliminates downtime of the production environment. First, create a duplicate boot environment. After that has been created, upgrade or install a Solaris Web Start Flash archive on the inactive boot environment. When you are ready, activate the inactive boot environment. During the next reboot, the inactive boot environment becomes the active boot environment. If there is a failure, you can recover your original boot environment by reactivating it and rebooting the system.

## Troubleshooting Solaris 10 OS Upgrades

Check the log files to troubleshoot any problems encountered during an upgrade of Solaris 10 OS. The output files and upgrade process results are stored in log files, for example:

- `/var/sadm/system/logs/upgrade_cleanup`
- `/var/sadm/system/logs/upgrade_log`

Following an upgrade, log files help you identify problems and determine if any cleanup is needed.

# Identifying the Solaris 10 OS Software Components

Solaris OS software is organized into three components:

- Software packages
- Software clusters
- Software groups

## Software Packages

A software package contains a group of software files and directories. The package also contains the related software installation scripts.

## Software Clusters

During the software installation process, software clusters group logical collections of software packages together. Table 15-1 shows the software packages that are grouped into the CDE software cluster.

**Table 15-1** Packages Included in the CDE Software Cluster

SUNWdtwm	SUNWdthez	SUNWdtbas	SUNWdtab
SUNWdtdst	SUNWdtjxt	SUNWdtdmr	SUNWdthed
SUNWdtscm	SUNWpdas	SUNWdtdmn	SUNWdtinc
SUNWdthe	SUNWdtim	SUNWdtdte	SUNWdtmad
SUNWdthev	SUNWdtezt	SUNWdtlog	SUNWdtma
SUNWdticn	SUNWscgui	SUNWdtdem	SUNWdtmaz

Some software clusters contain only one software package.

## Solaris OS Software Groups

Software groups are collections of Solaris OS software packages. Each software group includes support for different functions and hardware drivers. The Solaris OS is made up of seven software groups:

- Reduced Network Support Software Group
- Core System Support Software Group
- End User Solaris Software Group
- Developer Solaris Software Group
- Entire Solaris Software Group
- Entire Solaris Software Group Plus Original Equipment Manufacturer (OEM) Support

### Reduced Network Support Software Group

The Reduced Network Support software group (SUNWCrnet) contains the minimum software required to boot and run a Solaris system with limited network service support. The Reduced Network Support software group provides a multiuser text-based console and system administration utilities. This software group also enables the system to recognize network interfaces, but does not activate network services.

A system installed with the Reduced Network Support software group could, for example, be used as a *thin-client* host in a network. At least 2 Gbytes of space is recommended to install this software group.

### Core System Support Software Group

The Core System Support software group (SUNWCreq) contains the minimum software required to boot and run the Solaris OS in a minimum configuration, without the support to run many server applications. The Core System Support software group includes a minimum of networking software, including Telnet, file transfer protocol (FTP), network file system (NFS), network information service (NIS) clients, and domain name service (DNS). This software group also includes the drivers required to run the Common Desktop Environment (CDE) but does not include the CDE software. The Core System Support software group also does not include online manual pages. At least 2 Gbytes of space is recommended to install this software group.

## End User Solaris Software Group

The End User Solaris software group (SUNWCuser) contains the Core System Support software group, the recommended software for an end user, and the CDE. At least 5.3 Gbytes of space is recommended to install this software group.

## Developer Solaris Software Group

The Developer Solaris software group (SUNWCprog) contains the End User Solaris software group. It also contains the libraries, the include files, the online manual pages, and the programming tools for developing software. At least 6.6 Gbytes of space is recommended to install this software group.

## Entire Solaris Software Group

The Entire Solaris software group (SUNWCa11) contains the Developer Solaris software group. It also contains additional server software. The Entire Solaris software group is the entire Solaris OS software release minus OEM support. At least 6.7 Gbytes of space is recommended to install this software group.

## Entire Solaris Software Group Plus OEM Support

The Entire Solaris Software Group Plus OEM Support group (SUNWCXa11) contains the entire Solaris OS software release. It also contains additional hardware support for OEMs and hardware not on the system at the time of installation. This software group is recommended when you are installing the Solaris OS software on non-Sun servers that use UltraSPARC processors. At least 6.8 Gbytes of space is recommended to install this software group.

To view the names of the cluster configurations, perform the command:

```
grep METACLUSTER /var/sadm/system/admin/.clustertoc
METACLUSTER=SUNWCXall
METACLUSTER=SUNWCall
METACLUSTER=SUNWCprog
METACLUSTER=SUNWCuser
METACLUSTER=SUNWCreq
METACLUSTER=SUNWCrnet
METACLUSTER=SUNWCmreq
```

To determine which cluster configuration has been installed on the system, perform the command:

```
cat /var/sadm/system/admin/CLUSTER
CLUSTER=SUNWCXall
```

# Defining Guidelines for Installing Solaris 10 OS From a DVD

Following are some of the guidelines to be followed while planning to install the Solaris 10 OS from a DVD:

- Allocate space in the layout for a Live Upgrade slice.
- Allocate additional disk space for each language that you install.
- Allocate additional space in the /var file system if you plan to have your system support printing or mail.
- Allocate double the amount of physical memory in the /var file system if you plan to use the crash dump feature savecore on your system.
- Allocate additional space in the /export or /export/home file system if you plan to provide a home directory file system for users.
- Allocate space for the Solaris OS software group you want to install.
- Allocate an additional 30 percent more disk space for each file system that you create, and create a minimum number of file systems. This leaves room for upgrades to future software releases.

---

**Note** – By default, Solaris OS installation methods create only the / (root) file system, /export/home, and swap partitions.

---

- Allocate additional disk space for additional software or third-party software.

Before installing the Solaris OS software on a networked stand-alone system, you must provide the following information:

Host name	Determine a unique, and usually, short name for the networked system. You can use the command <code>uname -n</code> to find the host name on an existing system.
Host Internet Protocol (IP) address	Determine the software address that represents the host address and network address. You can use the <code>ifconfig interface</code> command (for example, <code>ifconfig hme0</code> ) to display your current IP address.



Name service type	Determine if the networked system is to be included in one of the following types of name service domains: Lightweight Directory Access Protocol (LDAP), NIS, Network Information Service Plus (NIS+), DNS, or none.
Subnet mask	Determine if the networked system is included in a particular subnet. The subnet mask is stored in the /etc/netmasks file.

**Note** – A subnet is used to partition network traffic. Segmenting network traffic over many different subnets increases the bandwidth available to each host.

Geographic location and time zone	Determine the specific region where the networked system physically resides.
root password	Determine a password assigned to the root user. Use the root password to gain access to root privileges on the networked system.
Language	Determine the language with which to install the Solaris OS. The installation software enables the user to choose from a list of languages. Prompts, messages, and other installation information are displayed in the chosen language. The language choices include English, German, Spanish, French, Italian, Japanese, Korean, Swedish, Simplified Chinese, and Traditional Chinese.

Make sure you have the required Solaris installation media. You will need the Solaris Operating System for SPARC Platforms DVD, or the Solaris Operating System for x86 Platforms DVD.

Before performing a software installation, always back up any modifications or data that exist in the previous version of the Solaris OS. Restore them after completing the installation.

## Secure by Default Installation Enhancement

Traditionally, Solaris systems have provided many network services by default. This open approach is convenient, but it also makes it easier to exploit vulnerabilities that may exist in the software that provides network services.

The Secure by Default (SBD) enhancement reduces this exposure by disabling as many network services as possible while still providing a useful system. The only network-listening service left enabled in the hardened configuration is SSH.

### Installation and Secure by Default

During an initial DVD-based installation of the Solaris 10 OS, as of the 11/06 release, a GUI screen or text-based panel displays that allows you to select the Secure by Default option.

Following is an example of the text-based prompt that displays during a new Solaris 10 installation. Selecting "Yes" enables an open configuration profile. Selecting "No" enables a limited profile.

Would you like to enable network services for use by remote clients?

- Yes
- No

Note: Selecting "No" provides a more secure configuration in which Secure Shell is the only network service provided to remote clients. Selecting "Yes" enables a larger set of services as in previous Solaris releases. If in doubt, it is safe to select "No" as any services can be individually enabled after installation.

Custom JumpStart and Flash Archive servers use the new `service_profile` keyword in their `sysidcfg` files to establish either a limited or open network service profile on installation clients.

## Simulation: Solaris 10 Operating System Installation

This module uses a simulation to present an example of the installation process.

This simulation describes installing the Solaris 10 OS on both SPARC and x86/x64-based systems, using DVD installation media.

Your instructor will provide instructions for using the installation simulation.

## Notes:

---

## Appendix A

---

# Backing Up a Mounted File System With a UFS Snapshot

---

## Objectives

Upon completion of this module, you should be able to:

- Create a UFS snapshot
- Back up the snapshot file

---

**Note** – For information on ZFS snapshots, see the “ZFS File System Introduction” module.

---

# Creating a UFS Snapshot

The UFS Copy on Write Snapshots feature provides administrators an online backup solution for ufs file systems. This utility enables you to use a point-in-time copy of a ufs file system, called a snapshot, to create an online backup. You can create the backup while the file system is mounted and the system is in multiuser mode.



**Note** – The UFS snapshots are similar to the Sun StorEdge™ Instant Image product. Instant Image allocates space equal to the size of the entire file system that is being captured. However, the file system data saved by UFS snapshots occupies only as much disk space as needed.

## Using the fssnap Command

You use the fssnap command to create, query, or delete temporary read-only snapshots of ufs file systems.

The format for the fssnap command is:

```
/usr/sbin/fssnap -F FSType -V -o special_option(s) mount-point | special
```

Table 15-2 shows some of the options for the fssnap command.

**Table 15-2** fssnap Command Options

Option	Description
-d	Deletes the snapshot associated with the given file system. If the -o unlink option was used when you built the snapshot, the backing-store file is deleted together with the snapshot. Otherwise, the backing-store file (which contains file system data) occupies disk space until you delete it manually.
-F <i>FSType</i>	Specifies the file system type to be used.
-i	Displays the state of an <i>FSType</i> snapshot.
-V	Echoes the complete command line but does not execute the command.
-o	Enables you to use <i>special_options</i> , such as the location and size of the backing-store (bs) file.

To create a UFS snapshot, specify a backing-store path and the actual file system to be captured. The following is the command format:

```
fssnap -F ufs -o bs=backing_store_path /file-system
```



---

**Note** – The *backing\_store\_path* can be a raw device, the name of an existing directory, or the name of a file that does not already exist.

---

The following example uses the fssnap command to create a snapshot of the /export/home file system.

```
fssnap -F ufs -o bs=/var/tmp /export/home
/dev/fssnap/0
```

The snapshot subsystem saves file system data in a file called a backing-store file before the data is overwritten. Some important aspects of a backing-store file are:

- A backing-store file is a bit-mapped file that takes up disk space until you delete the UFS snapshot.
- The size of the backing-store file varies with the amount of activity on the file system being captured.
- The destination path that you specify on the fssnap command line must have enough free space to hold the backing-store file.
- The location of the backing-store file must be different from that of the file system you want to capture in a UFS snapshot.
- A backing-store file can reside on different types of file systems, including another ufs file system or a mounted nfs file system.

The fssnap command creates the backing-store file and two read-only virtual devices. The block virtual device, /dev/fssnap/0, can be mounted as a read-only file system. The raw virtual device, /dev/rfssnap/0, can be used for raw read-only access to a file system.

These virtual devices can be backed up with any of the existing Solaris OS backup commands. The backup created from a virtual device is a backup of the original file system when the UFS snapshot was taken.



**Note** – When a UFS snapshot is first created, the file system locks temporarily. Users might notice a slight pause when writing to this file system. The length of the pause increases with the size of the file system. There is no performance impact when users are reading from the file system.

---

## Limiting the Size of the Backing-Store File

Before creating a UFS snapshot, use the `df -k` command to check that the backing-store file has enough disk space to grow. The size of the backing-store file depends on how much data has changed since the previous snapshot was taken.

You can limit the size of the backing-store file by using the `-o maxsize=n` option of the `fssnap` command, where `n` (`k`, `m`, or `g`) is the maximum size of the backing-store file specified in Kbytes, Mbytes, or Gbytes.



**Caution** – If the backing-store file runs out of disk space, the system automatically deletes the UFS snapshot, which causes the backup to fail. The active `ufs` file system is not affected. Check the `/var/adm/messages` file for possible UFS snapshot errors.

---



**Note** – You can force an unmount of an active `ufs` file system, for which a snapshot exists (for example, with the `umount -f` command). This action deletes the appropriate snapshot automatically.

---

The following example creates a snapshot of the `/export/home` file system, and limits the backing-store file to 500 Mbytes.

```
fssnap -F ufs -o bs=/var/tmp,maxsize=500m /export/home
/dev/fssnap/0
```

## Displaying Information for a `ufs` File System Snapshot

Use either `fssnap` command to display UFS snapshot information. The following example displays a list of all the current UFS snapshots on the system. The list also displays the corresponding virtual device for each snapshot.

```
fssnap -i
0 /export/home
1 /usr
2 /database
```

Use the **-i** option to the `/usr/lib/fs/ufs/fssnap` command to display detailed information for a specific UFS snapshot that was created by the `fssnap` command.

The following example shows the details for the `/export/home` snapshot.

```
/usr/lib/fs/ufs/fssnap -i /export/home
Snapshot number : 0
Block Device : /dev/fssnap/0
Raw Device : /dev/rfssnap/0
Mount point : /export/home
Device state : idle
Backing store path : /var/tmp/snapshot0
Backing store size : 0 KB
Maximum backing store size : 512000 KB
Snapshot create time : Mon Oct 11 08:58:33 2004
Copy-on-write granularity : 32 KB
```

## Backing Up the UFS Snapshot File

The virtual devices that contain the UFS snapshot act as standard read-only devices, which enable you to back up the virtual device in the same manner as you would back up a file system.

### Performing a Backup of a UFS Snapshot

Use the `tar` command or the `ufsdump` command to back up a UFS snapshot.

#### Using the `tar` Command to Back Up a Snapshot File

If you use the `tar` command to back up the UFS snapshot, mount the snapshot before backing it up. The following procedure demonstrates how to do this type of mount.

1. Create the mount point for the block virtual device.  
`# mkdir -p /backups/home.bkup`
2. Mount the block virtual device to the mount point.  
`# mount -F ufs -o ro /dev/fssnap/0 /backups/home.bkup`
3. Change directory to the mount point.  
`# cd /backups/home.bkup`
4. Use the `tar` command to write the data to tape.  
`# tar cvf /dev/rmt/0 .`

#### Using the `ufsdump` Command

If you use the `ufsdump` command to back up a UFS snapshot, you can specify the raw virtual device during the backup.

```
ufsdump 0uf /dev/rmt/0 /dev/rfssnap/0
```

Verify that the UFS snapshot is backed up.

```
ufsrestore tf /dev/rmt/0
```

## Performing an Incremental Backup Using a UFS Snapshot

Incremental backups of snapshots contain files that were modified since the last UFS snapshot. Use the `ufsdump` command with the `N` option to create an incremental UFS snapshot, which writes the name of the device being backed up, rather than the name of the snapshot device to the `/etc/dumpdates` file.

The following example shows how to use `ufsdump` to create an incremental file system backup.



**Note** – Use the `N` argument when backing up a snapshot. This argument ensures proper updates to the `/etc/dumpdates` file.

```
ufsdump 1ufN /dev/rmt/0 /dev/rdsck/c1t0d0s0 /dev/rfssnap/0
```

Next you would verify that the UFS snapshot is backed up to tape.

```
fsrestore tf /dev/rmt/0
```

To understand incremental backups of snapshots, consider the following:

1. Create a snapshot of the `/extra` file system that is going to be backed up while the file system is mounted.

```
fssnap -o bs=/var/tmp /extra
/dev/fssnap/0
#
```

2. Verify that the snapshot was successful, and view detailed information about the snapshot.

```
fssnap -i
0 /extra
/usr/lib/fs/ufs/fssnap -i /extra
Snapshot number : 0
Block Device : /dev/fssnap/0
Raw Device : /dev/rfssnap/0
Mount point : /extra
Device state : idle
Backing store path : /var/tmp/snapshot0
Backing store size : 0 KB
Maximum backing store size: Unlimited
Snapshot create time : Mon Oct 11 10:34:21 2004
```

## Backing Up the UFS Snapshot File

Copy-on-write granularity : 32 KB

3. Make a directory that will be used to mount and view the snapshot data.

```
mkdir /extrasnap
#
```

4. Mount the snapshot to the new mount point, and compare the size of the file system and the snapshot device.

```
mount -o ro /dev/fssnap/0 /extrasnap
df -k |grep extra
/dev/dsk/c1t0d0s0 1294023 9 1242254 1% /extra
/dev/fssnap/0 1294023 9 1242254 1% /extrasnap
```

5. Edit a file under the /extra directory and make it larger, and then compare the size of the file system and the snapshot device.

```
vi file1
(yank and put text, or read text in from another file)
df -k |grep extra
/dev/dsk/c1t0d0s0 1294023 20 1242243 1% /extra
/dev/fssnap/0 1294023 9 1242254 1% /extrasnap
```

Observe that the file system grew in size while the snapshot file did not.

6. Perform a full backup with the N option of the ufsdump command.

```
ufsdump 0ufN /dev/rmt/0 /dev/rdsk/c1t0d0s0 /dev/rfssnap/0
DUMP: Writing 32 Kilobyte records
DUMP: Date of this level 0 dump: Mon Oct 11 10:49:38 2004
DUMP: Date of last level 0 dump: the epoch
DUMP: Dumping /dev/rfssnap/0 (sys41:/extrasnap) to /dev/rmt/0.
DUMP: Mapping (Pass I) [regular files]
DUMP: Mapping (Pass II) [directories]
DUMP: Estimated 262 blocks (131KB).
DUMP: Dumping (Pass III) [directories]
DUMP: Dumping (Pass IV) [regular files]
DUMP: Tape rewinding
DUMP: 254 blocks (127KB) on 1 volume at 1814 KB/sec
DUMP: DUMP IS DONE
DUMP: Level 0 dump on Mon Oct 11 11:03:46 2004
```

7. Verify the backup.

```
ufsrestore tf /dev/rmt/0
2 .
3 ./file1
4 ./file2
5 ./file3
6 ./file4
```

- ```

#
8. Unmount the back up device and remove the snapshot.
# umount /extrasnap
# fssnap -d /extra
# rm /var/tmp/snapshot0
#
9. Make some changes to the /extra file system, such as copying some
files, and then re-create the snapshot.
# cp file1 file5
# cp file1 file6
# fssnap -o bs=/var/tmp /extra
/dev/fssnap/0
#
10. Re-mount the snapshot device, and compare the size of the file
system and the snapshot device.

# mount -o ro /dev/fssnap/0 /extrasnap
# df -k |grep extra
/dev/dsk/c1t0d0s0    1294023      46 1242217    1%   /extra
/dev/fssnap/0          1294023      46 1242217    1%   /extrasnap
#
11. Perform an incremental backup with the N option of the ufsdump
command.

# ufsdump 1ufN /dev/rmt/0 /dev/rdsck/c1t0d0s0 /dev/rfssnap/0
DUMP: Writing 32 Kilobyte records
DUMP: Date of this level 0 dump: Mon Oct 11 13:13:03 2004
DUMP: Date of last level 0 dump: Mon Oct 11 12:30:44 2004
DUMP: Dumping /dev/rfssnap/0 (sys41:/extrasnap) to /dev/rmt/0.
DUMP: Mapping (Pass I) [regular files]
DUMP: Mapping (Pass II) [directories]
DUMP: Estimated 294 blocks (147KB).
DUMP: Dumping (Pass III) [directories]
DUMP: Dumping (Pass IV) [regular files]
DUMP: Tape rewinding
DUMP: 254 blocks (127KB) on 1 volume at 1693 KB/sec
DUMP: DUMP IS DONE
DUMP: Level 1 dump on Mon Oct 11 13:22:36 2004
#
12. Verify the backup.

# ufsrestore tf /dev/rmt/0
2      .
7      ./file5
8      ./file6

```

```
#
```

The backup of the snapshot contains only the files that were added since the previous Level 0 backup.

Restoring Data From a UFS Snapshot Backup

The backup created from a virtual device is a backup of the original file system when the UFS snapshot was taken.

Restore a UFS snapshot from a backup tape in the same manner as you would the backup of an original file system. Data written to a tape by `ufsdump` is simply data, whether it is a snapshot or a file system.

To restore the `demo` directory from the snapshot backup of the `/usr` file system, complete the following steps:

1. Load the tape that contains the snapshot backup of the `/usr` file system into the tape drive.

2. Change to the `/usr` file system.

```
# cd /usr
```

3. Perform the a `ufsrestore` command.

```
# ufsrestore if /dev/rmt/0
```

```
ufsrestore > add demo
```

```
ufsrestore > extract
```

```
Specify next volume #: 1
```

```
set owner/mode for '.'? [yn] n
```

```
ufsrestore > quit
```

4. Verify that the `demo` directory exists, and eject the tape.

Deleting a UFS Snapshot

Deleting a UFS snapshot from the system is a multistep process and order-dependant. First, unmount the snapshot device, and then delete the snapshot. Finally, remove the backing-store file.

```
# umount /dev/fssnap/0
# fssnap -d /export/home
# rm /backing_store_file
```

Exercise: Working With UFS Snapshots

In this exercise, you complete the following tasks:

- Create a UFS snapshot of the /opt file system
- View the contents of the backing-store directory
- Display detailed information for the UFS snapshot
- Remove the snapshot and backing-store file

Preparation

Refer to the lecture notes as necessary to perform the tasks listed.

Task

Complete the following steps:

1. Create a snapshot of the /opt file system without specifying a file name for the backing-store file.
2. List the contents of the /var/tmp directory.

What is the default name assigned to the backing-store file?

3. Display the detailed information about the snapshot.

What is the maximum backing-store file size for the snapshot?

4. Delete the snapshot from the system.

5. List the contents of the /var/tmp directory. Has the backing-store file been removed?
-

6. Remove the backing-store file you created in Step 1.

Exercise Summary



Discussion – Take a few minutes to discuss what experiences, issues, or discoveries you had during the lab exercises.

- Experiences
- Interpretations
- Conclusions
- Applications

Exercise Solutions: Working With UFS Snapshots

This section contains solutions to the exercise.

Task

Complete the following steps:

1. Create a snapshot of the /opt file system without specifying a file name for the backing-store file.

```
# fssnap -F ufs -o bs=/var/tmp /export/home
/dev/fssnap/0
#
```

2. List the contents of the /var/tmp directory.

```
# ls -lh /var/tmp
total 862
(output omitted)
drwx----- 2 root      root          1.0K Oct  6 15:53 orbit-root/
drwx----- 2 1000     1000         1.0K Sep 25 14:03 orbit-student/
drwx----- 2 user5    staff         1.0K Oct  6 14:46 orbit-user5/
-rw----- 1 root      root        9.8G Oct  6 16:30 snapshot0
-rw----- 1 root      root        2.2K Oct  5 18:44 wscon-:0-M4ayyb
(output omitted)
```

What is the default name assigned to the backing-store file?

snapshot0

3. Display the detailed information about the snapshot.

```
# /usr/lib/fs/ufs/fssnap -i /export/home
Snapshot number           : 0
Block Device              : /dev/fssnap/0
Raw Device                : /dev/rfssnap/0
Mount point               : /export/home
Device state              : idle
Backing store path        : /var/tmp/snapshot0
Backing store size        : 0 KB
Maximum backing store size: Unlimited
Snapshot create time      : Tue Oct  6 16:30:04 2009
Copy-on-write granularity : 32 KB
```

What is the maximum backing-store file size for the snapshot?

Unlimited

4. Delete the snapshot from the system.

```
# fssnap -d /export/home  
Deleted snapshot 0.  
#
```

5. List the contents of the /var/tmp directory. Has the backing-store file been removed?

```
# ls -l /var/tmp/snapshot0  
-rw----- 1 root      root      10574196736 Oct  6 16:30  
/var/tmp/snapshot0  
#
```

No

6. Remove the backing-store file you created in Step 1.

```
# rm /var/tmp/snapshot0
```

Appendix B

Performing File System Backups

Objectives

Upon completion of this module, you should be able to:

- Identify backup fundamentals
- Back up an unmounted file system

Introducing Backup Fundamentals

Backing up file systems safeguards against data loss, damage, or corruption. Backup tapes are often referred to as dump tapes.

Routine File System Backups

To back up file systems, copy them to removable media, such as a tape. Perform backups regularly to prevent data loss from:

- Accidental file deletion
- Hardware failures
- Problems with re-installations or system upgrades
- System crashes
- System break-ins
- Natural disasters

Tape Drive Naming

All tape drives have logical device names that you use to reference the device on the command line. Figure 15-1 shows the format that all logical device names use.

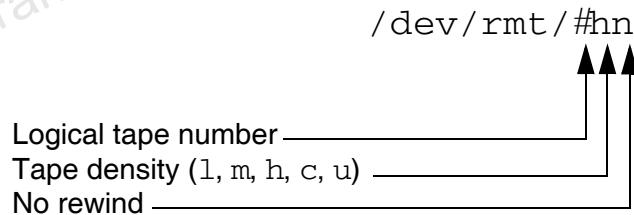


Figure 15-1 Logical Device Name Format

The logical tape numbers in the tape drive names always start with 0. For example:

- The first instance of a tape drive:
`/dev/rmt/0`
- The second instance of a tape drive:
`/dev/rmt/1`

- The third instance of a tape drive:

/dev/rmt/2

Two optional parameters further define the logical device name:

- Tape density – Five values can be given in the tape device name: l (low), m (medium), h (high), c (compressed), or u (ultra compressed).



Note – The tape density optional parameters h, c, and u reflect the same density. Refer to the man page `st(7D)` for more information.

- No rewind – The letter n at the end of a tape device name indicates that the tape should not be rewound when the current operation completes.

Tape densities depend on the tape drive. Check the manufacturer documentation to determine the correct densities for the tape media.



Note – The `st.conf` file contains configuration entries for supported devices of Sun Microsystems and many third-party devices.

Tape drives that support data compression contain internal hardware that performs the compression. If you back up a software-compressed file to a tape drive with hardware compression, the resulting file may be larger in size.

Tape Drive Control

Use the magnetic tape control `mt` command to send instructions to the tape drive. Not all tape drives support all `mt` commands.

The format for the `mt` command is:

`mt -f tape-device-name command count`

Use the `-f` option to specify the tape device name, typically a no-rewind device name. If no `-f` option is used, the default tape device file `/dev/rmt/0` is used.

Using the `mt` Command

Table 15-3 lists some of the `mt` commands that you can use to control a magnetic tape drive.

Table 15-3 Definitions of `mt` Commands

| Command | Definition |
|---------------------------|---|
| <code>mt status</code> | Displays status information about the tape drive |
| <code>mt rewind</code> | Rewinds the tape |
| <code>mt offline</code> | Rewinds the tape and, if appropriate, takes the drive unit offline and if the hardware supports it, unloads |
| <code>mt fsf count</code> | Moves the tape forward <i>count</i> records |

Assuming the tape was rewound to the start of tape, the following command positions the tape at the beginning of the third tape record.

```
# mt -f /dev/rmt/0n fsf 2
```

Scheduled Backup Strategies

Perform cumulative incremental backups daily.

To set up a backup schedule, determine:

- The file systems to back up
- A backup device (for example, tape drive)
- The number of tapes to use for the backup
- The type of backup (for example, full or incremental)
- The procedures for marking and storing tapes
- The time it takes to perform a backup

Determining File System Names to Back Up

1. Display the contents of the `/etc/vfstab` file.
2. View the `mount point` column to find the name of the file system that you want to back up.



Determining the Number of Tapes

1. Determine the number of tapes for a backup according to the size of the file system you are backing up.
2. Use `ufsdump`, with the `S` option, to determine the file system size as follows. The numeric option determines the appropriate dump level. The output is the estimated number of bytes that the system requires for a complete backup.

```
# ufsdump 0S filesystem_name
<number reported>
or
# ufsdump 3S filesystem_name
<number reported>
```

Note – To translate the output of the `ufsdump #S filesystem_name` to Mbytes, use the following calculation:

`# / (1024 x 1024)`.

3. Divide the reported bytes by the capacity of the tape to determine how many tapes you need to backup the file system.

Determining Back Up Frequency and Levels

Determine how often and at what level to backup each file system. The level of a backup refers to the amount of information that is backed up.

Identifying Incremental and Full Backups

You can perform a full backup or an incremental backup of a file system. A full backup is a complete file system backup. An incremental backup copies only files in the file system that have been added or modified since a previous lower-level backup.

Use dump level 0 to perform a full backup. You use dump levels 1 through 9 to schedule incremental backups. The level numbers have no meaning other than their relationship to each other as a higher or lower number.

Figure 15-2 shows *an example* of a file system backup performed in incremental levels.

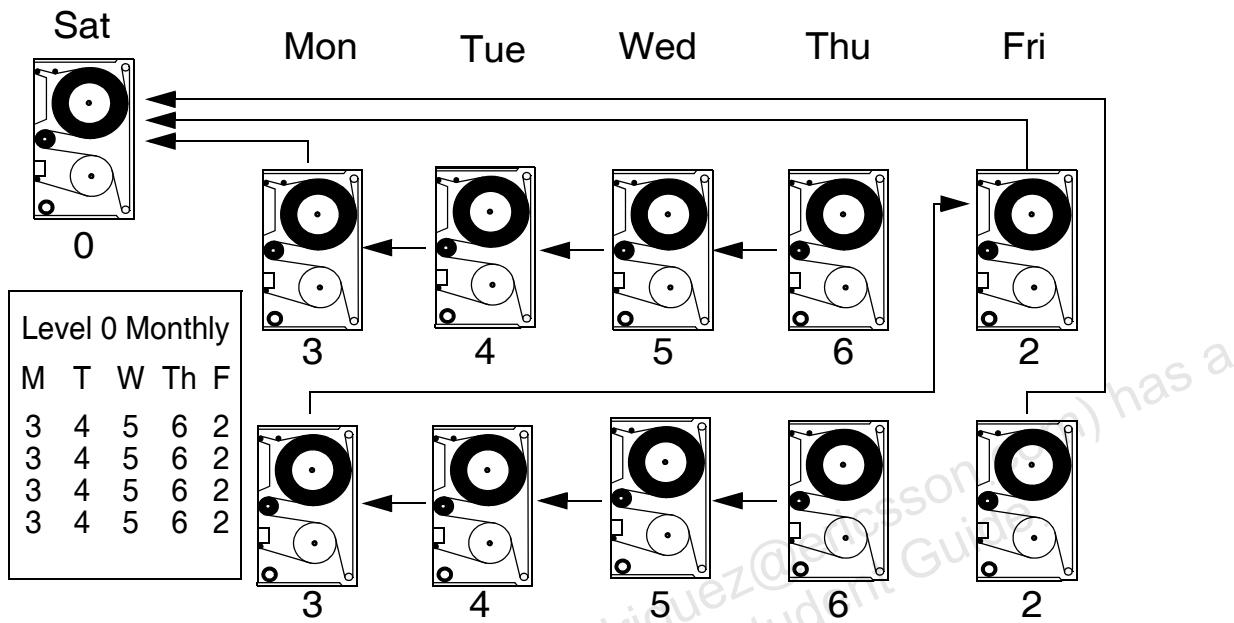


Figure 15-2 Incremental Backup Strategy Sample

Table 15-4 defines the elements of the sample incremental backup strategy shown in Figure 15-2 on page B-6. Use crontab to start a script that runs

Table 15-4 Incremental Back Up Level Definitions

Level	Example
0 (Full)	Performed once each month.
3	Performed every Monday. The backup copies new or modified files since the last lower-level backup (for example, 0).
4	Performed every Tuesday. The backup copies new or modified files since the last lower-level backup (for example, 3).
5	Performed every Wednesday. The backup copies new or modified files since the last lower-level backup (for example, 4).
6	Performed every Thursday. The backup copies new or modified files since the last lower-level backup (for example, 5).
2	Performed every Friday. The backup copies new or modified files since the last lower-level backup, which is the Level 0 backup at the beginning of the month.

The /etc/dumpdates File

Use `ufsdump -u`, to backup the `/etc/dumpdates` file records. Each line in the `/etc/dumpdates` file shows the file system that was backed up and the level of the last backup. It also shows the day, the date, and the time of the backup. The following is an example `/etc/dumpdates` file:

```
# cat /etc/dumpdates
/dev/rdsck/c0t2d0s6 0 Fri Nov 5 19:12:27 2004
/dev/rdsck/c0t2d0s0 0 Fri Nov 5 20:44:02 2004
/dev/rdsck/c0t0d0s7 0 Tue Nov 9 09:58:26 2004
/dev/rdsck/c0t0d0s7 1 Tue Nov 9 16:25:28 2004
```

When you perform an incremental backup, the `ufsdump` command consults the `/etc/dumpdates` file. It looks for the date of the next lower-level backup. Then, the `ufsdump` command copies to the backup media all of the files that were modified or added since the date of that lower-level backup.

When the backup is complete, the `/etc/dumpdates` file records a new entry that describes this backup. The new entry replaces the entry for the previous backup at that level.

View the `/etc/dumpdates` file to determine if the system is completing backups. If a backup does not complete because of equipment failure, the `/etc/dumpdates` file does not record the backup.



Note – When you restore an entire file system, check the `/etc/dumpdates` file for a list of the most recent dates and levels of backups. Use this list to determine which tapes are needed to restore the entire file system. The tapes should be physically marked with the dump level and date of the backup.

Backing Up an Unmounted File System

Determine if the file system is inactive, or unmounted, before you back up the file system. If the file system is active, the output of the backup can be inconsistent, and it may be impossible to restore some of the files correctly.

The `ufsdump` Command

Use `/usr/sbin/ufsdump` to back up a complete or a partial ufs file system. (Backups are often referred to as dumps.) The `ufsdump` command format is:

`ufsdump option(s) argument(s) filesystem_name`

ufsdump Command Options

Table 15-5 defines common options for the `ufsdump` command.

Table 15-5 `ufsdump` Command Options

Option	Description
0–9	Back up level. Level 0 is a full backup of the file system. Levels 1 through 9 are incremental backups of files that have changed since the last lower-level backup. When no backup level is given, the default is level 9.
v	Verify. After each tape is written, the system verifies the contents of the media against the source file system. If any discrepancies occur, the system prompts the operator to insert new media and repeat the process. Use this option only on an unmounted file system. Any activity in the file system causes the system to report discrepancies.
s	Size estimate. This option allows you to estimate the amount of space that will be needed on the tape to perform the level of backup you want.
l	Autoload. You use this option with an autoloading (stackloader) tape drive.
o	Offline. When the backup is complete, the system takes the drive offline, rewinds the tape (if you use a tape), and, if possible, ejects the media.

Table 15-5 ufsdump Command Options

Option	Description
u	Update. The system creates an entry in the /etc/dumpdates file with the device name for the file system disk slice, the backup level (0–9), and the date. If an entry already exists for a backup at the same level, the system replaces the entry.
n	Notify. The system sends messages to the terminals of all logged-in users who are members of the sys group to indicate that the ufsdump command requires attention.
f <i>device</i>	Specify. The system specifies the device name of the file system backup. When you use the default tape device, /dev/rmt/0, you do not need the -f option. The system assumes the default.

The ufsrestore Utility

The ufsrestore utility restores files from backup media created with the ufsdump command. ufsrestore actions are controlled by the key argument. The key is exactly one function letter (i, r, R, t, or x) and zero or more function modifiers (letters). The key string contains no SPACE characters. Function modifier arguments are listed on the command line in the same order as their corresponding function modifiers appear in the key string. For more information about the ufsrestore utility, see the ufsrestore man page.

Tape Backups

Use ufsdump to create file system backups to tape. Enter a dump level (0–9) to determine which files to back up.

Note – To backup file systems that you cannot unmount, such as the root (/) file system, it is recommended to shut down the system to single-user mode. For other file systems, unmounting the file system is sufficient.



Using the `ufsdump` Command

Perform the following steps to use the `ufsdump` command to back up an example file system (`/opt`) to tape:

1. Become the root user, and use the `shutdown` command to change the system to single-user mode.

```
# /usr/sbin/shutdown -y -g 300 "System is being shutdown for backup"
```

Shutdown started. Wed Mar 28 09:17:53 MDT 2007

```
Changing to init state s - please wait
Broadcast Message from root (console) on sys-05 Wed Mar 28 09:17:53...
THE SYSTEM sys-05 IS BEING SHUT DOWN NOW ! ! !
Log off now or risk your files being damaged
System is being shutdown for backup
(output omitted)
```

2. Check if the `/opt` file system is still mounted. If it is, unmount it manually.

```
# mount | grep /opt
/opt on /dev/dsk/c0t0d0s5
read/write/setuid/devices/intr/largefiles/logging/xattr/onerror=panic/dev
=220000d on Mon Mar  5 10:17:29 2007
# umount /opt
#
```

3. Check the integrity of the file system data with the `fsck` command.

```
# fsck /opt
** /dev/rdsk/c0t0d0s5
** Last Mounted on /opt
** Phase 1 - Check Blocks and Sizes
** Phase 2 - Check Pathnames
** Phase 3a - Check Connectivity
** Phase 3b - Verify Shadows/ACLs
** Phase 4 - Check Reference Counts
** Phase 5 - Check Cylinder Groups
320 files, 35595 used, 444924 free (436 frags, 55561 blocks, 0.1%
fragmentation)
#
```

4. Perform a full (Level 0) backup of the `/opt` file system.

```
# ufsdump 0uf/dev/rmt/0 /opt
DUMP: Date of this level 0 dump: Wed Mar 28 09:42:05 2007
DUMP: Date of last level 0 dump: the epoch
DUMP: Dumping /dev/rdsk/c0t0d0s5 (sys-05:/opt) to /dev/rmt/0.
DUMP: Mapping (Pass I) [regular files]
```

```
DUMP: Mapping (Pass II) [directories]
DUMP: Writing 32 Kilobyte records
DUMP: Estimated 71686 blocks (35.00MB).
DUMP: Dumping (Pass III) [directories]
DUMP: Dumping (Pass IV) [regular files]
DUMP: Tape rewinding
DUMP: 71678 blocks (35.00MB) on 1 volume at 811 KB/sec
DUMP: DUMP IS DONE
DUMP: Level 0 dump on Wed Mar 28 09:42:05 2007
#
```

5. Return the system to multi-user mode.

```
# exit
svc.startd: Returning to milestone all.
checking ufs filesystems
/dev/rdsck/c0t0d0s5: is logging.
(output omitted)
```

Remote Backups to a Tape

Use the `ufsdump` command to perform a backup to a remote tape device. The format for the `ufsdump` command is:

```
ufsdump options remotehost:tapedevice filesystem
```

To perform remote backups across the network, the system with the tape drive must have an entry in its `/.rhosts` file for every system that uses the tape drive.

Using the `ufsdump` Command

The following example shows how to perform a full (Level 0) backup of the `/opt` file system on the `sys-06` system, to the remote tape device on the `sys-05` system.

```
sys-06# ufsdump 0uf sys-05:/dev/rmt/0 /opt
DUMP: Date of this level 0 dump: Wed Mar 28 09:51:22 2007
DUMP: Date of last level 0 dump: the epoch
DUMP: Dumping /dev/rdsck/c1d0s5 (sys-06:/opt) to sys-05:/dev/rmt/0.
DUMP: Mapping (Pass I) [regular files]
DUMP: Mapping (Pass II) [directories]
DUMP: Writing 32 Kilobyte records
DUMP: Estimated 68668 blocks (33.53MB).
DUMP: Dumping (Pass III) [directories]
DUMP: Dumping (Pass IV) [regular files]
```

Backing Up an Unmounted File System

```
DUMP: Tape rewinding
DUMP: 68606 blocks (33.50MB) on 1 volume at 1455 KB/sec
DUMP: DUMP IS DONE
DUMP: Level 0 dump on Wed Mar 28 09:51:22 2007
sys-06#
```

Backups to Disk

Use `ufsdump` to create file system backups to files on disk. As with backups to tape, the dump level (0–9) specified in the `ufsdump` command determines which files to back up.

Use backups to disk if you need to create a backup quickly, for temporary use, or if you do not have a tape device available. A backup made to a file on disk is only as reliable as the disk it is on.

Use the `ufsdump -f` option to specify the name of the file that will receive the backup.

Using the `ufsdump` Command

Perform the following steps to use the `ufsdump` command to back up an example file system (`/opt`) to a file on disk:

1. Become the root user, and use the `shutdown` command to change the system to single-user mode.

```
# /usr/sbin/shutdown -y -g 300 "System is being shutdown for backup"
```

```
Shutdown started.      Wed Mar 28 09:17:53 MDT 2007
```

```
Changing to init state s - please wait
Broadcast Message from root (console) on sys-05 Wed Mar 28 09:17:53...
THE SYSTEM sys-05 IS BEING SHUT DOWN NOW ! ! !
Log off now or risk your files being damaged
System is being shutdown for backup
(output omitted)
```

2. Check if the `/opt` file system is still mounted. If it is, unmount it manually.

```
# mount | grep /opt
```

```
/opt on /dev/dsk/c0t0d0s5
read/write/setuid/devices/intr/largefiles/logging/xattr/onerror=panic/dev
=220000d on Mon Mar  5 10:17:29 2007
# umount /opt
```

#

3. Check the integrity of the file system data with the fsck command.

```
# fsck /opt
** /dev/rdsk/c0t0d0s5
** Last Mounted on /opt
** Phase 1 - Check Blocks and Sizes
** Phase 2 - Check Pathnames
** Phase 3a - Check Connectivity
** Phase 3b - Verify Shadows/ACLs
** Phase 4 - Check Reference Counts
** Phase 5 - Check Cylinder Groups
320 files, 35595 used, 444924 free (436 frags, 55561 blocks, 0.1%
fragmentation)
#
```

4. Check the space used on the file system you intend to back up, and the file system that will hold the backup file. Verify that enough space is available in the destination file system to hold the backup.

```
# df -h /opt /export/home
Filesystem           size   used  avail capacity  Mounted on
/dev/dsk/c0t0d0s5     470M   36M   387M    9%      /opt
/dev/dsk/c0t0d0s7      94G   64M   93G     1%      /export/home
#
```

5. Perform a full (Level 0) backup of the /opt file system to a file in the /export/home file system.

```
# ufsdump 0uf /export/home/dump_opt_level_0 /opt
DUMP: Date of this level 0 dump: Wed Mar 28 10:00:59 2007
DUMP: Date of last level 0 dump: the epoch
DUMP: Dumping /dev/rdsk/c0t0d0s5 (sys-05:/opt) to
/export/home/dump_opt_level_0.
DUMP: Mapping (Pass I) [regular files]
DUMP: Mapping (Pass II) [directories]
DUMP: Writing 32 Kilobyte records
DUMP: Estimated 71686 blocks (35.00MB).
DUMP: Dumping (Pass III) [directories]
DUMP: Dumping (Pass IV) [regular files]
DUMP: 71678 blocks (35.00MB) on 1 volume at 9977 KB/sec
DUMP: DUMP IS DONE
DUMP: Level 0 dump on Wed Mar 28 10:00:59 2007
#
```

6. Return the system to multi-user mode.

```
# exit
svc.startd: Returning to milestone all.
checking ufs filesystems
```

Backing Up an Unmounted File System

```
/dev/rdsk/c0t0d0s5: is logging.  
(output omitted)
```

Saving ZFS Data With Other Backup Products

In addition to the `zfs send` and `zfs receive` commands, you can use archive utilities, such as the `tar` and `cpio` commands, to save ZFS files. These utilities save and restore ZFS file attributes and ACLs. Check the appropriate options for both the `tar` and `cpio` commands.

For recent information about issues with ZFS and third-party backup products, see the Solaris 10 release notes or the ZFS FAQ, here:

<http://opensolaris.org/os/community/zfs/faq/#backupsoftware>

Exercise: Backing Up a File System to Tape

In this exercise, you complete the following tasks:

- Backup an available file system on your system to a tape drive

Preparation

This exercise requires a system that is configured with a tape drive and file system that is available to unmount. This exercise assumes that the /export/home file system exists on a separate partition from the / (root) file system and can be unmounted. Identify the slice on which the /export/home file system resides. Get a tape that is appropriate for your system from the instructor.

For classrooms without a tape drive, skip to the exercise: Backing Up a File System to an Extra Disk Drive.

Task

Complete the following steps:.

1. Recursively copy the contents of the /opt/ses directory to /export/home.
2. Unmount the /export/home file system. If your system reports that the /export/home file system is busy, use the `umount -f` command.
3. Insert a tape into your tape drive.
4. Use the `mt` command to rewind the tape to the beginning.
5. Use the `ufsdump` command to create a backup tape for the /export/home file system.
6. Mount the /export/home file system.
7. Copy the contents of the /etc/uucp directory to the /export/home directory.
8. Unmount the /export/home file system.
9. Move the tape to the next tape record.
10. Use the `ufsdump` command to create an incremental backup for the /export/home file system, using the non-rewinding tape device.
11. Rewind and eject the tape from the tape drive.

12. Review the contents of the /etc/dumpdates file.
13. Mount the /export/home file system.

Exercise Summary



Discussion – Take a few minutes to discuss what experiences, issues, or discoveries you had during the lab exercises.

- Experiences
- Interpretations
- Conclusions
- Applications

Exercise Solutions: Backing Up a File System to Tape

This section contains solutions to the exercise.

Task

Complete the following steps:

1. Recursively copy the contents of the /opt/ses directory to /export/home.
cp -r /opt/ses /export/home
2. Unmount the /export/home file system. If your system reports that the /export/home file system is busy, use the `umount -f` command.
umount /export/home
3. Insert a tape into your tape drive.
4. Use the `mt` command to rewind the tape to the beginning.
mt rewind
5. Use the `ufsdump` command to create a backup tape for the /export/home file system.

```
# ufsdump 0uf/dev/rmt/0 /export/home
```

You should see output similar to this:

```
DUMP: Date of this level 0 dump: Wed Mar 28 10:29:31 2007
DUMP: Date of last level 0 dump: the epoch
DUMP: Dumping /dev/rdsck/c0t0d0s7 (sys-05:/export/home) to /dev/rmt/0.
DUMP: Mapping (Pass I) [regular files]
DUMP: Mapping (Pass II) [directories]
DUMP: Writing 32 Kilobyte records
DUMP: Estimated 12130 blocks (5.92MB).
DUMP: Dumping (Pass III) [directories]
DUMP: Dumping (Pass IV) [regular files]
DUMP: Tape rewinding
DUMP: 12094 blocks (5.91MB) on 1 volume at 723 KB/sec
DUMP: DUMP IS DONE
DUMP: Level 0 dump on Wed Mar 28 10:29:31 2007
```

```
#
```

6. Mount the /export/home file system.
mount /export/home
7. Copy the contents of the /etc/uucp directory to the /export/home directory.

- # **cp -r /etc/uucp /export/home**
8. Unmount the /export/home file system.
umount /export/home
 9. Move the tape to the next tape record.
mt -f /dev/rmt/0n fsf 1
 10. Use the **ufsdump** command to create an incremental backup for the /export/home file system, using the non-rewinding tape device.
ufsdump 1uf /dev/rmt/0n /export/home

You should see output similar to this:

```
DUMP: Date of this level 1 dump: Wed Mar 28 10:31:39 2007
DUMP: Date of last level 0 dump: Wed Mar 28 10:29:31 2007
DUMP: Dumping /dev/rdsck/c0t0d0s7 (sys-05:/export/home) to /dev/rmt/0n.
DUMP: Mapping (Pass I) [regular files]
DUMP: Mapping (Pass II) [directories]
DUMP: Writing 32 Kilobyte records
DUMP: Estimated 6072 blocks (2.96MB).
DUMP: Dumping (Pass III) [directories]
DUMP: Dumping (Pass IV) [regular files]
DUMP: 6078 blocks (2.97MB) on 1 volume at 2284 KB/sec
DUMP: DUMP IS DONE
DUMP: Level 1 dump on Wed Mar 28 10:31:39 2007
```

#

11. Rewind and eject the tape from the tape drive.
mt -f /dev/rmt/0 offline
12. Review the contents of the /etc/dumpdates file.
more /etc/dumpdates

You should see one line showing information for the level 0 dump and another line for the level 1 dump, for example:

/dev/rdsck/c0t0d0s7	0 Wed Mar 28 10:29:31 2007
/dev/rdsck/c0t0d0s7	1 Wed Mar 28 10:31:39 2007

13. Mount the /export/home file system.
mount /export/home

Exercise: Backing Up a File System to an Extra Disk Drive

In this exercise, you complete the following tasks:

- Backup an available file system on your system to an extra disk drive

Preparation

This exercise requires a system that is configured with an extra disk drive and file system that is available to unmount. This exercise assumes that the /export/home file system exists on a separate partition from the / (root) file system and can be unmounted. Identify the slice on which the /export/home file system resides.

In this exercise, you use the format utility to partition your extra hard drive so that it has one partition that uses all of the disk space. You then create a file system on that partition, and mount the file system.

Task

Complete the following steps:

1. Use the format utility to partition the extra hard drive and create one large partition to use as a backup device. Use the All Free Hog method to set partition 0 to use all disk space.
2. Create a file system on the large partition on the extra disk.
3. Run a file system check on the new file system.
4. Create a mount point directory and mount the new file system.
5. Recursively copy the contents of the /opt/ses directory to /export/home.
6. Unmount the /export/home file system. If your system reports that the /export/home file system is busy, use the `umount -f` command.
7. Use the `ufsdump` command to backup the /export/home file system to a file below the newly created mount point. Review the man page for `ufsdump` for information on how to use a dump file instead of a tape device.
8. Mount the /export/home file system.
9. Copy the contents of the /etc/uucp directory to the /export/home directory.

Exercise: Backing Up a File System to an Extra Disk Drive

10. Unmount the /export/home file system.
11. Use the `ufsdump` command to create an incremental backup for the /export/home file system. Save this data in a separate file on the spare disk.
12. Review the contents of the /etc/dumpdates file.
13. Mount the /export/home file system.

Exercise Summary



Discussion – Take a few minutes to discuss what experiences, issues, or discoveries you had during the lab exercises.

- Experiences
- Interpretations
- Conclusions
- Applications

Exercise Solutions: Backing Up File System to Extra Disk Drive

This section contains solutions to the exercise.

Task

Complete the following steps:

1. Use the format utility to partition the extra hard drive and create one large partition to use as a backup device. Use the All Free Hog method to set partition 0 to use all disk space.

This is an example from a SPARC system:

```
sys-05# format
Searching for disks...done
```

AVAILABLE DISK SELECTIONS:

- 0. c0t0d0 <HDS722512VLAT80 cyl 57459 alt 2 hd 16 sec 255>
/pci@1e,600000/ide@d/dad@0,0
- 1. c0t1d0 <DEFAULT cyl 57459 alt 2 hd 16 sec 255>
/pci@1e,600000/ide@d/dad@1,0

Specify disk (enter its number) : **1**
selecting c0t1d0

When you are through partitioning the extra disk drive, your partition table should look something like this:

Part	Tag	Flag	Cylinders	Size	Blocks
0	root	wm	0 - 57458	111.79GB	(57459/0/0) 234432720
1	swap	wu	0	0	(0/0/0) 0
2	backup	wu	0 - 57458	111.79GB	(57459/0/0) 234432720
3	unassigned	wm	0	0	(0/0/0) 0
4	unassigned	wm	0	0	(0/0/0) 0
5	unassigned	wu	0	0	(0/0/0) 0
6	unassigned	wu	0	0	(0/0/0) 0
7	unassigned	wu	0	0	(0/0/0) 0

This is an example from an x86/x64 system:

```
sys-05# format
Searching for disks...done
```

AVAILABLE DISK SELECTIONS:

- 0. c1d0 <DEFAULT cyl 39887 alt 2 hd 64 sec 63>
/pci@0,0/pci-ide@7/ide@0/cmdk@0,0

```

1. c2d0 <DEFAULT cyl 10008 alt 2 hd 255 sec 63>
   /pci@0,0/pci-ide@7/ide@1/cmdk@0,0
Specify disk (enter its number): 1
selecting c2d0

```

When you are through partitioning the extra disk drive, your partition table should look something like this:

Part	Tag	Flag	Cylinders	Size	Blocks
0	reserved	wm	3 - 10007	76.64GB	(10005/0/0) 160730325
1	swap	wu	0	0	(0/0/0) 0
2	backup	wu	0 - 10007	76.67GB	(10008/0/0) 160778520
3	unassigned	wu	0	0	(0/0/0) 0
4	unassigned	wu	0	0	(0/0/0) 0
5	unassigned	wu	0	0	(0/0/0) 0
6	unassigned	wu	0	0	(0/0/0) 0
7	unassigned	wu	0	0	(0/0/0) 0
8	boot	wu	0 - 0	7.84MB	(1/0/0) 16065
9	alternates	wm	1 - 2	15.69MB	(2/0/0) 32130

2. Create a file system on the large partition on the extra disk. For example, on a SPARC system:


```
# newfs /dev/rdsk/c0t1d0s0
```

 On an x86/x64 system:


```
# newfs /dev/rdsk/c2d0s0
```
3. Run a file system check on the new file system. For example, on a SPARC system:


```
# fsck /dev/rdsk/c0t1d0s0
```

 On an x86/x64 system:


```
# fsck /dev/rdsk/c2d0s0
```
4. Create a mount point directory and mount the new file system. For example:


```
# mkdir /extra
# mount /dev/dsk/c0t1d0s0 /extra
```
5. Recursively copy the contents of the /opt/ses directory to /export/home.


```
# cp -r /opt/ses /export/home
```
6. Unmount the /export/home file system. If your system reports that the /export/home file system is busy, use the umount -f command.


```
# umount /export/home
```
7. Use the ufsdump command to backup the /export/home file system to a file below the newly created mount point. Review the man page for ufsdump for information on how to use a dump file instead of a tape device.

```
# ufsdump 0uf/extra/dump_export_home_0/export/home
DUMP: Date of this level 0 dump: Wed Mar 28 11:07:04 2007
DUMP: Date of last level 0 dump: the epoch
DUMP: Dumping /dev/rdsk/c0t0d0s7 (sys-05:/export/home) to
/extradata/dump_export_home_0.
DUMP: Mapping (Pass I) [regular files]
DUMP: Mapping (Pass II) [directories]
DUMP: Writing 32 Kilobyte records
DUMP: Estimated 12130 blocks (5.92MB).
DUMP: Dumping (Pass III) [directories]
DUMP: Dumping (Pass IV) [regular files]
DUMP: 12094 blocks (5.91MB) on 1 volume at 8084 KB/sec
DUMP: DUMP IS DONE
DUMP: Level 0 dump on Wed Mar 28 11:07:04 2007
```

#

8. Mount the /export/home file system.
mount/export/home
9. Copy the contents of the /etc/uucp directory to the /export/home directory.
cp -r /etc/uucp/export/home
10. Unmount the /export/home file system.
umount/export/home
11. Use the ufsdump command to create an incremental backup for the /export/home file system. Save this data in a separate file on the spare disk.

```
# ufsdump 1uf/extra/dump_export_home_1/export/home
```

```
DUMP: Date of this level 1 dump: Wed Mar 28 11:08:35 2007
DUMP: Date of last level 0 dump: Wed Mar 28 11:07:04 2007
DUMP: Dumping /dev/rdsk/c0t0d0s7 (sys-05:/export/home) to
/extradata/dump_export_home_1.
DUMP: Mapping (Pass I) [regular files]
DUMP: Mapping (Pass II) [directories]
DUMP: Writing 32 Kilobyte records
DUMP: Estimated 6072 blocks (2.96MB).
DUMP: Dumping (Pass III) [directories]
DUMP: Dumping (Pass IV) [regular files]
DUMP: 6078 blocks (2.97MB) on 1 volume at 5525 KB/sec
DUMP: DUMP IS DONE
DUMP: Level 1 dump on Wed Mar 28 11:08:35 2007
```

#

12. Review the contents of the /etc/dumpdates file.
more /etc/dumpdates

You should see one line showing information for the Level 0 dump and another line for the Level 1 dump, for example:

```
/dev/rdsk/c0t0d0s7          0 Wed Mar 28 11:07:04 2007  
/dev/rdsk/c0t0d0s7          1 Wed Mar 28 11:08:35 2007
```

13. Mount the /export/home file system.

```
# mount /export/home
```

Notes:

Appendix C

Performing File System Restores

Objectives

Upon completion of this module, you should be able to restore ufs file systems. For ZFS file system boot information, see the ZFS File System Introduction module.

Note – Solaris releases that support the SPARC Newboot feature might require SPARC systems with at least 512 Mbytes of memory.

Restoring a ufs File System

Restore a file system to rebuild a damaged file system, to reinstall or upgrade the Solaris OS software, or to reorganize file systems on existing or new disks.

Restoring a Regular File System

When you restore system data, consider the following questions:

- Can the system boot on its own (regular file system restore)?
- Do you need to boot the system from CD-ROM, DVD, or network (critical file system restore)?
- Do you need to boot the system from CD-ROM, DVD, or network and repair the boot drive (special case recovery)?

To restore files or file systems, determine the following:

- The file system backup tapes that are needed
- The device name to which you will restore the file system
- The name of the temporary directory to which you will restore individual files
- The type of backup device to be used (local or remote)
- The backup device name (local or remote)

To restore a regular file system, such as the /export/home or /opt file system, back up to the disk, you use the `ufsrestore` command. The `ufsrestore` command copies files to the disk, relative to the current working directory, from backup tapes that were created by the `ufsdump` command.

You can use the `ufsrestore` command to reload an entire file system hierarchy from a Level 0 backup and related incremental backups. You can also restore one or more single files from any backup tape.

The format for the `ufsrestore` command is:

```
ufsrestore option(s) argument(s) filesystem
```

Table 15-6 describes some options that you can use with the `ufsrestore` command.

Table 15-6 Options for the `ufsrestore` Command

Option	Description
<code>t</code>	Lists the table of contents of the backup media.
<code>r</code>	Restores the entire file system from the backup media.
<code>x file1 file2</code>	Restores only the files named on the command line.
<code>i</code>	Invokes an interactive restore.
<code>v</code>	Specifies verbose mode. This mode displays the path names to the terminal screen as each file is restored.
<code>f device</code>	Specifies the tape device name. When not specified, the <code>/dev/rmt/0</code> device file is used.

When you restore an entire file system from a backup tape, the system creates a `restoresymtable` file. The `ufsrestore` command uses the `restoresymtable` file for check-pointing or passing information between incremental restores. You can remove the `restoresymtable` file when the restore is complete.

Using the `ufsrestore` Command to Restore a Regular File System

The following procedure demonstrates how to use the `ufsrestore` command to restore the `/opt` file system on the `c0t0d0s5` slice.

1. Create the new file system structure.
`# newfs /dev/rdsk/c0t0d0s5`
2. Mount the file system to the `/opt` directory, and change to that directory.
`# mount /dev/dsk/c0t0d0s5 /opt`
`# cd /opt`
3. Restore the entire `/opt` file system from the backup tape.
`# ufsrestore rf /dev/rmt/0`



Note – Always restore a file system by starting with the Level 0 backup tape, continuing with the next-lower-level tape, and continuing through the highest-level tape.

4. Remove the `restoresymtable` file.
`# rm restoresymtable`
5. Unmount the new file system.
`# cd /`
`# umount /opt`
6. Use the `fsck` command to check the restored file system.
`# fsck /dev/rdsk/c0t0d0s5`
7. Perform a full backup of the file system.
`# ufsdump 0uf /dev/rmt/0 /dev/rdsk/c0t0d0s5`

Note – The system administrator should always back up the newly created file system because the `ufsrestore` command repositions the files and changes the inode allocation.

8. Reboot the system.

```
# init 6
```

Restoring the `/usr` File System

To restore the `/usr` file system, boot from the Solaris 10 OS DVD, and then use the `ufsrestore` command to restore files back to the `/usr` partition.

Note – If the `/` (root), `/usr`, or `/var` file systems are unusable because of some type of corruption or damage, the system will not boot.



Using the `ufsrestore` Command to Restore a Critical File System

The following procedure demonstrates how to restore the `/usr` file system on Slice 6 of the boot disk.

To boot a SPARC-based system from DVD, perform the following:

1. Insert the Solaris 10 OS DVD.
2. From the OBP level, boot from DVD with the single-user mode option.

```
ok boot cdrom -s
```

To boot an x86/x64-based system from DVD, perform the following:

1. Insert the Solaris 10 OS DVD.
2. Use the reset button, or required key sequence (e.g., Control-Alt-Delete) to reset the system. You can power-cycle the system as an alternative. Your system needs to be configured to boot from DVD, which might require changes to BIOS settings before booting.

The system begins its boot process from the Solaris miniroot on DVD. The following list of boot choices displays:

1. Solaris Interactive (default)
2. Custom JumpStart
3. Solaris Interactive Text (Desktop session)
4. Solaris Interactive Text (Console session)
5. Apply driver updates
6. Single user shell

Enter the number of your choice.

3. Select boot choice 6, the single user shell. If an instance of Solaris is found on disk, the system asks if you want to mount it on /a. For example:

```
Solaris 10 5/09 s10x_u7wos_08 x86 was found on /dev/dsk/c1d0s0.
```

```
Do you wish to have it mounted read-write on /a? [y,n,?]
```

4. Enter n to avoid mounting the root file system. You will manually mount the file system you want to restore.

To restore the file system, do the following:

1. Create the new file system structure. For example, on a SPARC system:

```
# newfs /dev/rdsk/c0t0d0s6
```

Device names used for disk slices on x86/x64 systems differ slightly, for example, /dev/dsk/c1d0s6.

2. Mount the file system to the mount point /a, and change to that directory.

```
# mount /dev/dsk/c0t0d0s6 /a
```



- ```
cd /a
```
3. Restore the entire /usr file system from the backup tape.
- ```
# ufsrestore rf /dev/rmt/0
```

Note – Restore a file system by starting with the Level 0 backup tape, continuing with the next-lower-level tape, and continuing through the highest-level tape.

4. Remove the restoresymtable file.
- ```
rm restoresymtable
```
5. Unmount the new file system.
- ```
# cd /
# umount /a
```
6. Use the fsck command to check the restored file system.
- ```
fsck /dev/rdsck/c0t0d0s6
```
7. Perform a full backup of the file system.
- ```
# ufsdump 0uf /dev/rmt/0 /dev/rdsck/c0t0d0s6
```
8. Reboot the system. On x86/x64 systems, eject the DVD to prevent the system from booting from it.
- ```
init 6
```

## Performing Special Case Recovery of / (Root) on SPARC

Perform a special case recovery of the / (root) file system if there is damage to the boot block. To restore / (root), boot from the Solaris 10 OS DVD, and use ufsrestore.

The following demonstrates restoring the / (root) file system on slice 0 of the boot disk.

1. Insert the Solaris 10 OS DVD, and boot with the single-user mode option.
- ```
ok boot cdrom -s
```
2. Create the new file system structure.
- ```
newfs /dev/rdsck/c0t0d0s0
```
3. Mount the file system to the mount point /a and change to that directory.



```
mount /dev/dsk/c0t0d0s0 /a
cd /a
4. Restore the / (root) file system from the backup tape.
ufsrestore rf /dev/rmt/0
```

**Note** – Restore a file system by starting with the Level 0 backup tape, and continuing with the next-lower-level tape, and continuing through the highest-level tape.

5. Remove the **restoresymtable** file.
 

```
rm restoresymtable
```
6. Install the bootblk in Sectors 1 through 15 of the boot disk. To do this, change to the directory that contains the bootblk, and enter the **installboot** command.
 

```
cd /usr/platform/`uname -m`/lib/fs/ufs
installboot bootblk /dev/rdsck/c0t0d0s0
```
7. Unmount the new file system.
 

```
cd /
umount /a
```
8. Use the **fsck** command to check the restored file system.
 

```
fsck /dev/rdsck/c0t0d0s0
```
9. Perform a full backup of the file system.
 

```
ufsdump 0uf /dev/rmt/0 /dev/rdsck/c0t0d0s0
```
10. Reboot the system.
 

```
init 6
```

## Performing Special Case Recovery of /(Root) on x86/x64

If there is damage to the partitions within the active Solaris partition on an x86/x64 system, perform a special case recovery to recover the / (root) file system. If you reconstruct the entire boot disk, establish **fdisk** partitions before restoring the Solaris root file system or any additional file systems. To restore the / (root) file system, boot from the Solaris 10 OS DVD, and use the **ufsrestore** command.

The following procedure demonstrates how to restore the / (root) file system on slice 0 of the active Solaris partition. This procedure describes performing the file system restore, and installing the GRUB stage1 and stage2 programs on the active Solaris partition

1. Insert the Solaris 10 OS DVD.
2. Use the reset button or required key sequence (e.g., Ctrl-Alt-Delete) to reset the system. Power-cycle the system as an alternative.

---

**Note** – Your system needs to be configured to boot from a DVD. This might require changes to BIOS settings before booting.

---

3. The system begins its boot process from the Solaris miniroot on DVD. The following list of boot choices displays:
  1. Solaris Interactive (default)
  2. Custom JumpStart
  3. Solaris Interactive Text (Desktop session)
  4. Solaris Interactive Text (Console session)
  5. Apply driver updates
  6. Single user shell

Enter the number of your choice.

- 4. Select boot choice 6, “Single user shell.” If an instance of Solaris is found on the disk, the system asks if you want to mount it on /a. For example:

Solaris 10 5/09 s10x\_u7wos\_08 X86 was found on /dev/dsk/c1d0s0.

Do you wish to have it mounted read-write on /a? [y,n,?]

5. Enter n to avoid mounting the root file system. You will manually mount the file system you want to restore.
6. Create the new file system structure.  
`# newfs /dev/rdsk/c1d0s0`
7. Mount the file system to the mount point /a and change to that directory.  
`# mount /dev/dsk/c1d0s0 /a`  
`# cd /a`
8. Restore the / (root) file system from the backup tape.  
`# ufsrestore rf /dev/rmt/0`



**Note** – Always restore a file system by starting with a Level 0 backup tape, and continuing with the next-lower-level tape. Then continue through the highest-level tape.

9. Remove the **restoresymtable** file.

```
rm restoresymtable
```

10. Install the GRUB stage1 and stage2 programs on the Solaris partition. The **installgrub** command installs the GRUB stage1 program in sector 0, and GRUB stage2 program beginning at sector 50, (using 233 sectors) of the partition you specify. These sectors fall within the reserved first cylinder of the partition, and depending on the nature of the damage, may not require re-installation.

To install the stage1 and stage2 programs found in miniroot from the Solaris OS DVD, use the following command:

```
/sbin/installgrub /boot/grub/stage1 /boot/grub/stage2 /dev/rdsck/c1d0s0
```

To install the stage1 and stage2 programs found in the restored root file system, use the following command:

```
/sbin/installgrub /a/boot/grub/stage1
/a/boot/grub/stage2 /dev/rdsck/c1d0s0
```

11. Unmount the new file system.

```
cd /
umount /a
```

12. Use the **fsck** command to check the restored file system.

```
fsck /dev/rdsck/c0t0d0s0
```

13. Perform a full backup of the file system.

```
ufsdump 0uf /dev/rmt/0 /dev/rdsck/c0t0d0s0
```

14. Reboot the system. Eject the DVD to prevent the system from booting from it.

```
init 6
```

## Invoking an Interactive Restore

The **ufsrestore i** command invokes an interactive interface. Through the interface, you can browse the directory hierarchy of the backup tape and select individual files to extract. The term *volume* is used by **ufsrestore** and should be considered a single tape.

## Using the ufsrestore i Command

The following procedure demonstrates how to use the `ufsrestore i` command to extract individual files from a backup tape.

1. Become the root user, and change to the temporary directory that you want to receive the extracted files.

```
cd /export/home/tmp
```

2. Perform the `ufsrestore i` command.

```
ufsrestore ivf /dev/rmt/0
```

Verify volume and initialize maps

Media block size is 64

Dump date: Mon Oct 11 12:30:44 2004

Dumped from: the epoch

Level 0 dump of /export/home on sys43:/dev/dsk/c0t0d0s7

Label: none

Extract directories from tape

Initialize symbol table.

3. Display the contents of the directory structure on the backup tape.

```
ufsrestore > ls
```

..

|        |               |               |          |
|--------|---------------|---------------|----------|
| 2 *./  | 13 directory1 | 15 directory3 | 11 file2 |
| 2 *../ | 14 directory2 | 10 file1      | 12 file3 |

4. Change to the target directory on the backup tape.

```
ufsrestore > cd directory1
```

```
ufsrestore > ls
```

./directory1:

|         |        |            |            |            |
|---------|--------|------------|------------|------------|
| 3904 ./ | 2 *../ | 3905 file1 | 3906 file2 | 3907 file3 |
|---------|--------|------------|------------|------------|

5. Add the files you want to restore to the extraction list.

```
ufsrestore > add file1 file2
```

Make node ./directory1

Files you want to restore are marked with an asterisk (\*) for extraction. If you extract a directory, all of the directory contents are marked for extraction.

In this example, two files are marked for extraction. The `ls` command displays an asterisk in front of the selected file names, `file1` and `file2`.

```
ufsrestore > ls
```

./directory1:

|          |        |             |             |            |
|----------|--------|-------------|-------------|------------|
| 3904 *./ | 2 *../ | 3905 *file1 | 3906 *file2 | 3907 file3 |
|----------|--------|-------------|-------------|------------|

6. To delete a file from the extraction list, use the **delete** command.

```
ufsrestore > delete file1
```

The **ls** command displays the **file1** file without an asterisk.

```
ufsrestore > ls
```

```
./directory1:
```

```
3904 *./ 2 *.../ 3905 file1 3906 *file2 3907 file3
```

7. To view the files and directories marked for extraction, use the **marked** command.

```
ufsrestore > marked
```

```
./directory1:
```

```
3904 *./ 2 *.../ 3906 *file2
```

8. To restore the selected files from the backup tape, perform the command:

```
ufsrestore > extract
```

Extract requested files

You have not read any volumes yet.

Unless you know which volume your file(s) are on you should start with the last volume and work towards the first.

Specify next volume #: **1**




---

**Note** – The **ufsrestore** command has to find the selected files. If you used more than one tape for the backup, first insert the tape with the highest volume number and type the appropriate number. Then repeat this step, working towards Volume #1 until all files are restored.

---

```
extract file ./directory1/file2
```

Add links

Set directory mode, owner, and times.

set owner/mode for '.'? [yn] **n**




---

**Note** – Answering **y** sets ownership and permissions of the temporary directory to those of the mount point on the tape.

---

9. To exit the interactive restore after the files are extracted, perform the command:

```
ufsrestore> quit
```

10. Move the restored files to their original or permanent directory location, and delete the files from the temporary directory.

```
mv /export/home/tmp/directory1/file2 /export/home
rm -r /export/home/tmp/directory1
```



**Note** – You can use the help command in an interactive restore to display a list of available commands.

## Performing an Incremental Restore

When performing incremental restores, start with the last volume and work towards the first. The system uses information in the restoresymtable file to restore incremental backups on top of the latest full backup. The following procedure demonstrates how to restore the /export/home file system from incremental tapes.



**Note** – This procedure makes use of the interactive restore to assist in showing the concept of incremental restores. You would typically use a command, such as ufsrestore rf, for restoring entire file systems.

1. View the contents of the /etc/dumpdates file for information about the /export/home file system.

```
more /etc/dumpdates |grep c0t0d0s7
/dev/rdsck/c0t0d0s7 0 Wed Apr 7 09:55:34 2004
/dev/rdsck/c0t0d0s7 1 Web Apr 7 09:57:30 2004
```

2. Create the new file system structure for the /export/home file system.

```
newfs /dev/rdsck/c0t0d0s7
```

3. Mount the file system and change to that directory.

```
mount /dev/dsk/c0t0d0s7 /export/home
cd /export/home
```

4. Insert the Level 0 backup tape.

5. Restore the /export/home file system from the backup tapes.

```
ufsrestore rvf /dev/rmt/0
Verify volume and initialize maps
Media block size is 64
Dump date: Wed Apr 7 09:55:34 2004
Dumped from: the epoch
Level 0 dump of /export/home on sys41:/dev/dsk/c0t0d0s7
Label: none
Begin level 0 restore
Initialize symbol table.
Extract directories from tape
```

```
Calculate extraction list.
Make node ./directory1
Make node ./directory2
Make node ./directory3
Extract new leaves.
Check pointing the restore
extract file ./file1
extract file ./file2
extract file ./file3
Add links
Set directory mode, owner, and times.
Check the symbol table.
Check pointing the restore

6. Load the next lower-level tape into the tape drive.
ufsrestore rvf /dev/rmt/0
Verify volume and initialize maps
Media block size is 64
Dump date: Wed Apr 07 09:57:30 2004
Dumped from: Wed Apr 07 09:55:34 2004
Level 1 dump of /export/home on sys41:/dev/dsk/c0t0d0s7
Label: none
Begin incremental restore
Initialize symbol table.
Extract directories from tape
Mark entries to be removed.
Calculate node updates.
Make node ./directory4
Make node ./directory5
Make node ./directory6
Find unreferenced names.
Remove old nodes (directories).
Extract new leaves.
Check pointing the restore
extract file ./file4
extract file ./file5
extract file ./file6
Add links
Set directory mode, owner, and times.
Check the symbol table.
Check pointing the restore
#
```

## Alternative Steps

The following steps are alternatives to the previous Steps 5 and 6.

5. Restore the /export/home file system from the backup tapes. (This example uses an interactive, verbose restore to provide more detailed information.)

```
ufsrestore ivf /dev/rmt/0
Verify volume and initialize maps
Media block size is 64
Dump date: Mon Oct 11 13:10:12 2004
Dumped from: the epoch
Level 0 dump of /export/home on sys41:/dev/dsk/c0t0d0s7
Label: none
Extract directories from tape
Initialize symbol table.
ufsrestore > ls
.:
2 *./ 8 directory2 5 file2
2 *../ 9 directory3 6 file3
7 directory1 4 file1 3 lost+found/
```

The system lists files from the last Level 0 backup.

```
ufsrestore > add *
Warning: ./lost+found: File exists
ufsrestore > extract
Extract requested files
You have not read any volumes yet.
Unless you know which volume your file(s) are on you should start
with the last volume and work towards the first.
Specify next volume #: 1
extract file ./file1
extract file ./file2
extract file ./file3
extract file ./directory1
extract file ./directory2
extract file ./directory3
Add links
Set directory mode, owner, and times.
set owner/mode for '.'? [yn] n
Directories already exist, set modes anyway? [yn] n
ufsrestore > q
#
```

6. The information in the /etc/dumpdates file shows an incremental backup that was taken after the Level 0 backup. Load the next tape and perform the incremental restore.

```
ufsrestore iv
Verify volume and initialize maps
Media block size is 64
Dump date: Wed Apr 07 09:57:30 2004
Dumped from: Wed Apr 07 09:55:34 2004
Level 1 dump of /export/home on sys41:/dev/dsk/c0t0d0s7
Label: none
Extract directories from tape
Initialize symbol table.
ufsrestore > ls
.:
 2 *./ 13 directory4 15 directory6 11 file5
 2 *.../ 14 directory5 10 file4 12 file6
ufsrestore > add *
ufsrestore > extract
Extract requested files
You have not read any volumes yet.
Unless you know which volume your file(s) are on you should start
with the last volume and work towards the first.
Specify next volume #: 1
extract file ./file4
extract file ./file5
extract file ./file6
extract file ./directory4
extract file ./directory5
extract file ./directory6
Add links
Set directory mode, owner, and times.
set owner/mode for '.'? [yn] n
ufsrestore > q
#
```

# Exercise: Restoring Files and File Systems

In this exercise, you complete the following tasks:

- Backup the root (/) file system to tape
- Backup the root (/) file system to a file on disk
- Restore a selective file from a `ufsdump` archive
- Restore the root (/) file system from tape or disk

## Preparation

Some steps in this lab use a tape drive. If your system has an attached tape drive, obtain an appropriate tape from your instructor. A tape drive is not necessary to complete the exercise.

This exercise requires a system with a / (root) file system that is separate from the /usr and /var file systems, and that has an unused spare disk with the capacity to store the entire root (/) file system.

Identify the device name of the / (root) file system. For example:

```
df -h /
Filesystem size used avail capacity Mounted on
/dev/dsk/c0t0d0s0 4.8G 175M 4.6G 4% /
```

Record the device name of the root file system, `/dev/dsk/c0t0d0s0` in this example.

## Exercise Summary

This exercise is composed of four tasks that illustrate how to back up and restore the root (/) file system on SPARC and x86/x64 systems. The tasks in this exercise describe the following concepts:

- Back Up the root (/) File System to Tape - demonstrates the method used to create a `ufsdump` archive of the root file system on a tape device.
- Back Up the root (/) File System to a File on Disk - demonstrates the method used to create a `ufsdump` archive of the root file system on a file in a UFS file system.

- Restore the /etc/inet/hosts File From Tape or From a File in Disk - describes how to use an interactive restore session to select and restore a file from a ufsdump archive found either on tape, or in a file on disk.
- Destroy and Restore the root (/) File System - demonstrates destroying the root (/) file system by removing critical system files, booting SPARC and x86/x64 systems from DVD, creating a new root (/) file system, restoring the root file system from a ufsdump archive on tape or on disk, and installing boot programs.

### Task 1- Back Up the Root (/) File System to Tape

Complete the following steps:

If your system has an attached tape drive, perform the following steps to back up the root (/) file system to tape. Obtain a tape for your tape drive from your instructor.

1. Log in as the root user, and open a terminal window. Shut down the system to run level 0. Then boot the system to run level S (single-user milestone). Supply the root password as required to enter run level S.
2. Verify that a tape is in your tape drive.
3. Use the ufsdump command to create a backup tape for the root (/) file system.
4. Verify that the / (root) file system is on the tape.
5. Allow the system to remain at the single-user milestone.

### Task 2 - Back up the / (Root) File System to a File on Disk

Perform the following steps to back up your root (/) file system to a file on your spare disk.

1. Use the format utility to partition the extra hard drive and create one large partition to use as a backup device. Use the All Free Hog method to set partition 0 to use all disk space.



**Note** – This procedure assumes that on x86/x64 systems, the spare disk has one Solaris fdisk partition that consumes the entire disk.

---

2. Create a file system on the large partition on the extra disk.
3. Mount the new file system on /mnt.
4. Use the `ufsdump` command to backup the root (/) file system to a file below the /mnt directory. For example:
5. Verify that the / (root) file system has been backed up to the file on disk.
6. Allow the system to continue to boot to run level 3 (multi-user or all milestone).

## Task 3 – Restore /etc/inet/hosts File From Tape or File on Disk

Complete the following steps:

1. Log in as the root user, and open a terminal window. Change to the /var/tmp directory.
2. Verify that the new file system that contains your backup file is still mounted on /mnt, and that the backup file exists in it.
3. If the file system is no longer mounted, mount the new file system on /mnt.
4. Enter the `ufsrestore if` command to retrieve the /etc/inet/hosts file from the tape or the backup file on disk.
5. Change to the /etc/inet directory in the `ufsdump` archive, and list the files in the directory.
6. Add the hosts file to the list of files to extract, and display the list.
7. Extract the hosts file from the `ufsdump` archive. Specify volume number 1. Do not set the owner and mode for ., and then quit the `ufsrestore` command.
8. Verify that the etc/inet/hosts file exists below the /var/tmp directory.
9. Unmount the file system mounted on /mnt.



**Note** – Be certain to unmount the /mnt file system before you destroy your root file system in the next task; especially if your only root (/) file system backup is in a file below /mnt.

## Task 4 – Destroy and Restore the Root (/) File System

Complete the following steps:

### Destroy the Root File System

1. Change to the root (/) directory, and remove the following critical system directories and their contents: /platform and /devices.
2. Attempt to shut down the system to run-state 0.
3. If the system fails to shut down:
  - a. On SPARC systems, press the Stop-A key sequence to abort the operating system.
  - b. On x86/x64 systems, press the reset button, or use a key sequence (such as Control-Alt-Delete) to reset the system.  
Power-cycle the system as an alternative.
4. Attempt to boot the system from the boot disk.

What happens?

---

---

### Boot the SPARC system from DVD

1. Insert the Solaris 10 OS DVD
2. Boot the system from the DVD with the single-user mode option.

### Boot the x86/x64 System from DVD

1. Insert the Solaris 10 OS DVD.
2. If the Press any key to continue prompt displays, press any key, and allow the system to boot from DVD.

---

**Note** – Your system needs to be configured to boot from DVD, which might require changes to BIOS settings before booting.

---

3. If the system is not responsive, use the reset button, or required key sequence (e.g., Control-Alt-Delete) to reset the system. You can power-cycle the system as an alternative.



4. The system begins its boot process from the Solaris miniroot on DVD. The following list of boot choices displays:

1. Solaris Interactive (default)
2. Custom JumpStart
3. Solaris Interactive Text (Desktop session)
4. Solaris Interactive Text (Console session)
5. Apply driver updates
6. Single user shell

Enter the number of your choice.

5. Select boot choice 6, the single user shell. If an instance of Solaris is found on disk, the system asks if you want to mount it on /a. For example:

Solaris 10 5/09 s10x\_u7wos\_08 x86 was found on /dev/dsk/c1d0s0.

Do you wish to have it mounted read-write on /a? [y,n,?]

6. Enter n to avoid mounting the root file system. You will manually mount the file system you want to restore.

### Create a New Root File System and Prepare for the Restore

1. Use the newfs command to create a new file system on the / (root) slice.

---

**Note** – Be certain to specify the correct device name for the root slice. Be sure you do not specify a slice on the spare disk.

---

2. Mount the new file system on the /a directory.
3. Verify that your root backup tape is in the tape drive.
4. Mount the file system on your spare disk that contains your root backup as /mnt.
5. Verify that the file system that contains your backup file is mounted on /mnt, and that the backup file exists in it.

### Restore the Root (/) File System from Tape or Disk

1. Change directory to /a.
2. To restore the root (/) file system from tape, use the ufsrestore command and specify the tape device.
3. To restore the root (/) file system from disk, use the ufsrestore command and specify the backup file on disk.



## Exercise: Restoring Files and File Systems

---

4. Remove the `restoresymtable` file.
5. Change directory to root (/).
6. Unmount the file systems mounted on /a and /mnt.

### Install the Boot Block or GRUB Boot Programs

Depending on the damage incurred on disk, it may or may not be necessary to install the boot block on SPARC systems, or the GRUB boot programs (stage1 and stage2) on x86/x64 systems. If it is necessary to install these programs, perform the following steps:

1. On a SPARC system, to install a new boot block in sectors 1 through 15 of the root (/) slice, change directory to the location of the boot block, and use the `installboot` command.
2. On an x86/x64 system, to install the GRUB stage1 and stage2 programs found in miniroot from the Solaris OS DVD, use the `installgrub` command.

### Reboot the System

1. On x86/x64 systems, eject the Solaris OS DVD to prevent the system from booting from it.
2. Reboot the system.
3. Login as root and verify that the system is operational.

## Exercise Summary



**Discussion** – Take a few minutes to discuss what experiences, issues, or discoveries you had during the lab exercises.

- Experiences
- Interpretations
- Conclusions
- Applications

# Exercise Solutions: Restoring Files and File Systems

This section contains solutions to the exercise.




---

**Note** – Disk and file system device names used in these procedures represent devices found on both SPARC and x86/x64 systems, and are provided as examples only. Use correct device names for your specific system.

---

## Task 1- Back Up the Root (/) File System to Tape

Complete the following steps:

If your system has an attached tape drive, perform the following steps to back up the root (/) file system to tape. Obtain a tape for your tape drive from your instructor.

1. Log in as the root user, and open a terminal window. Shut down the system to run level 0. Then boot the system to run level S (single-user milestone). Supply the root password as required to enter run level S.

```
init 0
(output omitted)
```

On a SPARC system, use the `boot -s` command at the `ok` prompt:

```
ok boot -s
```

On an x86/x64 system, perform the following steps:

- a. Press any key at the `Press any key to reboot` prompt.
- b. When the GRUB main menu displays, enter `e` to display the list of commands for the default boot choice. For example:

```
GNU GRUB version 0.95 (572K lower / 3380808K upper memory)
+-----+
| kernel /platform/i86pc/multiboot
| module /platform/i86pc/boot_archive
|
```

- c. Highlight the `kernel` command in the command list, and enter `e` to edit the command. Add the `-s` option to the end of the `kernel` command line, and press Return. For example:

```
grub edit> kernel /platform/i86pc/multiboot -s
```

- Unauthorized reproduction or distribution prohibited. Copyright© 2011, Oracle and/or its affiliates.
- d. Enter b to boot the system.
  - 2. Verify that a tape is in your tape drive.
  - 3. Use the `ufsdump` command to create a backup tape for the root (/) file system.  
`# ufsdump 0uf /dev/rmt/0 /`
  - 4. Verify that the / (root) file system is on the tape.  
`# ufsrestore tvf /dev/rmt/0`  
*The screen should list directory structures under root (/) first, followed by files.*
  - 5. Allow the system to remain at the single-user milestone.

## Task 2 - Back up / (Root) File System to a File on Disk

Perform the following steps to back up your root (/) file system to a file on your spare disk.

- 1. Use the `format` utility to partition the extra hard drive and create one large partition to use as a backup device. Use the All Free Hog method to set partition 0 to use all disk space. This is an example from a SPARC system:

```
format
Searching for disks...done
```

AVAILABLE DISK SELECTIONS:

- 0. c0t0d0 <HDS722512VLAT80 cyl 57459 alt 2 hd 16 sec 255>  
`/pci@1e,600000/ide@d/dad@0,0`
- 1. c0t1d0 <DEFAULT cyl 57459 alt 2 hd 16 sec 255>  
`/pci@1e,600000/ide@d/dad@1,0`

Specify disk (enter its number): **1**

selecting c0t1d0

When you are through partitioning the extra disk drive, your partition table should look something like this:

| Part | Tag        | Flag | Cylinders | Size     | Blocks                |
|------|------------|------|-----------|----------|-----------------------|
| 0    | root       | wm   | 0 - 57458 | 111.79GB | (57459/0/0) 234432720 |
| 1    | swap       | wu   | 0         | 0        | (0/0/0) 0             |
| 2    | backup     | wu   | 0 - 57458 | 111.79GB | (57459/0/0) 234432720 |
| 3    | unassigned | wm   | 0         | 0        | (0/0/0) 0             |
| 4    | unassigned | wm   | 0         | 0        | (0/0/0) 0             |

## Exercise Solutions: Restoring Files and File Systems

|              |    |   |   |         |   |
|--------------|----|---|---|---------|---|
| 5 unassigned | wu | 0 | 0 | (0/0/0) | 0 |
| 6 unassigned | wu | 0 | 0 | (0/0/0) | 0 |
| 7 unassigned | wu | 0 | 0 | (0/0/0) | 0 |

This is an example from an x86/x64 system:



**Note** – This procedure assumes that on x86/x64 systems, the spare disk has one Solaris fdisk partition that consumes the entire disk.

### # format

Searching for disks...done

#### AVAILABLE DISK SELECTIONS:

0. c1d0 <DEFAULT cyl 39887 alt 2 hd 64 sec 63>  
/pci@0,0/pci-ide@7/ide@0/cmdk@0,0
1. c2d0 <DEFAULT cyl 10008 alt 2 hd 255 sec 63>  
/pci@0,0/pci-ide@7/ide@1/cmdk@0,0

Specify disk (enter its number) : **1**

selecting c2d0

When you are through partitioning the extra disk drive, your partition table should look something like this:

| Part | Tag        | Flag | Cylinders | Size    | Blocks                |
|------|------------|------|-----------|---------|-----------------------|
| 0    | reserved   | wm   | 3 - 10007 | 76.64GB | (10005/0/0) 160730325 |
| 1    | swap       | wu   | 0         | 0       | (0/0/0) 0             |
| 2    | backup     | wu   | 0 - 10007 | 76.67GB | (10008/0/0) 160778520 |
| 3    | unassigned | wu   | 0         | 0       | (0/0/0) 0             |
| 4    | unassigned | wu   | 0         | 0       | (0/0/0) 0             |
| 5    | unassigned | wu   | 0         | 0       | (0/0/0) 0             |
| 6    | unassigned | wu   | 0         | 0       | (0/0/0) 0             |
| 7    | unassigned | wu   | 0         | 0       | (0/0/0) 0             |
| 8    | boot       | wu   | 0 - 0     | 7.84MB  | (1/0/0) 16065         |
| 9    | alternates | wm   | 1 - 2     | 15.69MB | (2/0/0) 32130         |

2. Create a file system on the large partition on the extra disk. For example, on a SPARC system:

**# newfs /dev/rdsck/c0t1d0s0**

On an x86/x64 system:

**# newfs /dev/rdsck/c2d0s0**

3. Mount the new file system on /mnt. For example, on a SPARC system:

**# mount /dev/dsk/c0t1d0s0 /mnt**

On an x86/x64 system:

- # **mount /dev/dsk/c2d0s0 /mnt**
4. Use the **ufsdump** command to backup the root (/) file system to a file below the /mnt directory. For example:  
# **ufsdump 0uf /mnt/dump\_root\_0 /**
  5. Verify that the / (root) file system has been backed up to the file on disk. For example:  
# **ufsrestore tvf /mnt/dump\_root\_0**  
*The screen should list directory structures under root (/) first, followed by files.*
  6. Allow the system to continue to boot to run level 3 (multi-user or all milestone).  
# <Control-D>

## Task 3 – Restore /etc/inet/hosts From Tape or File on Disk

Complete the following steps:

1. Log in as the root user, and open a terminal window. Change to the /var/tmp directory.  
# **cd /var/tmp**
2. Verify that the new file system that contains your backup file is still mounted on /mnt, and that the backup file exists in it.  
# **ls /mnt**  
dump\_root\_0 lost+found  
#
3. If the file system is no longer mounted, mount the new file system on /mnt. For example, on a SPARC system:  
# **mount /dev/dsk/c0t1d0s0 /mnt**  
On an x86/x64 system:  
# **mount /dev/dsk/c2d0s0 /mnt**
4. Enter the **ufsrestore if** command to retrieve the /etc/inet/hosts file from the tape or the backup file on disk.  
For example, to restore the file from tape:  
# **ufsrestore if /dev/rmt/0**  
For example, to restore the file from a file on disk:  
# **ufsrestore if /mnt/dump\_root\_0**

```
ufsrestore > ls
```

*See files and directories for the / (root) file system.*

Change to the /etc/inet directory in the ufsdump archive, and list the files in the directory. See files and directories for the /etc/inet file system.

```
ufsrestore > cd /etc/inet
```

```
ufsrestore > ls
```

Add the hosts file to the list of files to extract, and display the list.

See the hosts file listed.

```
ufsrestore > add hosts
```

```
ufsrestore > marked
```

```
./etc/inet:
```

```
*hosts
```

```
ufsrestore >
```

5. Extract the hosts file from the ufsdump archive. Specify volume number 1. Do not set the owner and mode for ., and then quit the ufsrestore command.

```
ufsrestore > extract
```

Extract requested files

You have not read any volumes yet.

Unless you know which volume your file(s) are on you should start with the last volume and work towards the first.

Specify next volume #: **1**

set owner/mode for '.'? [yn] **n**

```
ufsrestore > q
```

6. Verify that the etc/inet/hosts file exists below the /var/tmp directory.

```
ls etc/inet/hosts
```

```
etc/inet/hosts
```

7. Unmount the file system mounted on /mnt.

```
umount /mnt
```

```
#
```

---

**Note –** Unmount the /mnt file system before you destroy your root file system in the next task, especially if your only root (/) file system backup is in a file below /mnt.

---



## Task 4 – Destroy and Restore the Root (/) File System

Complete the following steps:

### Destroy the Root File System

1. Change to the root (/) directory, and remove the following critical system directories and their contents: /platform and /devices.

```
cd /
rm -r /platform /devices
```

2. Attempt to shut down the system to run-state 0.

```
init 0
```

3. If the system fails to shut down:

- a. On SPARC systems, press the Stop-A key sequence to abort the operating system.
- b. On x86/x64 systems, press the reset button, or use a key sequence (e.g., Ctrl-Alt-Delete) to reset the system. Power-cycle the system as an alternative.

4. Attempt to boot the system from the boot disk. You will see that the system fails to boot. SPARC systems display the message:

Boot load failed.

The file just loaded does not appear to be executable  
x86/x64 systems display the message:

Error 20: Multiboot kernel must be loaded before  
modules

### Boot the SPARC system from DVD

1. Insert the Solaris 10 OS DVD
2. Boot the system from the DVD with the single-user mode option.  
ok **boot cdrom -s**

### Boot the x86/x64 System from DVD

1. Insert the Solaris 10 OS DVD.
2. If the Press any key to continue prompt displays, press any key, and allow the system to boot from DVD.



**Note** – Your system needs to be configured to boot from DVD, which might require changes to BIOS settings before booting.

---

3. If the system is not responsive, use the reset button, or required key sequence (e.g., Control-Alt-Delete) to reset the system. You can power-cycle the system as an alternative.
4. The system begins its boot process from the Solaris miniroot on DVD. The following list of boot choices displays:
  1. Solaris Interactive (default)
  2. Custom JumpStart
  3. Solaris Interactive Text (Desktop session)
  4. Solaris Interactive Text (Console session)
  5. Apply driver updates
  6. Single user shell

Enter the number of your choice.

5. Select boot choice 6, the single user shell. If an instance of Solaris is found on disk, the system asks if you want to mount it on /a. For example:

Solaris 10 5/09 s10x\_u7wos\_08 X86 was found on /dev/dsk/c1d0s0.

Do you wish to have it mounted read-write on /a? [y,n,?]

6. Enter n to avoid mounting the root file system. You will manually mount the file system you want to restore.

### Create a New Root File System and Prepare for the Restore

1. Use the newfs command to create a new file system on the / (root) slice.

**Note** – Be certain to specify the correct device name for the root slice. Be sure you do not specify a slice on the spare disk.

---

For example, on a SPARC system:

```
newfs /dev/rdsk/c0t0d0s0
```

On an x86/x64 system:

```
newfs /dev/rdsk/c1d0s0
```

2. Mount the new file system on the /a directory.

For example, on a SPARC system:

```
mount /dev/dsk/c0t0d0s0 /a
```

On an x86/x64 system:

```
mount /dev/dsk/c1d0s0 /a
```

3. Verify that your root backup tape is in the tape drive.
4. Mount the file system on your spare disk that contains your root backup as /mnt.

For example, on a SPARC system:

```
mount /dev/dsk/c0t1d0s0 /mnt
```

On an x86/x64 system:

```
mount /dev/dsk/c2d0s0 /mnt
```

5. Verify that the file system that contains your backup file is mounted on /mnt, and that the backup file exists in it.

```
ls /mnt
```

```
dump_root_0 lost+found
```

## Restore the Root (/) File System from Tape or Disk

1. Change directory to /a.

```
cd /a
```

2. To restore the root (/) file system from tape, use the ufsrestore command and specify the tape device:

```
ufsrestore rf /dev/rmt/0
```

3. To restore the root (/) file system from disk, use the ufsrestore command and specify the backup file on disk:

```
ufsrestore rf /mnt/dump_root_0
```

4. Remove the restoresymtable file.

```
rm restoresymtable
```

5. Change directory to root (/).

```
cd /
```

6. Unmount the file systems mounted on /a and /mnt.

```
umount /a
```

```
umount /mnt
```

## Install the Boot Block or GRUB Boot Programs

Depending on the damage incurred on disk, it may or may not be necessary to install the boot block on SPARC systems, or the GRUB boot programs (stage1 and stage2) on x86/x64 systems. If it is necessary to install these programs, perform the following steps:

1. On a SPARC system, to install a new boot block in sectors 1 through 15 of the root (/) slice, change directory to the location of the boot block, and use the `installboot` command. For example:  

```
cd /usr/platform/`uname -m`/lib/fs/ufs
installboot bootblk /dev/rdsk/c0t0d0s0
```
2. On an x86/x64 system, to install the GRUB stage1 and stage2 programs found in miniroot from the Solaris OS DVD, use the `installgrub` command.

```
/sbin/installgrub /boot/grub/stage1 /boot/grub/stage2 /dev/rdsk/c1d0s0
```

## Reboot the System

1. On x86/x64 systems, eject the Solaris OS DVD to prevent the system from booting from it.
2. Reboot the system.  

```
init 6
```
3. Login as root and verify that the system is operational.